

Model Compression: Pruning and Quantization


Pooyan Jamshidi
UofSC

The slides are mainly based on a NeurIPS'15 tutorial by William Dally



**Reducing Size of
Network Reduces Work
and Storage**

Prune Unneeded Connections

Pruning




Retrain to Recover Accuracy





Han et al. Learning both Weights and Connections for Efficient Neural Networks, NIPS 2015





Pruning of AlexNet



Pruning of VGG-16



Pruning Neural Talk and LSTM



Pruning Neural Talk and LSTM



- **Original:** a basketball player in a white uniform is playing with a ball
- **Pruned 90%:** a basketball player in a white uniform is playing with a basketball



- **Original :** a brown dog is running through a grassy field
- **Pruned 90%:** a brown dog is running through a grassy area



- **Original :** a man is riding a surfboard on a wave
- **Pruned 90%:** a man in a wetsuit is riding a wave on a beach



- **Original :** a soccer player in red is running in the field
- **Pruned 95%:** a man in a red shirt and black and white black shirt is running through a field

Speedup of Pruning on CPU/GPU




Figure 9: Compared with the original network, pruned network layer achieved 3 \times speedup on CPU, 3.5 \times on GPU and 4.2 \times on mobile GPU on average. Batch size = 1 targeting real time processing. Performance number normalized to CPU.

Intel Core i7 5930K: MKL CBLAS GEMV, MKL SPBLAS CSRMV
NVIDIA GeForce GTX Titan X: cuBLAS GEMV, cuSPARSE CSRMV
NVIDIA Tegra K1: cuBLAS GEMV, cuSPARSE CSRMV


History of Pruning

Yann LeCun, John S. Denker, and Sara A. Solla. Optimal Brain Damage. In *Advances in Neural Information Processing Systems*, pages 598–605. Morgan Kaufmann, 1990.


Babak Hassibi, David G Stork, et al. Second order derivatives for network pruning: Optimal brain surgeon. *Advances in neural information processing systems*, pages 164–164, 1993.

**Reduce Storage for
Each Remaining Weight**


Trained Quantization (Weight Sharing)





Weight Sharing via K-Means




Trained Quantization



Bits per Weight



Pruning + Trained Quantization



Summary of Compression

Table 1: The compression pipeline can save $35\times$ to $49\times$ parameter storage with no loss of accuracy.

| Network | Top-1 Error | Top-5 Error | Parameters | Compress Rate |
|--------------------------|-------------|-------------|----------------|---------------|
| LeNet-300-100 Ref | 1.64% | - | 1070 KB | 40× |
| LeNet-300-100 Compressed | 1.58% | - | 27 KB | |
| LeNet-5 Ref | 0.80% | - | 1720 KB | 39× |
| LeNet-5 Compressed | 0.74% | - | 44 KB | |
| AlexNet Ref | 42.78% | 19.73% | 240 MB | 35× |
| AlexNet Compressed | 42.78% | 19.70% | 6.9 MB | |
| VGG-16 Ref | 31.50% | 11.32% | 552 MB | 49× |
| VGG-16 Compressed | 31.17% | 10.91% | 11.3 MB | |

Compress neural networks without affecting accuracy by:

1. Pruning the unimportant connections =>
2. Quantizing the network and enforce weight sharing =>
3. Apply Huffman encoding

30x – 50x Compression Means

- Complex DNNs can be put in mobile applications (<100MB total):
 - 1GB network (250M Weights) become 20-30 MB
- Memory bandwidth reduced by 30-50x:
 - Particularly for FC layers in real-time applications with no reuse
- Memory working set fits in on-chip SRAM
 - 5pJ/word access vs 640pJ/word