

# Machine Learning Systems

Lecture 9: Convolutional Neural Networks (CNNs, ConvNets)

Pooyan Jamshidi



# Convolutional Neural Networks (CNNs, ConvNets)

Pooyan Jamshidi  
UofSC

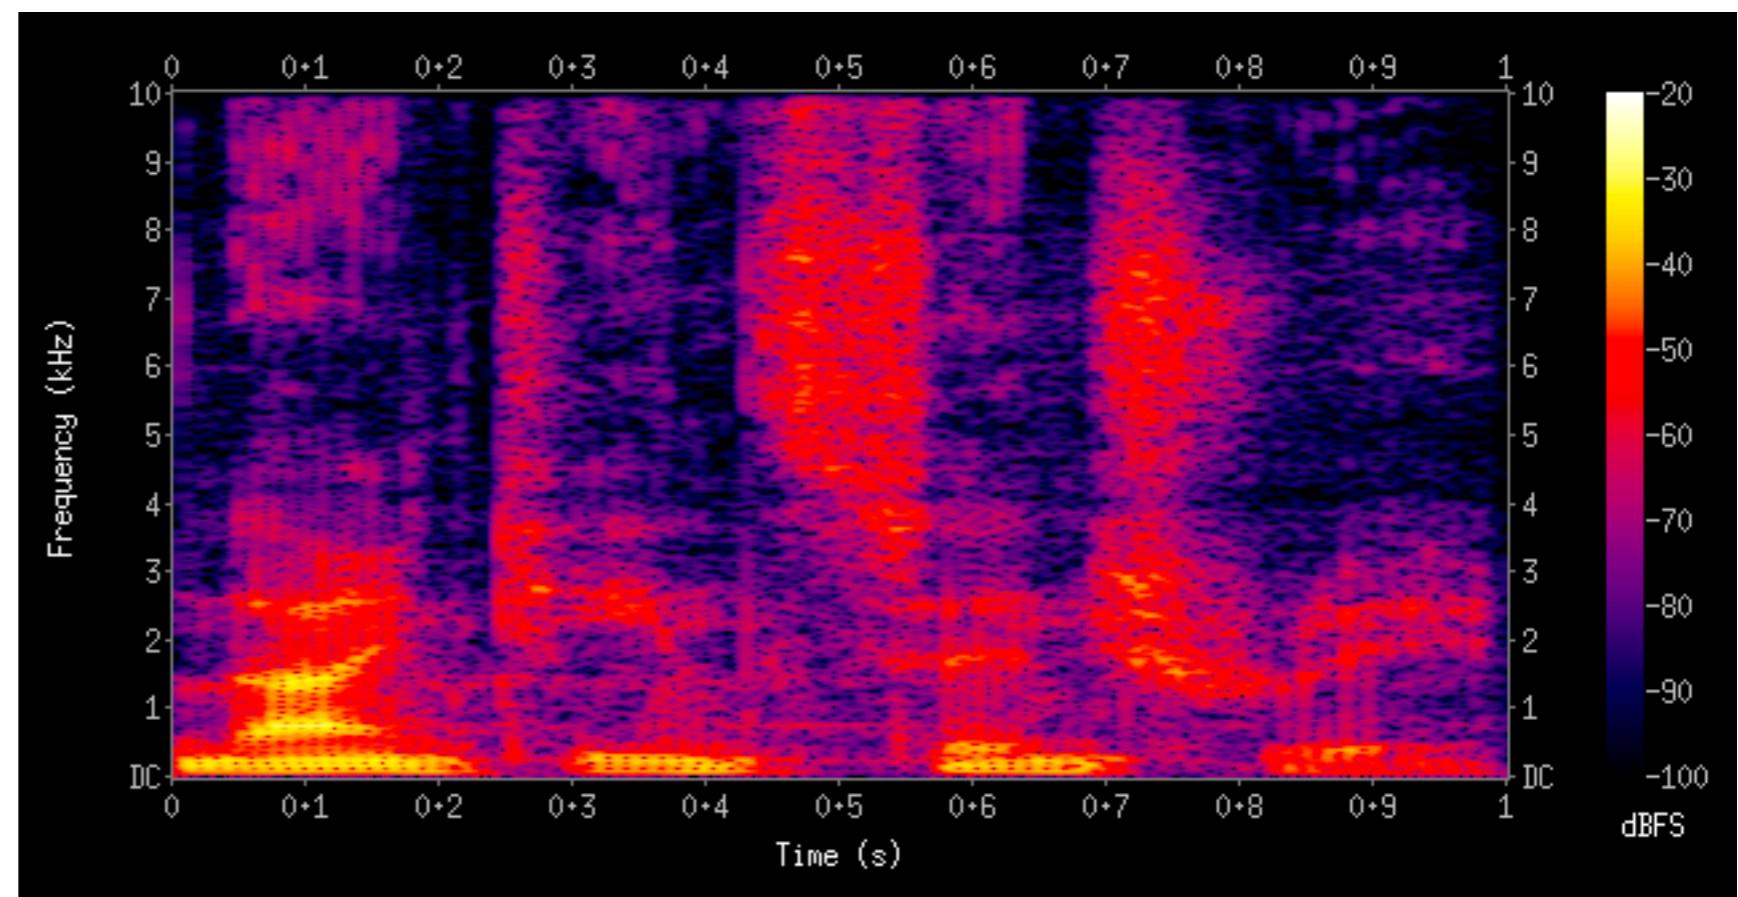
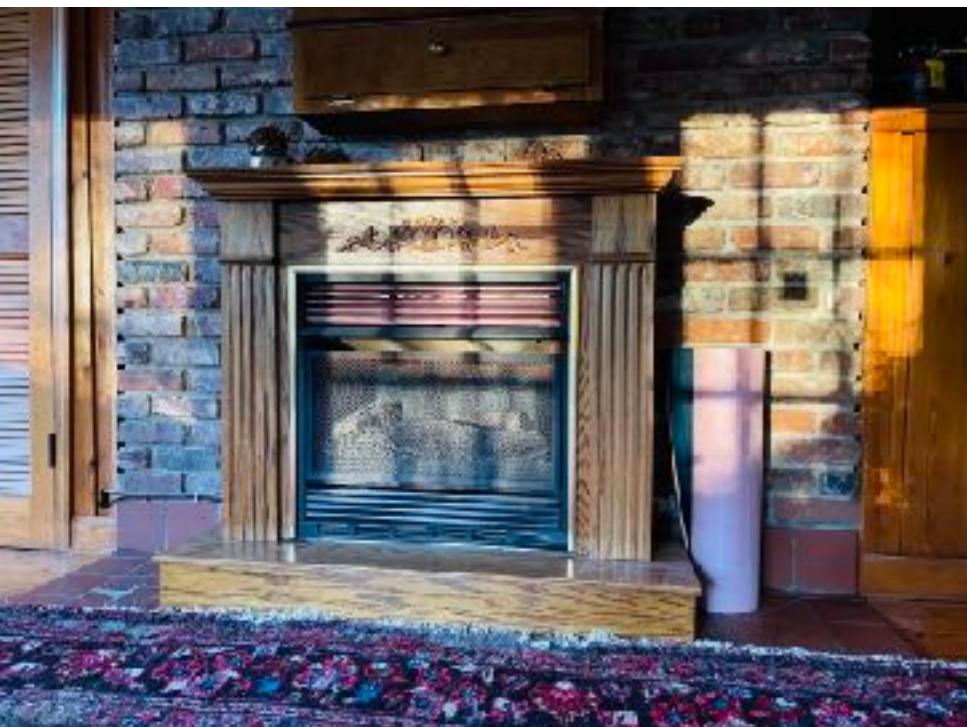
Partially based on:

- Chapter 9 of the *Deep Learning Book*: <https://www.deeplearningbook.org/>
- CS231n Convolutional Neural Networks for Visual Recognition

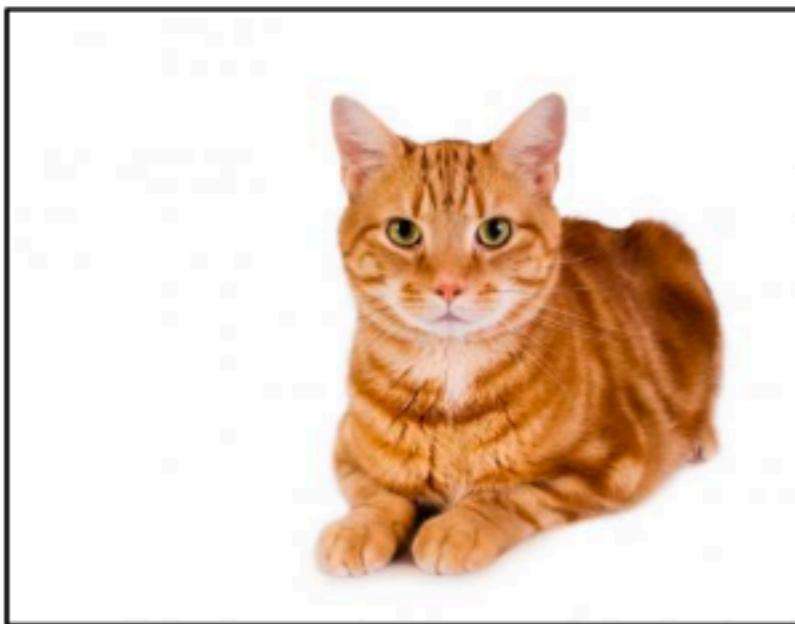
# Convolutional Networks

- Scale up neural networks to process very large images / video sequences
  - Sparse connections
  - Parameter sharing
- Automatically generalize across spatial translations of inputs
- Applicable to any input that is laid out on a grid (1-D, 2-D, 3-D, ...)

# Examples of Grid Structures



# Shift Invariance in Convolutional Neural Networks



Cat



Cat

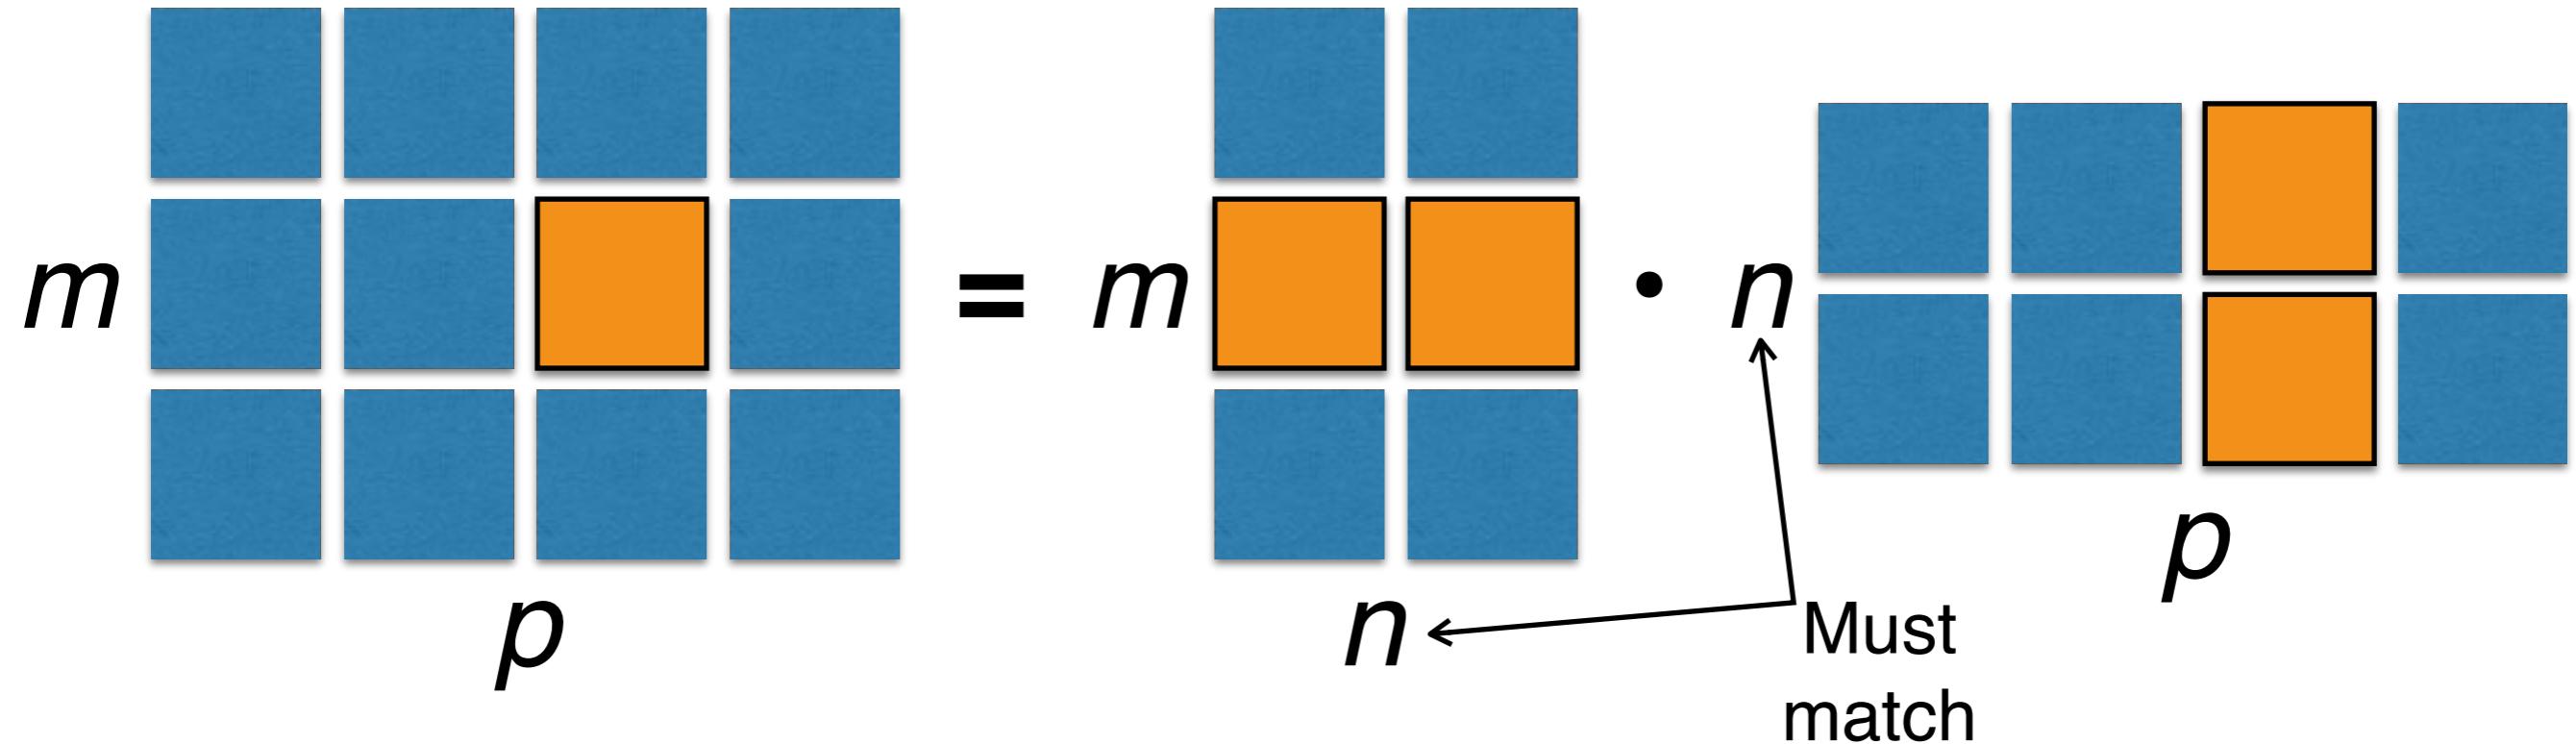
# Key Idea

- Replace matrix multiplication in neural nets with convolution
- Everything else stays the same
  - Maximum likelihood
  - Back-propagation
  - etc.

# Matrix (Dot) Product

$$C = AB. \quad (2.4)$$

$$C_{i,j} = \sum_k A_{i,k} B_{k,j}. \quad (2.5)$$



# Matrix Transpose

$$(\mathbf{A}^\top)_{i,j} = A_{j,i}. \quad (2.3)$$

The diagram shows a 3x2 matrix  $A$  with elements  $A_{1,1}, A_{1,2}, A_{2,1}, A_{2,2}, A_{3,1}, A_{3,2}$ . A curved arrow points from the element  $A_{1,2}$  to its transpose position  $A_{2,1}$ , indicating the swap of indices. To the right, the transpose matrix  $\mathbf{A}^\top$  is shown as a 2x3 matrix with elements  $A_{1,1}, A_{2,1}, A_{3,1}, A_{1,2}, A_{2,2}, A_{3,2}$ .

$$\mathbf{A} = \begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \\ A_{3,1} & A_{3,2} \end{bmatrix} \Rightarrow \mathbf{A}^\top = \begin{bmatrix} A_{1,1} & A_{2,1} & A_{3,1} \\ A_{1,2} & A_{2,2} & A_{3,2} \end{bmatrix}$$

Figure 2.1: The transpose of the matrix can be thought of as a mirror image across the main diagonal.

$$(\mathbf{AB})^\top = \mathbf{B}^\top \mathbf{A}^\top. \quad (2.9)$$

# 2D Convolution

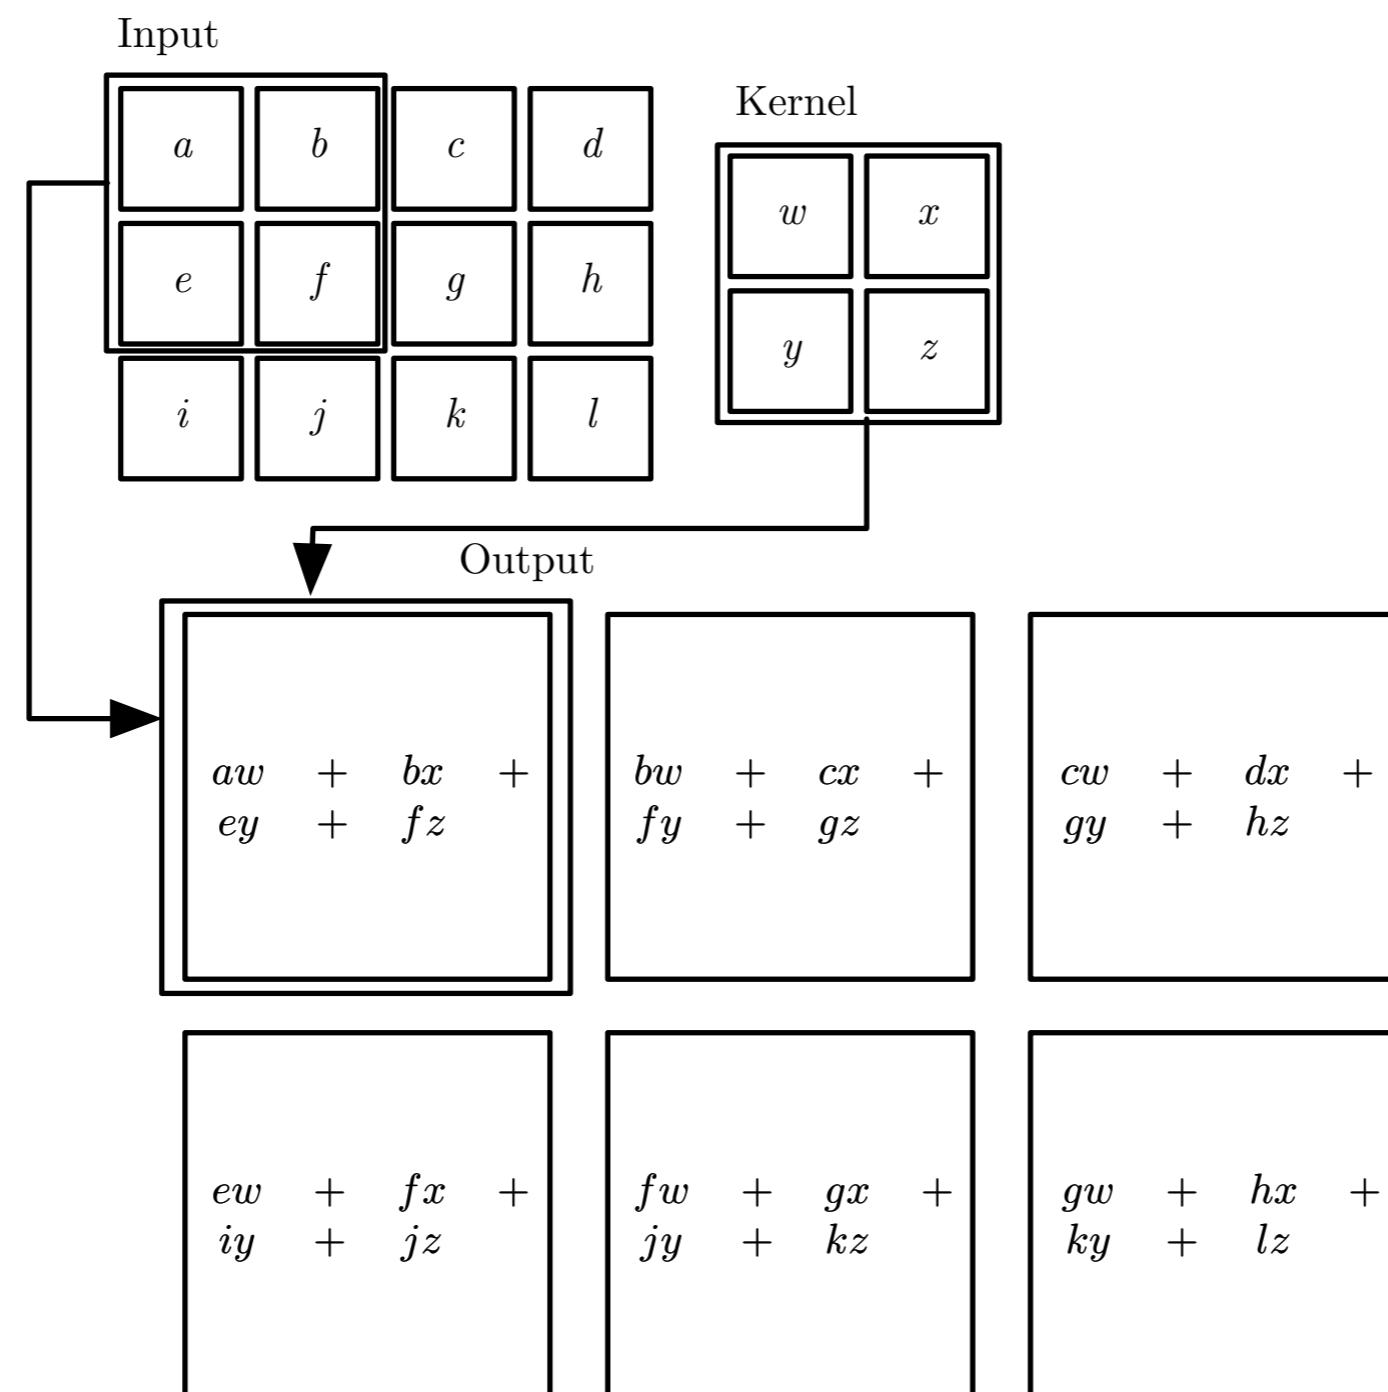


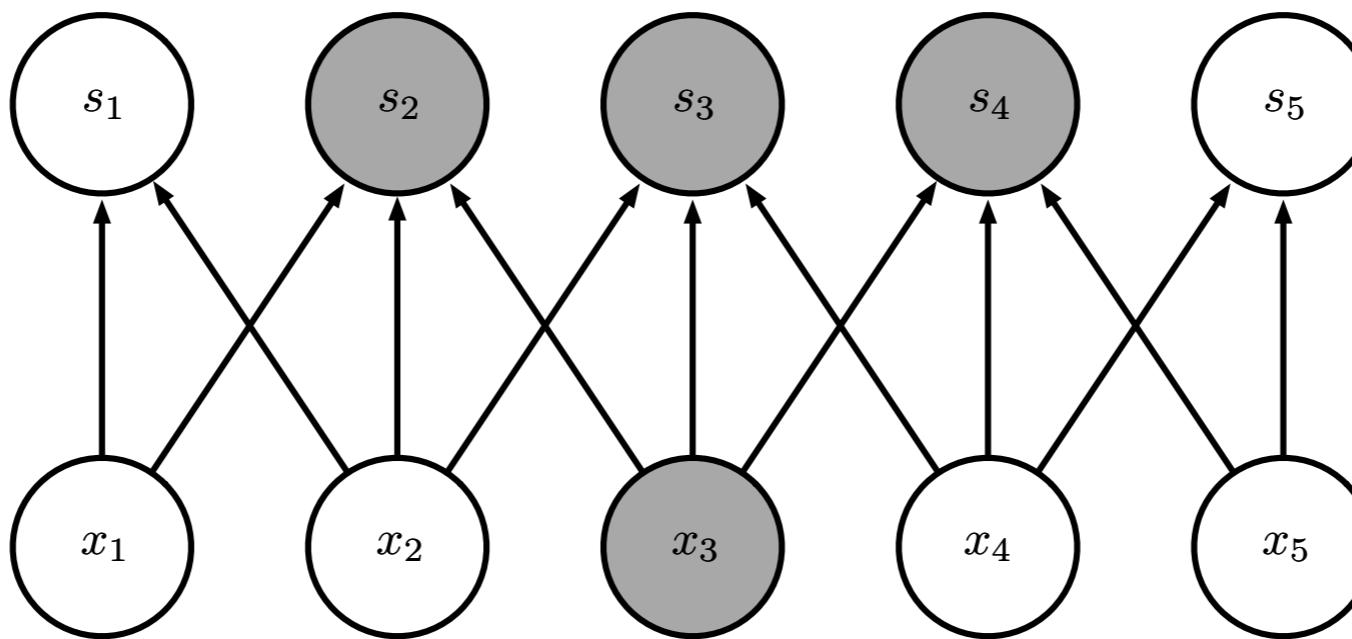
Figure 9.1

# Three Operations

- Convolution: like matrix multiplication
  - Take an input, produce an output (hidden layer)
- “Deconvolution”: like multiplication by transpose of a matrix
  - Used to back-propagate error from output to input
  - Reconstruction in autoencoder / RBM
- Weight gradient computation
  - Used to backpropagate error from output to weights
  - Accounts for the parameter sharing

# Sparse Connectivity

Sparse  
connections  
due to small  
convolution  
kernel



Dense  
connections

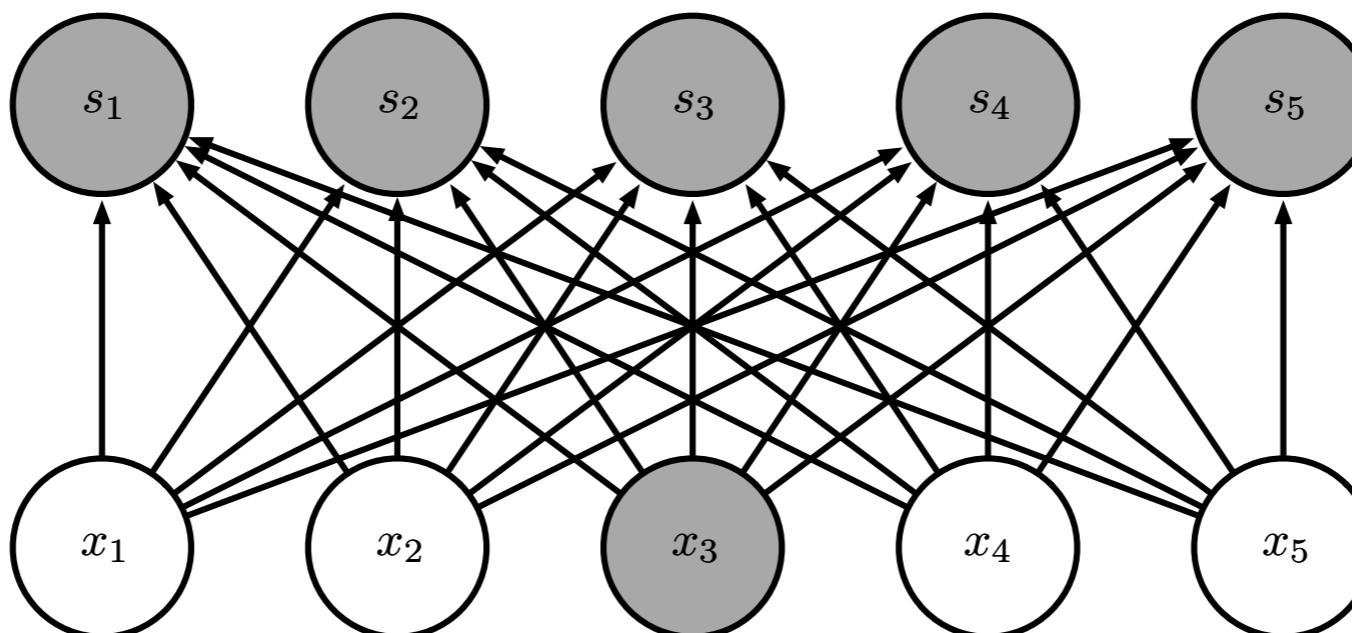
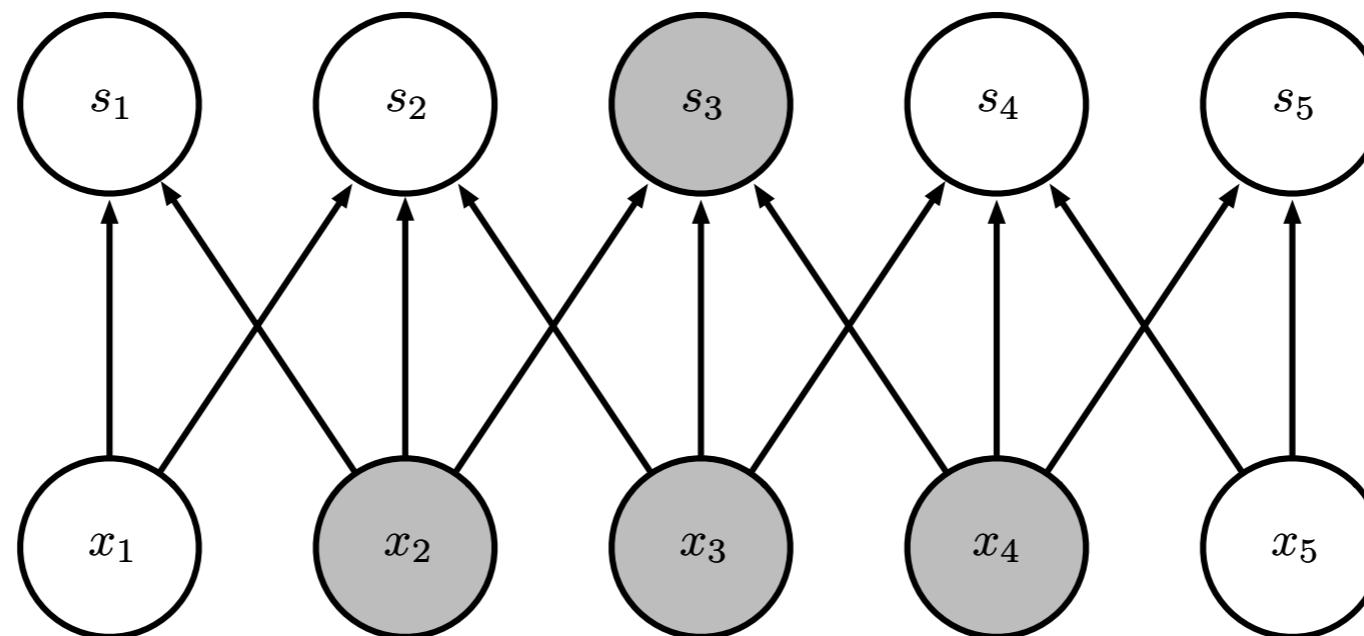


Figure 9.2

# Sparse Connectivity

Sparse  
connections  
due to small  
convolution  
kernel



Dense  
connections

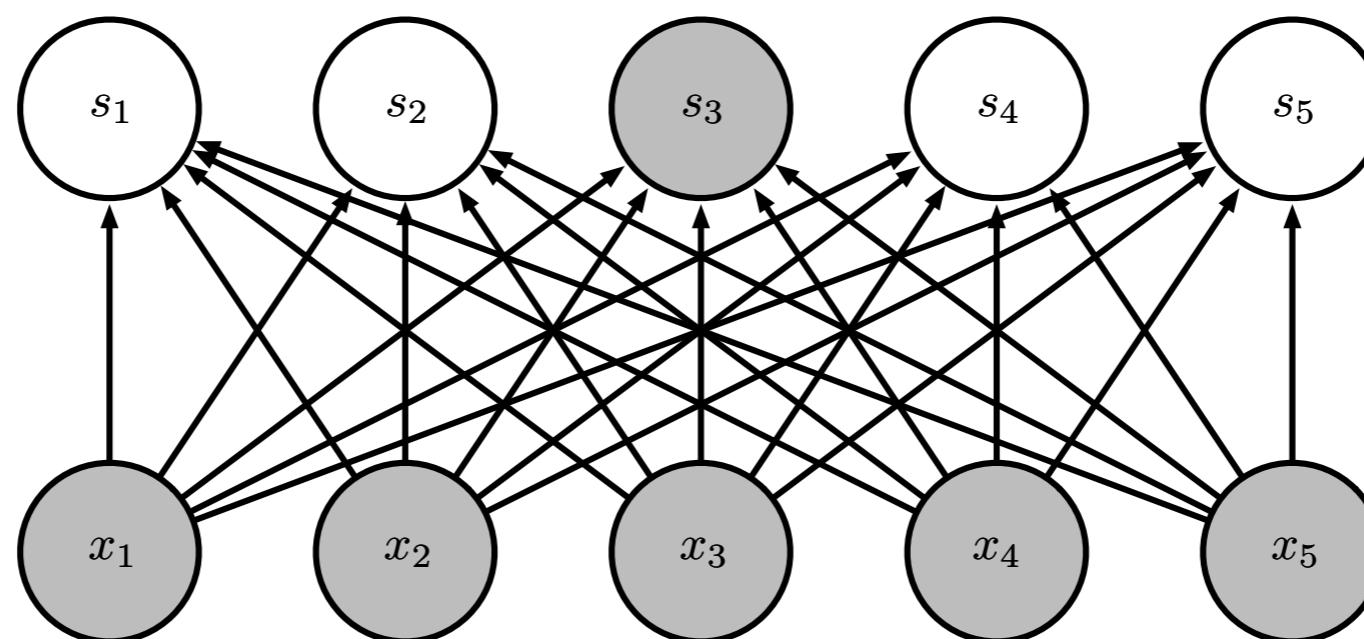


Figure 9.3

# Growing Receptive Fields

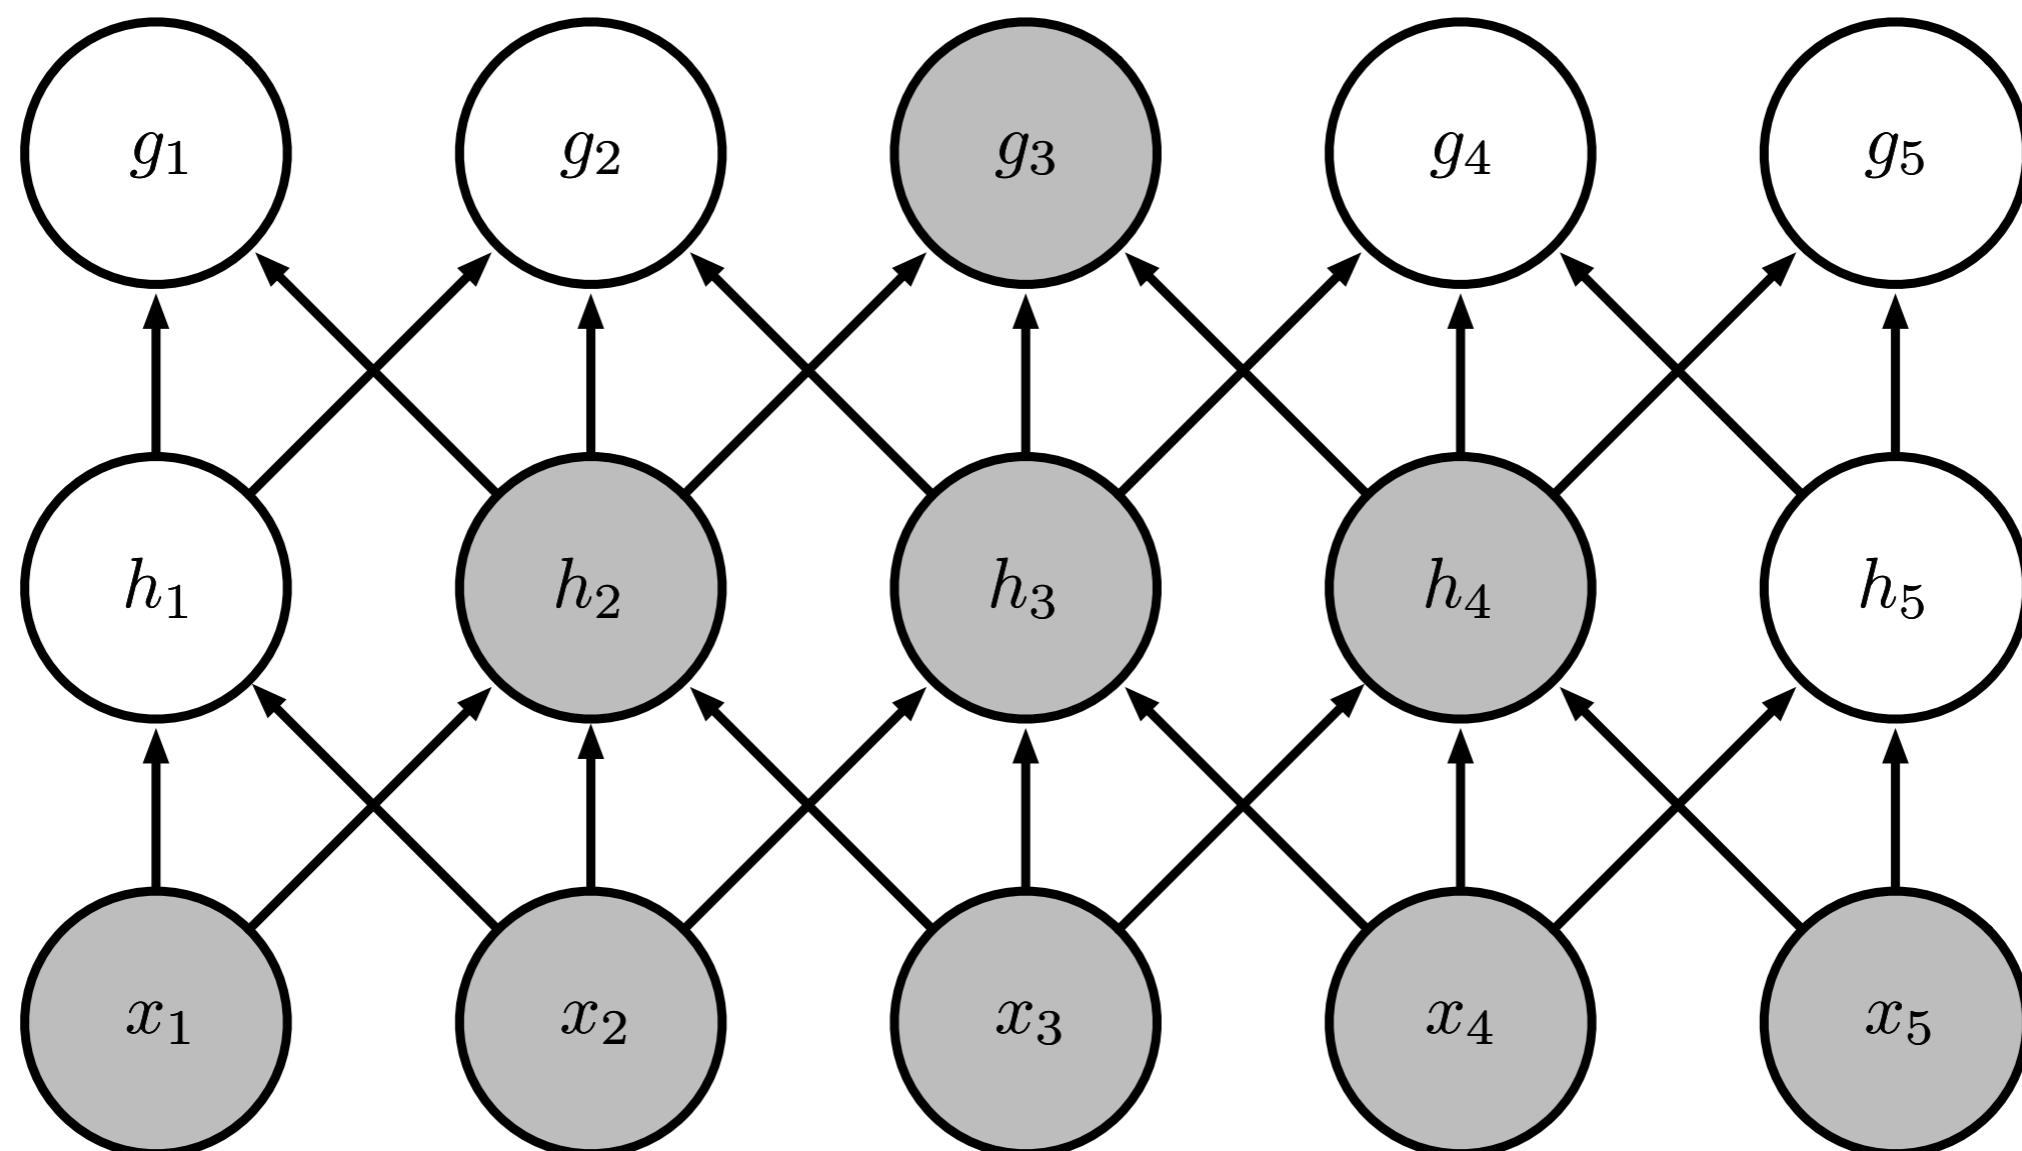
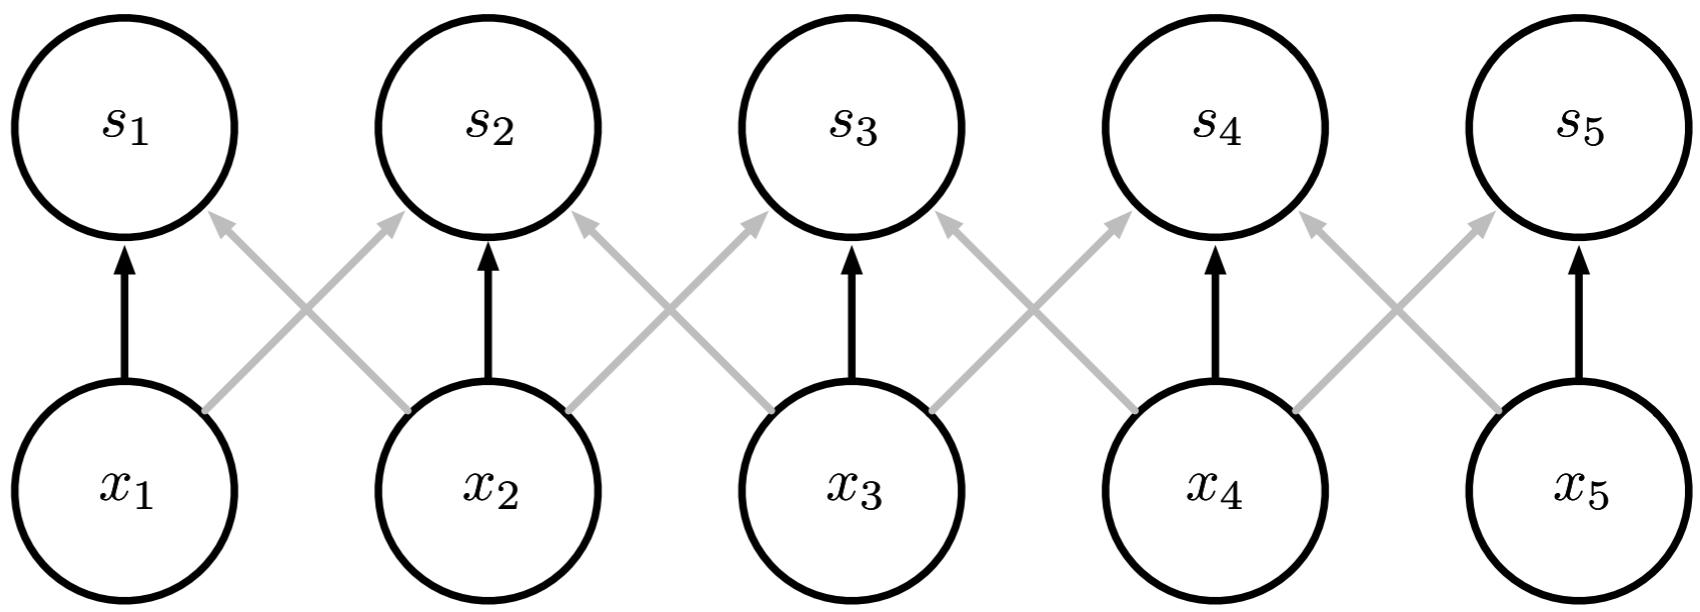


Figure 9.4

# Parameter Sharing

Convolution  
shares the same  
parameters  
across all spatial  
locations



Traditional  
matrix  
multiplication  
does not share  
any parameters

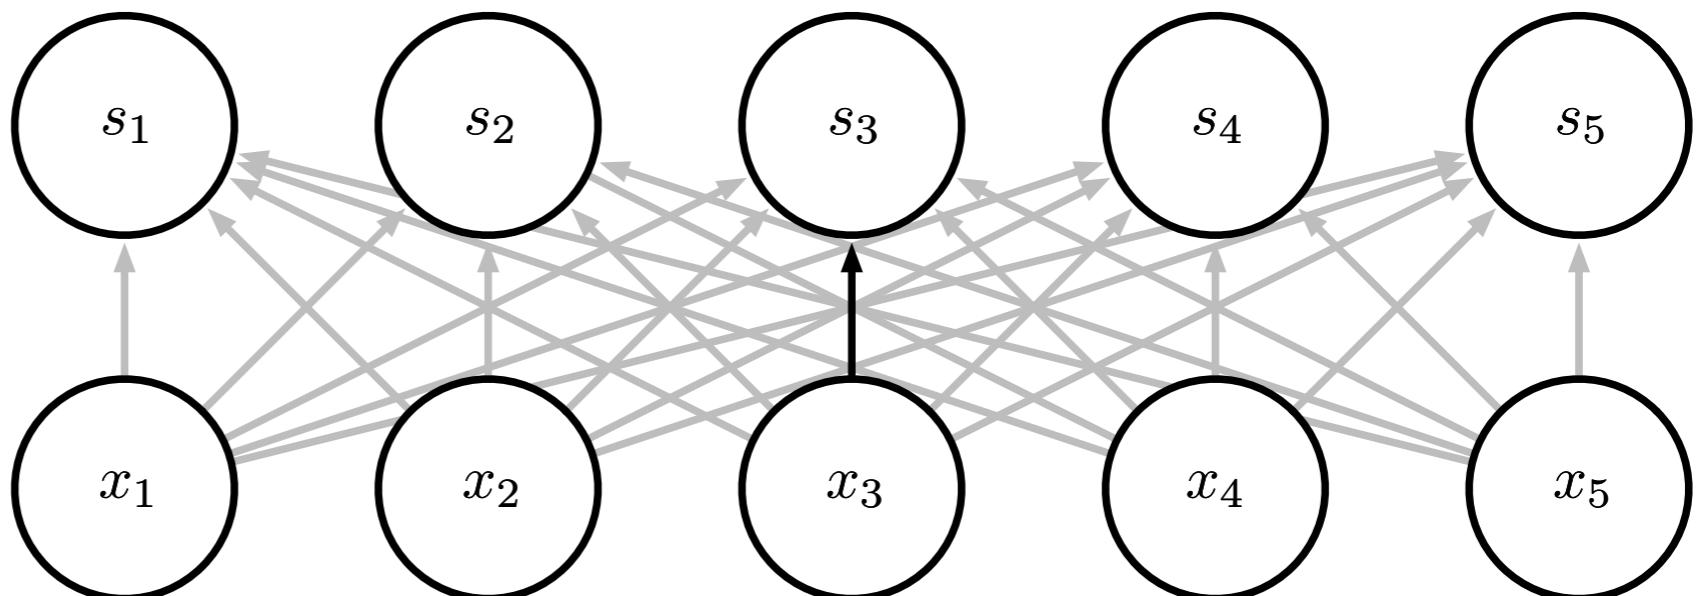


Figure 9.5

# Edge Detection by Convolution

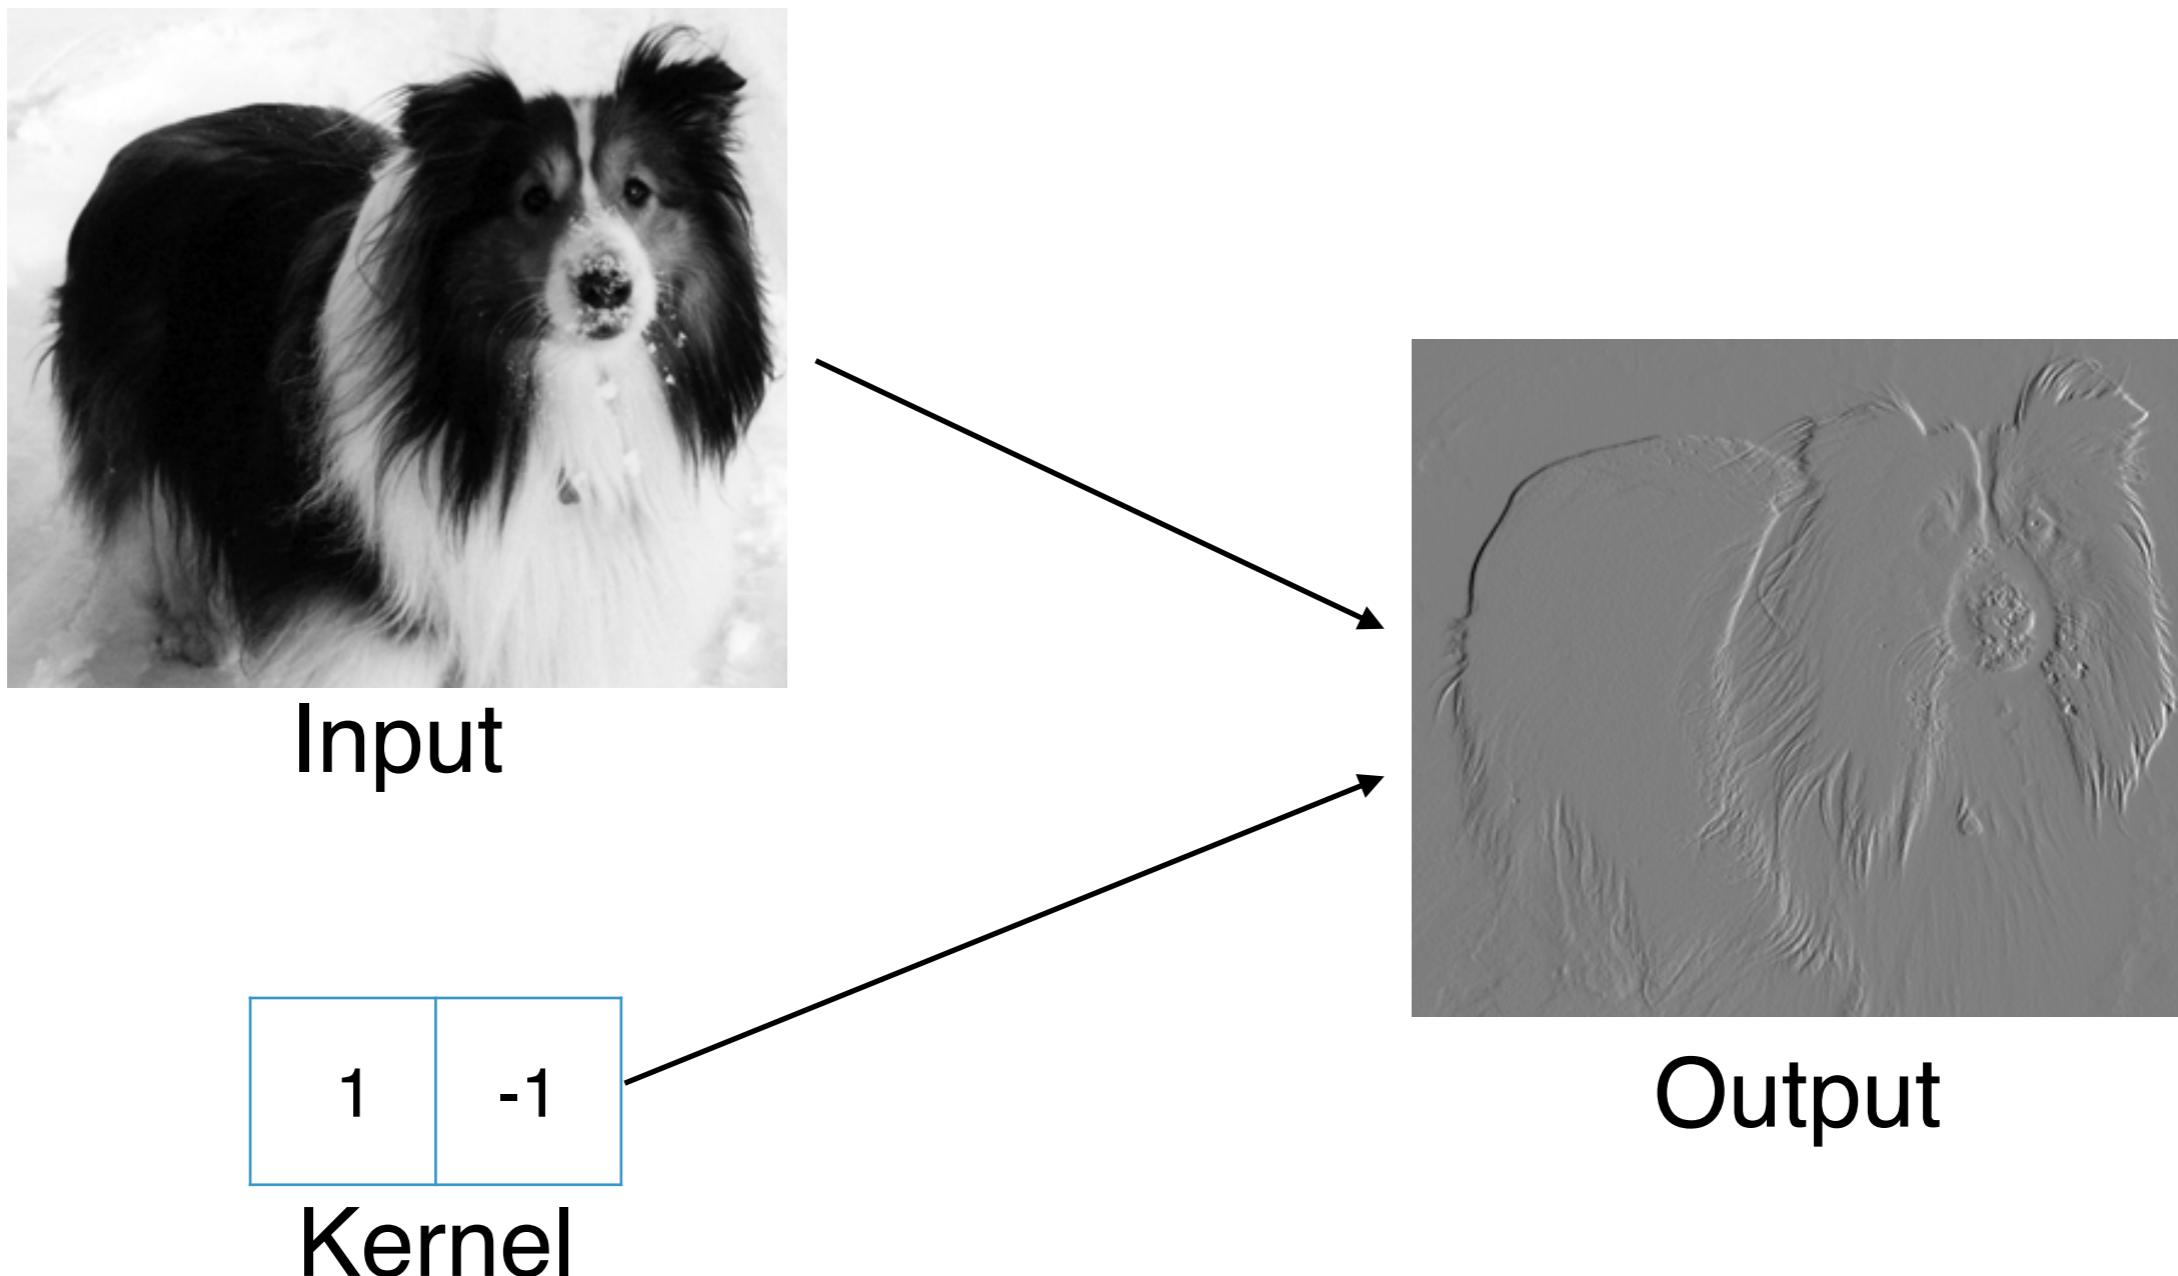


Figure 9.6

# Efficiency of Convolution

Input size: 320 by 280

Kernel size: 2 by 1

Output size: 319 by 280

	Convolution	Dense matrix	Sparse matrix
Stored floats			
Float muls or adds			

# Efficiency of Convolution

Input size: 320 by 280

Kernel size: 2 by 1

Output size: 319 by 280

	Convolution	Dense matrix	Sparse matrix
Stored floats	2	$319 * 280 * 320 * 28 > 8e9$	$2 * 319 * 280 = 178,640$
Float muls or adds	$319 * 280 * 3 = 267,960$	$> 16e9$	Same as convolution (267,960)

# Convolutional Network Components

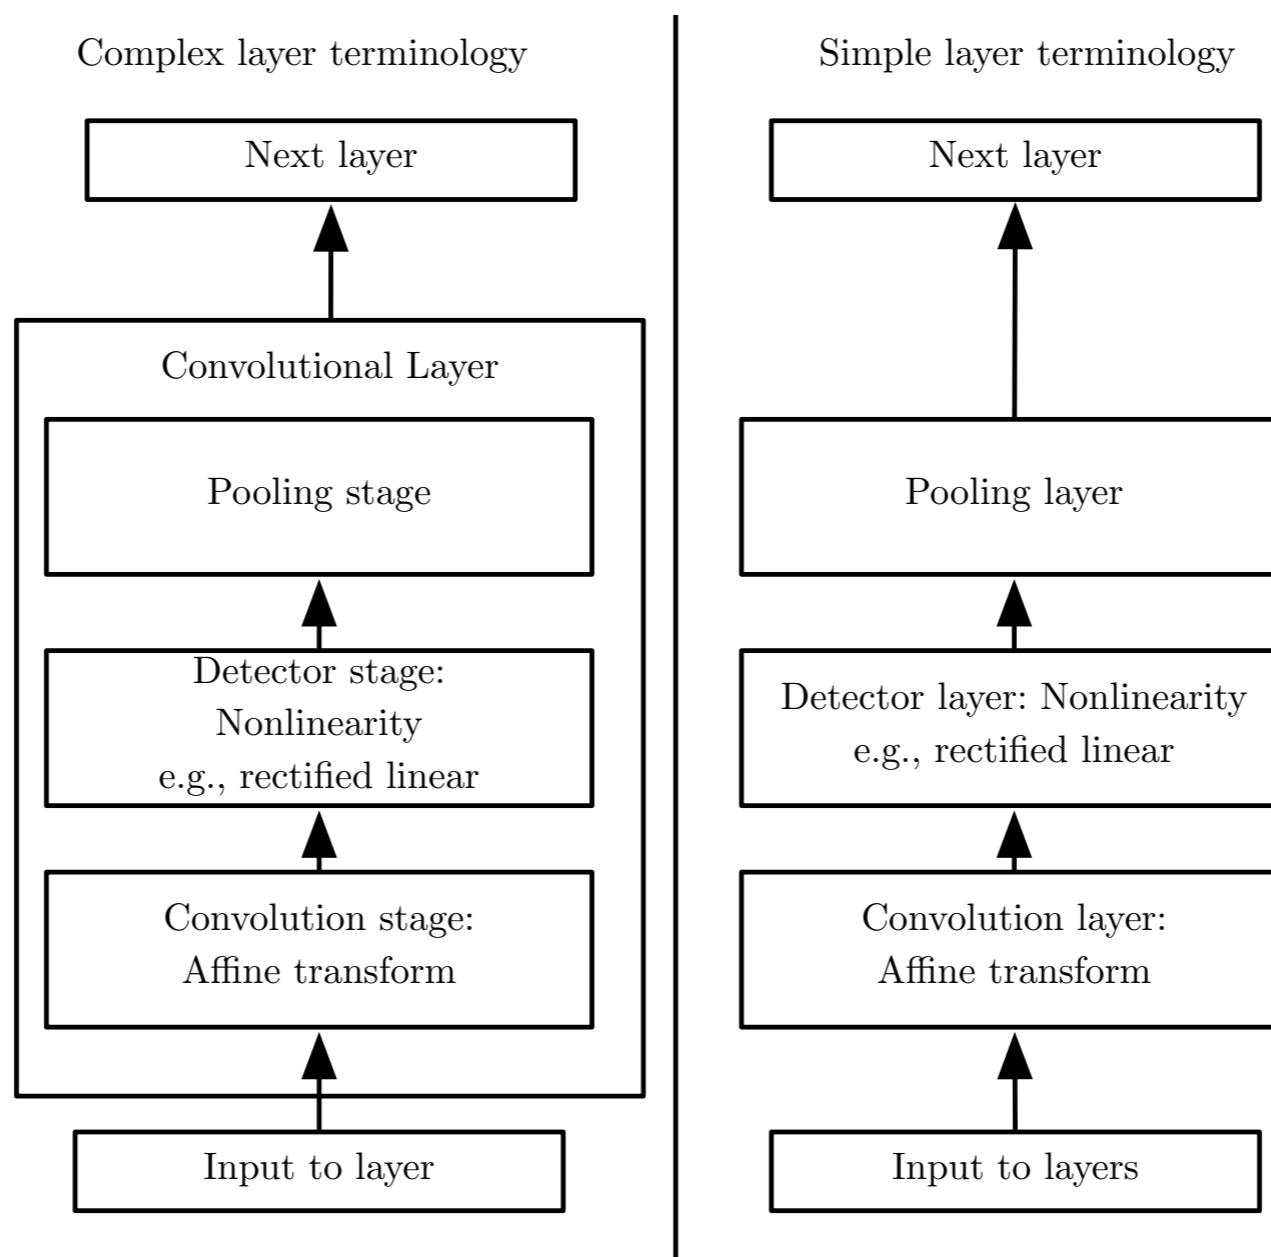
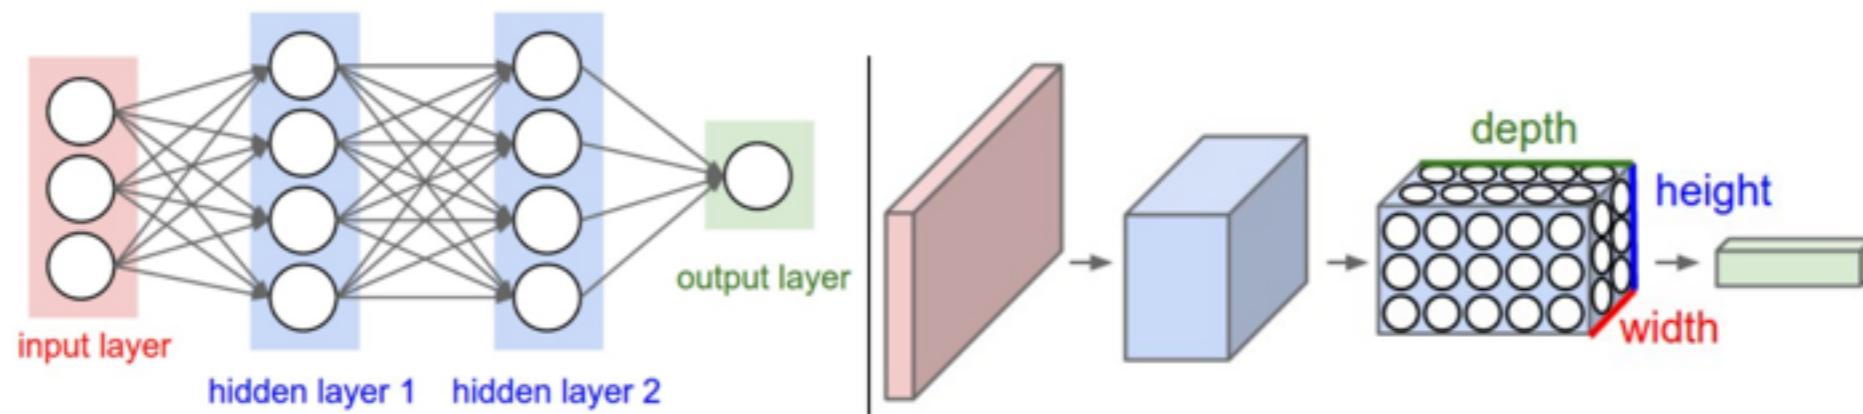


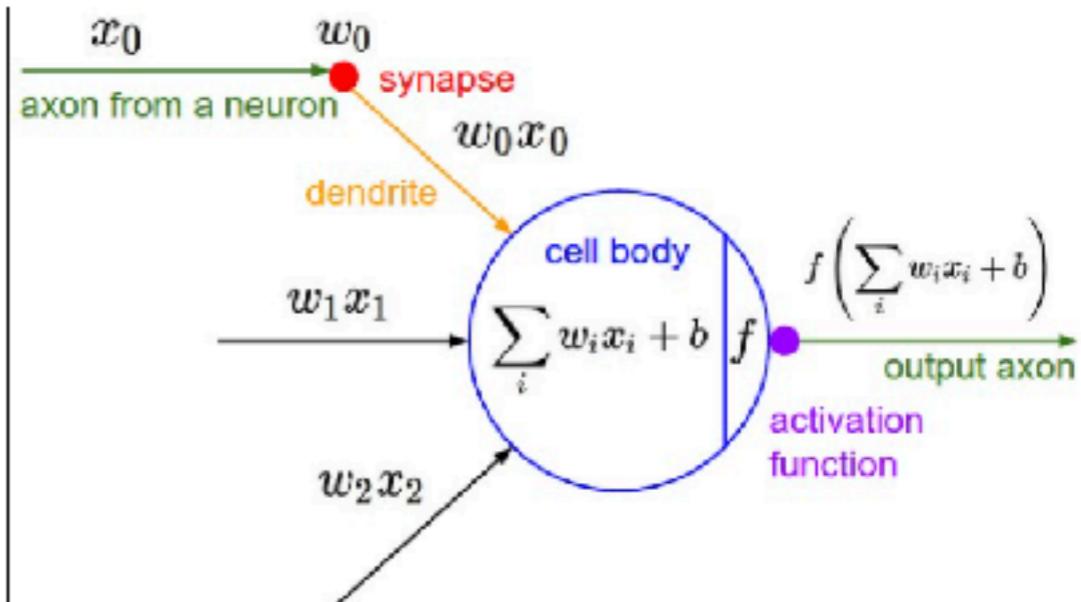
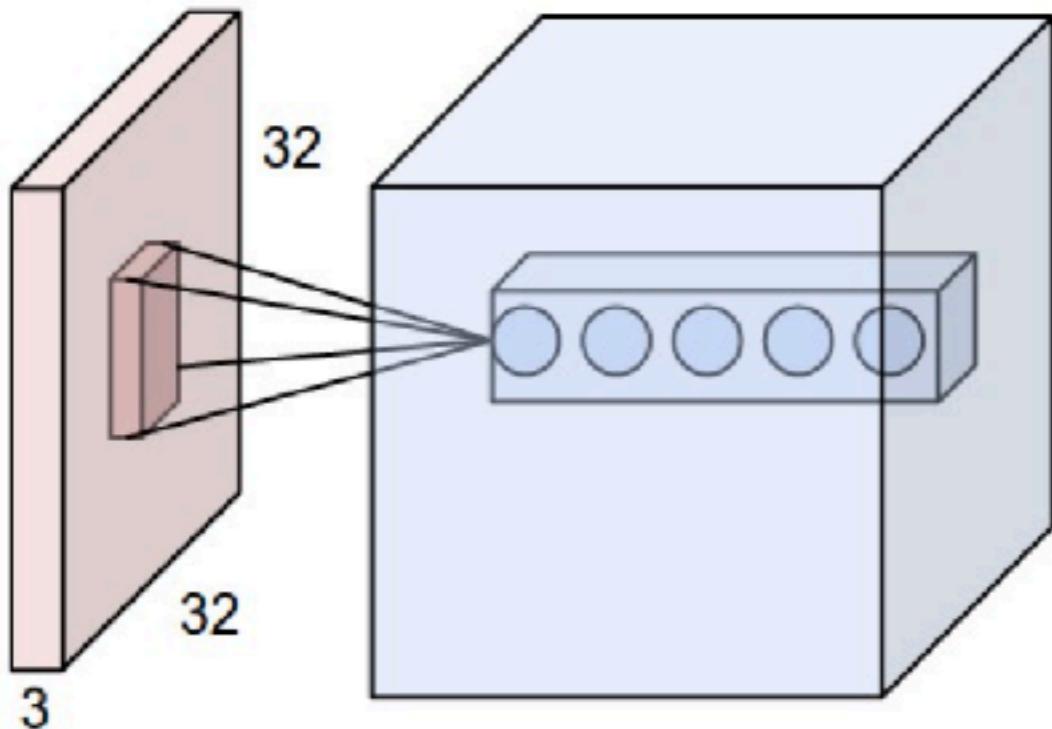
Figure 9.7

# Regular fully-connected NN vs ConvNet



A ConvNet is made up of Layers. Every Layer has a simple API: It transforms an input 3D volume to an output 3D volume with some differentiable function that may or may not have parameters.

# Local Connectivity



*Example 1.* For example, suppose that the input volume has size [32x32x3], (e.g. an RGB CIFAR-10 image). If the receptive field (or the filter size) is 5x5, then each neuron in the Conv Layer will have weights to a [5x5x3] region in the input volume, for a total of  $5 \times 5 \times 3 = 75$  weights (and +1 bias parameter). Notice that the extent of the connectivity along the depth axis must be 3, since this is the depth of the input volume.

*Example 2.* Suppose an input volume had size [16x16x20]. Then using an example receptive field size of 3x3, every neuron in the Conv Layer would now have a total of  $3 \times 3 \times 20 = 180$  connections to the input volume. Notice that, again, the connectivity is local in space (e.g. 3x3), but full along the input depth (20).

# Example of learned kernels



# Max Pooling and Invariance to Translation

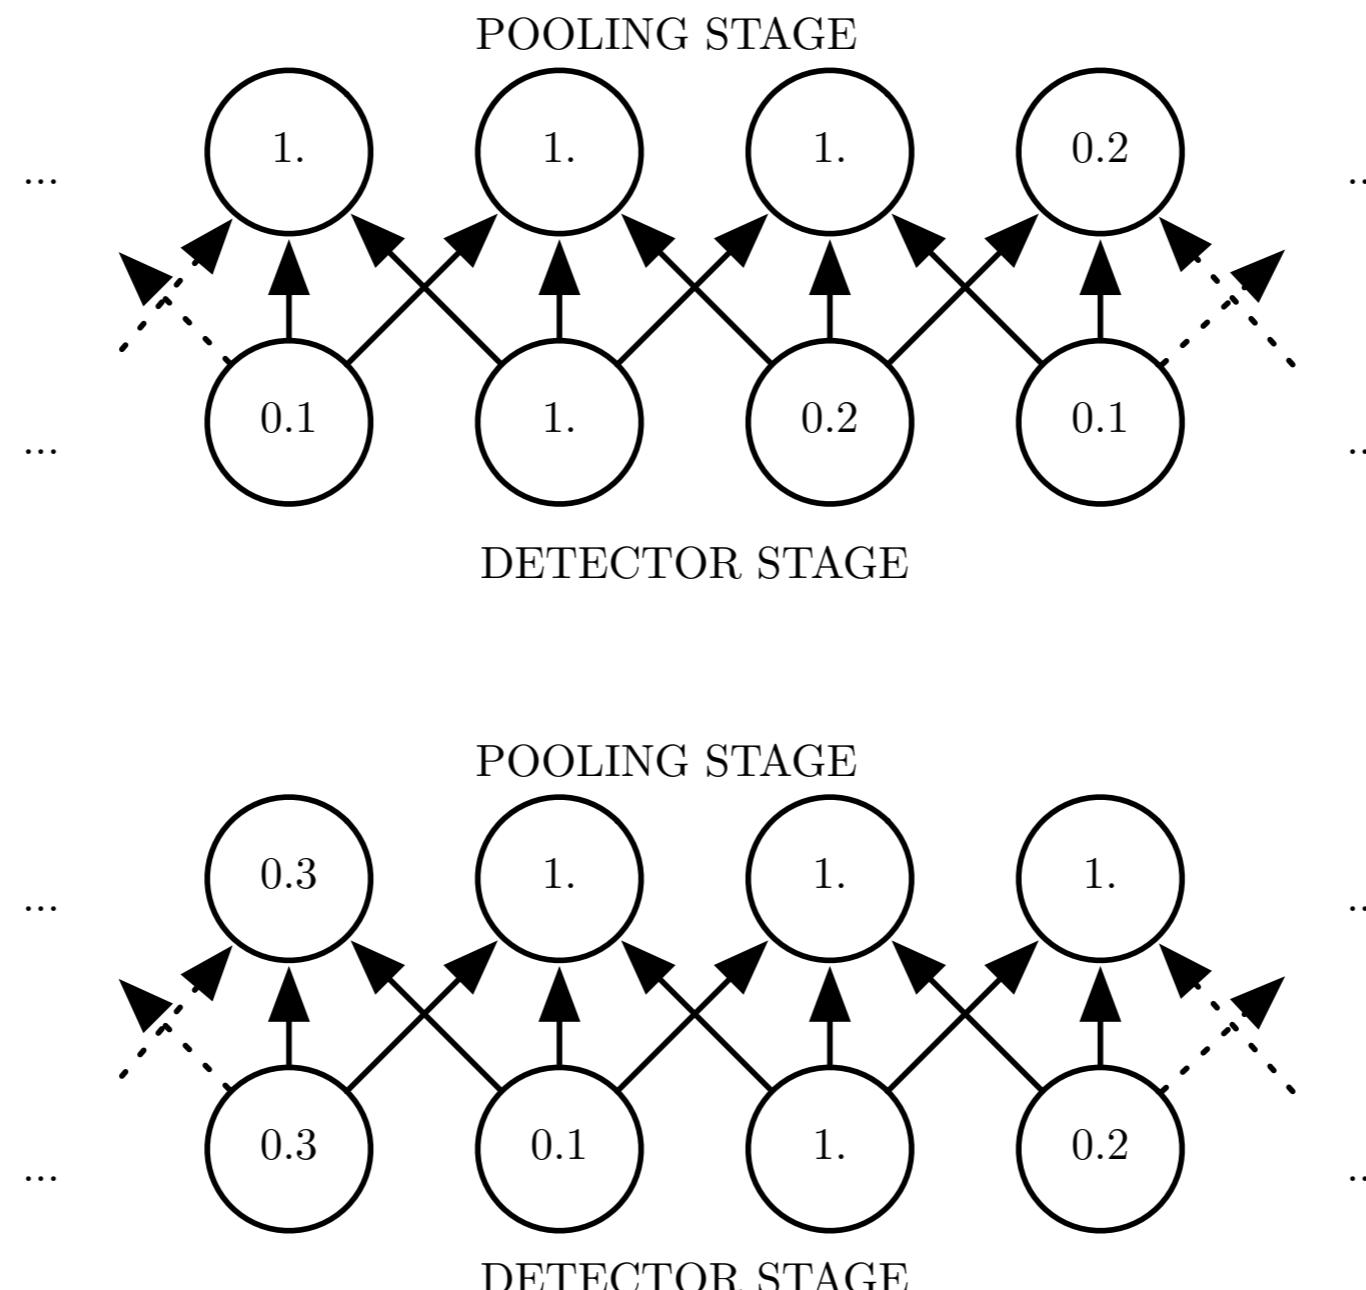


Figure 9.8

# Cross-Channel Pooling and Invariance to Learned Transformations

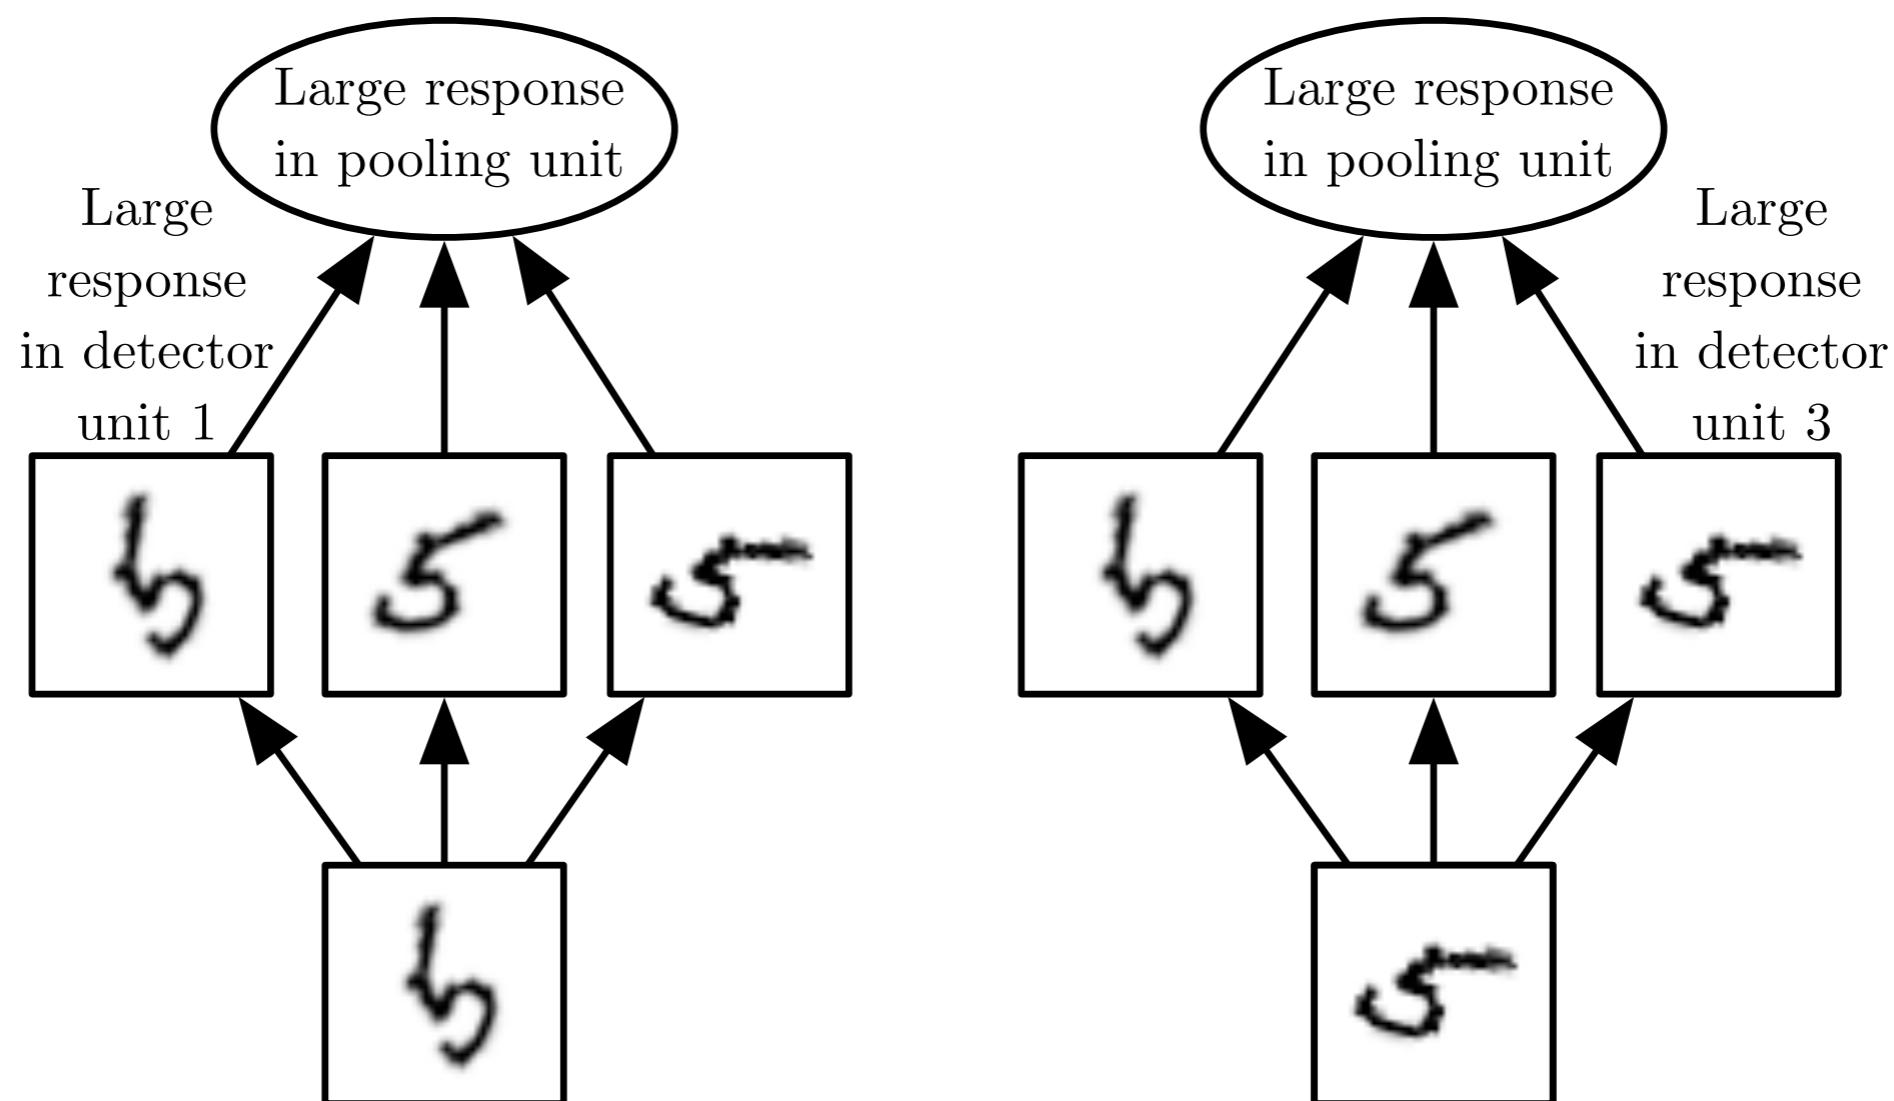


Figure 9.9

# Pooling with Downsampling

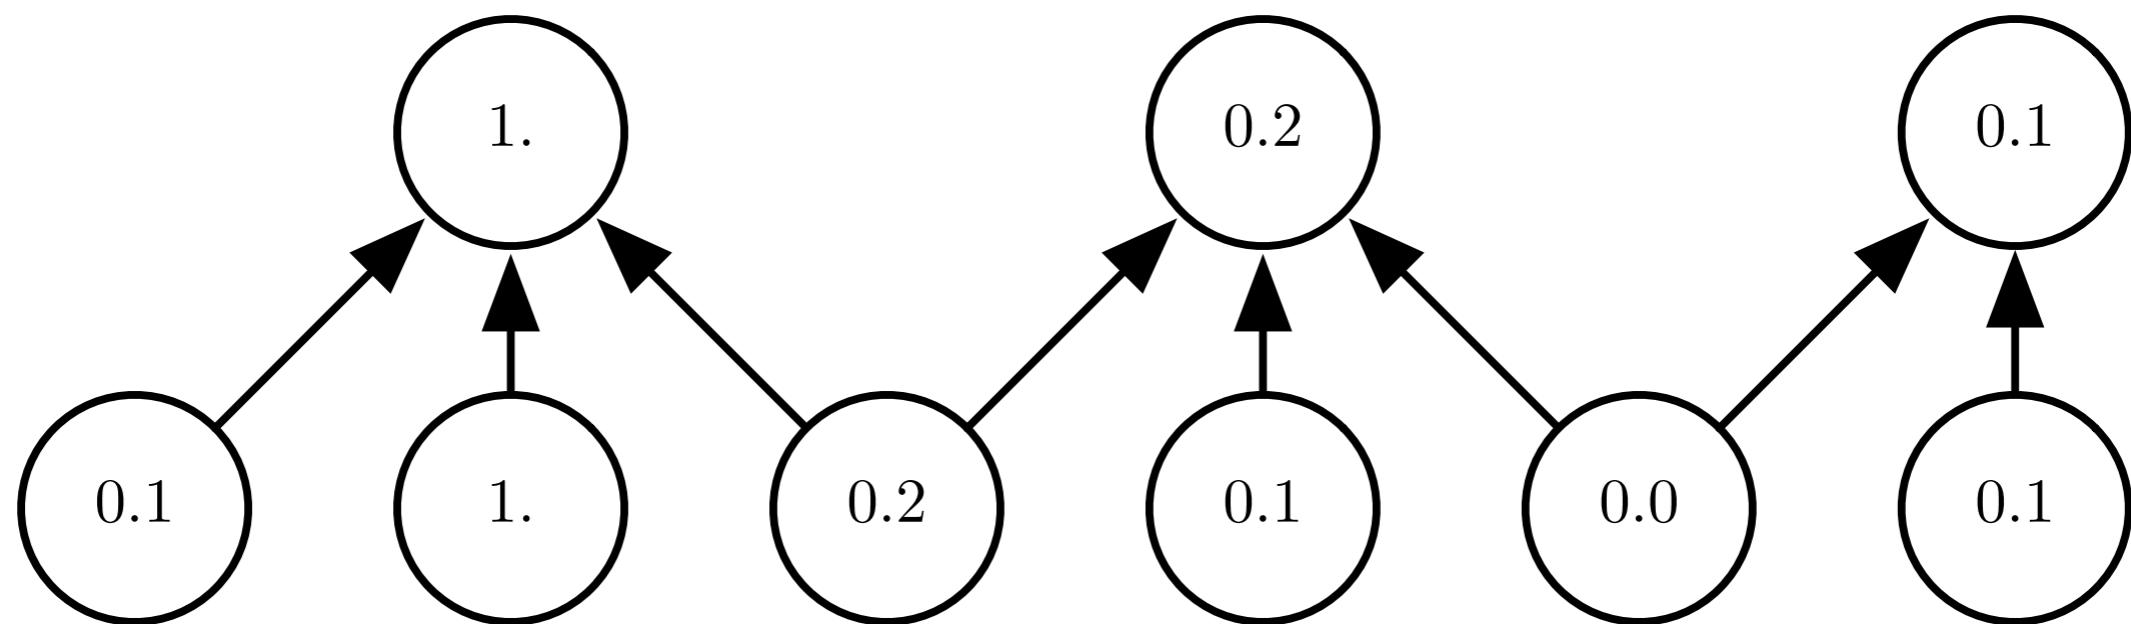
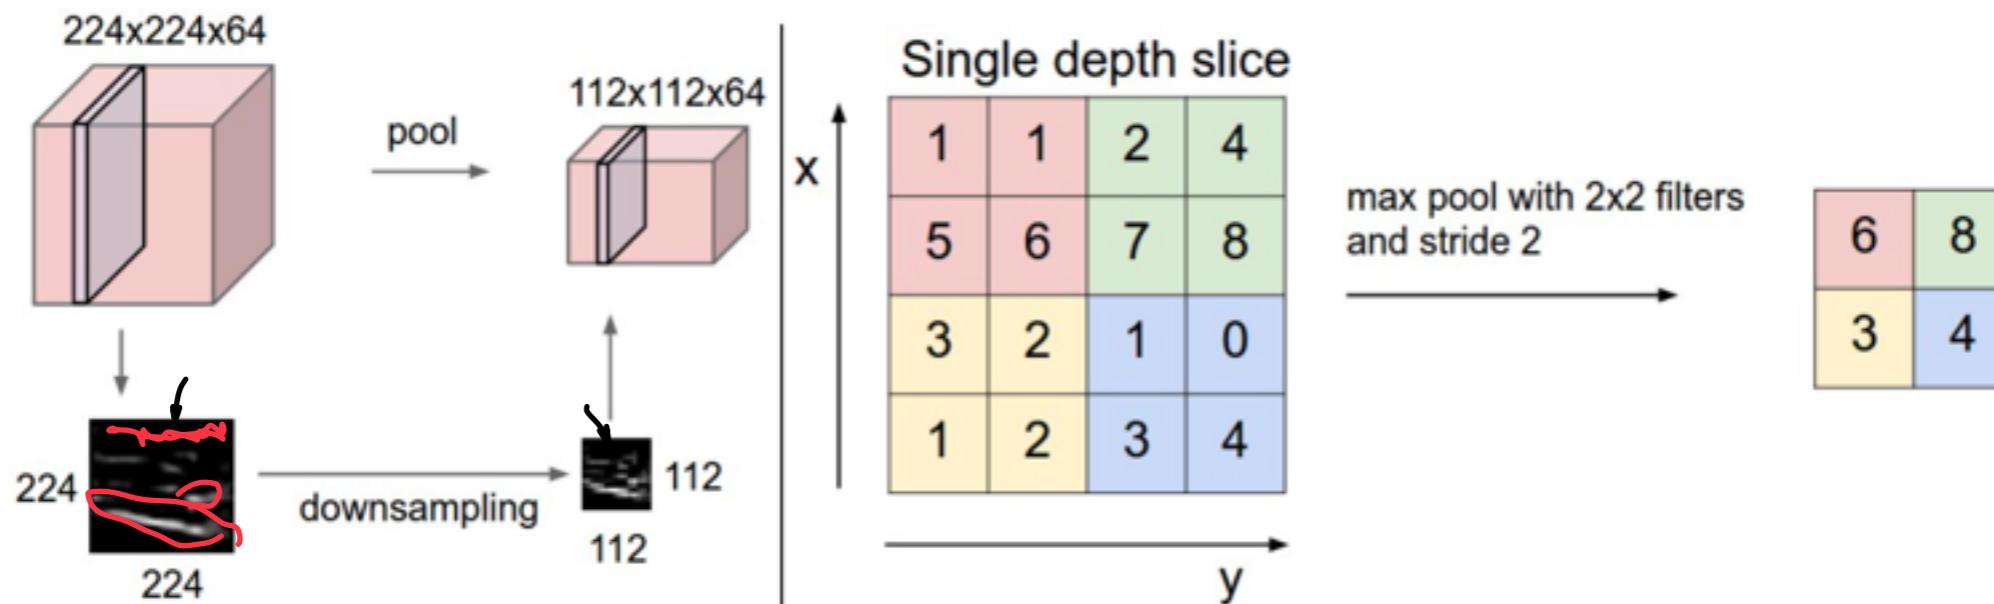


Figure 9.10

# Pooling layer downsamples the volume spatially



# Example Classification Architectures

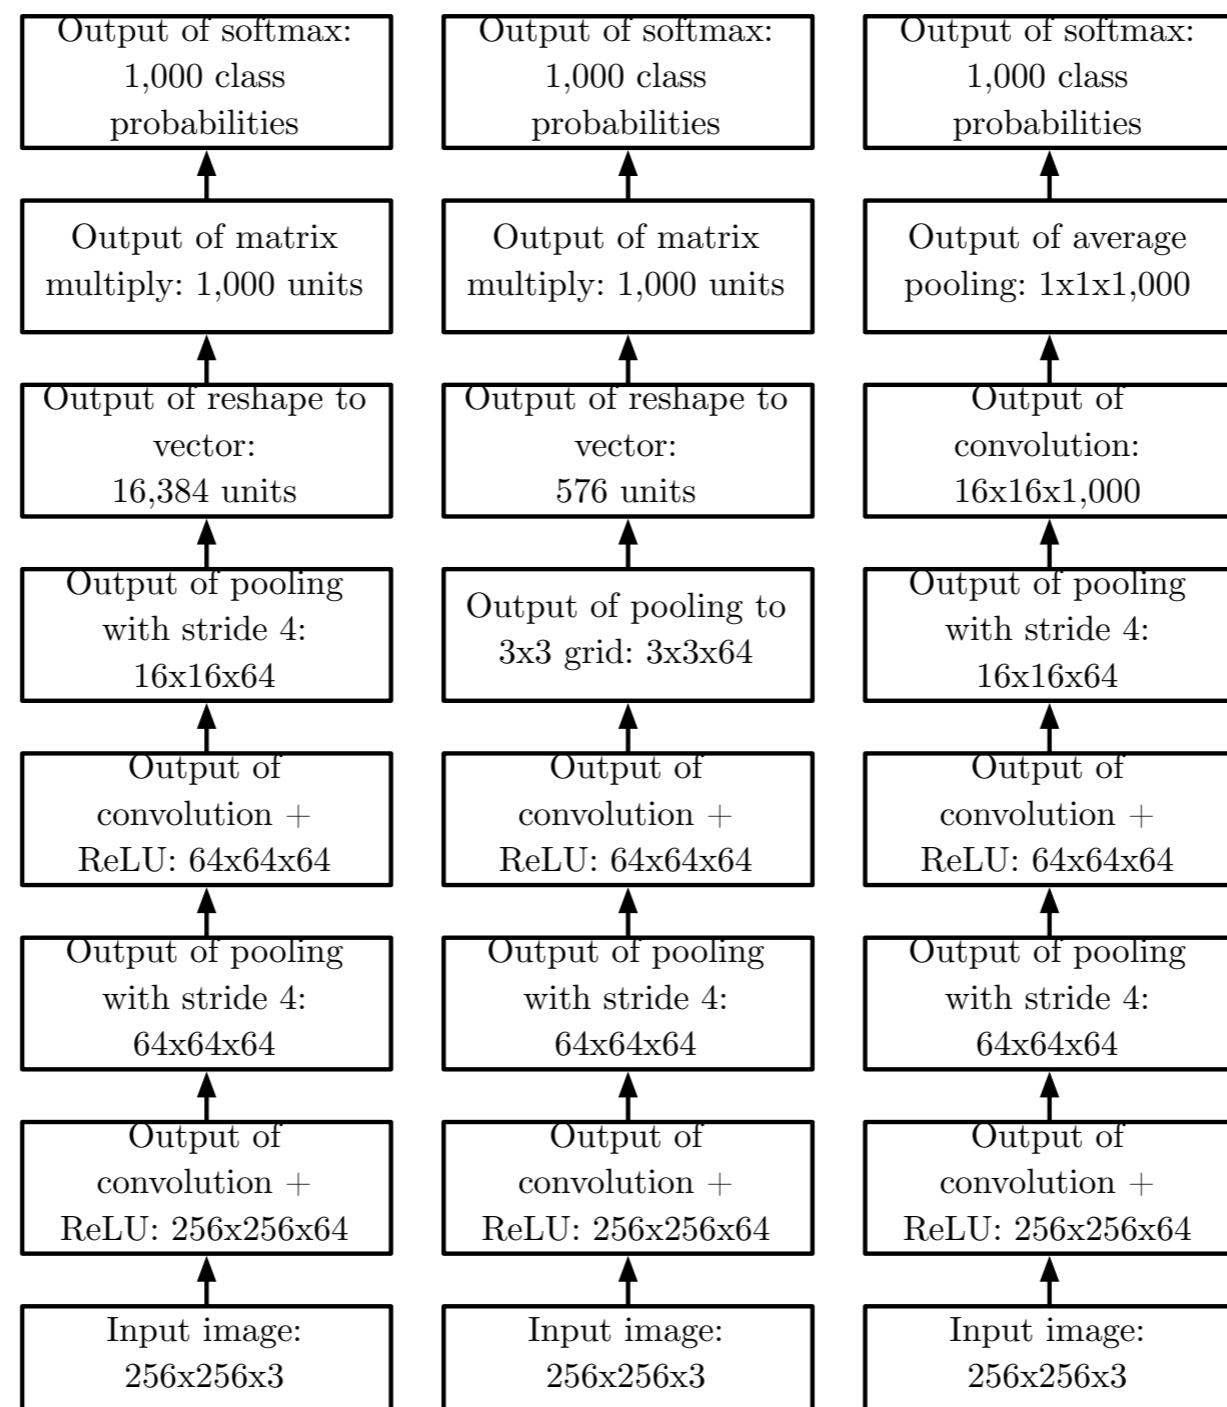
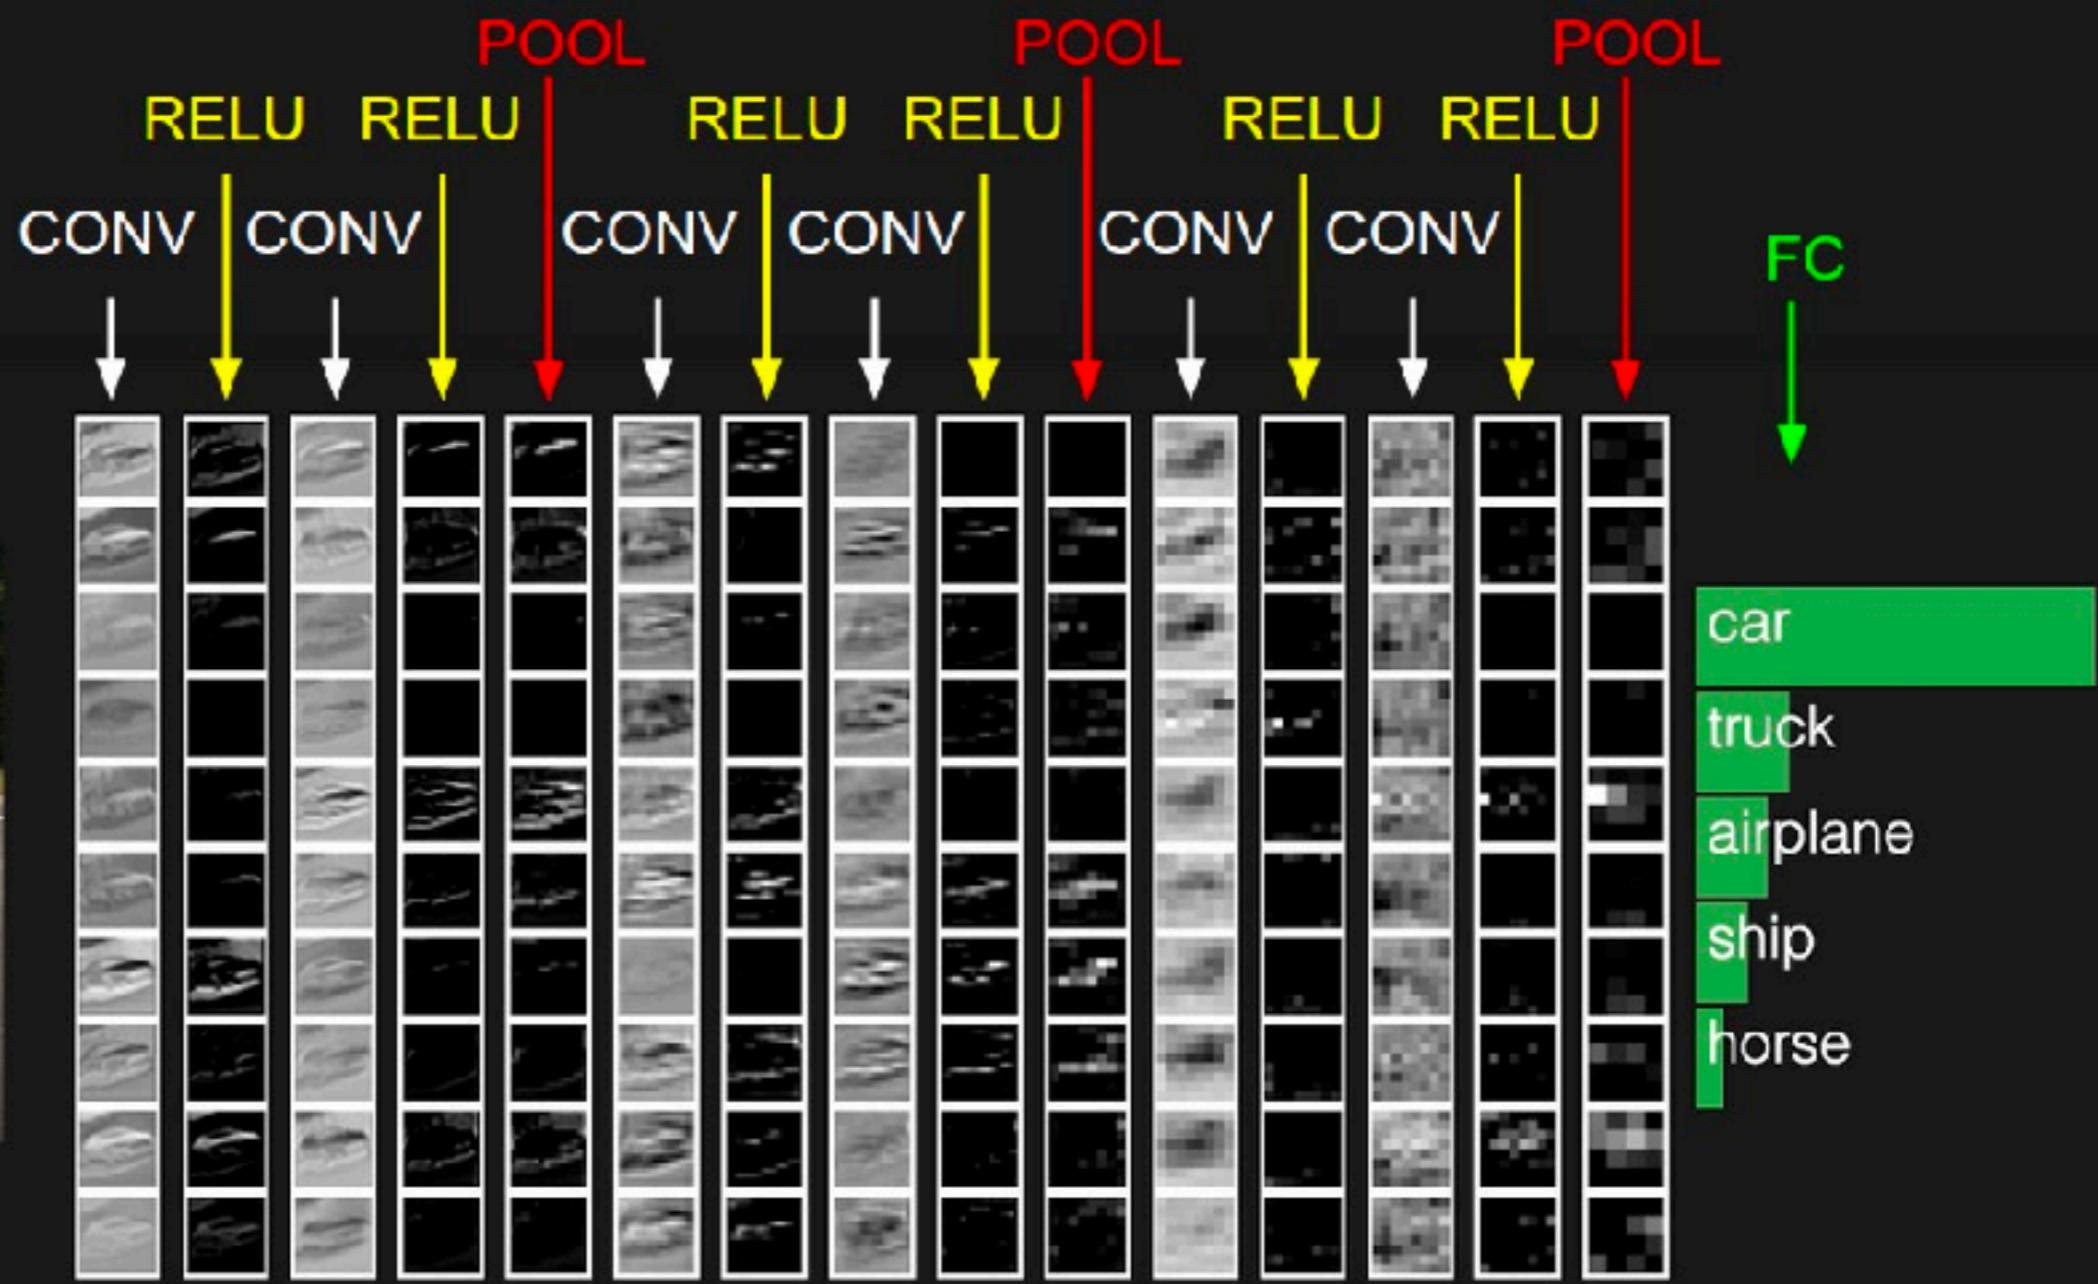


Figure 9.11

# ConvNet architecture



# Architecture Overview of ConvNets

[INPUT - CONV - RELU - POOL - FC]

- INPUT [32x32x3] will hold the raw pixel values of the image, in this case an image of width 32, height 32, and with three color channels R,G,B.
- CONV layer will compute the output of neurons that are connected to local regions in the input, each computing a dot product between their weights and a small region they are connected to in the input volume. This may result in volume such as [32x32x12] if we decided to use 12 filters.
- RELU layer will apply an elementwise activation function, such as the  $\max(0, x)$  thresholding at zero. This leaves the size of the volume unchanged ([32x32x12]).
- POOL layer will perform a downsampling operation along the spatial dimensions (width, height), resulting in volume such as [16x16x12].
- FC (i.e. fully-connected) layer will compute the class scores, resulting in volume of size [1x1x10], where each of the 10 numbers correspond to a class score, such as among the 10 categories of CIFAR-10. As with ordinary Neural Networks and as the name implies, each neuron in this layer will be connected to all the numbers in the previous volume.

# Architecture Overview of ConvNets

- A ConvNet architecture is in the simplest case a list of Layers that transform the image volume into an output volume (e.g. holding the class scores)
- There are a few distinct types of Layers (e.g. CONV/FC/RELU/POOL are by far the most popular)
- Each Layer accepts an input 3D volume and transforms it to an output 3D volume through a differentiable function
- Each Layer may or may not have parameters (e.g. CONV/FC do, RELU/POOL don't)
- Each Layer may or may not have additional hyperparameters (e.g. CONV/FC/POOL do, RELU doesn't)

# Convolution with Stride

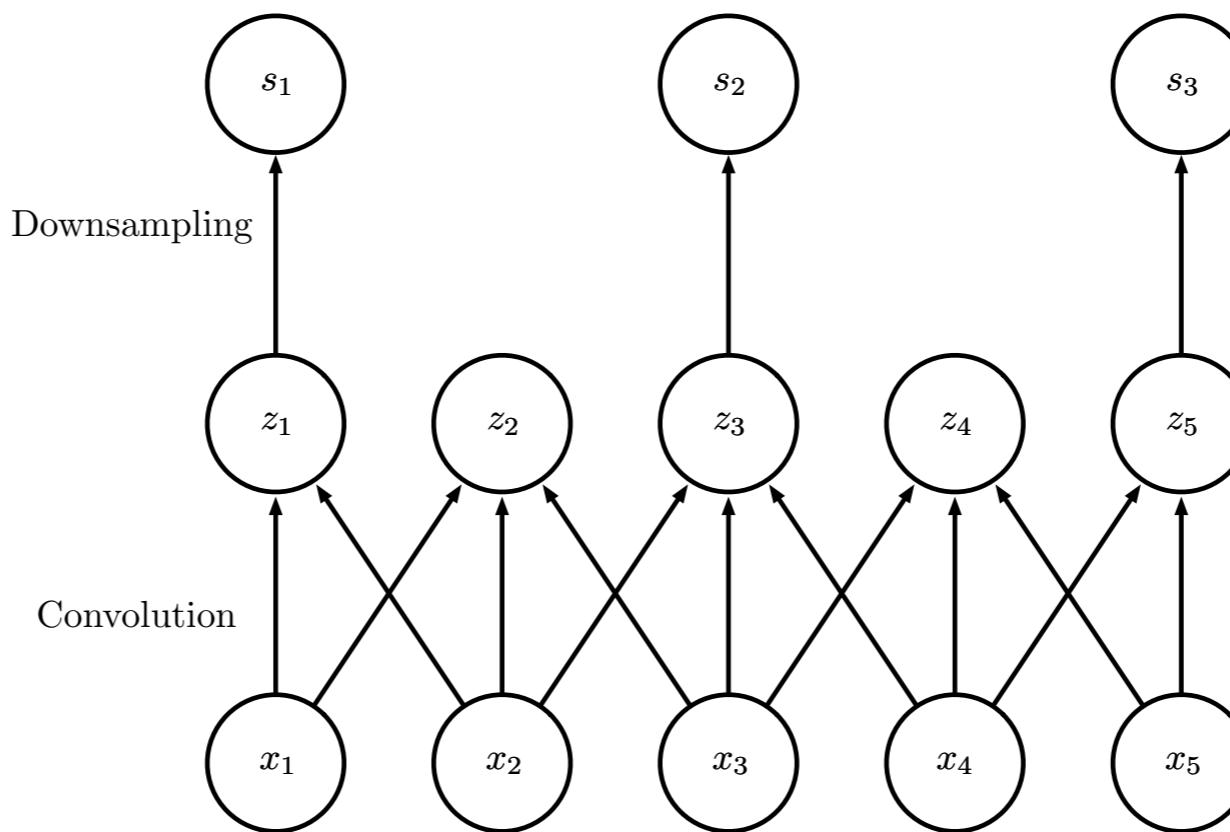
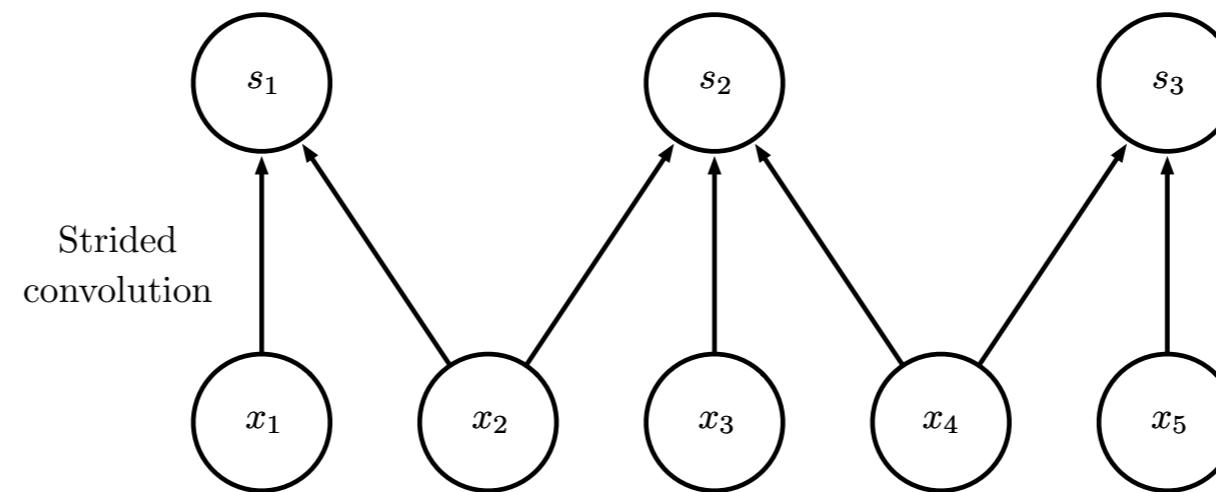


Figure 9.12

# Zero Padding Controls Size

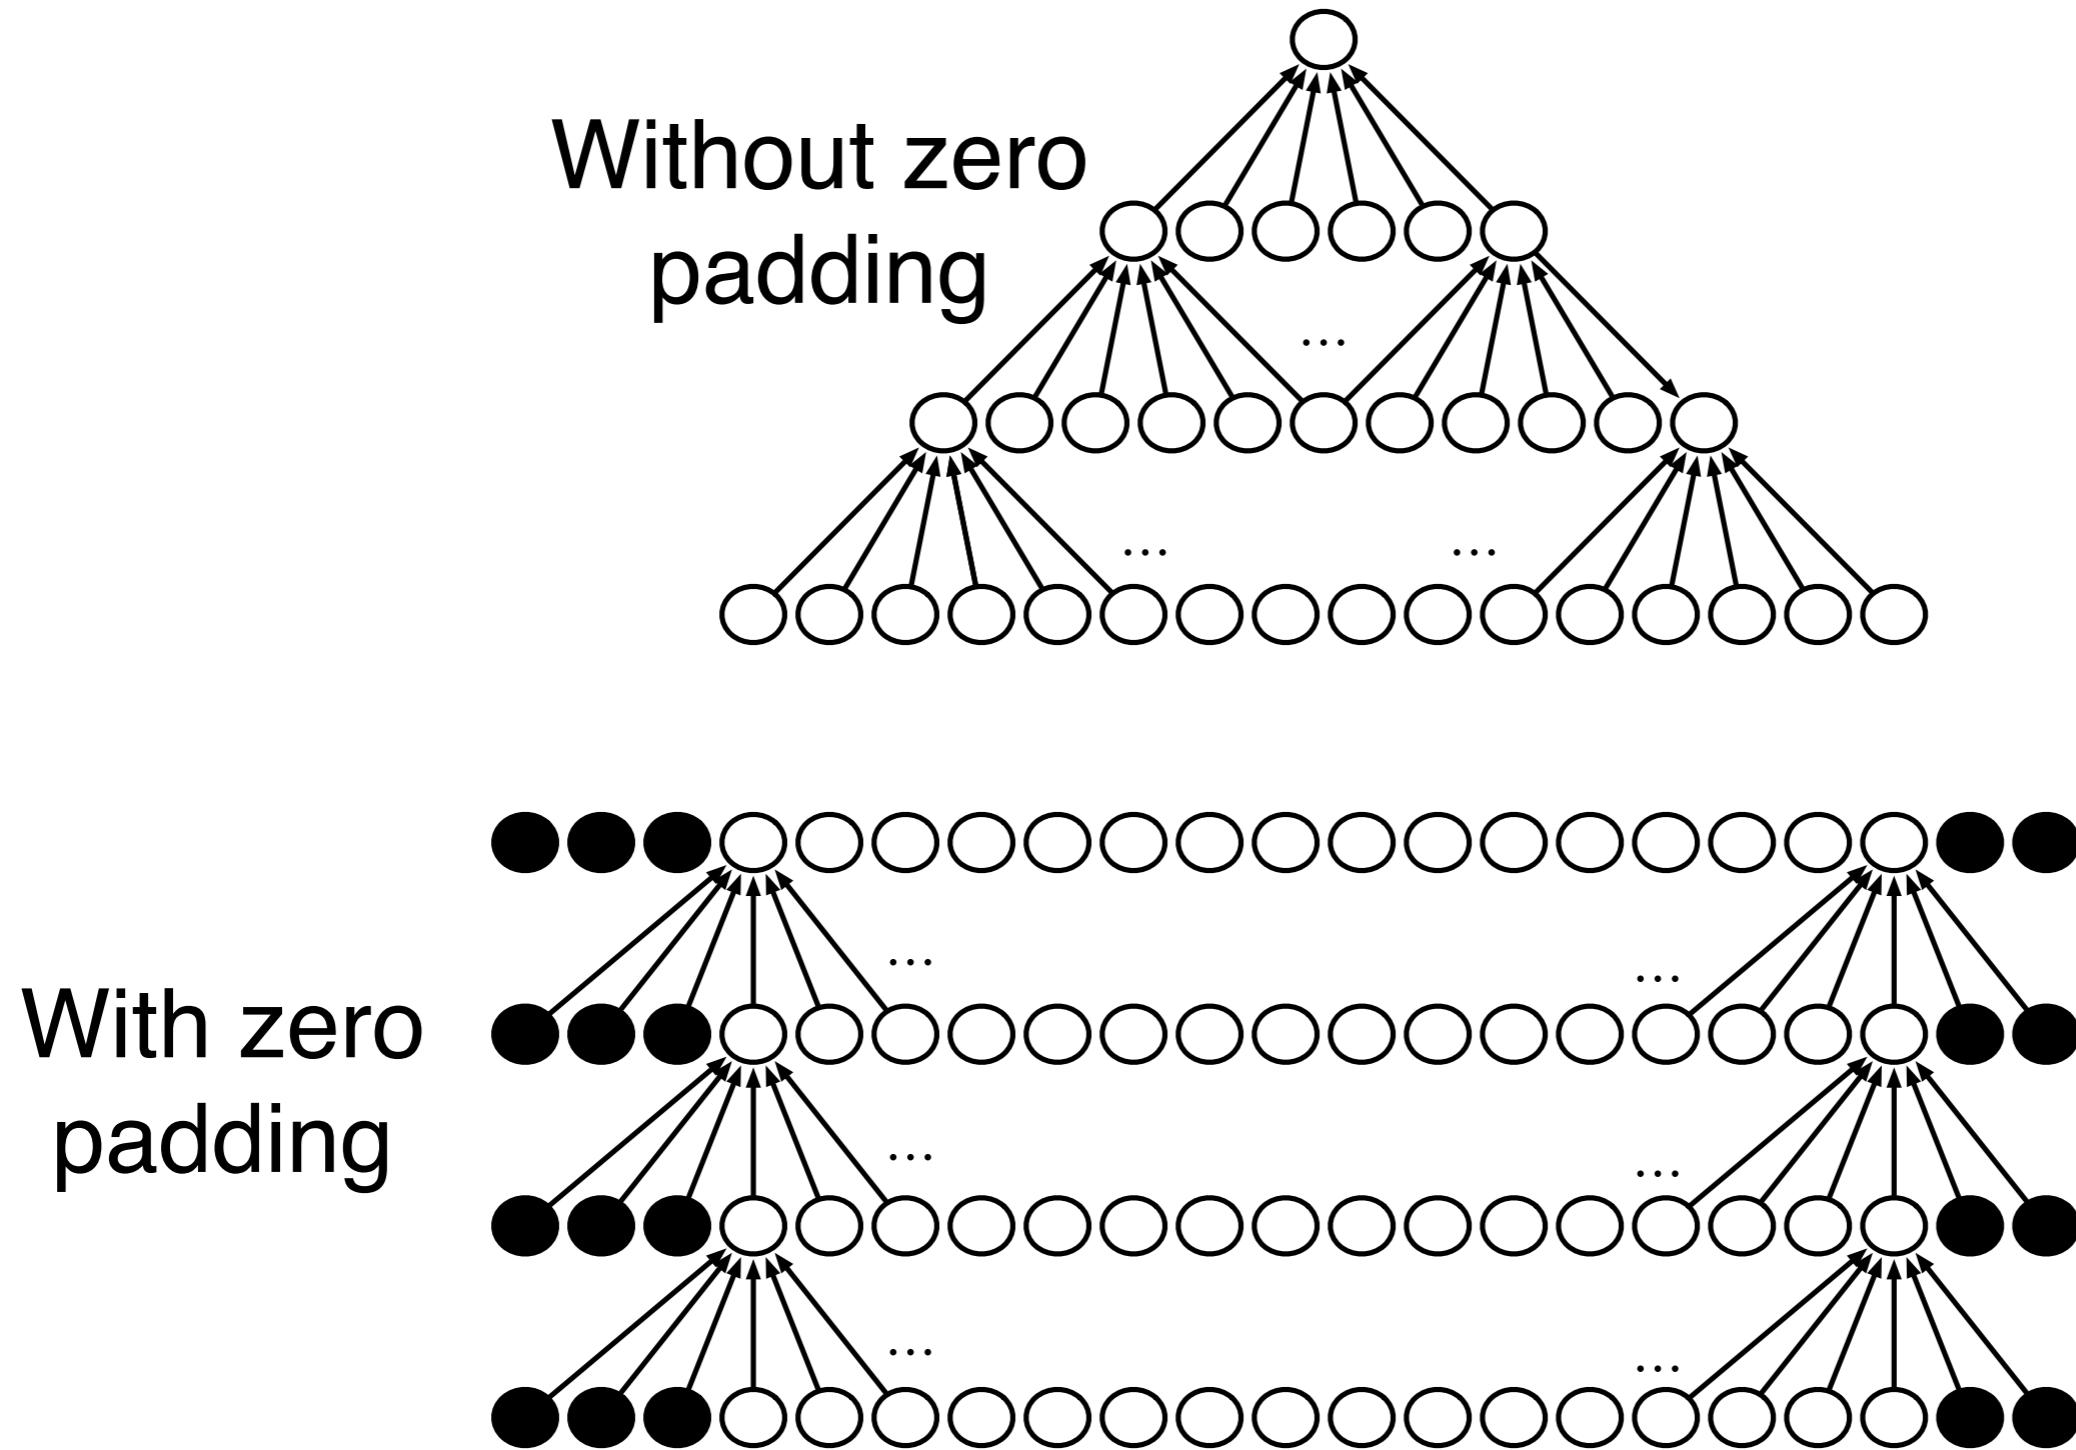
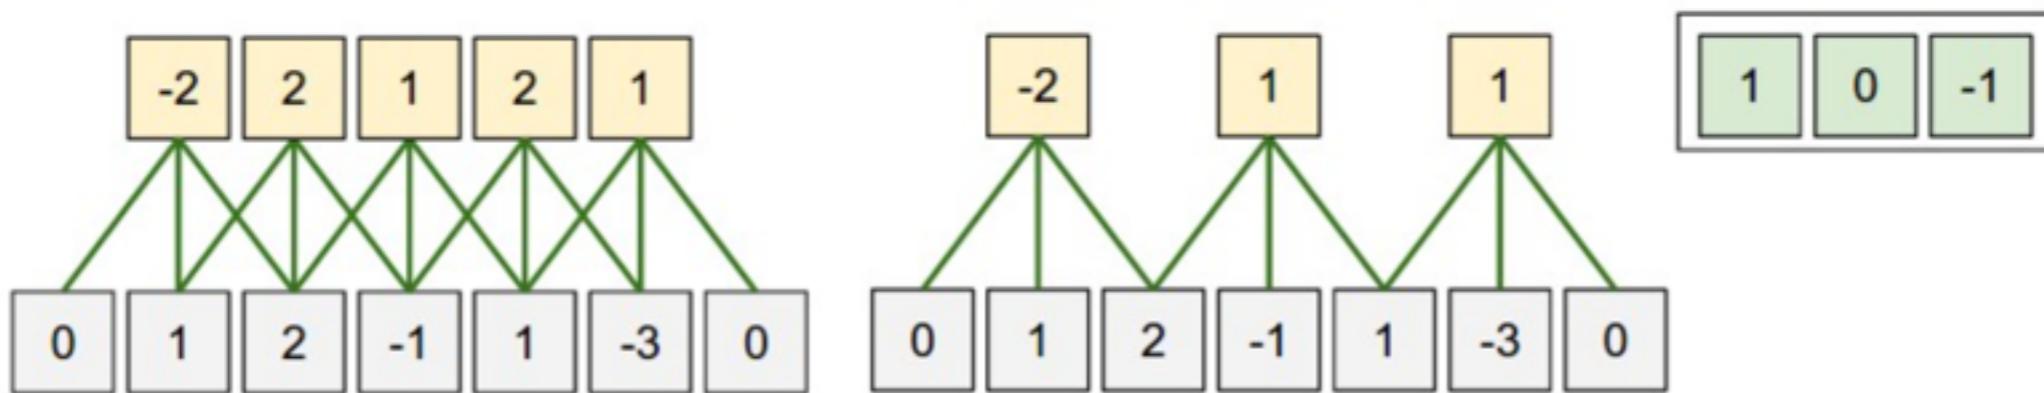
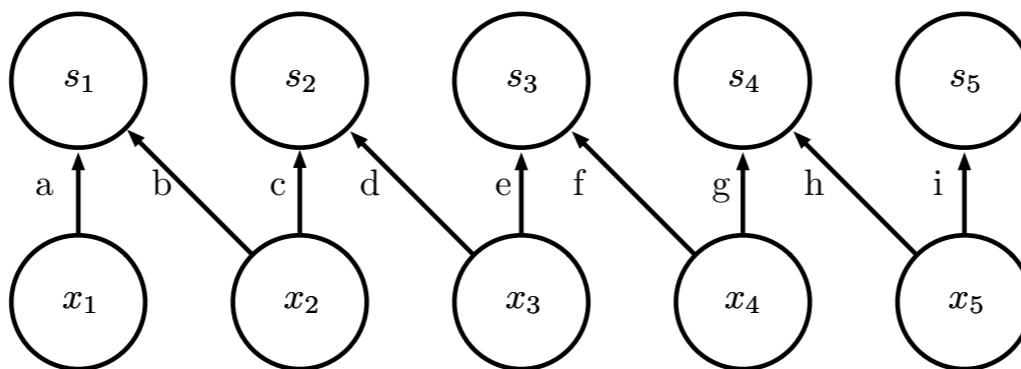


Figure 9.13

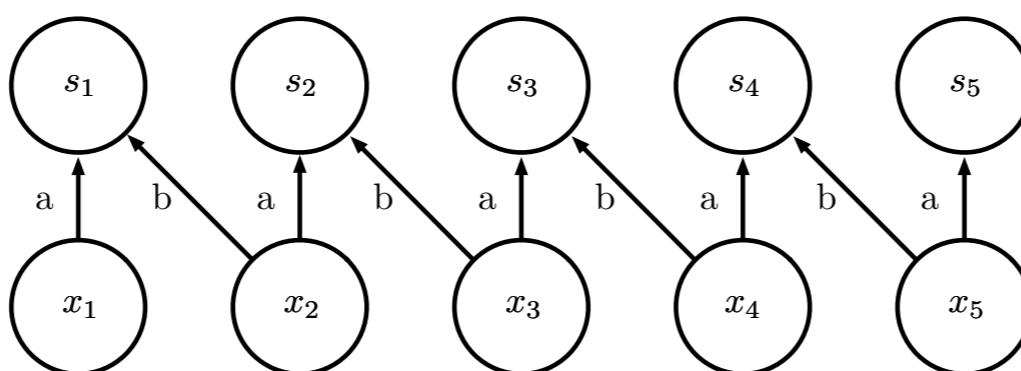
# Output size with zero padding and stride



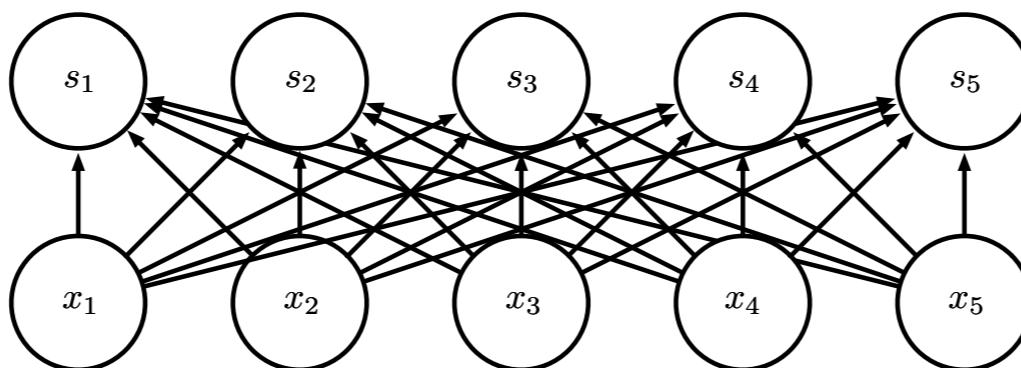
# Kinds of Connectivity



Local connection:  
like convolution,  
but no sharing



Convolution



Fully connected

Figure 9.14

# Partial Connectivity Between Channels

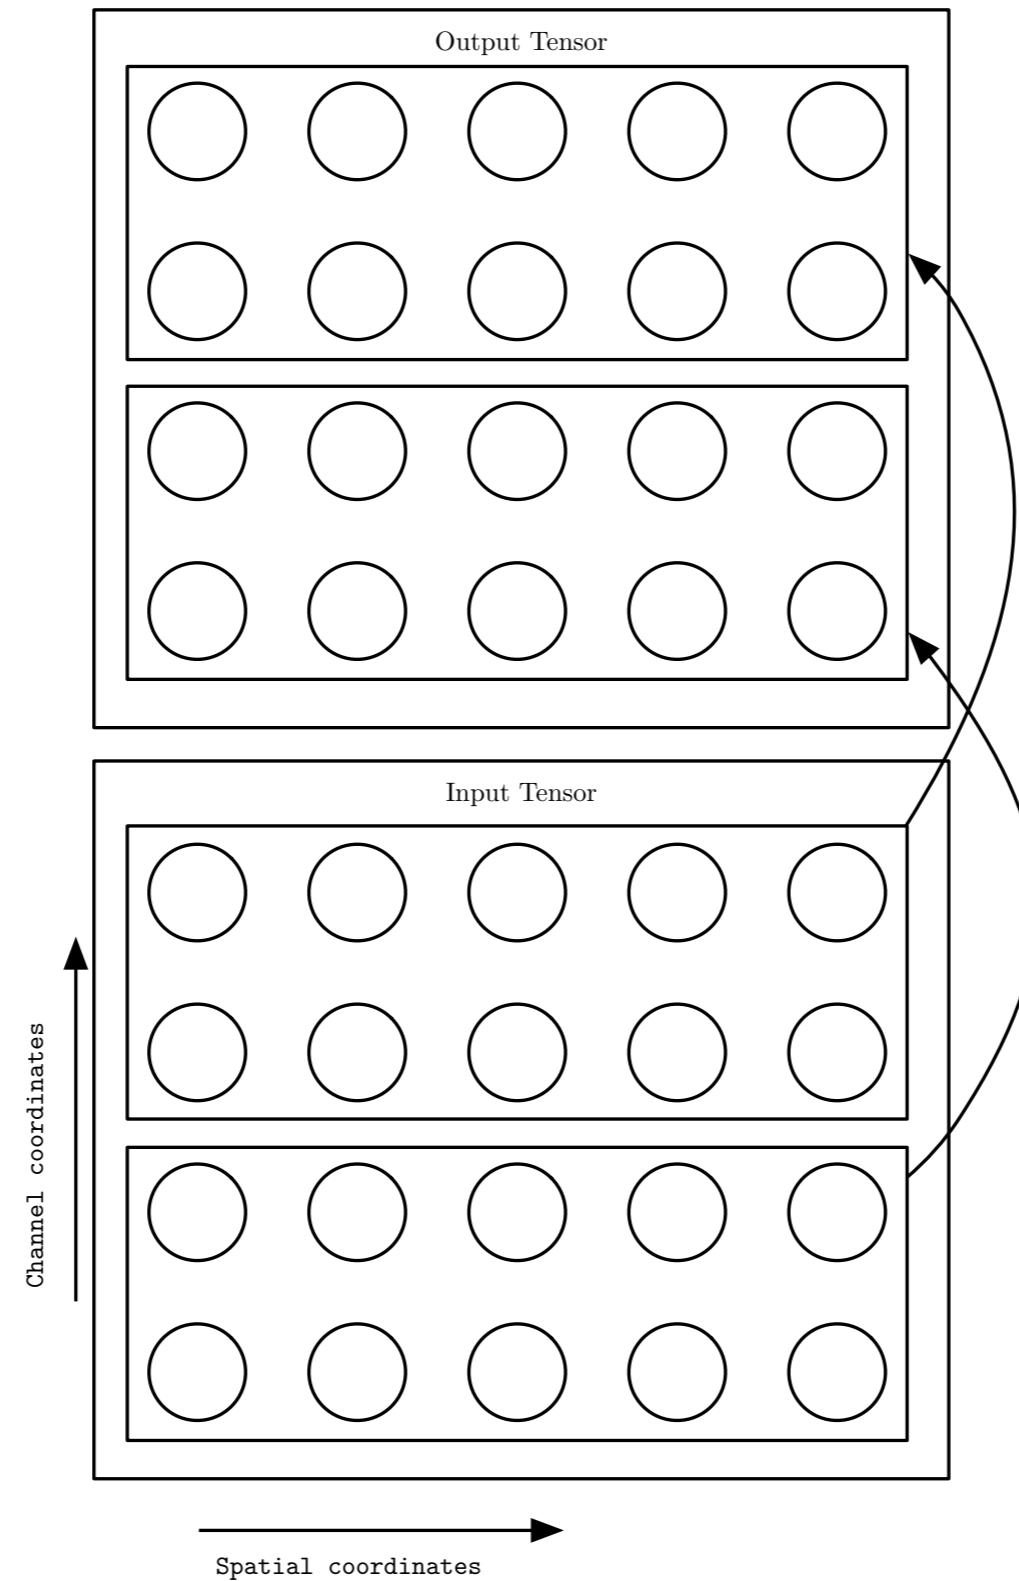
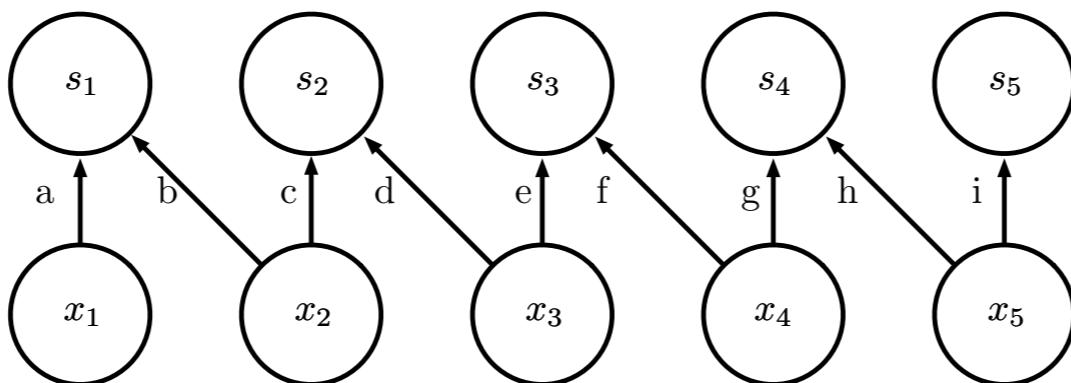
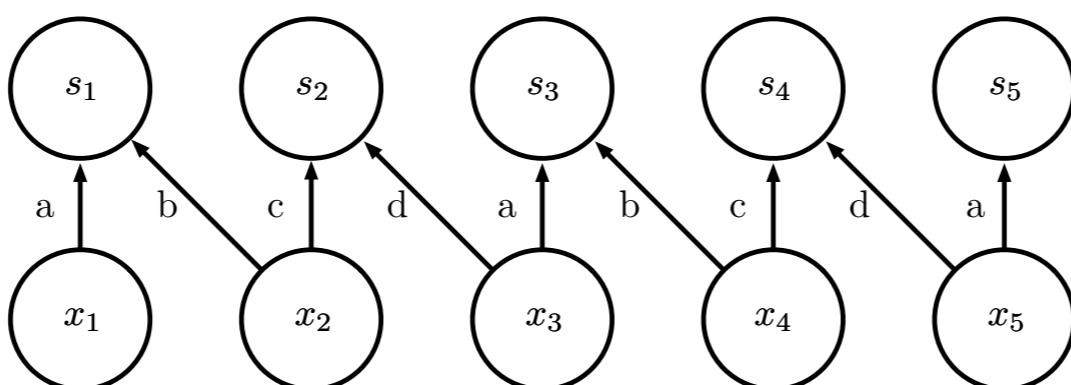


Figure 9.15

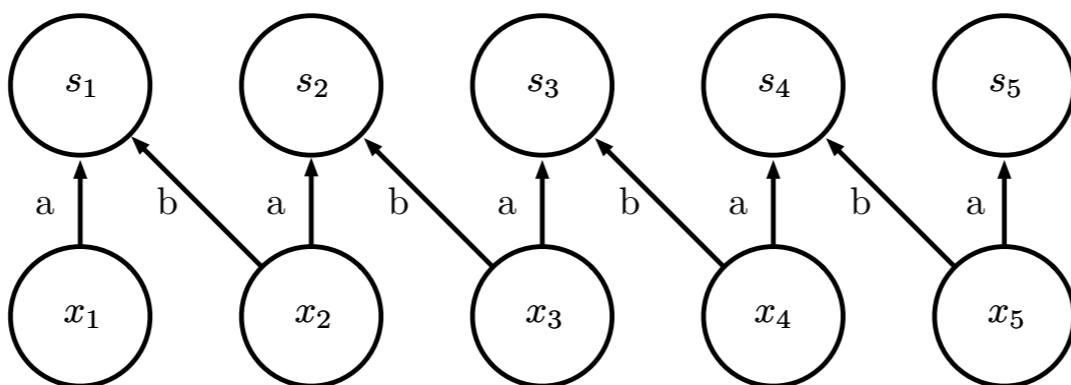
# Tiled convolution



Local connection  
(no sharing)



Tiled convolution  
(cycle between  
groups of shared  
parameters)



Convolution  
(one group shared  
everywhere)

Figure 9.16

# Recurrent Pixel Labeling

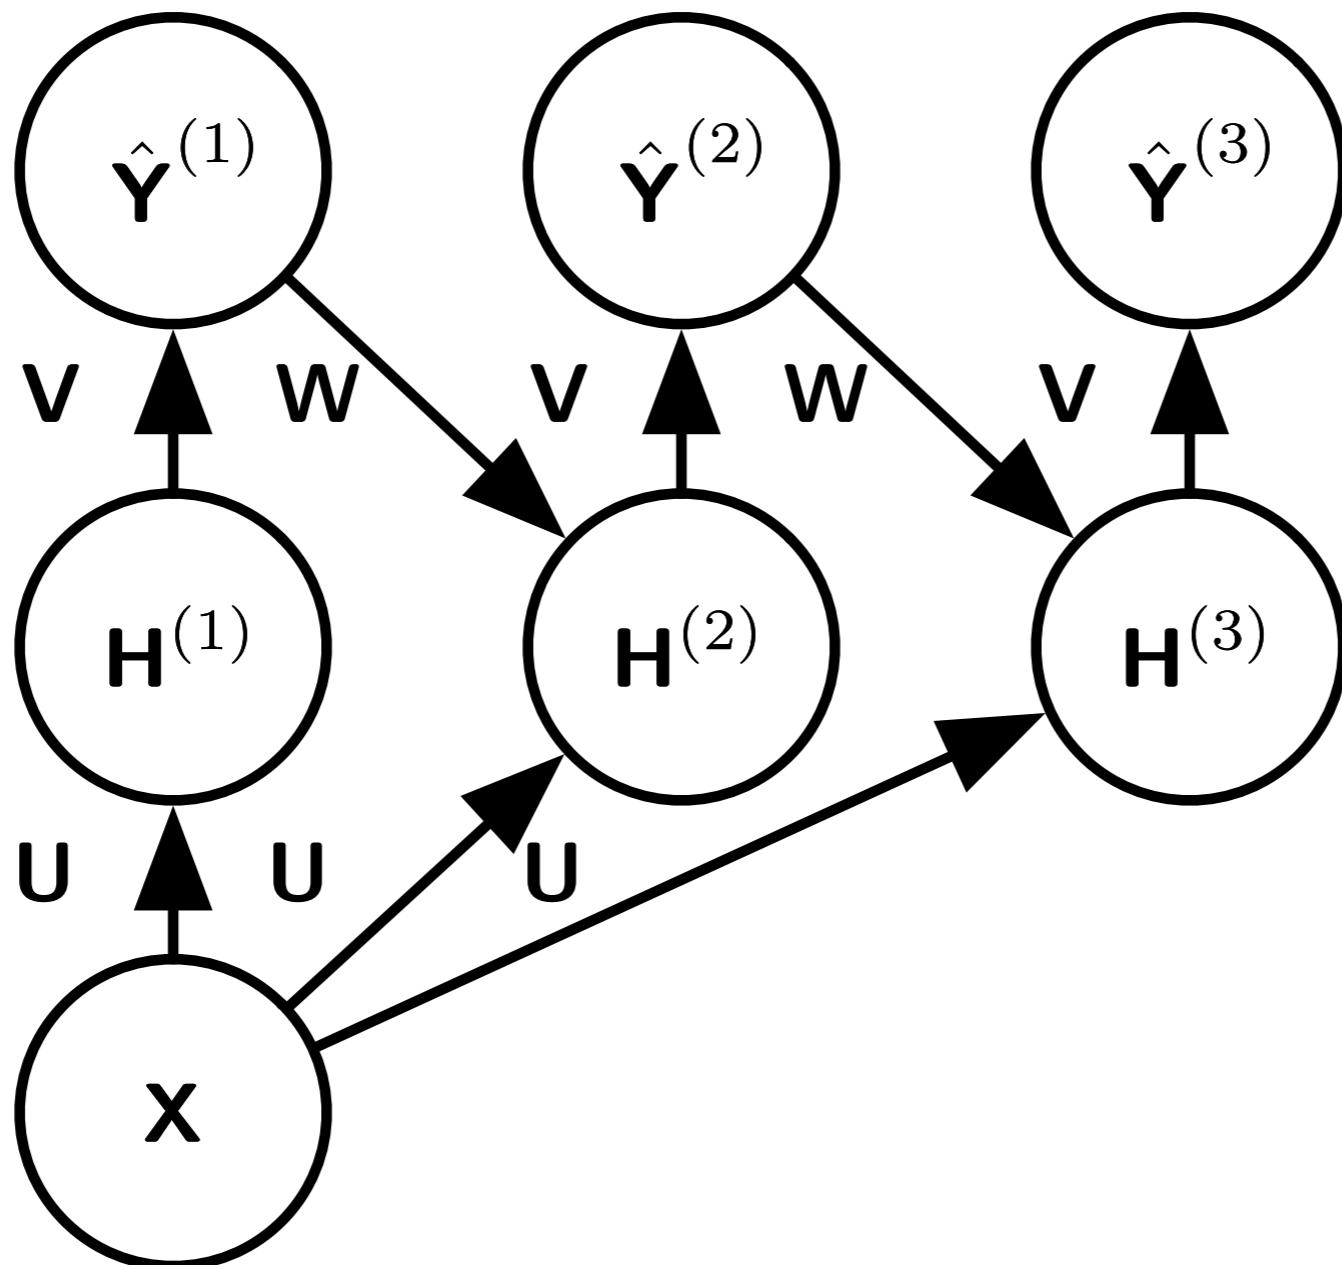


Figure 9.17

# Gabor Functions

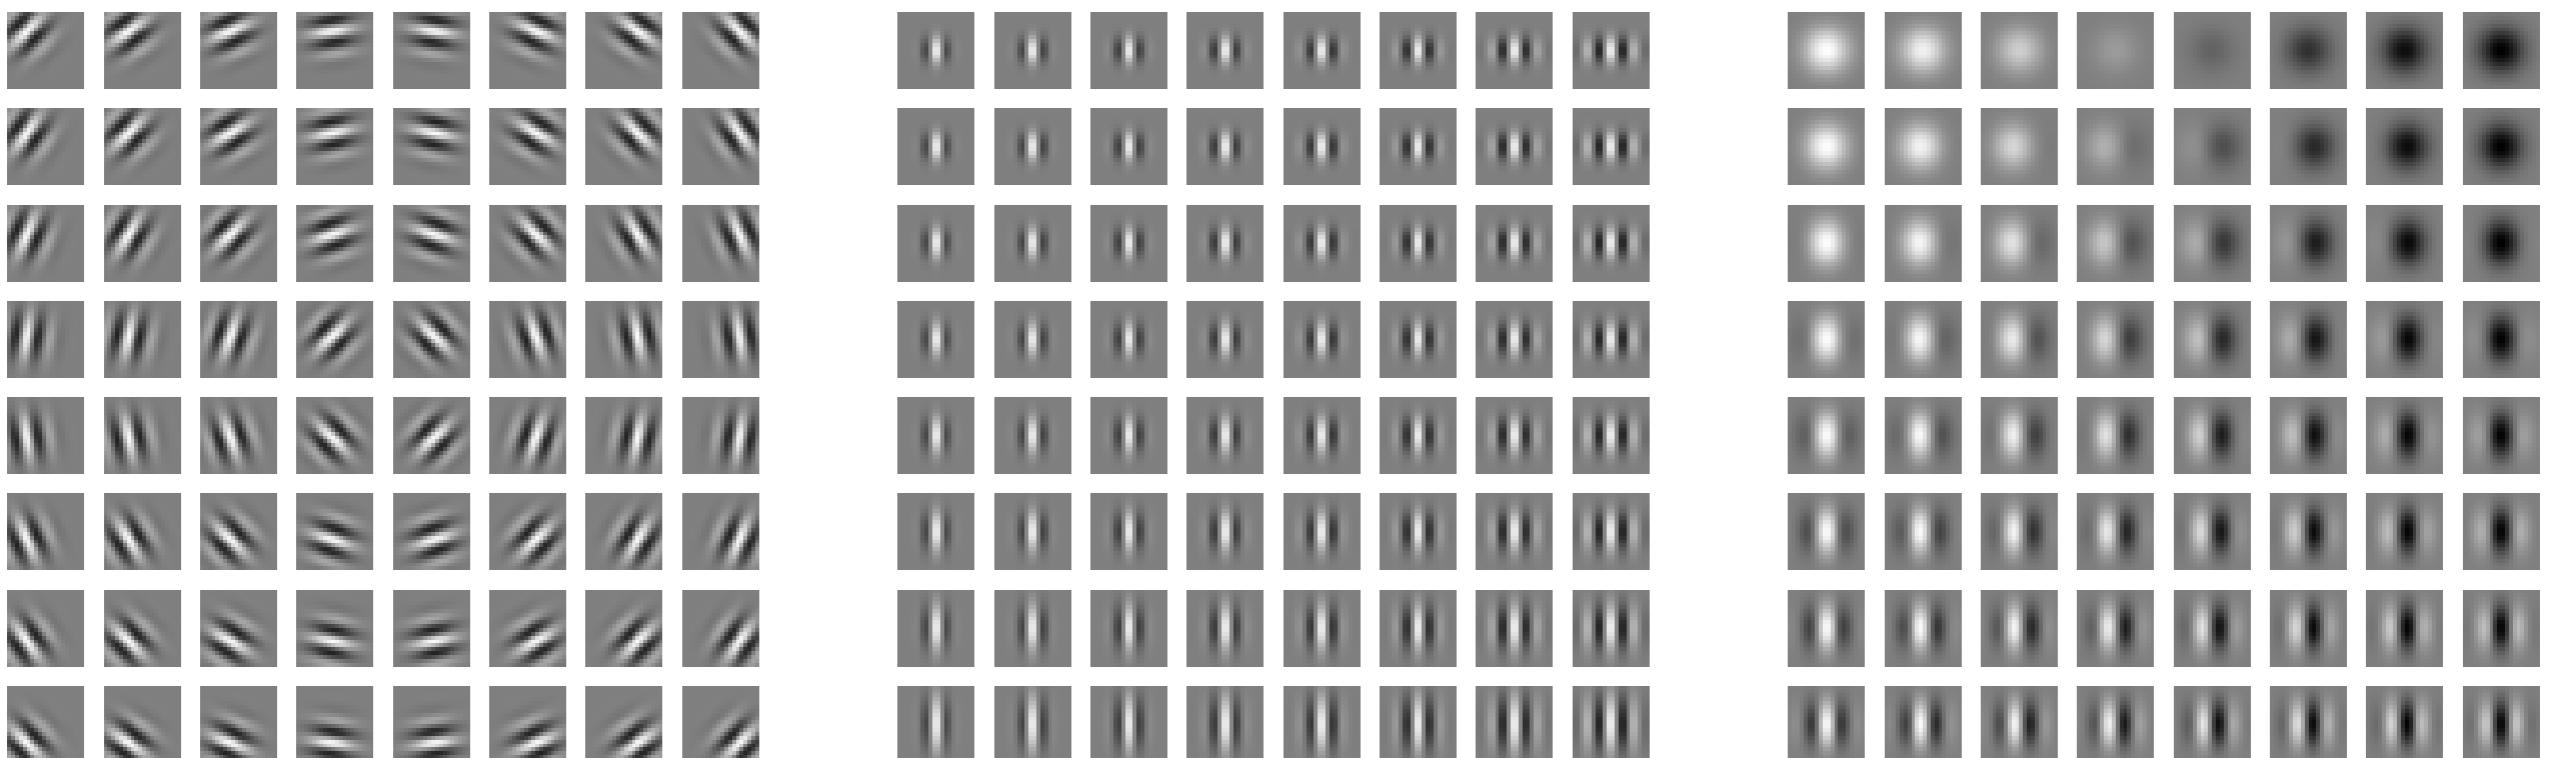


Figure 9.18

# Gabor-like Learned Kernels

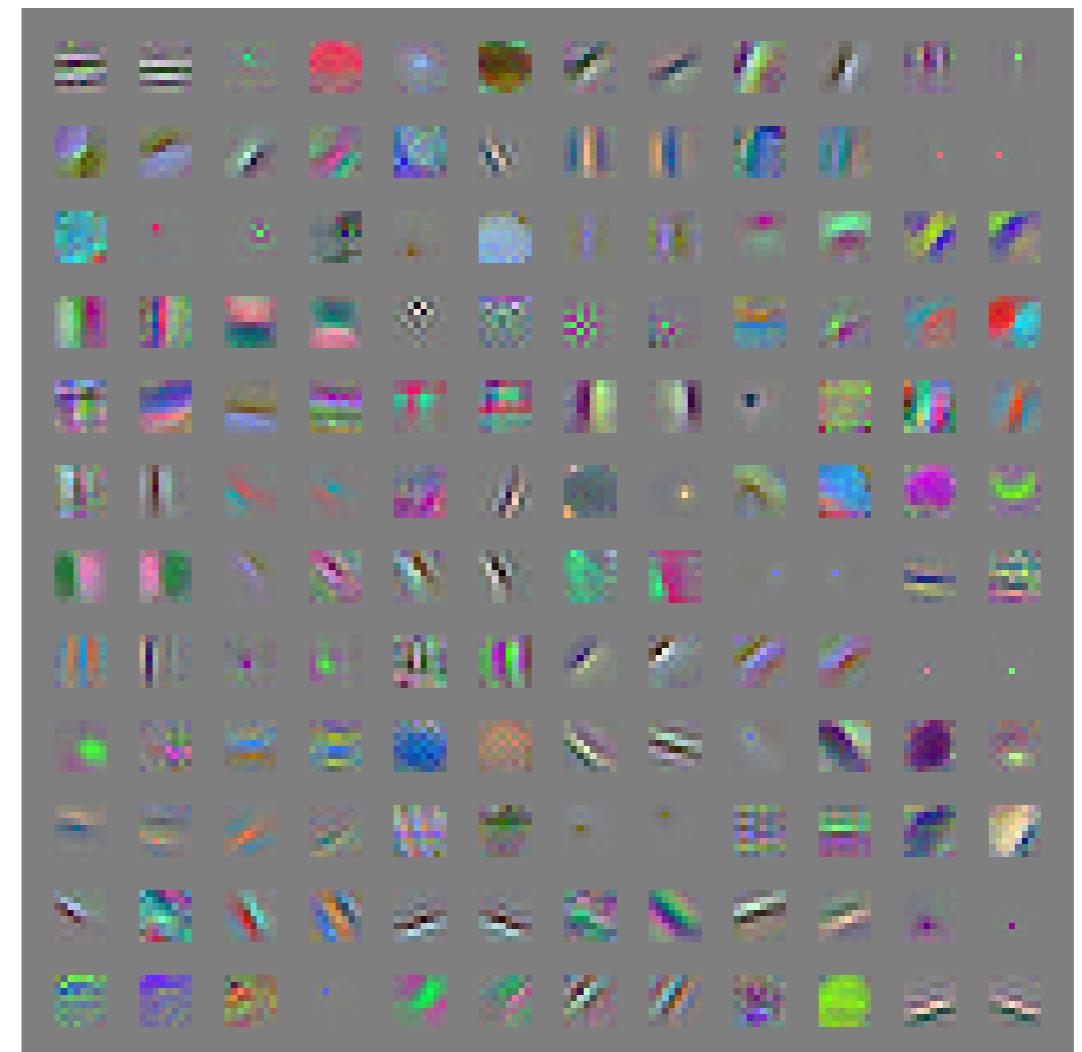
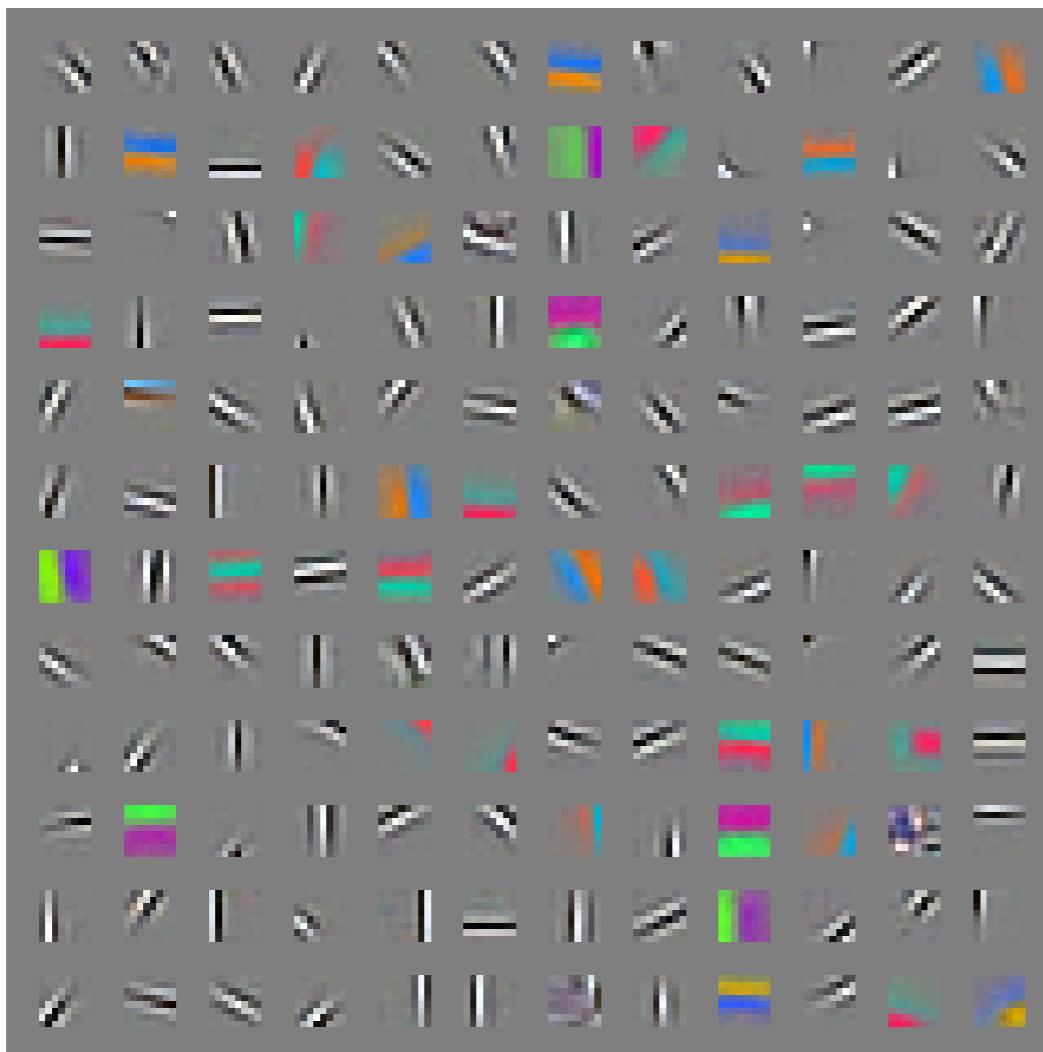


Figure 9.19

# Major Architectures

- Spatial Transducer Net: input size scales with output size, all layers are convolutional
- All Convolutional Net: no pooling layers, just use strided convolution to shrink representation size
- Inception: complicated architecture designed to achieve high accuracy with low computational cost
- ResNet: blocks of layers with same spatial size, with each layer's output added to the same buffer that is repeatedly updated. Very many updates = very deep net, but without vanishing gradient.