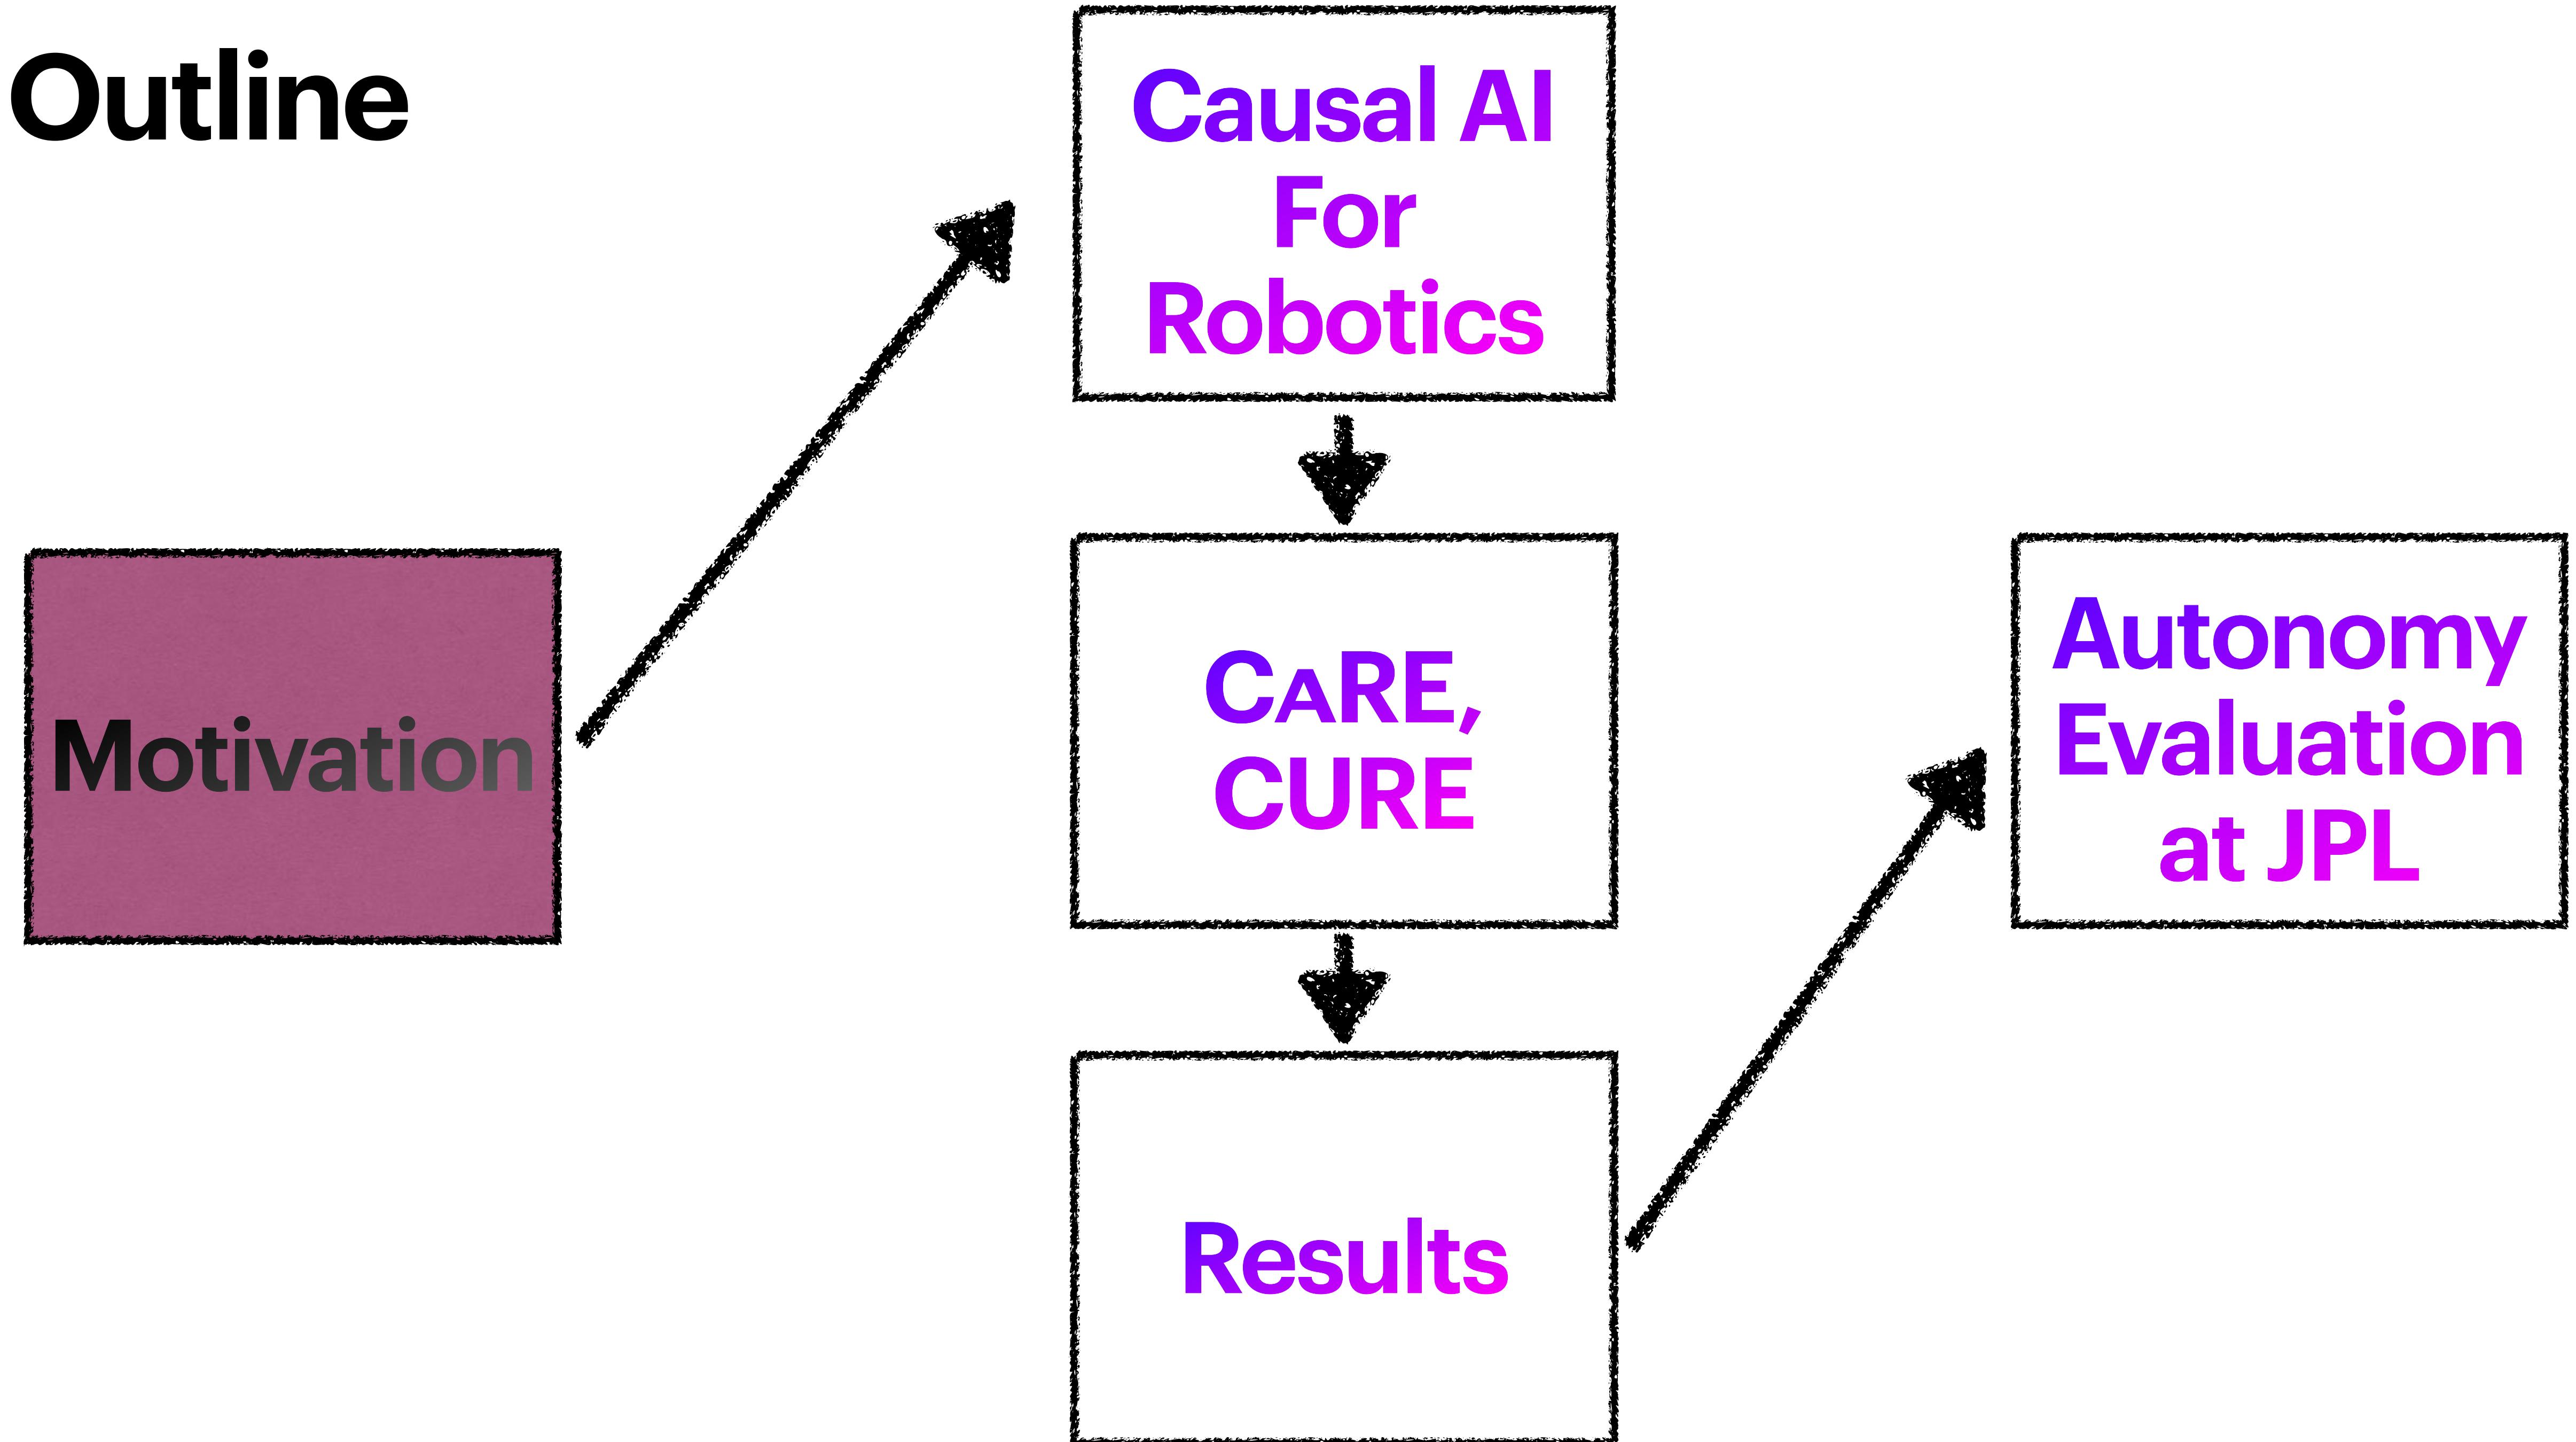


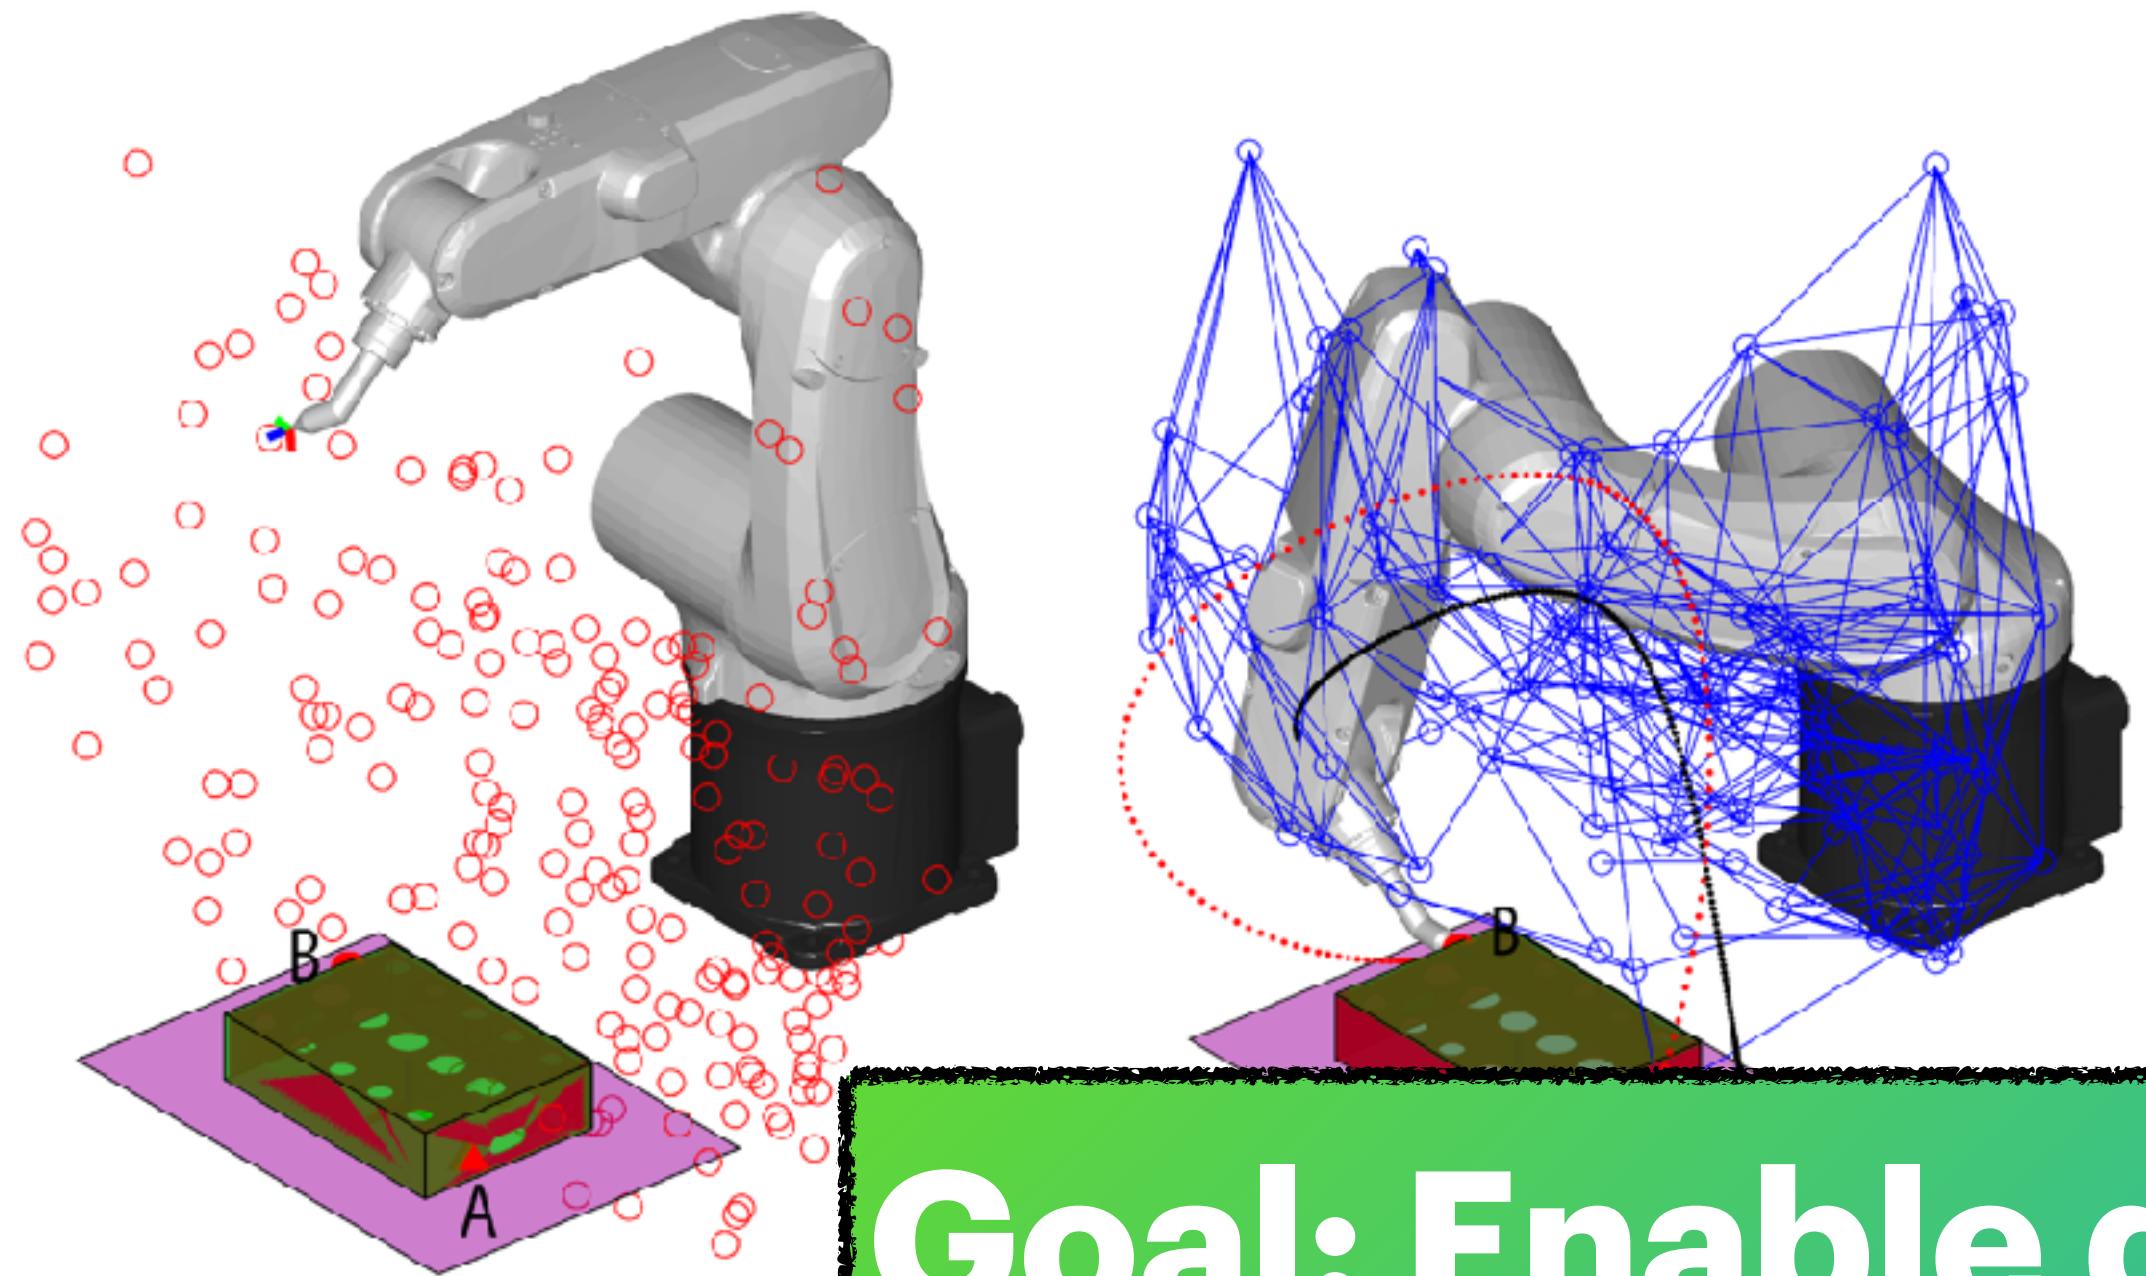
Performance Modeling, Debugging, and Optimization of Highly Configurable Robotic Systems with Causal AI

A Path towards Building Dependable Robotic Systems

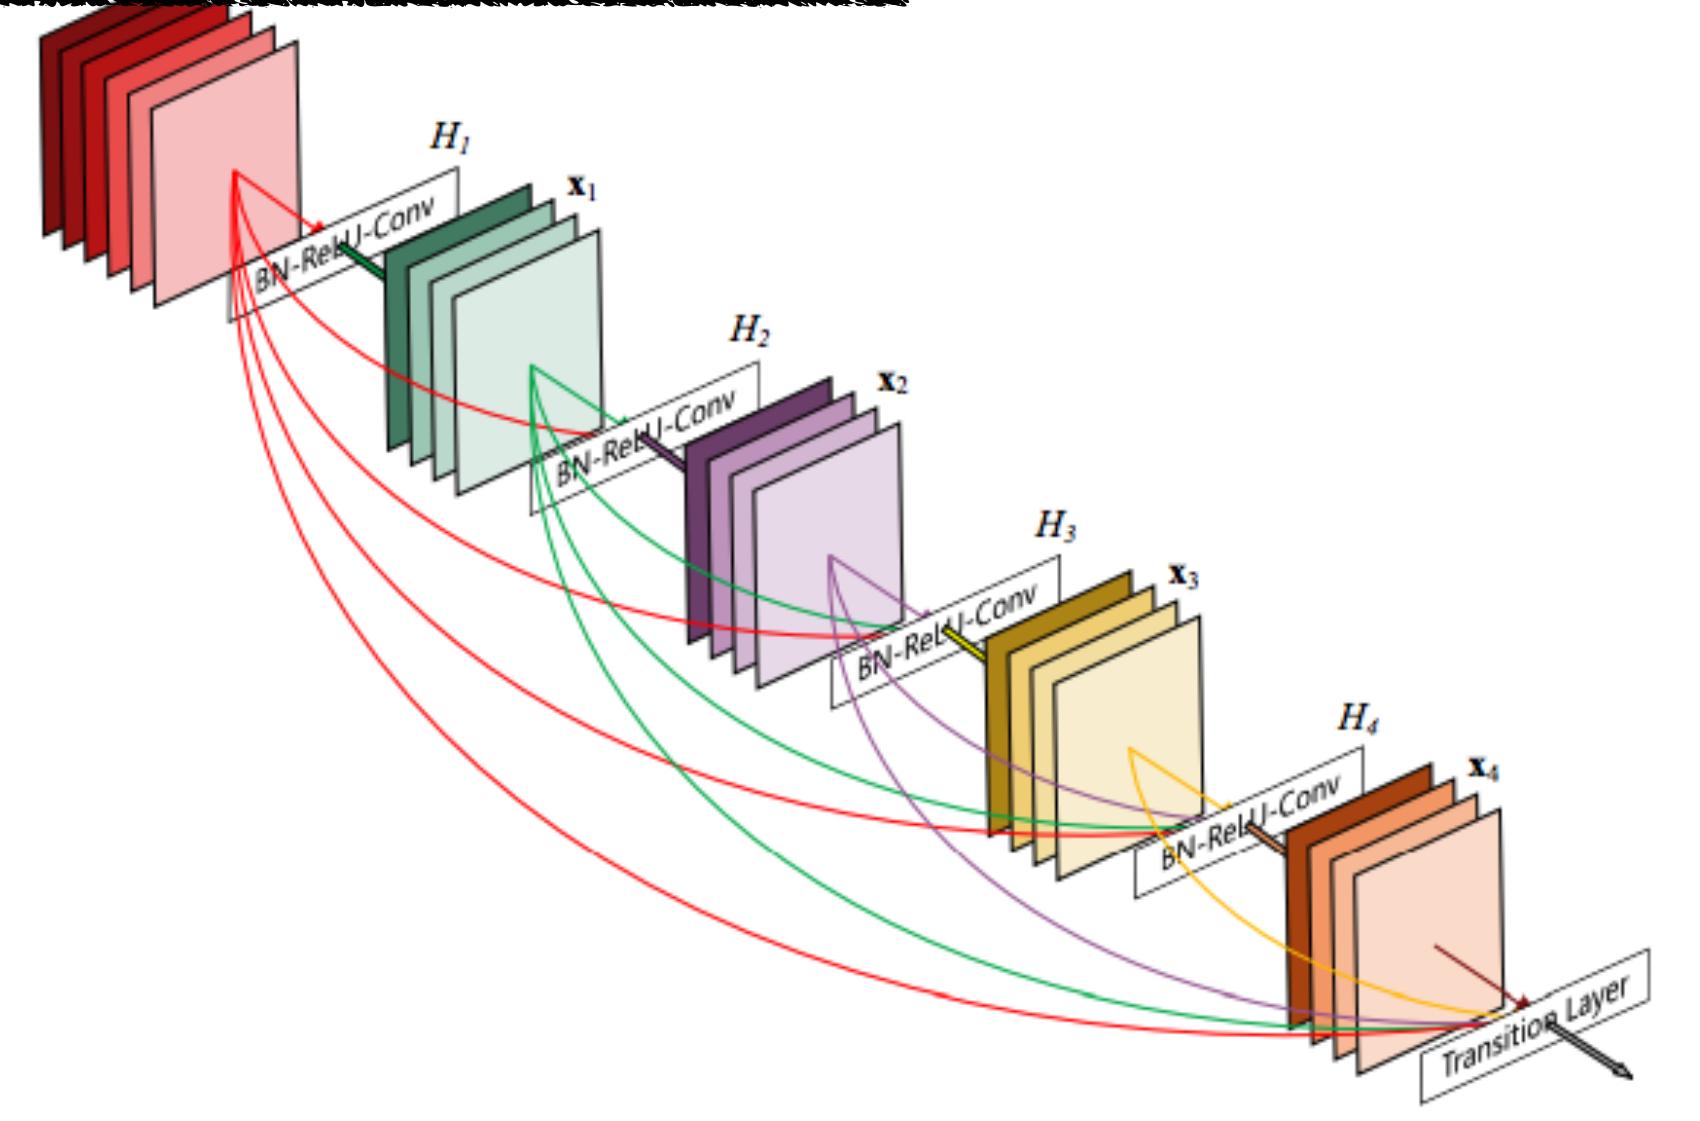
Md Abir Hossen
University of South Carolina

Outline





Goal: Enable developers/users
to find the right quality tradeoff



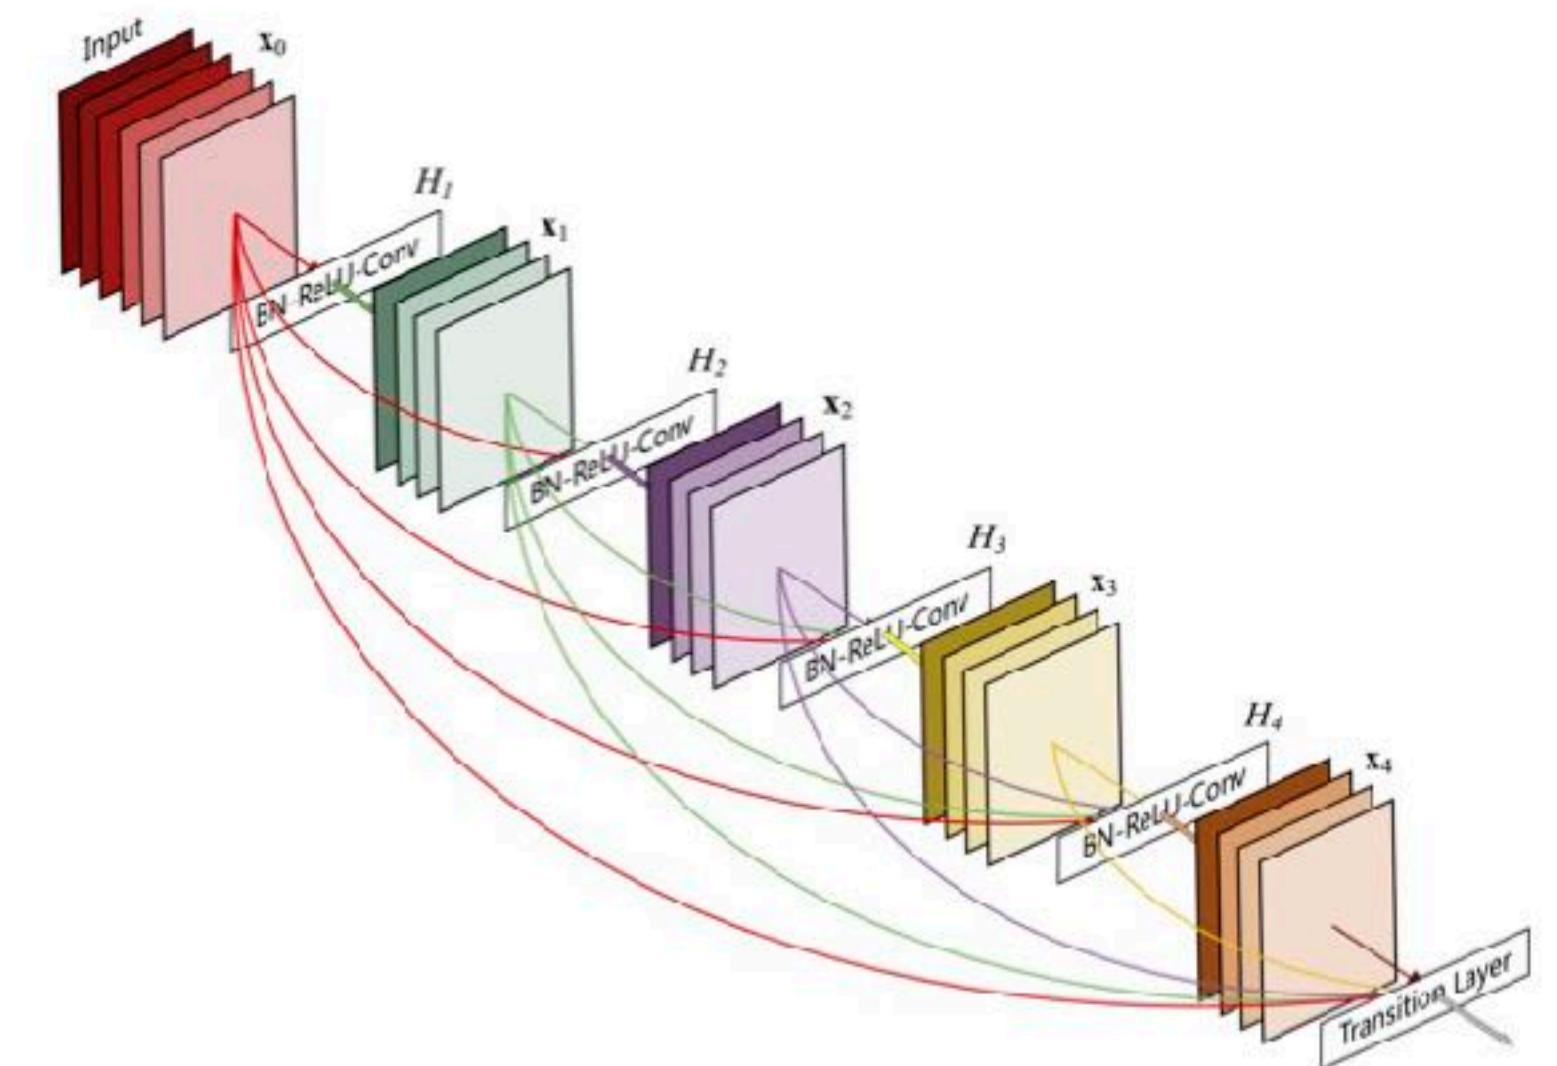
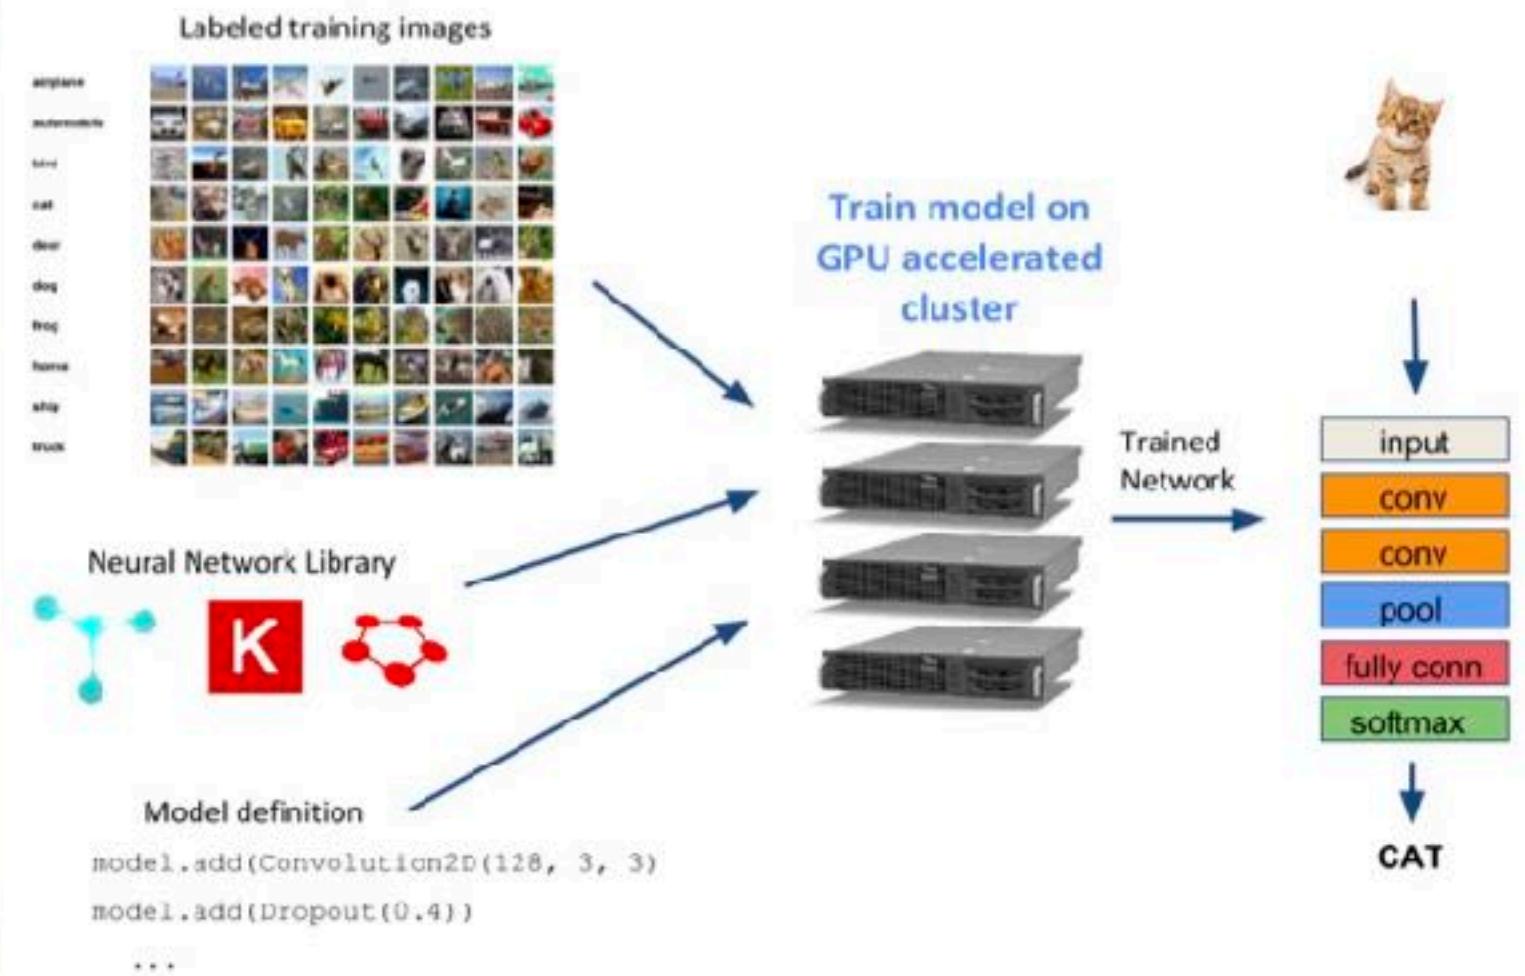
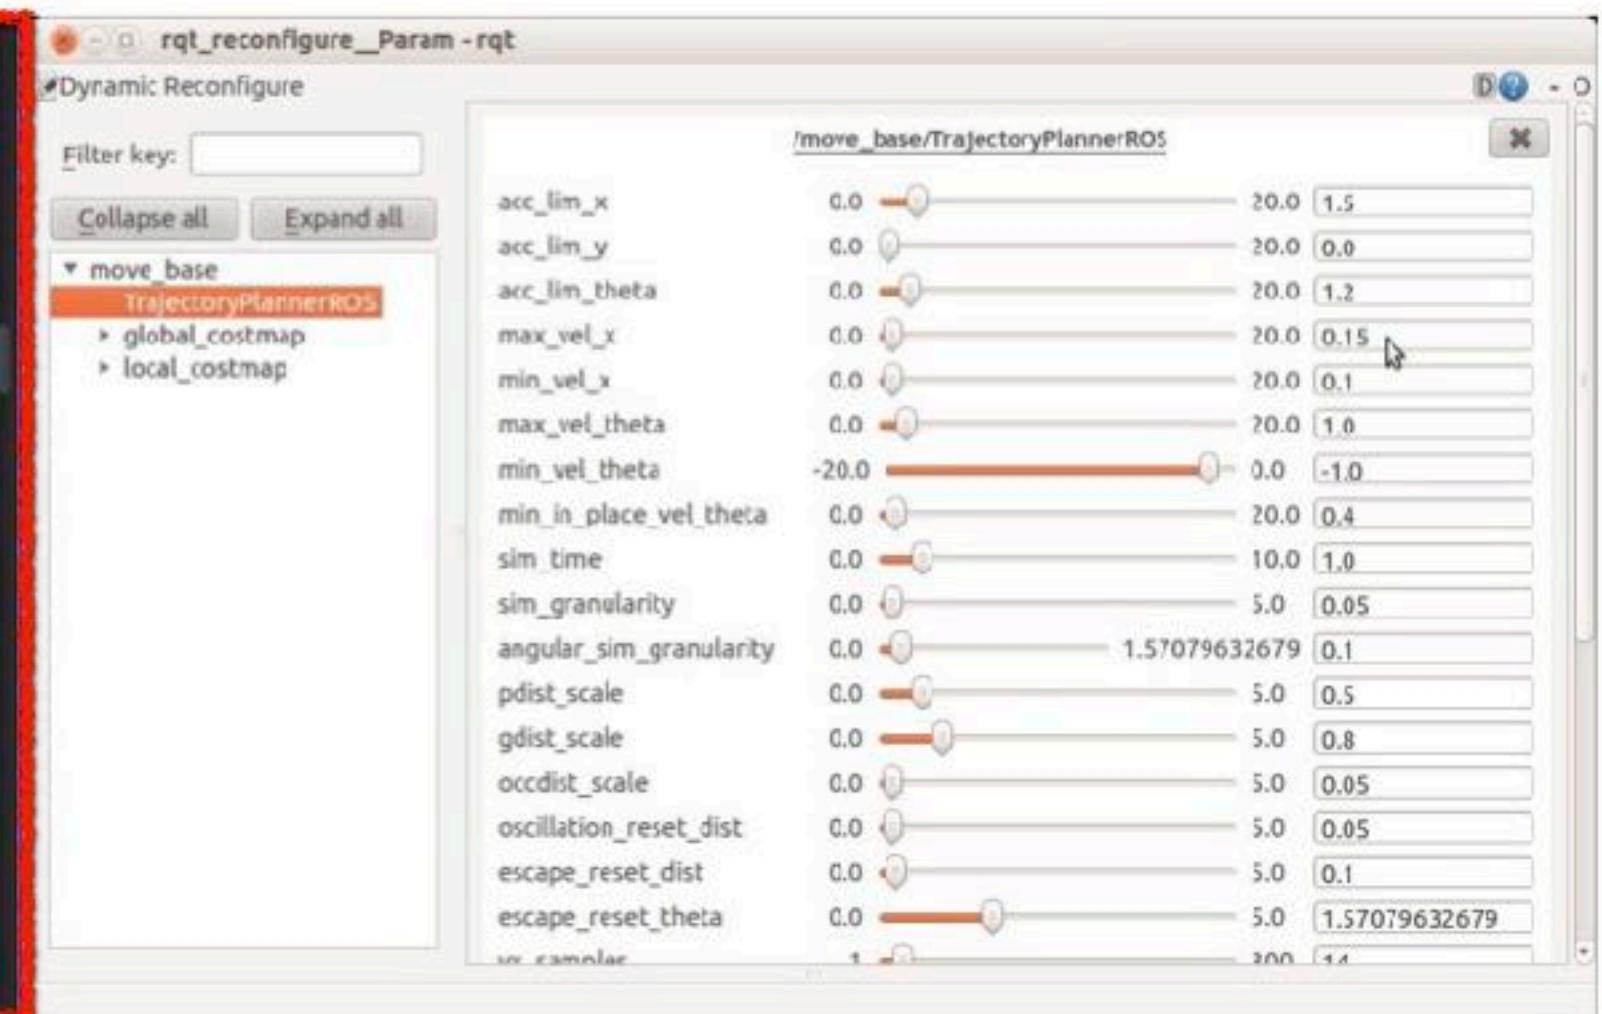
built

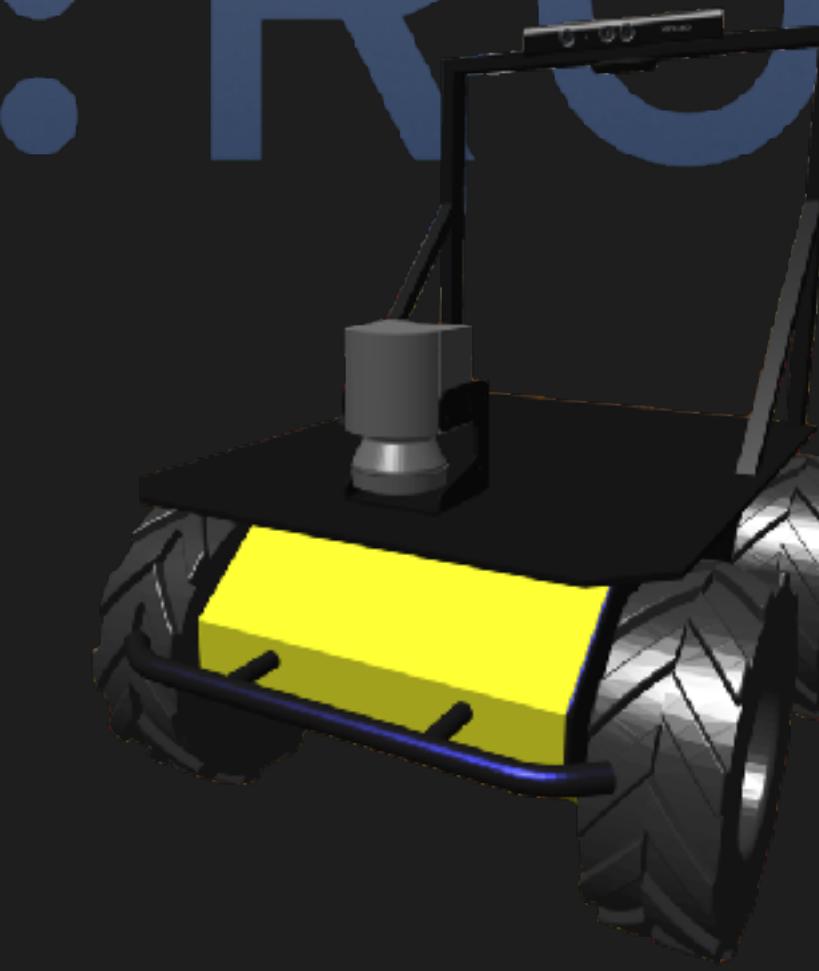
Today's most popular systems are configurable

```

102
103 drpc.port: 3772
104 drpc.worker.threads: 64
105 drpc.max_buffer_size: 1048576
106 drpc.queue.size: 128
107 drpc.invocations.port: 3773
108 drpc.invocations.threads: 64
109 drpc.request.timeout.secs: 600
110 drpc.childopts: "-Xmx768m"
111 drpc.http.port: 3774
112 drpc.https.port: -1
113 drpc.https.keystore.password: ""
114 drpc.https.keystore.type: "JKS"
115 drpc.http.creds.plugin: org.apache.storm.security.auth.DefaultHttpCredentialsPlugin
116 drpc.authorizer.acl.filename: "drpc-auth-acl.yaml"
117 drpc.authorizer.acl.strict: false
118
119 transactional.zookeeper.root: "/transactional"
120 transactional.zookeeper.servers: null
121 transactional.zookeeper.port: null
122
123 ## blobstore configs
124 supervisor.blobstore.class: "org.apache.storm.blobstore.NimbusBlobStore"
125 supervisor.blobstore.download.thread.count: 5
126 supervisor.blobstore.download.max_retries: 3
127 supervisor.localizer.cache.target.size.mb: 10240
128 supervisor.localizer.cleanup.interval.ms: 600000
...

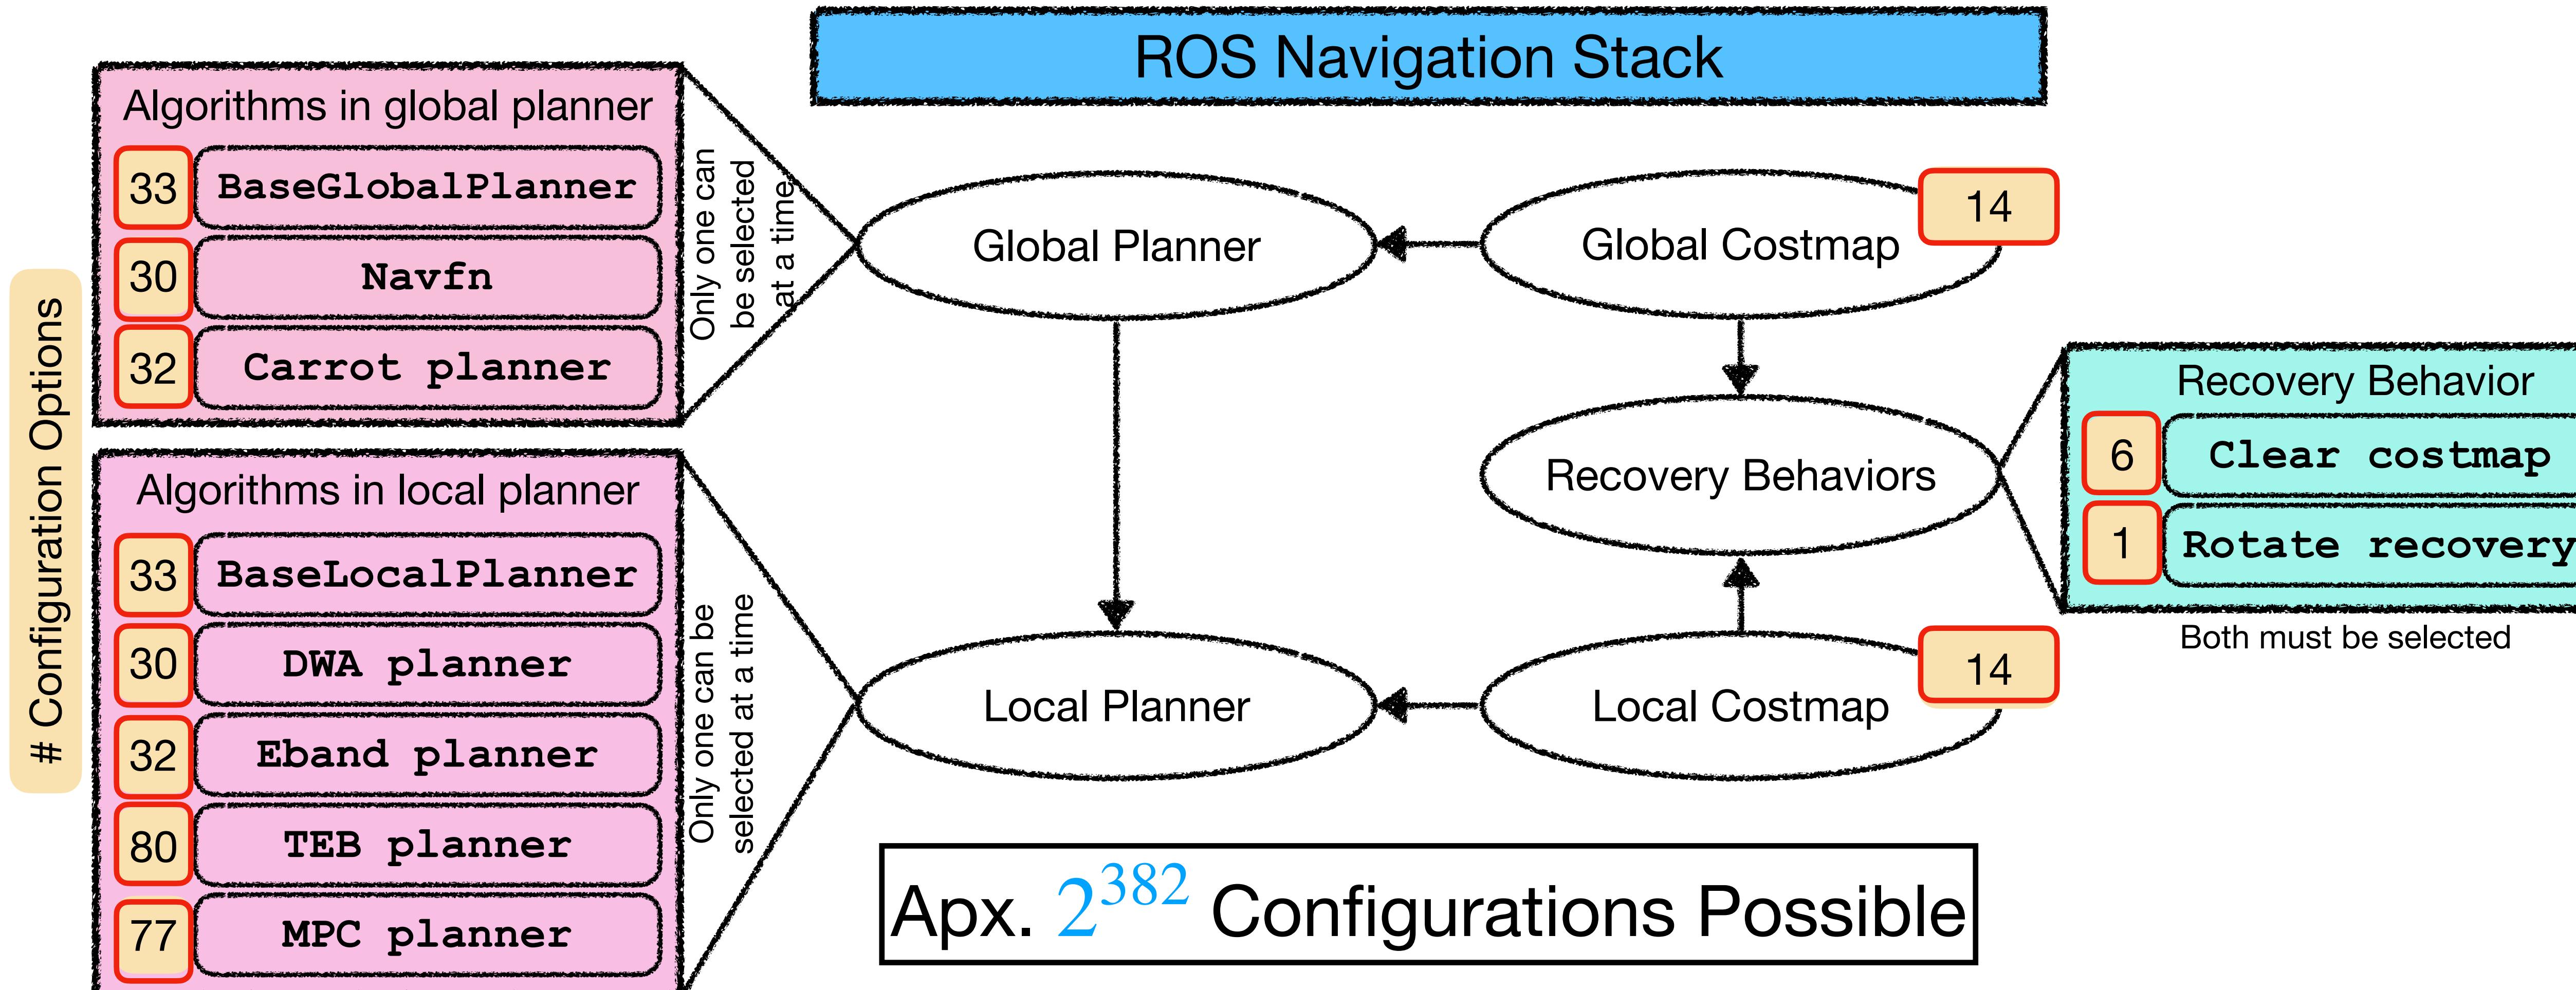
```



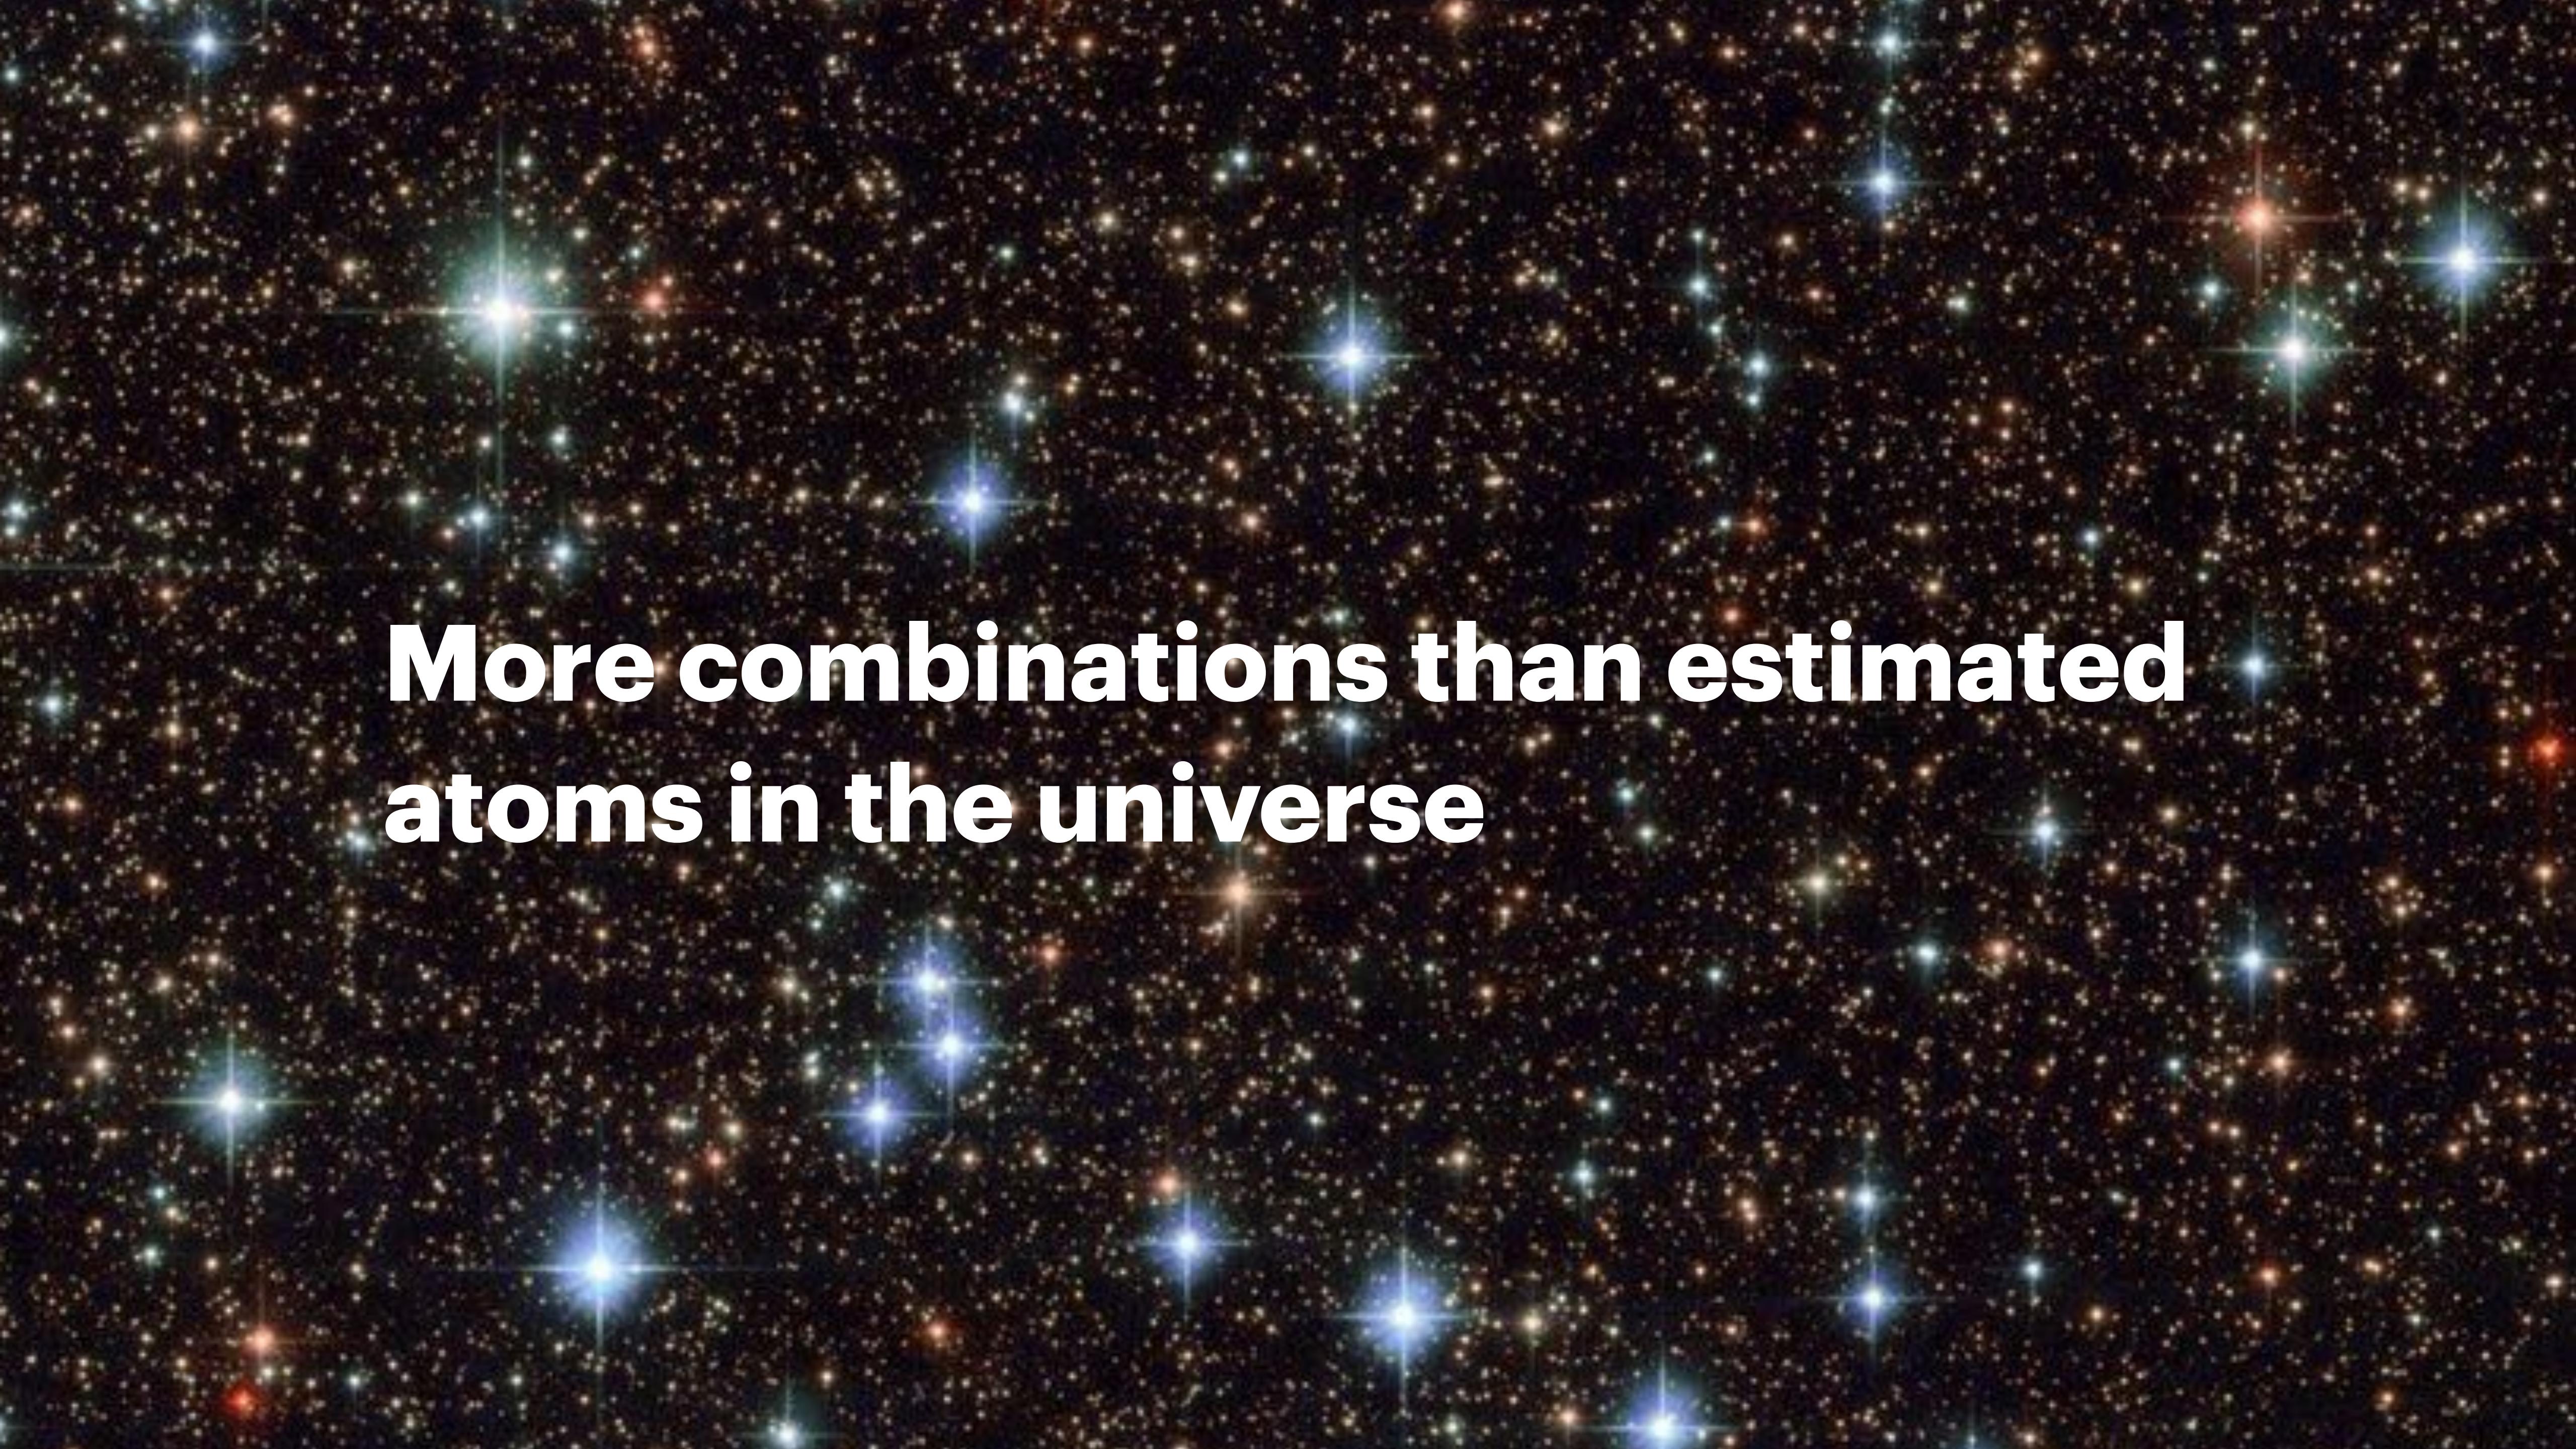


```
23 # Goal Tolerance Parameters
24 yaw_goal_tolerance: 0.1
25 xy_goal_tolerance: 0.2
26 latch_xy_goal_tolerance: false
27
28 # Forward Simulation Parameters
29 sim_time: 2.0
30 sim_granularity: 0.02
31 angular_sim_granularity: 0.02
32 vx_samples: 6
33 vtheta_samples: 20
34 controller_frequency: 20.0
35
36 # Trajectory scoring parameters
37 meter_scoring: true # Whether the gdist_scale and pdist_scale parameters should assume that goal_distance and path_distance are in meters. default true
38 occdist_scale: 0.1 #The weighting for how much the controller should attempt to avoid obstacles. default 0.01
39 pdist_scale: 0.75 # The weighting for how much the controller should stay close to the path it was given . default 0.6
40 gdist_scale: 1.0 # The weighting for how much the controller should attempt to reach its local goal, also controls speed
41
42 heading_lookahead: 0.325 #How far to look ahead in meters when scoring different in-place-rotation trajectories
43 heading_scoring: false #Whether to score based on the robot's heading to the path or its distance from the path. default false
44 heading_scoring_timestep: 0.8 #How far to look ahead in time in seconds along the simulated trajectory when using heading scoring
45 dwa: true #Whether to use the Dynamic Window Approach (DWA) or whether to use Trajectory Rollout
46 simple_attractor: false
47 publish_cost_grid_pc: true
48
49 # Oscillation Prevention Parameters
50 oscillation_reset_dist: 0.25 #How far the robot must travel in meters before oscillation flags are reset (double, default: 0.25)
51 escape_reset_dist: 0.1
52 escape_reset_theta: 0.1
```

Configuration Space in Robotic Systems



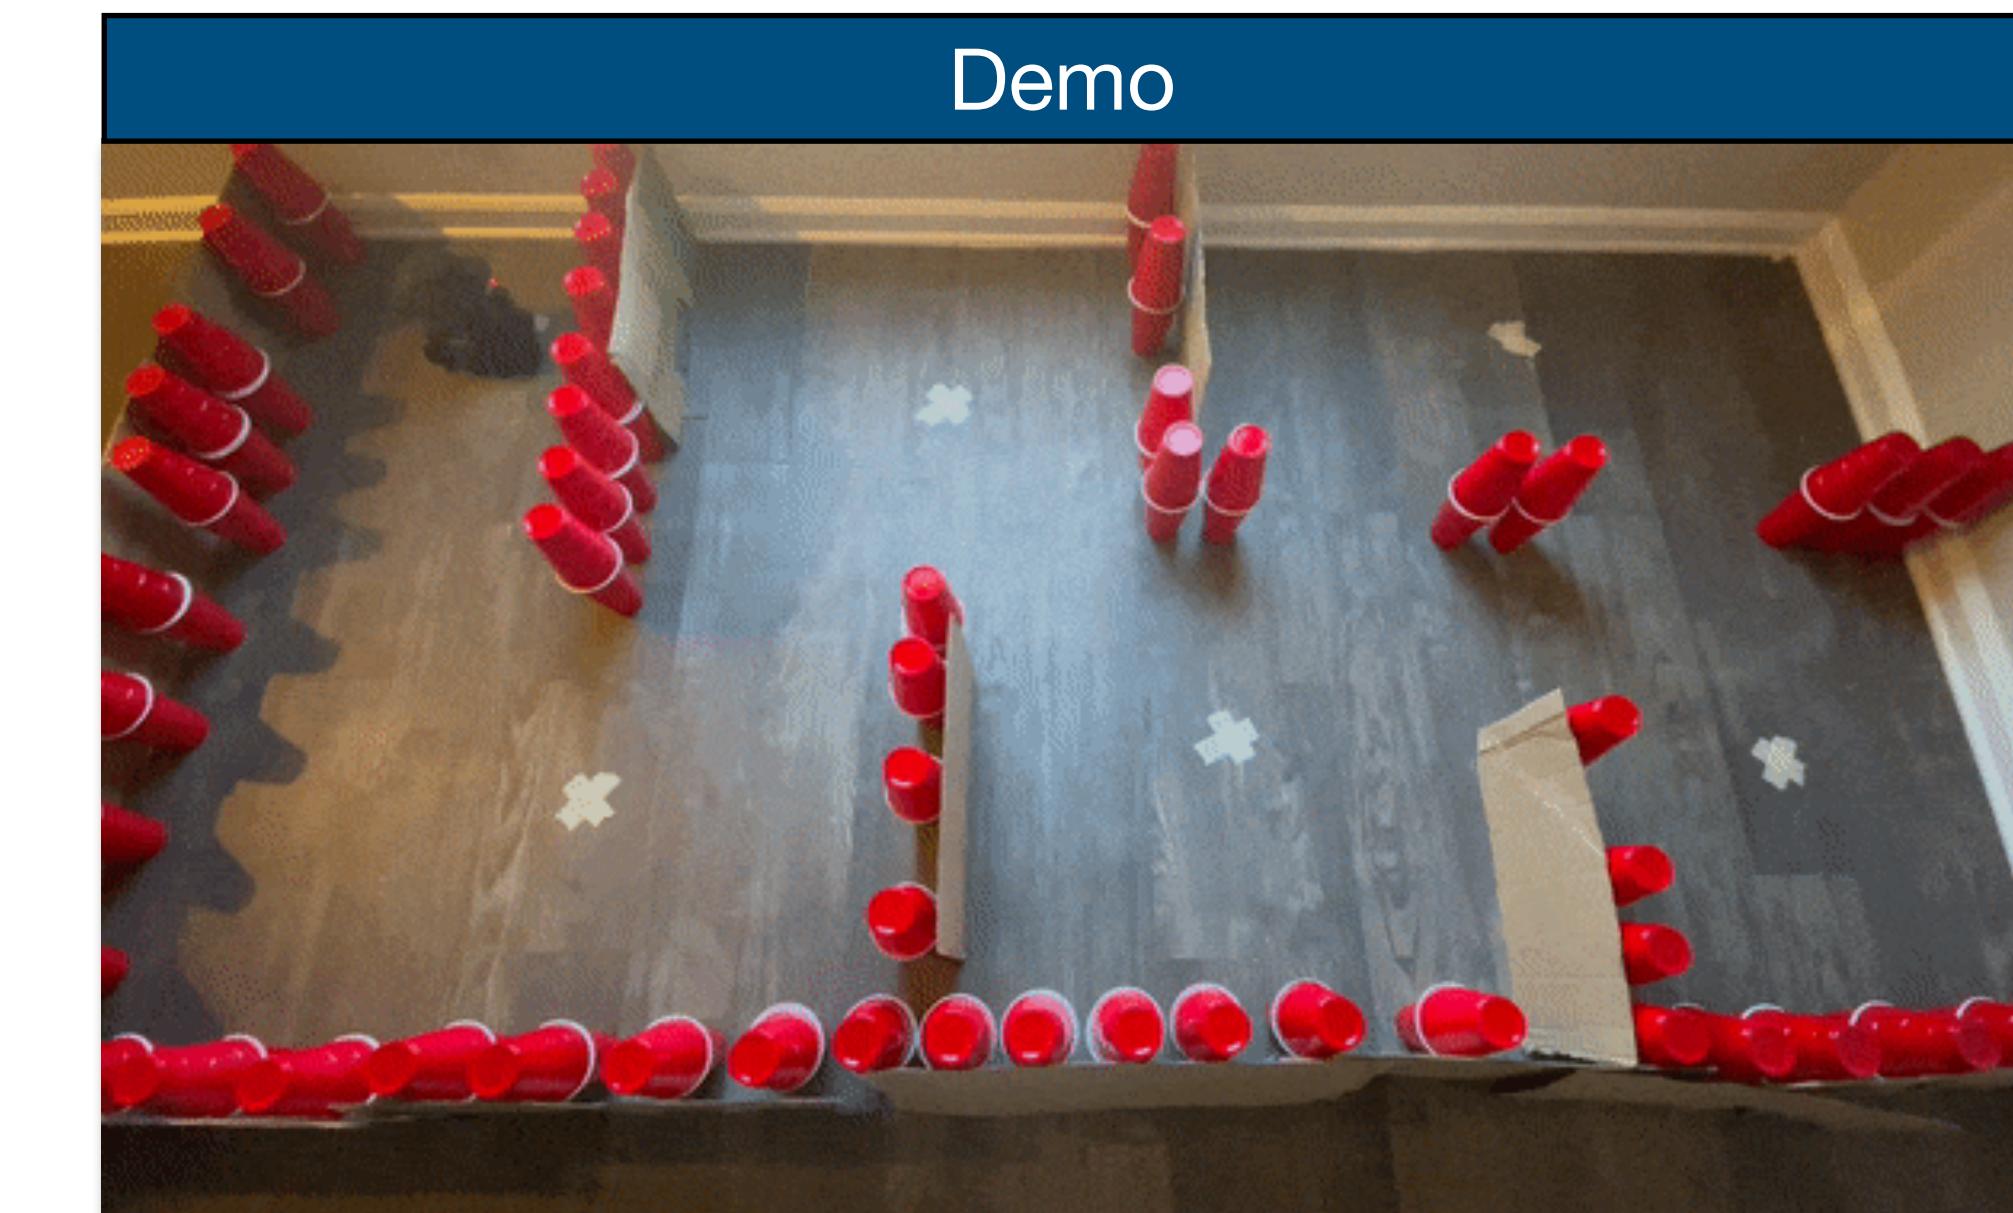
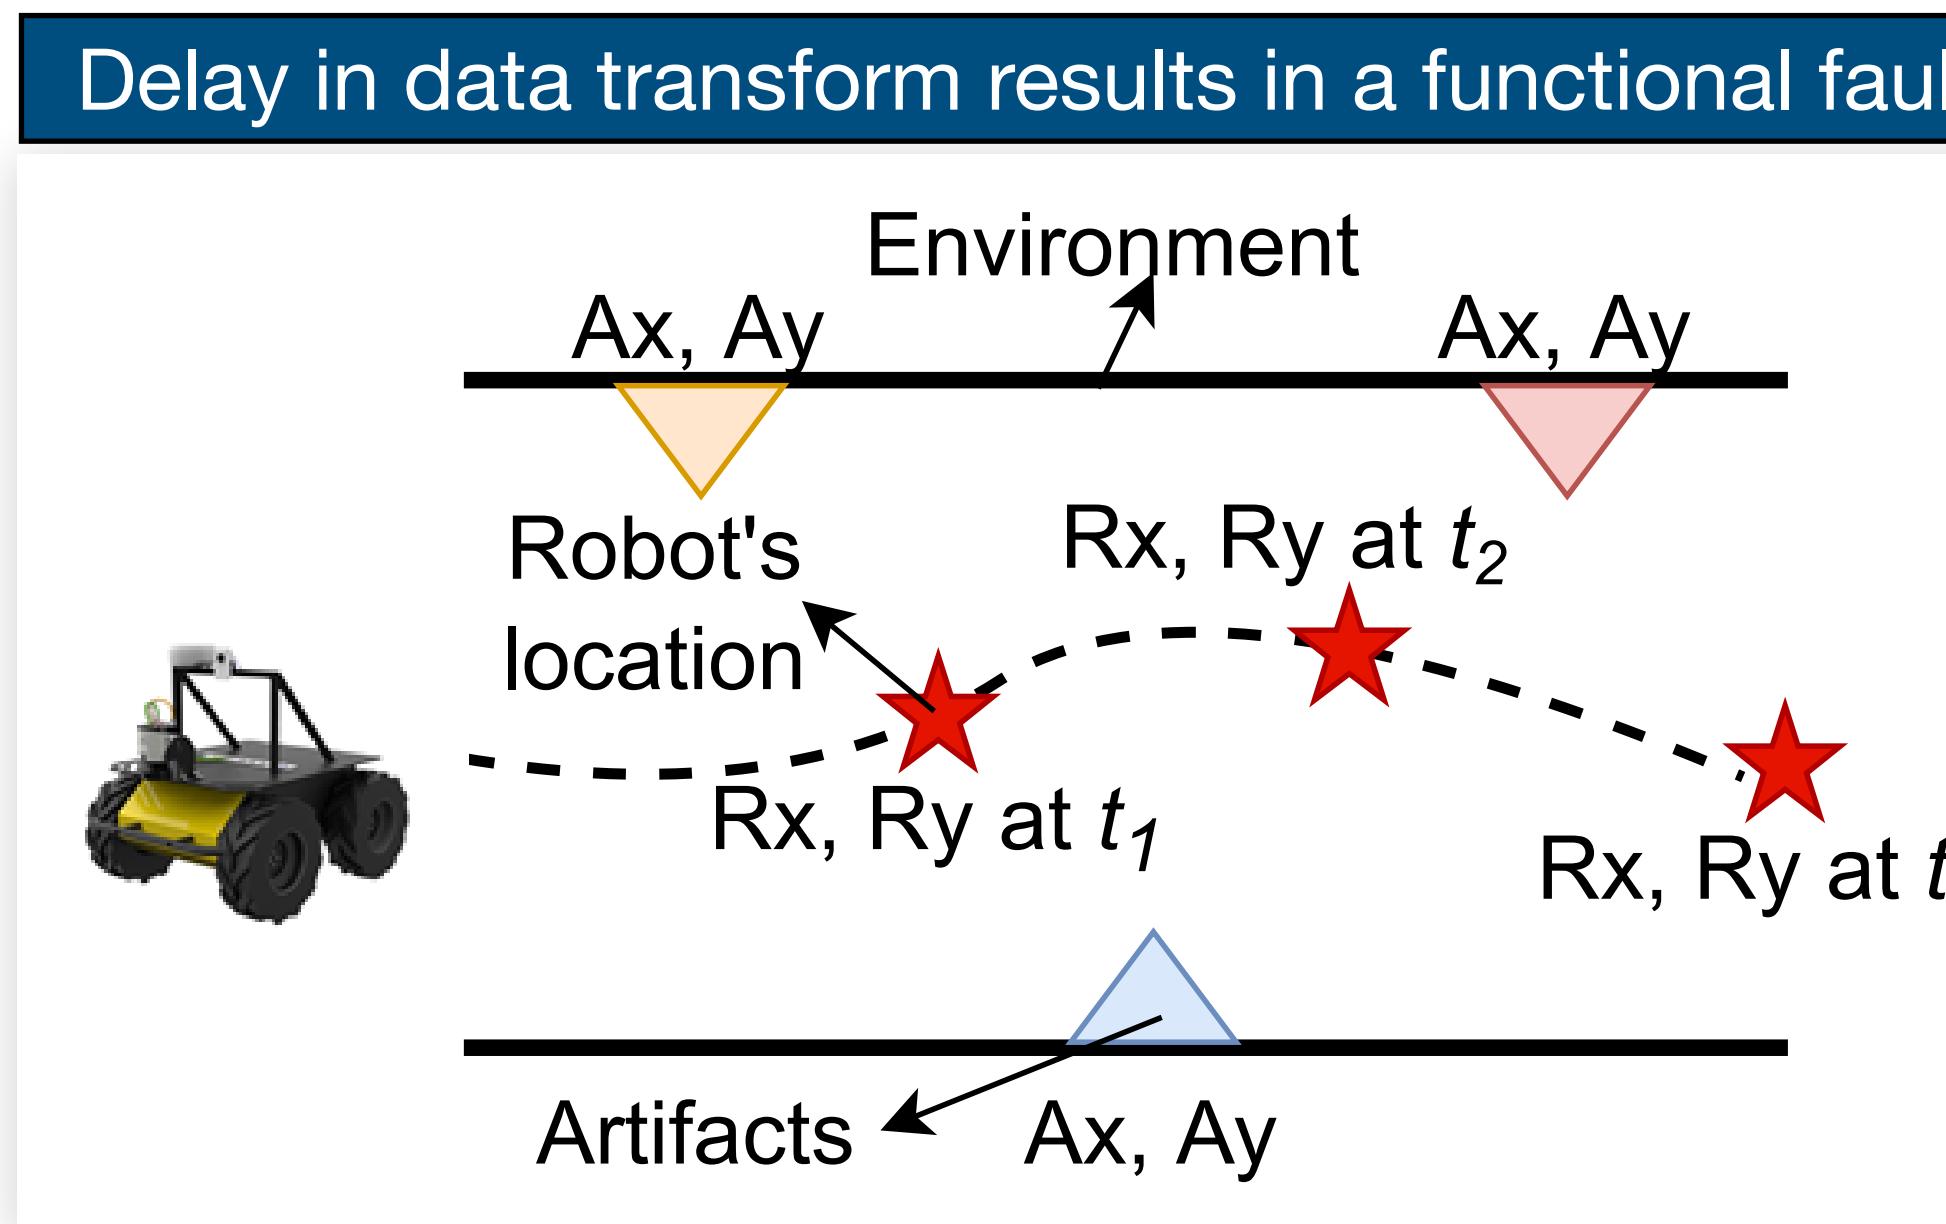
Complex interactions between options (intra or inter components) give rise to a combinatorially large configuration space.



**More combinations than estimated
atoms in the universe**

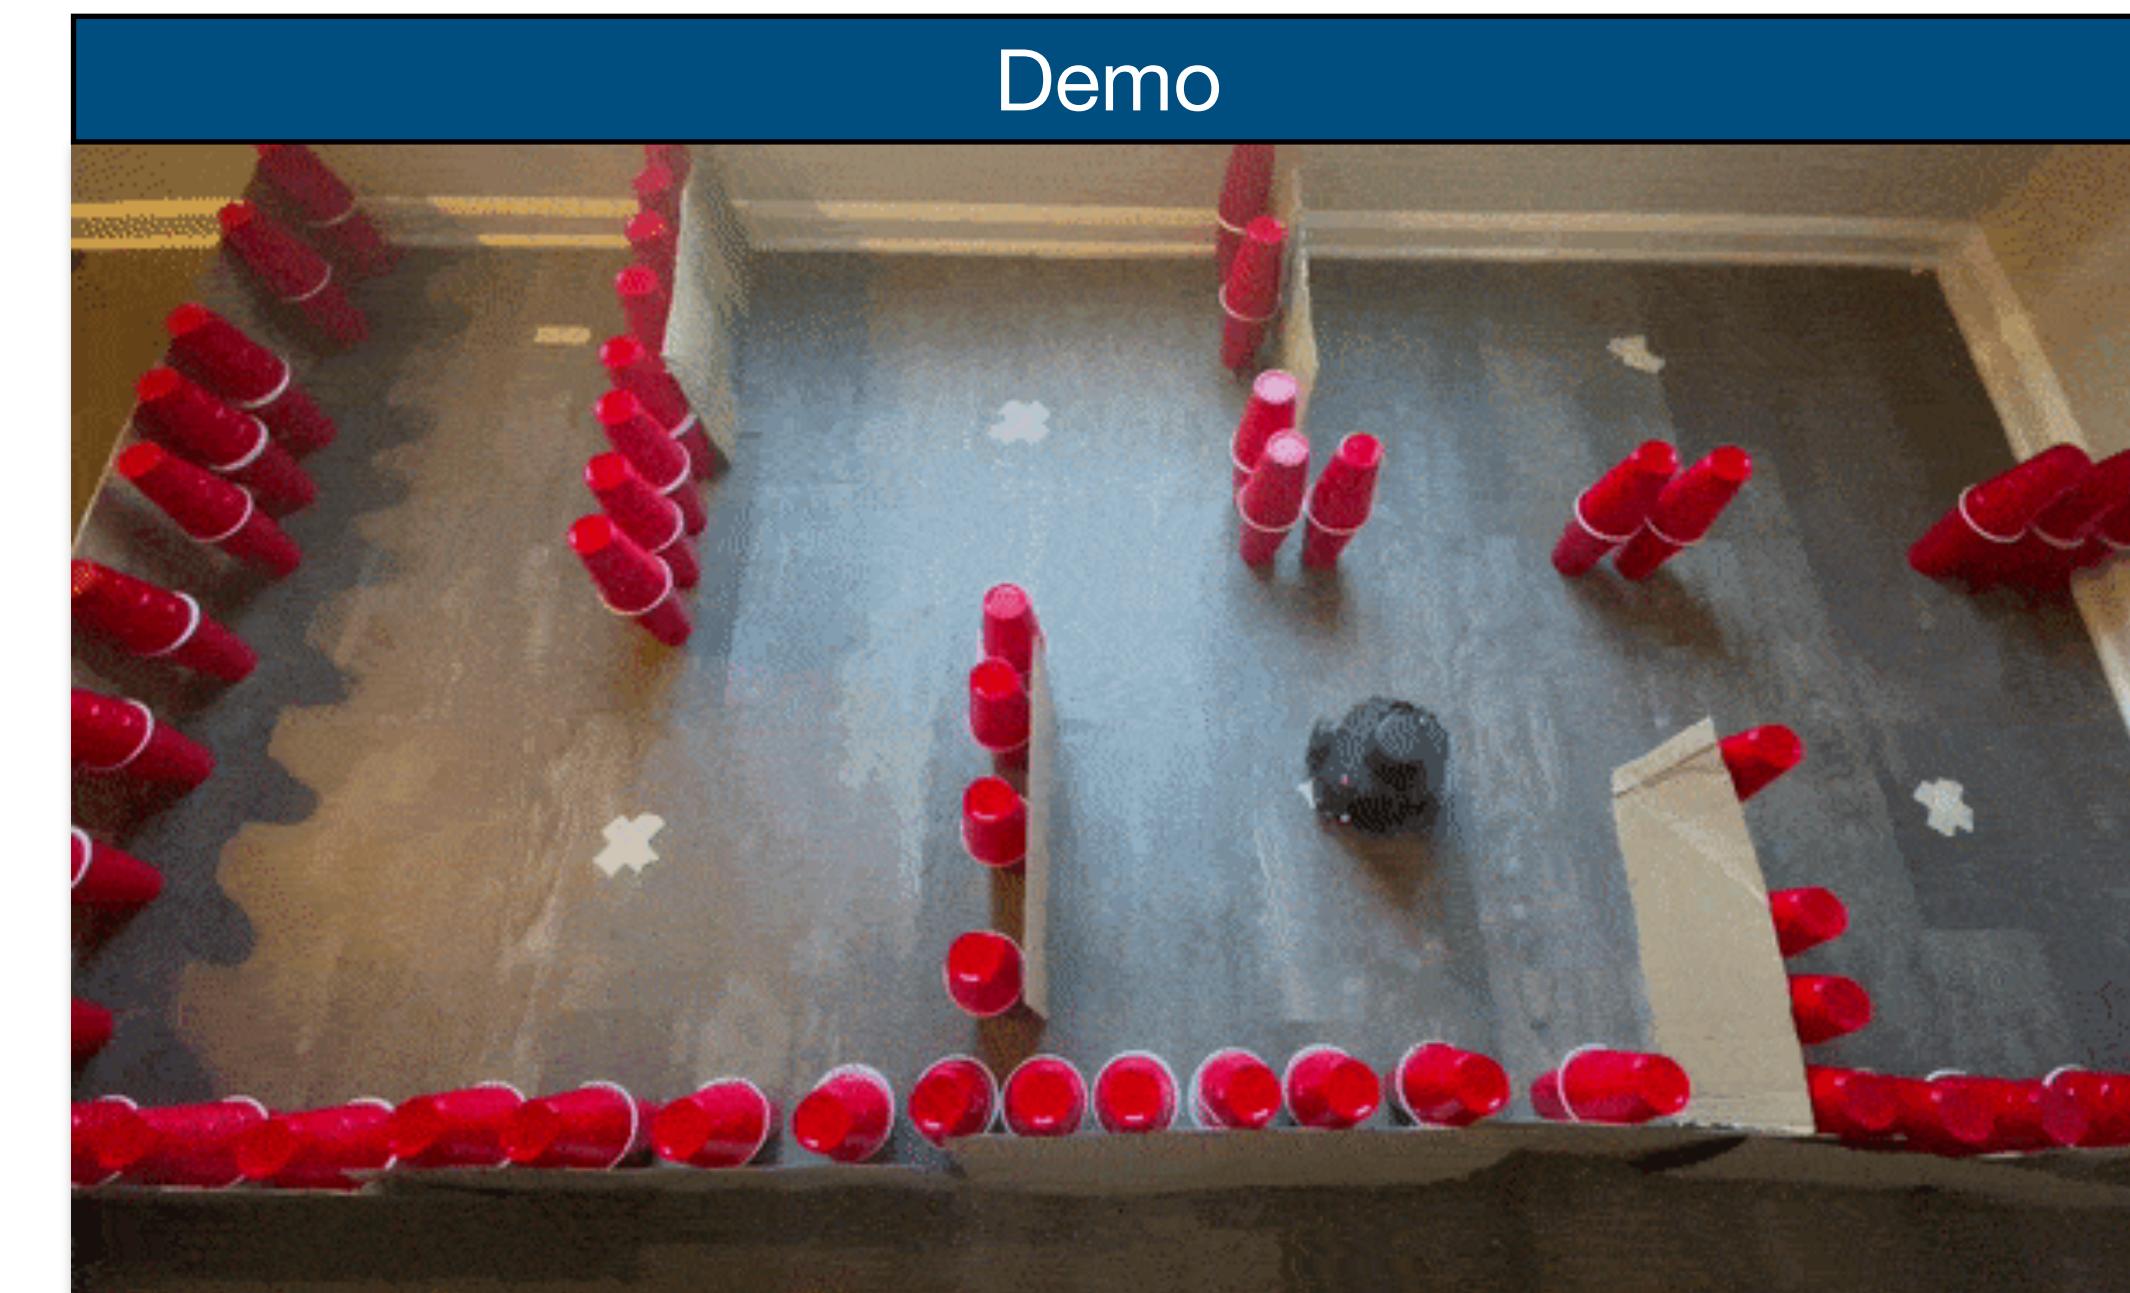
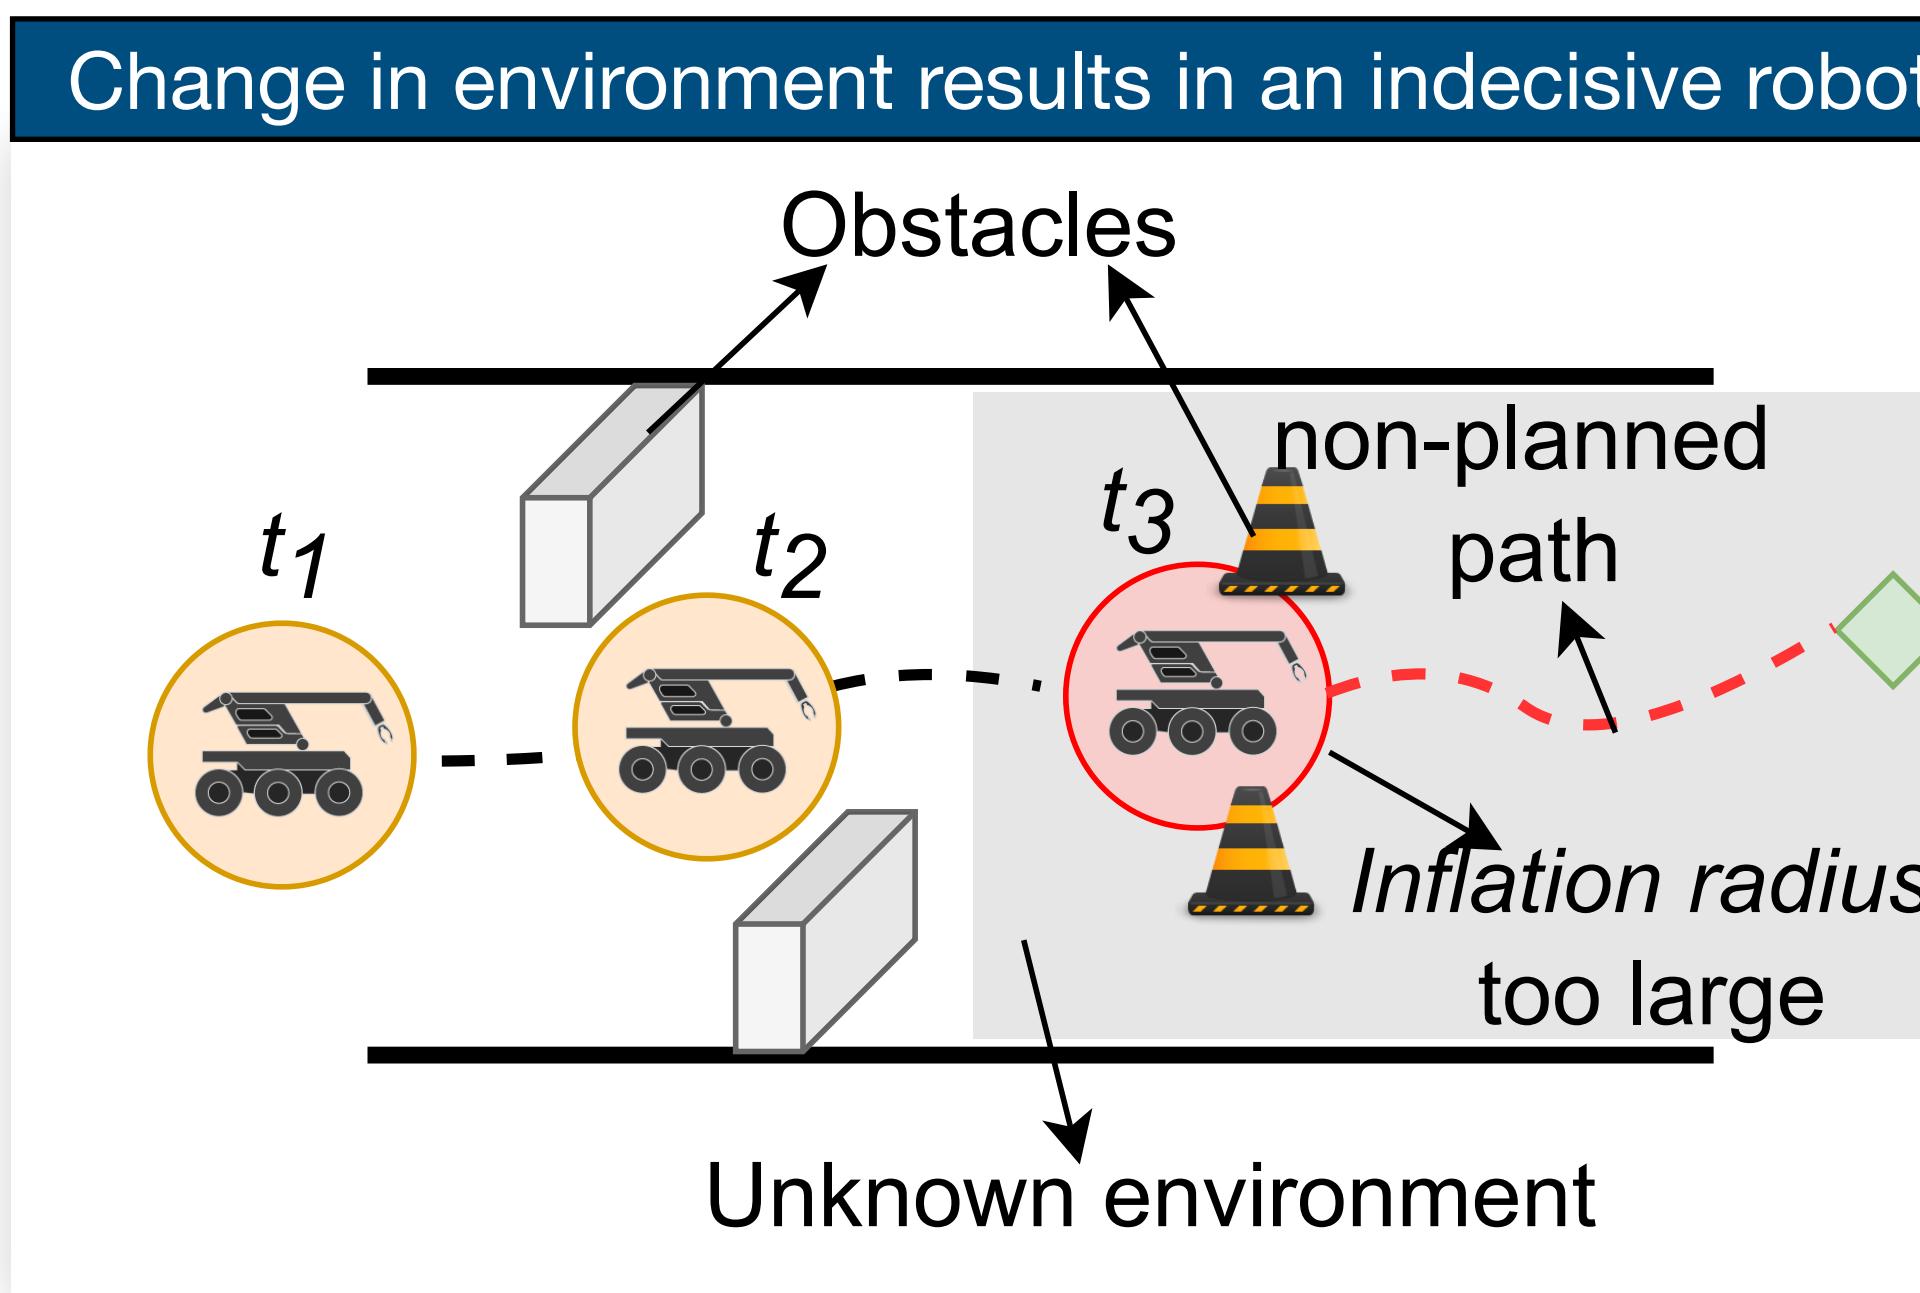
Motivating Scenarios

- Robot *stops 0.5 m away* from the target location and transmits *incorrect artifact locations* to the control station.
- A cause for this fault might be *a delay in data transformation* (the sensor transmits data at a rate of 1 Hz and the robot travels at 0.5 m/s).
- As a result, when the *costmap* receives data from the sensor, it is *a second old* and the robot has already traveled *0.5 m* away from that position.

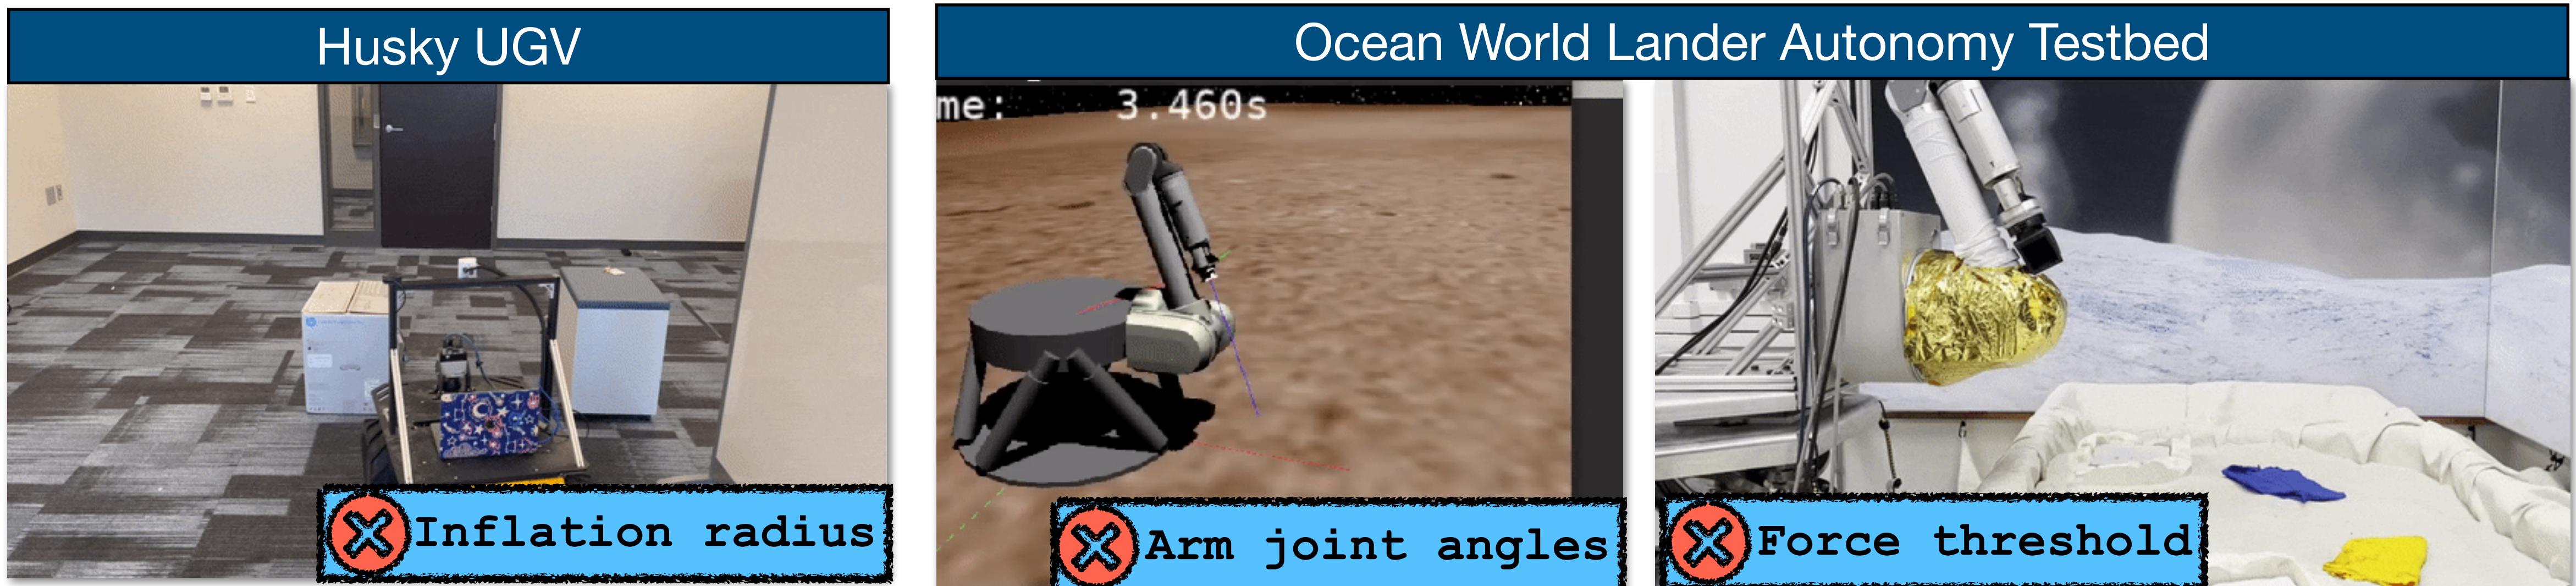
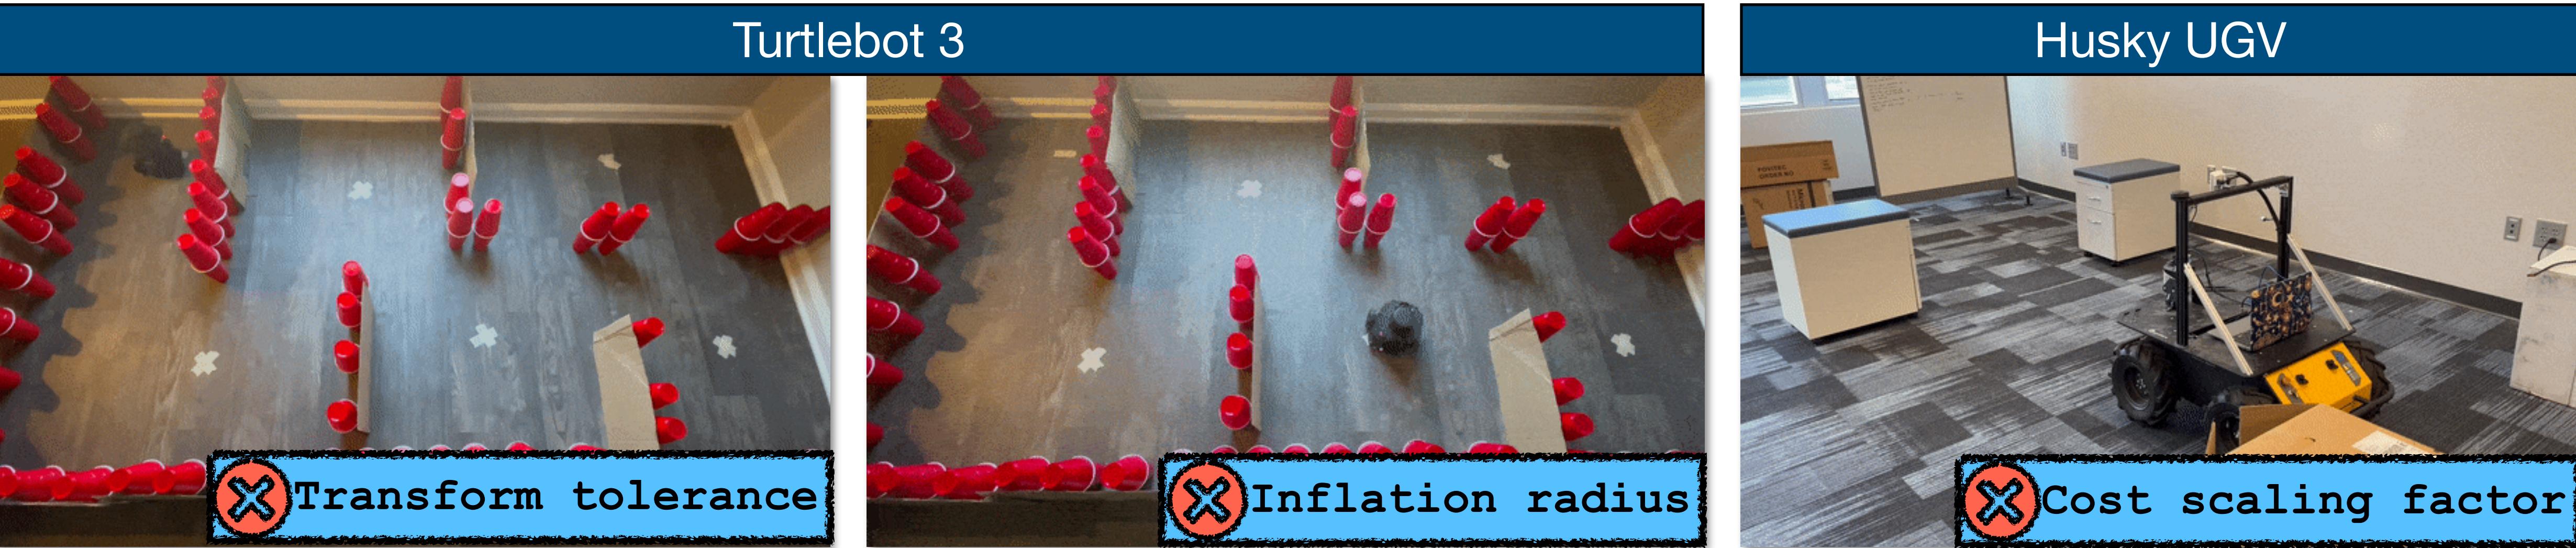


Motivating Scenarios

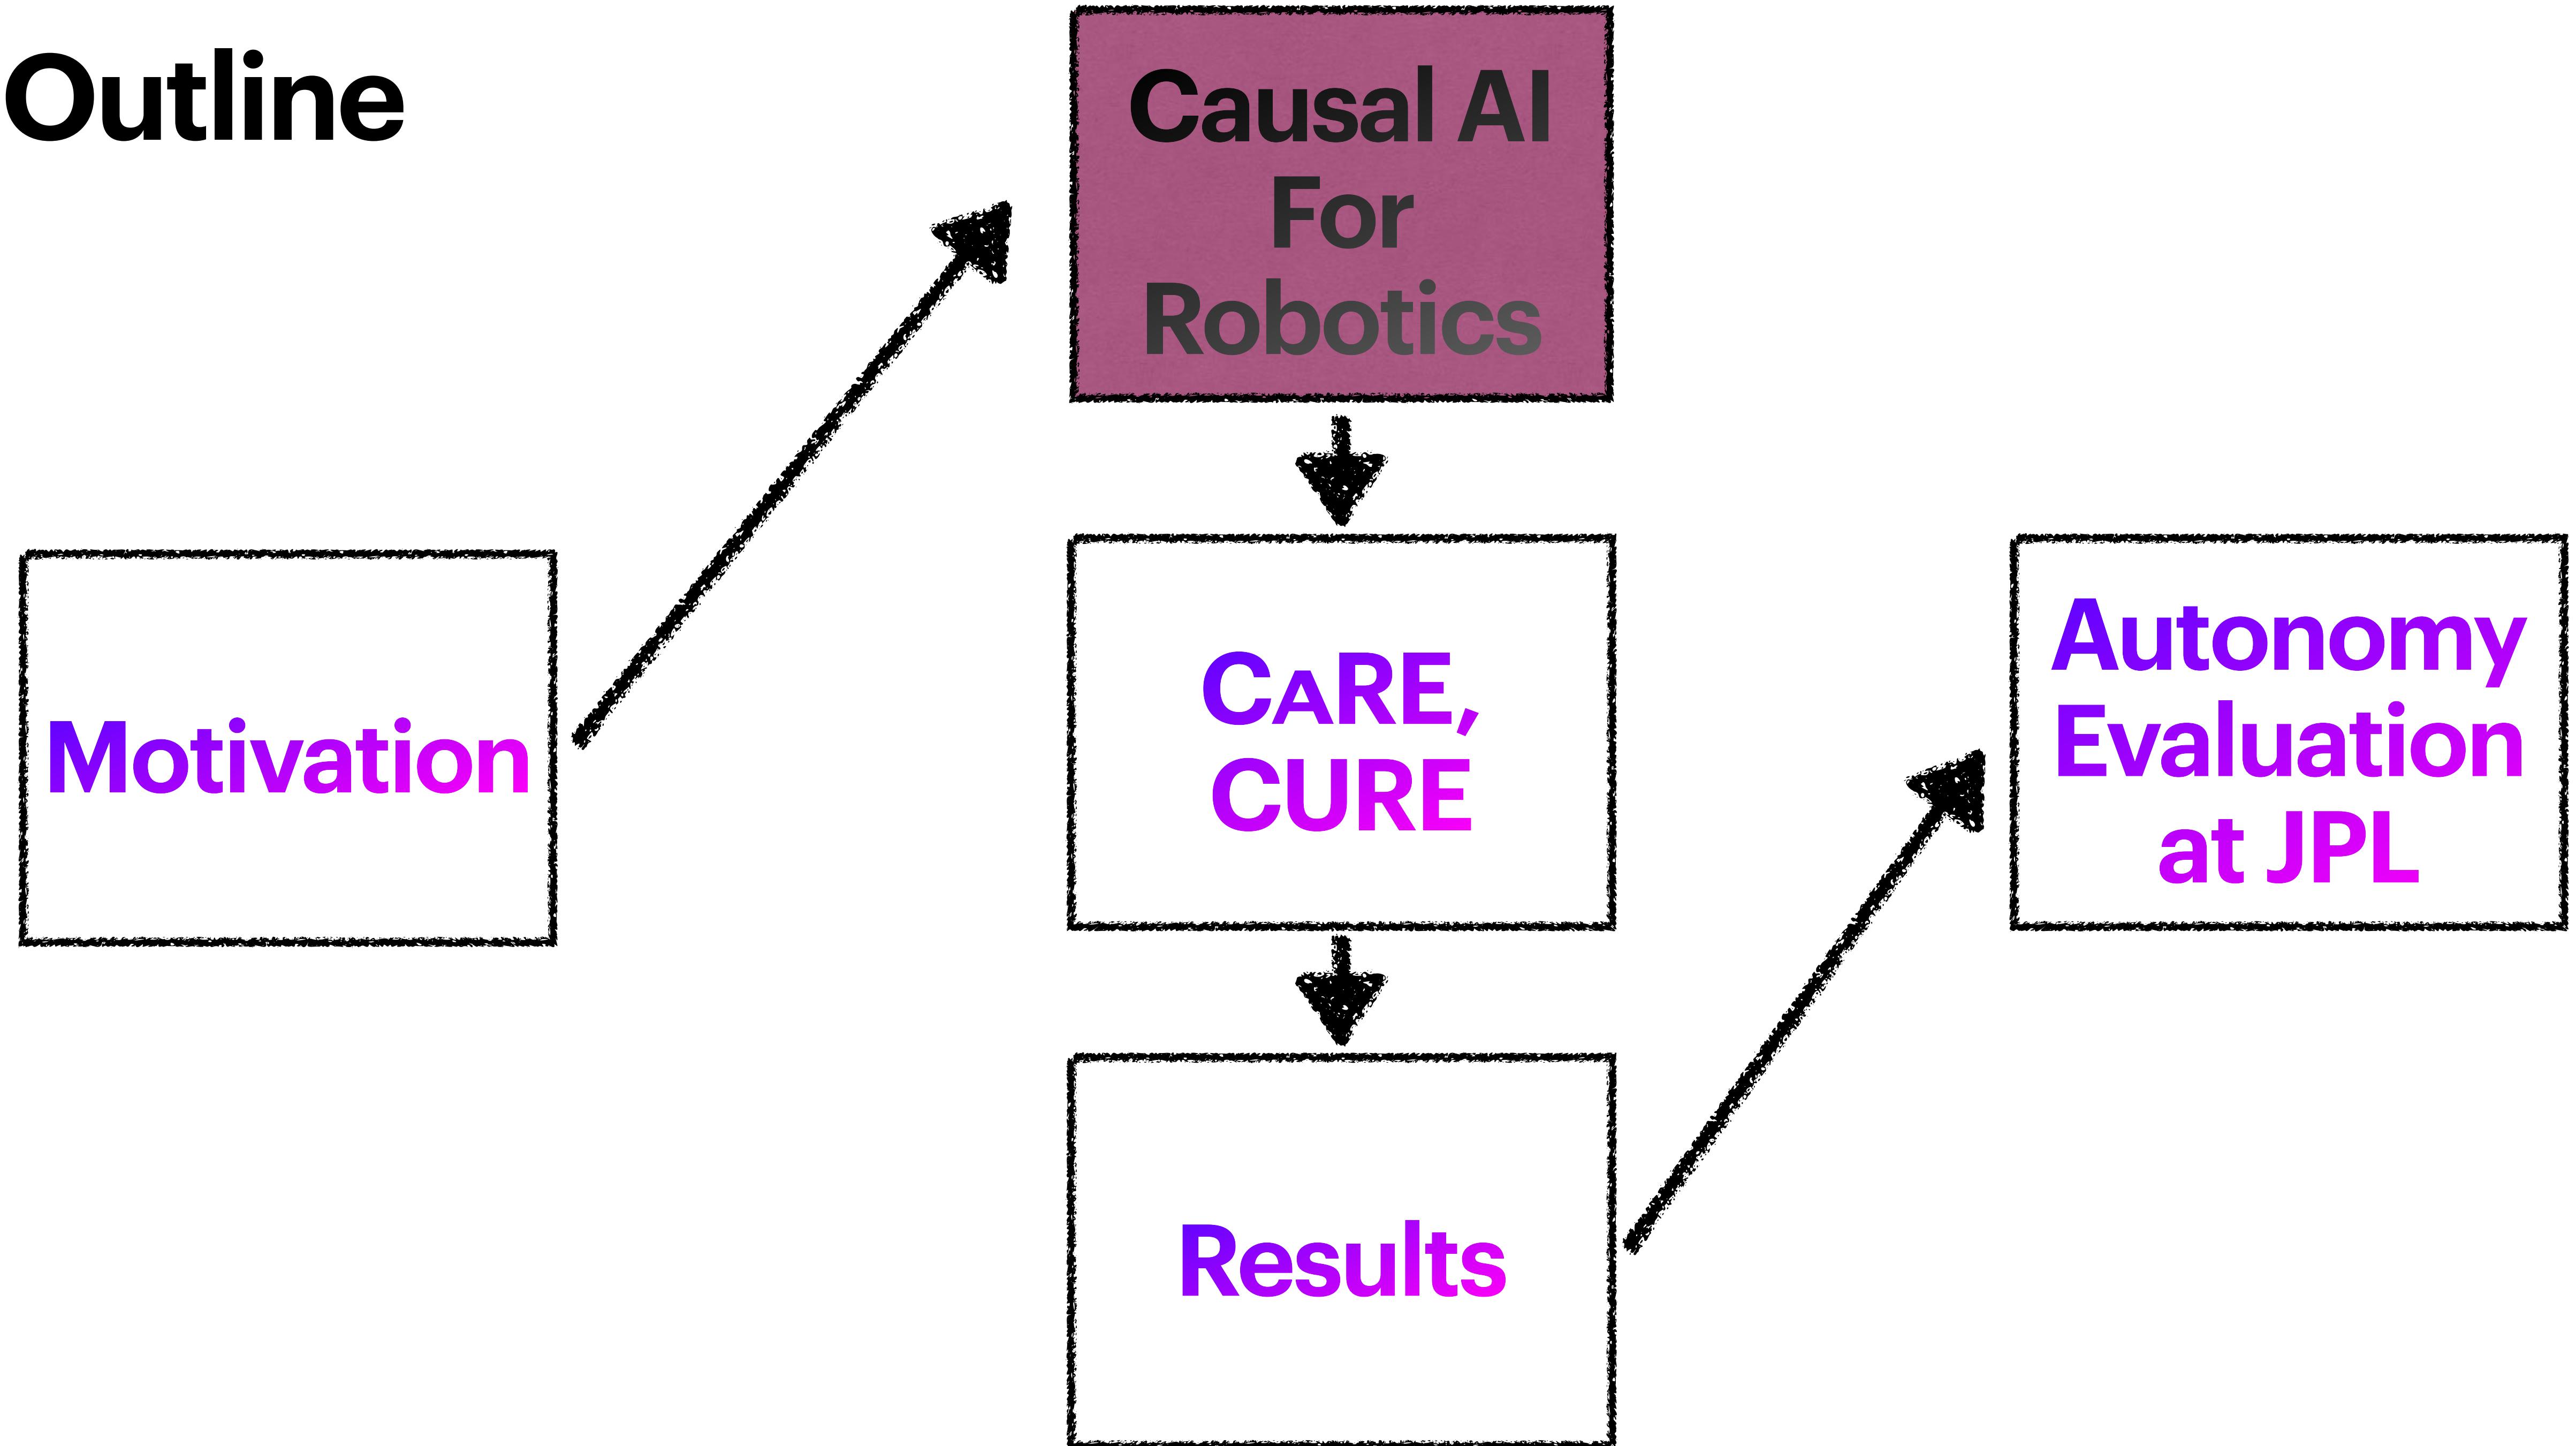
- At t_3 the robot encounters unique obstacles that too close together **violating the inflation radius**.
- Resulting in an **indecisive robot** that is **stuck in place**.



Motivating Scenarios



Outline



Causal AI in Systems and Software

Robotics

IEEE ROBOTICS AND AUTOMATION LETTERS, VOL. 7, NO. 4, OCTOBER 2022

Why Did I Fail? A Causal-Based Method to Find Explanations for Robot Failures

Maximilian Diehl , Graduate Student Member, IEEE, and Karinne Ramirez-Amaro , Member, IEEE

UAV

SWARMBUG: Debugging Configuration Bugs in Swarm Robotics

Chijung Jung

University of Virginia
Charlottesville, Virginia, USA
cj5kd@virginia.edu

Ali Ahad

University of Virginia
Charlottesville, Virginia, USA
aa5rn@virginia.edu

Jinho Jung

Georgia Institute of Technology
Atlanta, Georgia, USA
jinho.jung@gatech.edu

Sebastian Elbaum

University of Virginia
Charlottesville, Virginia, USA
selbaum@virginia.edu

Yonghwi Kwon

University of Virginia
Charlottesville, Virginia, USA
yongkwon@virginia.edu

Software Engineering

Causal Testing: Understanding Defects' Root Causes

Brittany Johnson

University of Massachusetts Amherst
Amherst, MA, USA
bjohnson@cs.umass.edu

Yuriy Brun

University of Massachusetts Amherst
Amherst, MA, USA
brun@cs.umass.edu

Alexandra Meliou

University of Massachusetts Amherst
Amherst, MA, USA
ameli@cs.umass.edu

Operating System

COZ: Finding Code that Counts with Causal Profiling

Charlie Curtsinger Emery D. Berger

School of Computer Science
University of Massachusetts Amherst
{charlie, emery}@cs.umass.edu

Big Data

Causality-Guided Adaptive Interventional Debugging

Anna Fariha

University of Massachusetts Amherst
afariha@cs.umass.edu

Suman Nath

Microsoft Research
Suman.Nath@microsoft.com

Alexandra Meliou

University of Massachusetts Amherst
ameli@cs.umass.edu

Database

Demonstration of Inferring Causality from Relational Databases with CaRL

Moe Kayali, Babak Salimi, Dan Suciu

{kayali, bsalimi, suciu}@cs.washington.edu
University of Washington

Misconfiguration and its Effects

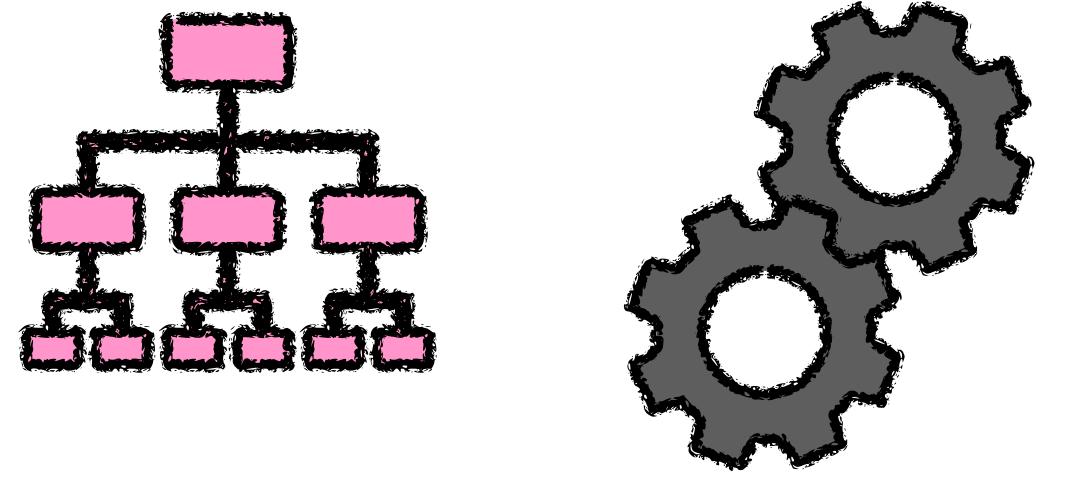
- Misconfigurations can elicit unexpected interactions between software and hardware
- These can result in non-functional faults
 - Affecting non-functional system properties like latency, throughput, energy consumption, etc.



The system doesn't crash or exhibit obvious misbehavior



Systems are still operational but with a degraded performance, e.g., high latency, low throughput, high energy consumption, high heat dissipation, or a combination of several



My research goal is to improve performance by considering configuration level changes.

Motivating Example

<https://github.com/ros-navigation/navigation2/issues/2439>

Error while starting navigation using local planner(DWB) for hardware

Closed

Sooifyan opened this issue on Jul 7, 2021 · 5 comments



ros-navigation / navigation2



Sooifyan commented on Jul 7, 2021 · edited

Bug report

Required Info:

- Operating System:
 - Ubuntu 20.04
- ROS2 Version:
 - Foxy
- Version or commit hash:
 - Latest of foxy-devel branch(2nd July 2021)
- DDS implementation:

Steps to reproduce issue

This error is seen in simulation when the robot collides with obstacles(similar issue).

Expected behaviour

After giving a goal, the global planner is properly planning the path towards the goal. The expected behaviour is that the path must follow this path using a local planner.

Actual behavior

Here the path is always failing and starts to spin, And the bot is not hitting any obstacle, and transformations are perfect. The robot is not even moving a little bit; it directly starts the recoveries servers as soon as a global planner is planned.

This user is trying to fix an error encountered by local planner. The local planner is **not following** the **path** planned by the global planner.

The user expects the local planner to follow the path planned by the global planner.

Performance debugging aims at finding the **root cause** of the misconfiguration and **fix it**.

Motivating Example

The screenshot shows a GitHub issue thread titled "ros-navigation / navigation2". The first comment, by SteveMacenski, is highlighted with a green border and contains a link to a search result: <https://github.com/ros-planning/navigation2/search?q=Trajectory+goes+off>. The second comment, by Soofiyan, is highlighted with a red border. The third comment, also by SteveMacenski, is highlighted with a green border at the bottom.

SteveMacenski commented on Jul 7, 2021 · edited

You can certainly tune the DWB controller into behavior where it has a hard time moving. Have you tuned the parameters of DWB and local costmap from the defaults in nav2_bringup? (not that those are perfect by any stretch either, but its a baseline we can use to test that hypothesis) The error also looks to be pretty straight forward: <https://github.com/ros-planning/navigation2/search?q=Trajectory+goes+off>, is your costmap reasonably sized? Please provide your configuration of DWB / local costmap. I also don't see a footprint in rviz, are you sure that is set correctly?

Soofiyan commented on Jul 7, 2021

Thank you for replying to my query. I have tuned this for almost 5-6 hours. Sometimes it is going towards the goal but still failing in the middle of the trajectory. 9/10 times it is still going into the recovery mode without hitting any obstacle, as mentioned earlier. If you can suggest any change in a parameter, it will be beneficial.
Here is my nav2_param file, which I am using

SteveMacenski commented on Jul 8, 2021

Please ask on ROS Answers -- I can't provide personalized tuning assistance to every user. If you identify a tangible bug or issue, I'm happy to discuss a resolution to that.

The developer is asking if the **configuration** of local planner was **set correctly?**

The user tuned the parameters for **almost 5-6 hours**, yet **9/10 times the robot is failing in the middle of the trajectory.**

The developer concluded that he **cannot provided personalized tuning assistance to every user.**

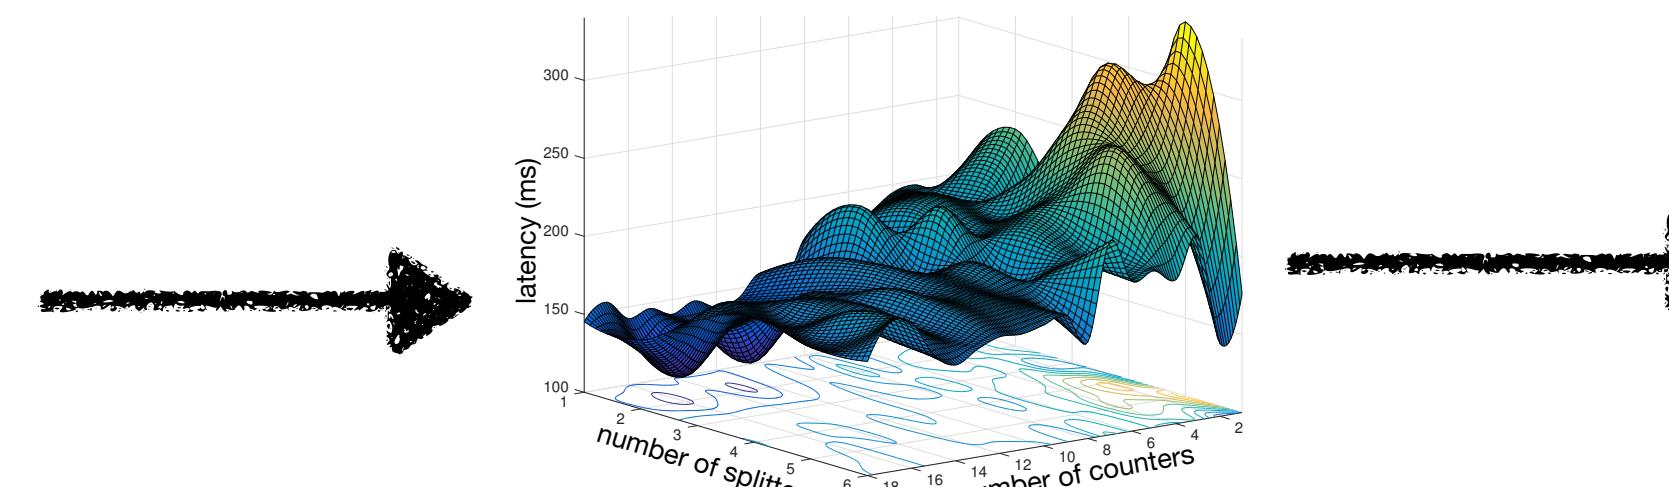
**How to resolve these issues
in an automatic manner?**

Performance Influence Models

	Path dist bias	Transform tolerance	...	Traveled distance	...	Mission success
c_1	0.75	0.5	...	8.32 m	...	1
c_2	1.5	1.2	...	12.6 m	...	0
...
c_n	0.1	0.2	...	7.43 m	...	0

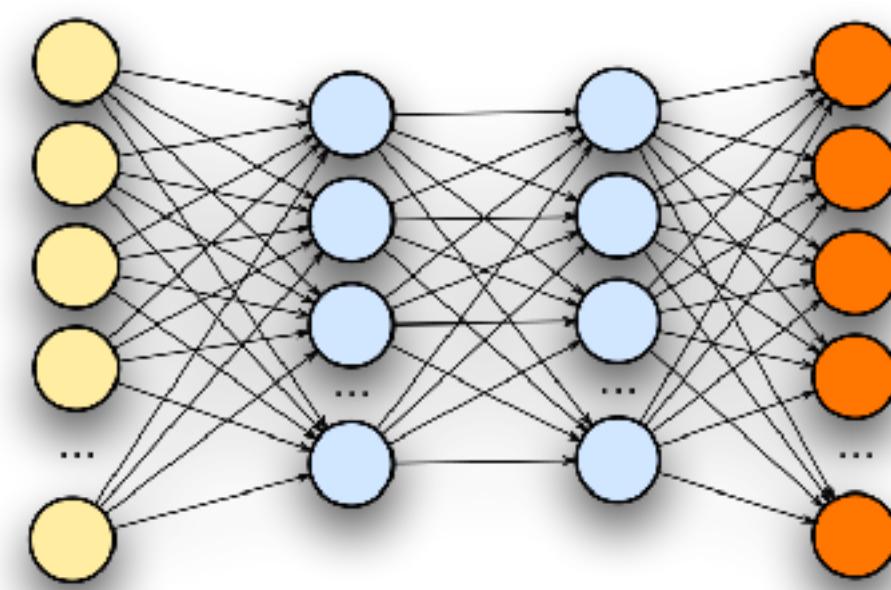


Observational Data



$$Y_i = \beta_0 + \beta_1 X_i + \epsilon$$

Regression Equation



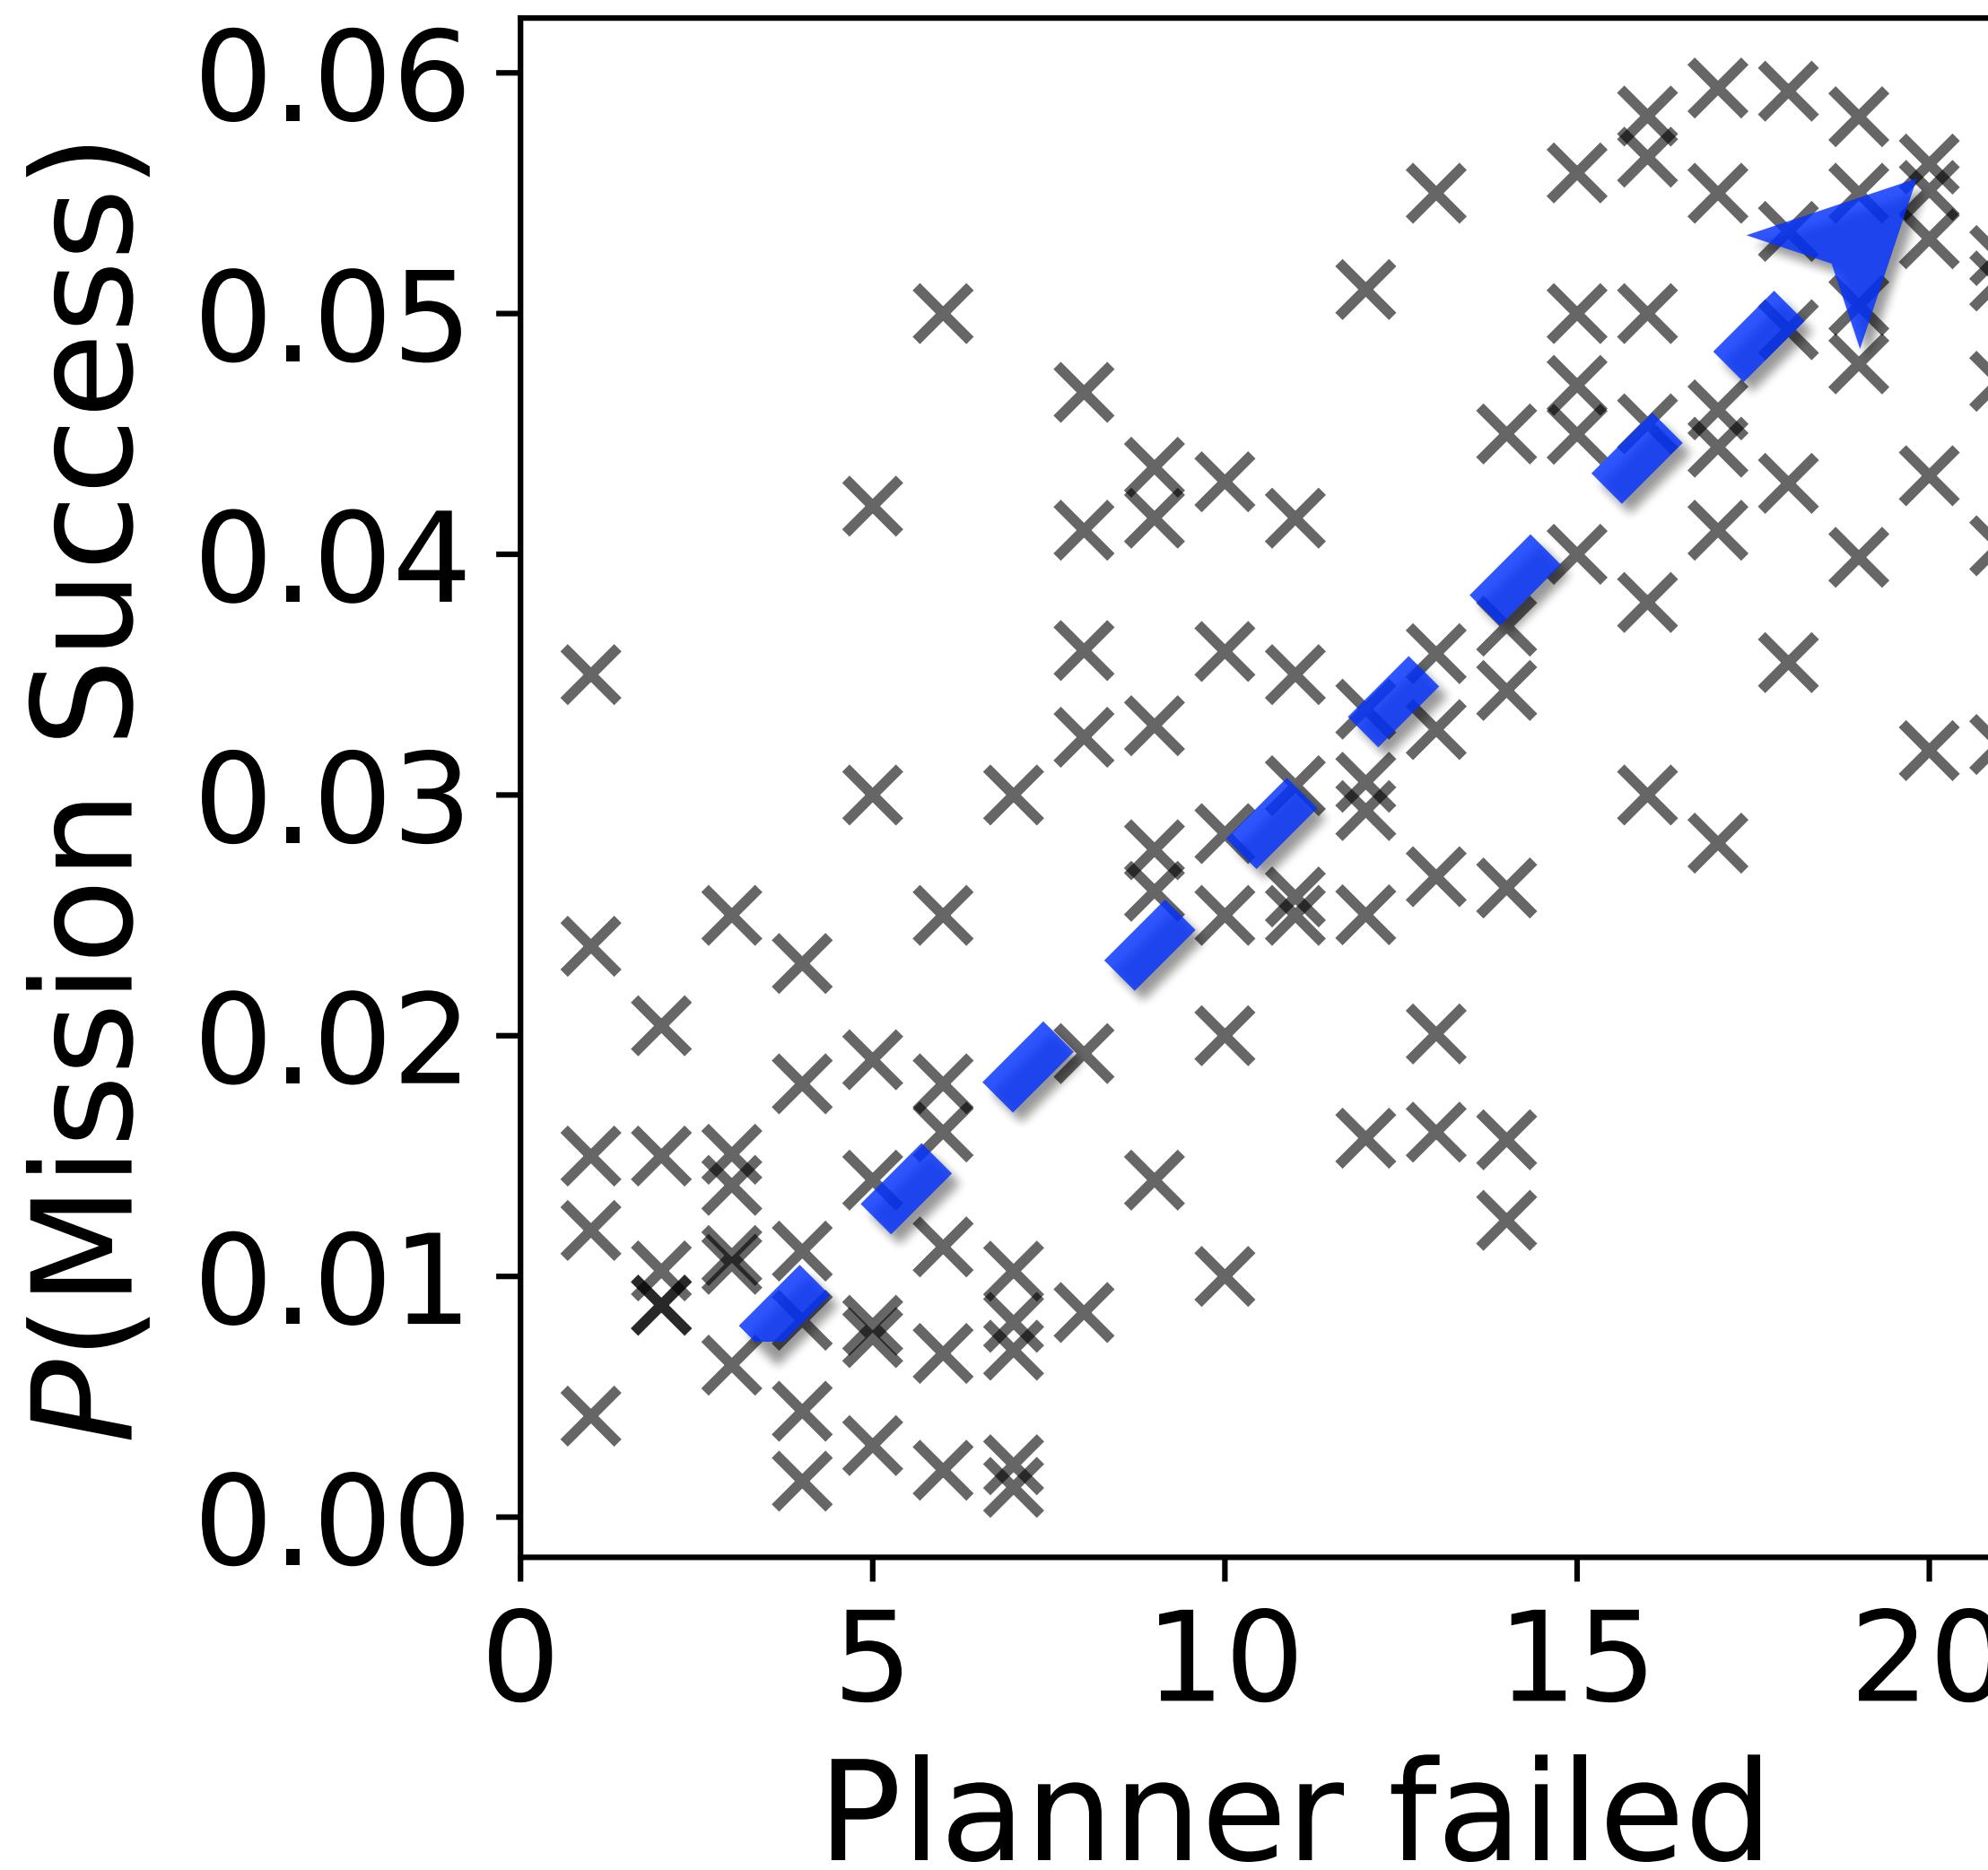
Black-box models

These methods rely on **statistical correlations** to **extract meaningful information** required for performance tasks.

Challenges

- **Trial-and-error** requires **non-trivial human efforts** due to the **large configuration space**.
- **Even after finding the optimal fix**, the new fix is **not guaranteed** to function in **different environments**.
- **Performance influence models** suffer from several shortcomings:
 - **Non-transferable** across different robotic platforms and environments
 - Could produce **Incorrect explanations**
 - Could produce **Unreliable predictions**
 - Data collection is **expensive** from physical hardware

Incorrect Reasoning About The Robot's Behavior



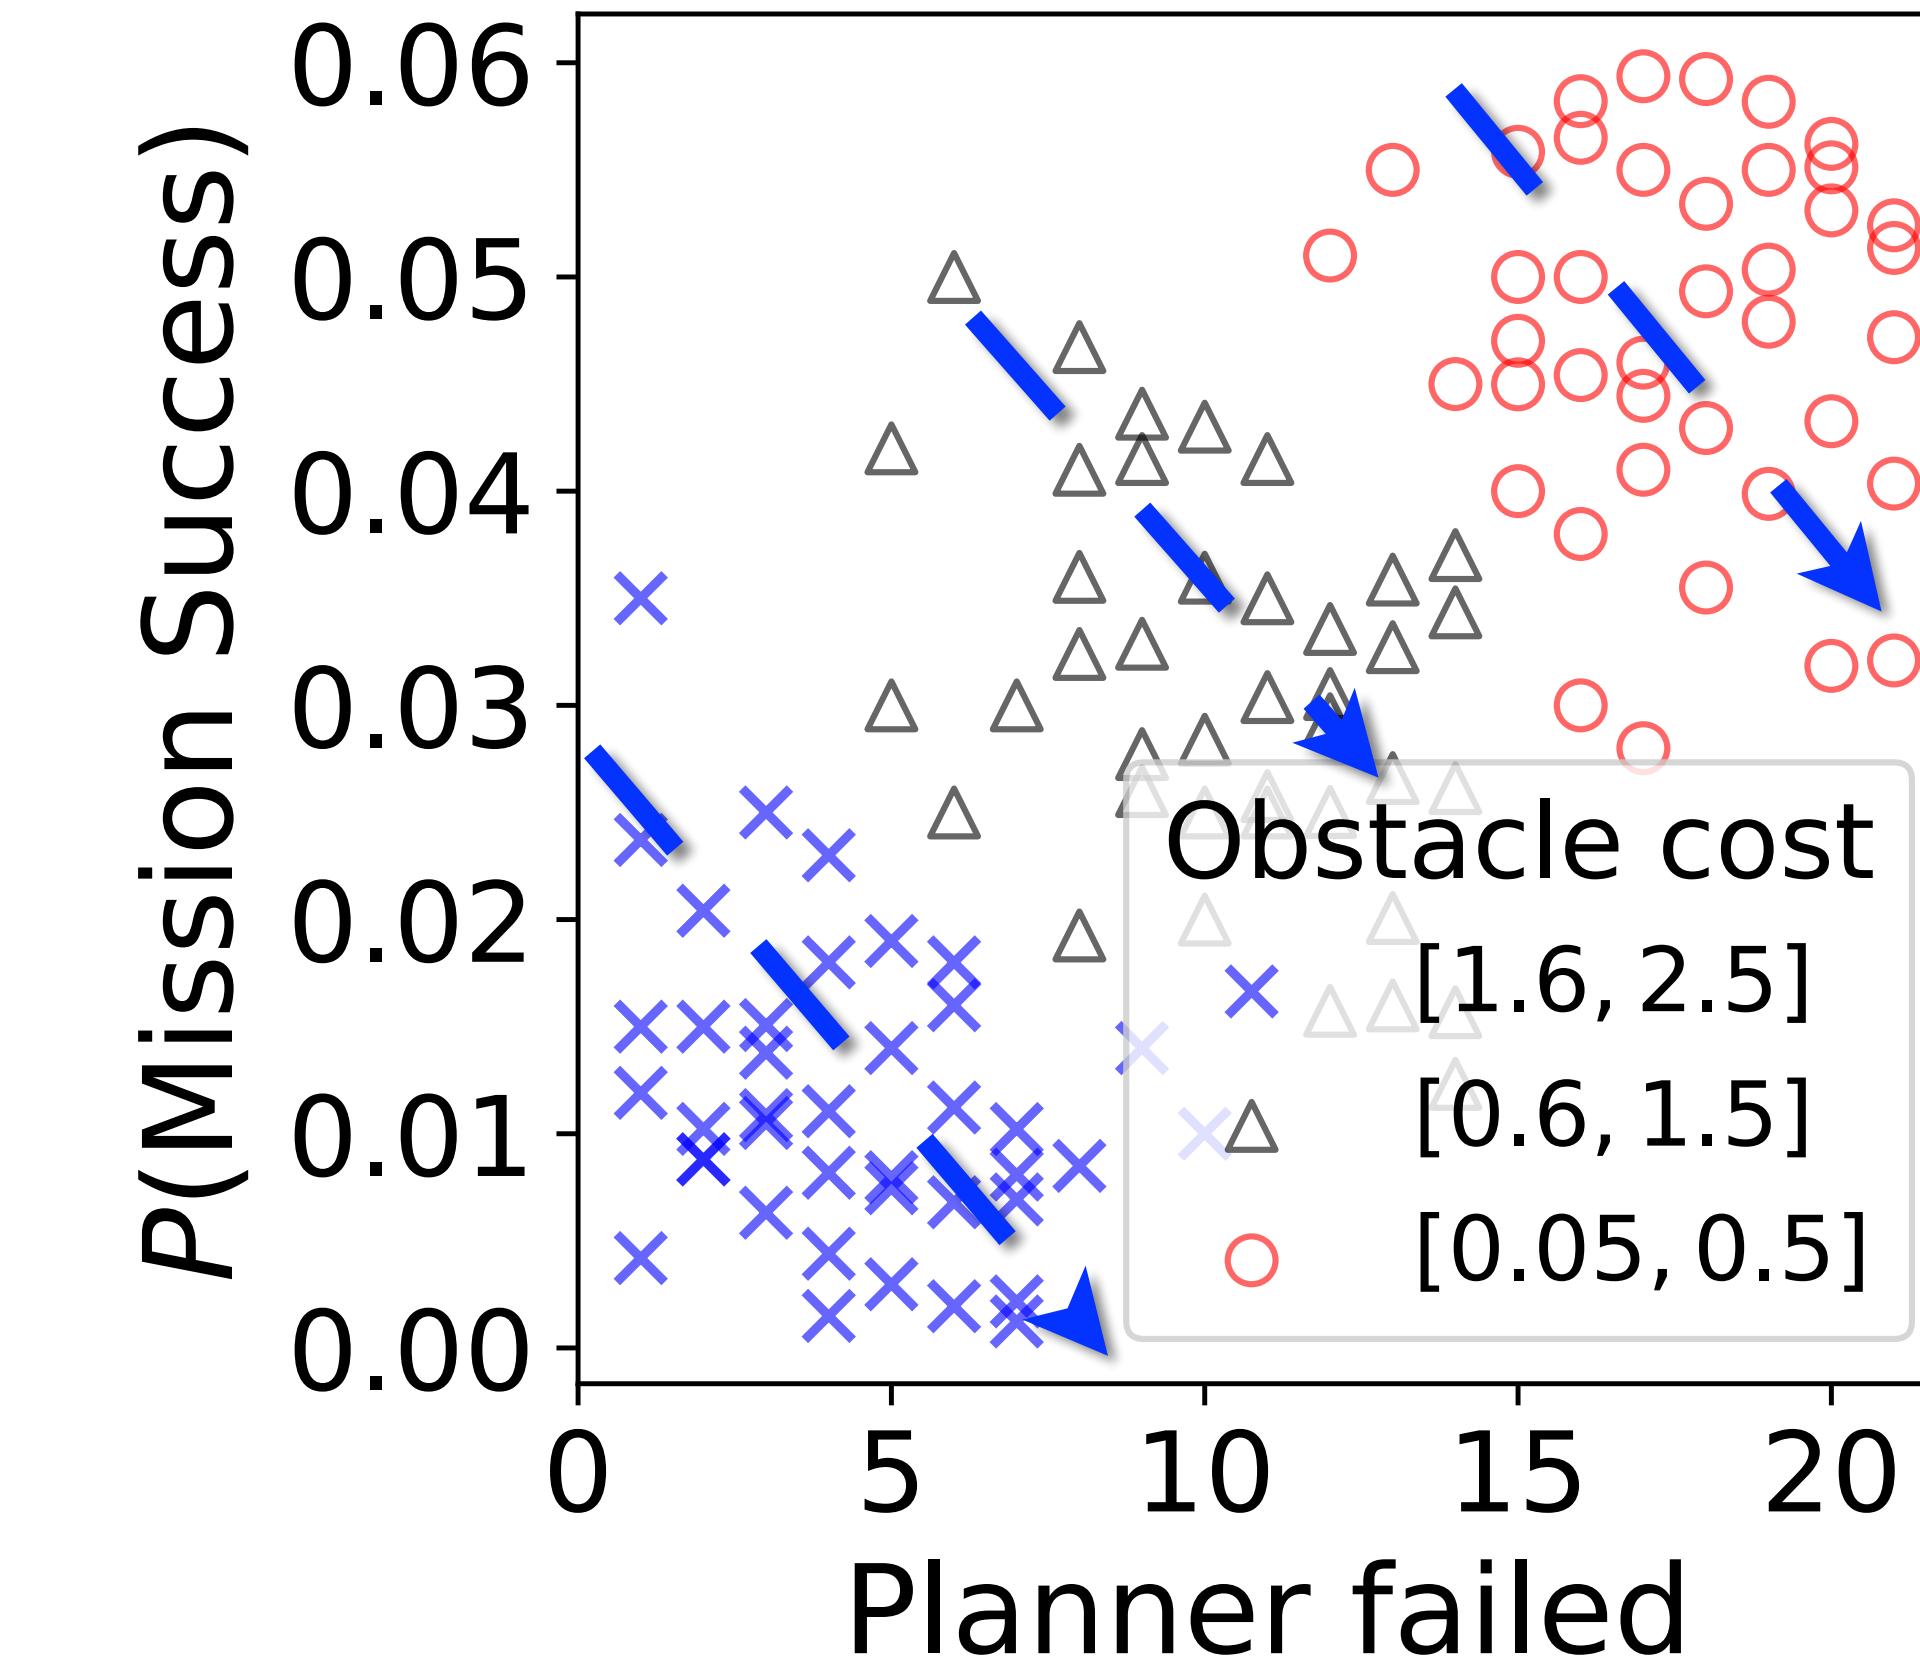
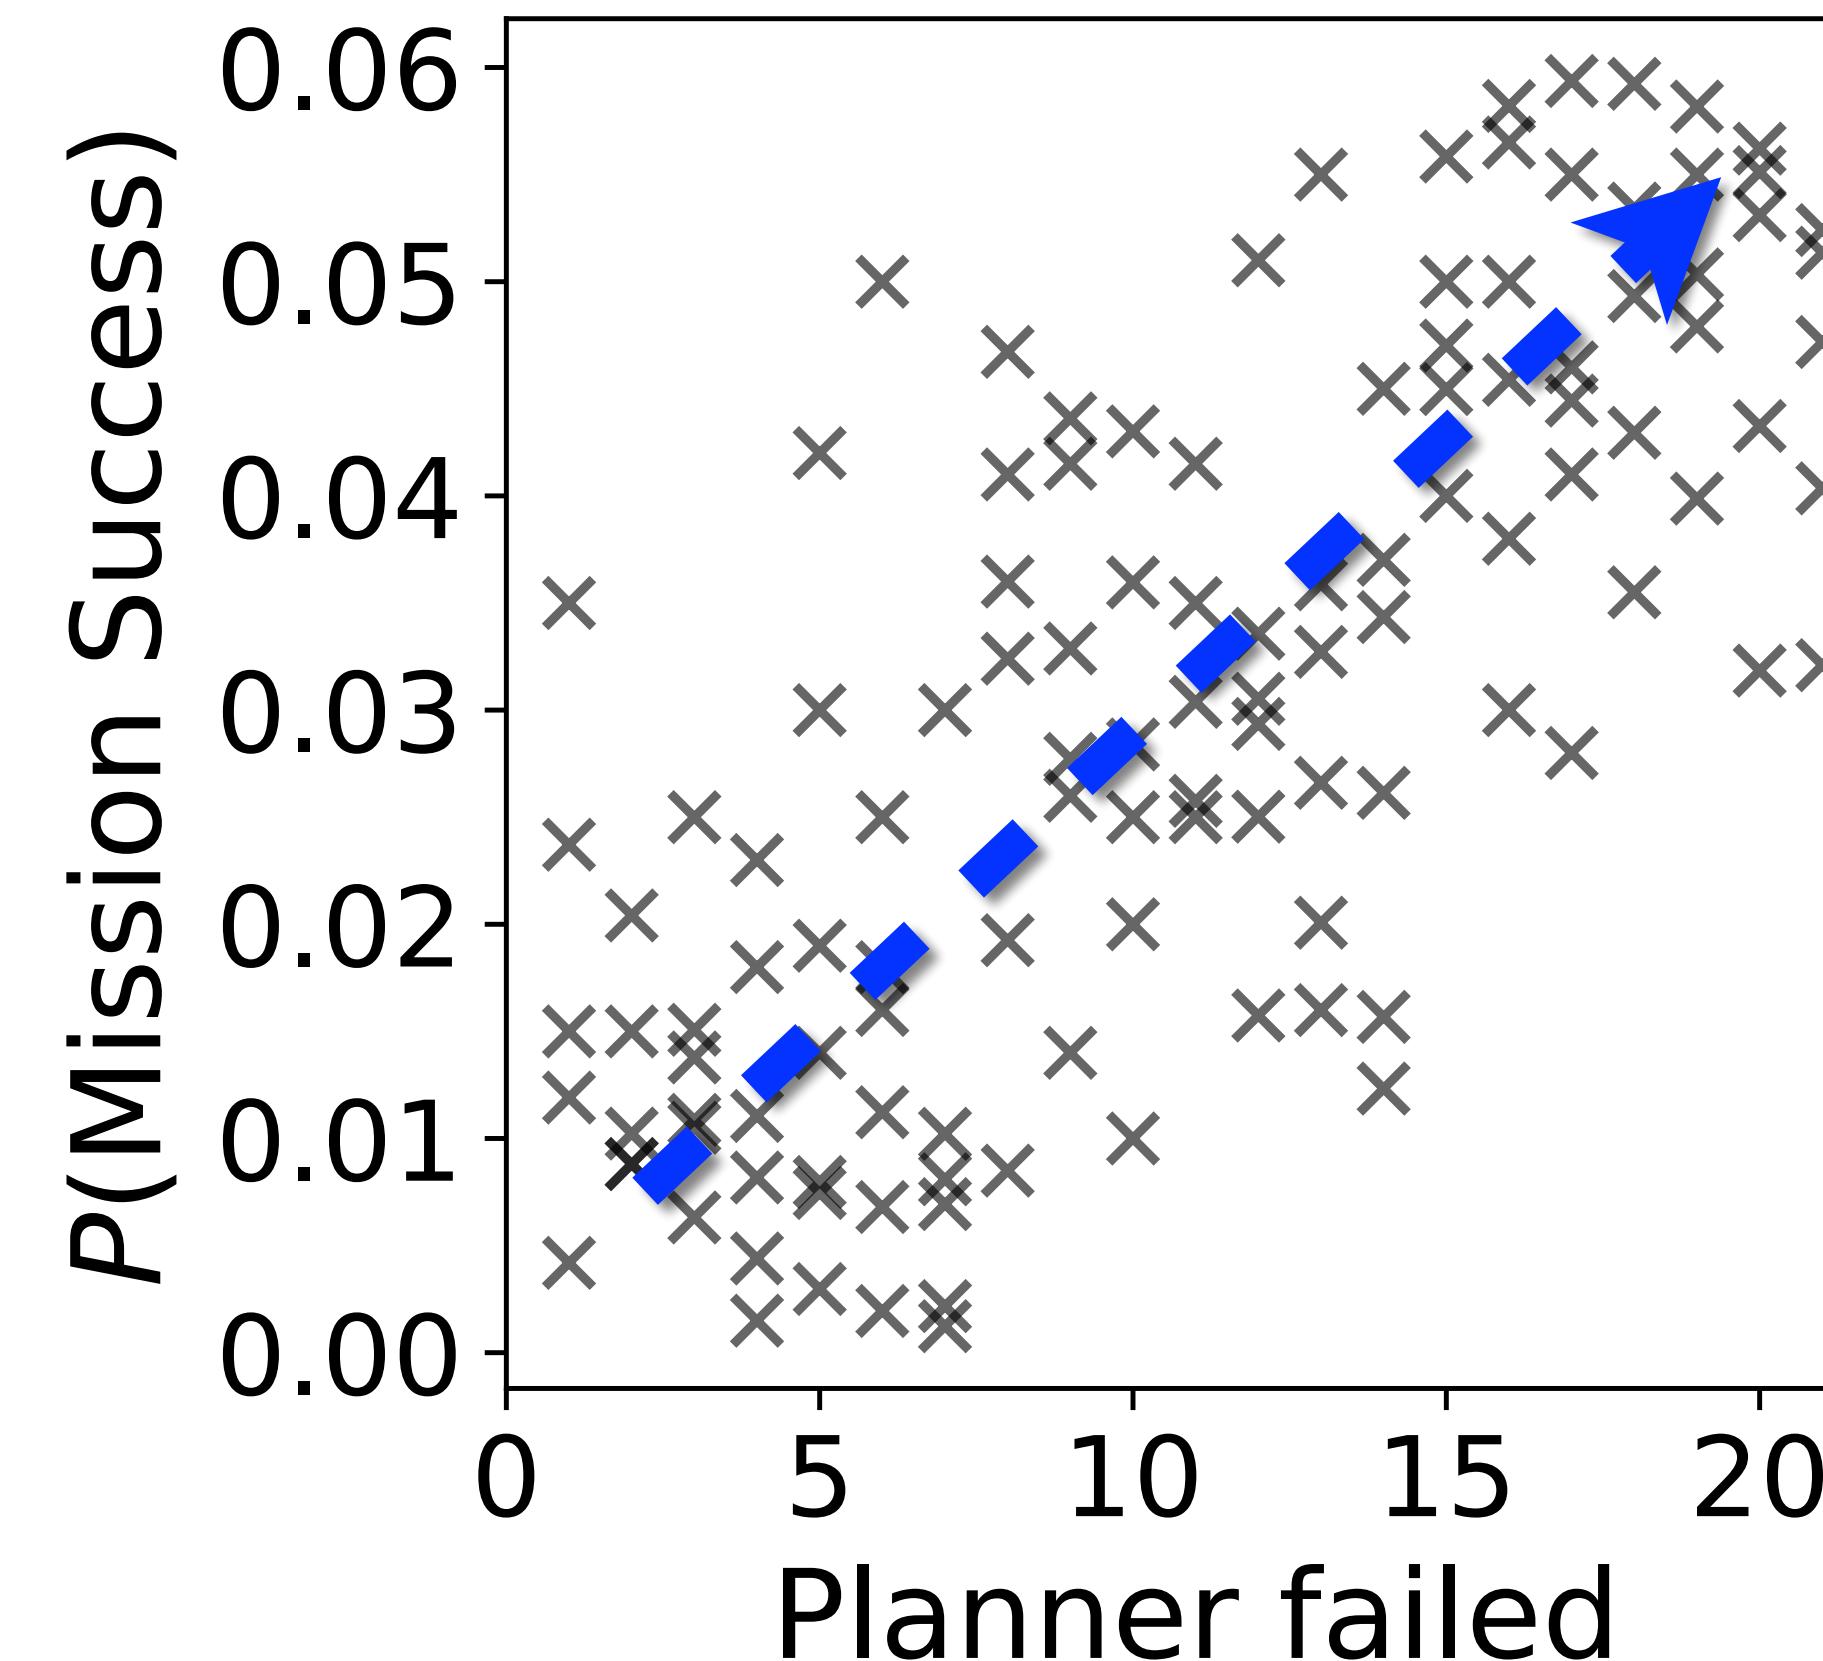
Increasing Planner failed increases $P(\text{Mission success})$

This is counter-intuitive

More Planner failed should reduce $P(\text{Mission success})$ not increase it

Purely statistical models built on this data will be unreliable.

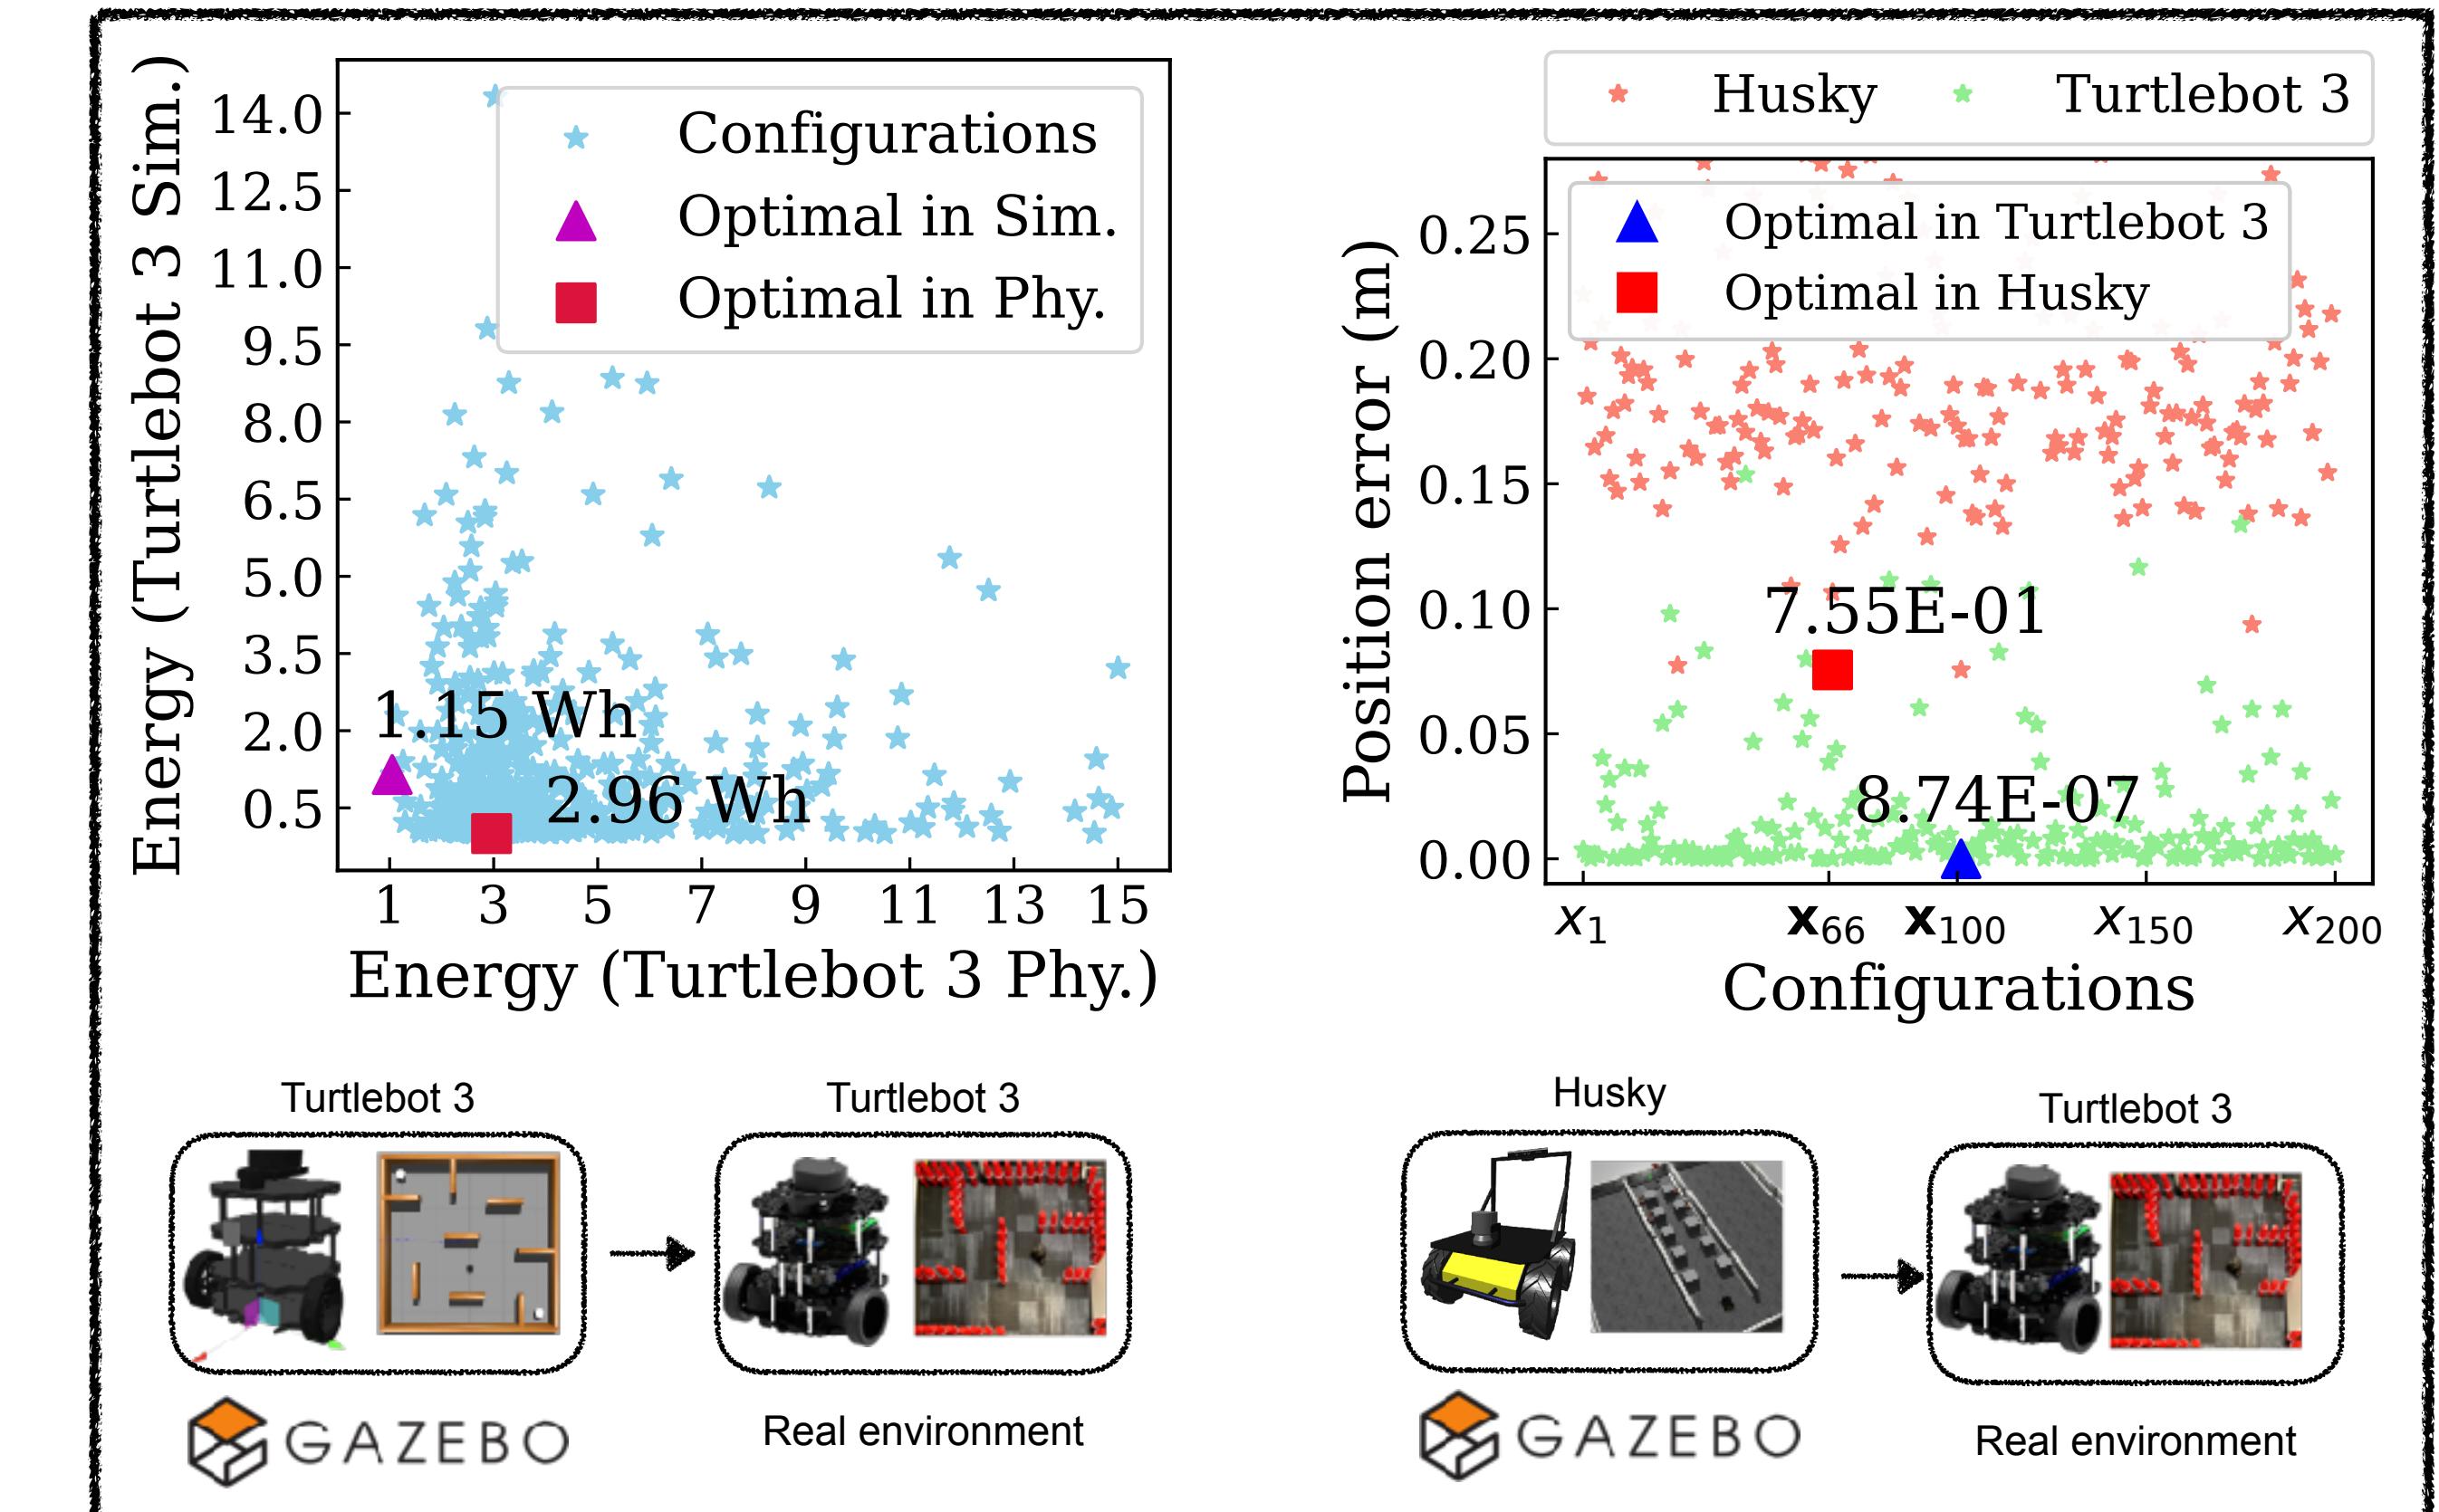
Performance Influence Models might be Unreliable



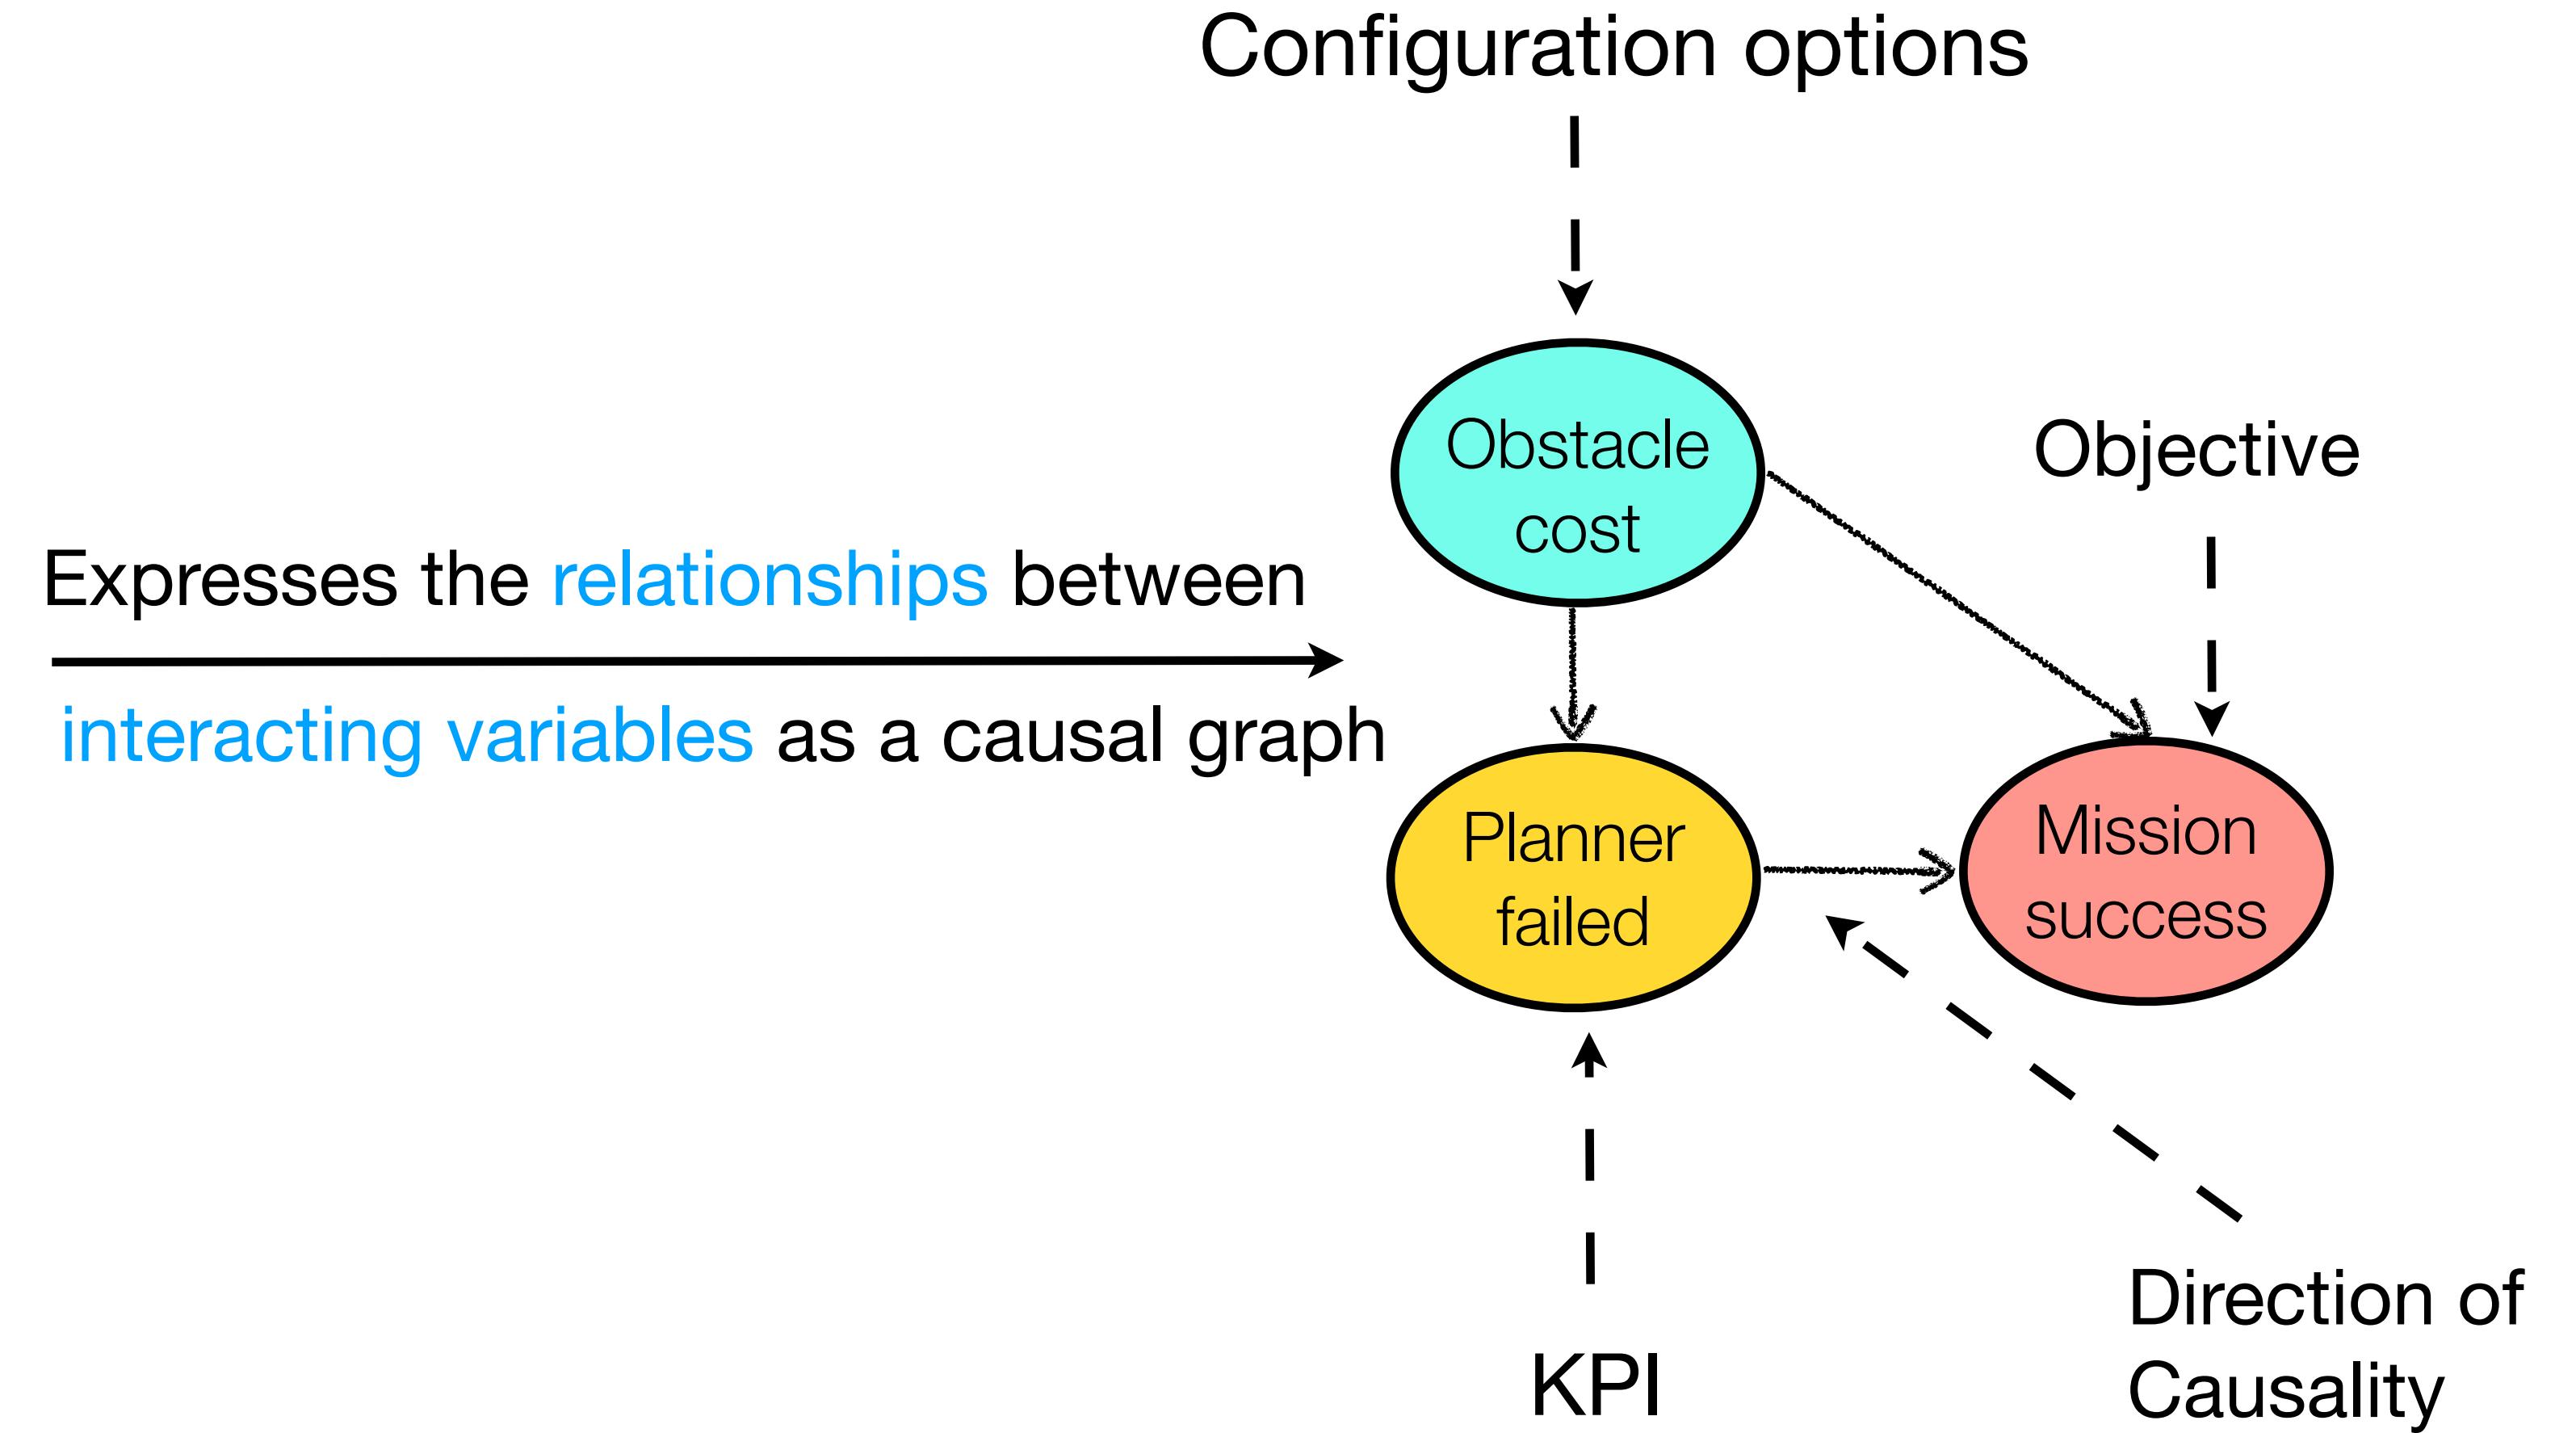
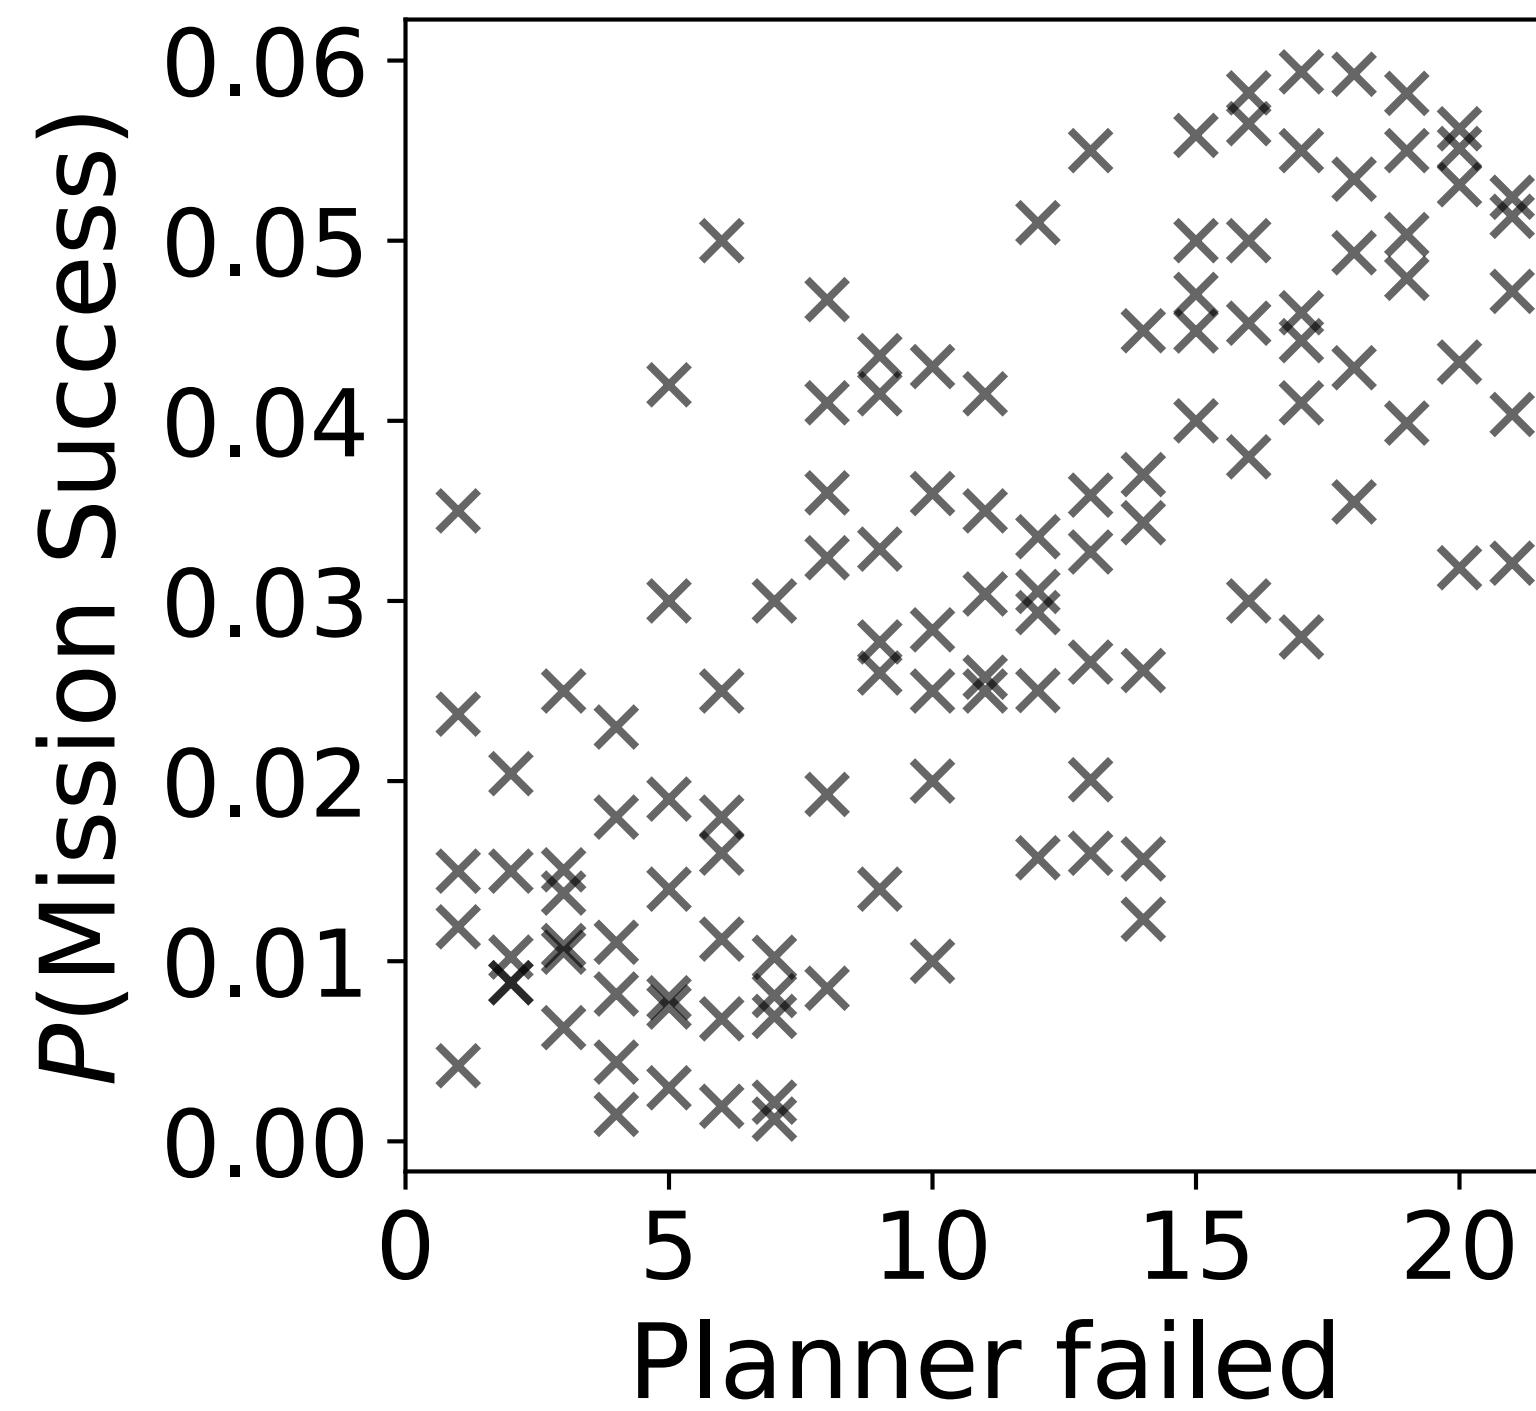
Segregating data on **Obstacle Cost** indicates that within **each group** increase of **Planner Failed** result in a **decrease in $P(\text{Mission success})$** .

Optimal configuration in one deployment environment does not remain optimal in another.

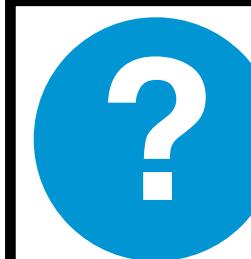
- Optimal configuration for *Turtlebot 3* in simulation differs from its physical counterpart (energy consumption is increased by 2.57x).
- Optimal configuration for *Turtlebot 3* is not suitable in *Husky* (significant increase in position error by 8.64×10^5 times).
- Machine learning models employed in the transfer learning approaches are vulnerable to spurious correlations.



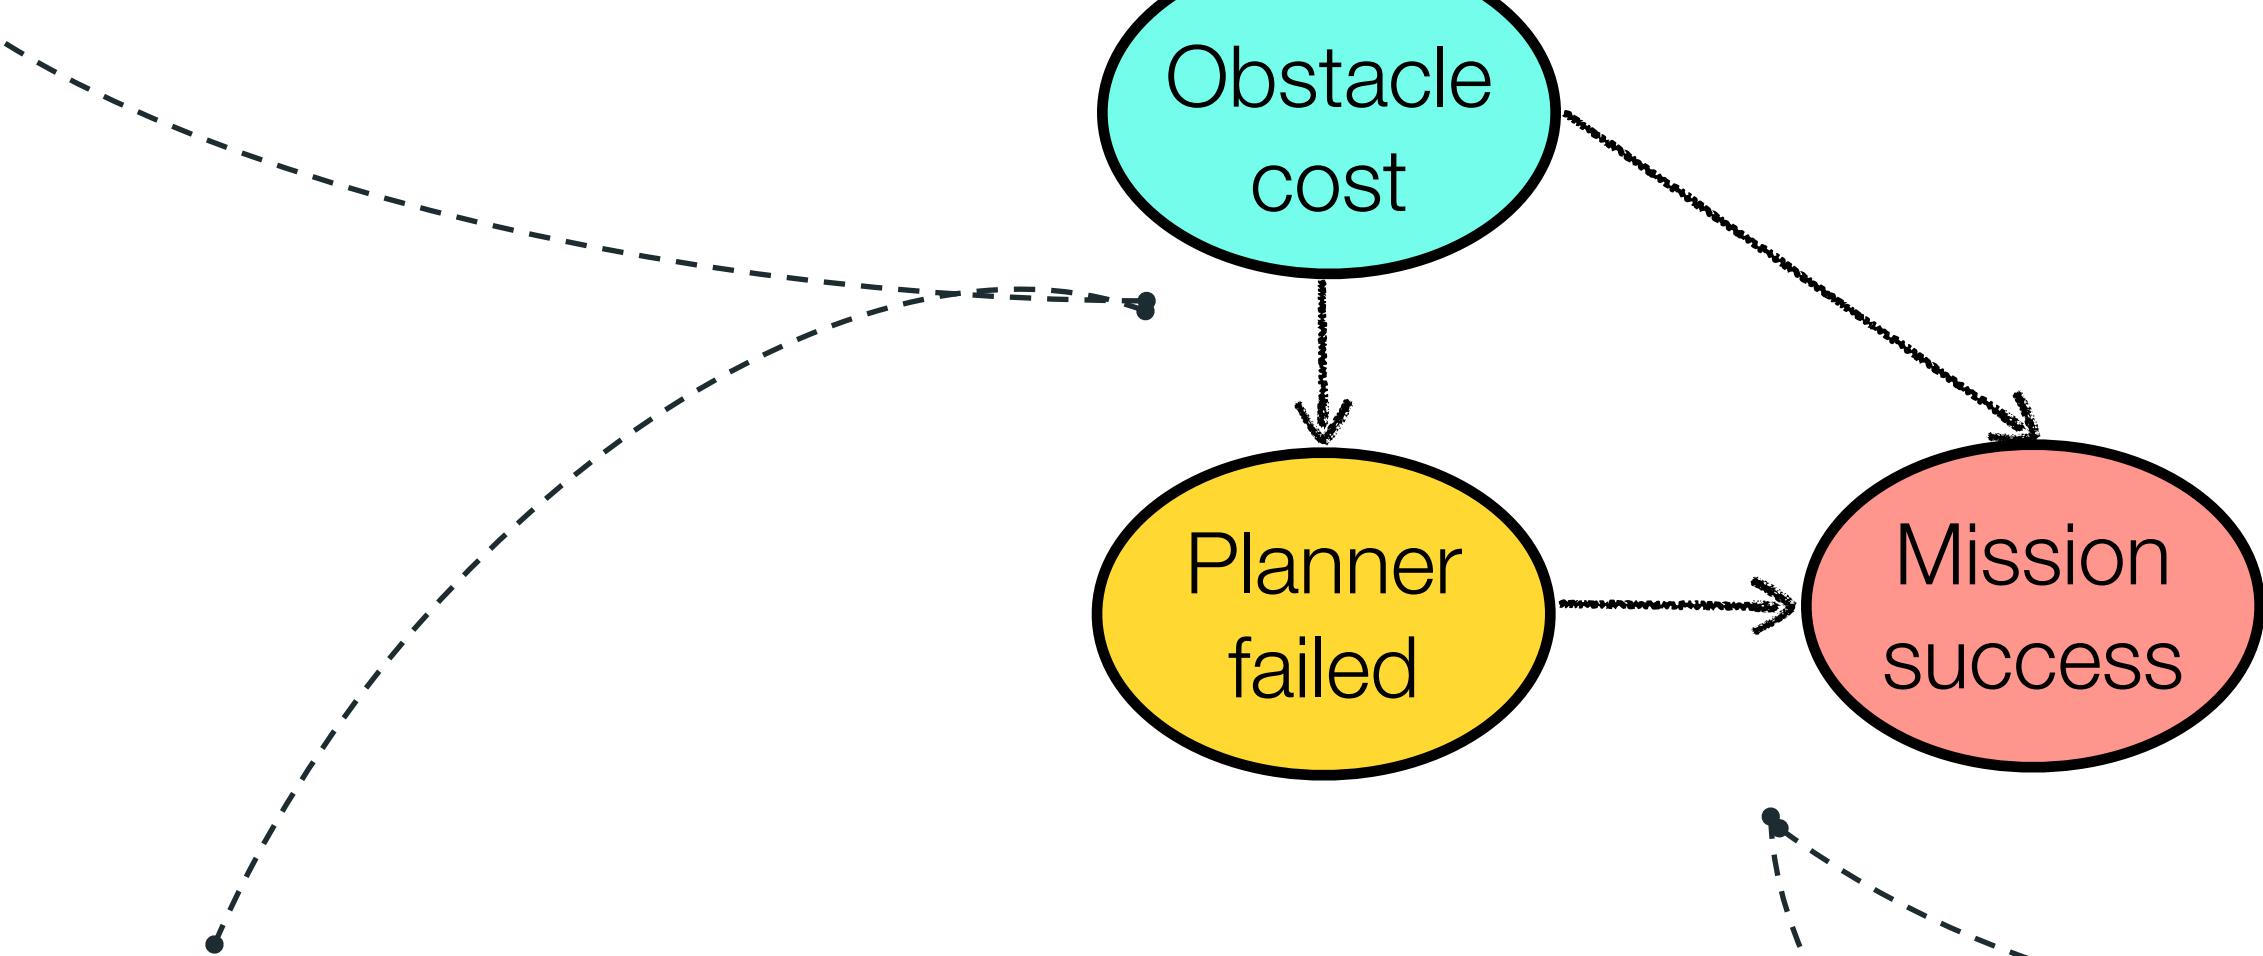
Our Key Idea: Building Causal Performance Model instead of Performance Influence Models



How to use Causal Models?



How to generate a causal model?



Learning

Given the observational data D , recover the causal graph $C_{\mathcal{G}}$ that encodes the dependency structure between \mathcal{X} , \mathcal{S} and \mathcal{Y} of \mathcal{V} such that the following structural constraints are satisfied.

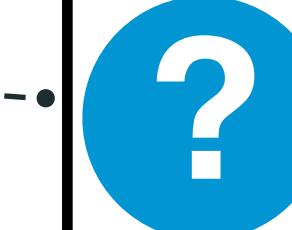
$$v_i \not\perp v_j \quad \forall v_i \in \mathcal{X} \subset \mathcal{V}, \forall v_j \in \mathcal{Y} \subset \{\mathcal{V} \setminus \mathcal{X} \setminus \mathcal{S}\}$$

Inference

Given the causal graph $C_{\mathcal{G}}$, determine the configuration option in \mathcal{X} which is the root cause for the observed functional fault characterized by performance objectives \mathcal{Y} as follows:

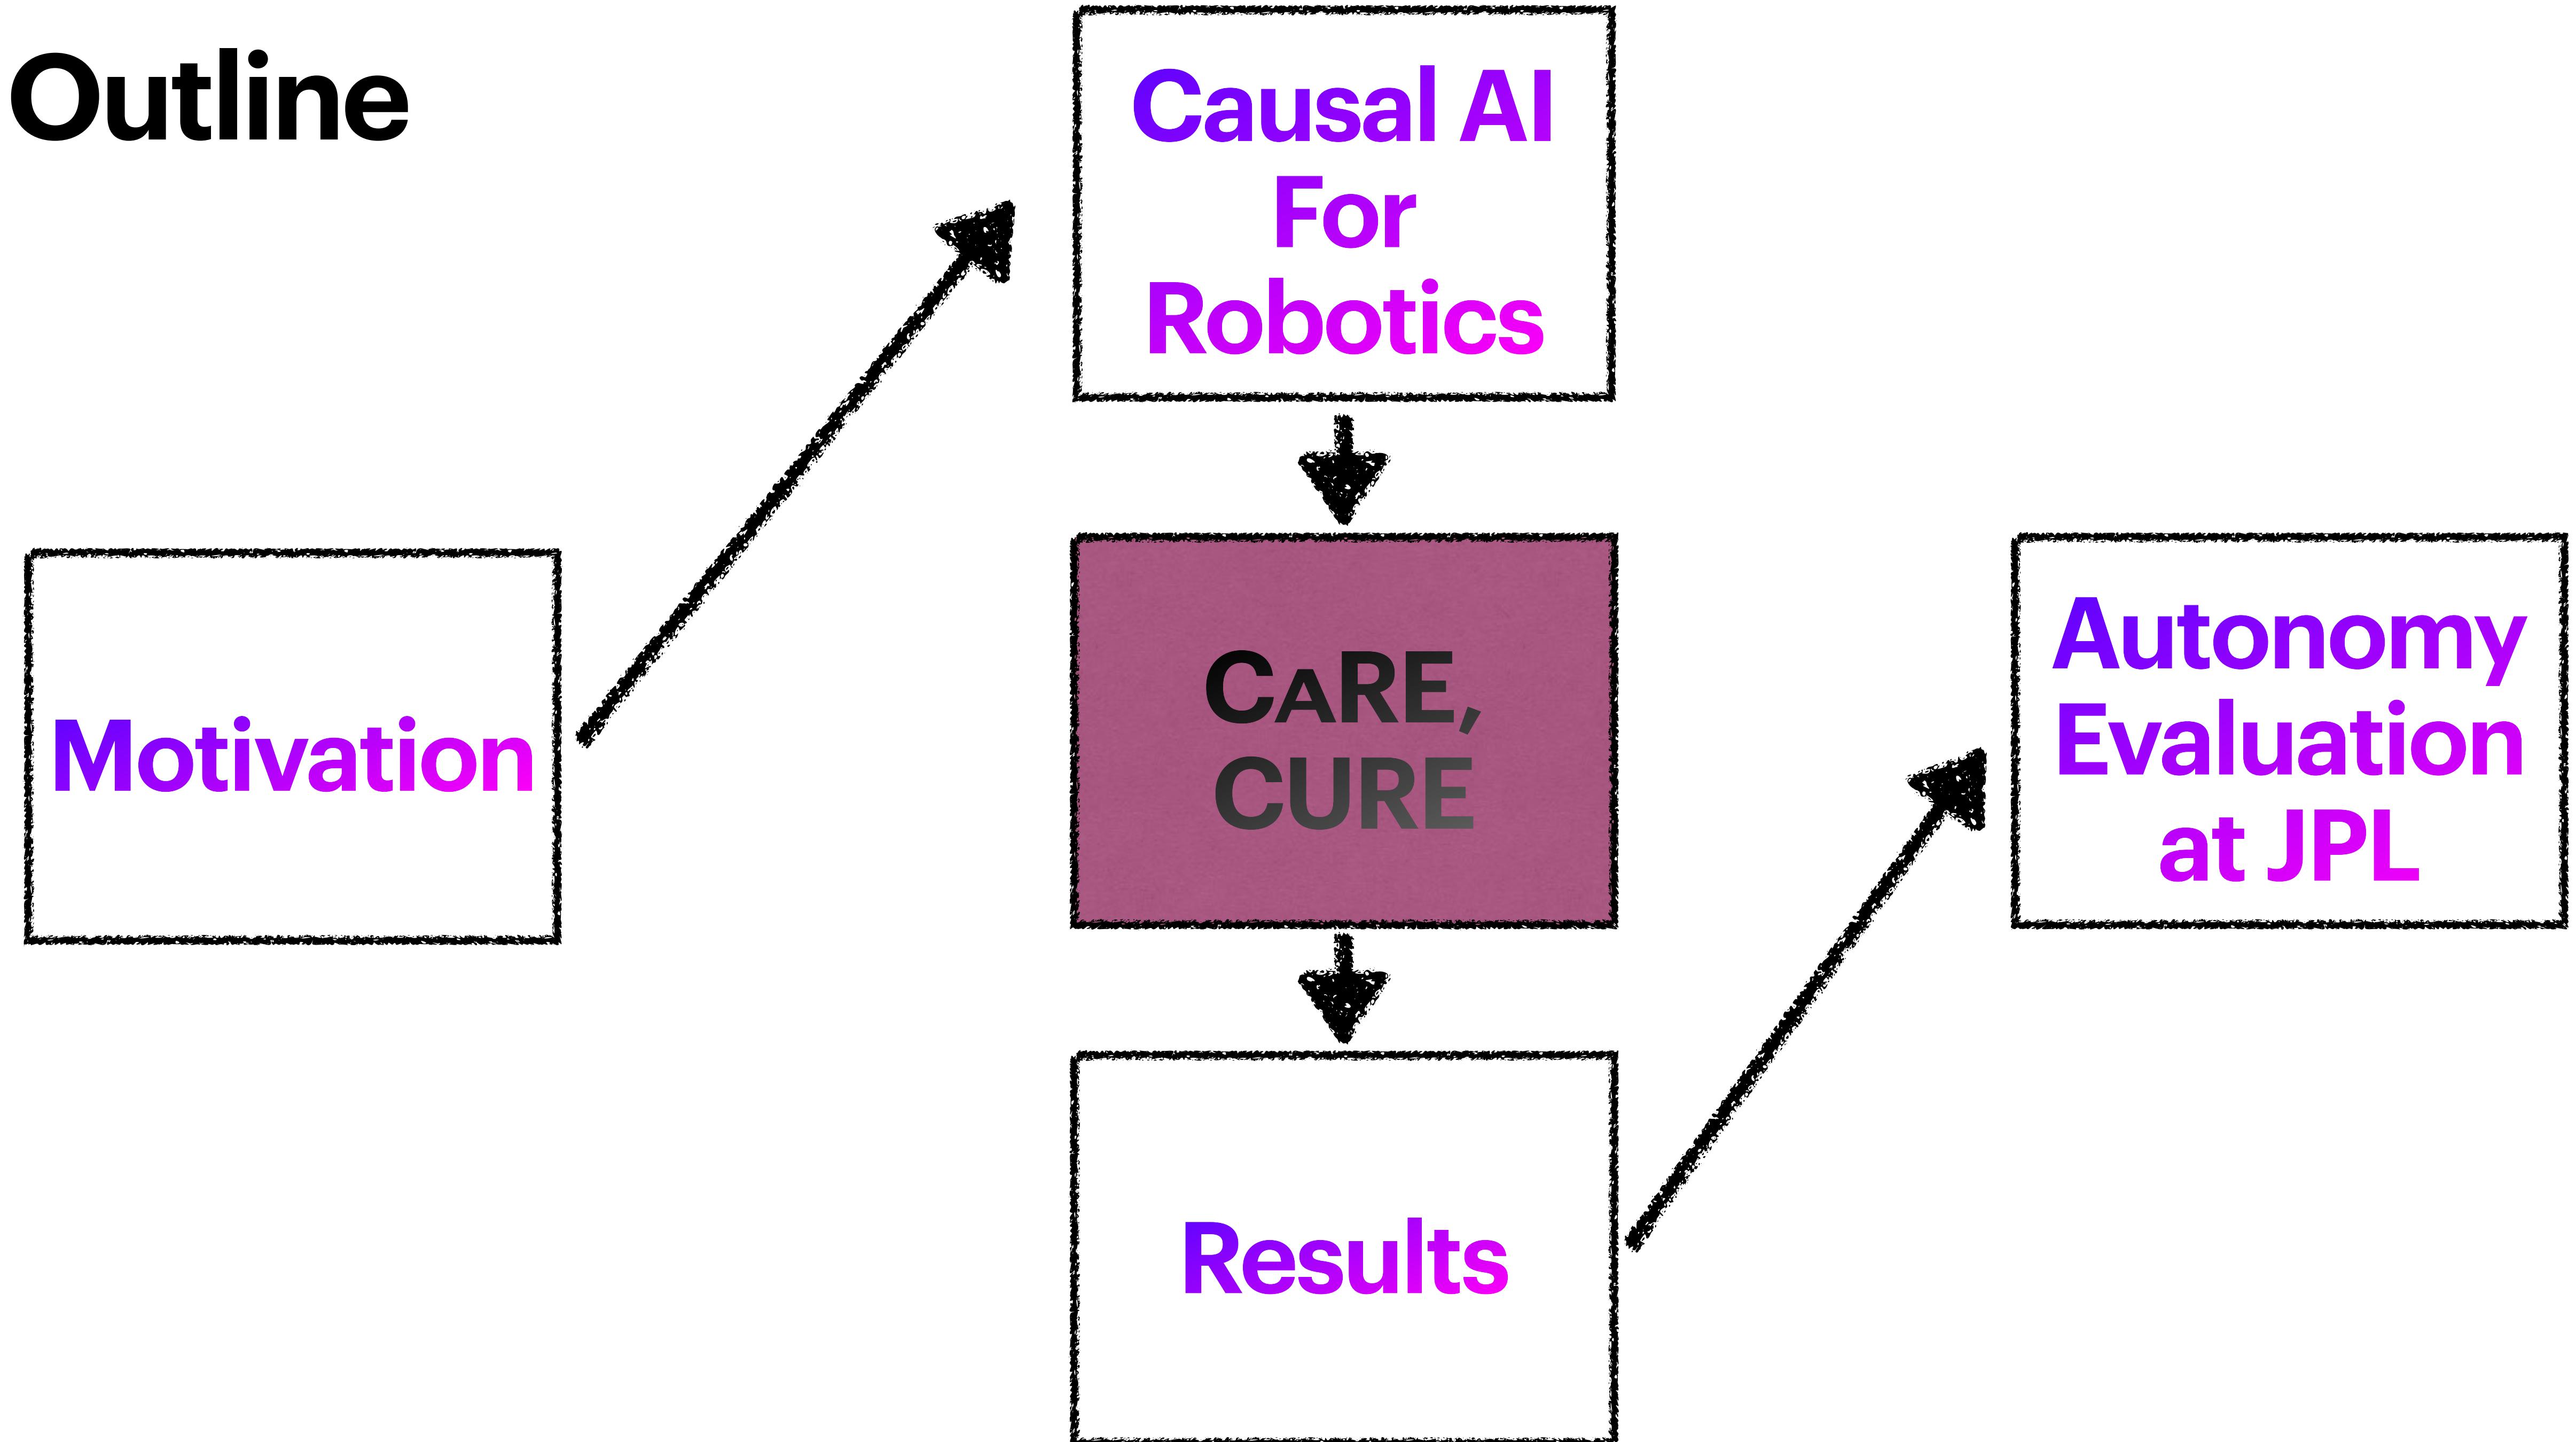
$$\{v_i^*\} = \underset{v_i}{\operatorname{argmax}} ACE(v_i, v_j^*)$$

where $\{v_i^*\} \subset \mathcal{X}$ is the set of root causes, $\{v_j^*\} \subset \mathcal{Y}$ are the performance objectives characterizing the functional fault, and ACE represents the average causal effect.



How to use the causal model for diagnosing misconfigurations?

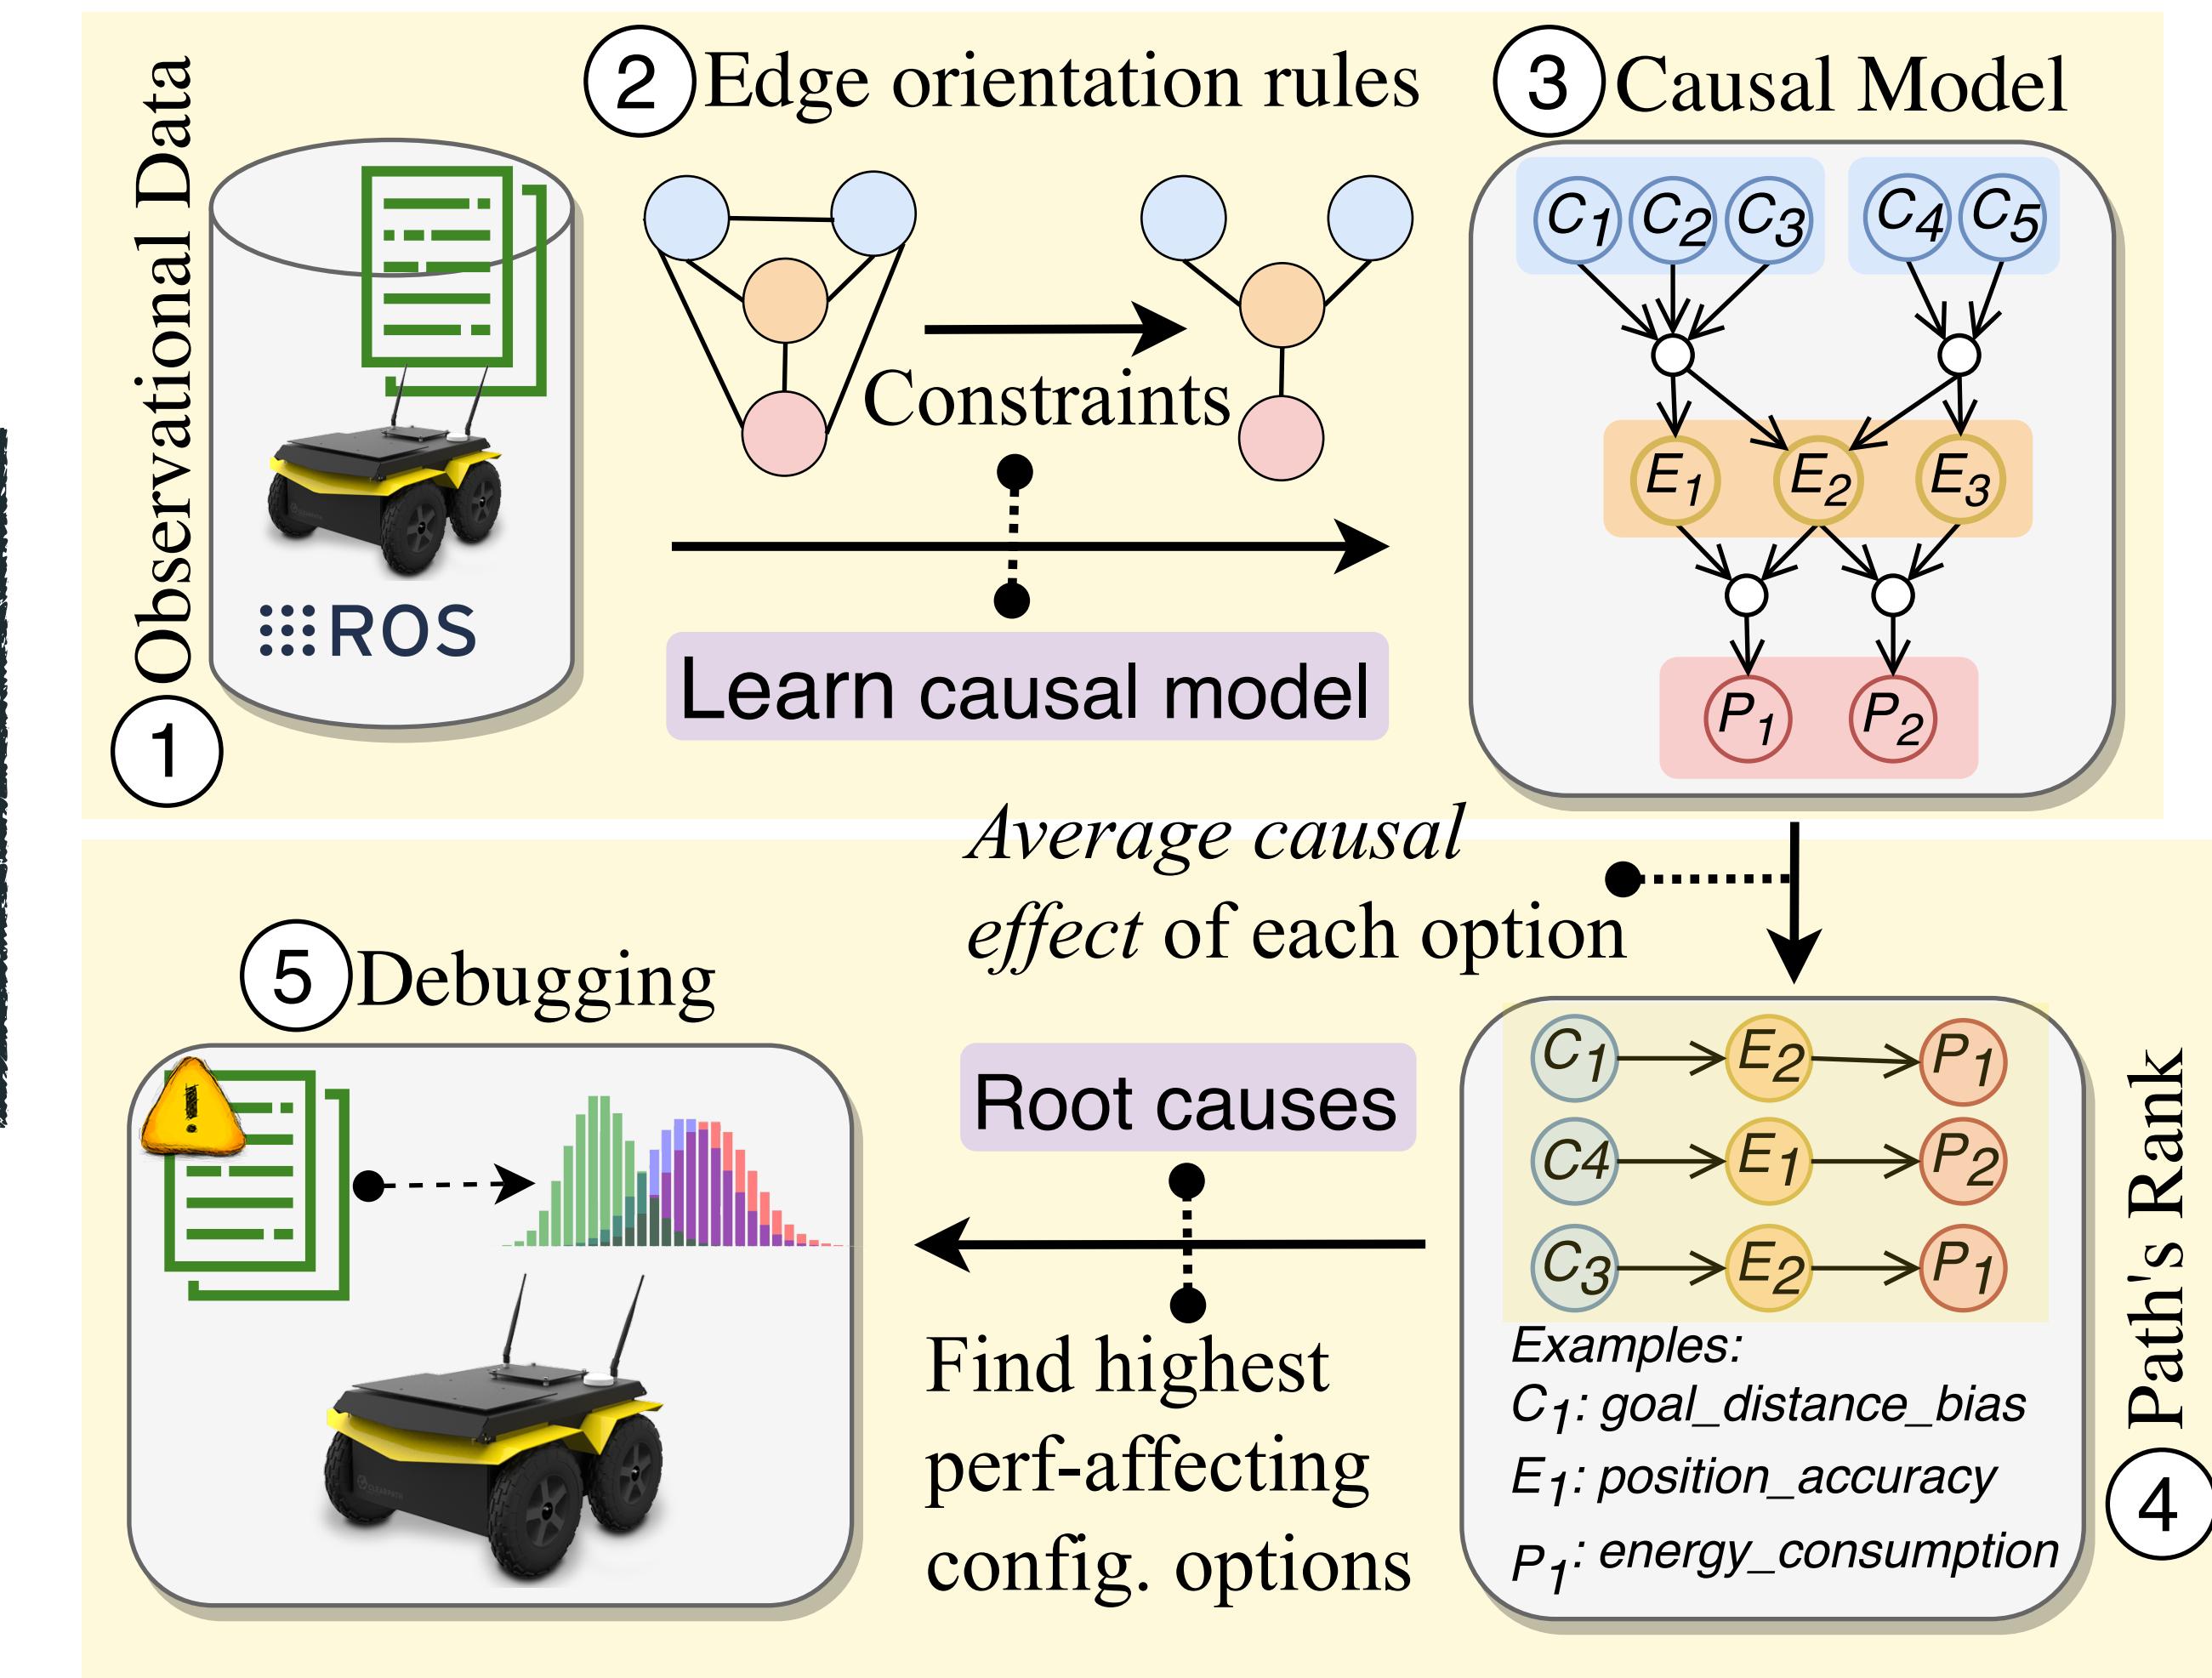
Outline



CaRE: End-to-end Pipeline

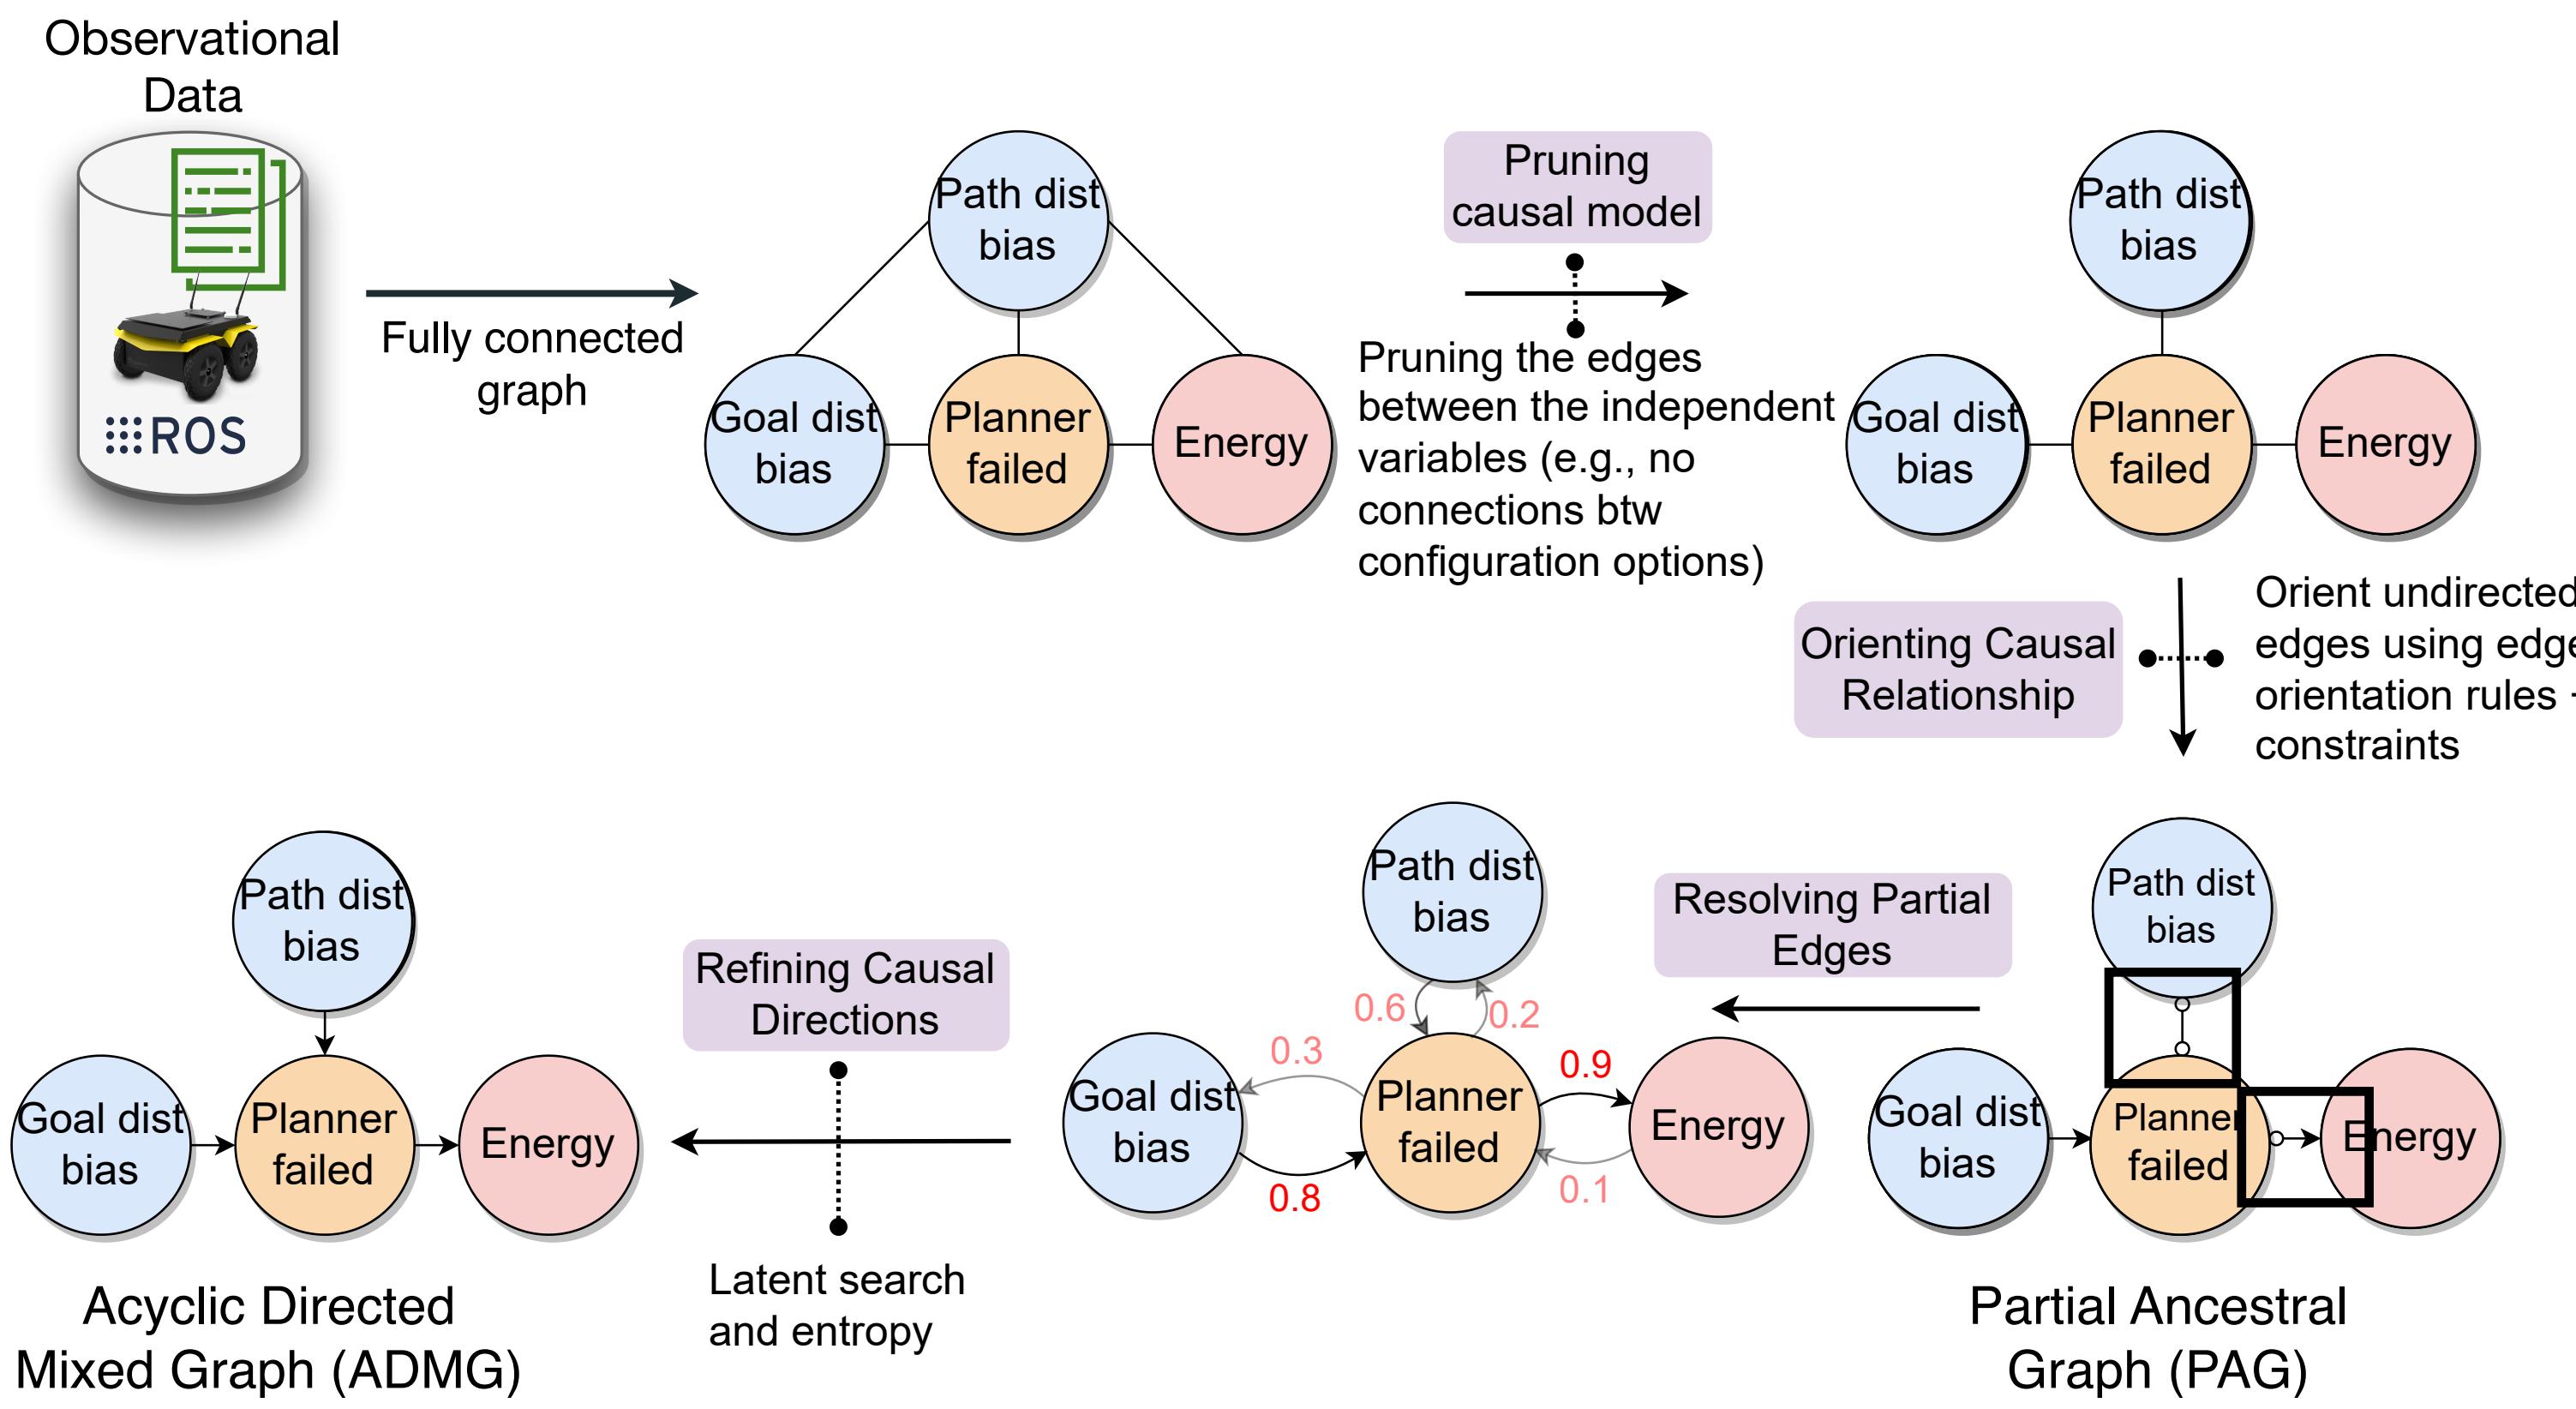
Stages

1. Learn causal model from observational data
2. Identify the causal paths
3. Average causal effect estimation



Learning and Estimating Causal Effects

Learning the Causal Model



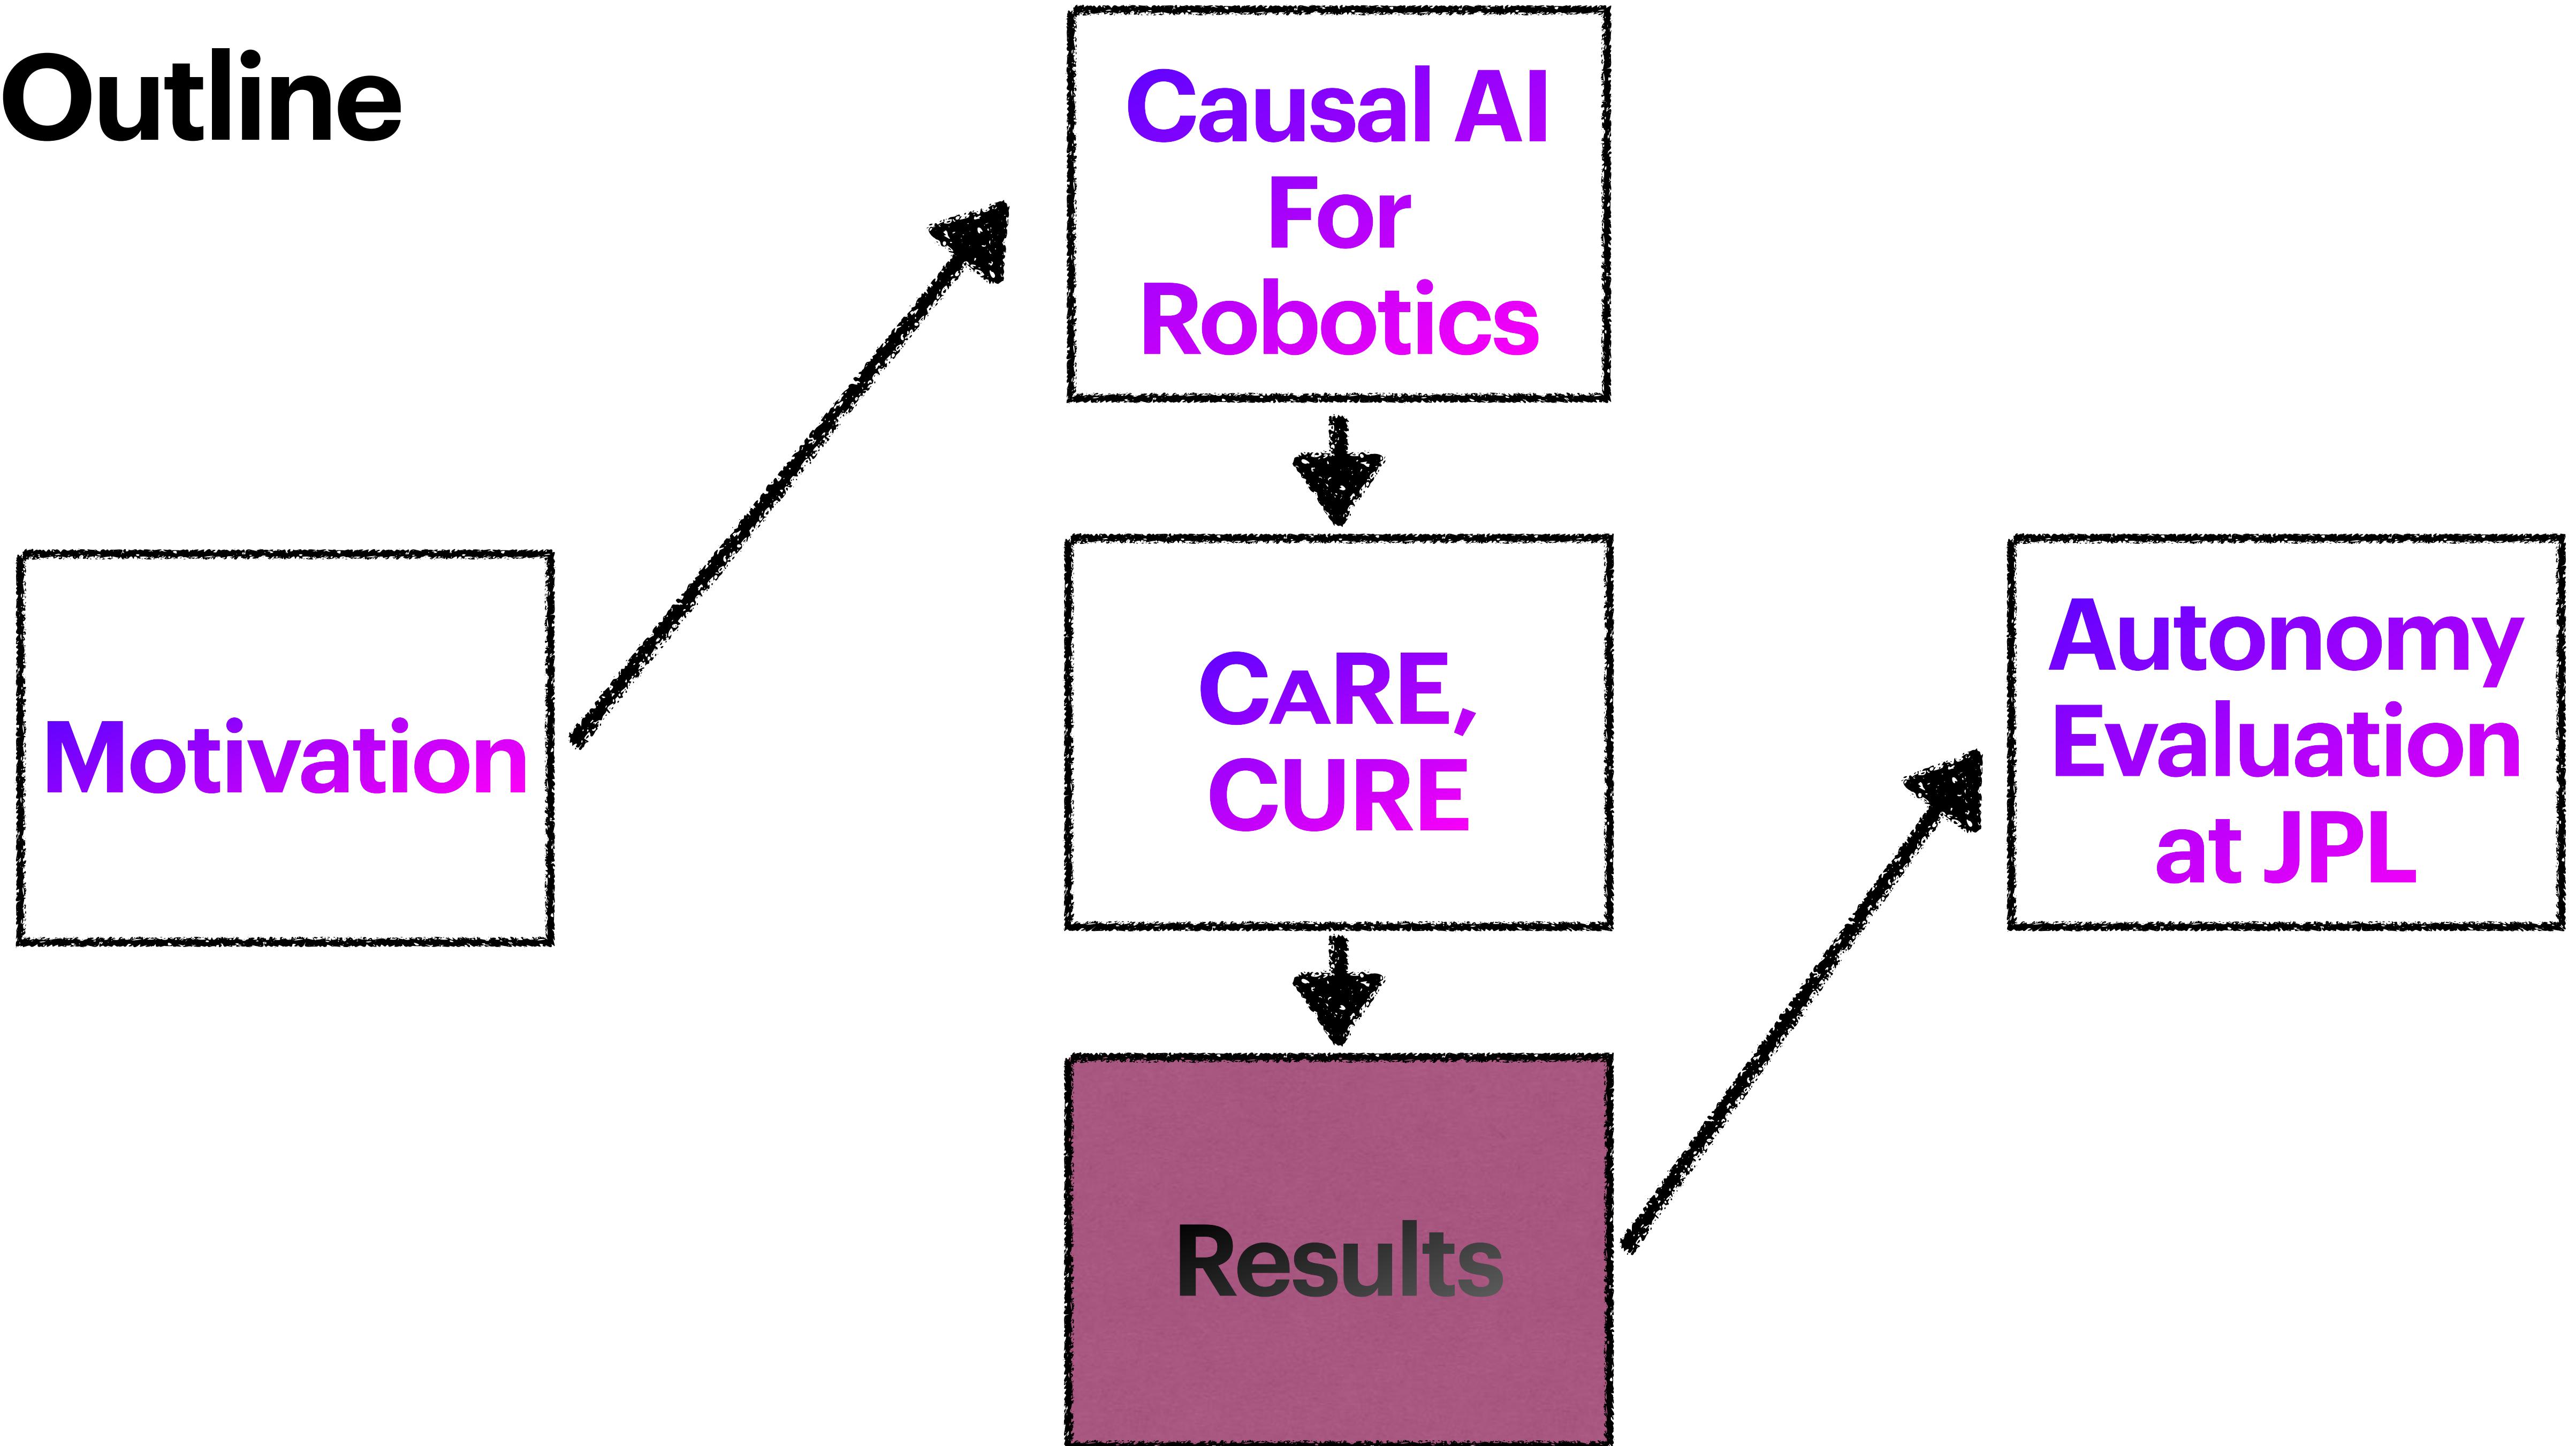
Causal Effect Estimation

- Use **do-calculus** to measure average causal effect

Estimate the probability of satisfying energy given $\text{path_dist_bias}=20.0$

$$P(\text{Energy} < 25\text{Wh} | \text{do}(\text{Path_dist_bias} = 20.0))$$

Outline



Research Questions

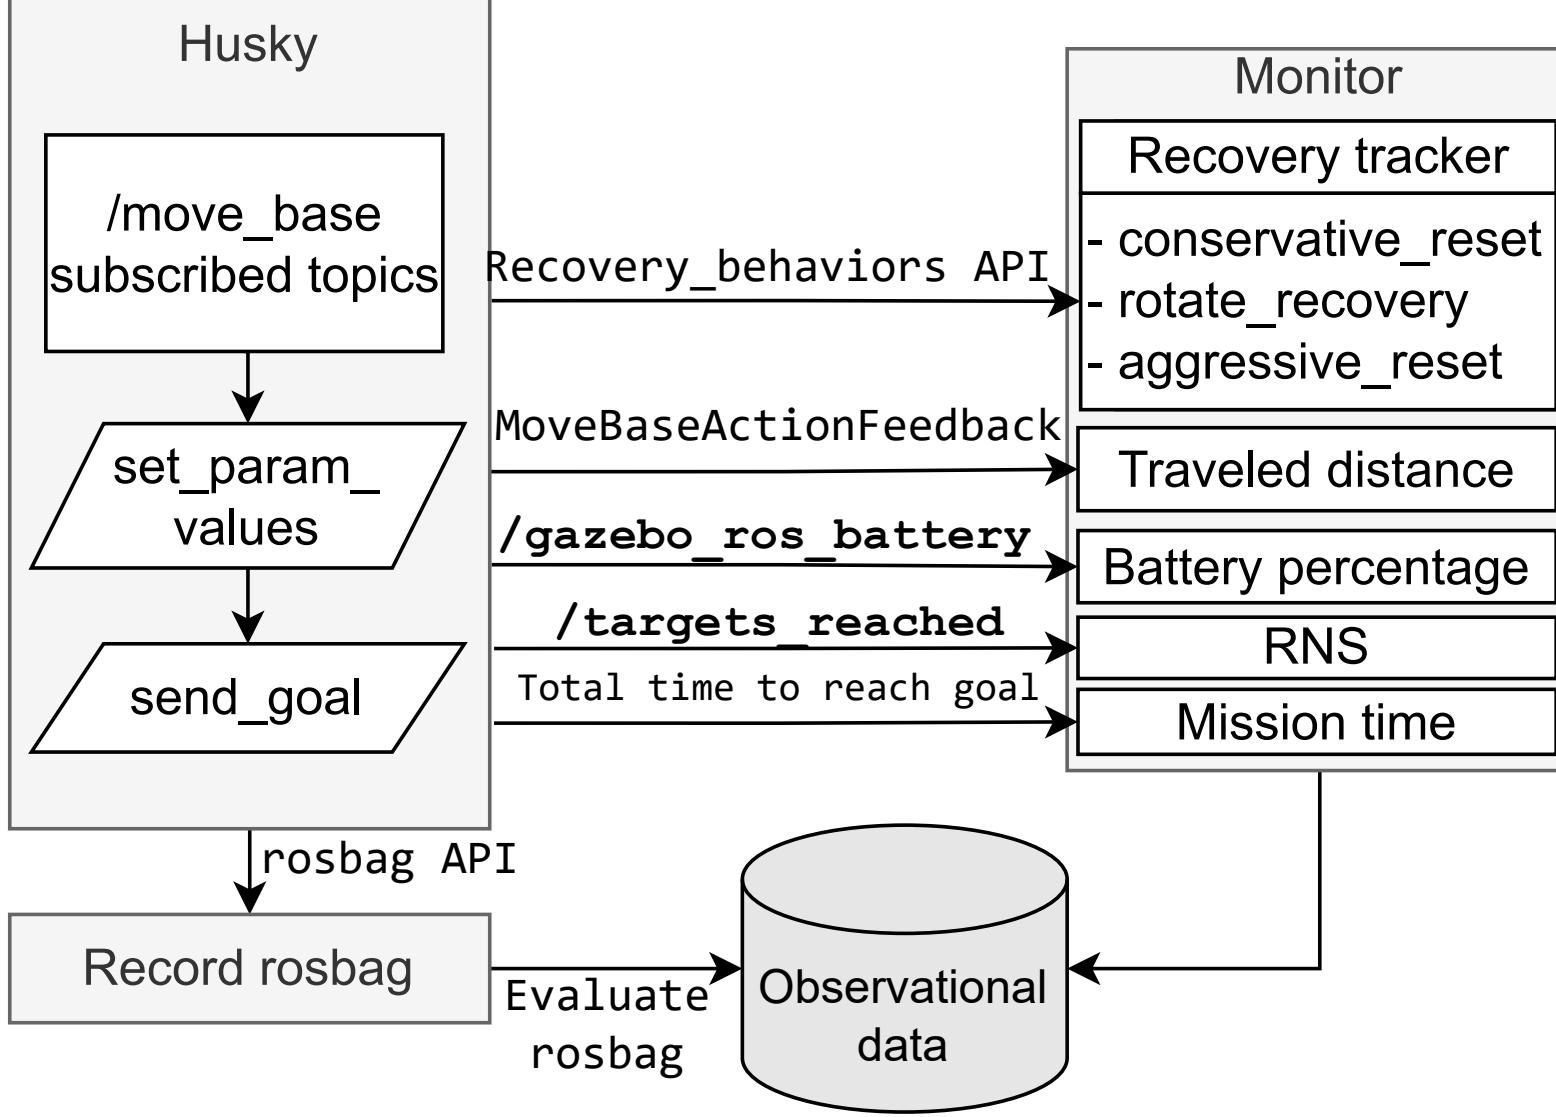
- **Research Question 1 (Accuracy).** To what extent are the root causes determined by CaRE the true root causes of the observed functional faults?
- **Research Question 2 (Transferability).** To what extent can CaRE accurately detect misconfigurations when deployed in a different platform?

Results

Experimental Setup

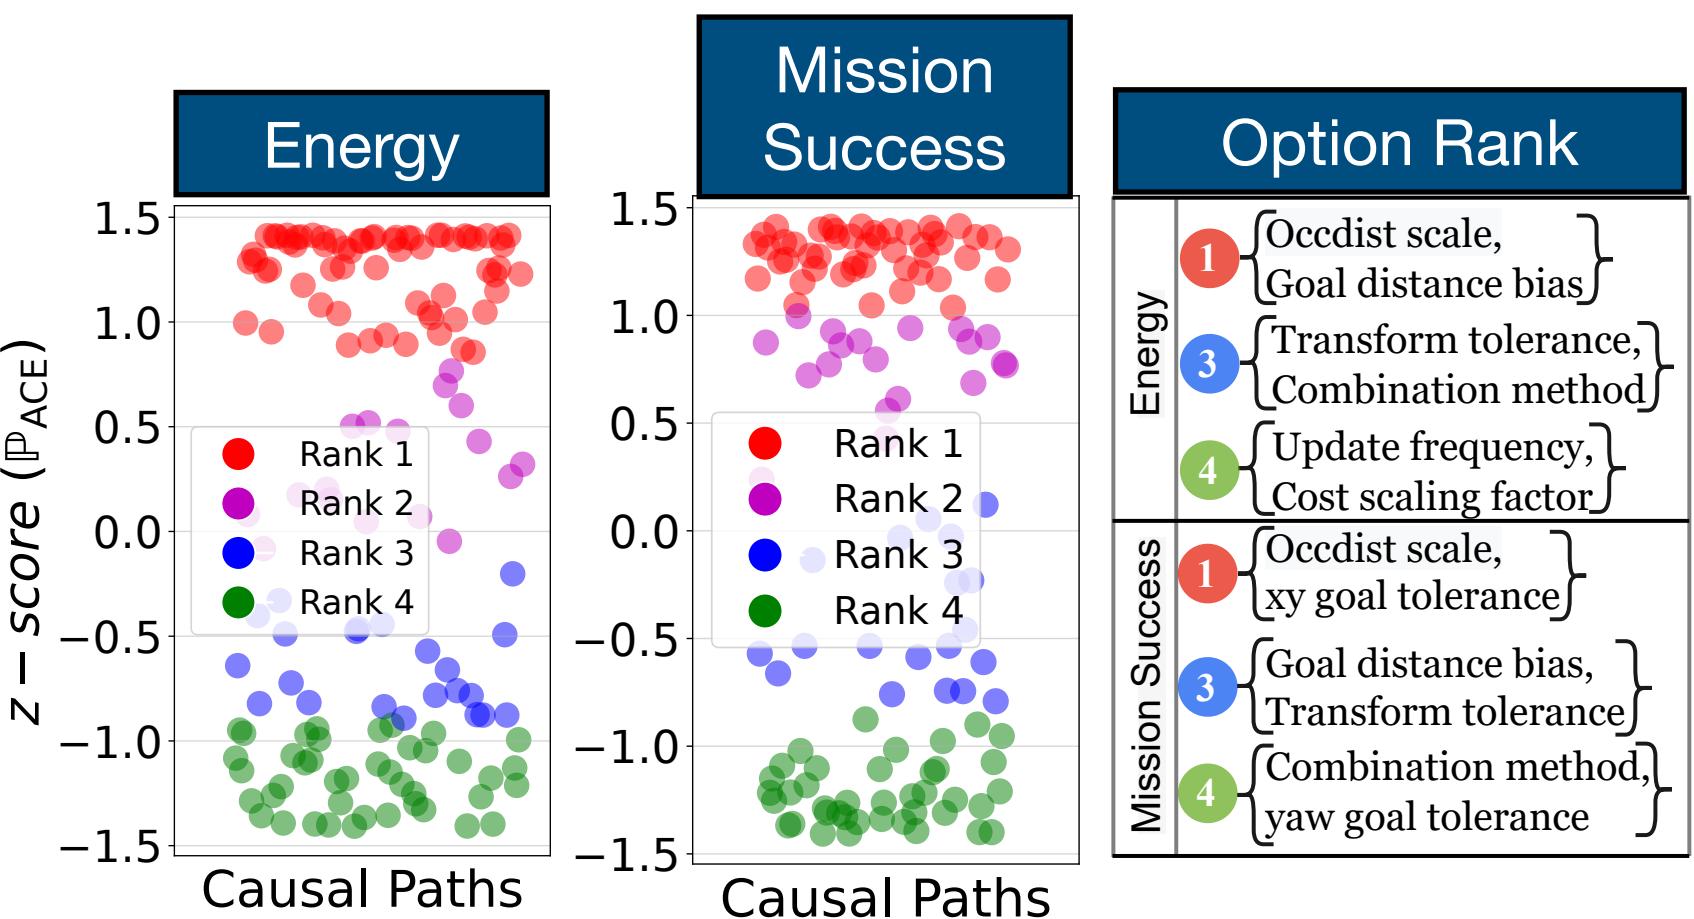
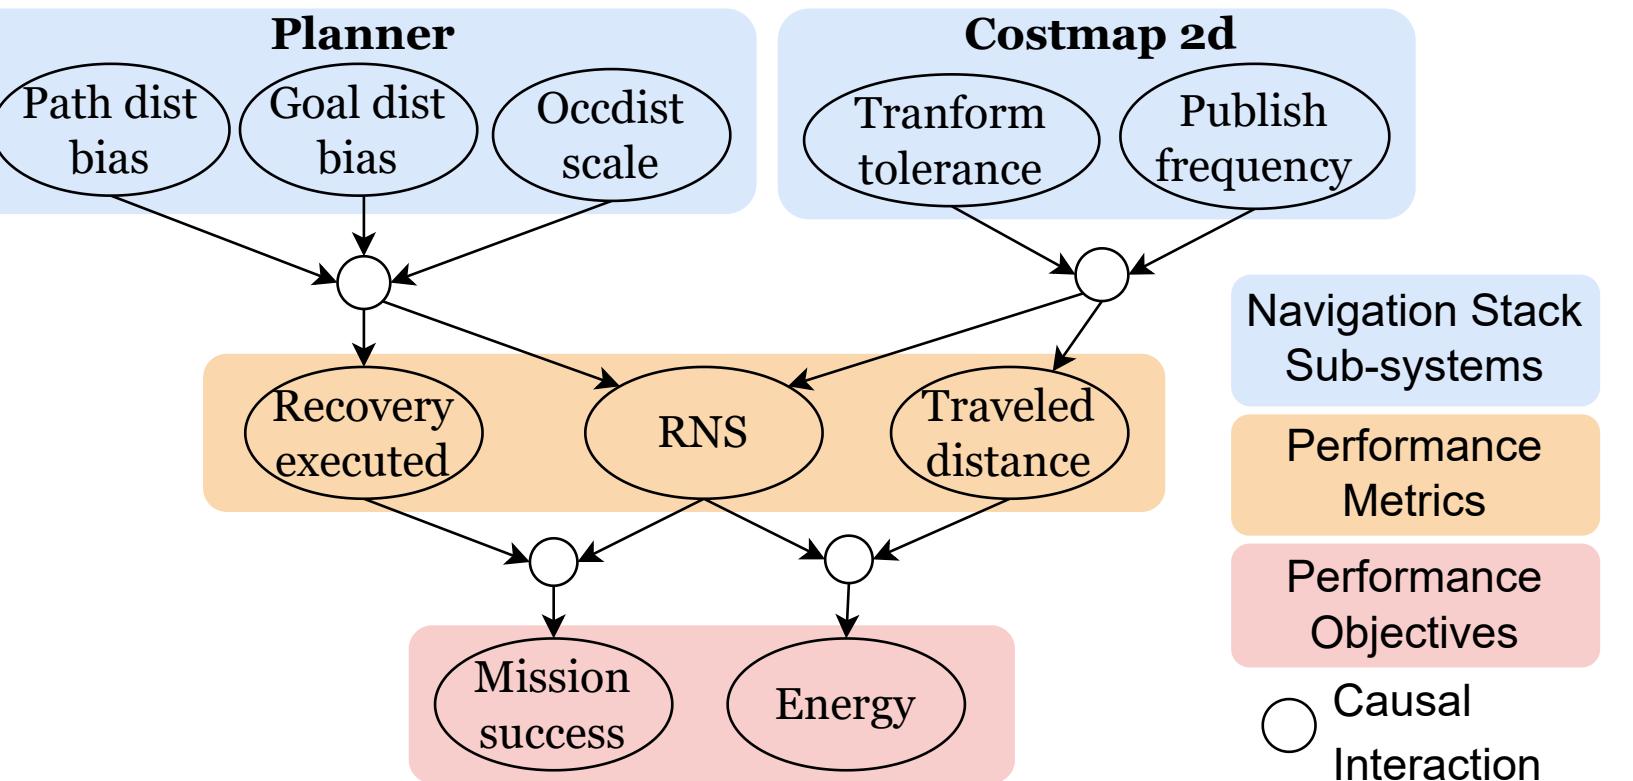


Observational Data Collection

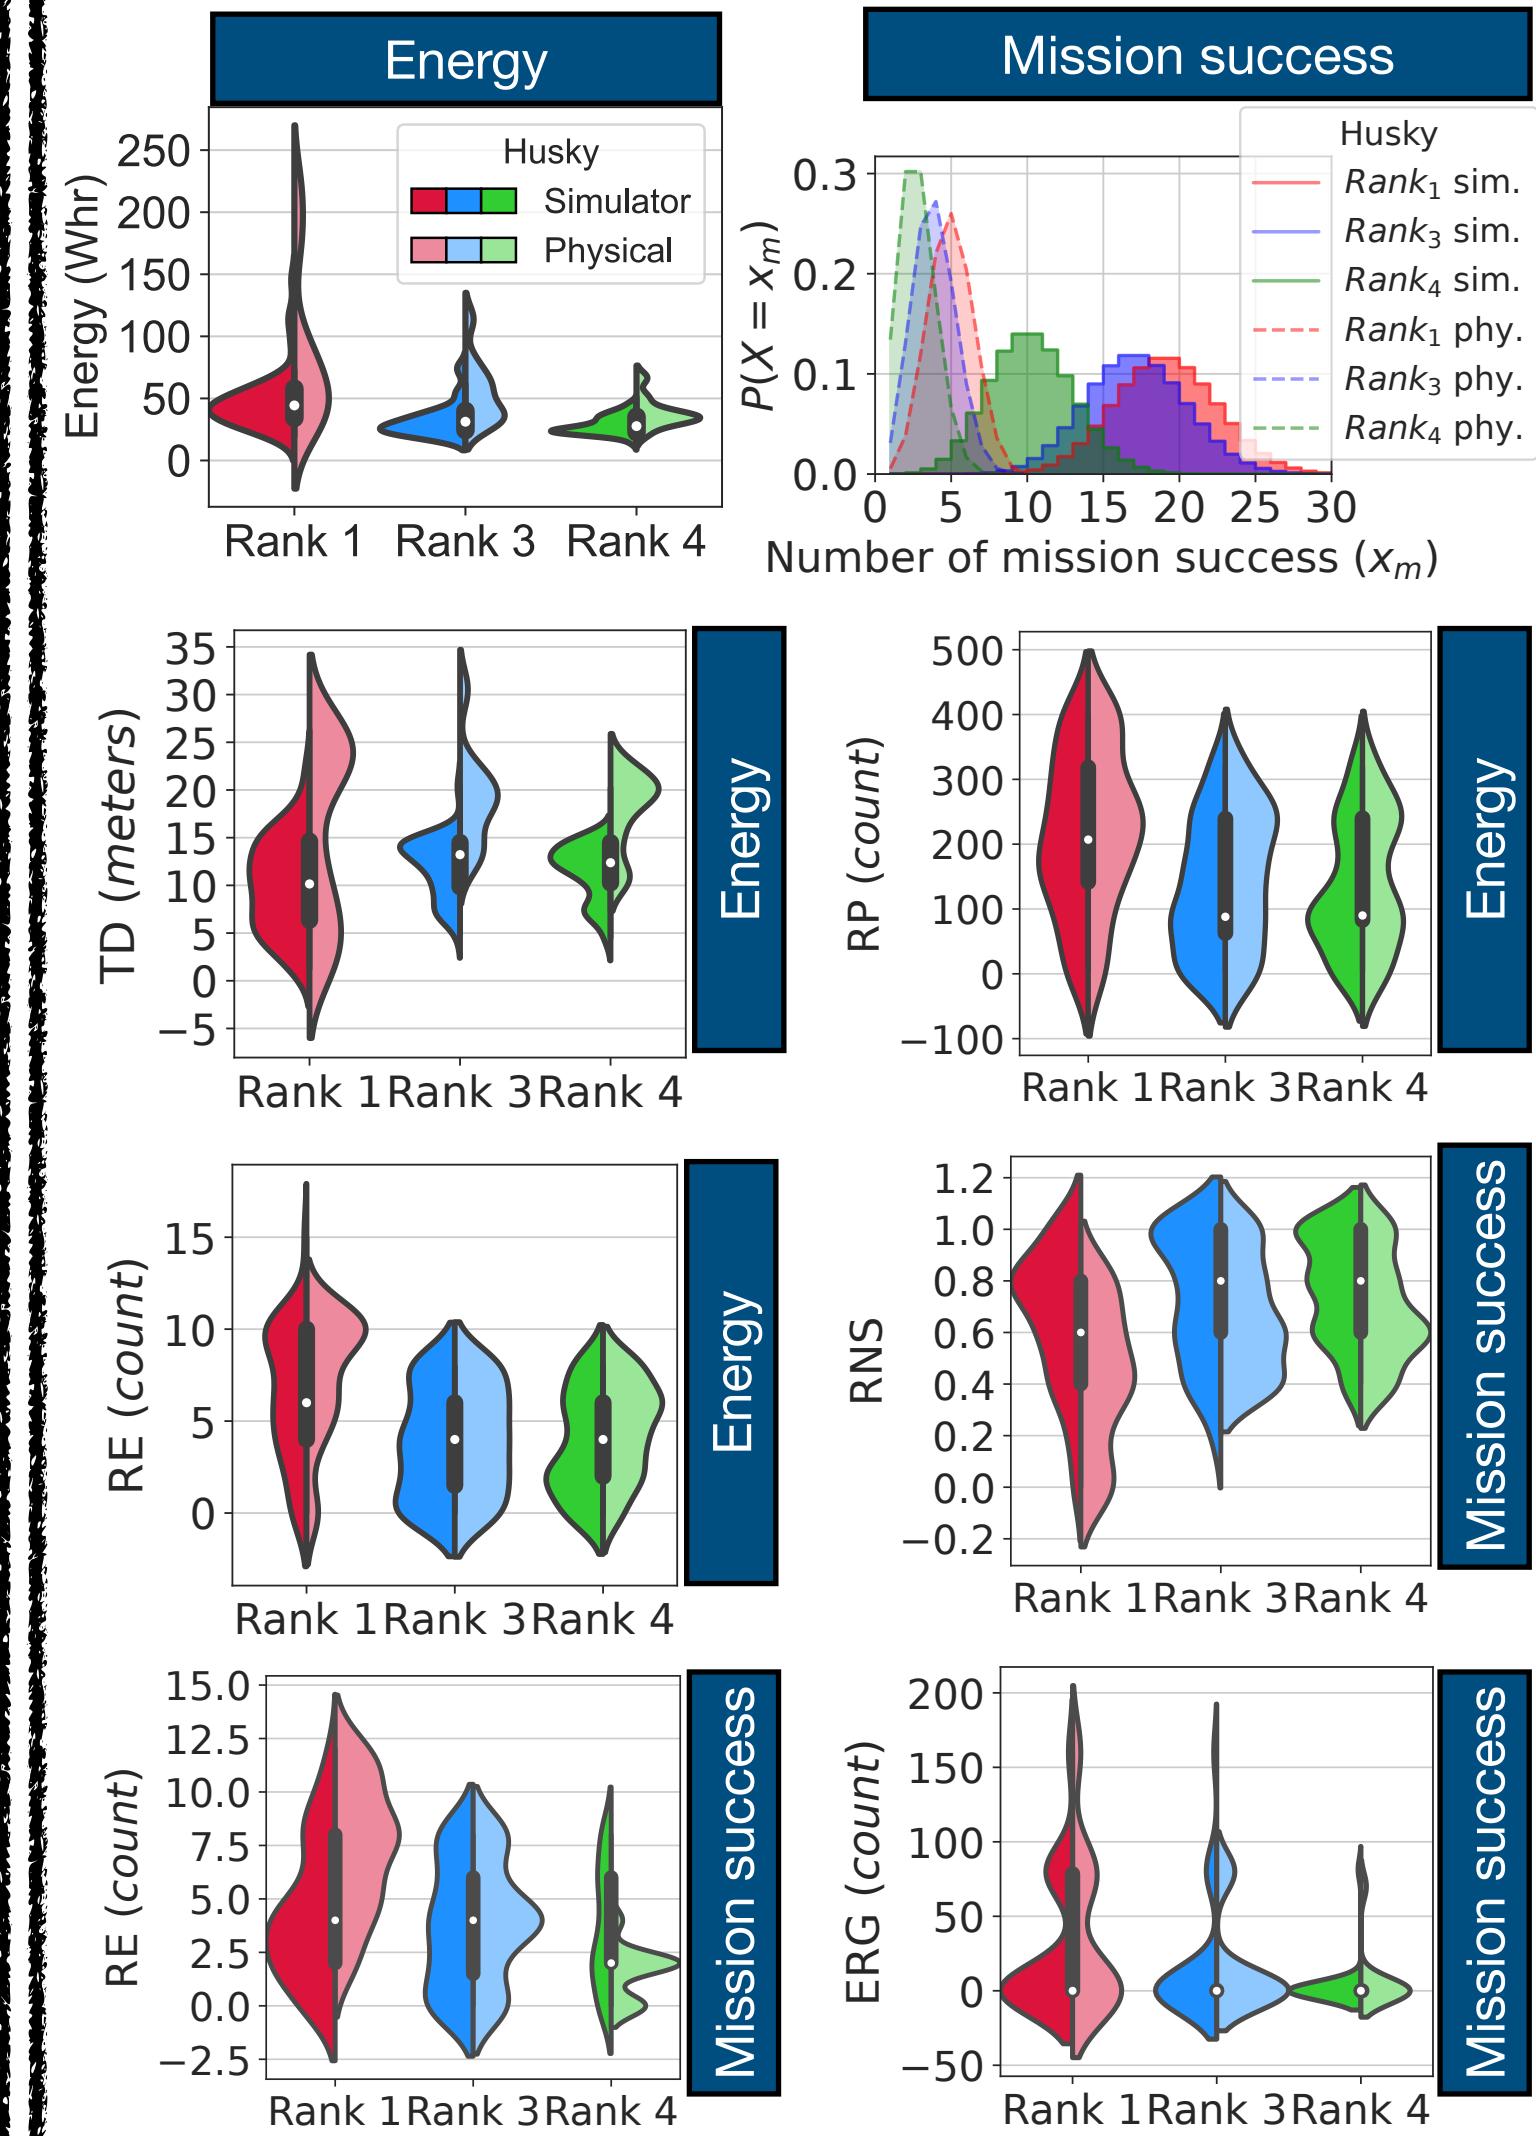


Applying CaRE

A partial causal model discovered in our experiments using *Husky* in simulation



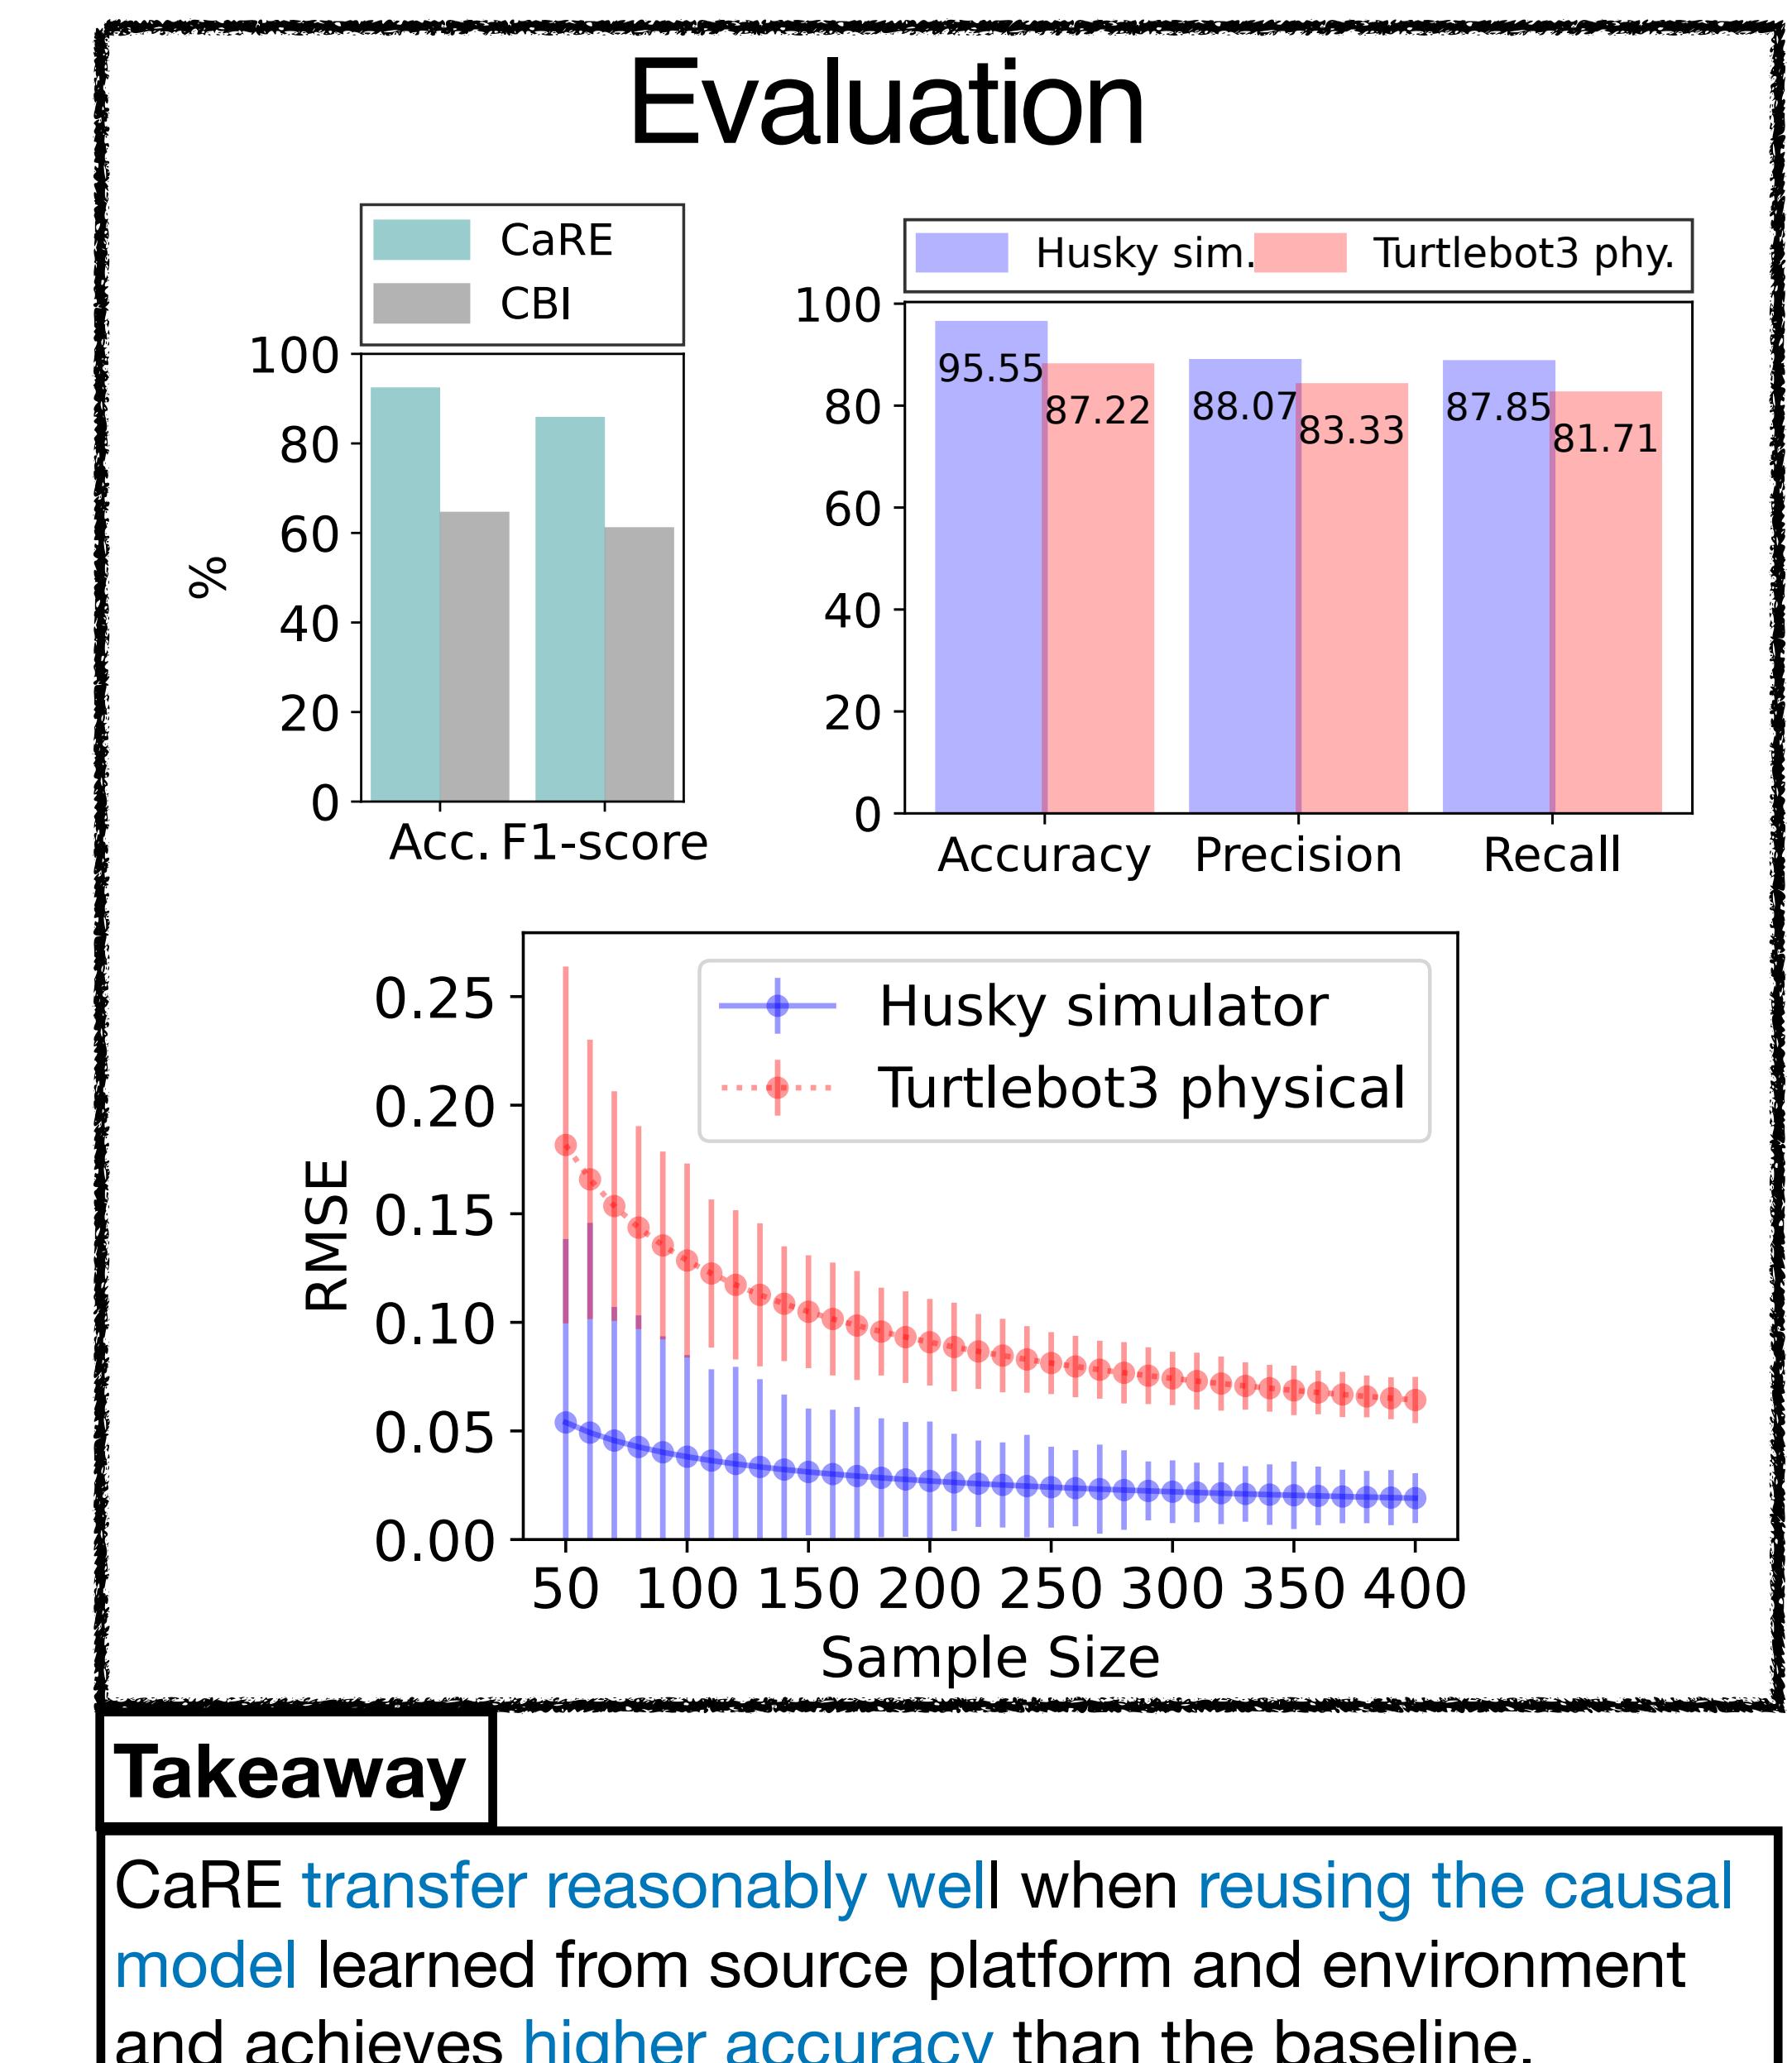
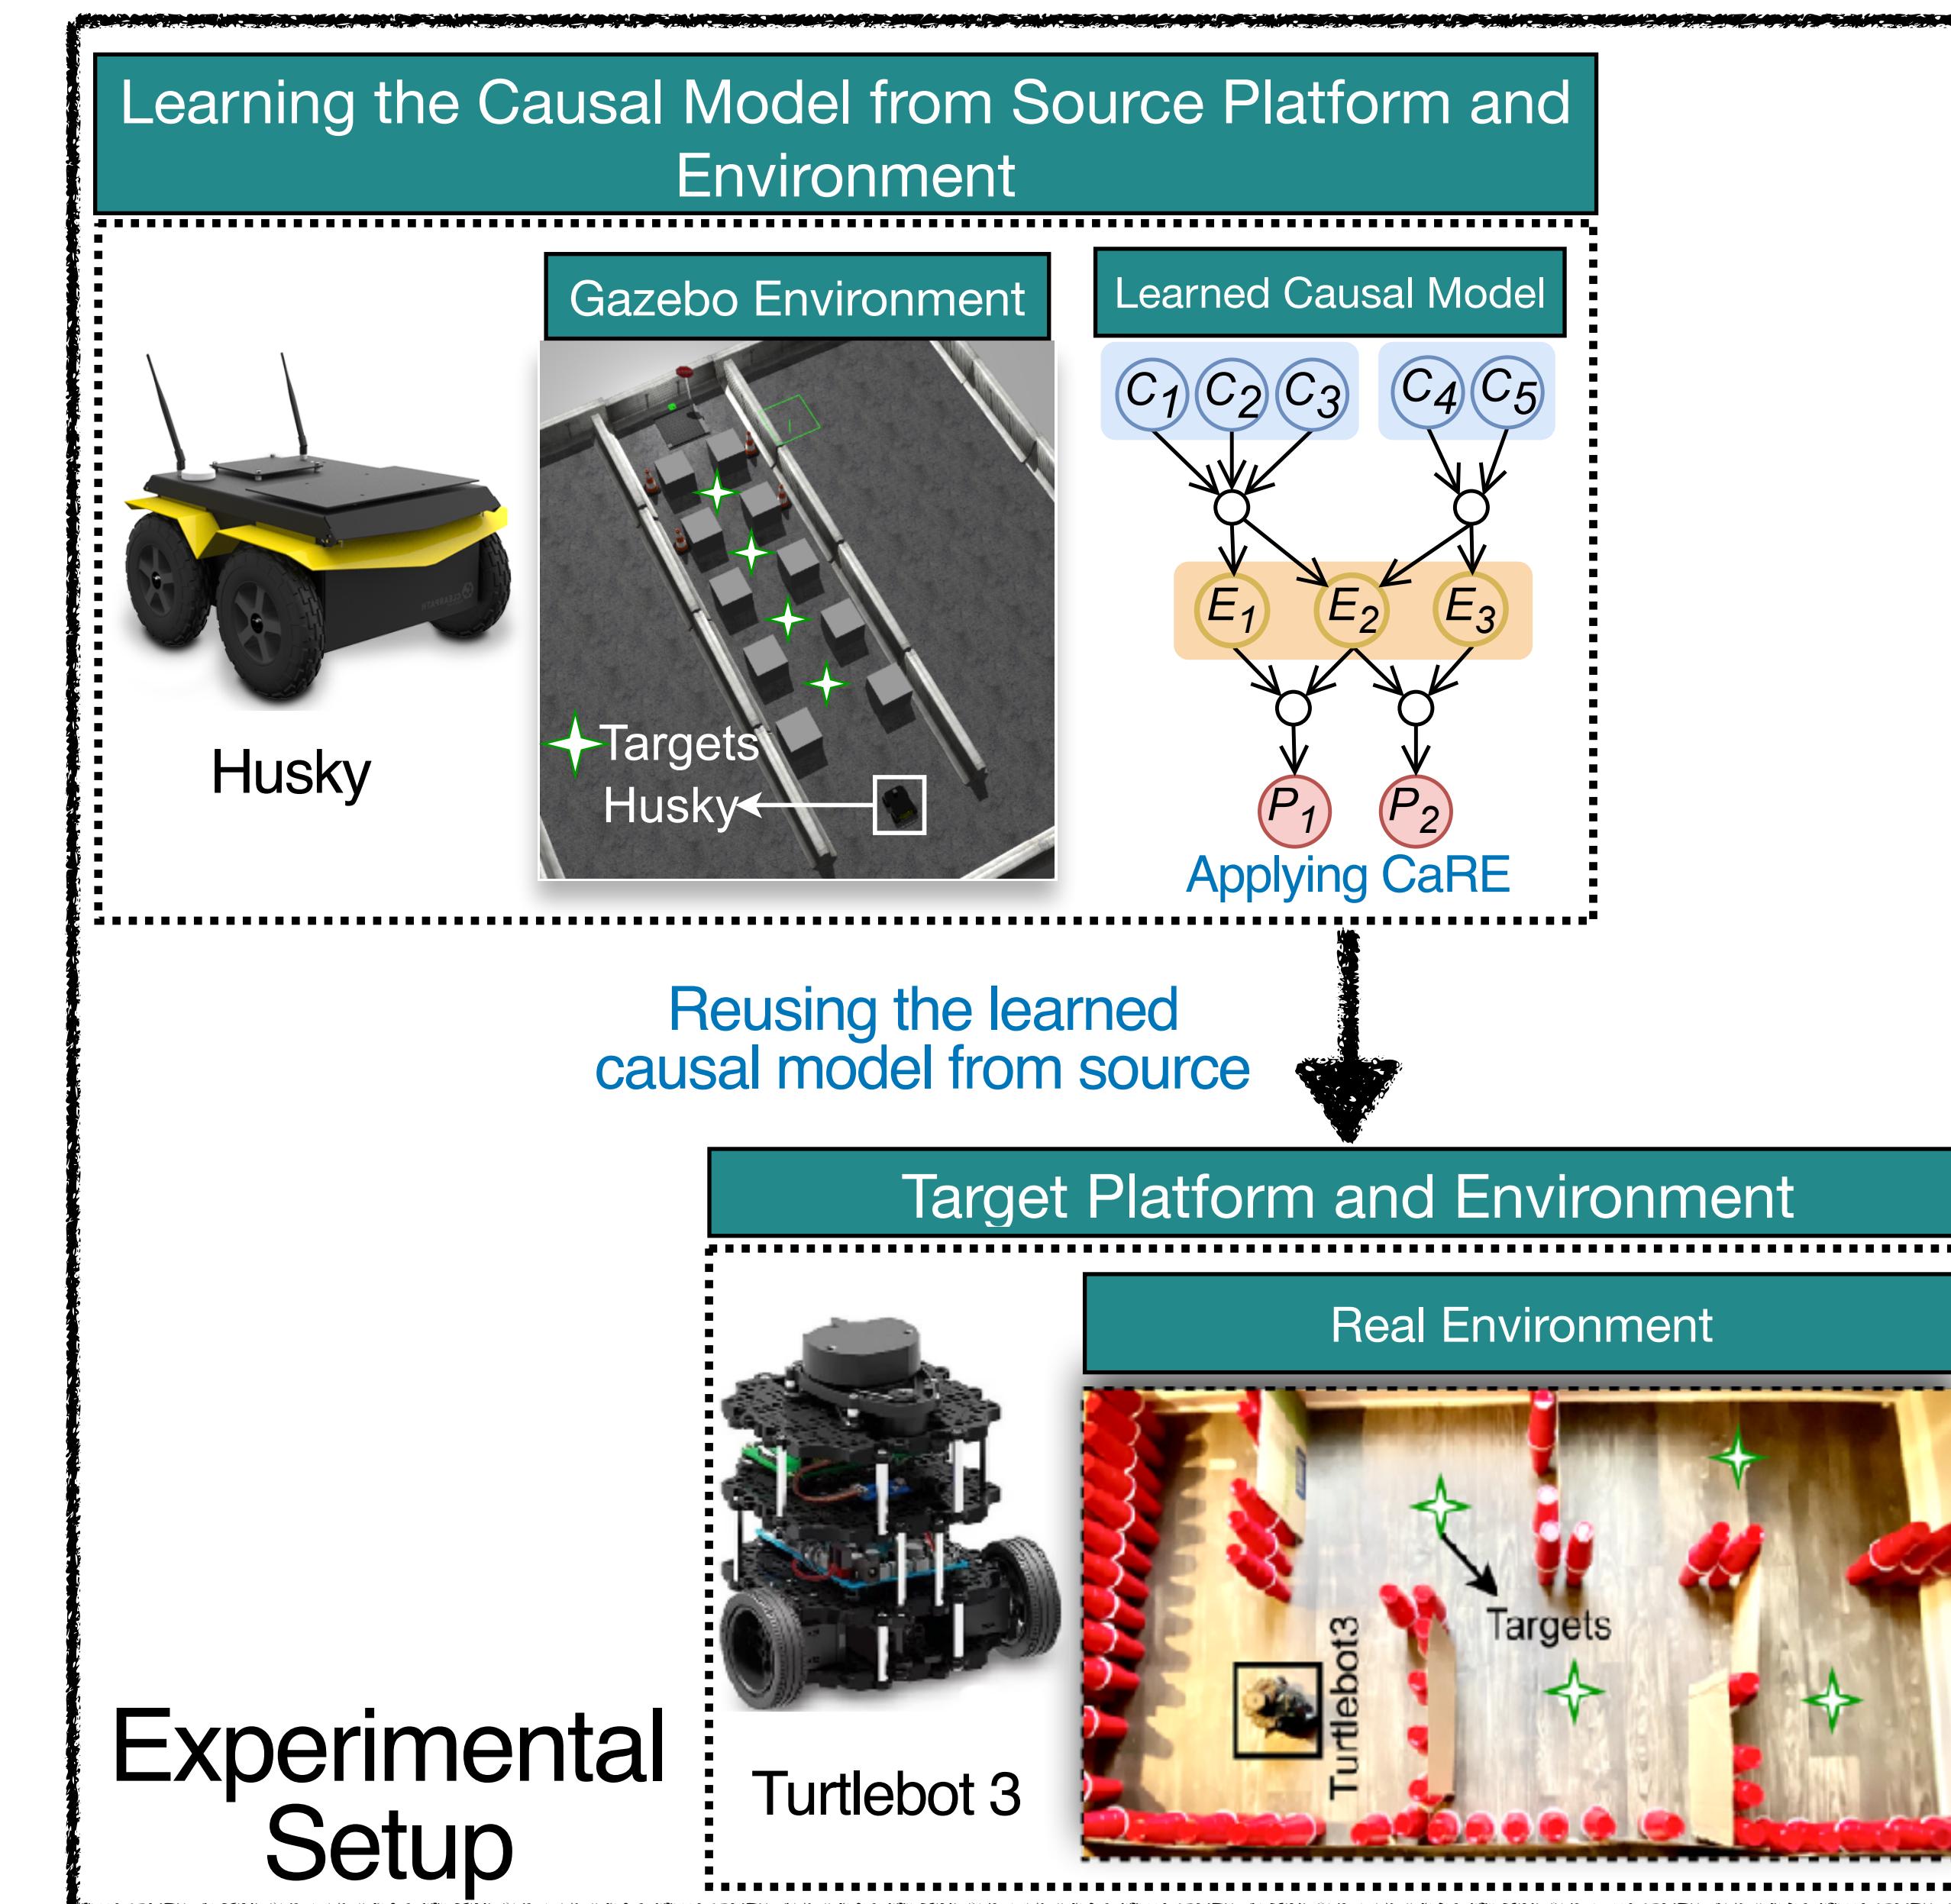
Evaluation



Takeaway

Configuration options which **rank higher** have the **strongest influence** on the performance objectives.

Results

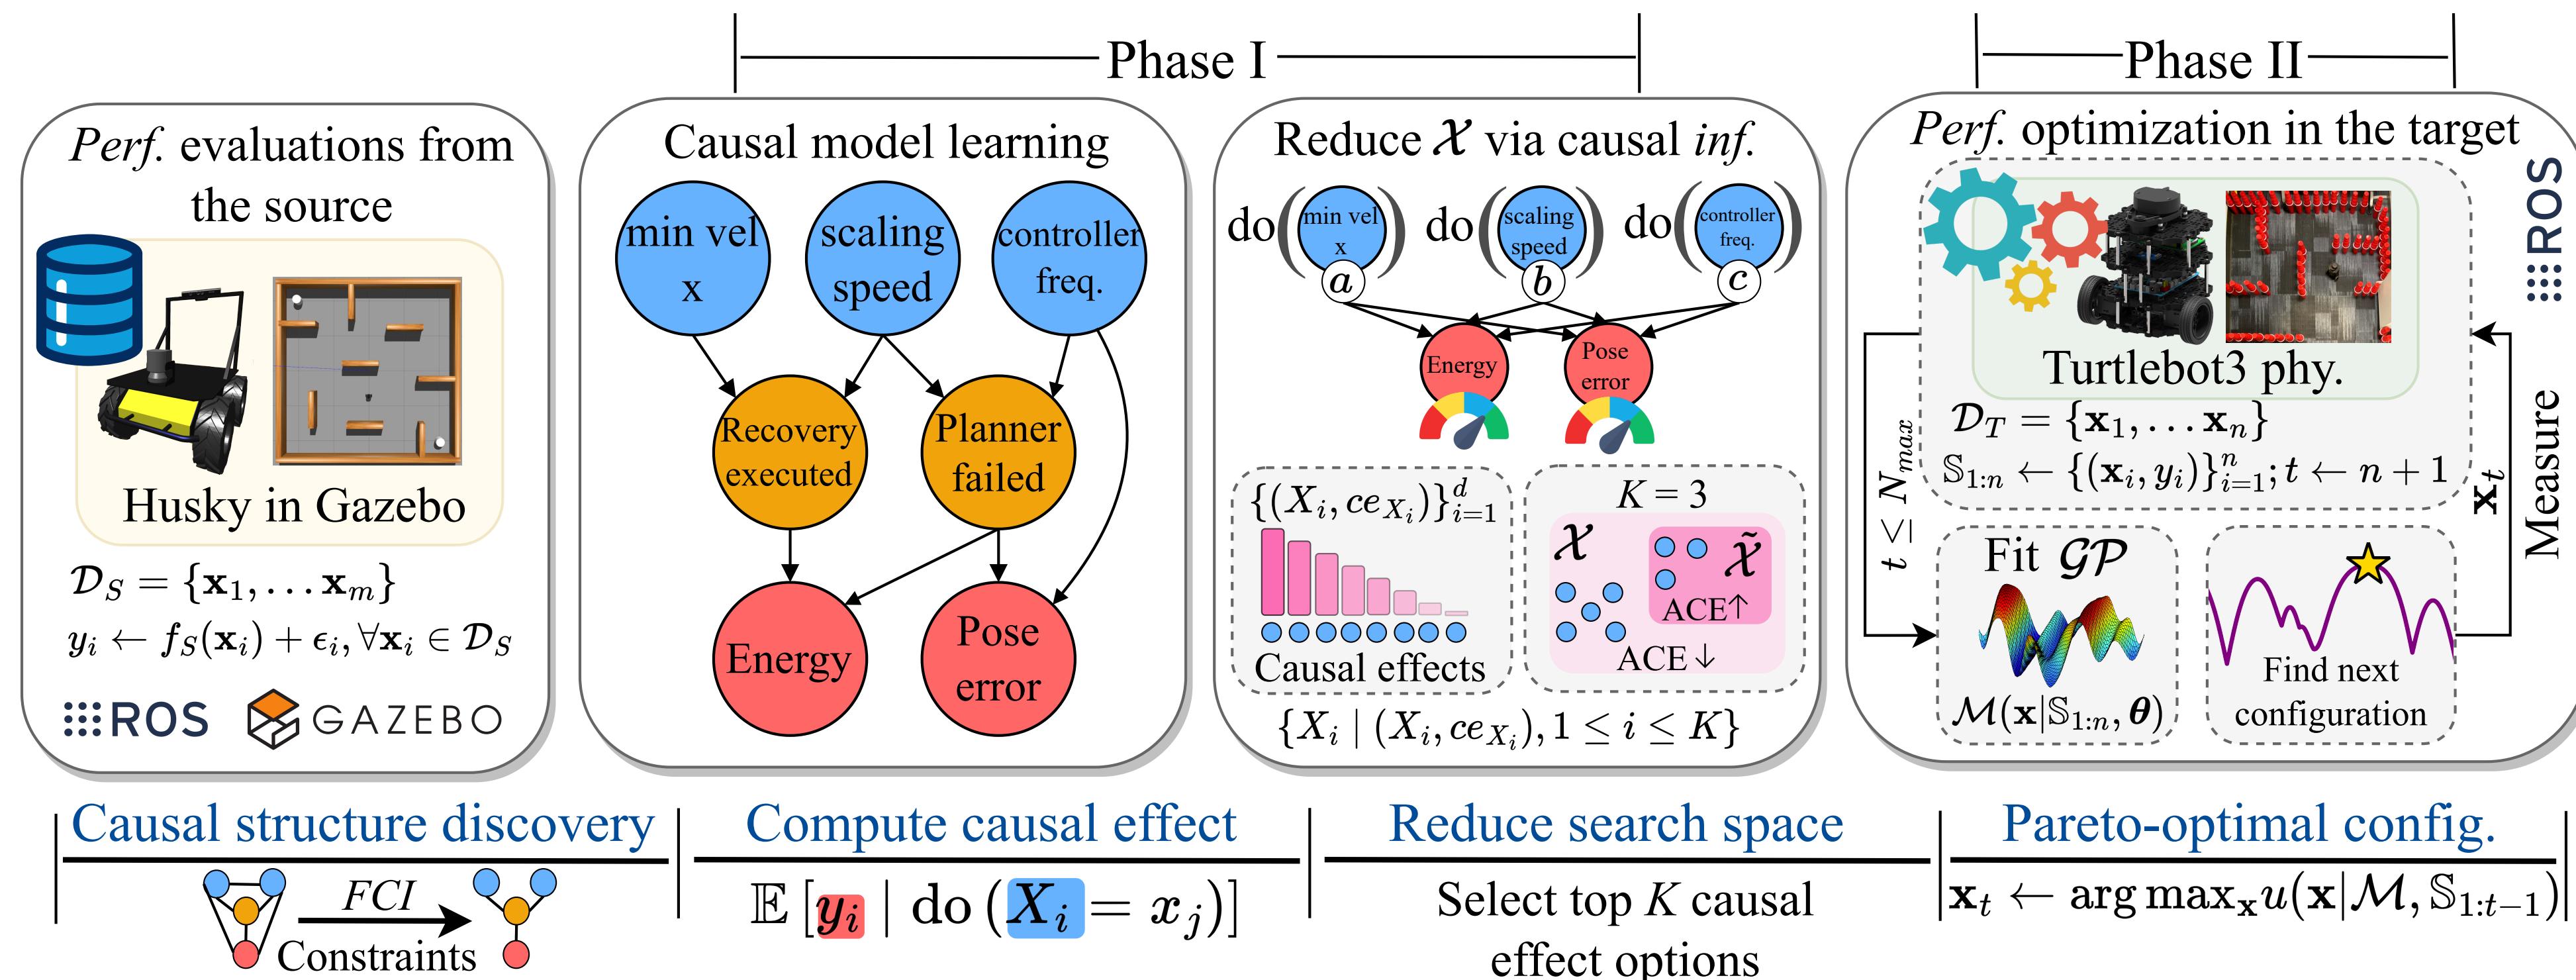


CURE: Simulation Augmented Auto-tuning in Robotics

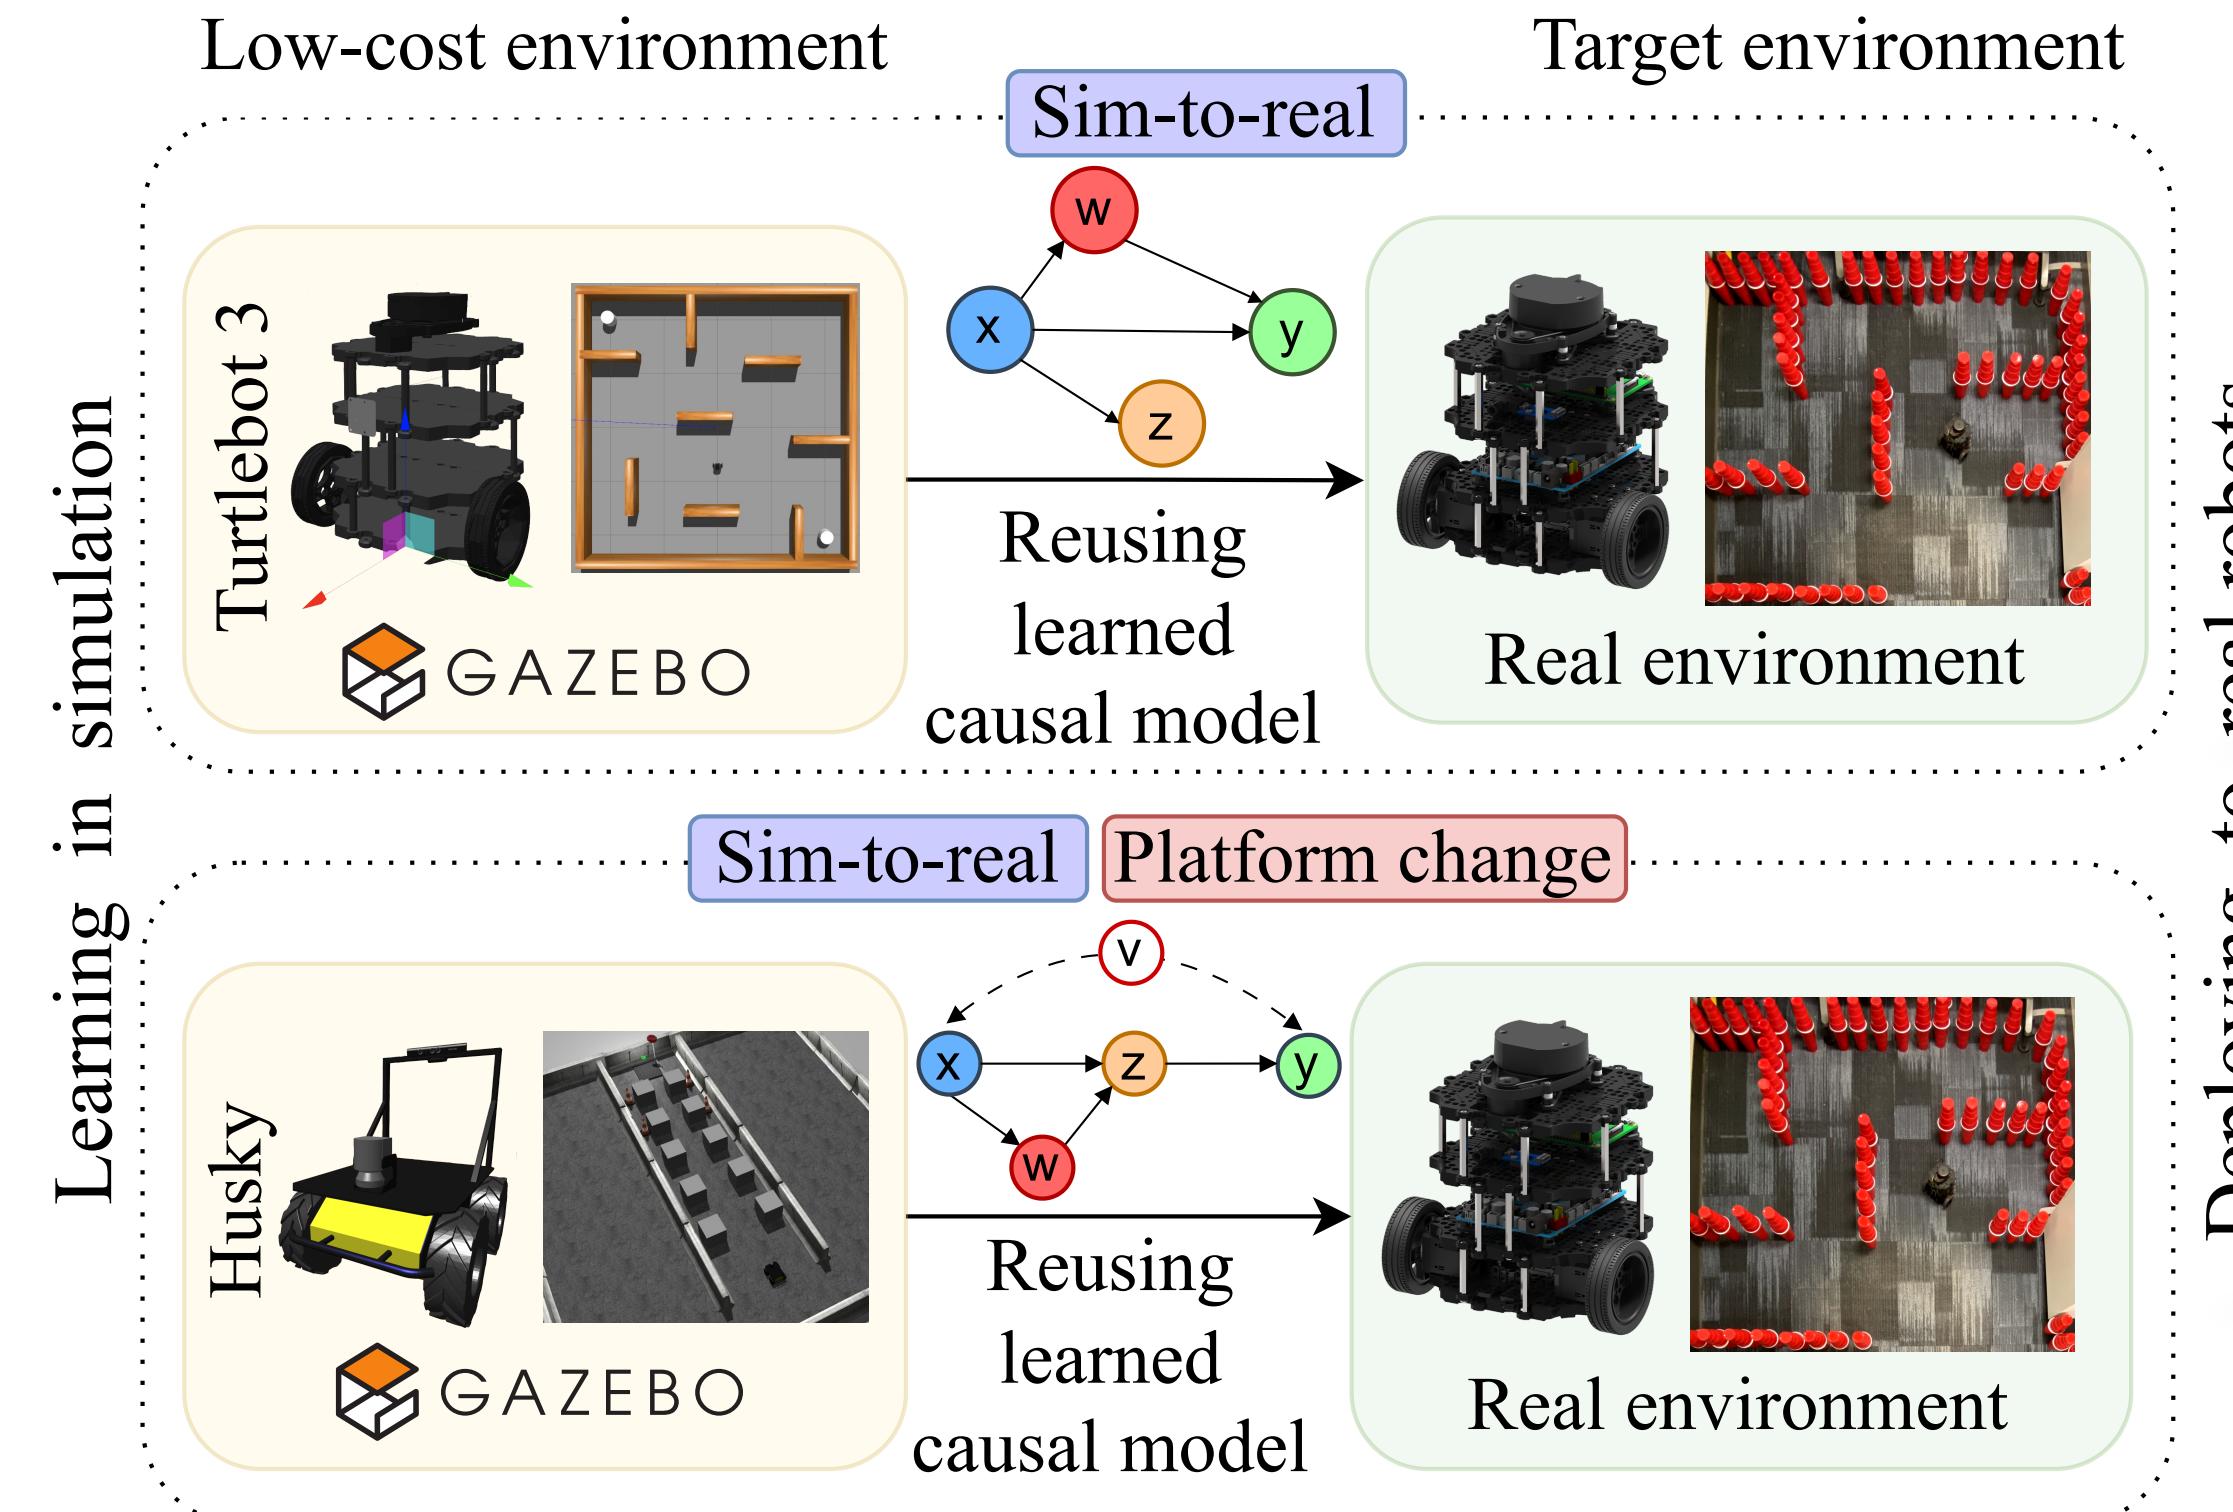
Goal

$$\mathbf{x}^* = \arg \min_{\mathbf{x} \in \mathcal{X}} f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_m(\mathbf{x}), \text{s.t. } h(\mathbf{x}) \geq 0 \quad (1)$$

where $\mathbf{x}^* \in \mathcal{X}$ is one of the Pareto-optimal configurations and adhere to the safety constraints.



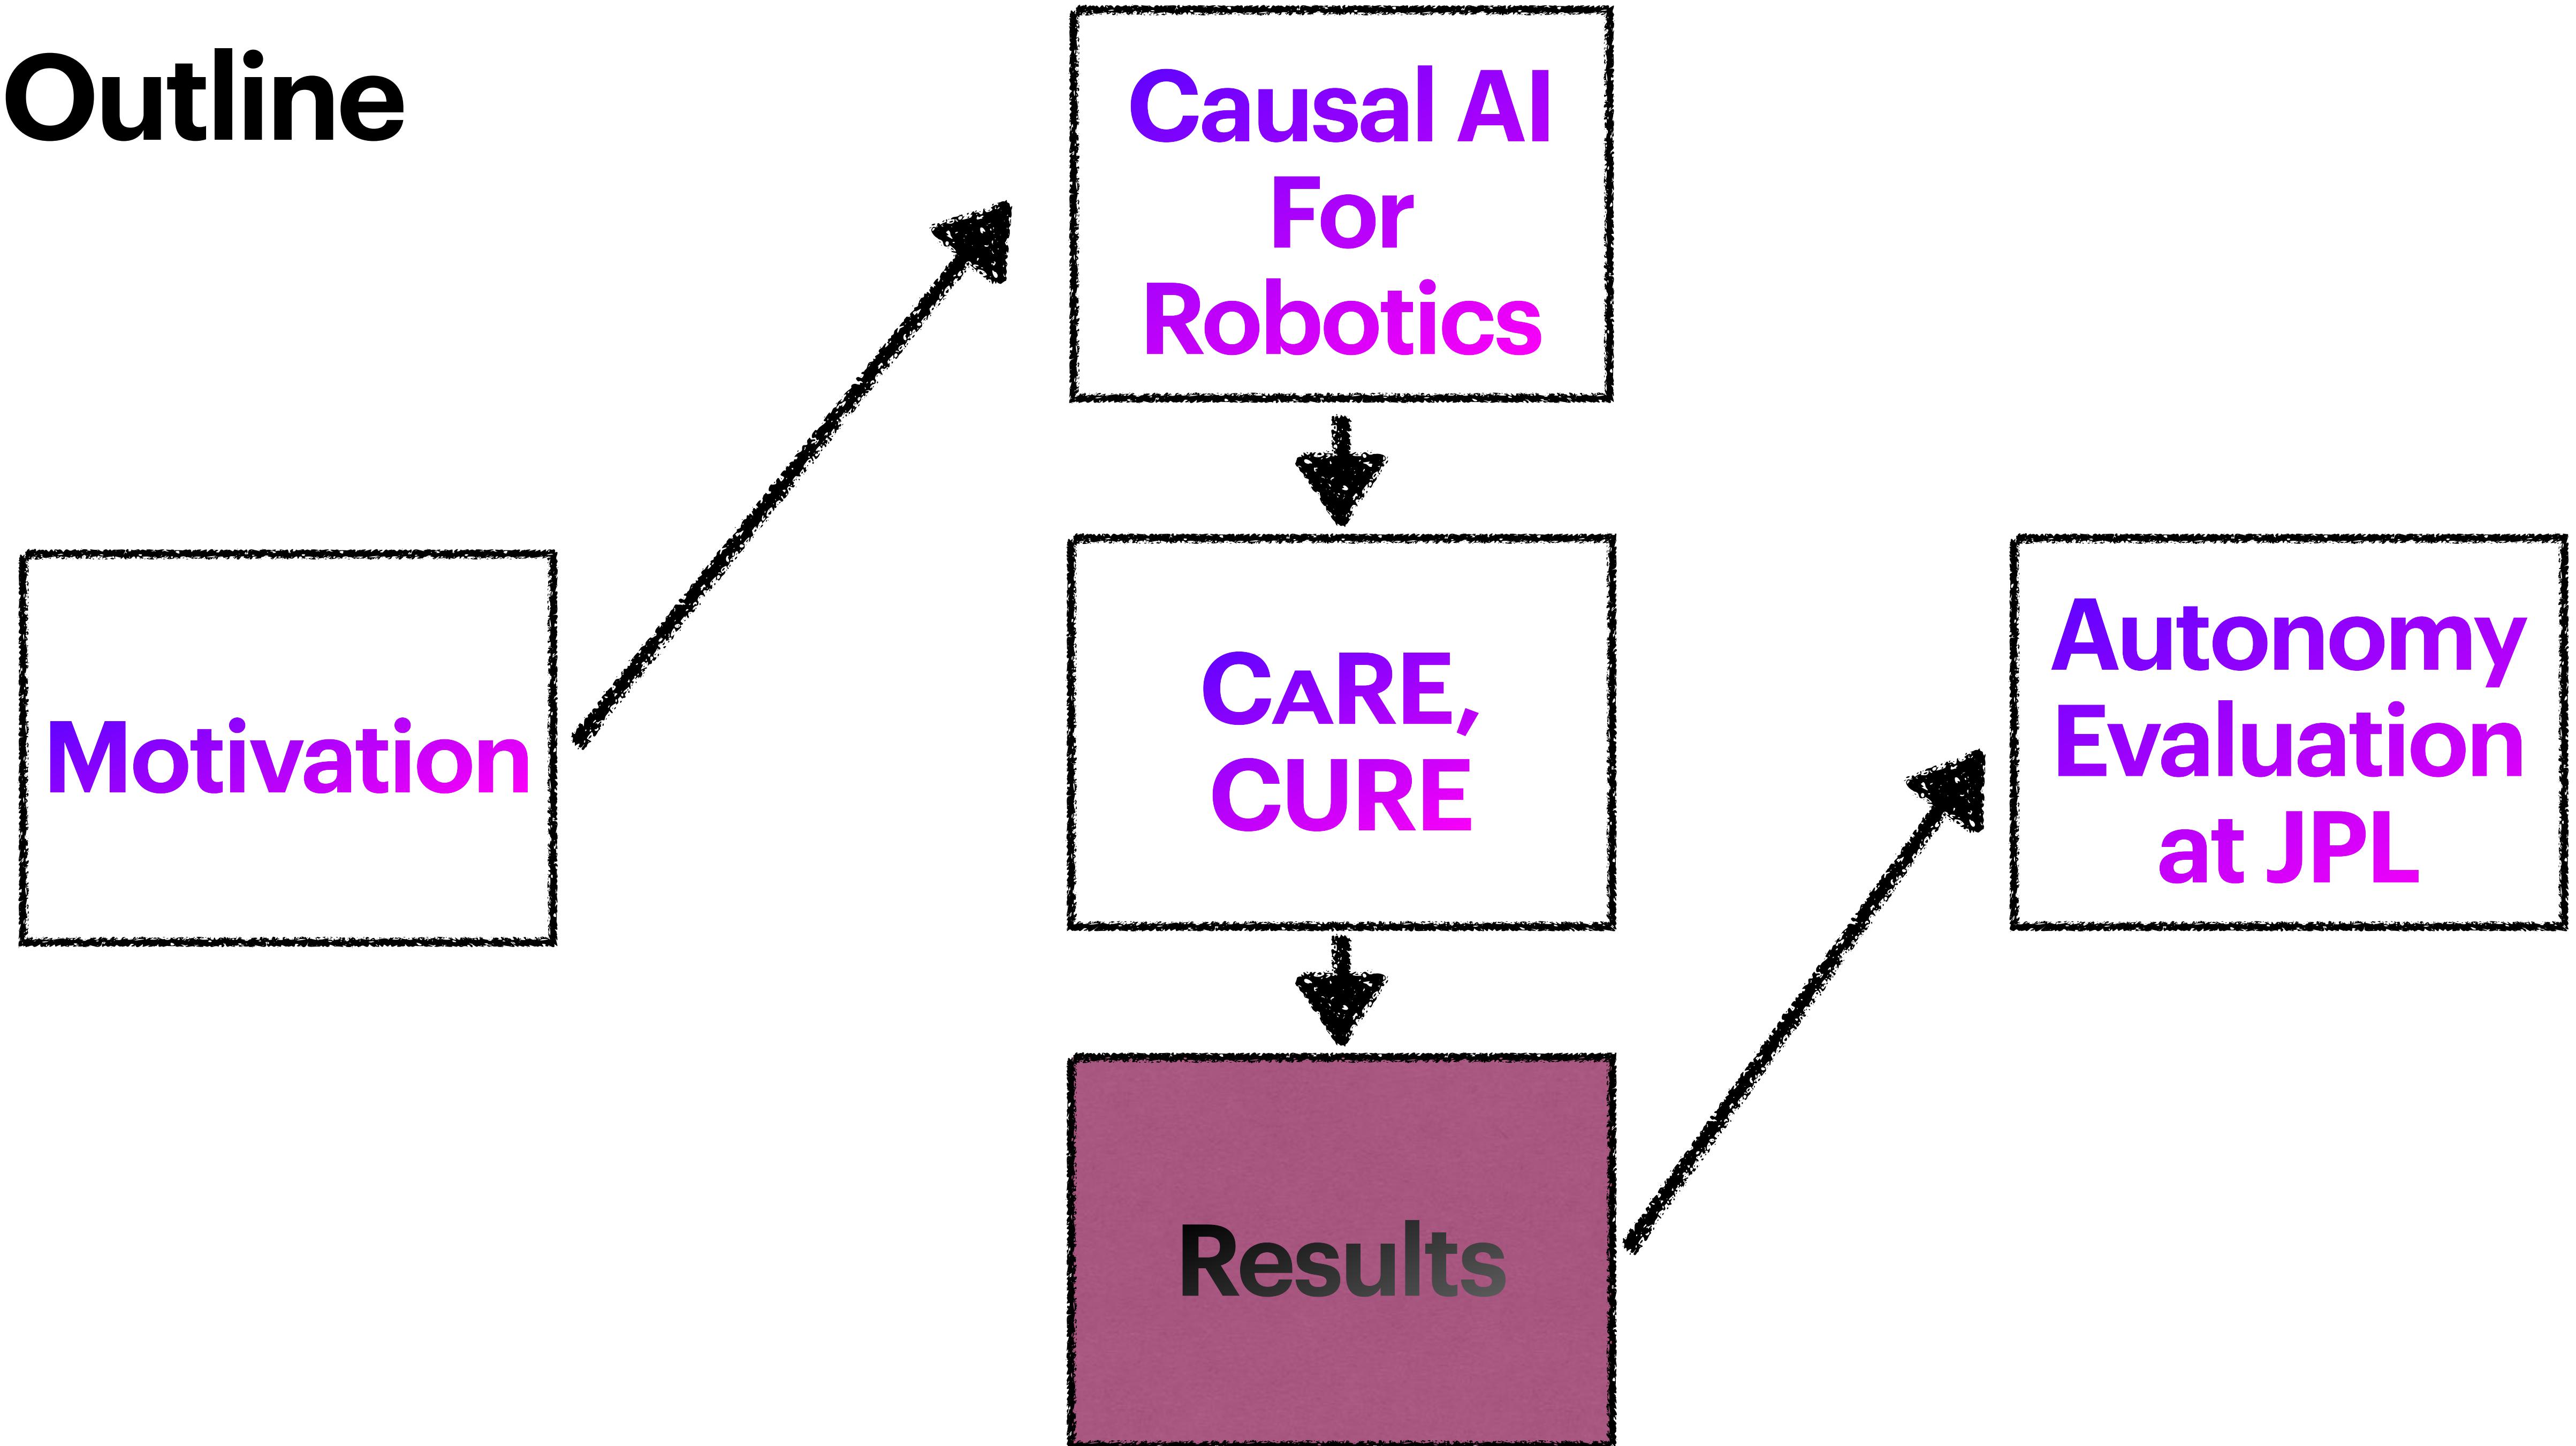
Experimental Setup



RQ1 (Effectiveness): How effective is CURE in (i) ensuring optimal performance; (ii) utilizing the budget; and (iii) respecting the safety constraints compared to the baselines?

RQ2 (Transferability): How does the effectiveness of CURE change when the severity of deployment changes varies (e.g., environment and platform change)?

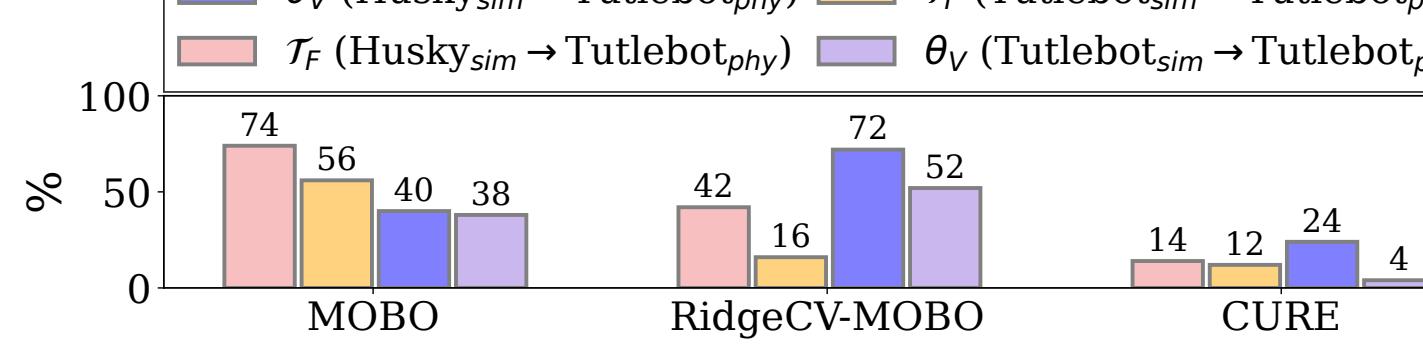
Outline



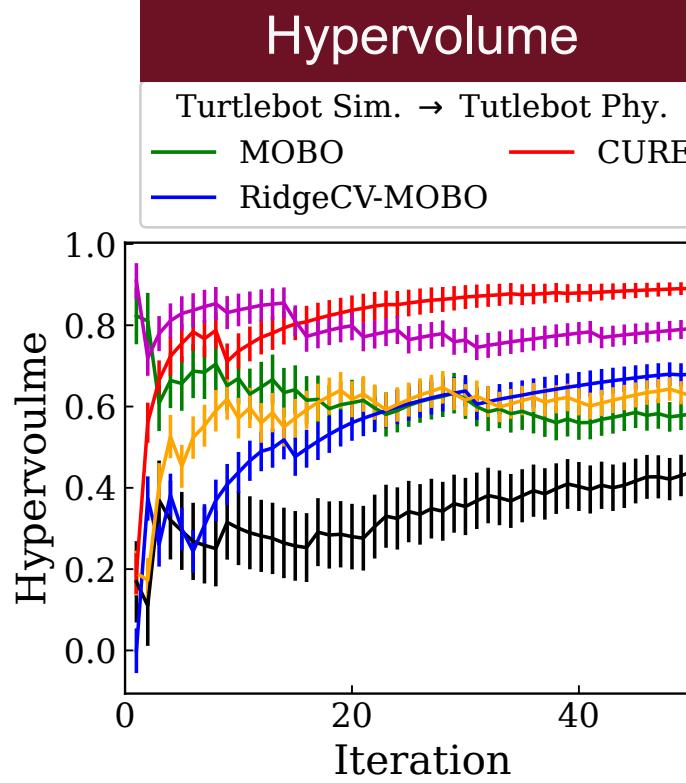
Results

Transferability and Effectiveness

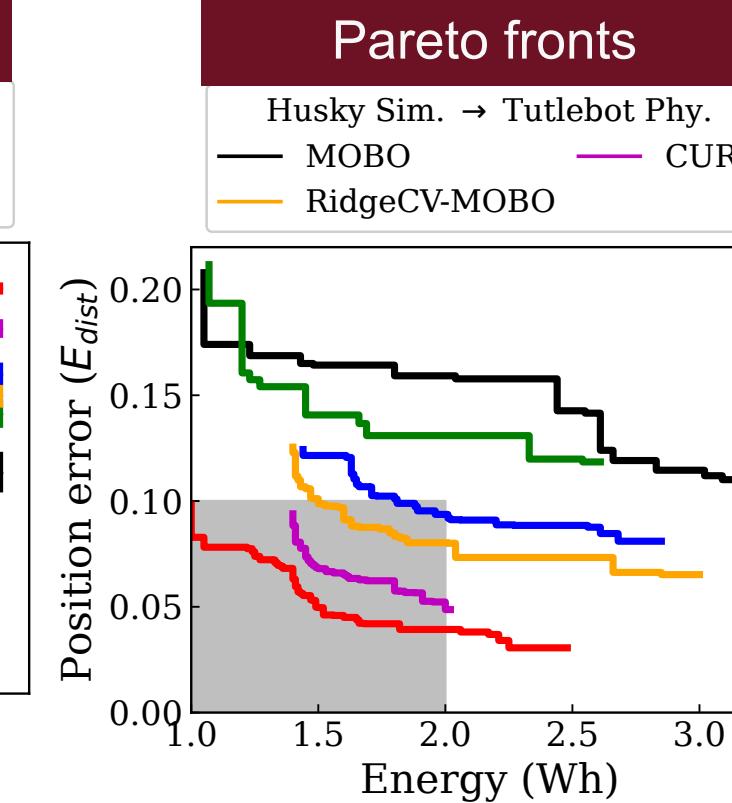
Constraint violation (θ_V) and task failure (\mathcal{T}_F) for different severity change.



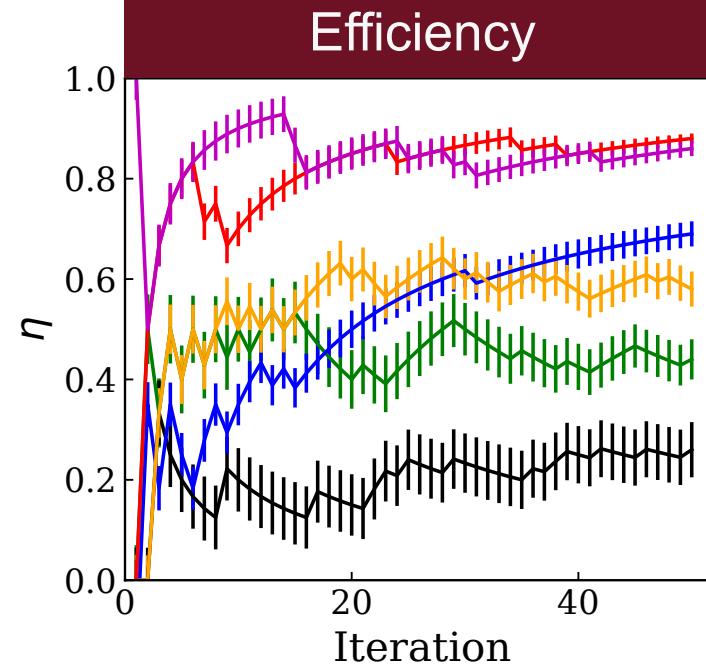
Hypervolume



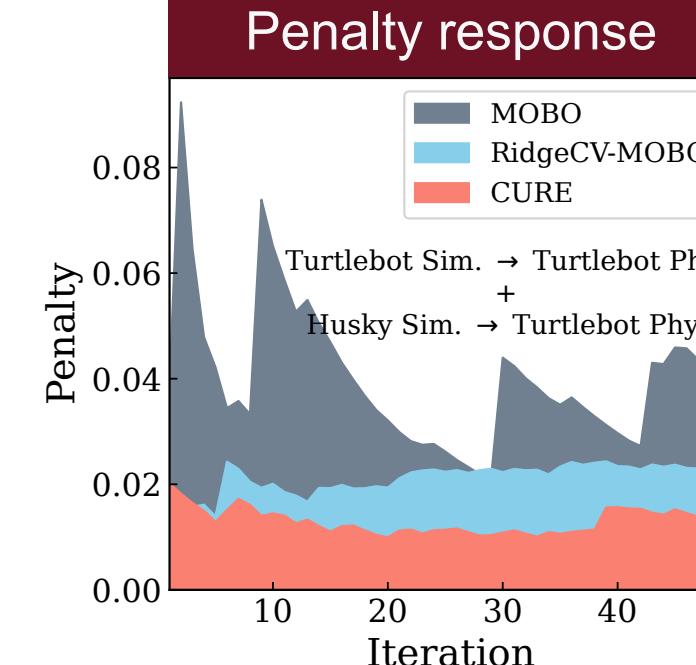
Pareto fronts



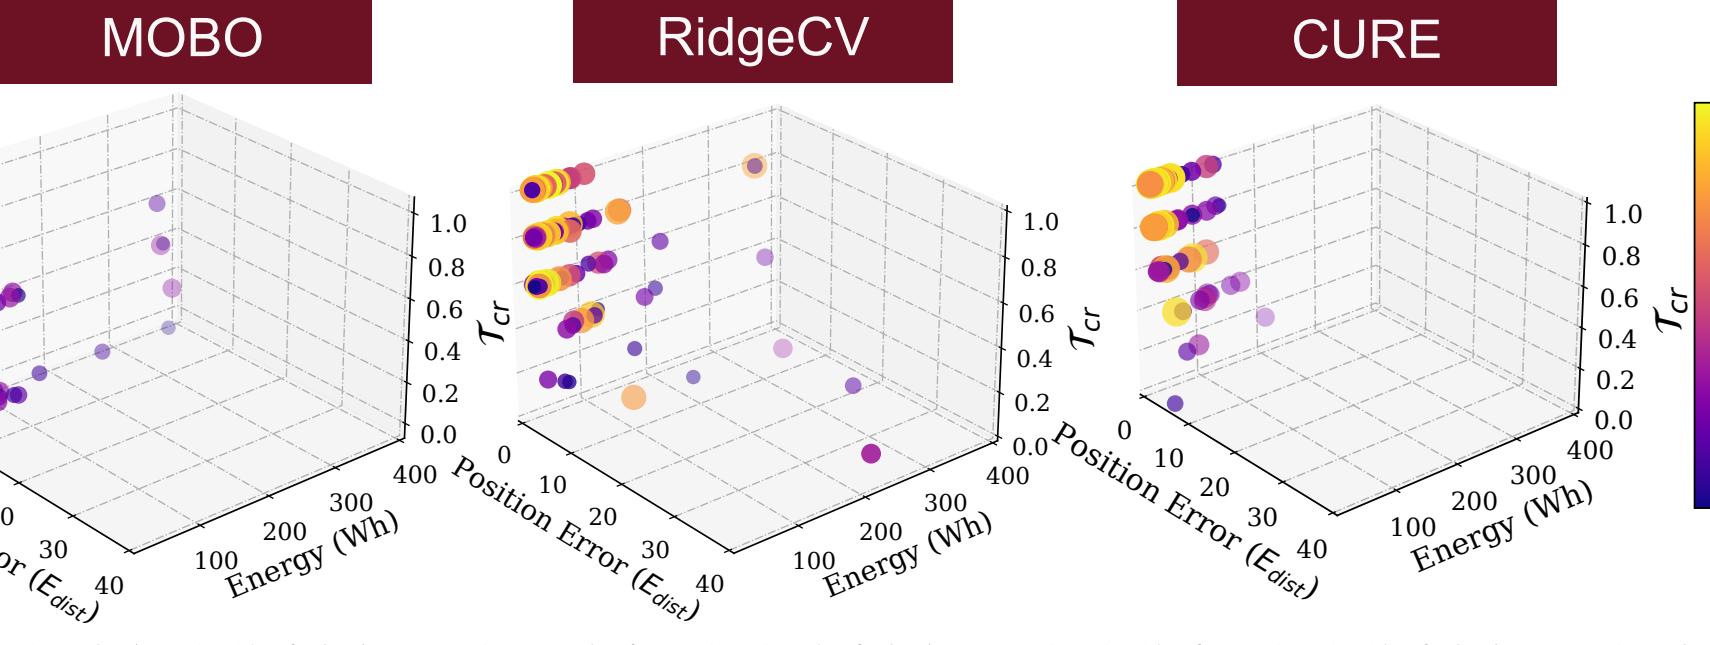
Efficiency



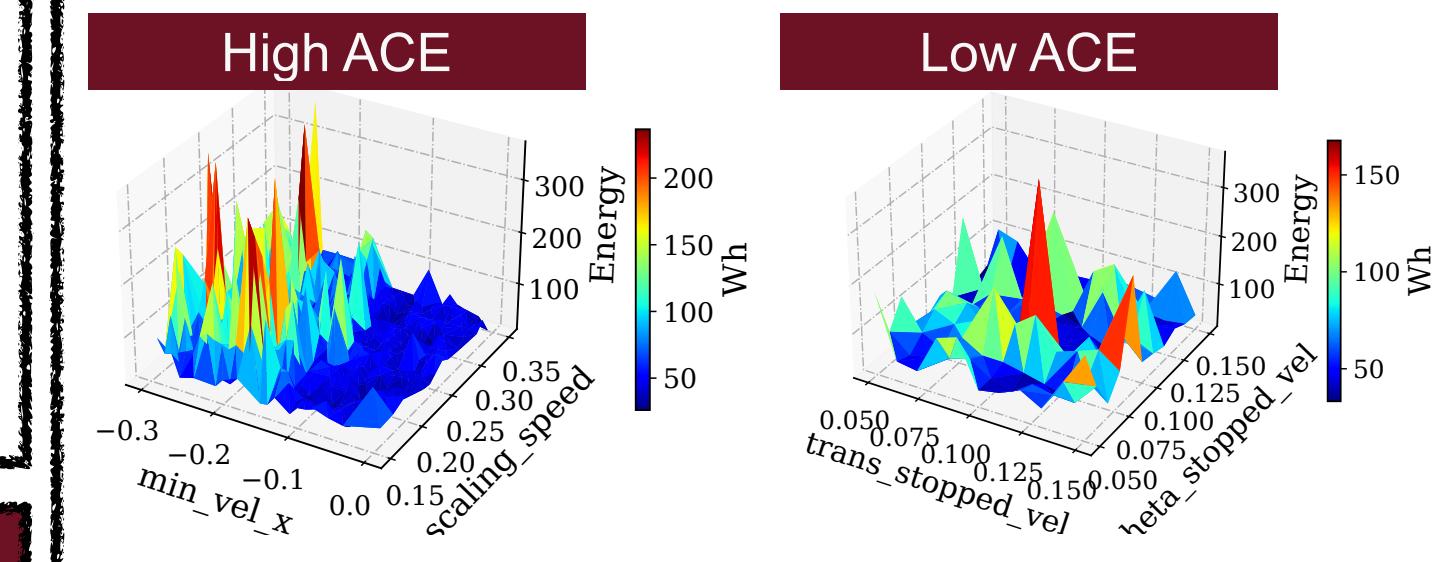
Penalty response



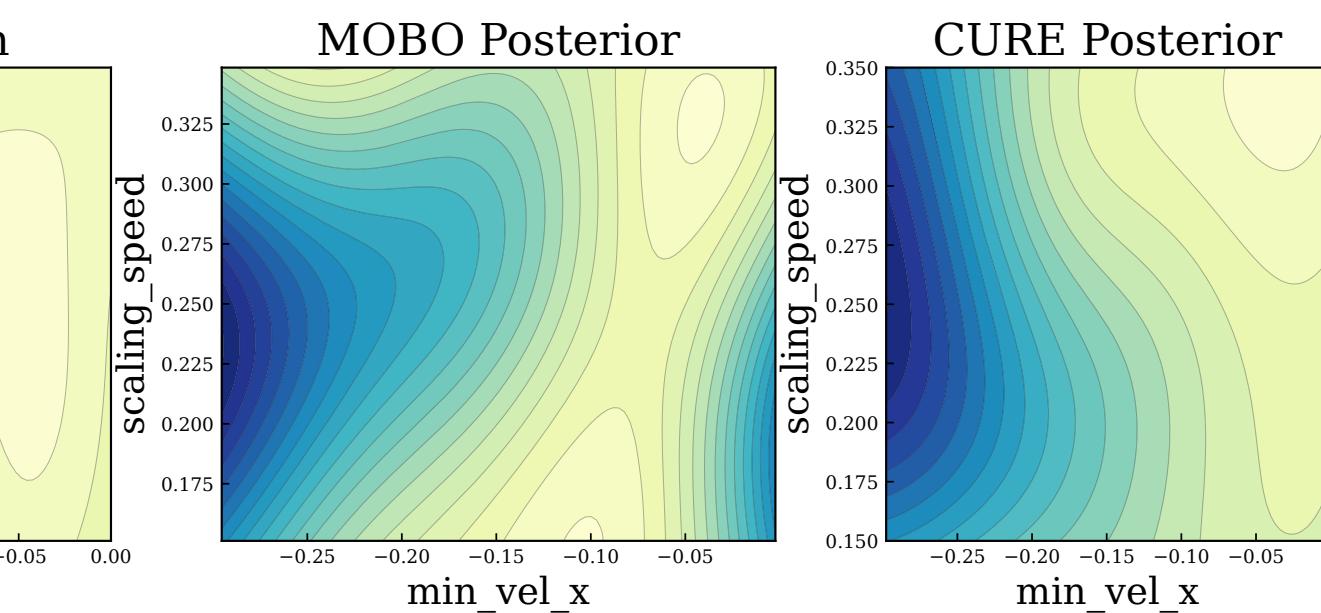
CURE demonstrates a denser surface response near the Pareto front and achieved higher task success rate (\mathcal{T}_{cr}) in fewer iterations.



Pairwise interaction between configuration options.



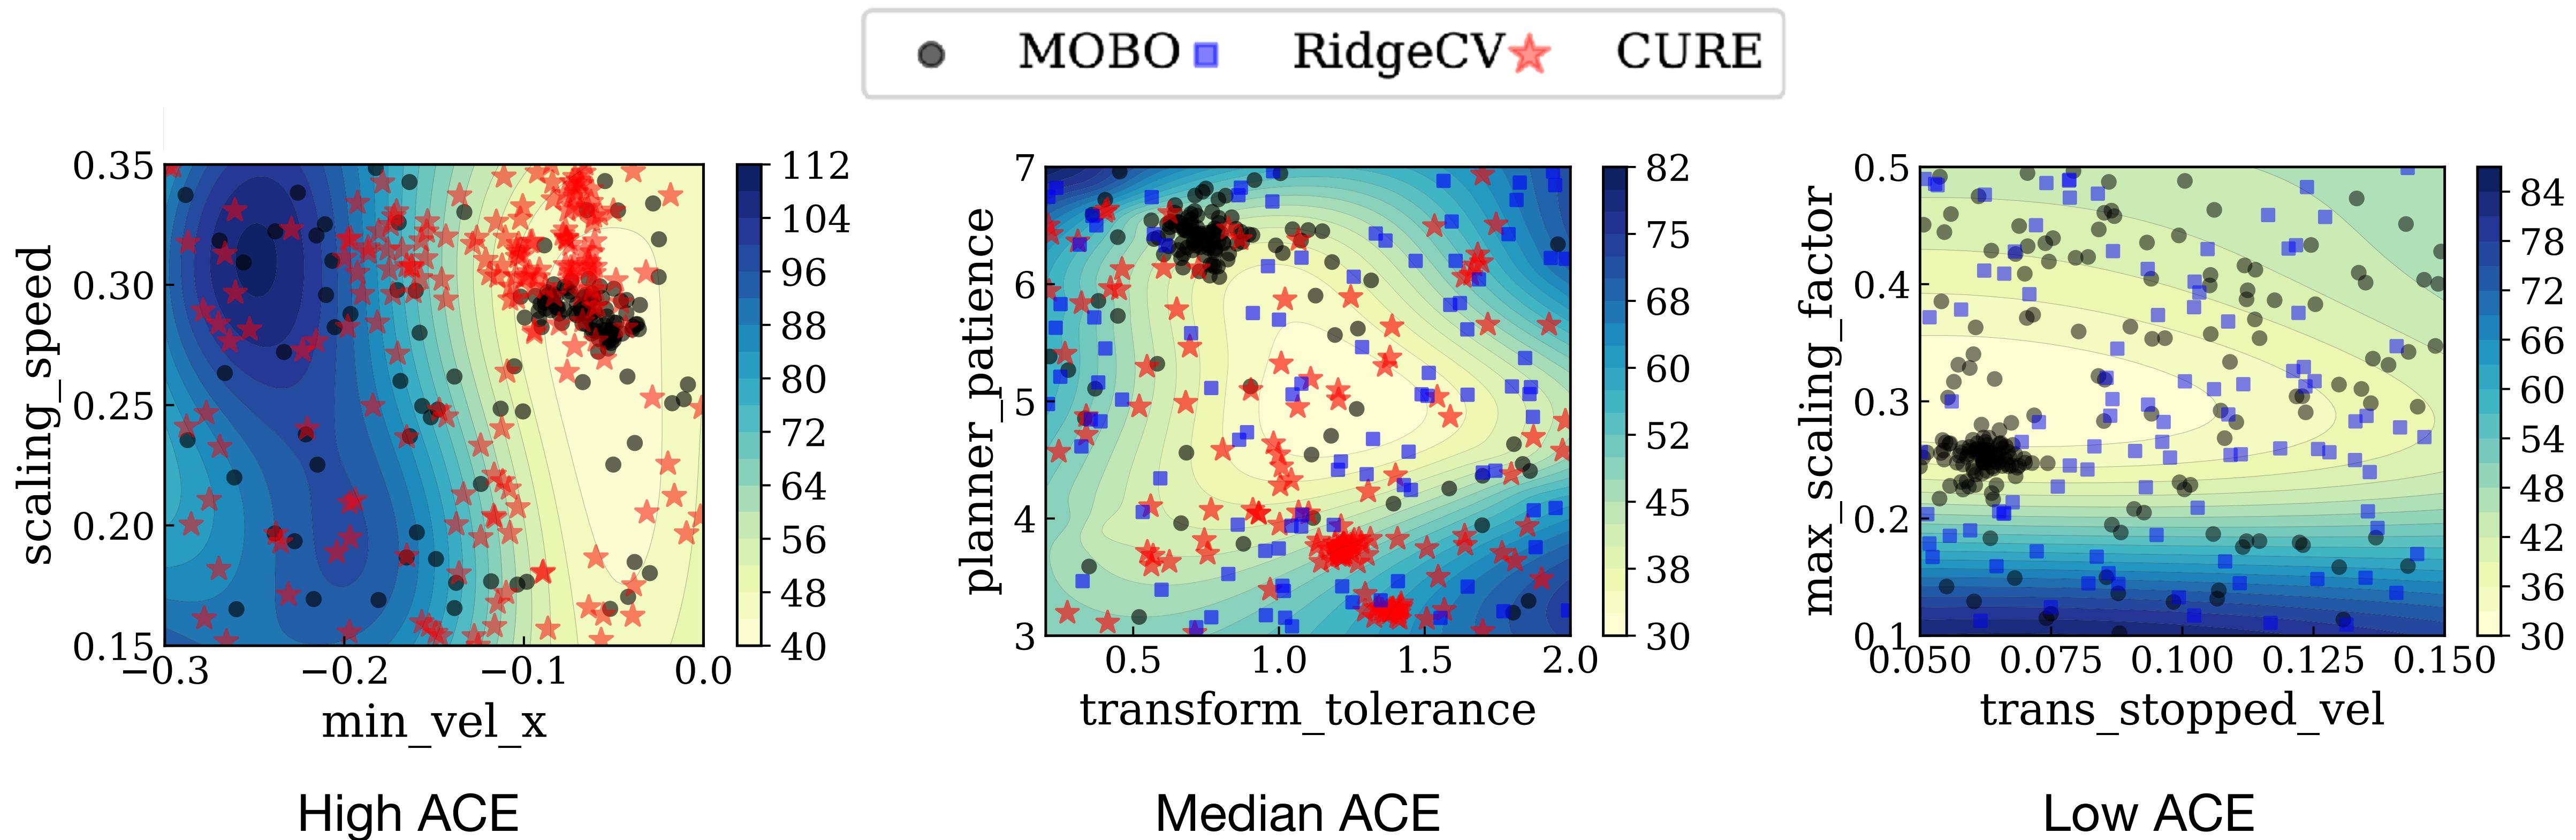
CURE demonstrates a better understanding about the performance behavior compared to multi-objective Bayesian optimization (MOBO).



Takeaway

- CURE finds a configuration that improves performance by 2x.
- CURE achieves this improvement with gains in efficiency of 4.6x when we transfer the knowledge learned from *Husky* in simulation to the *Turtlebot 3* physical robot.
- The learned causal model is transferable across similar but different settings, encompassing different environments, mission/tasks, and new robotic systems.

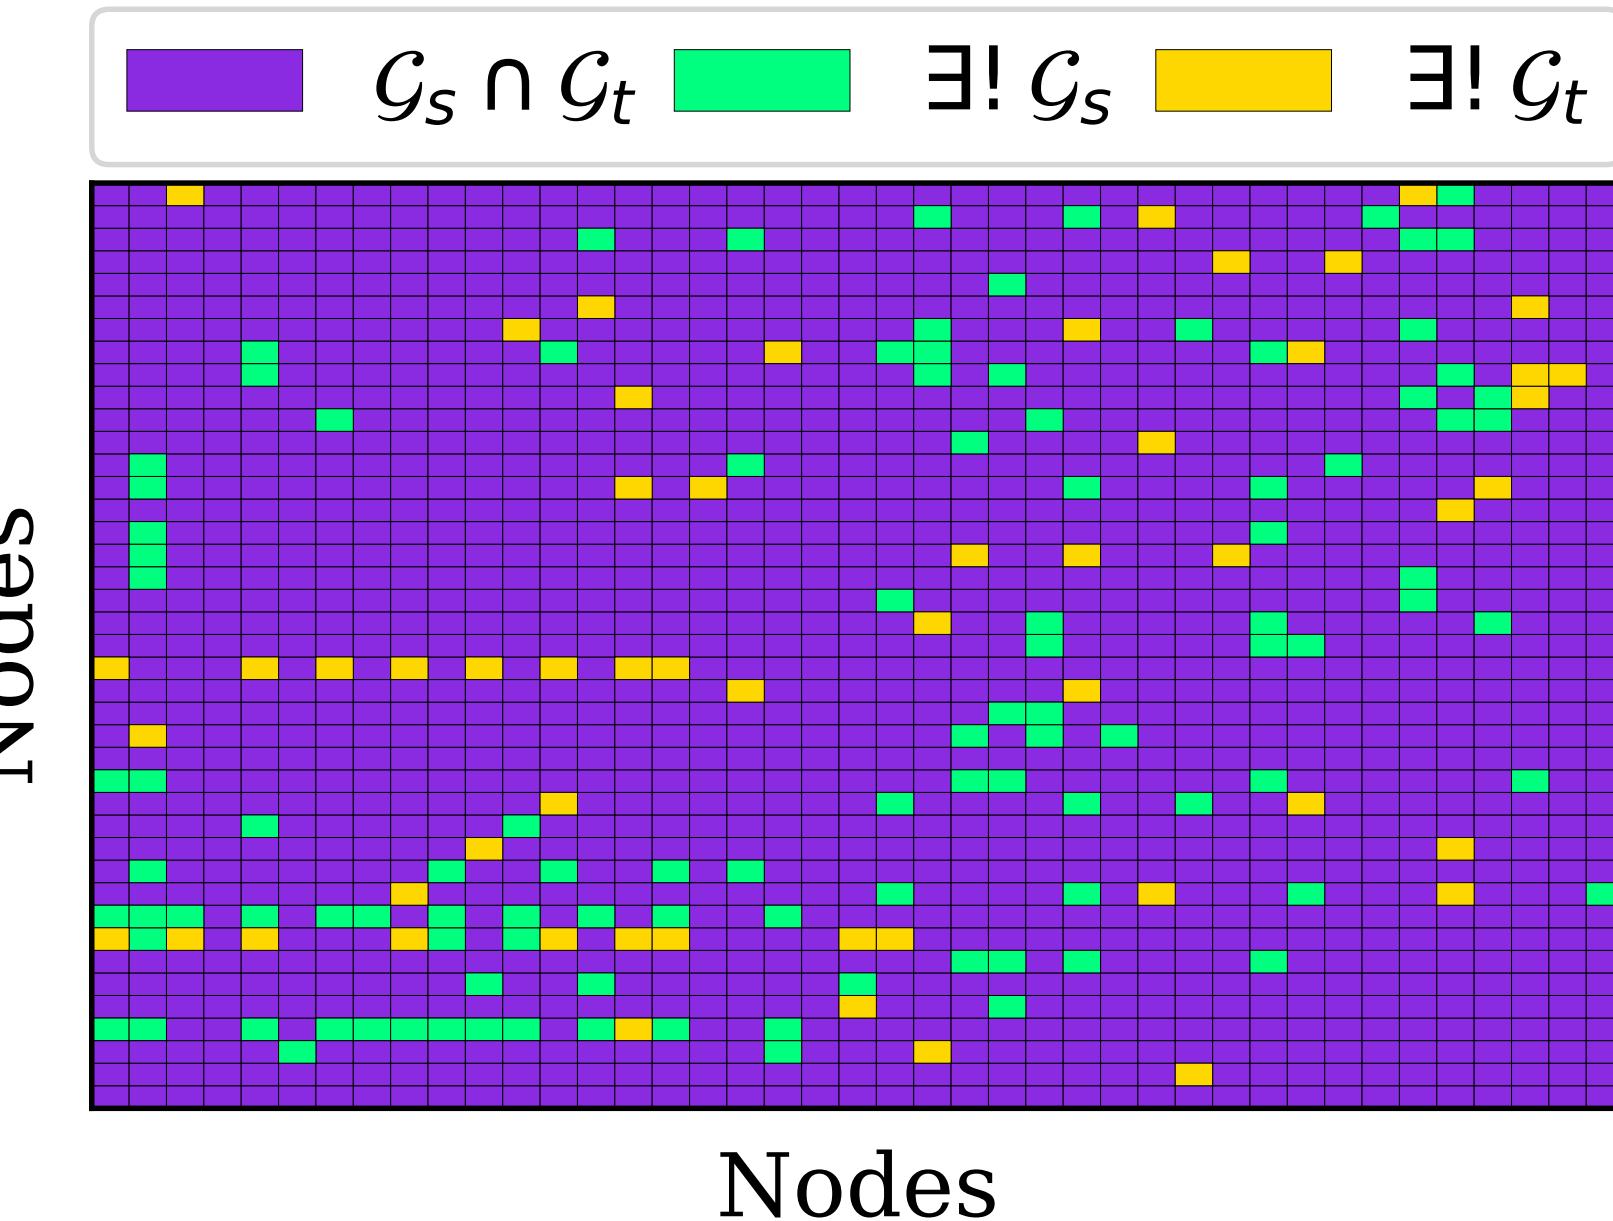
Why causal reasoning based optimization works better?



CURE's **efficient budget utilization** is attributed to a comprehensive evaluation of the **core configuration options**.

Why causal reasoning based optimization works better?

- Significant overlap between causal structures (common edges are represented as blue squares) developed in *Husky* (\mathcal{G}_s) and *Turtlebot 3* (\mathcal{G}_t).



Unique edges are represented as green and yellow squares in \mathcal{G}_s and \mathcal{G}_t , respectively.

CURE leverages the knowledge derived from the causal model learned on the source platform.

Limitations

- **Causal model error:** The NP-hard complexity of causal discovery introduces a challenge, implying that the identified causal model may not always represent the ground-truth causal relationships among variables.
 - Discrepancies between the causal structure discovered and the actual structures.
- **Potential biases when transferring the causal model:** Caution must be exercised when reusing the entire causal graph learned from the source platform.

Code and Data

IEEE Robotics and Automation Letters, 2023 → Presented in IROS 2023

IEEE ROBOTICS AND AUTOMATION LETTERS, VOL. 8, NO. 7, JULY 2023

4115

CARE: Finding Root Causes of Configuration Issues in Highly-Configurable Robots

Md Abir Hossen[✉], Sonam Kharade, Bradley Schmerl, Senior Member, IEEE, Javier Cámará[✉], Jason M. O’Kane[✉], Ellen C. Czaplinski[✉], Katherine A. Dzurilla, David Garlan[✉], Life Fellow, IEEE, and Pooyan Jamshidi

Abstract—Robotic systems have subsystems with a combinatorially large configuration space and hundreds or thousands of possible software and hardware configuration options interacting non-trivially. The configurable parameters are set to target specific objectives, but they can cause functional faults when incorrectly configured. Finding the root cause of such faults is challenging due to the exponentially large configuration space and the dependencies between the robot’s configuration settings and performance. This paper proposes CARE—a method for diagnosing the root cause of functional faults through the lens of causality. CARE abstracts the causal relationships between various configuration options and the robot’s performance objectives by learning a causal structure and estimating the causal effects of options on robot performance indicators. We demonstrate CARE’s efficacy by finding the root cause of the observed functional faults and validating the diagnosed root cause by conducting experiments in both physical robots (*Husky* and *Turtlebot 3*) and in simulation (*Gazebo*). Furthermore, we demonstrate that the causal models learned from robots in simulation (e.g., *Husky* in *Gazebo*) are transferable to physical robots across different platforms (e.g., *Husky* and *Turtlebot 3*).

I. INTRODUCTION

ROBOTIC systems are highly configurable, typically composed of multiple subsystems (e.g., localization, navigation), each of which has numerous configurable components (e.g., selecting path planning algorithms in the planner). Once an algorithm has been selected for a component, its associated parameters must be set to the appropriate values (e.g., use grid path = True). The configuration space in such robotic systems is combinatorially large, with hundreds if not thousands of software and hardware configuration choices that interact non-trivially with one another. Indeed, incorrectly specified configuration options are one of the most common causes of system failure [1]. The configuration space in robotic systems directly impacts mission objectives (e.g., navigating from one place to another), enabling trade-offs in the objective space (e.g., the time that it takes to reach the target location(s) vs. the energy consumption for the task). The magnitude of the trade-off (even for the same configuration option) is dictated by the characteristics

IEEE Transactions on Robotics, 2024 (under review)

CURE: Simulation-Augmented Auto-Tuning in Robotics

Md Abir Hossen[✉], Sonam Kharade[✉], Jason M. O’Kane[✉], Senior Member, IEEE, Bradley Schmerl[✉], Senior Member, IEEE, David Garlan[✉], Life Fellow, IEEE, and Pooyan Jamshidi[✉]

Abstract—Robotic systems are typically composed of various subsystems, such as localization and navigation, each encompassing numerous configurable components (e.g., selecting different planning algorithms). Once an algorithm has been selected for a component, its associated configuration options must be set to the appropriate values. Configuration options across the system stack interact non-trivially. Finding optimal configurations for highly configurable robots to achieve desired performance poses a significant challenge due to the interactions between configuration options across software and hardware that result in an exponentially large and complex configuration space. These challenges are further compounded by the need for transferability between different environments and robotic platforms. Data efficient optimization algorithms (e.g., Bayesian optimization) have been increasingly employed to automate the tuning of configurable parameters in cyber-physical systems. However, such optimization algorithms converge at later stages, often after exhausting the allocated budget (e.g., optimization steps, allotted time) and lacking transferability. This paper proposes CURE—a method that identifies causally relevant configuration options, enabling the optimization process to operate in a reduced search space, thereby enabling faster optimization.

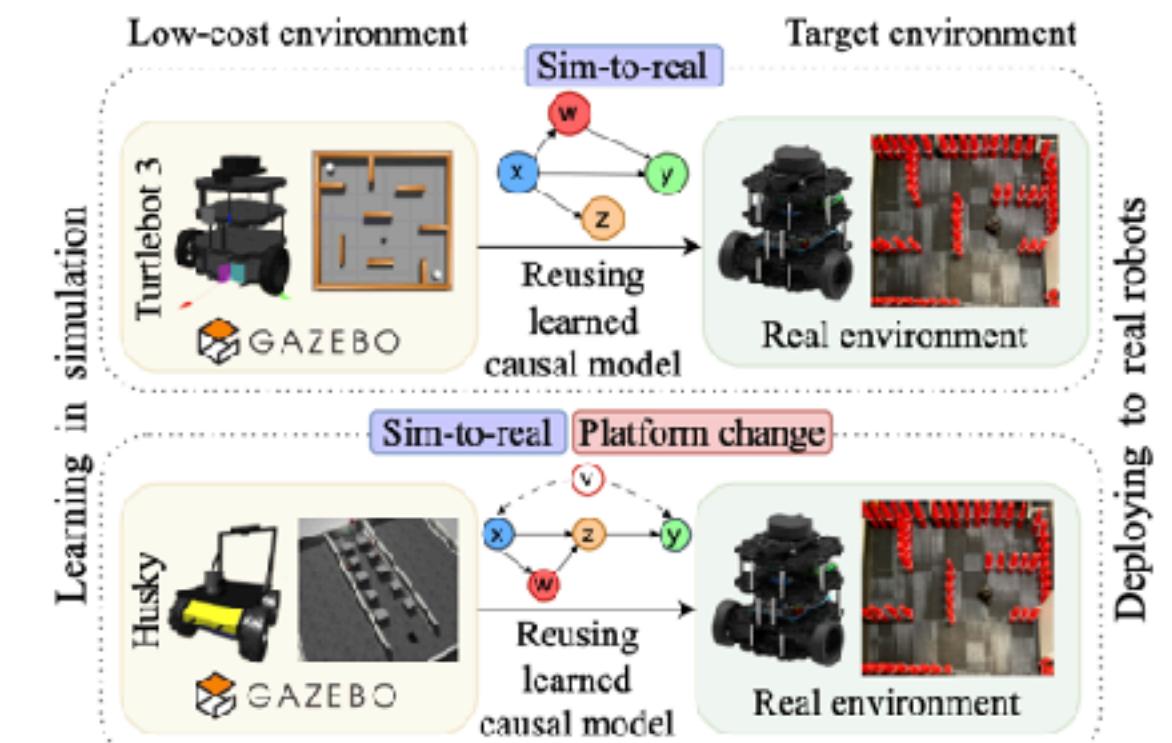


Fig. 1: Sim-to-real: applying the knowledge of the learned causal model using *Turtlebot 3* in simulation to the *Turtlebot 3* physical robot. Sim-to-real & Platform change: transferring the causal model learned using *Husky* in simulation to the *Turtlebot 3* physical robot.



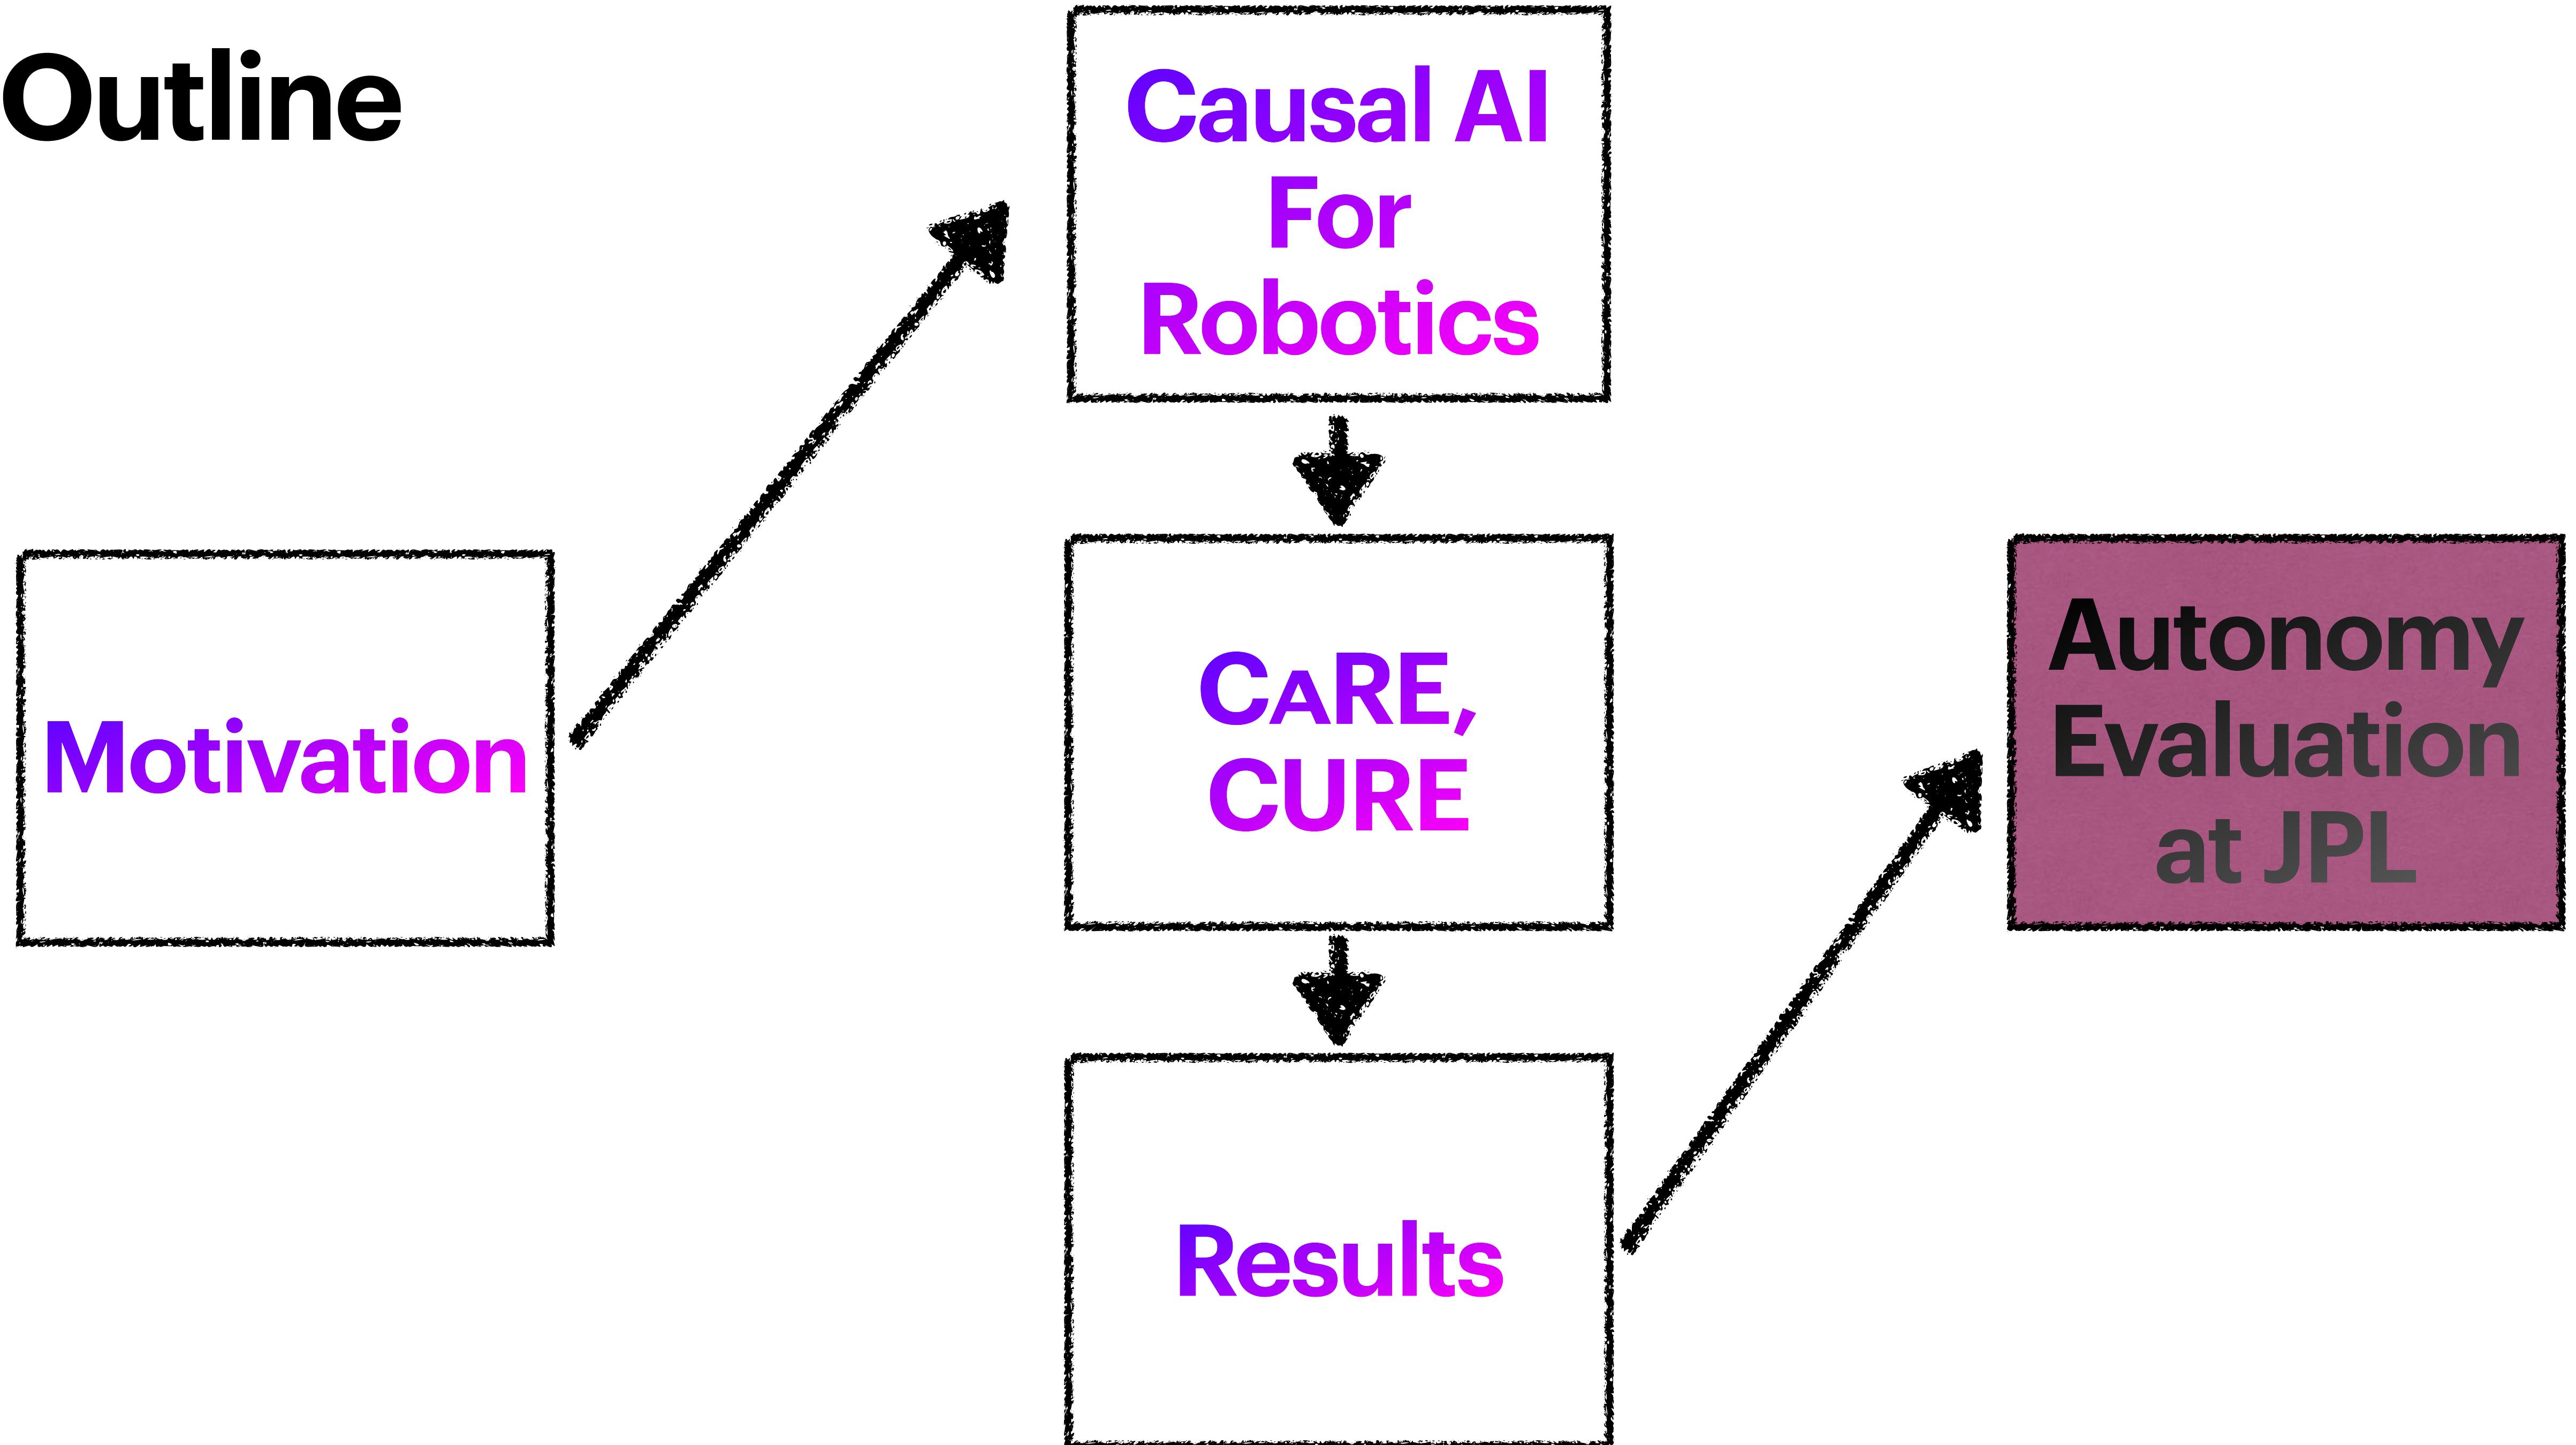
<https://github.com/softsys4ai/care>



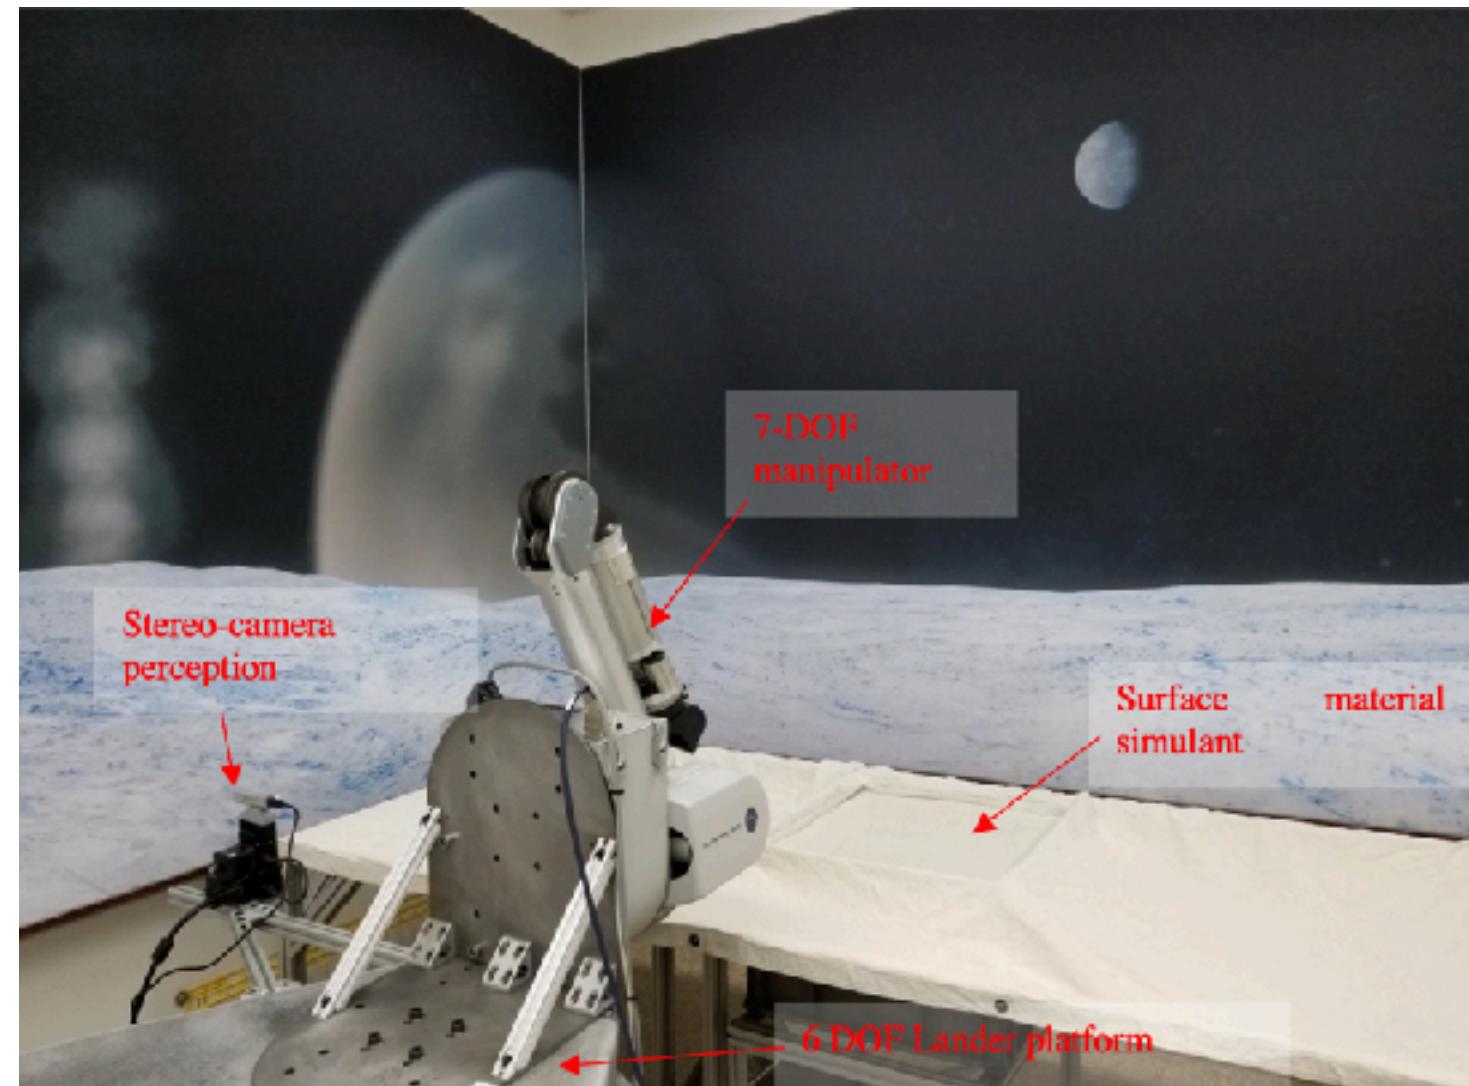
<https://github.com/softsys4ai/cure>



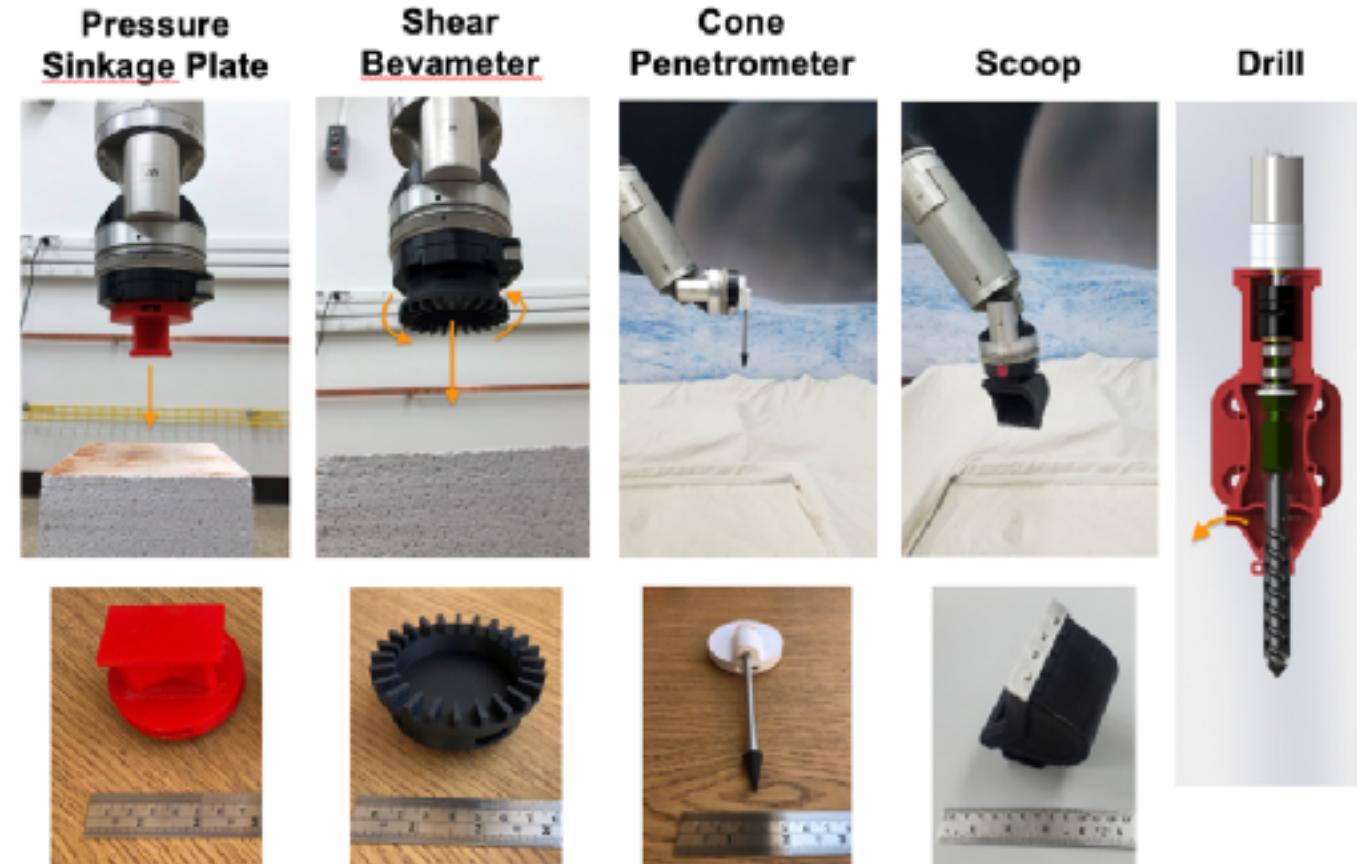
Outline



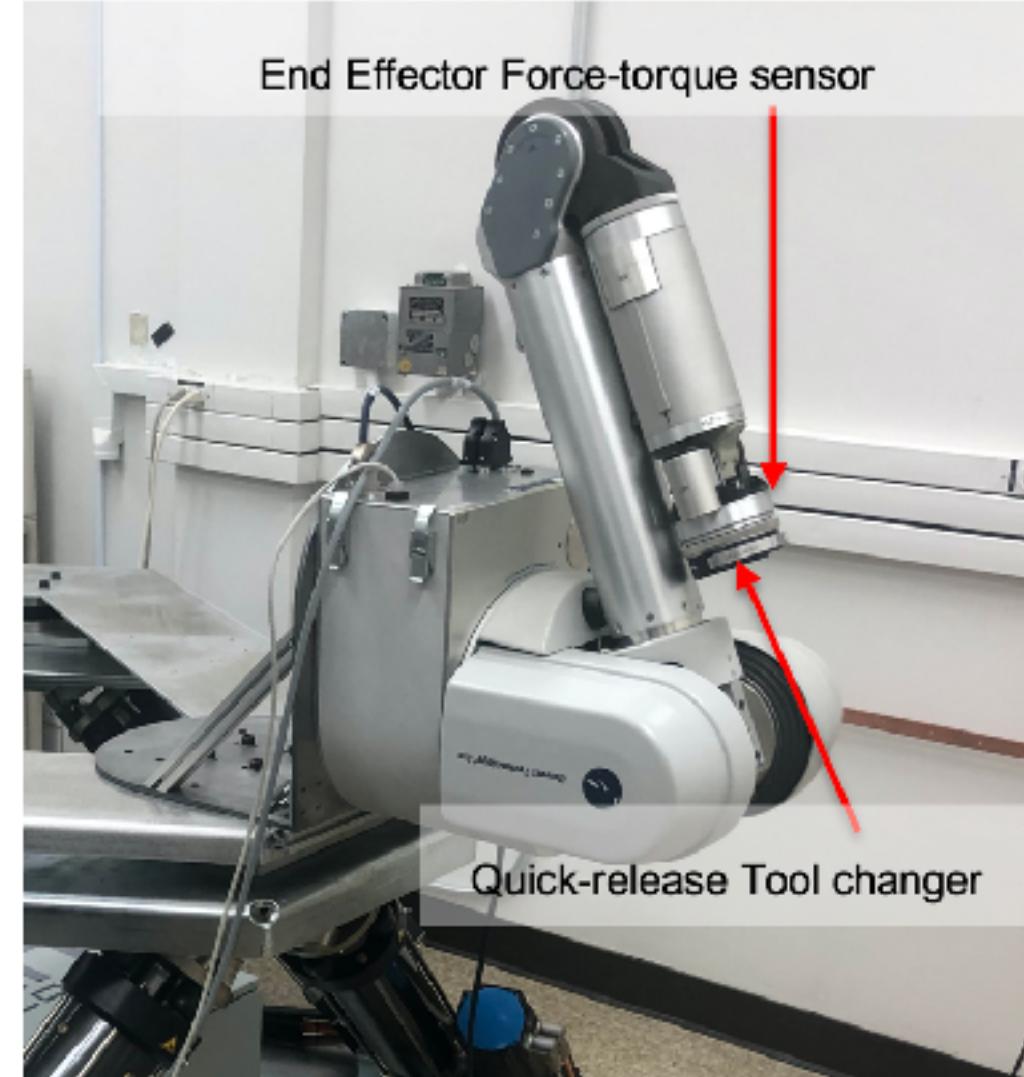
Physical Space Lander Testbed at JPL



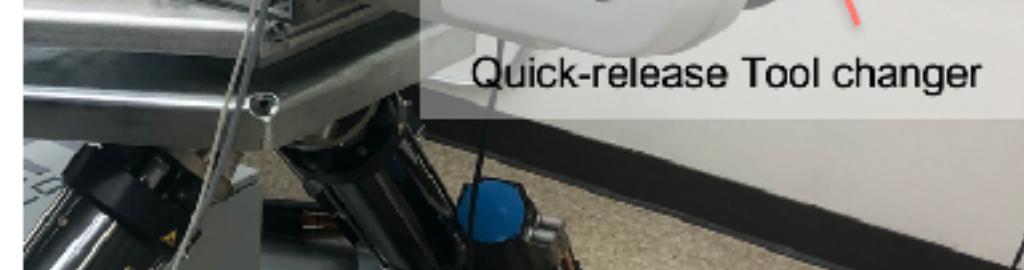
Testbed setup and major components



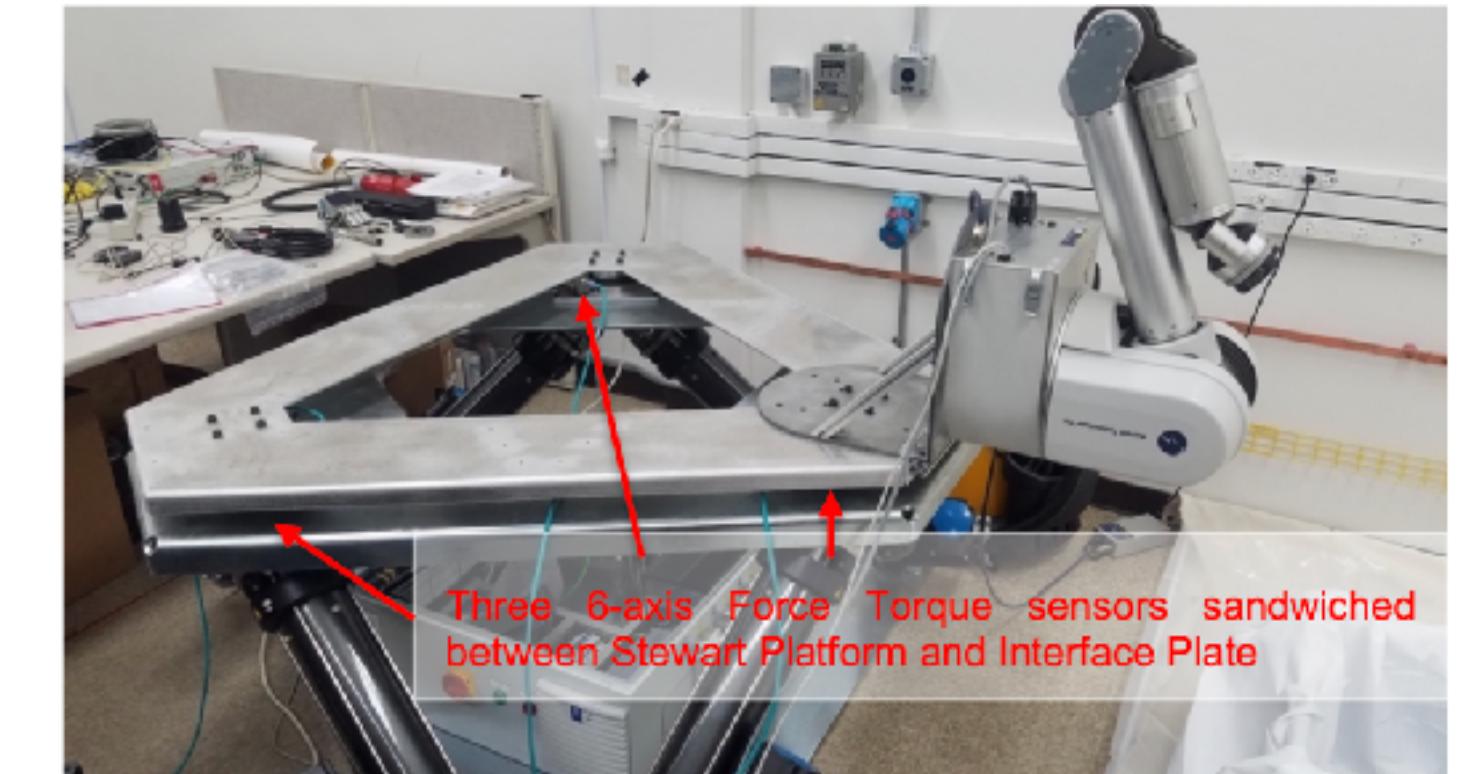
Modular instruments to be mounted on robot arm



Barrett WAM seven DOF manipulator arm mounted to lander with wrist FTS and tool changer



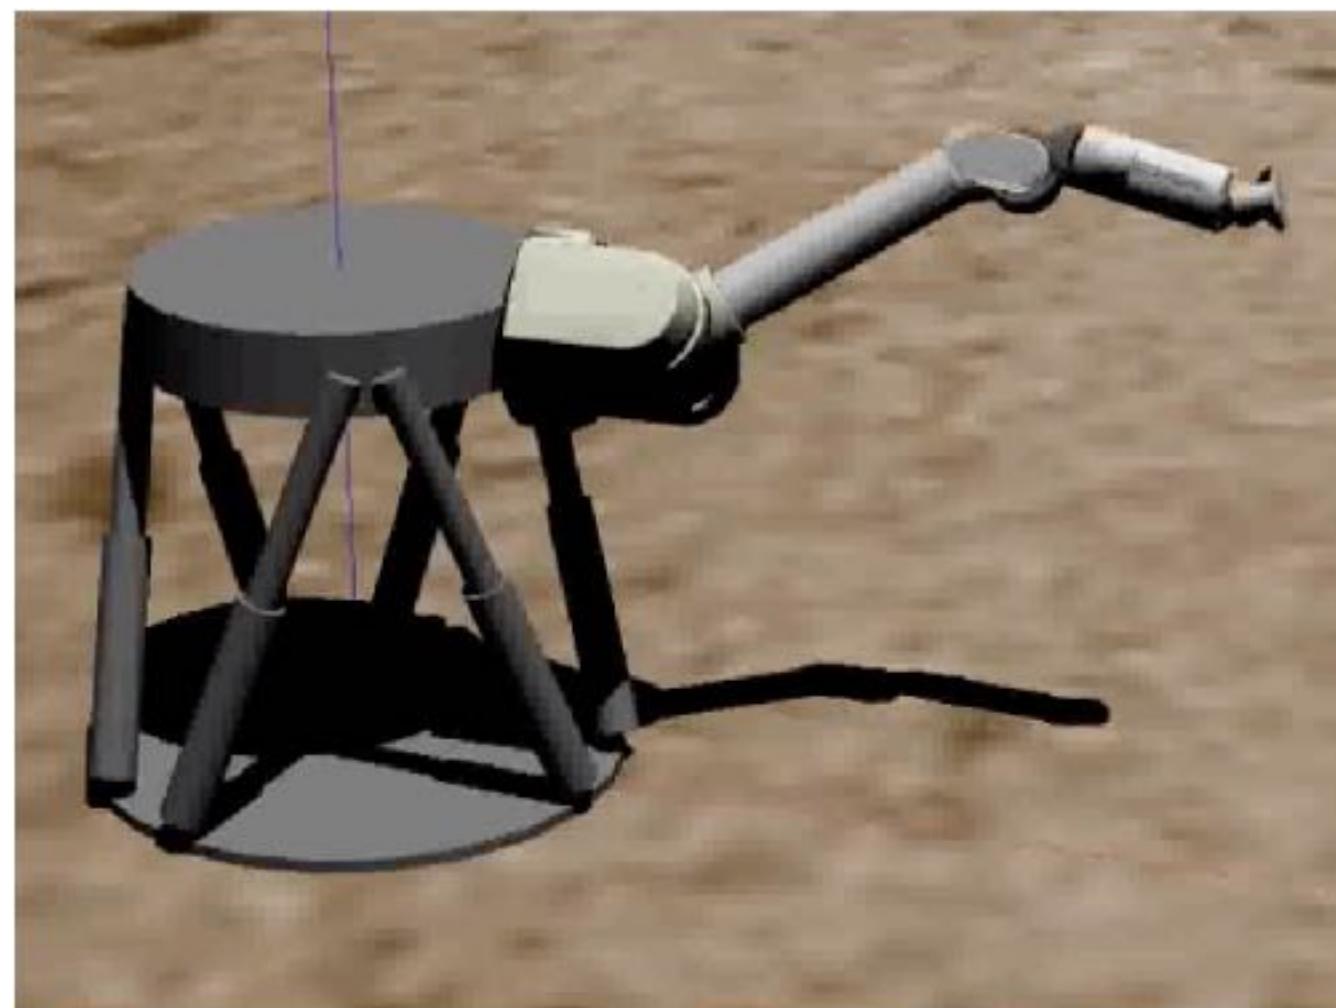
HITL simulator of lander and manipulator



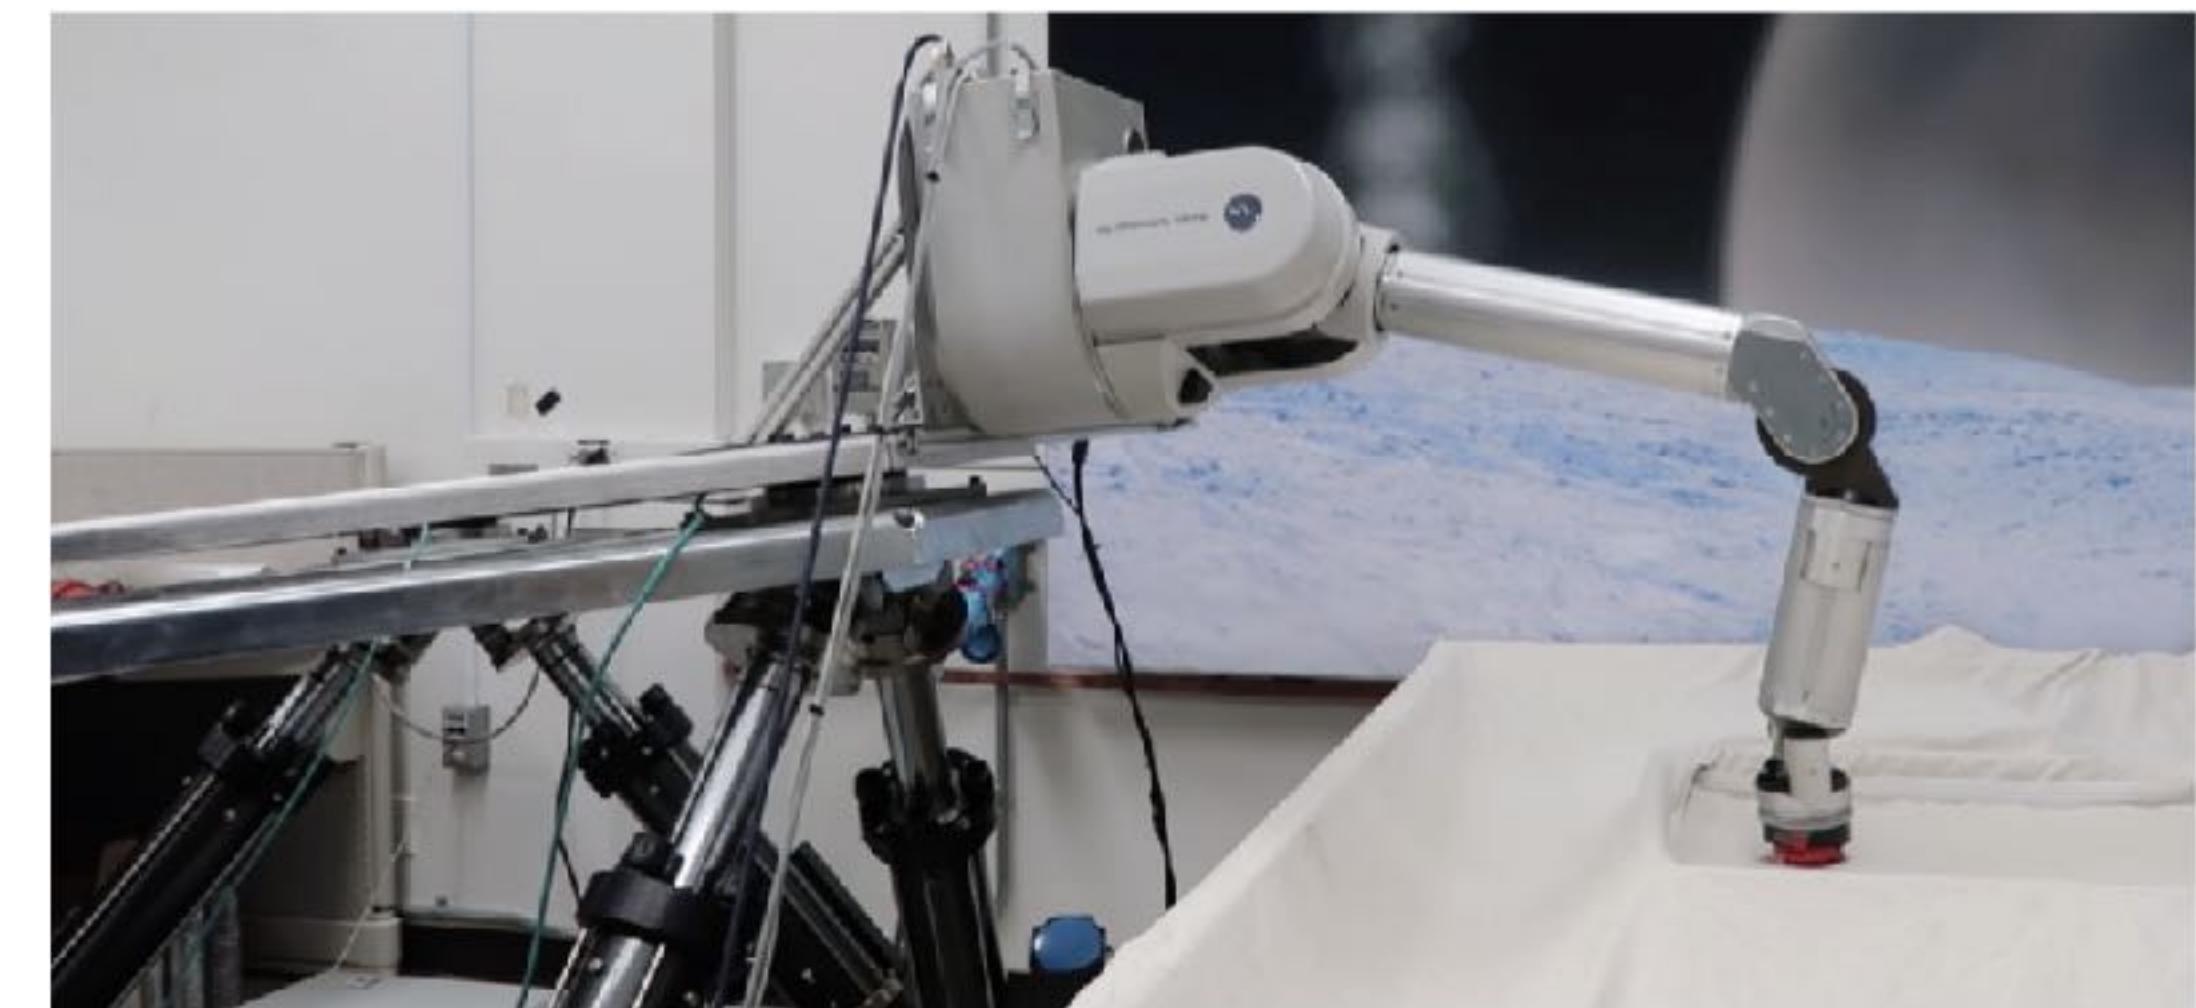
E2M Technologies six DOF Stewart Platform representing spacecraft lander

Learning in Simulation for Transfer Learning to Physical Testbed

Simulation Environment



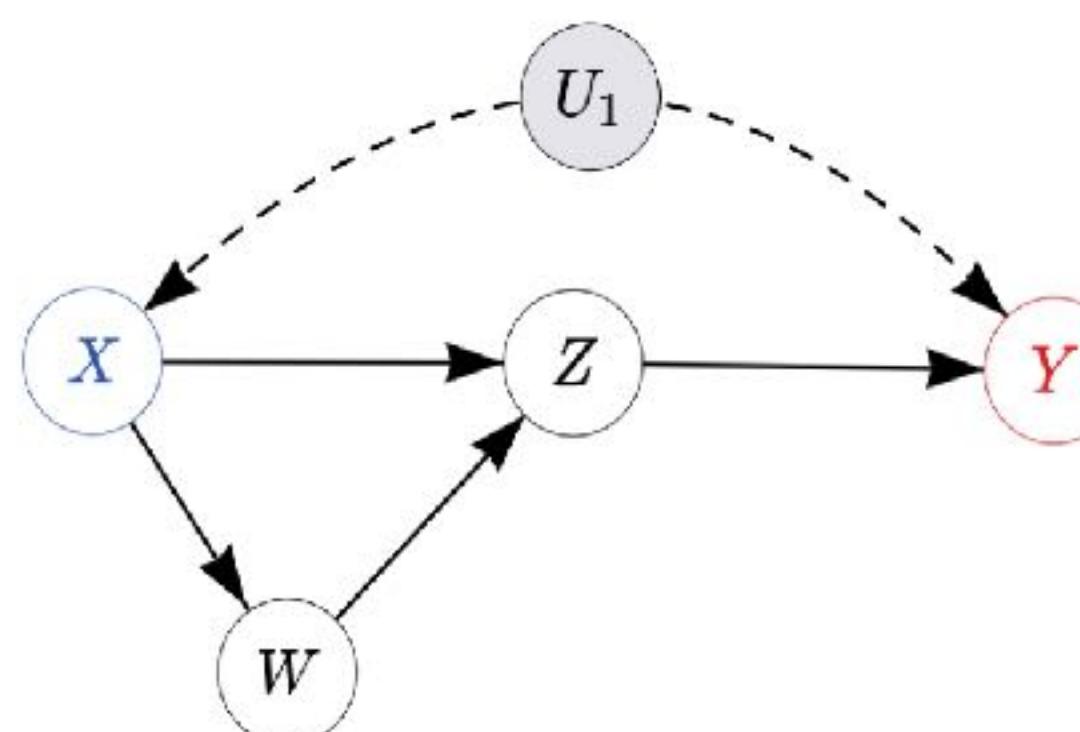
Physical Testbed



Transfer



Causal Invariances



OWLAT-sim

OWLAT

Demo



Extending CURE: Causal Multi-information Source Optimization

- Configuration: $\mathbf{x} = \langle X_1 = \mathbf{x}_1, X_2 = \mathbf{x}_2, \dots, X_d = \mathbf{x}_d \rangle$
- Set of configurations: $\{\mathbf{x}_1, \dots, \mathbf{x}_n\}$
- Target functions: $\mathbf{f}^{(t)} = \langle f_1^{(t)}, \dots, f_m^{(t)} \rangle$ s.t. q constraints, $h_1(\mathbf{x}) \geq 0, \dots, h_q(\mathbf{x}) \geq 0$
- Access to k bias and/or noisy sources: S_1, \dots, S_k , and $\mathbf{f}^{(S_j)} = \langle f_1^{(S_j)}, \dots, f_m^{(S_j)} \rangle$
- $\mathbf{f}^{(S_j)}$ approximates the target $\mathbf{f}^{(t)}$ with variable bias $\delta_j(\mathbf{x}) = \mathbf{f}^{(t)}(\mathbf{x}) - \mathbf{f}^{(S_j)}(\mathbf{x})$
- Intervention cost: $\mathcal{C} = \sum_{e=1}^{T^{(t)}} C_t(\mathbf{x}_e^{(t)}) + \sum_{j=1}^k \sum_{e=1}^{T^{(S_j)}} C_s(\mathbf{x}_e^{(s_j)})$

Our goal is to find a configuration \mathbf{x}^* in the target environment that results in Pareto-optimal performance with minimal cost, \mathcal{C} :

$$\mathbf{x}^* = \operatorname{argmax}_{\mathbf{x} \in \mathcal{X}} \left(\mathbf{f}^{(t)}(\mathbf{x}), -\mathcal{C} \right), \text{ s.t. } h_l(\mathbf{x}) \geq 0.$$

Causal Multi-information Source Optimization: Method

Causal GP (CGP)

$$f(x) \sim GP(m(x)), K(x, x')$$

$$m(x) = \mathbb{E}[Y | do(X = x)]$$

Causal prior

$$K(x, x') = K_{RBF}(x, x') + \sigma(x)\sigma(x')$$

$$\text{where, } \sigma(x) = \sqrt{\mathbb{V}[Y | do(X = x)]}$$

Variance estimated
from observational data

Causal multi-source multi-output GP

$$\mathbf{f}^{(S_j)} \sim CGP(\mathbf{m}^{(S_j)}(x), K^{(S_j)}([x, i], [x', j]))$$

$$\mathbf{f}^{(t)} = \rho \mathbf{f}^{(S_j)} + \delta(x)$$

ρ is a scaling factor, aligns the $\mathbf{f}^{(S_j)}$ and $\mathbf{f}^{(t)}$

$$\delta(x) \sim CGP(\mathbf{m}_\delta(x), K_\delta([x, i], [x', j]))$$

Captures discrepancy between $\mathbf{f}^{(S_j)}$ and $\mathbf{f}^{(t)}$

$$K[x, i], [x', j] = K_{inputs}(x, x') \cdot K_{objectives}(i, j) + \sigma(x)\sigma(x')$$

Covariance between two datapoints.
Inputs x and x' , and objectives i and j



**“Ideas don’t come out fully formed,
they only become clear as you work
on them.”**

- Mark Zuckerberg



- Your first paper is not your last paper.
- Don't try to solve all research questions in a single paper.
- Research is an iterative process.

Lessons Learned

- Establishing clear and regular communication channels.
- Defined Roles and Responsibilities.
- Successful collaboration often requires adaptability to different working styles.

Trust your expertise—reviews are guidance, not gospel.



Awesome[∞]Collaborators

Pooyan Jamshidi



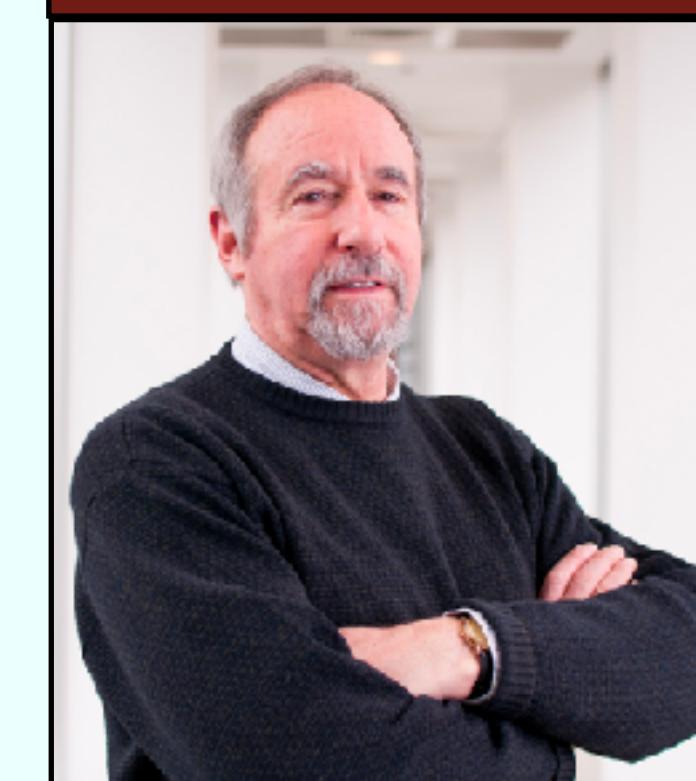
Bradley Schmerl



Jason M. O'Kane



David Garlan



Sonam Kharade



Javier Cámarra

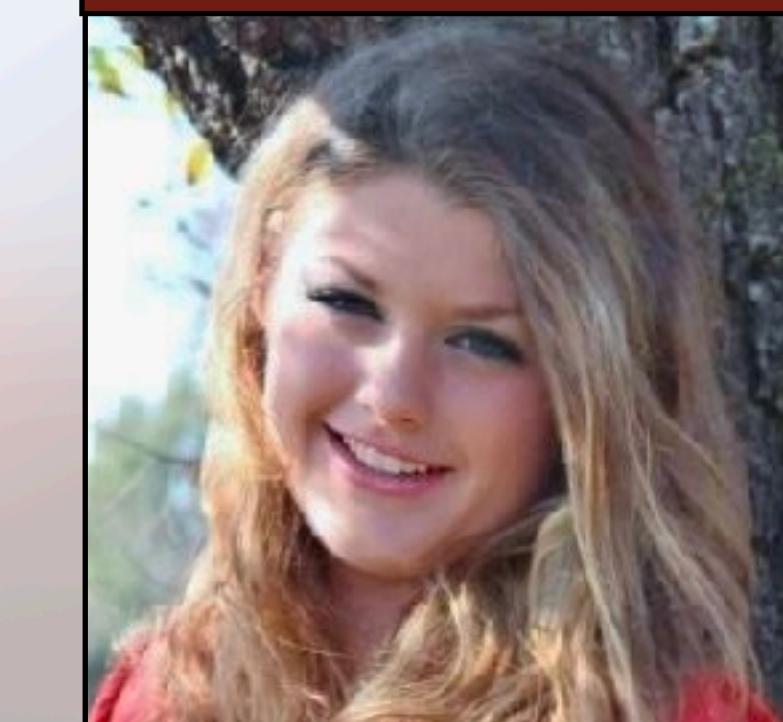


Ellen C. Czaplinski



Jet Propulsion Laboratory
California Institute of Technology

Katherine A. Dzurilla



Jet Propulsion Laboratory
California Institute of Technology

Thank you!

