

Transfer Learning for Performance Analysis of Highly-Configurable Systems

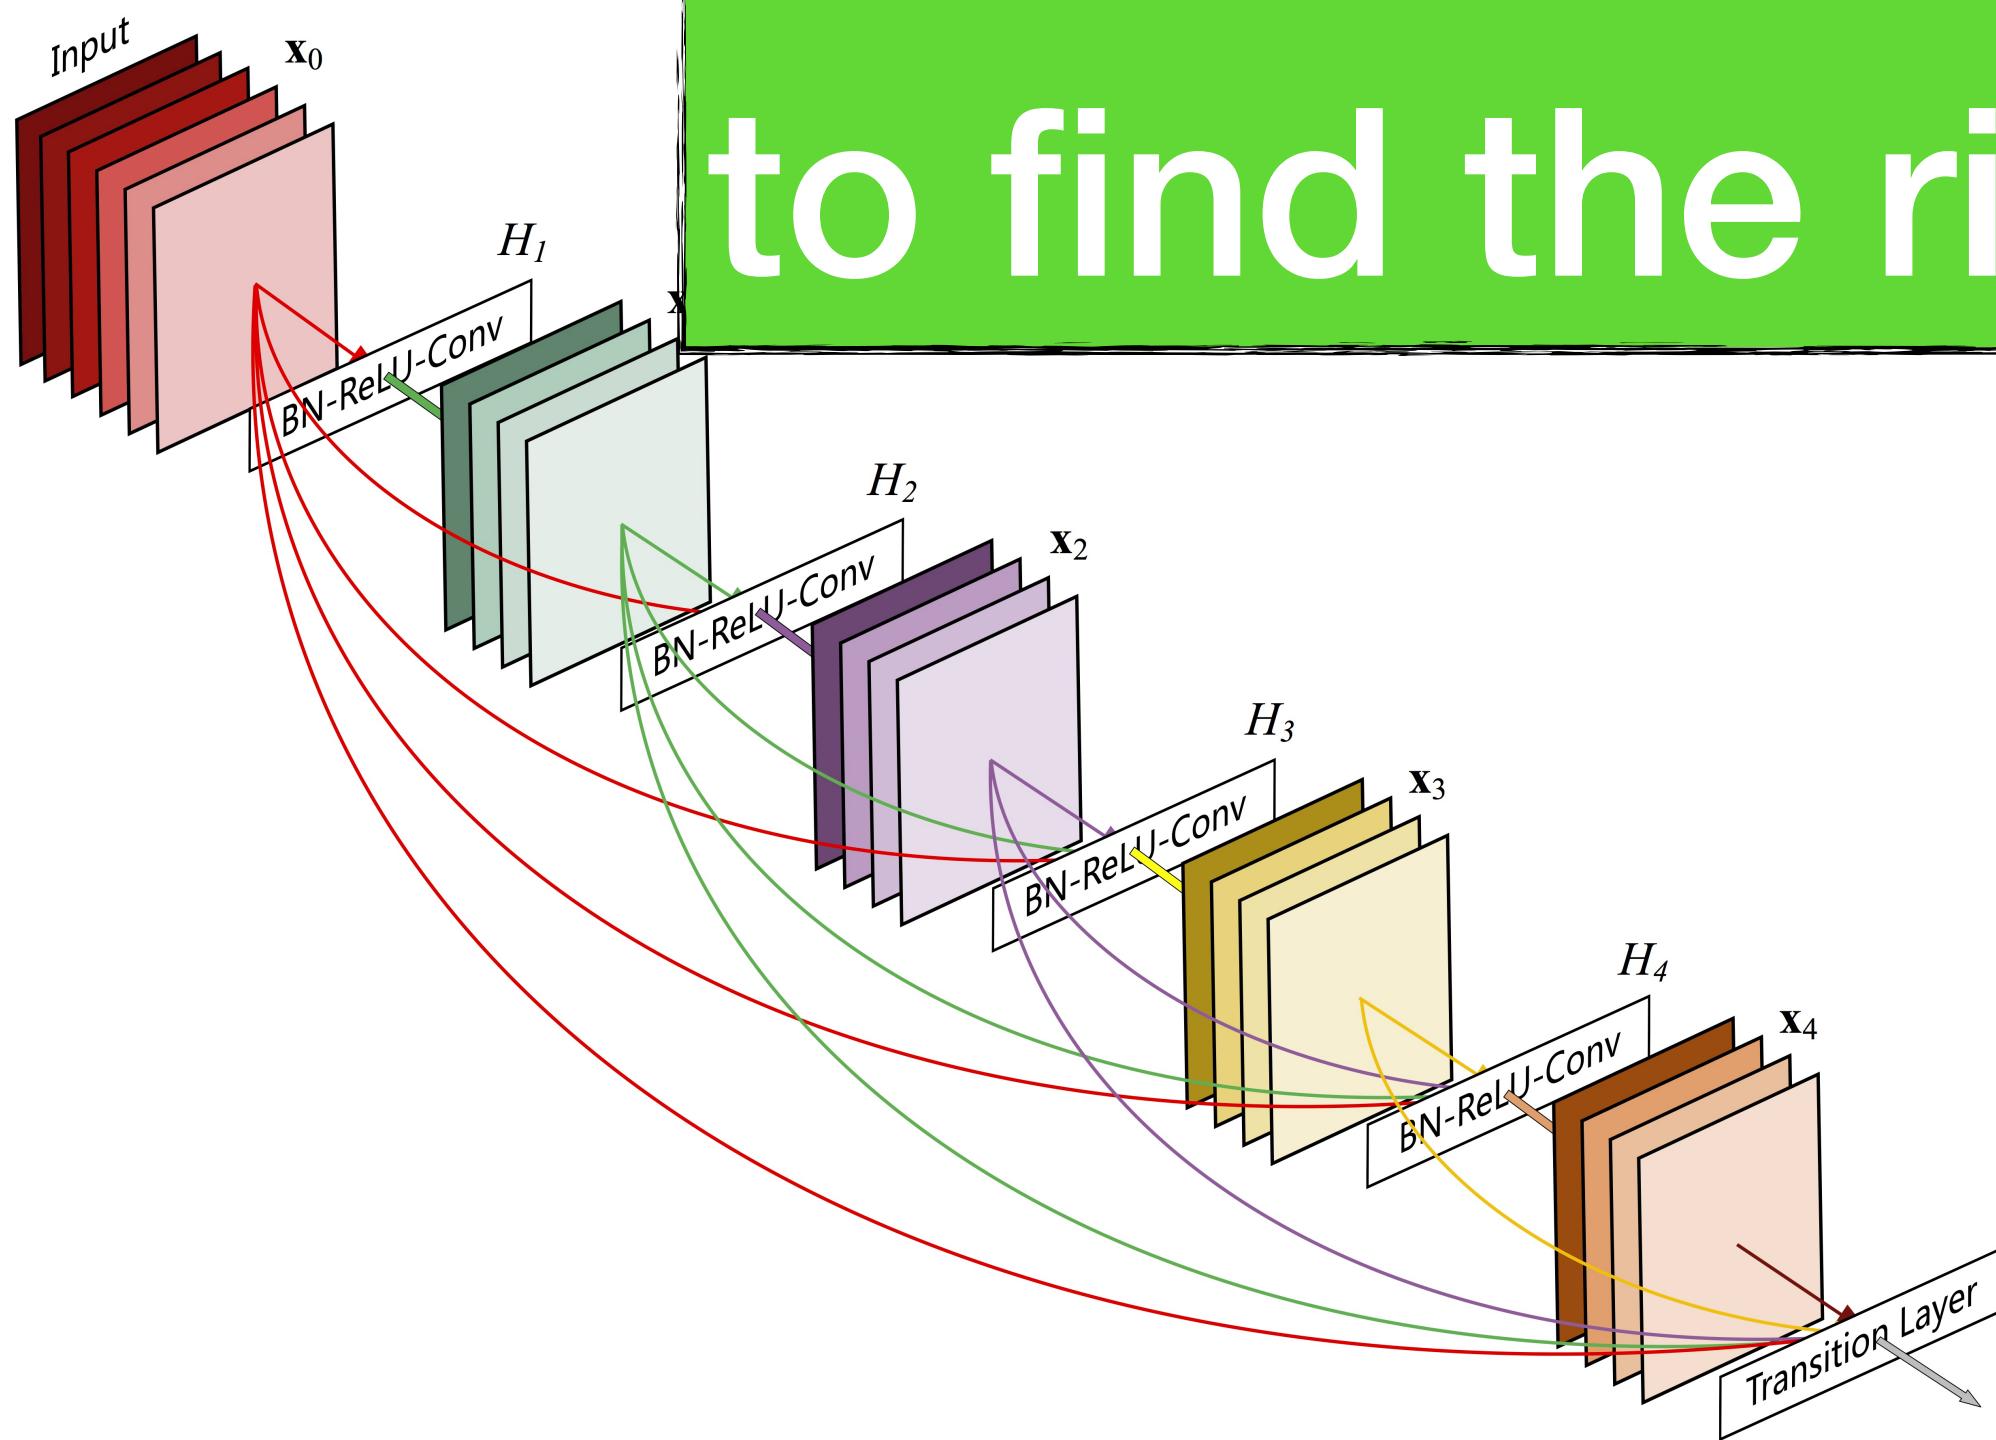
Pooyan Jamshidi
University of South Carolina



[@pooyanjamshidi](https://twitter.com/pooyanjamshidi)



Goal: Enable developers/users
to find the right quality tradeoff



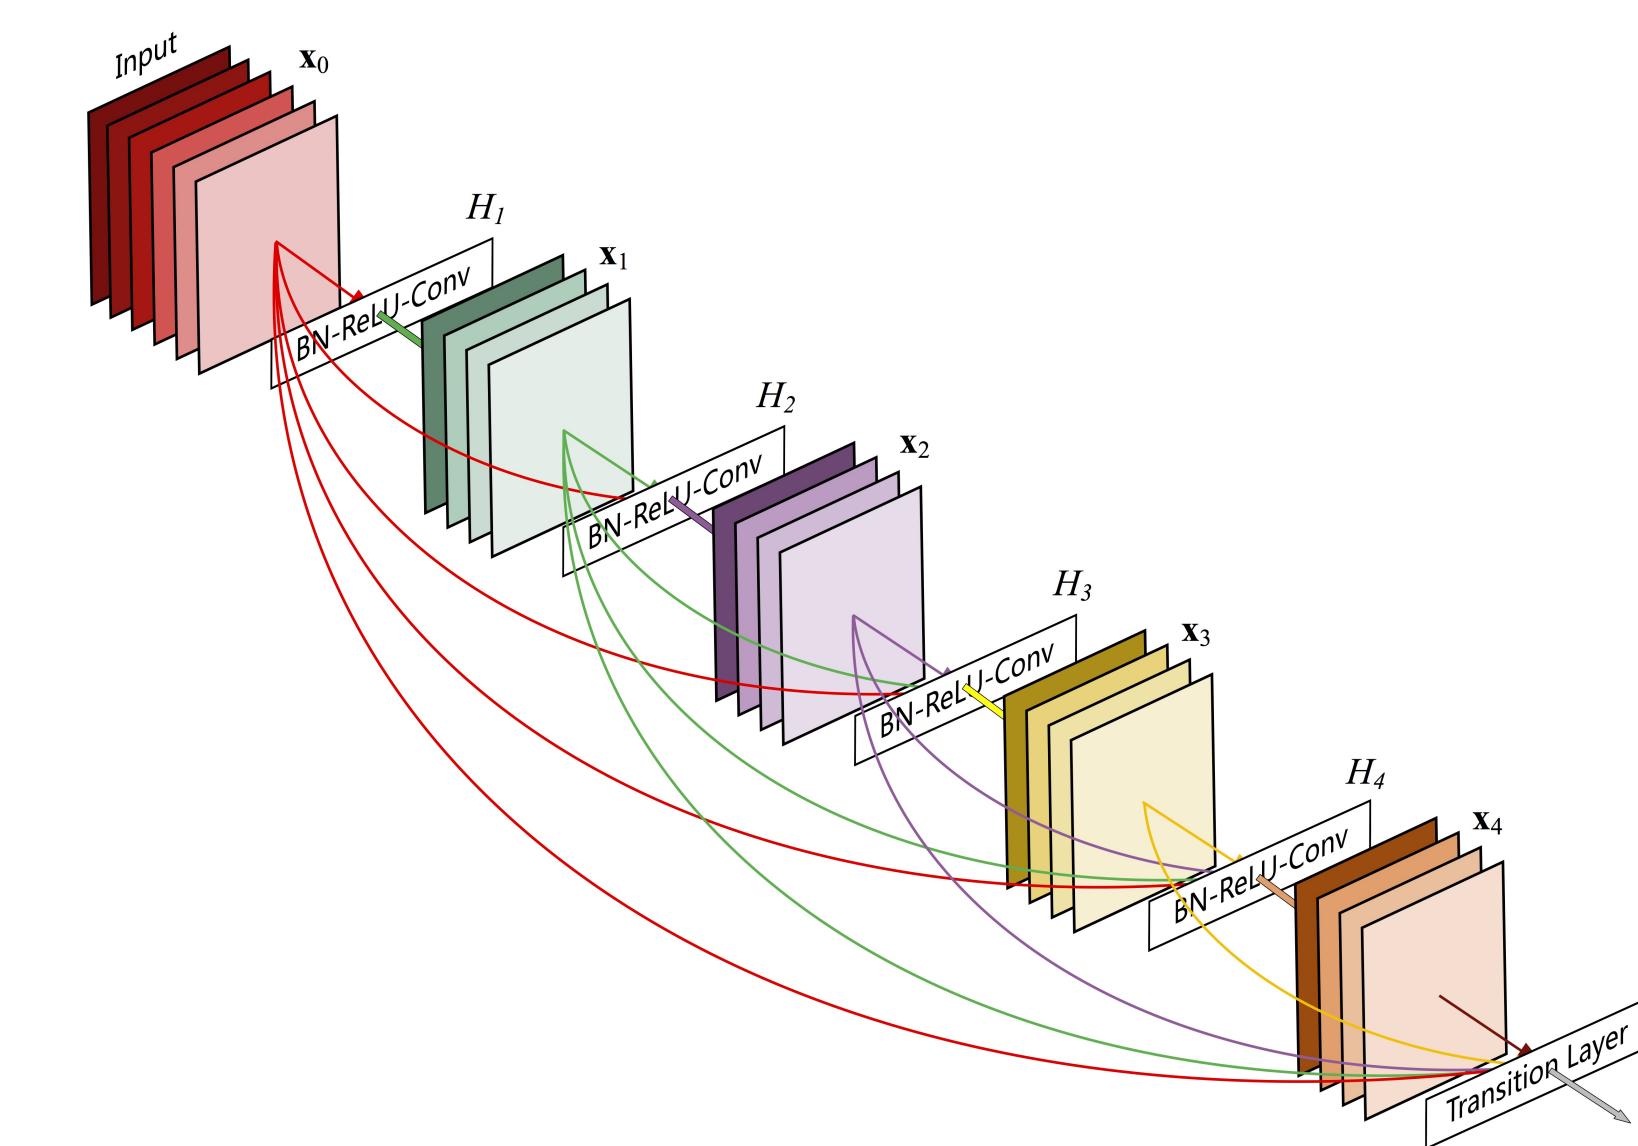
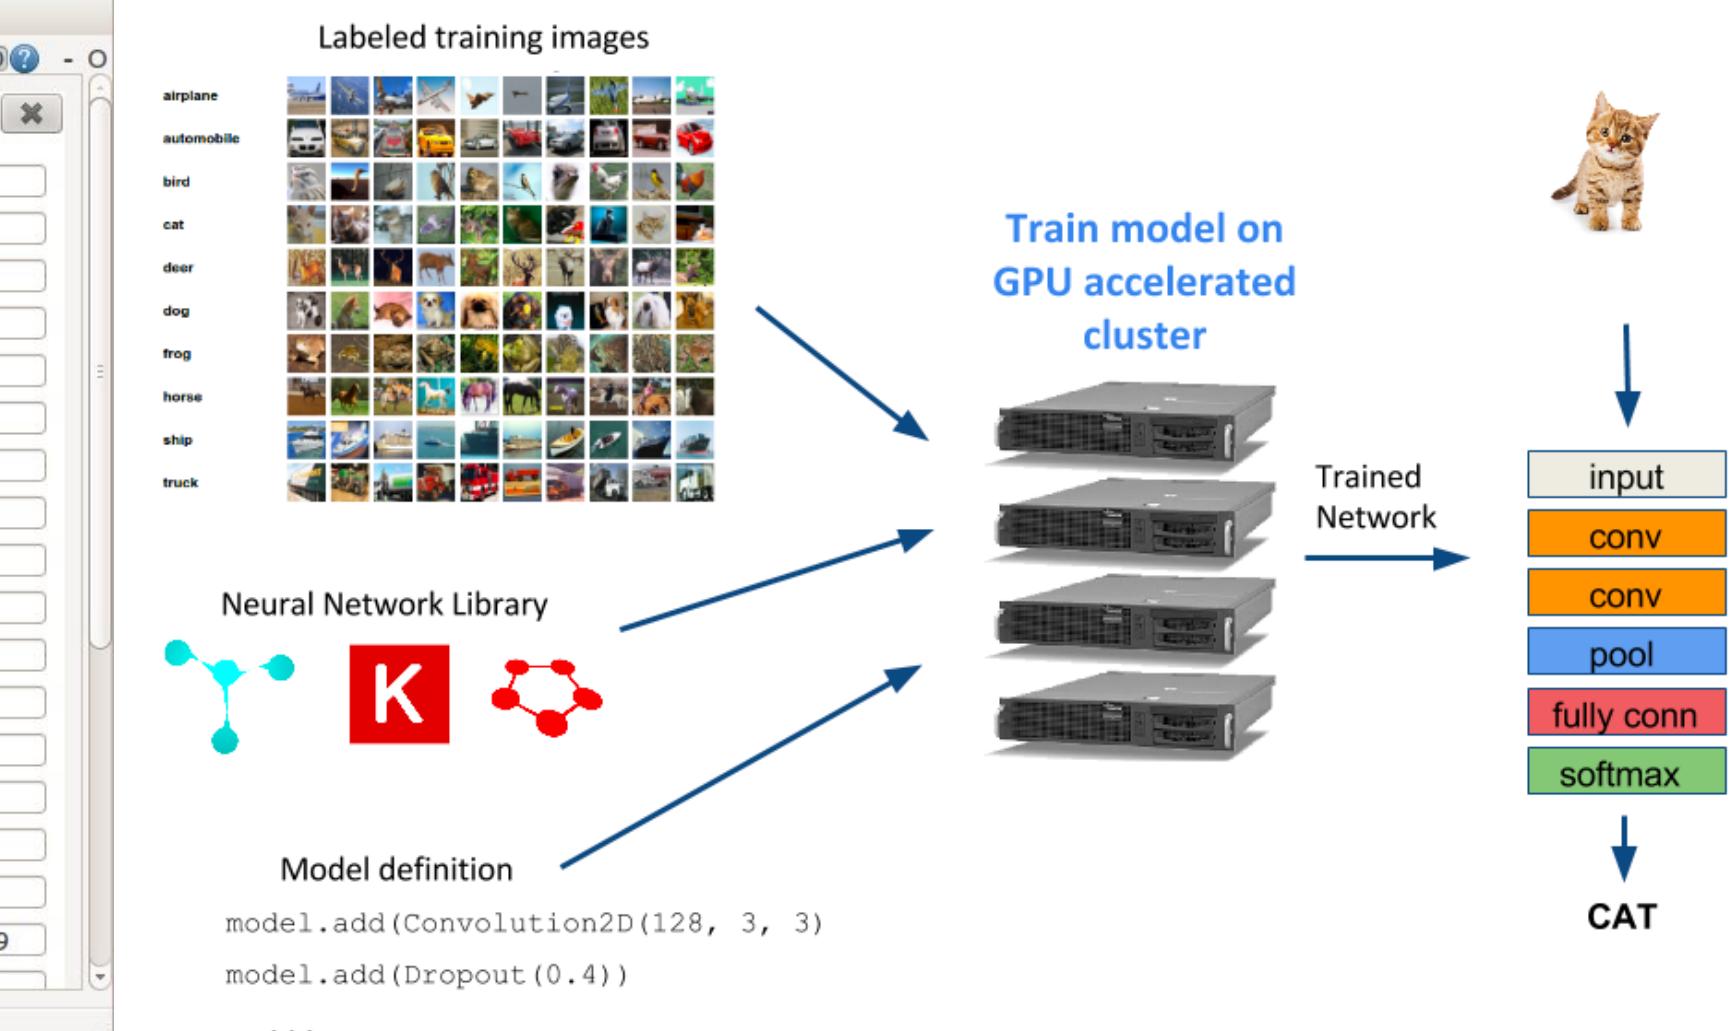
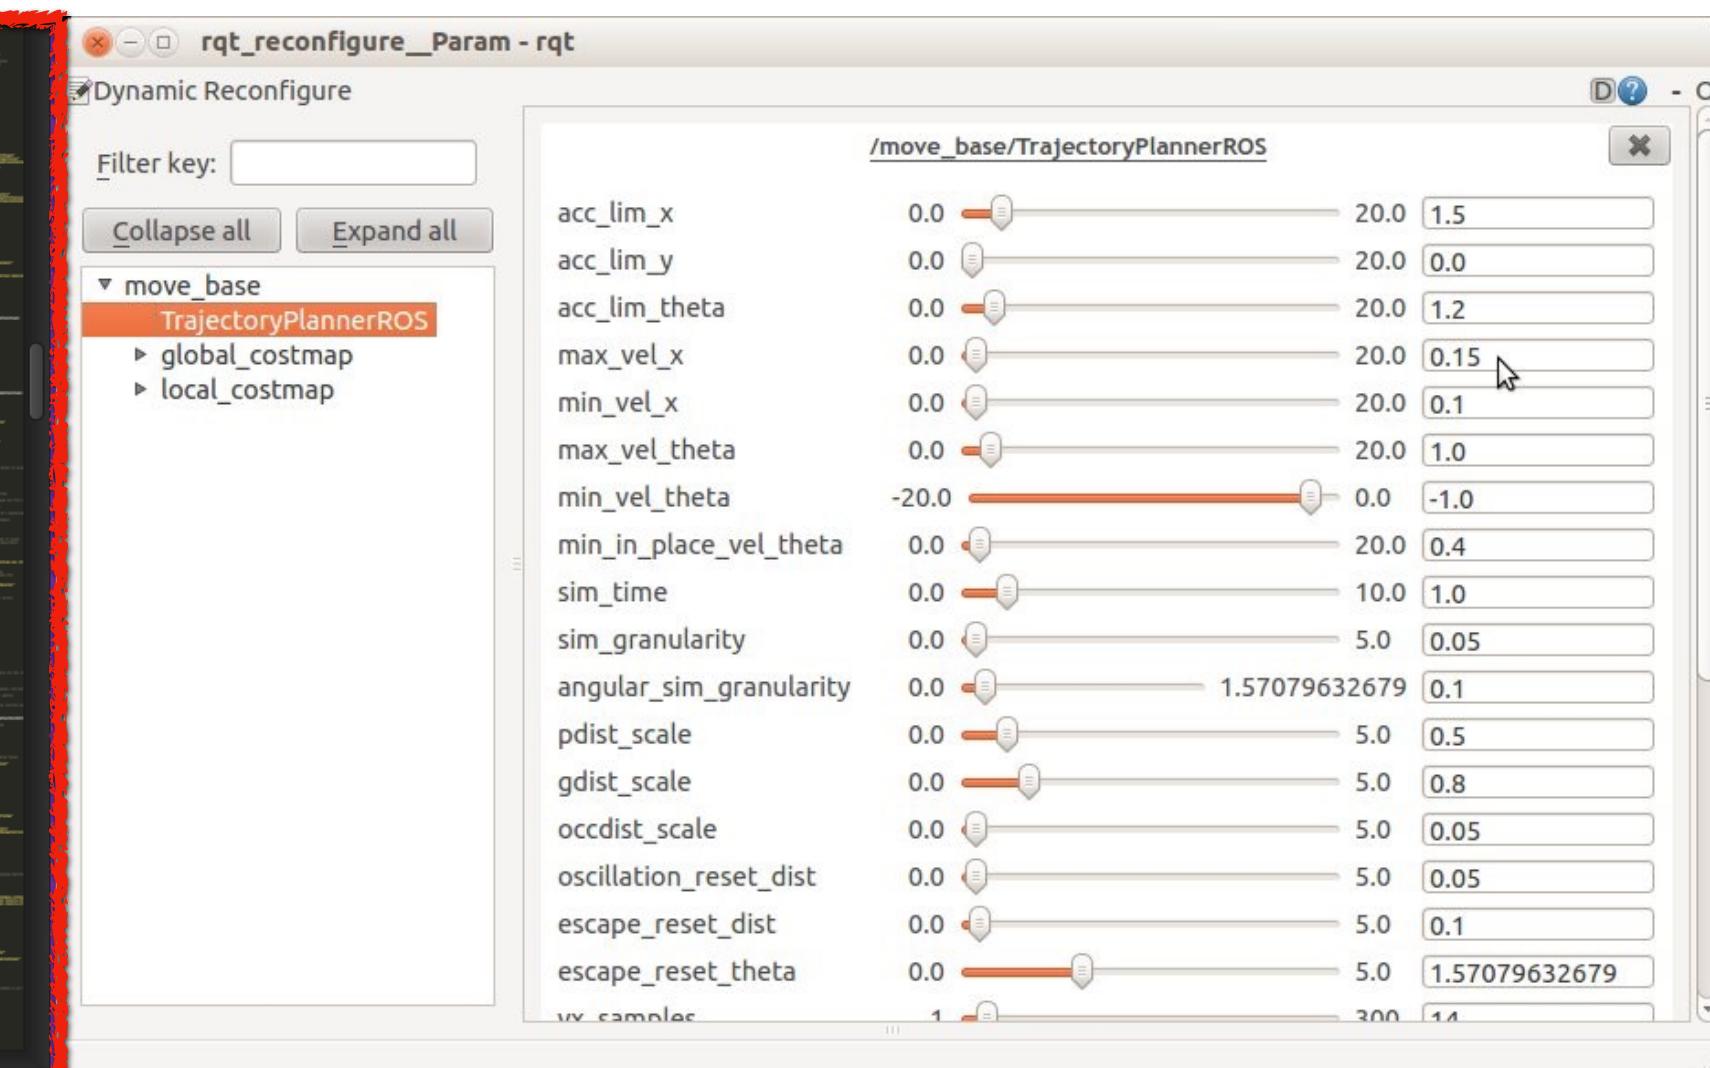
built

Today's most popular systems are configurable

```

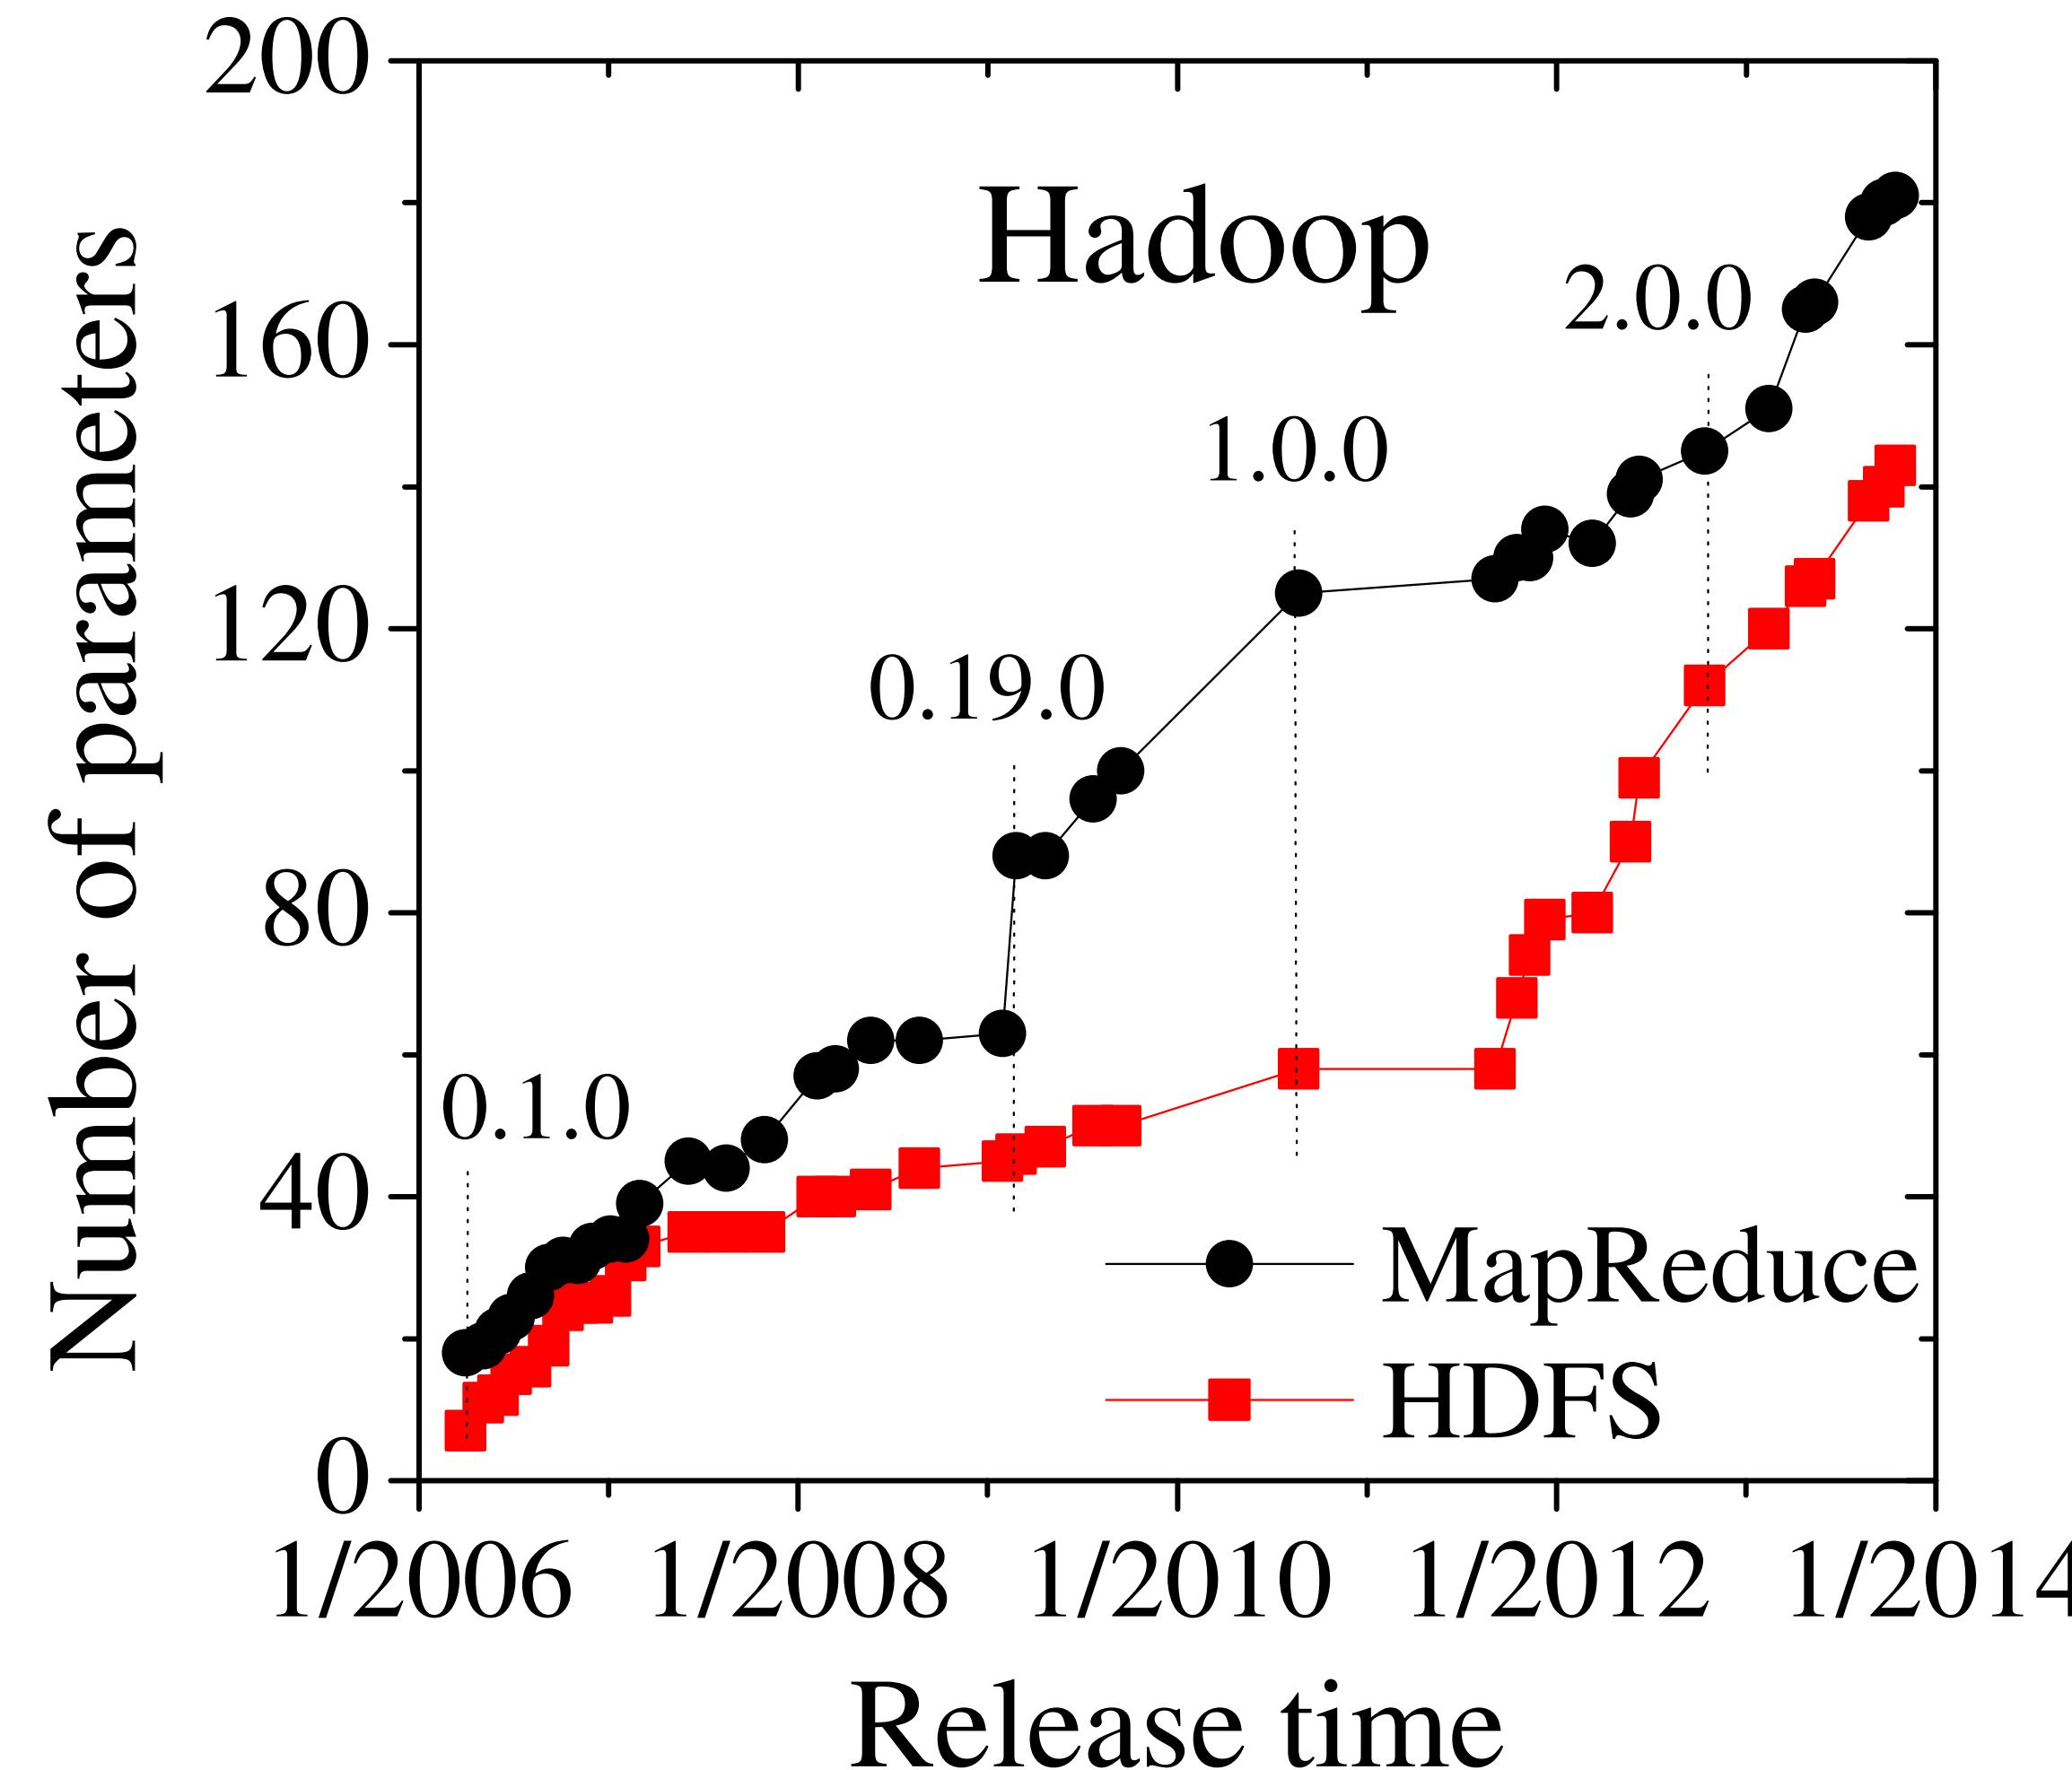
102
103 drpc.port: 3772
104 drpc.worker.threads: 64
105 drpc.max_buffer_size: 1048576
106 drpc.queue.size: 128
107 drpc.invocations.port: 3773
108 drpc.invocations.threads: 64
109 drpc.request.timeout.secs: 600
110 drpc.childopts: "-Xmx768m"
111 drpc.http.port: 3774
112 drpc.https.port: -1
113 drpc.https.keystore.password: ""
114 drpc.https.keystore.type: "JKS"
115 drpc.http.creds.plugin: org.apache.storm.security.auth.DefaultHttpCredentialsPlugin
116 drpc.authorizer.acl.filename: "drpc-auth-acl.yaml"
117 drpc.authorizer.acl.strict: false
118
119 transactional.zookeeper.root: "/transactional"
120 transactional.zookeeper.servers: null
121 transactional.zookeeper.port: null
122
123 ## blobstore configs
124 supervisor.blobstore.class: "org.apache.storm.blobstore.NimbusBlobStore"
125 supervisor.blobstore.download.thread.count: 5
126 supervisor.blobstore.download.max_retries: 3
127 supervisor.localizer.cache.target.size.mb: 10240
128 supervisor.localizer.cleanup.interval.ms: 600000
129

```

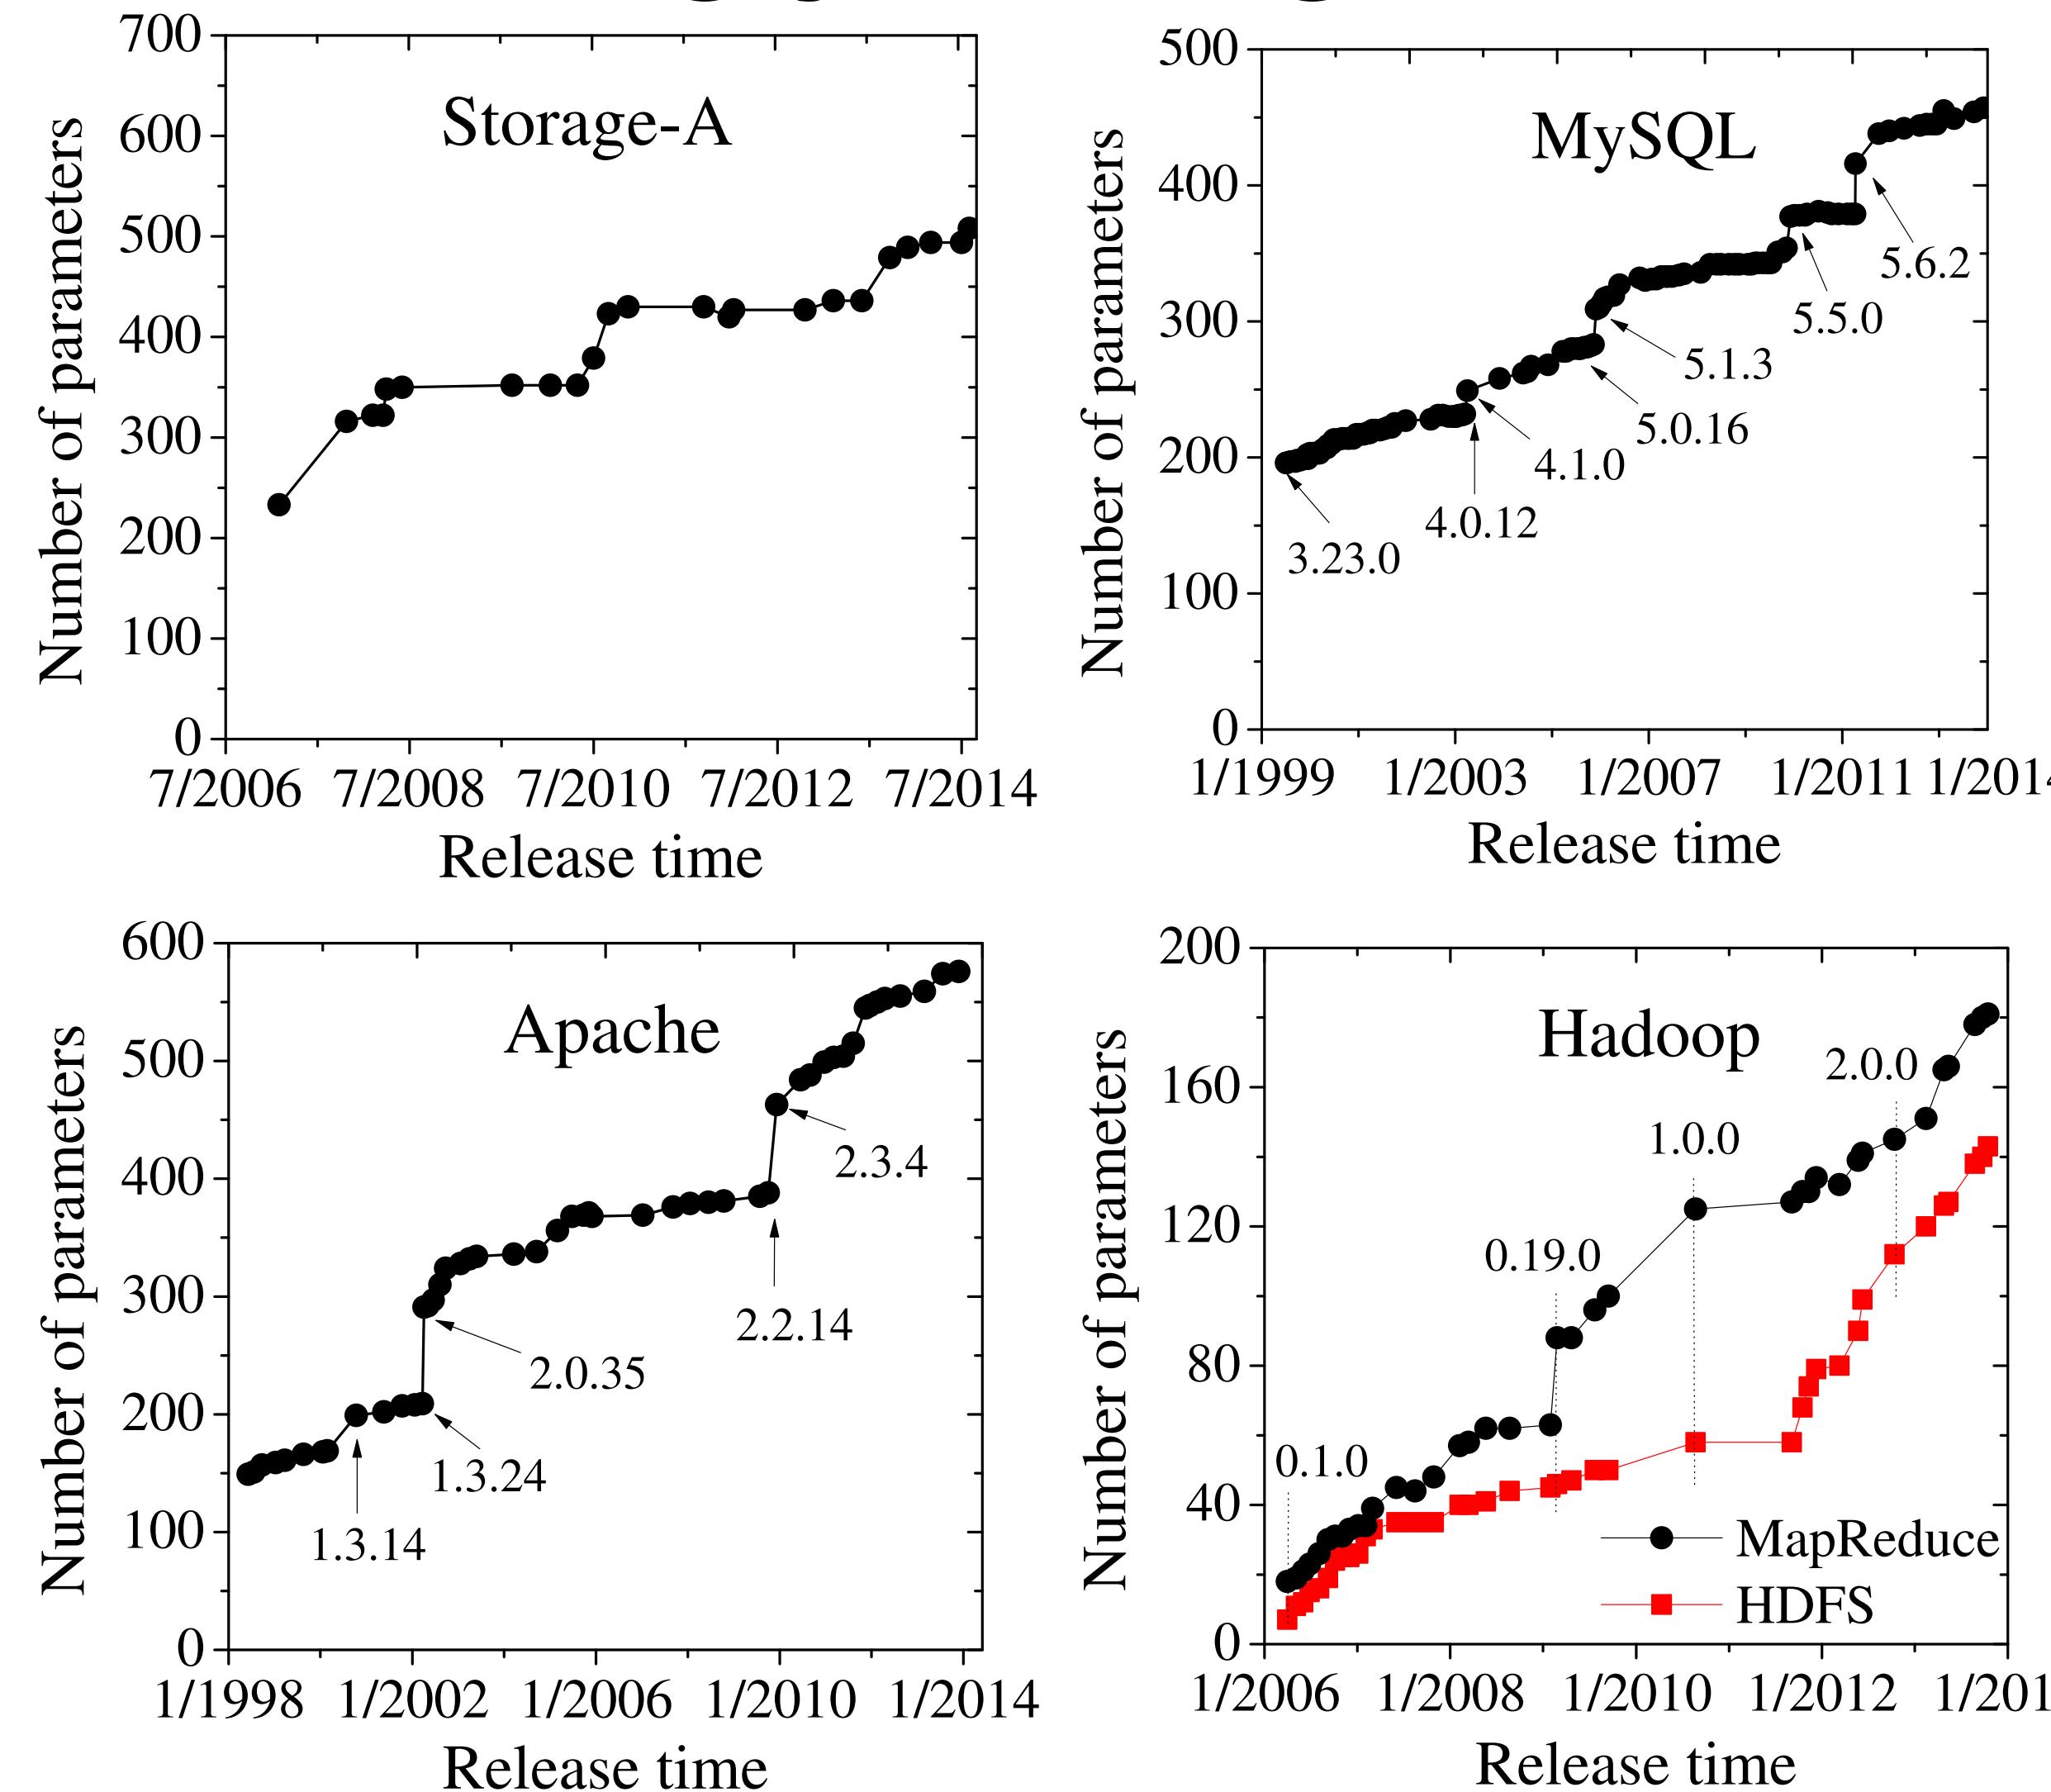


```
102  
103 drpc.port: 3772  
104 drpc.worker.threads: 64  
105 drpc.max_buffer_size: 1048576  
106 drpc.queue.size: 128  
107 drpc.invocations.port: 3773  
108 drpc.invocations.threads: 64  
109 drpc.request.timeout.secs: 600  
110 drpc.childopts: "-Xmx768m"  
111 drpc.http.port: 3774  
112 drpc.https.port: -1  
113 drpc.https.keystore.password: ""  
114 drpc.https.keystore.type: "JKS"  
115 drpc.http.creds.plugin: org.apache.storm.security.auth.DefaultHttpCredentialsPlugin  
116 drpc.authorizer.acl.filename: "drpc-auth-acl.yaml"  
117 drpc.authorizer.acl.strict: false  
118  
119 transactional.zookeeper.root: "/transactional"  
120 transactional.zookeeper.servers: null  
121 transactional.zookeeper.port: null  
122  
123 ## blobstore configs  
124 supervisor.blobstore.class: "org.apache.storm.blobstore.NimbusBlobStore"  
125 supervisor.blobstore.download.thread.count: 5  
126 supervisor.blobstore.download.max_retries: 3  
127 supervisor.localizer.cache.target.size.mb: 10240  
128 supervisor.localizer.cleanup.interval.ms: 600000  
129
```

Empirical observations confirm that systems are becoming increasingly configurable



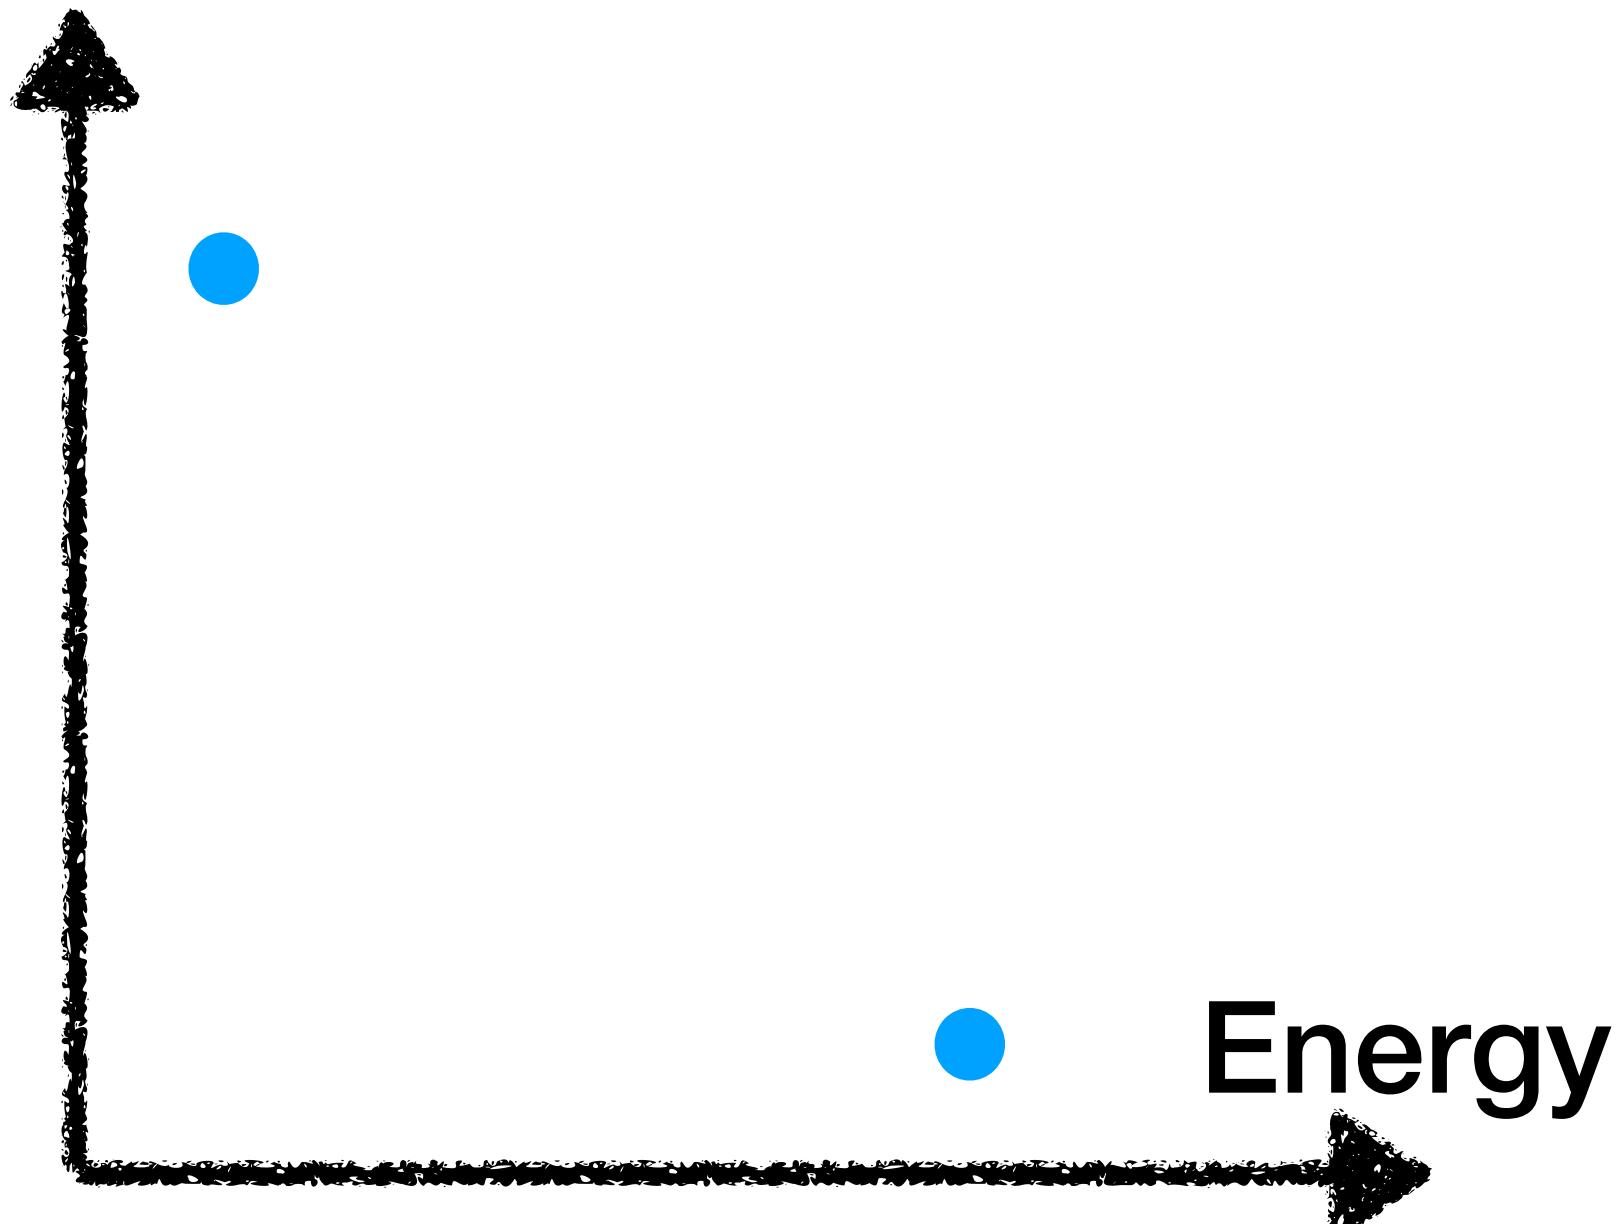
Empirical observations confirm that systems are becoming increasingly configurable



Configurations determine the performance behavior

```
void Parrot_setenv( . . . name, . . . value){  
#ifdef PARROT HAS SETENV  
    my_setenv(name, value, 1);  
#else  
    int name_len=strlen(name);  
    int val_len=strlen(value);  
    char* envs=glob_env;  
    if(envs==NULL){  
        return;  
    }  
    strcpy(envs,name);  
    strcpy(envs+name_len,"=");  
    strcpy(envs+name_len + 1,value);  
    putenv(envs);  
#endif  
}
```

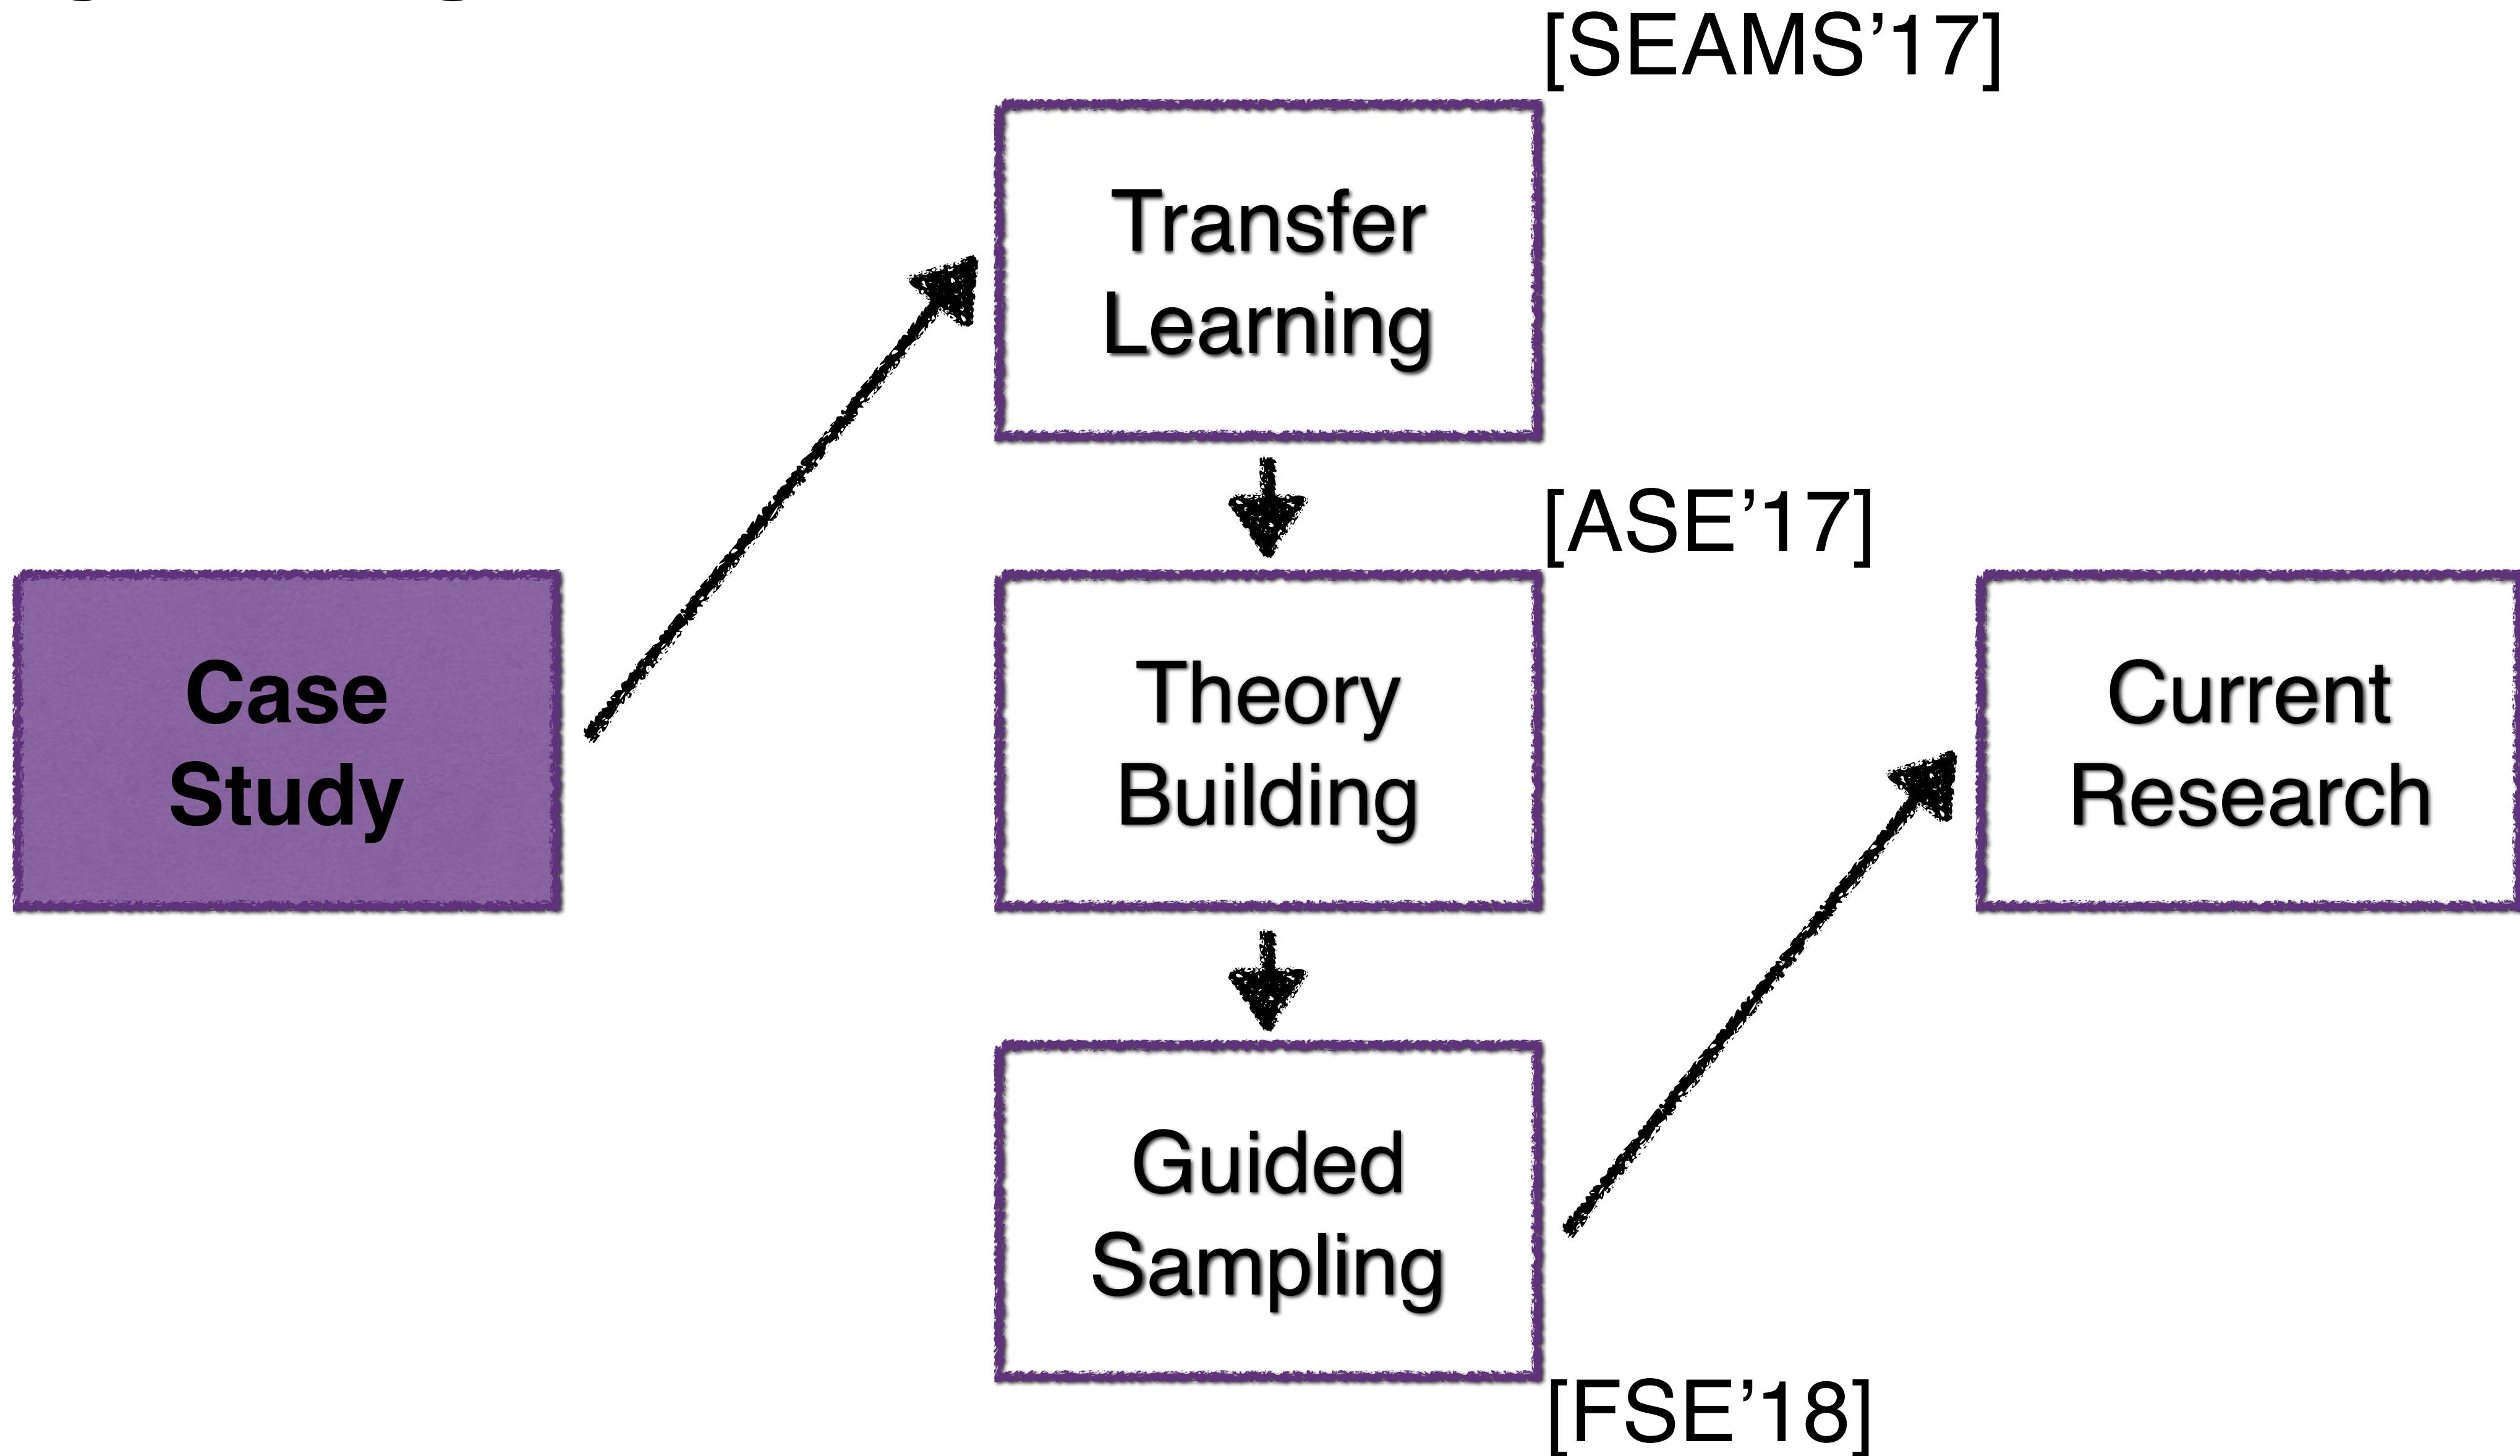
Speed



How do we **understand** performance behavior of
real-world highly-configurable systems that **scale** well...

... and enable developers/users to **reason** about
qualities (performance, energy) and to make **tradeoff?**

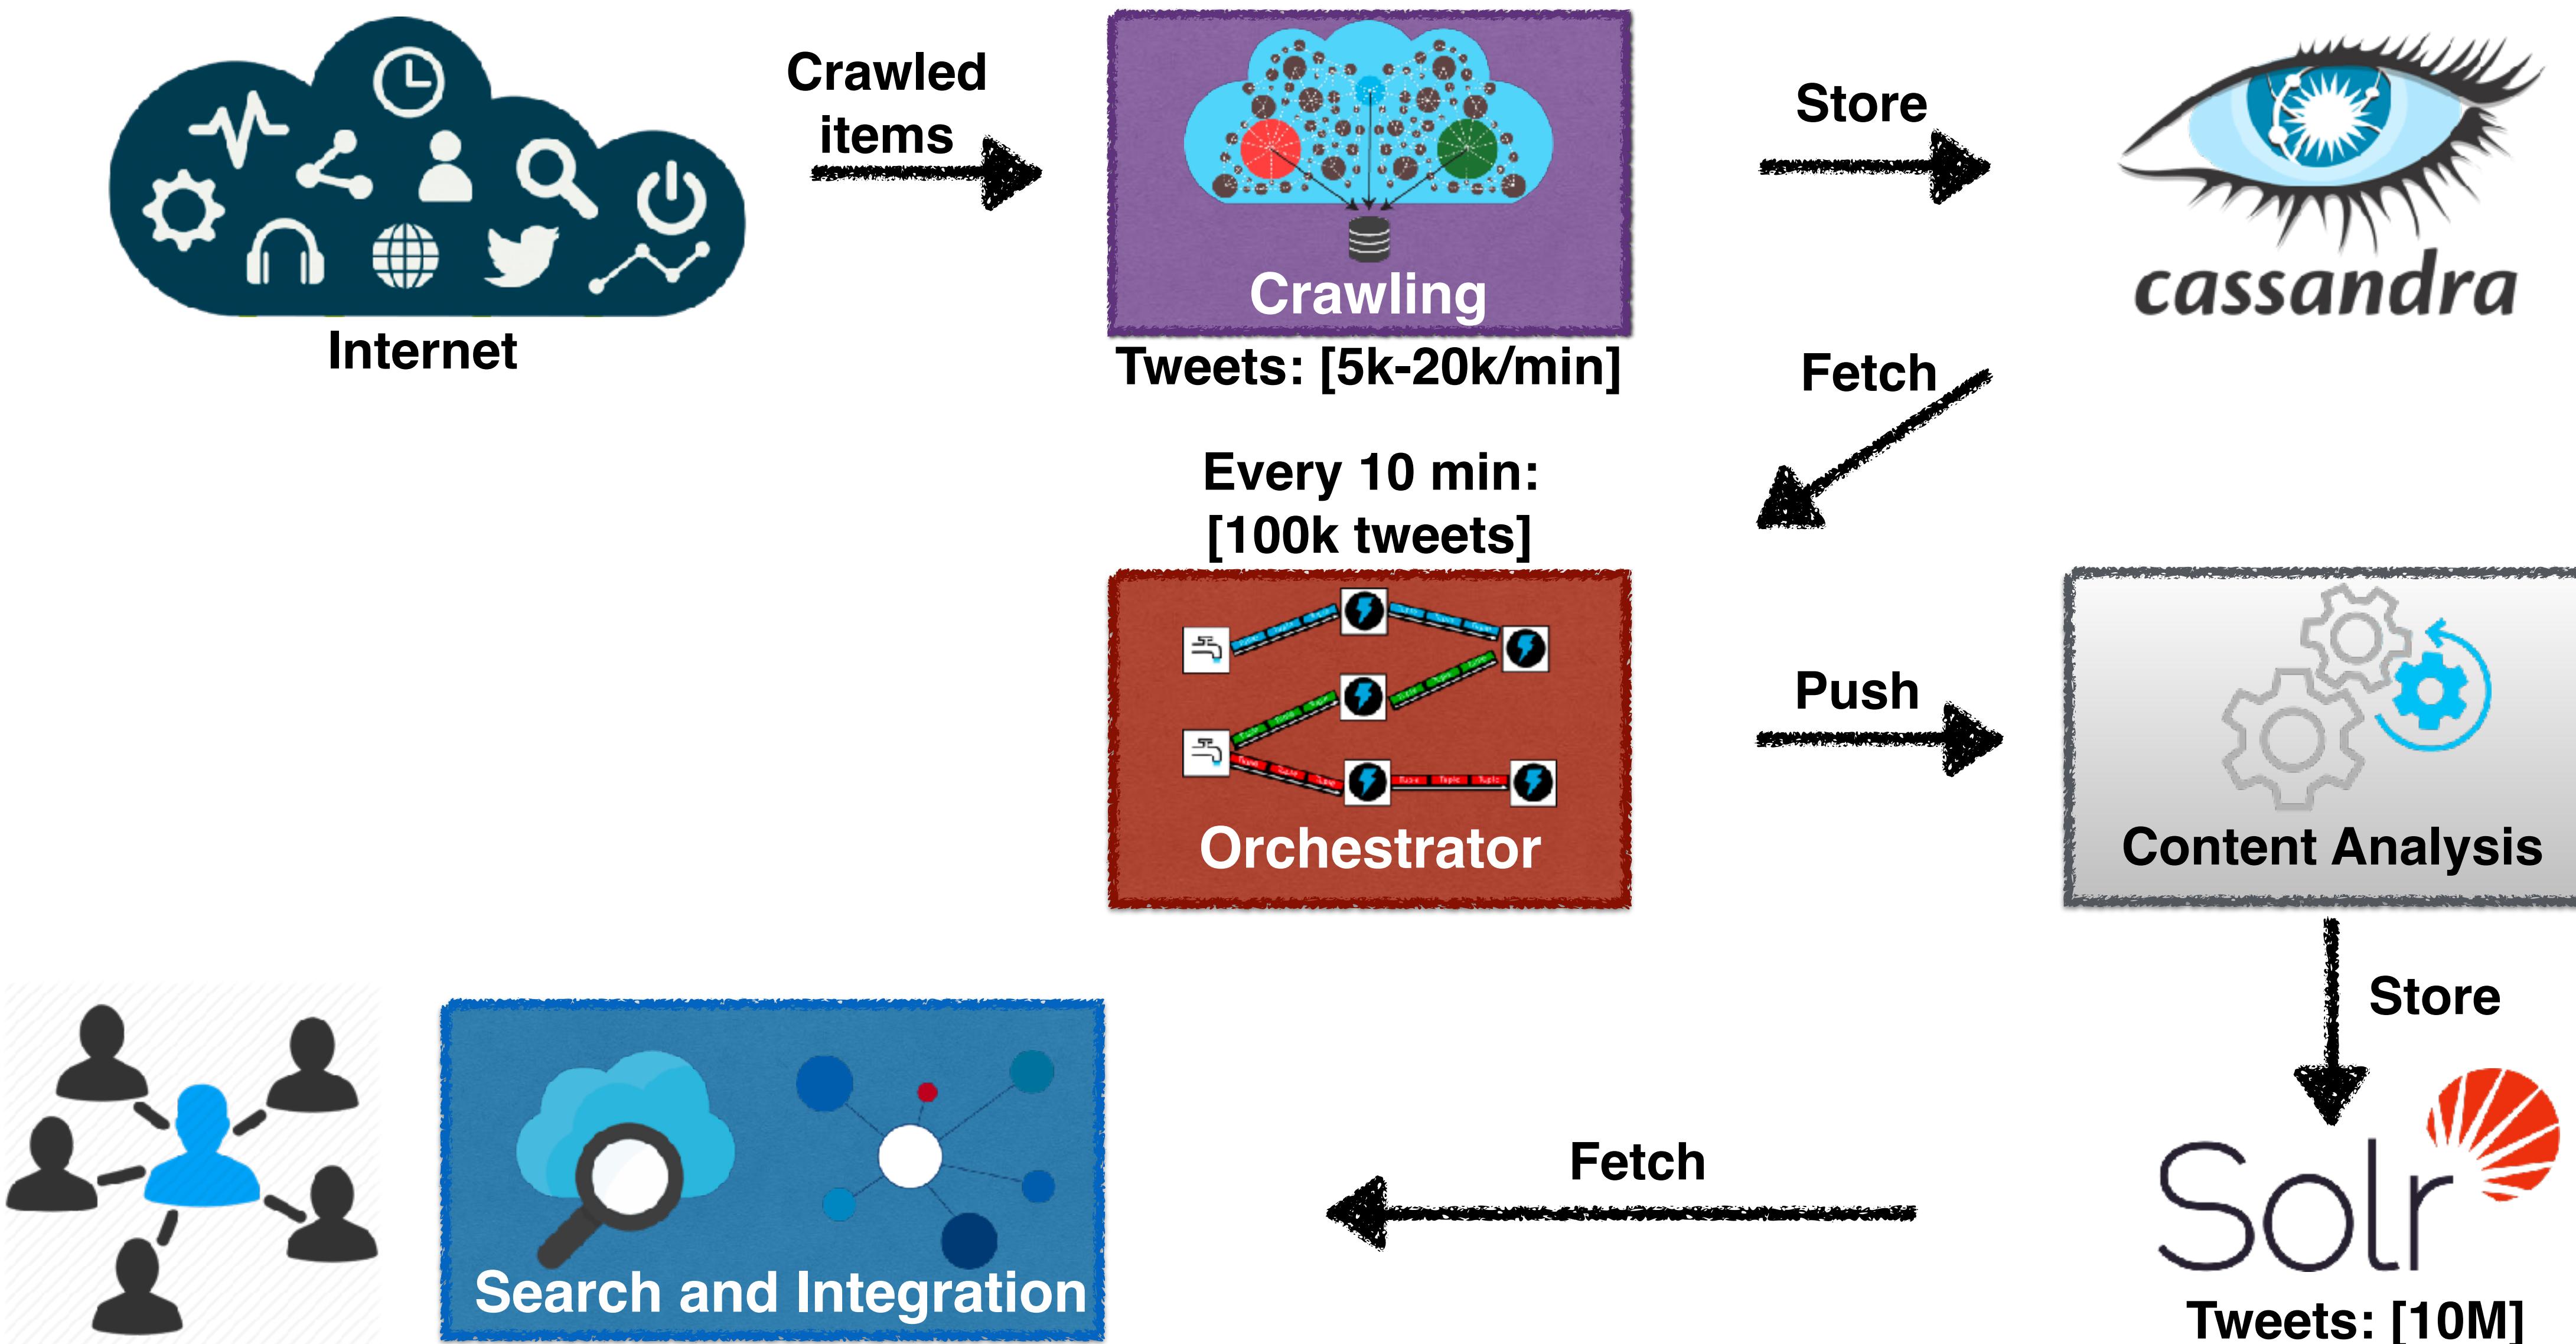
Outline



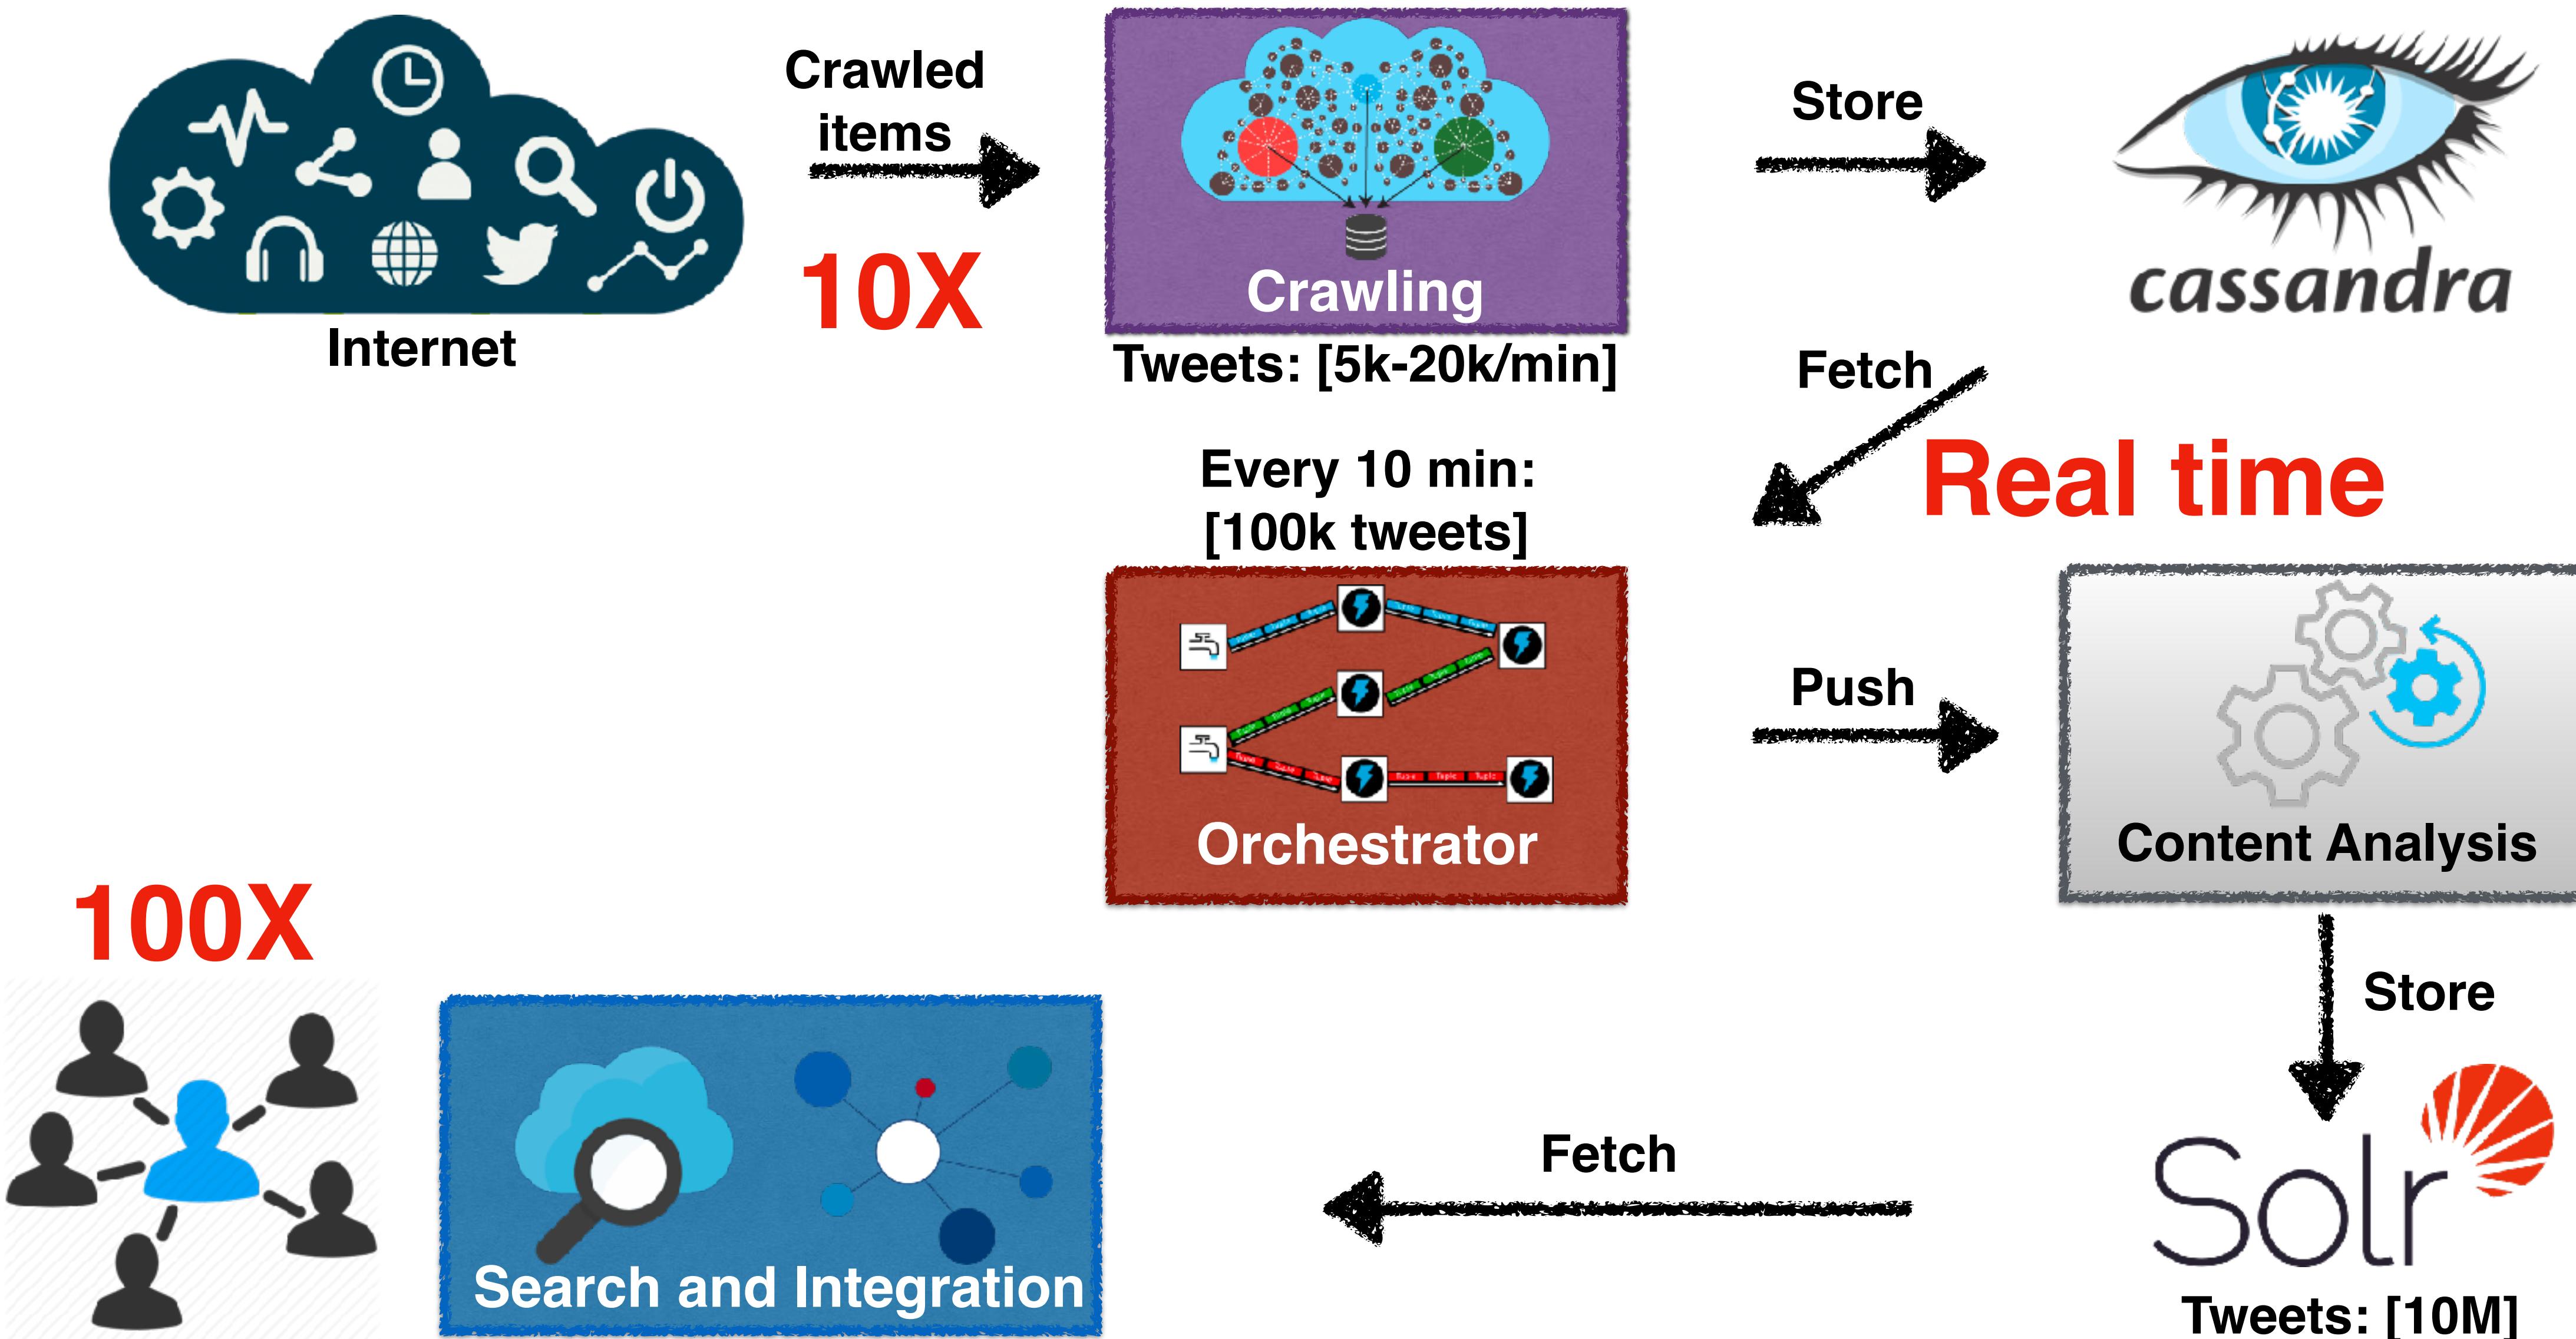
SocialSensor

- Identifying trending topics
- Identifying user defined topics
- Social media search

SocialSensor



Challenges



How can we gain a better performance without using more resources?

Let's try out different system configurations!

Opportunity: Data processing engines in the pipeline were all configurable



> 100



STORM

> 100

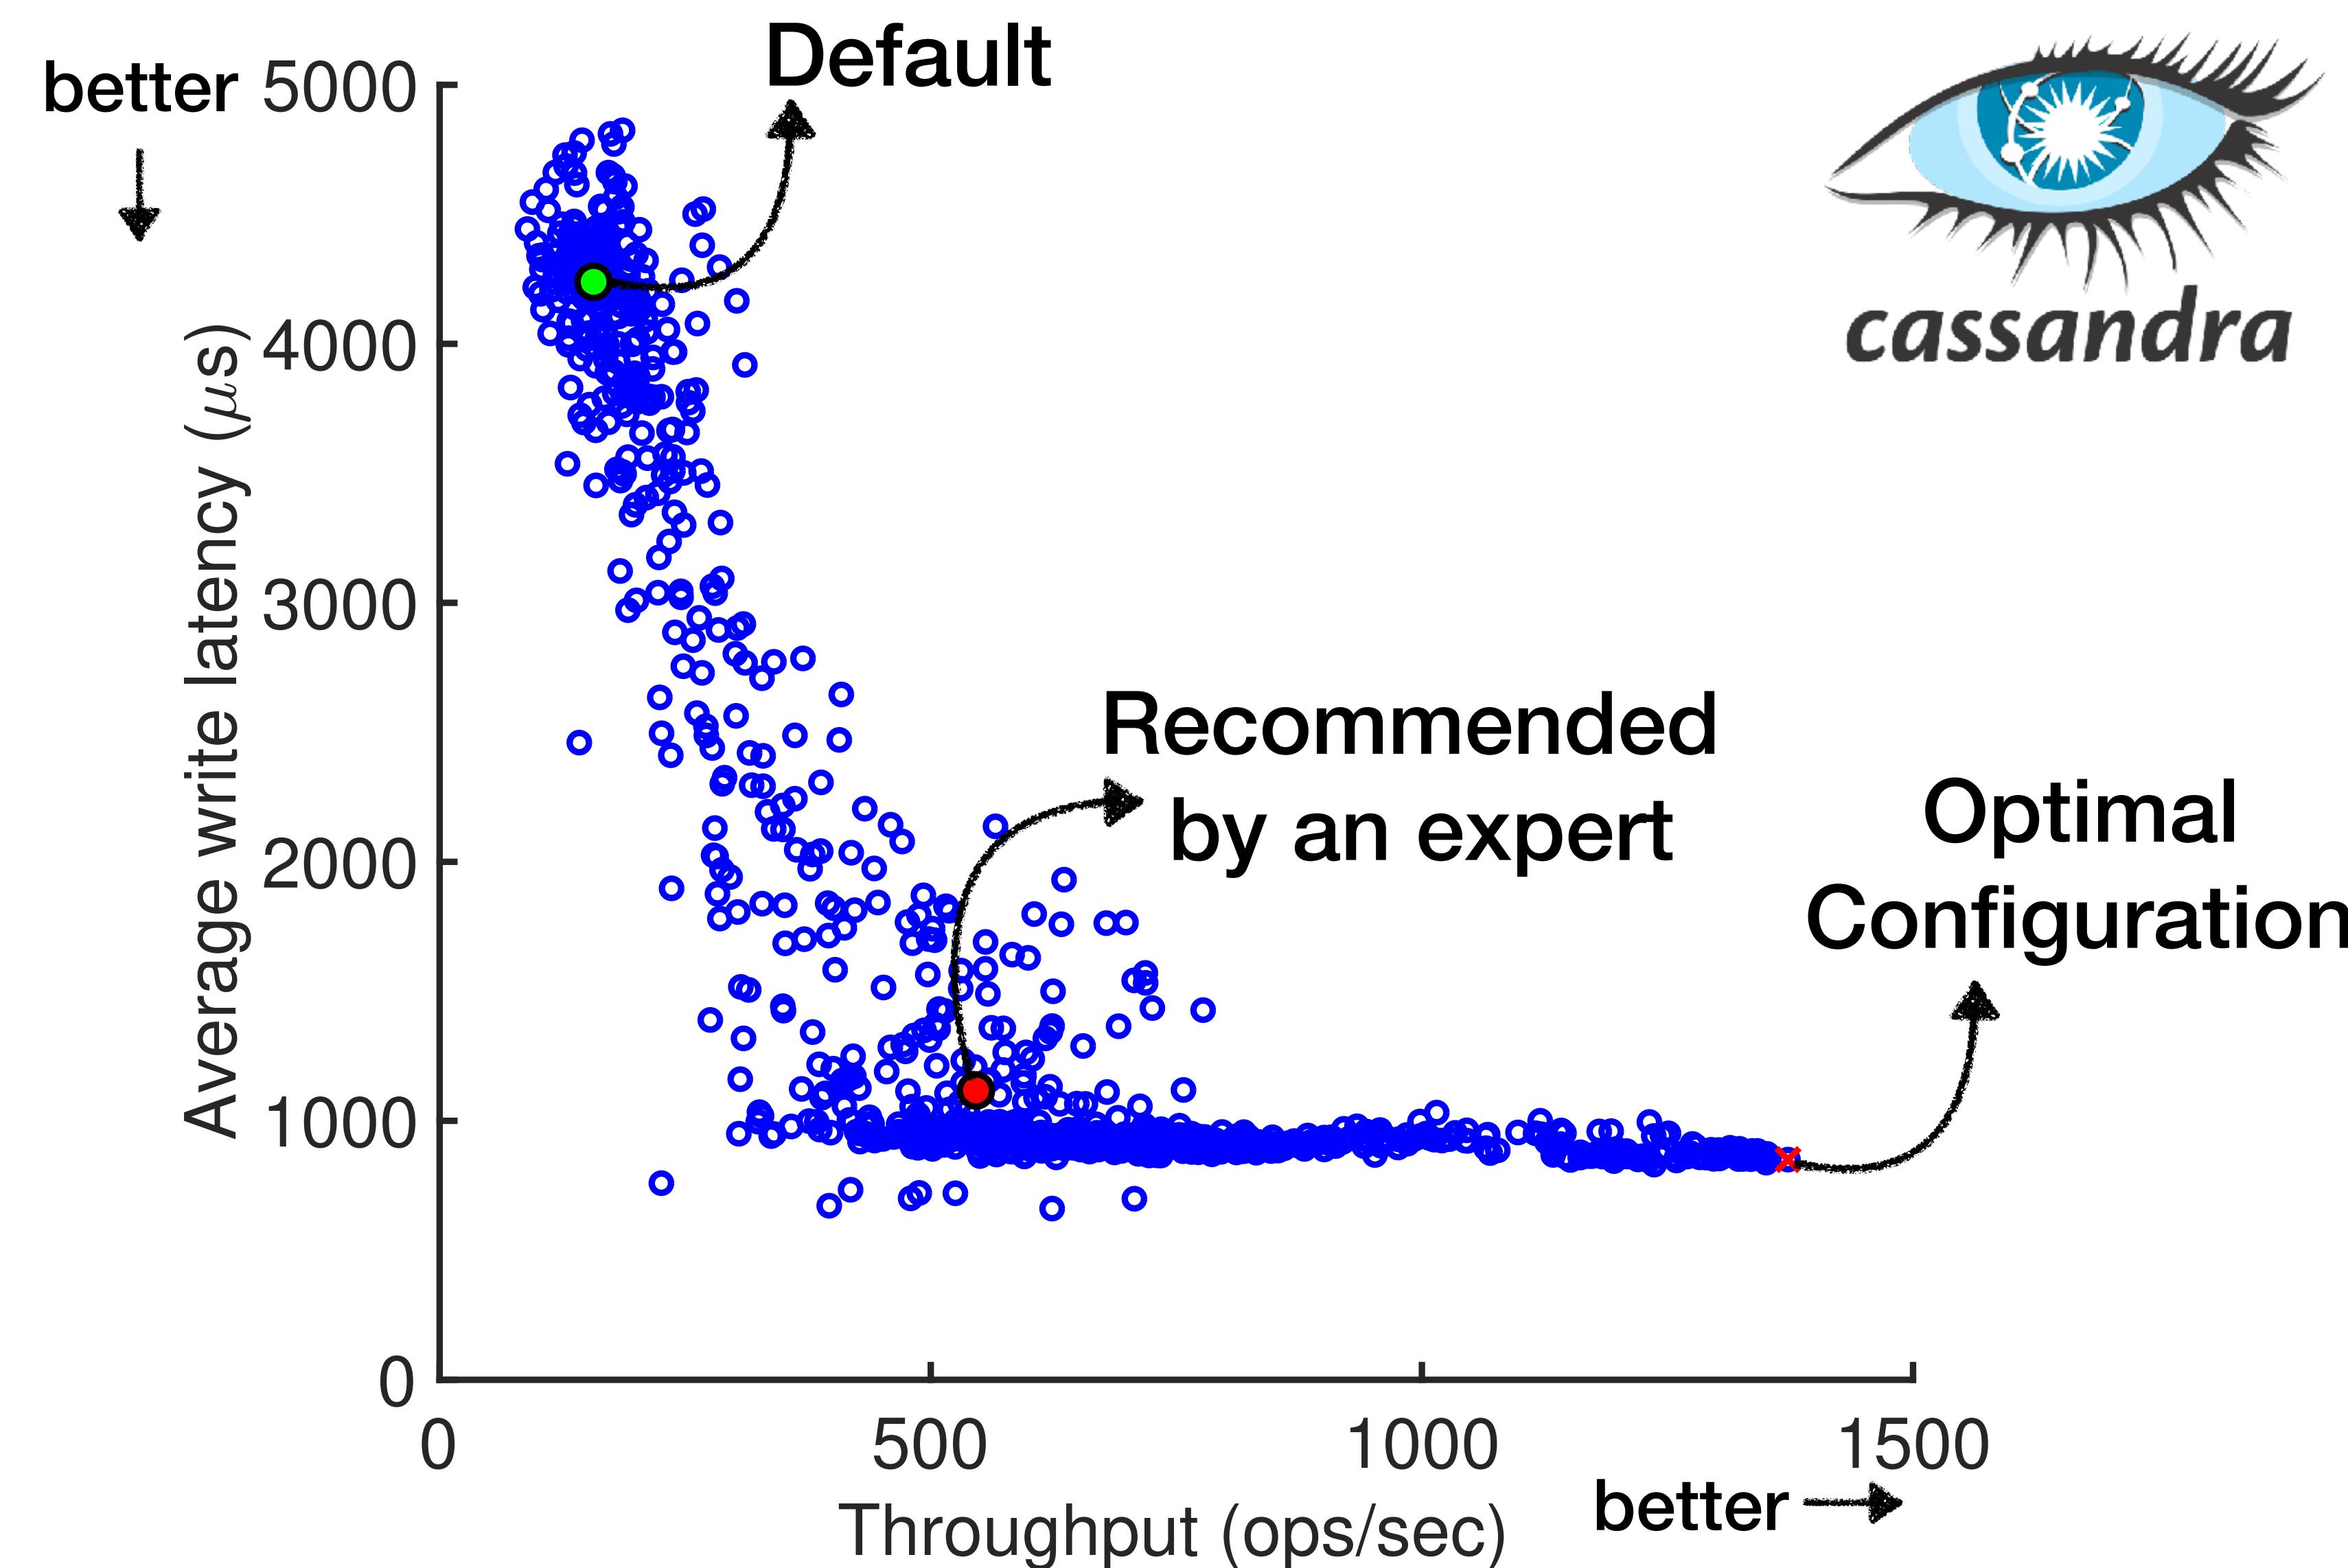
2^{300}



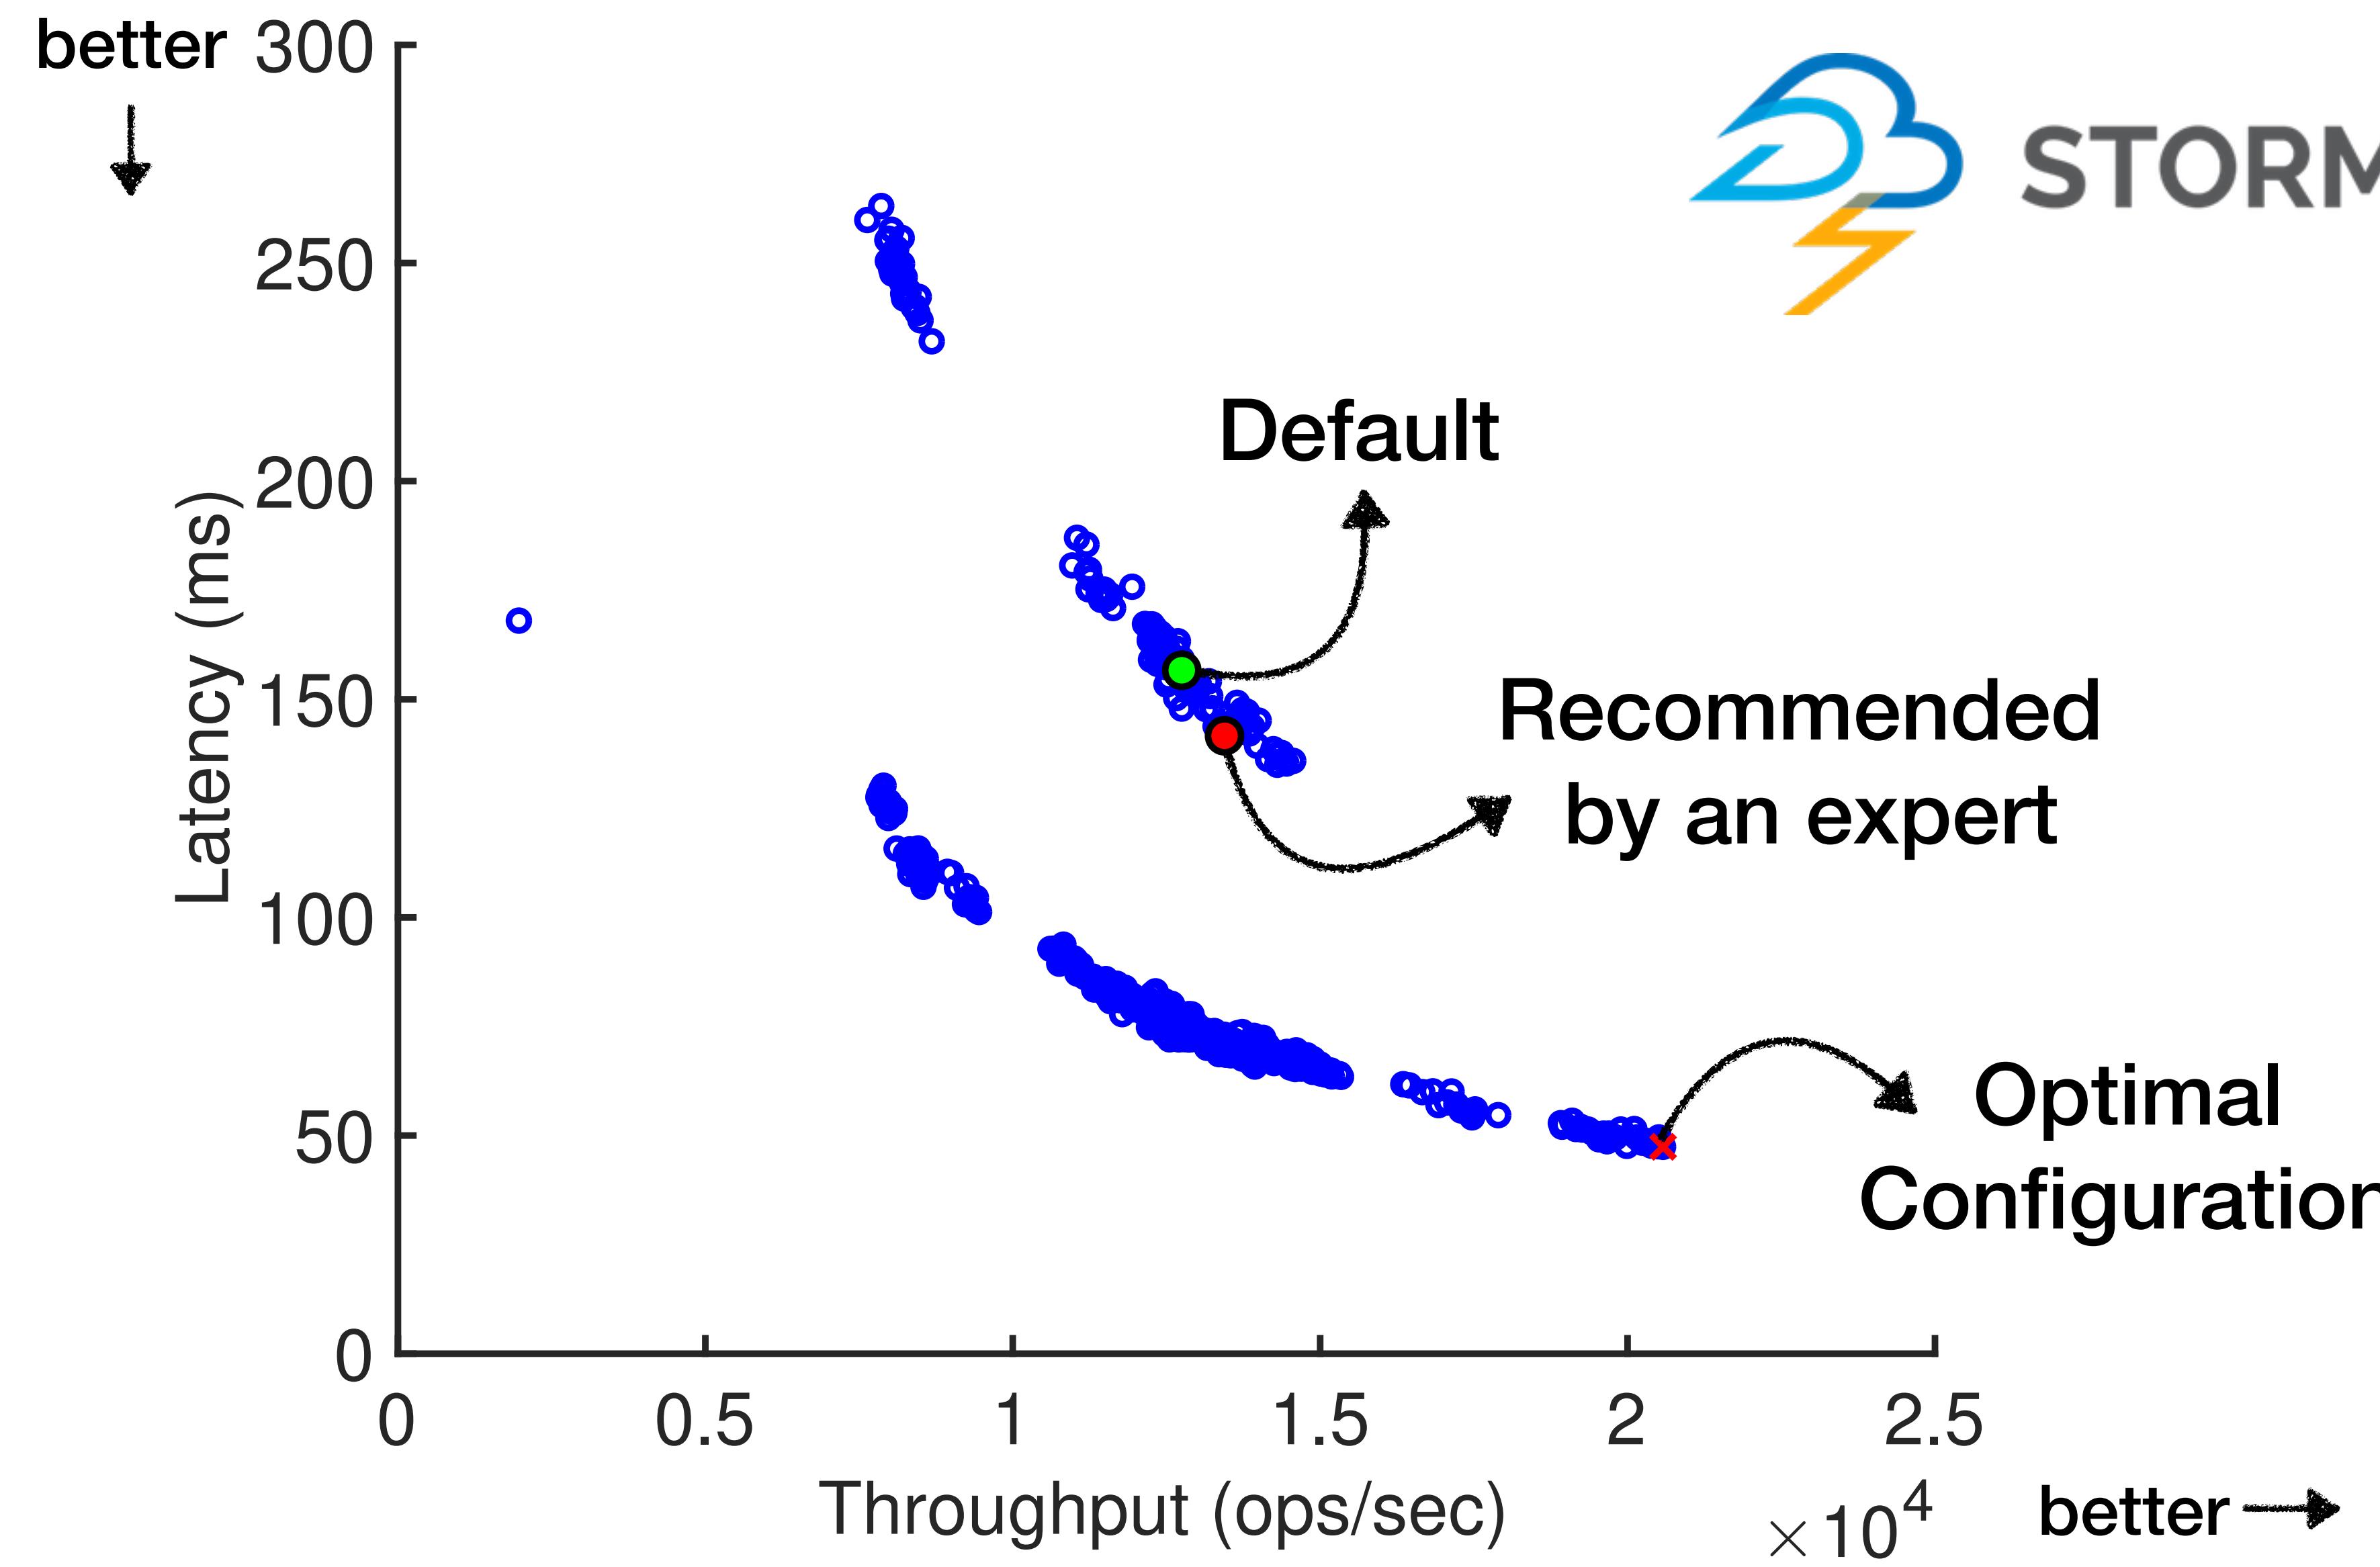
> 100



Default configuration was bad, so was the expert'



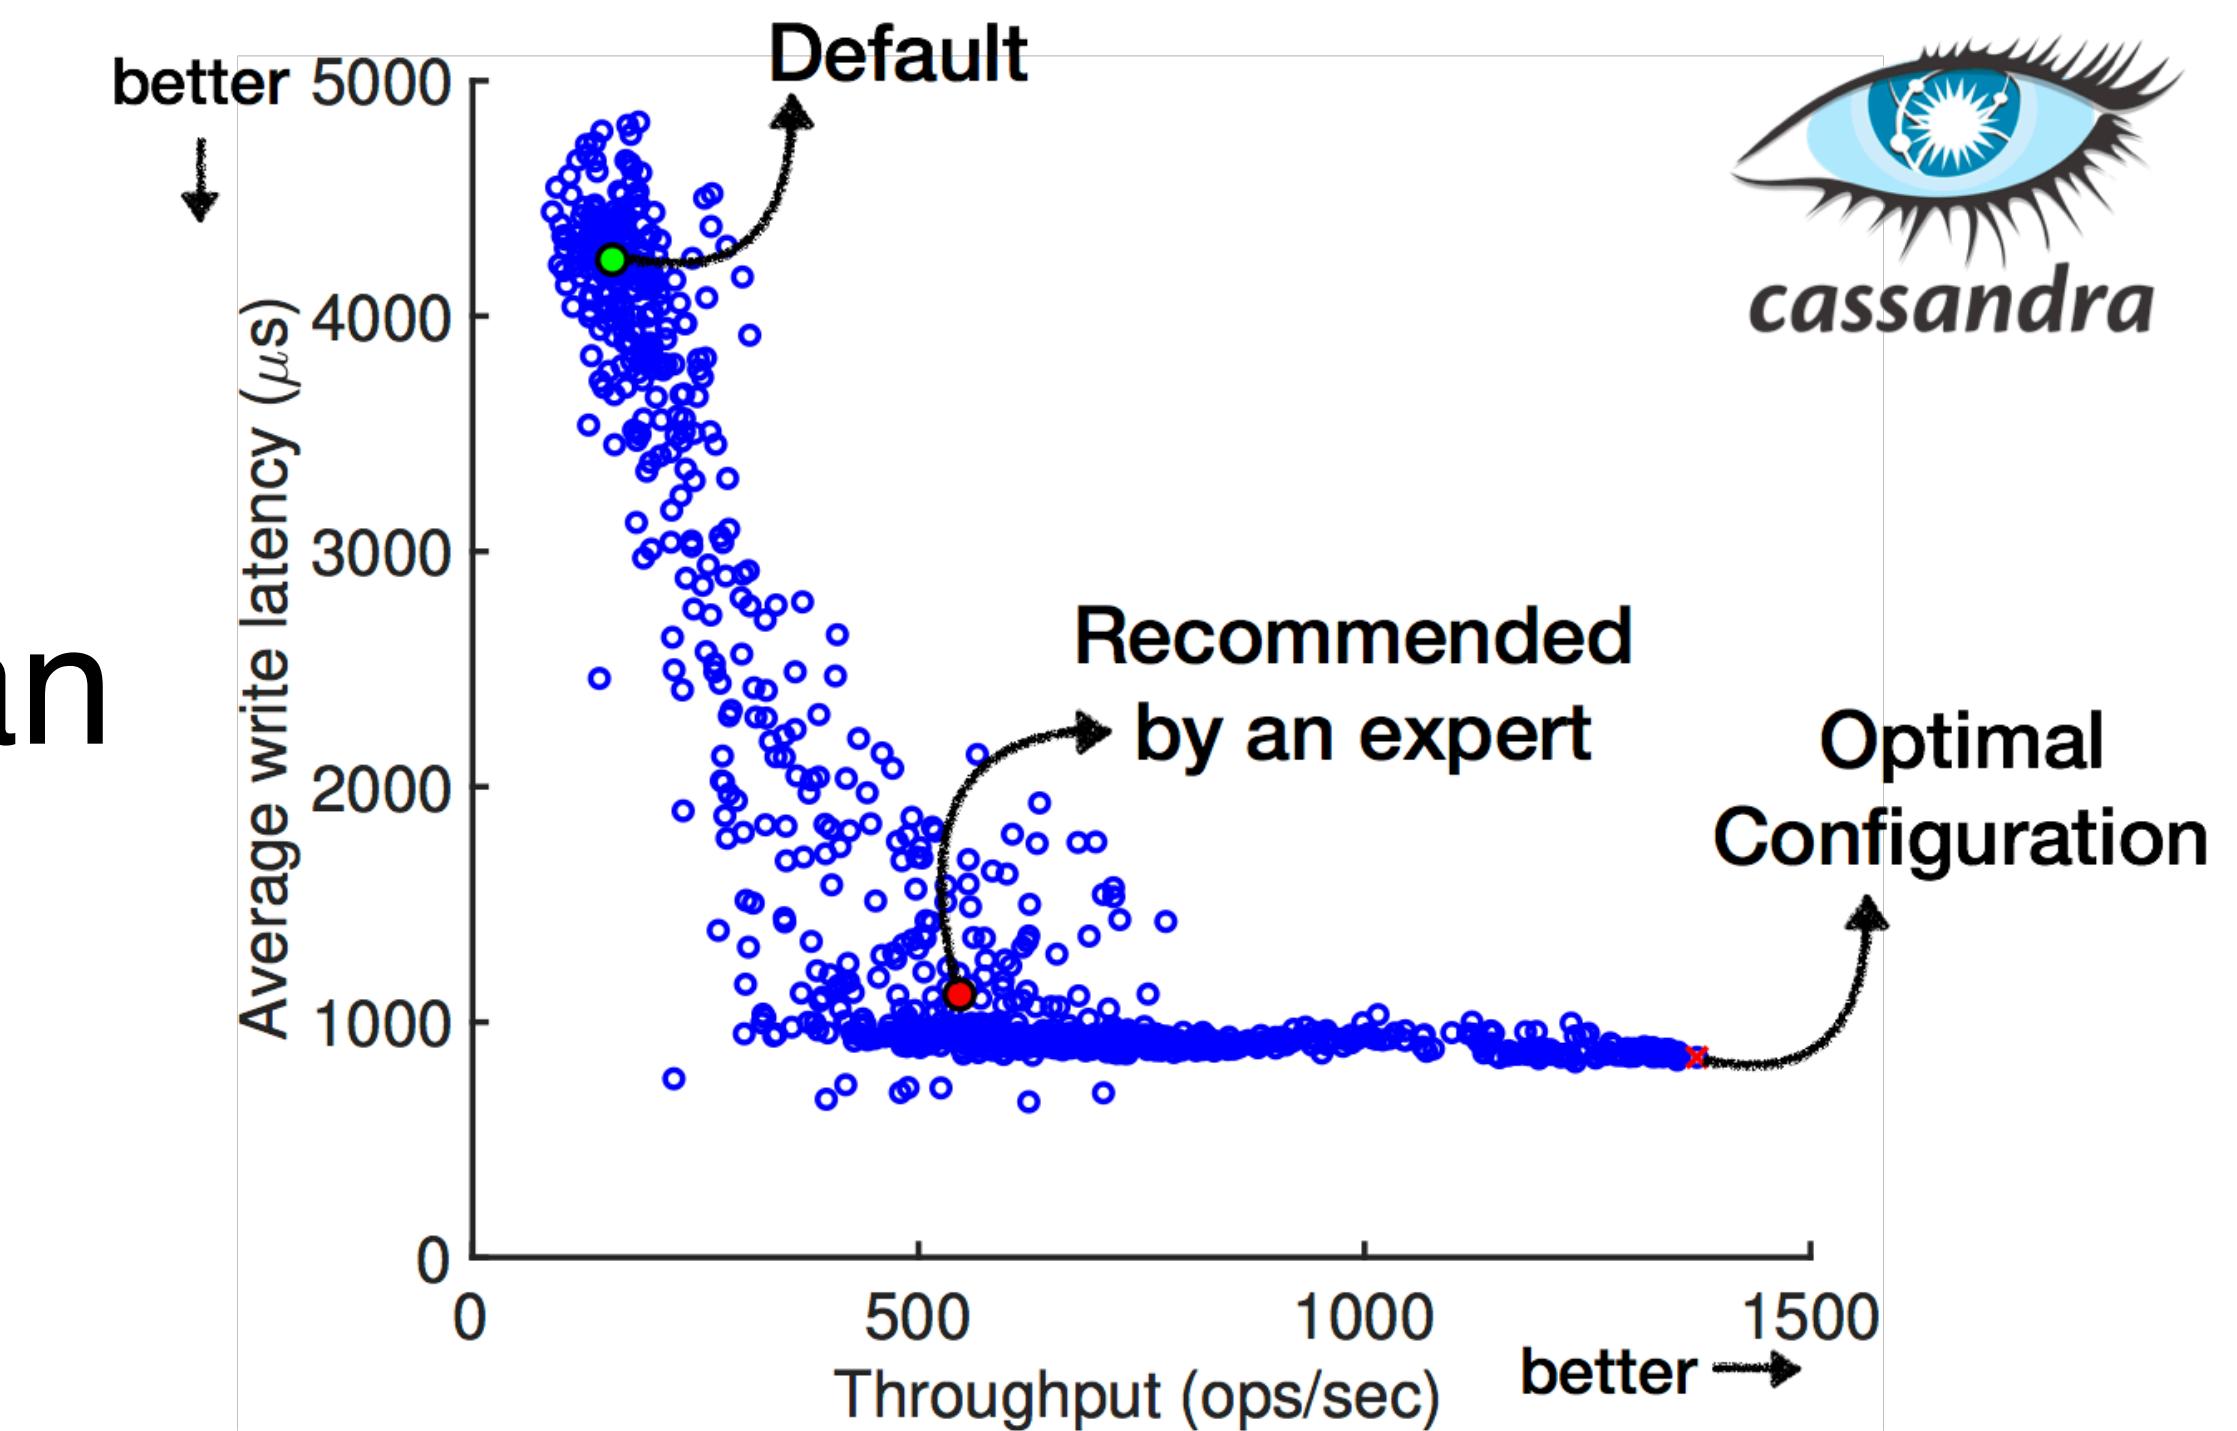
Default configuration was bad, so was the expert'



Why this is an important problem?

Significant time saving

- 2X-10X faster than worst
- Noticeably faster than median
- Default is bad
- Expert's is not optimal



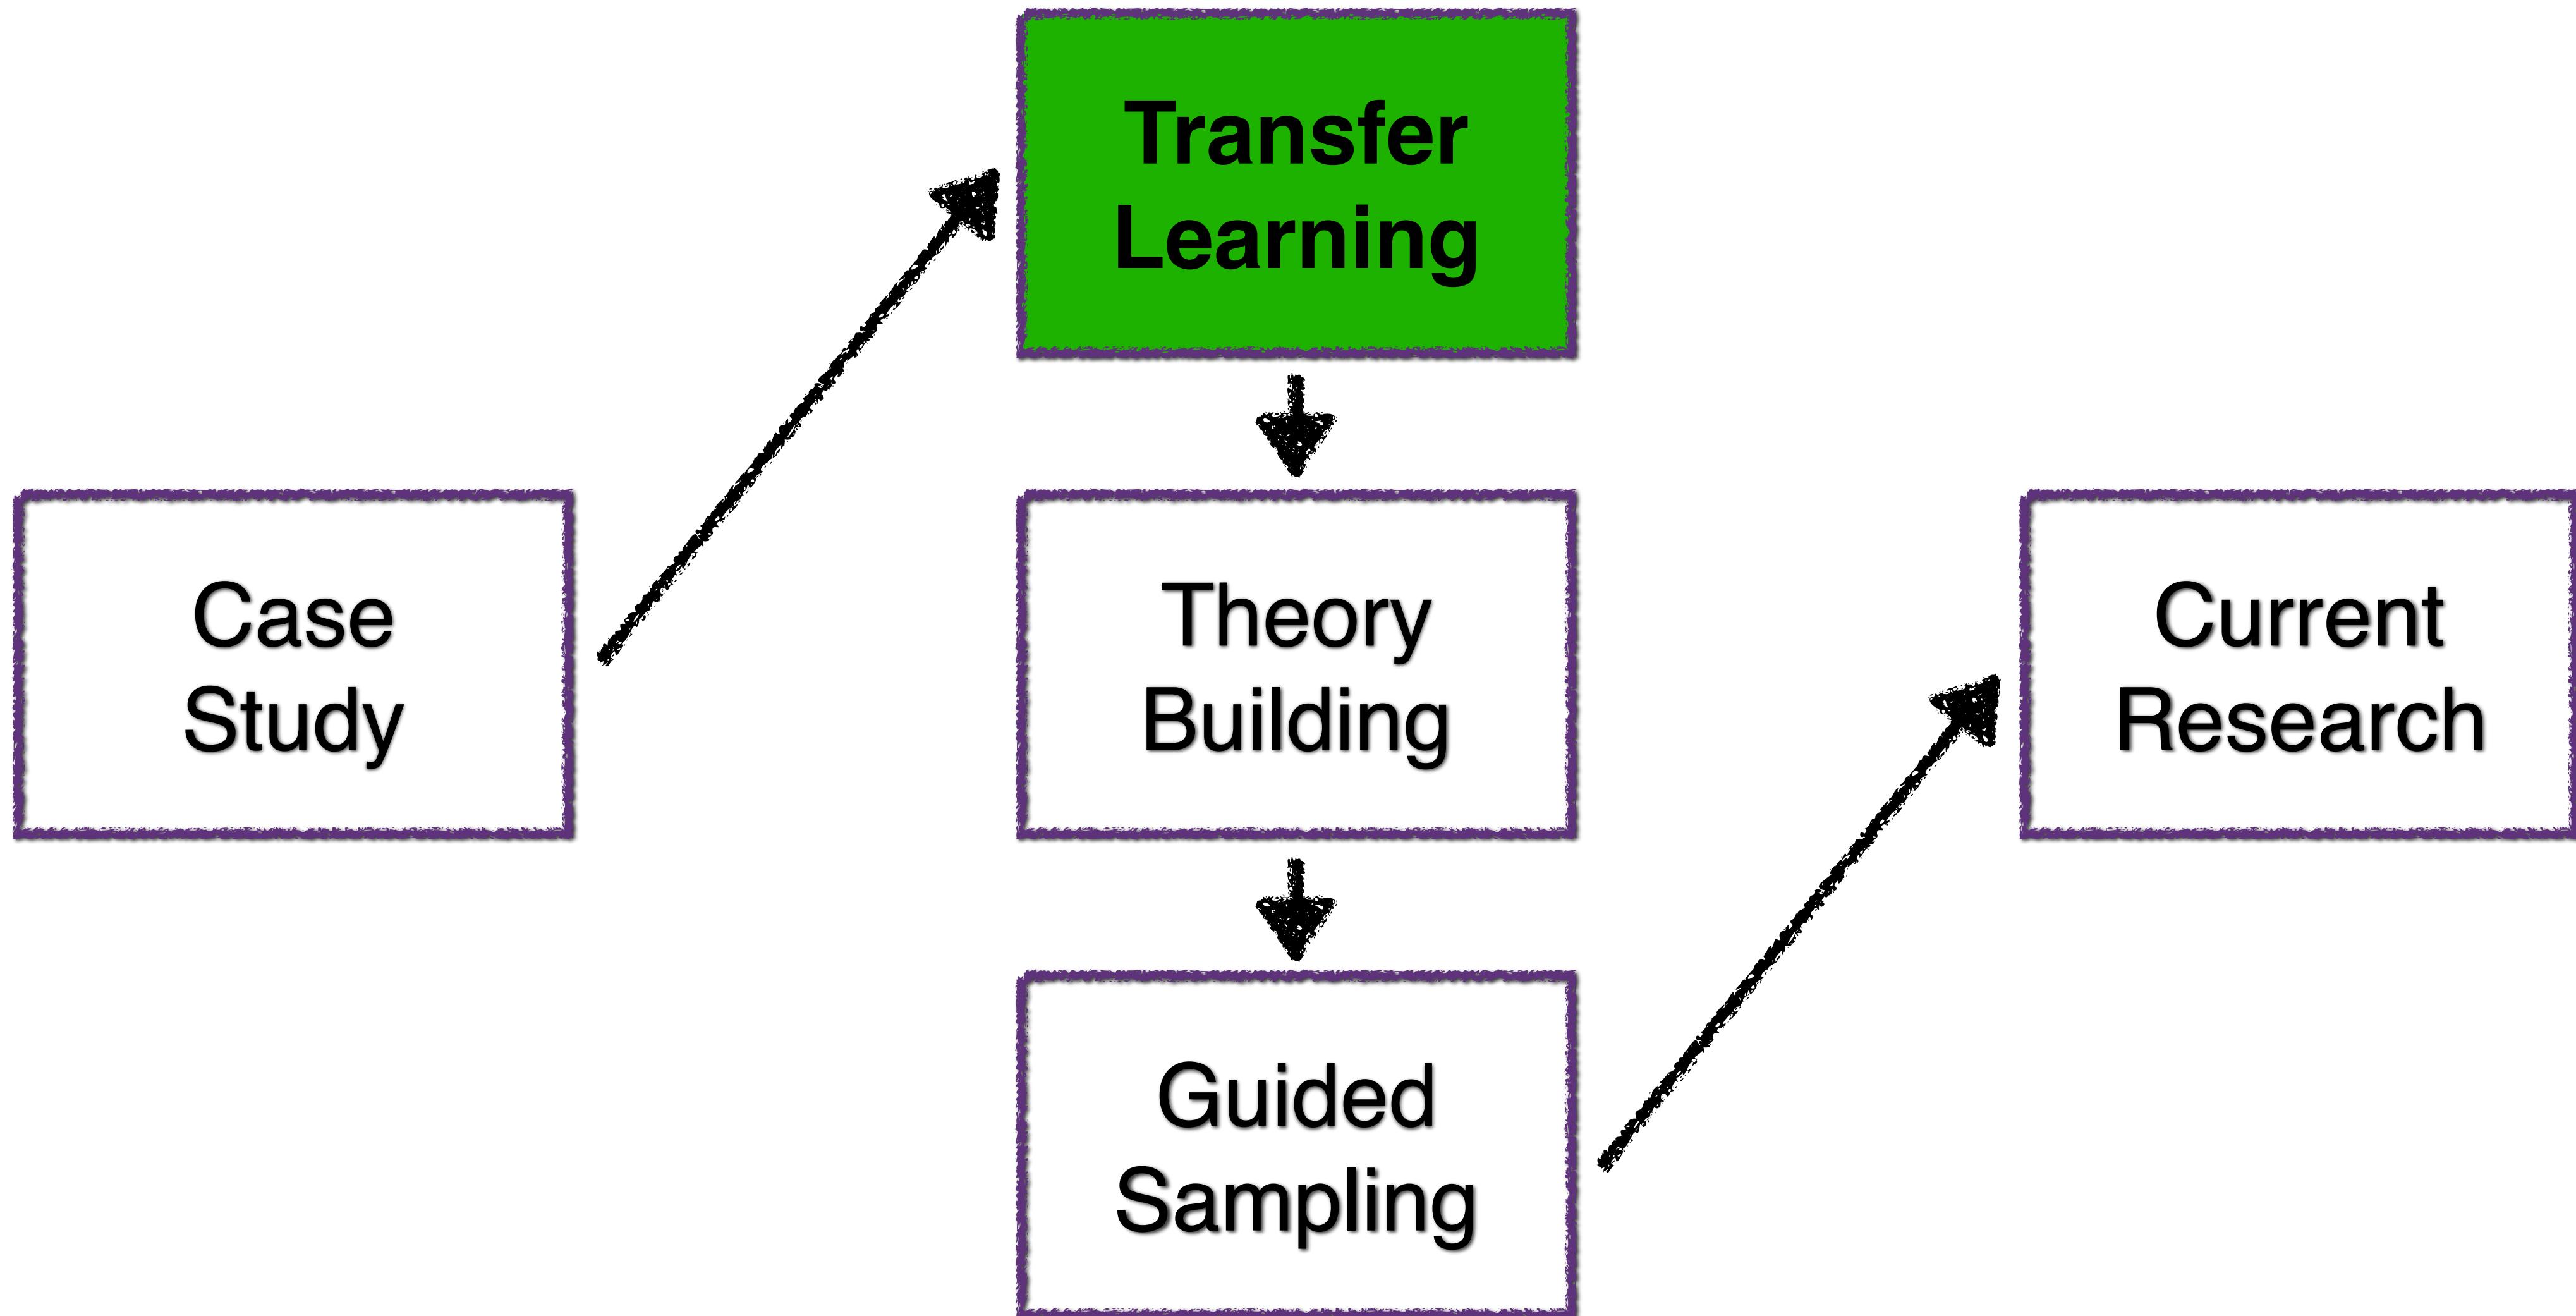
Large configuration space

- Exhaustive search is expensive
- Specific to hardware/workload/version

What did happen at the end?

- Achieved the objectives (**100X** user, same experience)
- Saved money by reducing cloud resources up to **20%**
- Our tool was able to identify configurations that was consistently better than expert recommendation

Outline



To enable performance tradeoff, we need a model to reason about qualities

```
void Parrot_setenv(. . . name,. . . value){  
#ifdef PARROT HAS SETENV  
    my_setenv(name, value, 1);  
#else  
    int name_len=strlen(name);  
    int val_len=strlen(value);  
    char* envs=glob_env;  
    if(envs==NULL){  
        return;  
    }  
    strcpy(envs,name);  
    strcpy(envs+name_len,"=");  
    strcpy(envs+name_len + 1,value);  
    putenv(envs);  
#endif  
}
```

Execution time (s)

$$f(\cdot) = 5 + 3 \times o_1$$

$$f(o_1 := 1) = 8$$

$$f(o_1 := 0) = 5$$

What is a performance model?

$$f : \mathbb{C} \rightarrow \mathbb{R}$$

$$f(o_1, o_2) = 5 + \boxed{3o_1} + \boxed{15o_2} - \boxed{7o_1 \times o_2}$$

$$c = \langle o_1, o_2 \rangle$$

$$c = \langle o_1, o_2, \dots, o_{10} \rangle$$

⋮

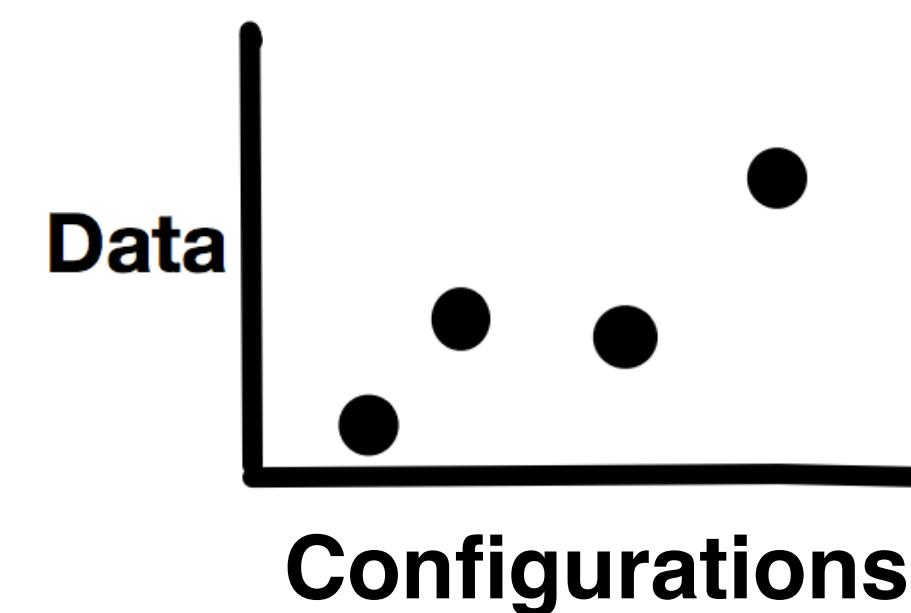
$$c = \langle o_1, o_2, \dots, o_{100} \rangle$$

How do we learn performance models?



TurtleBot

Measure



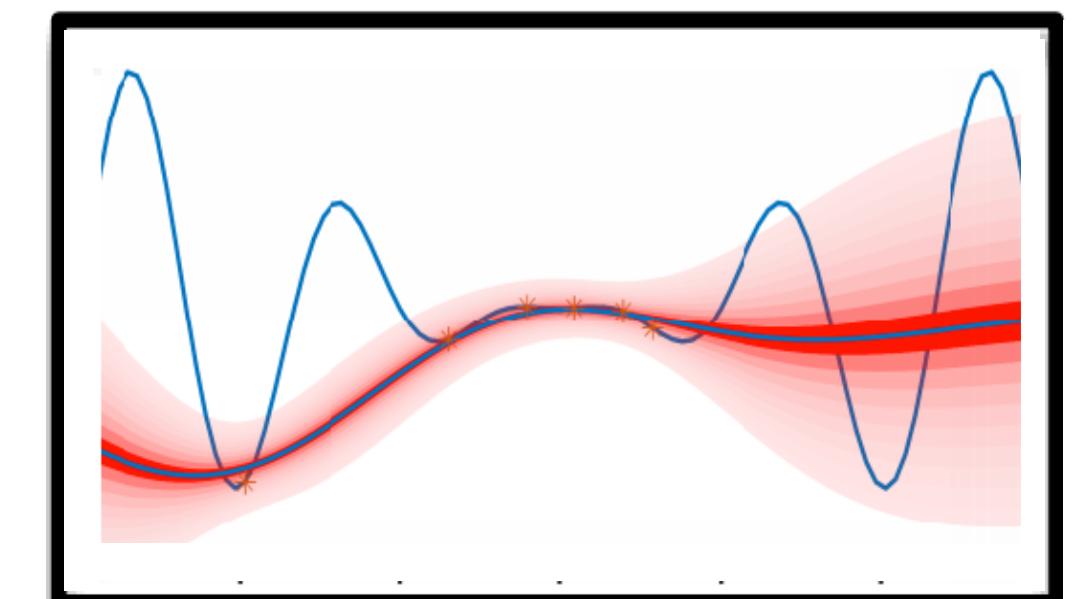
Learn



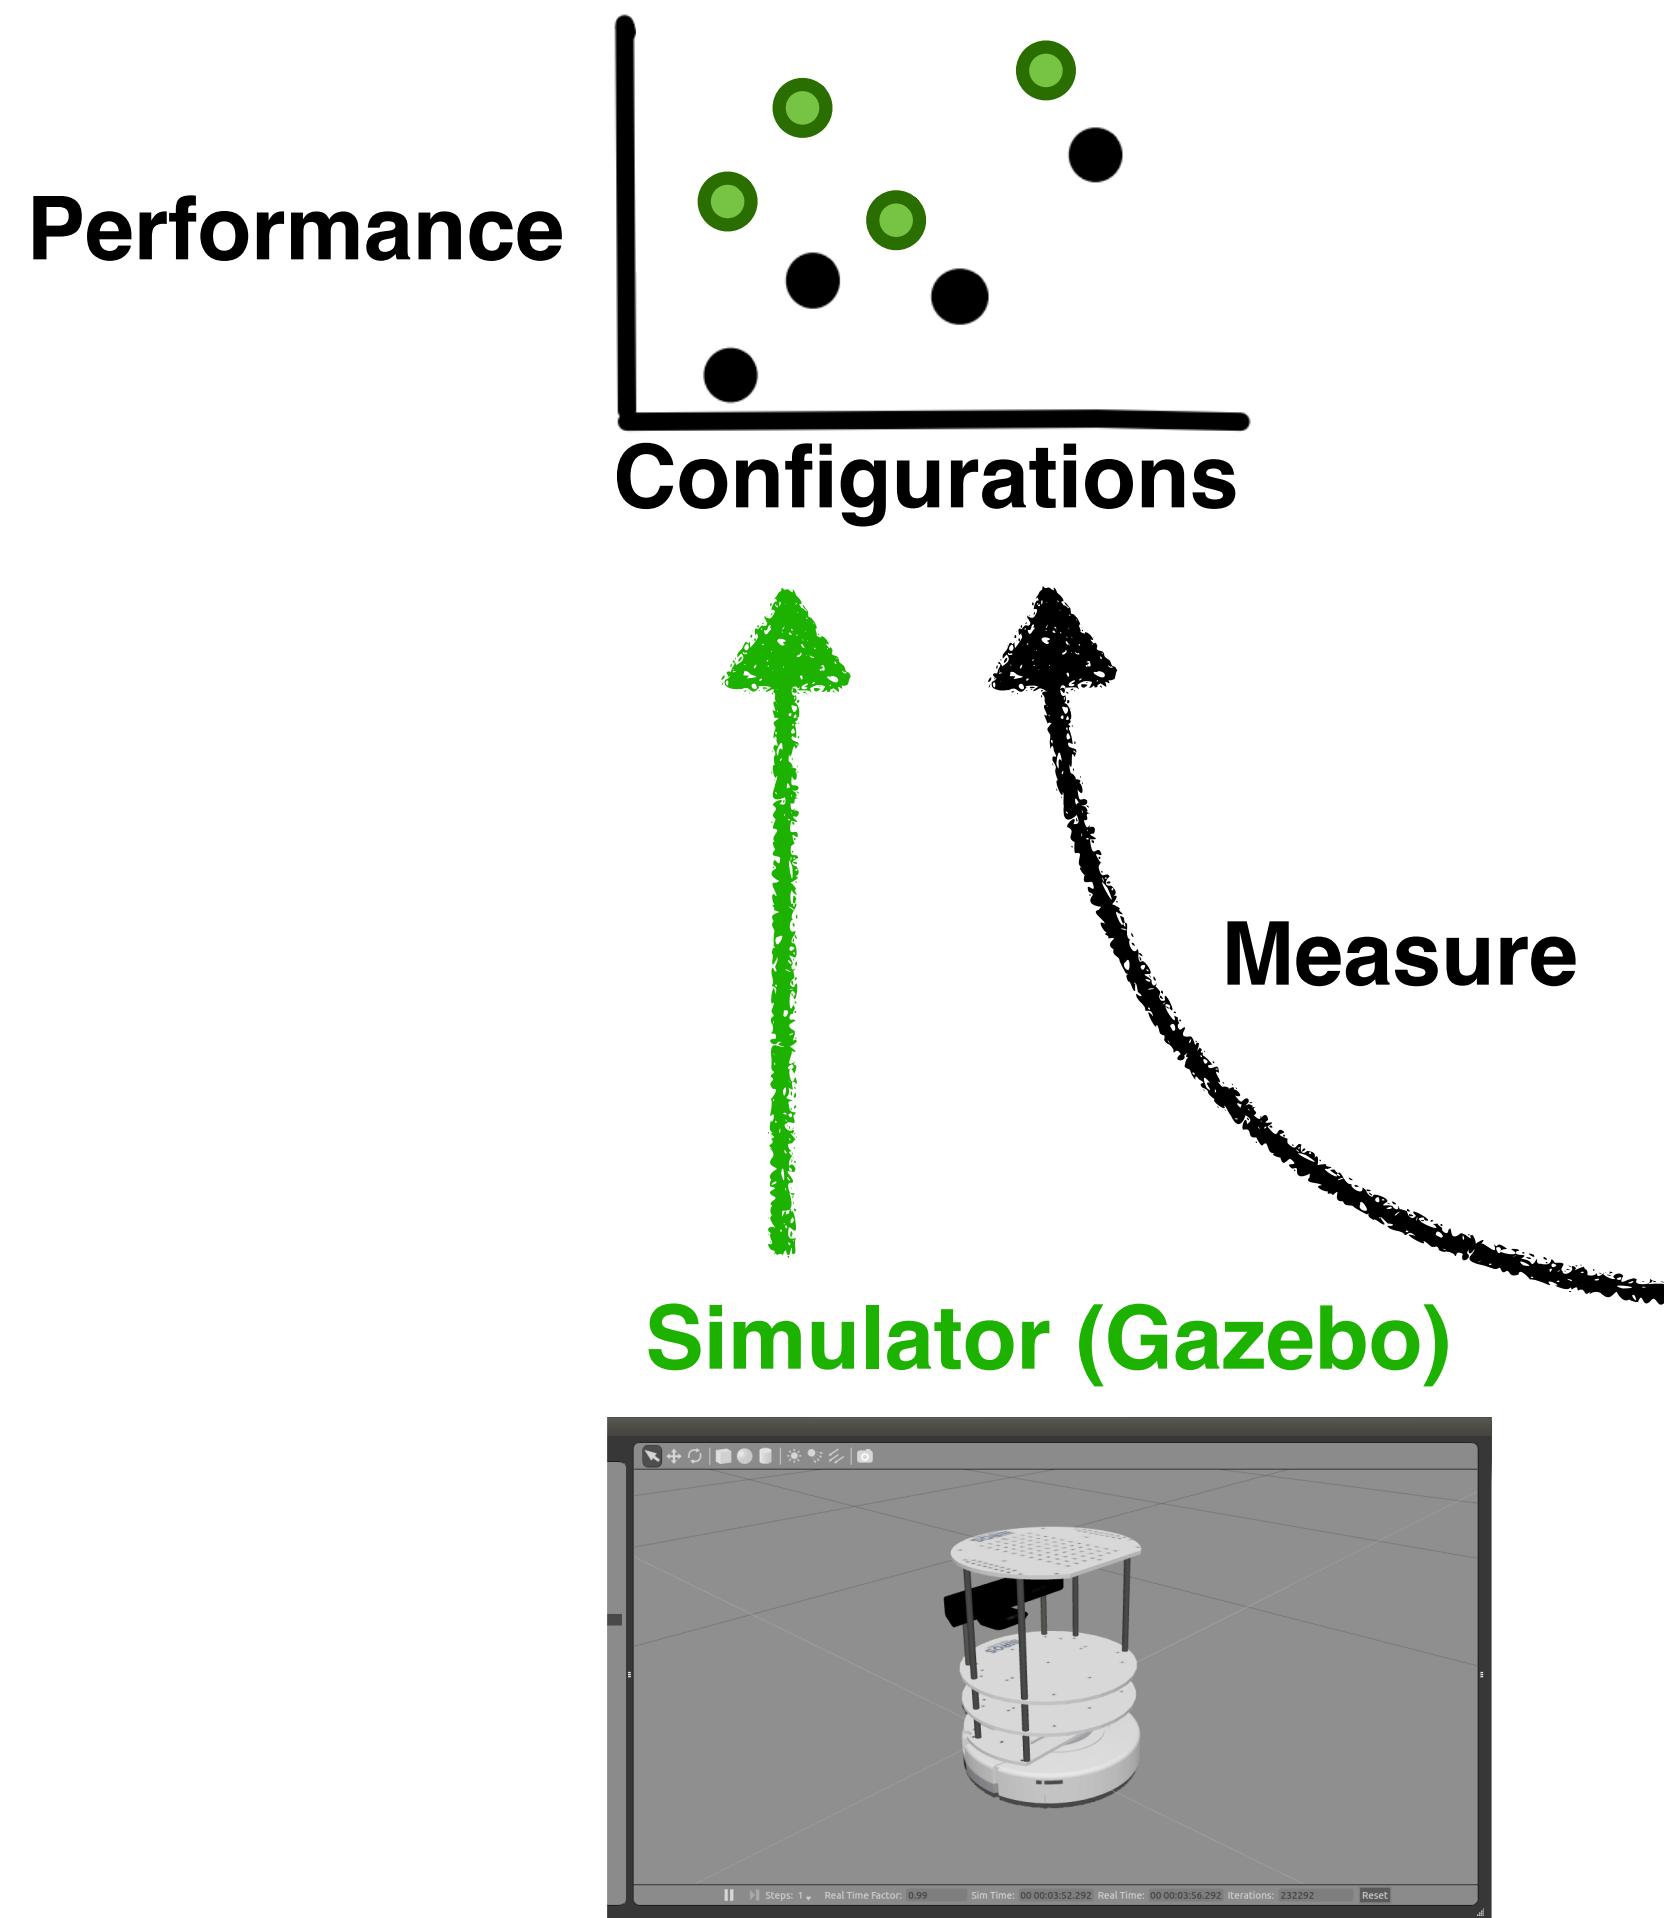
Optimization
Reasoning
Debugging



$$f(o_1, o_2) = 5 + 3o_1 + 15o_2 - 7o_1 \times o_2$$



Insight: Performance measurements of the real system is “similar” to the ones from the simulators

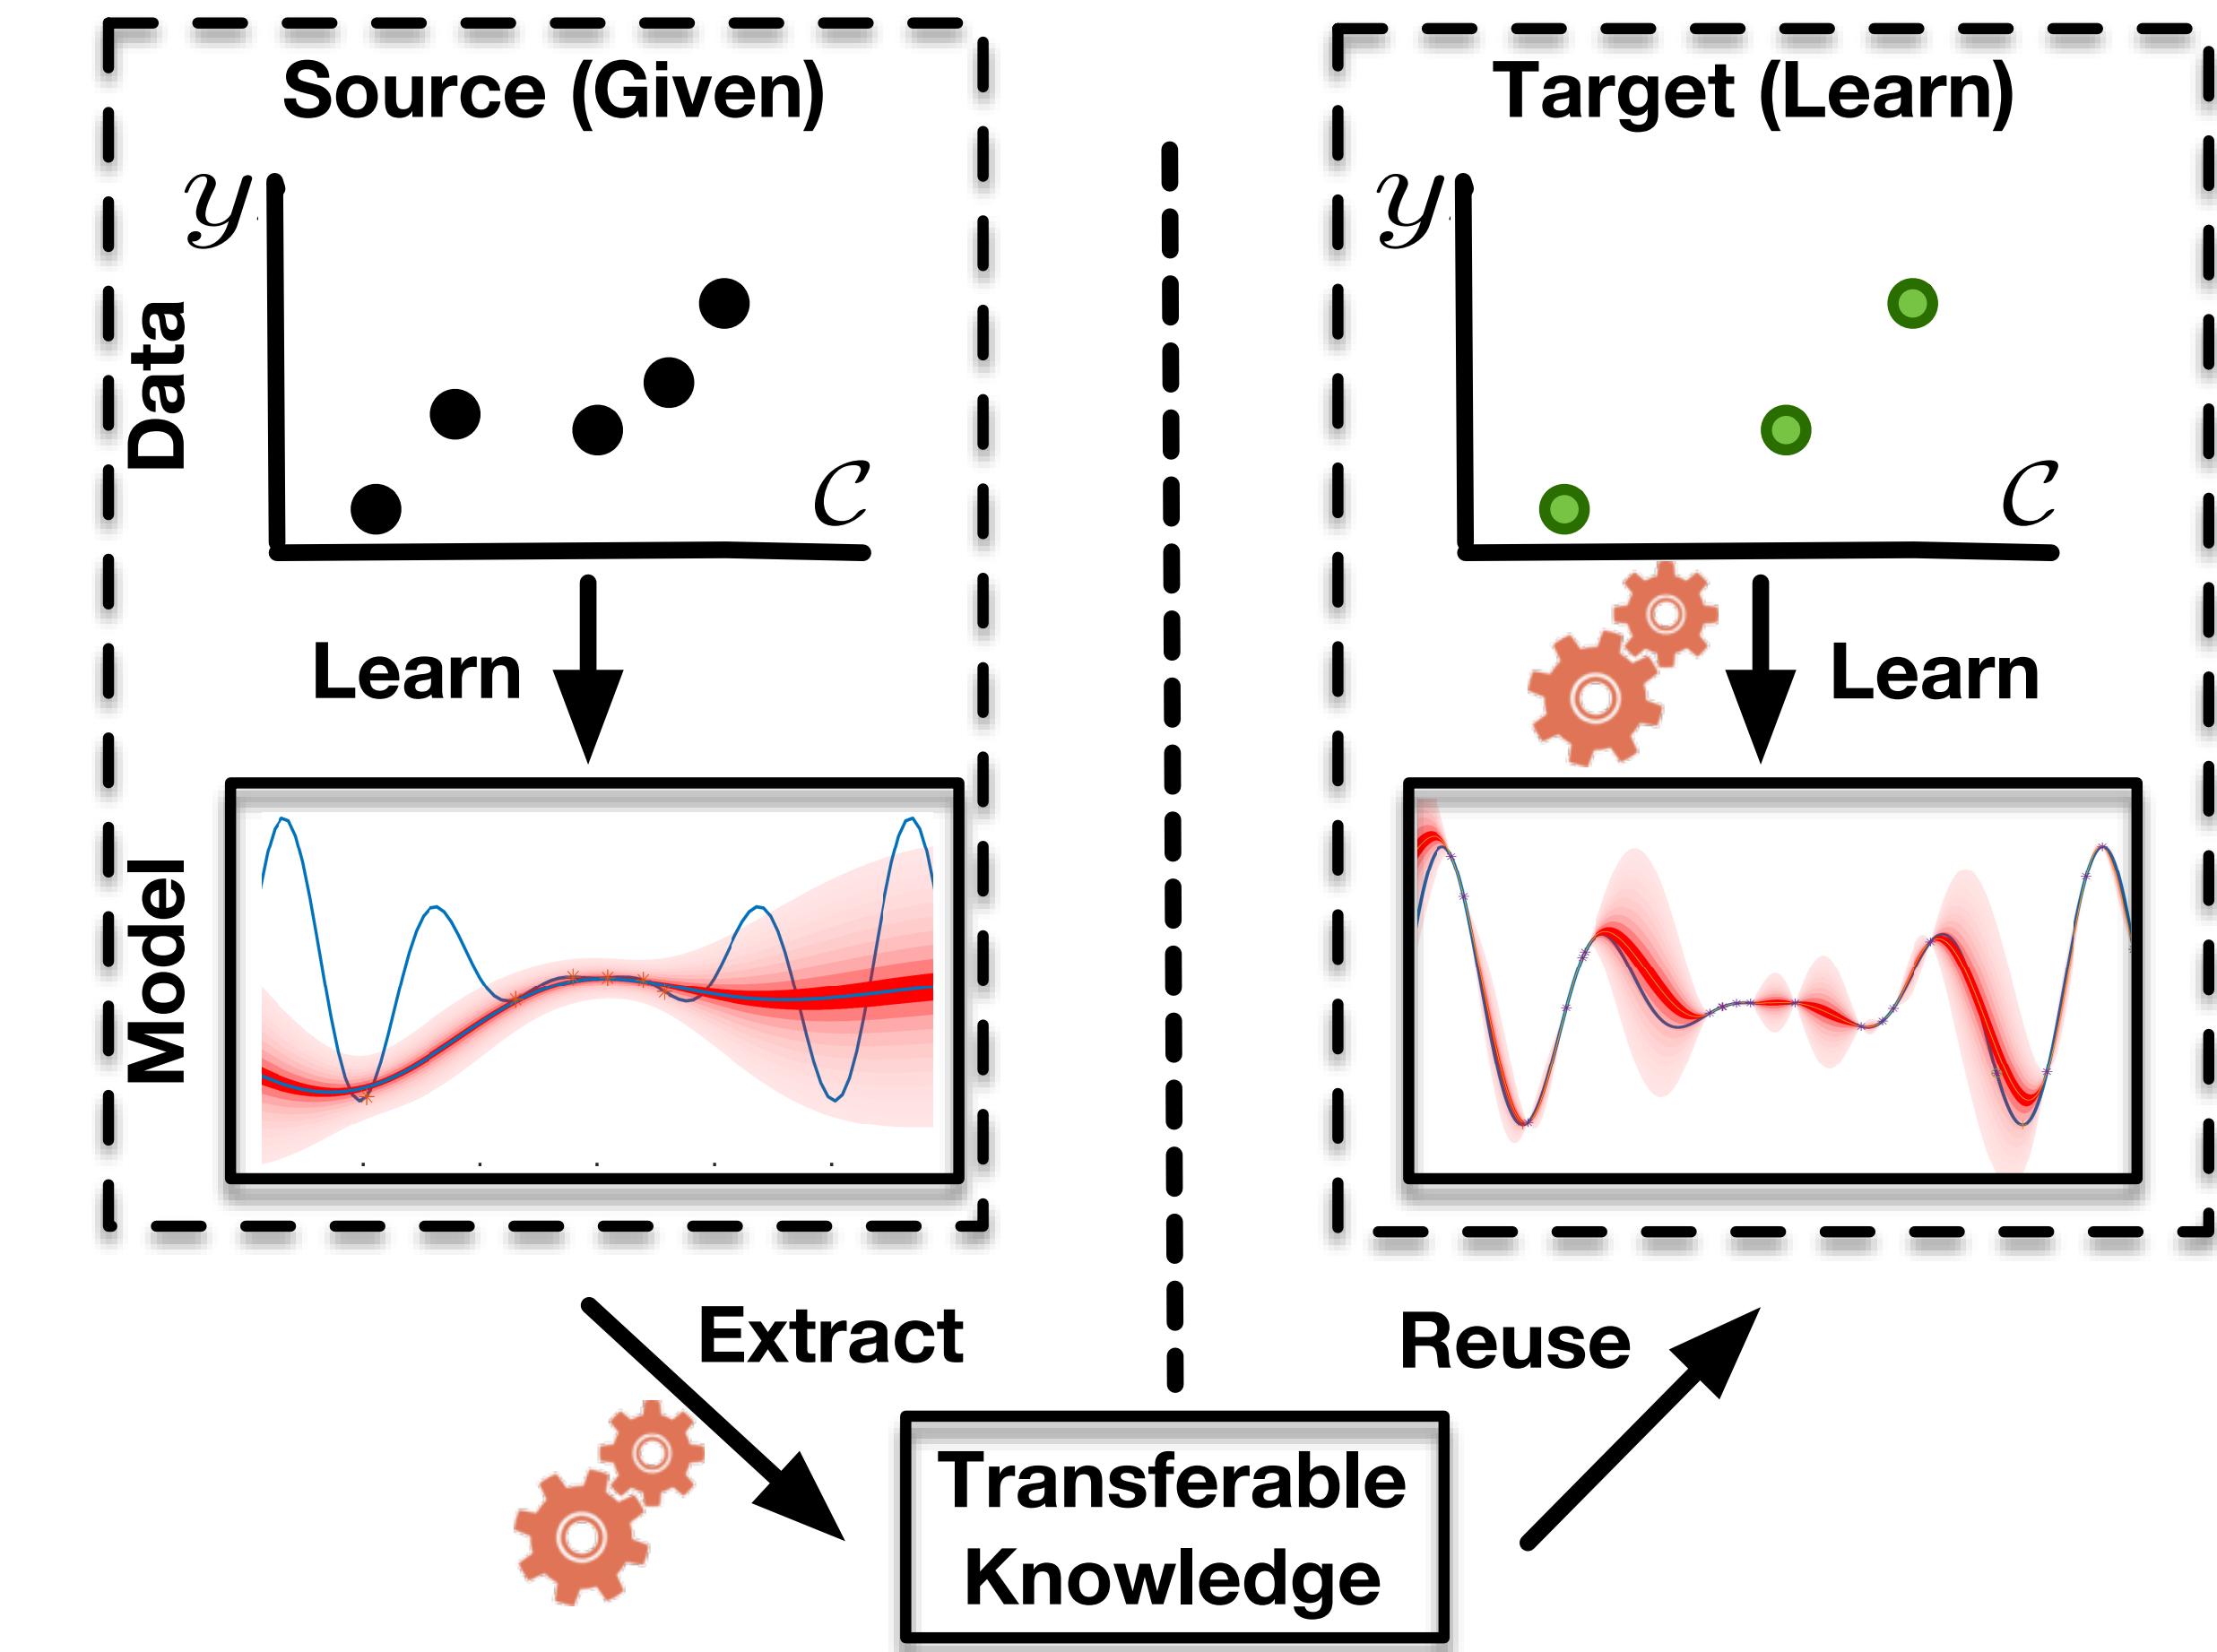


So why not reuse these data,
instead of measuring on real robot?

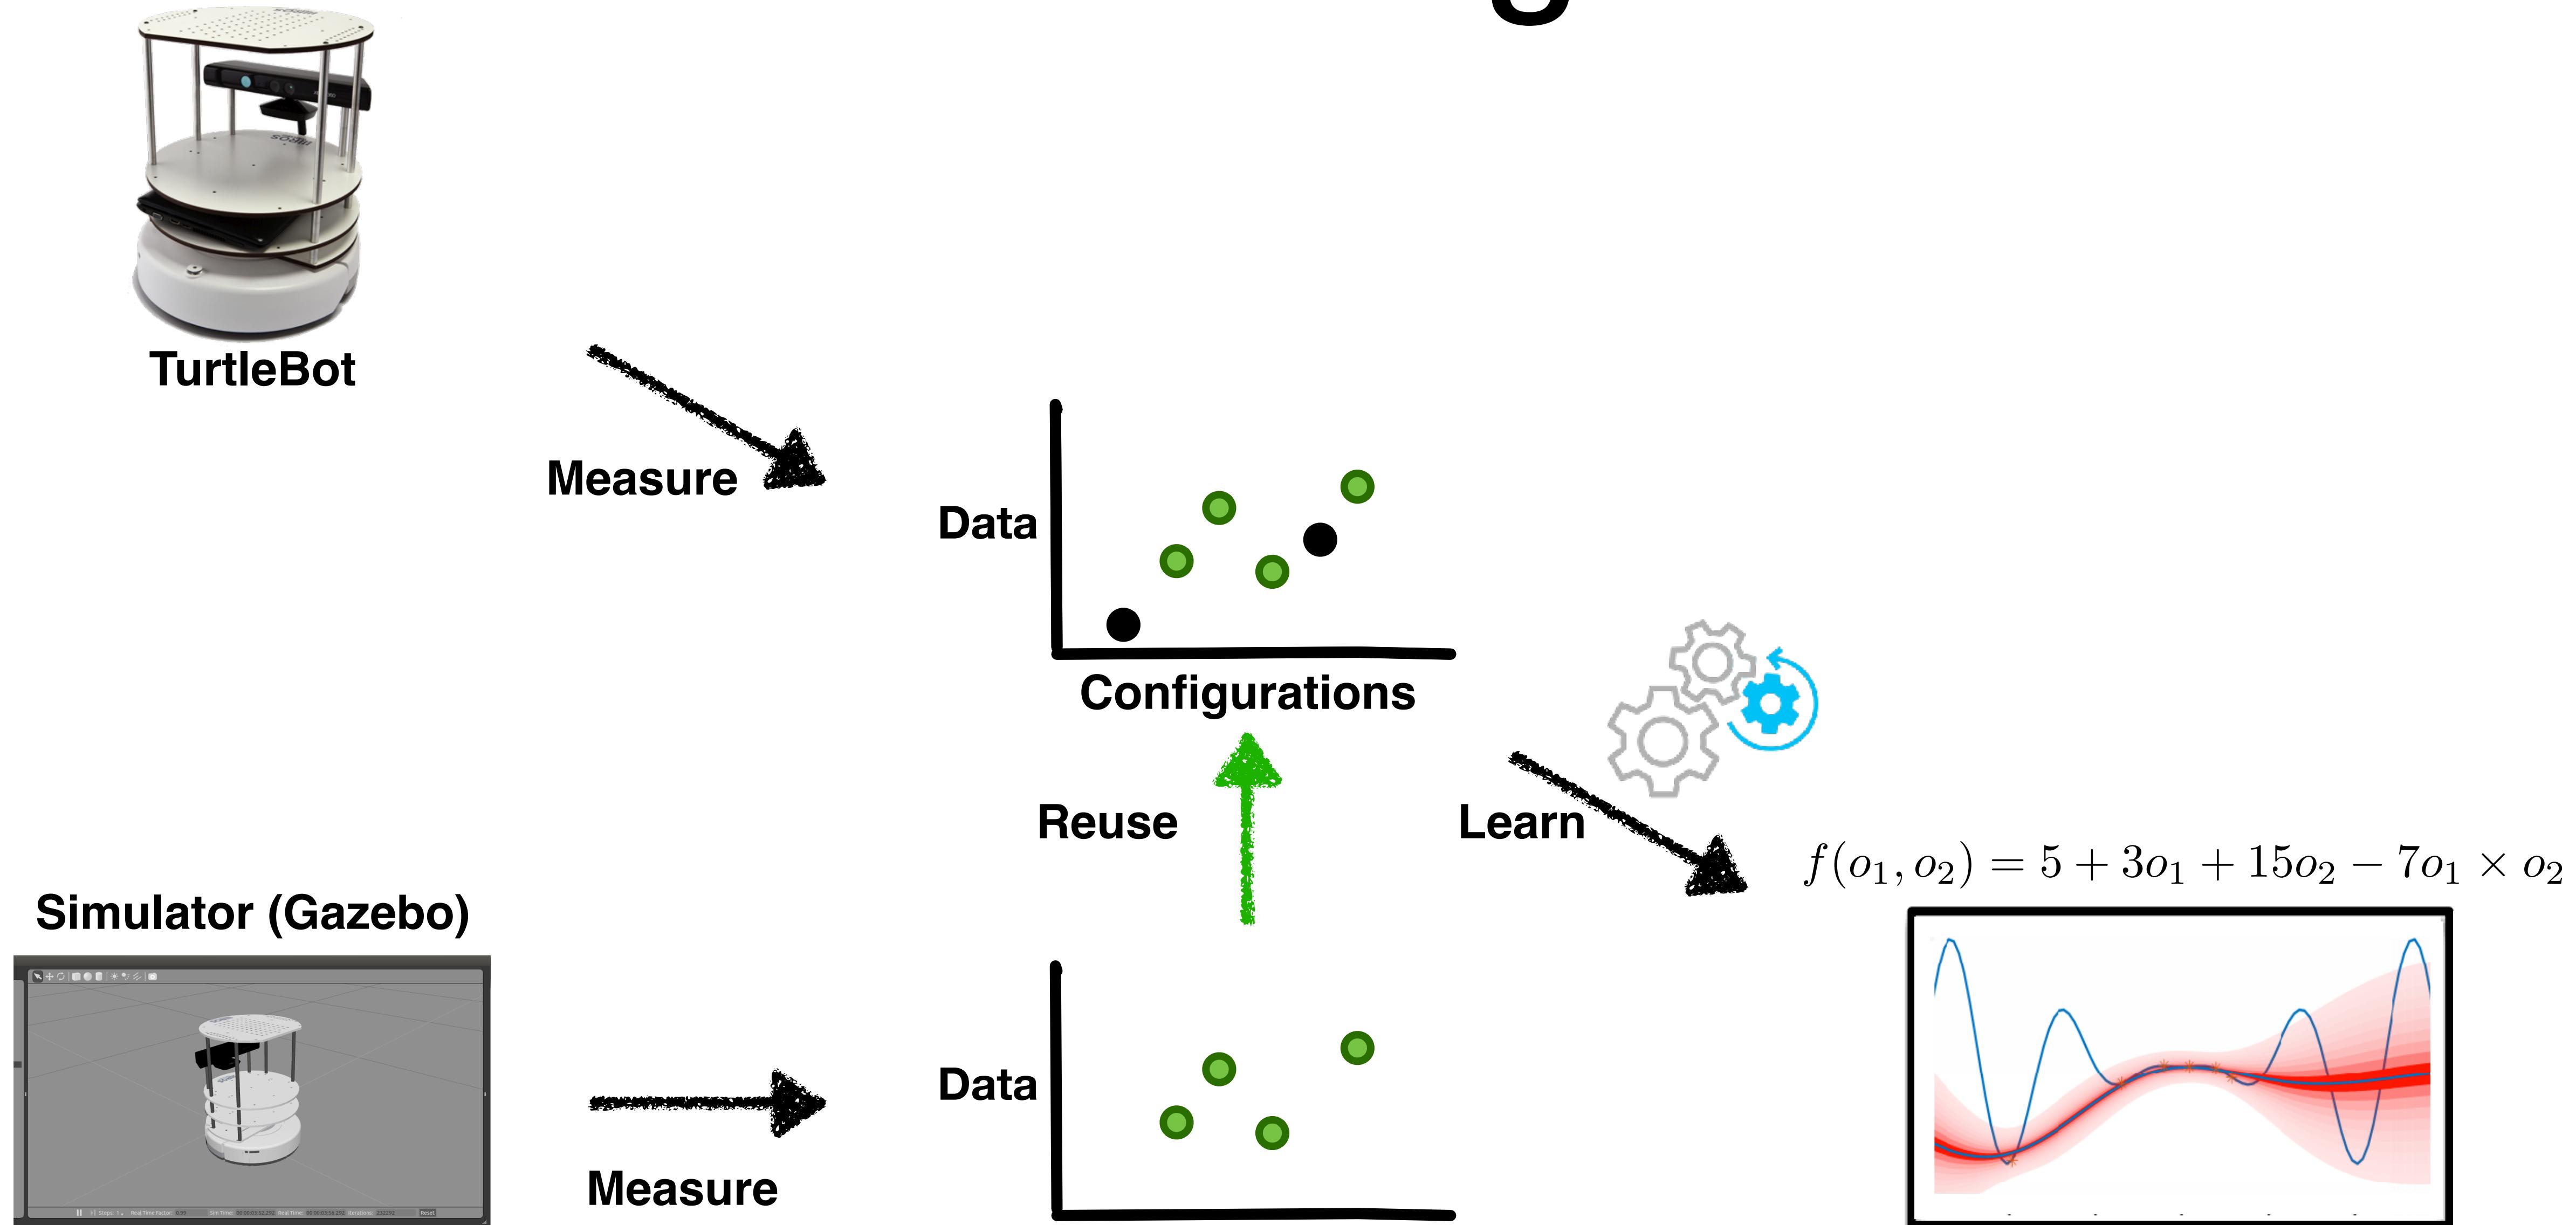


We developed methods to make learning cheaper via transfer learning

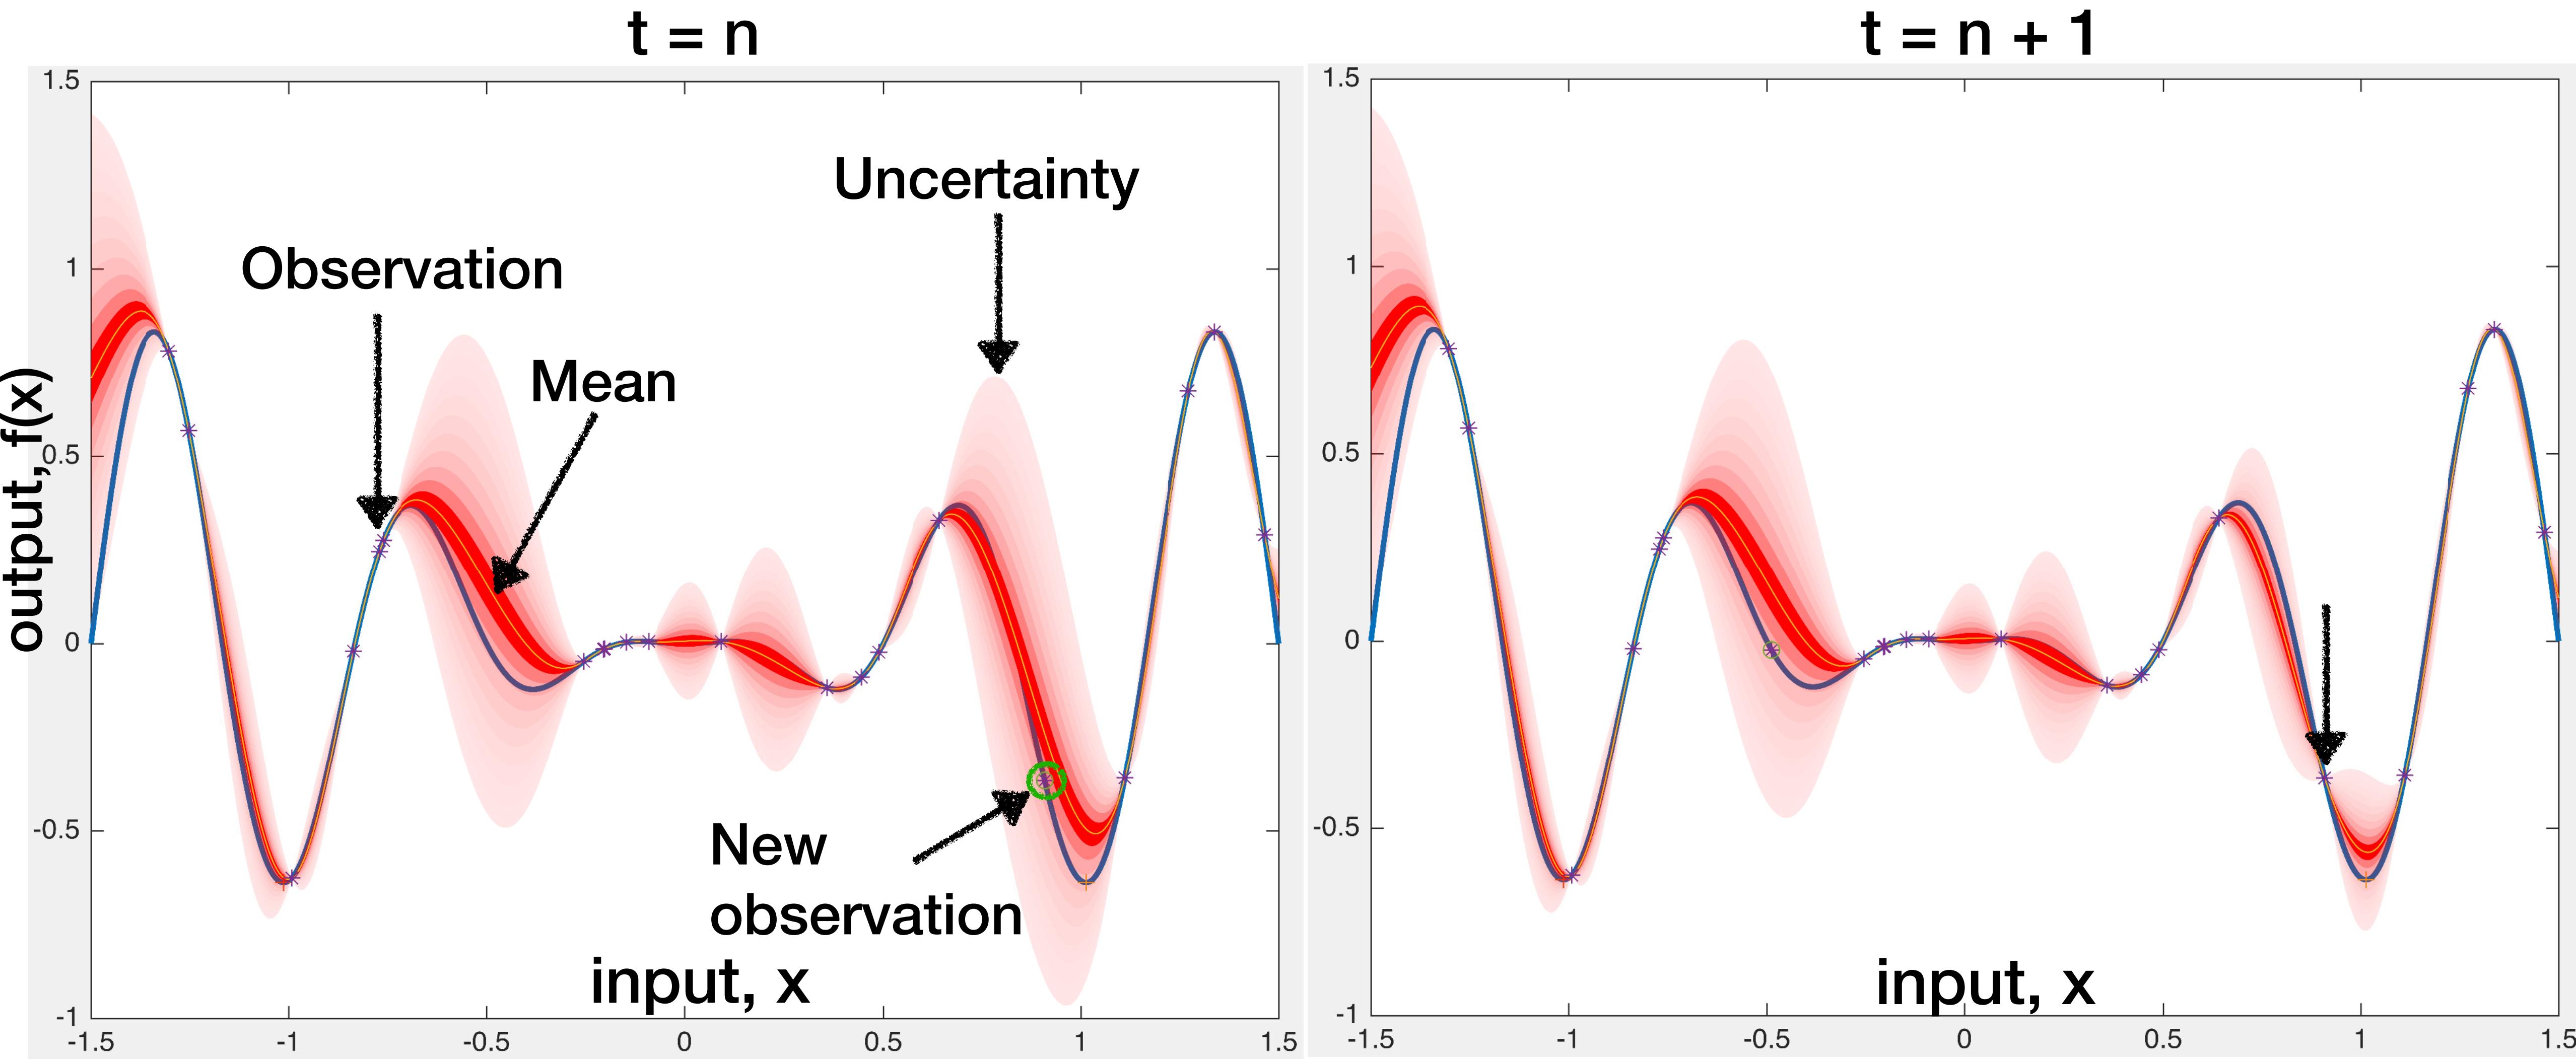
Goal: Gain strength by transferring information across environments



Our transfer learning solution



Gaussian processes for performance modeling



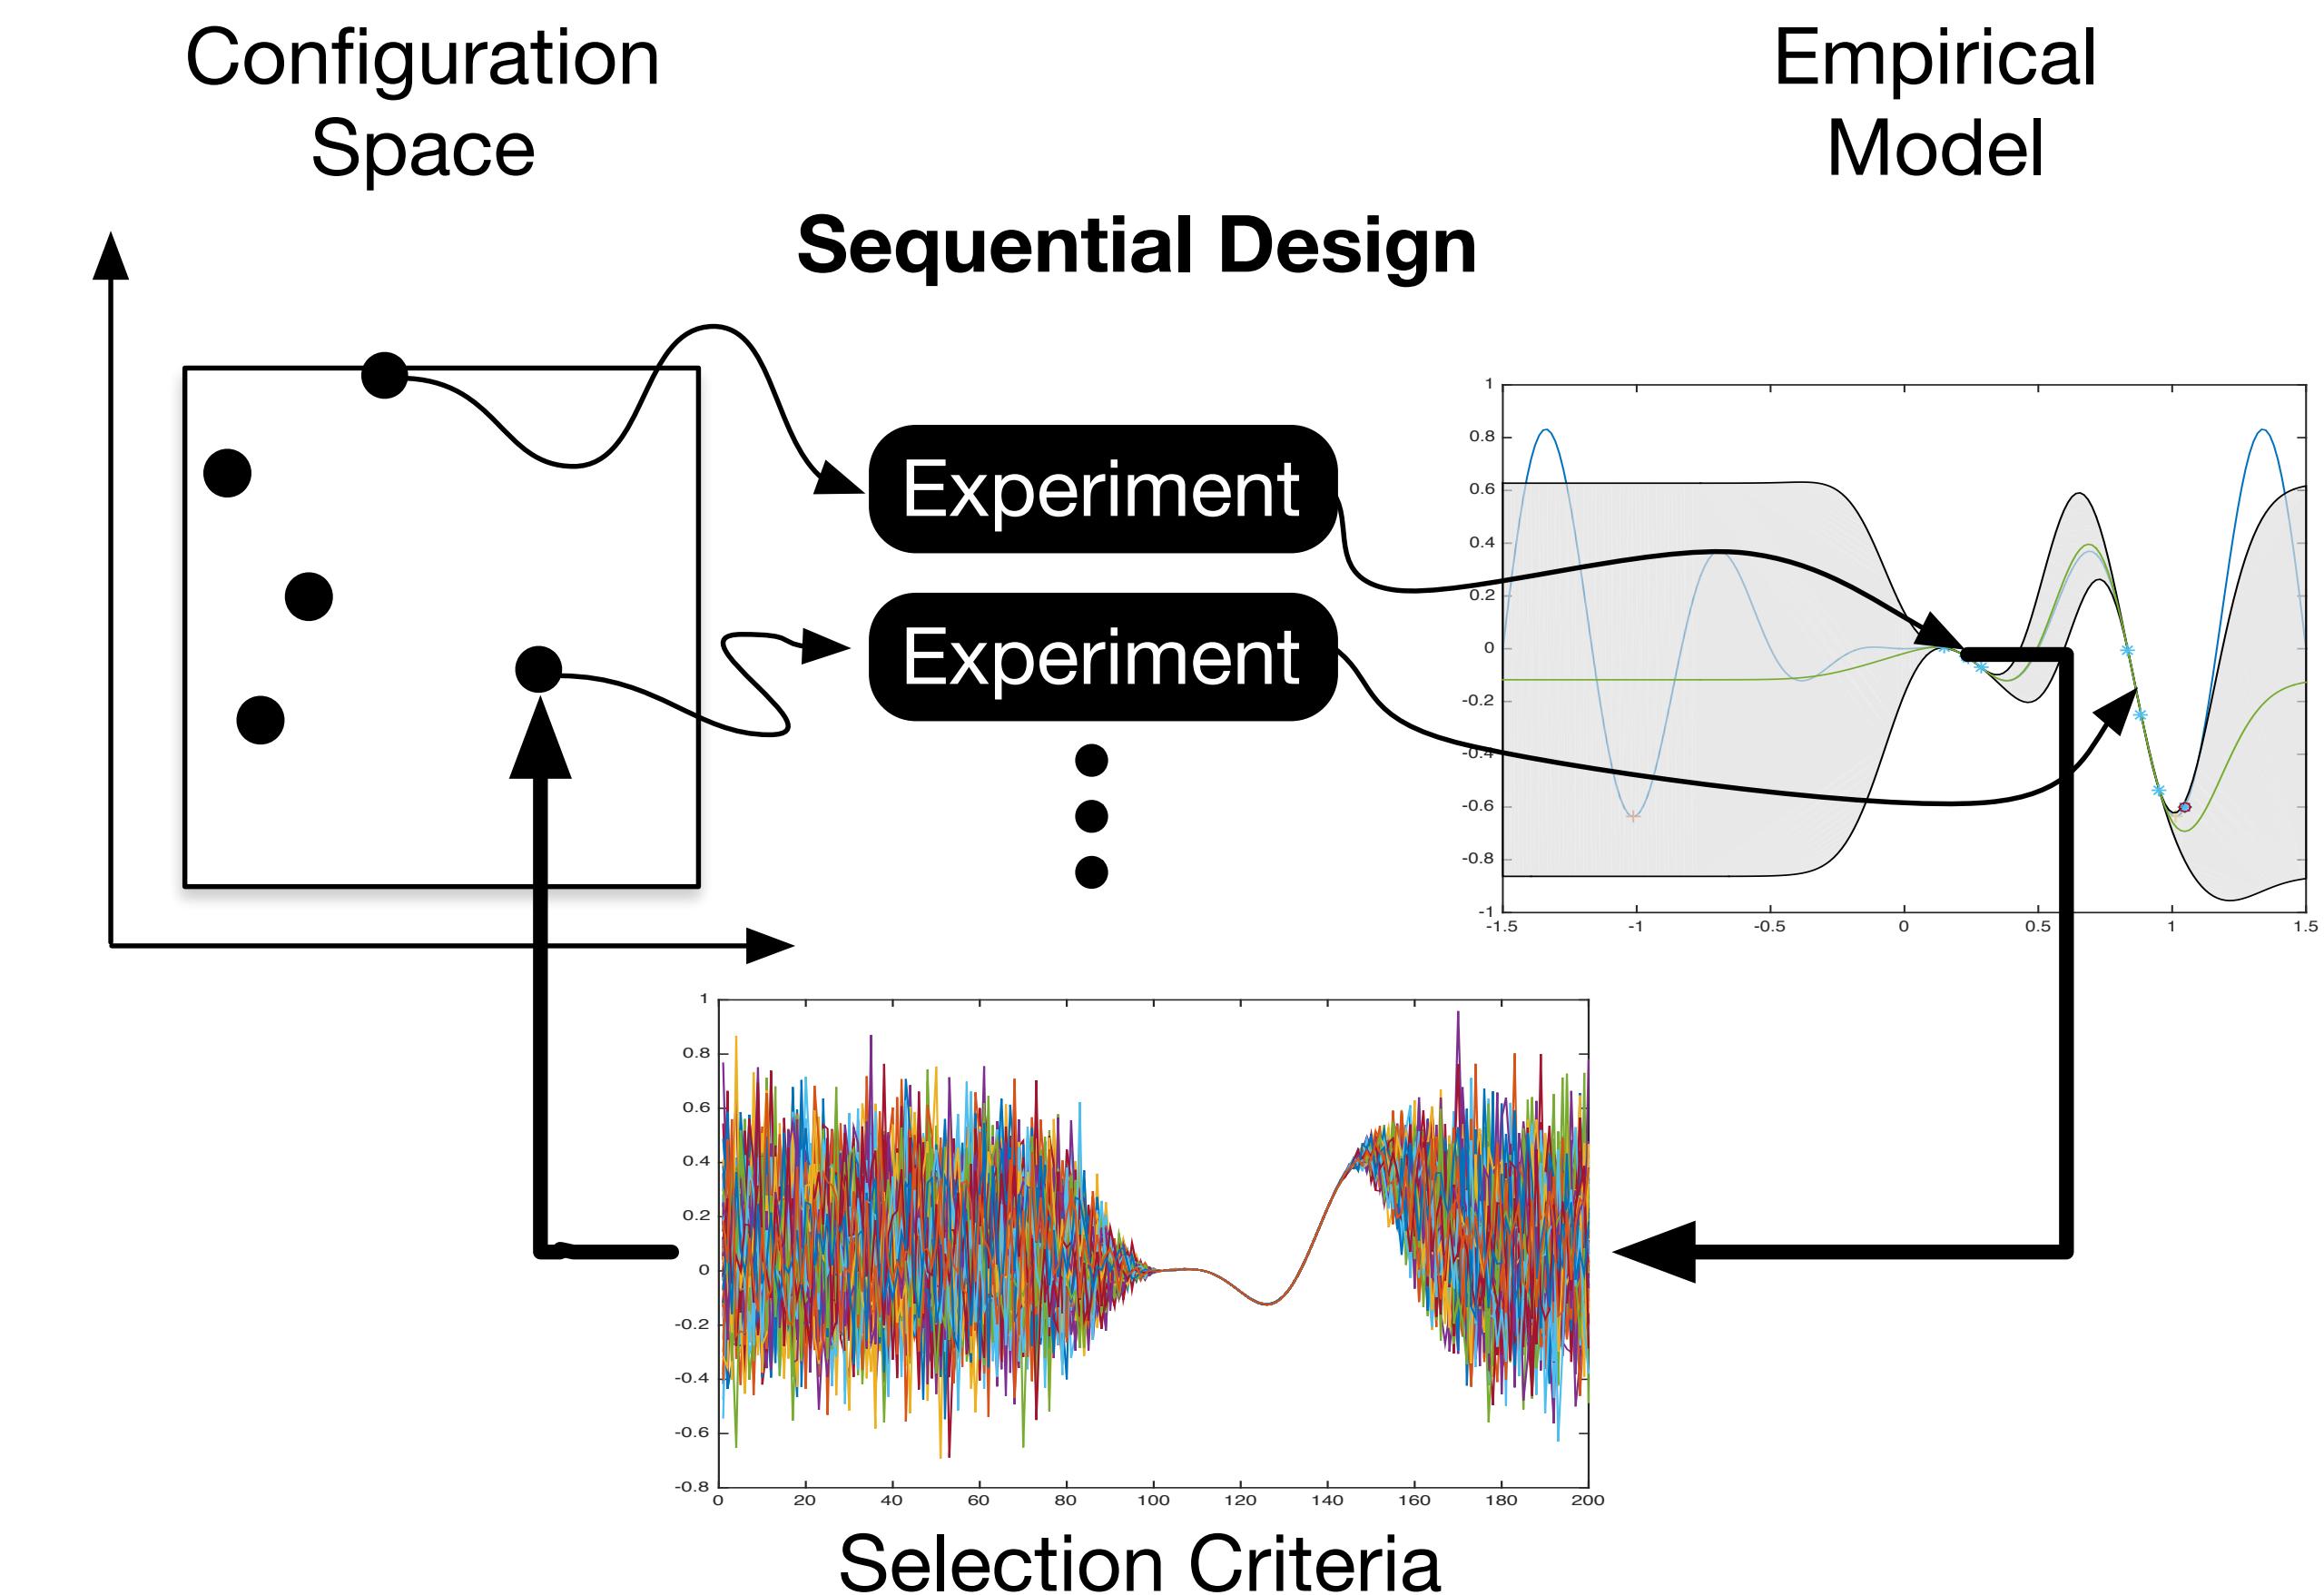
Gaussian Processes enables reasoning about performance

Step 1: Fit GP to the data seen so far

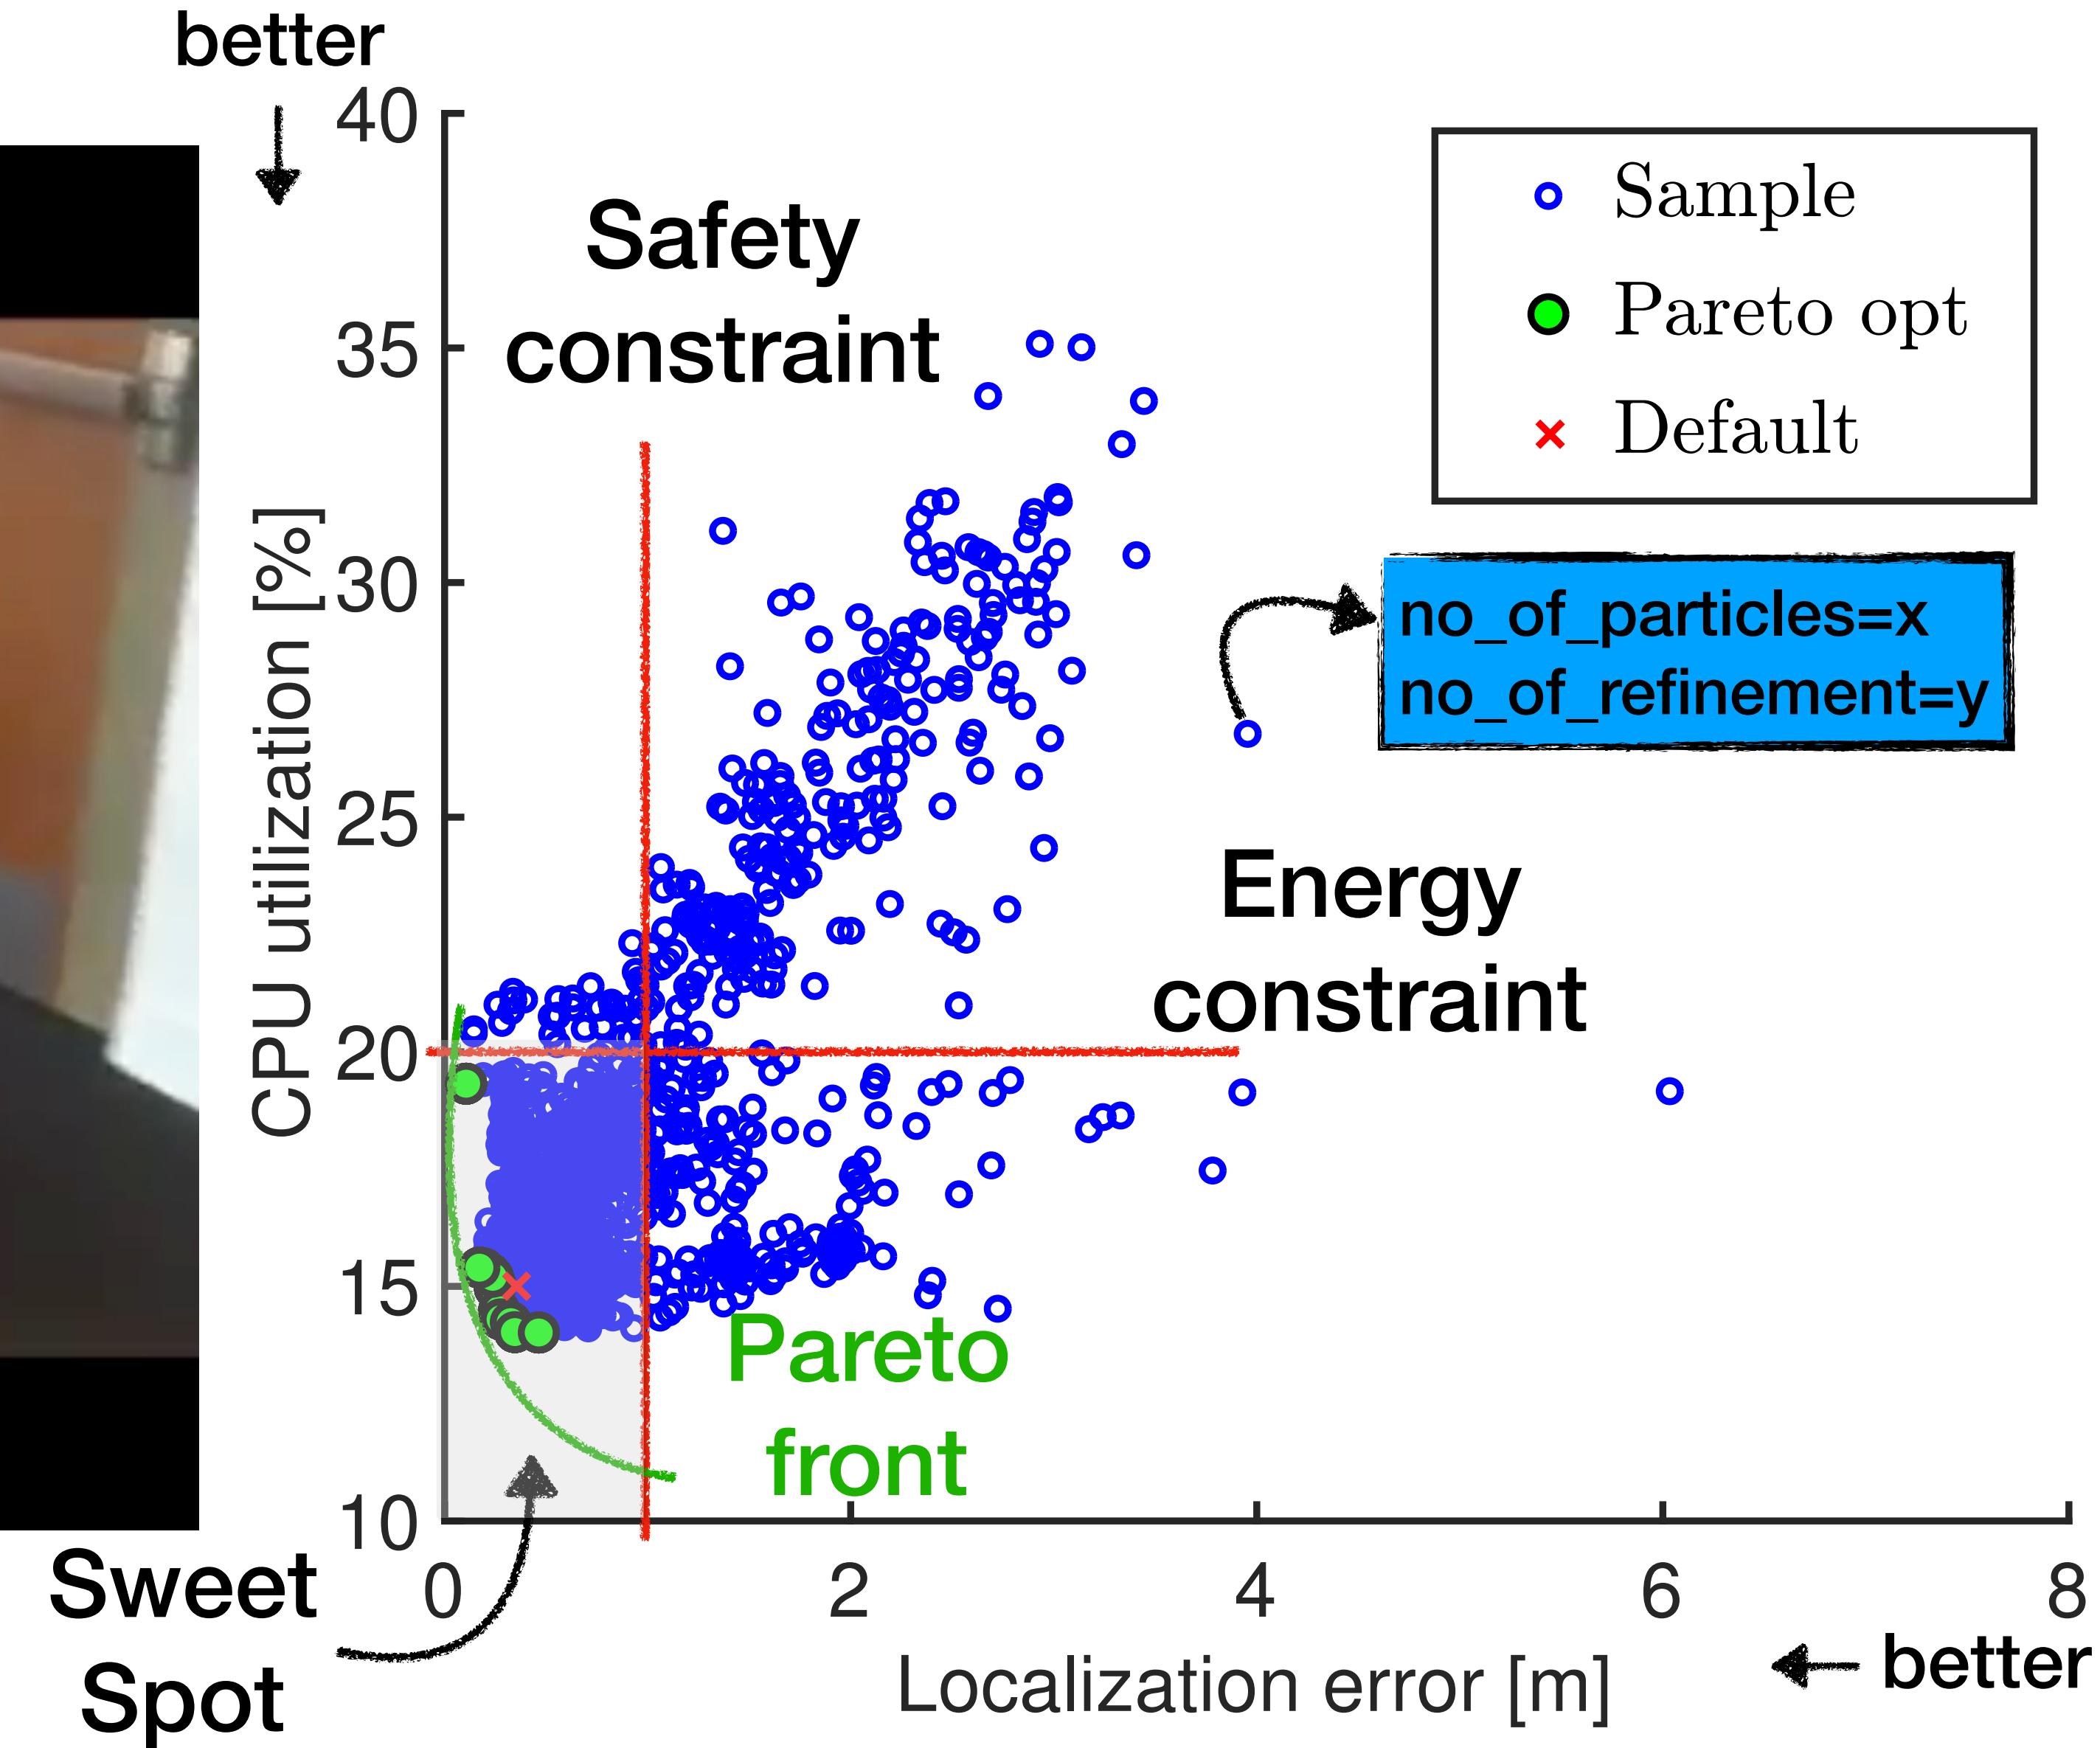
Step 2: Explore the model for regions of most variance

Step 3: Sample that region

Step 4: Repeat



CoBot experiment: DARPA BRASS



Details: [SEAMS '17]

Transfer Learning for Improving Model Predictions in Highly Configurable Software

Pooyan Jamshidi, Miguel Velez, Christian Kästner
Carnegie Mellon University, USA
{pjamshid,mvelezce,kaestner}@cs.cmu.edu

Norbert Siegmund
Bauhaus-University Weimar, Germany
norbert.siegmund@uni-weimar.de

Prasad Kawthekar
Stanford University, USA
pkawthek@stanford.edu

Abstract—Modern software systems are built to be used in dynamic environments using configuration capabilities to adapt to changes and external uncertainties. In a self-adaptation context, we are often interested in reasoning about the performance of the systems under different configurations. Usually, we learn a black-box model based on real measurements to predict the performance of the system given a specific configuration. However, as modern systems become more complex, there are many configuration parameters that may interact and we end up learning an exponentially large configuration space. Naturally, this does not scale when relying on real measurements in the actual changing environment. We propose a different solution:

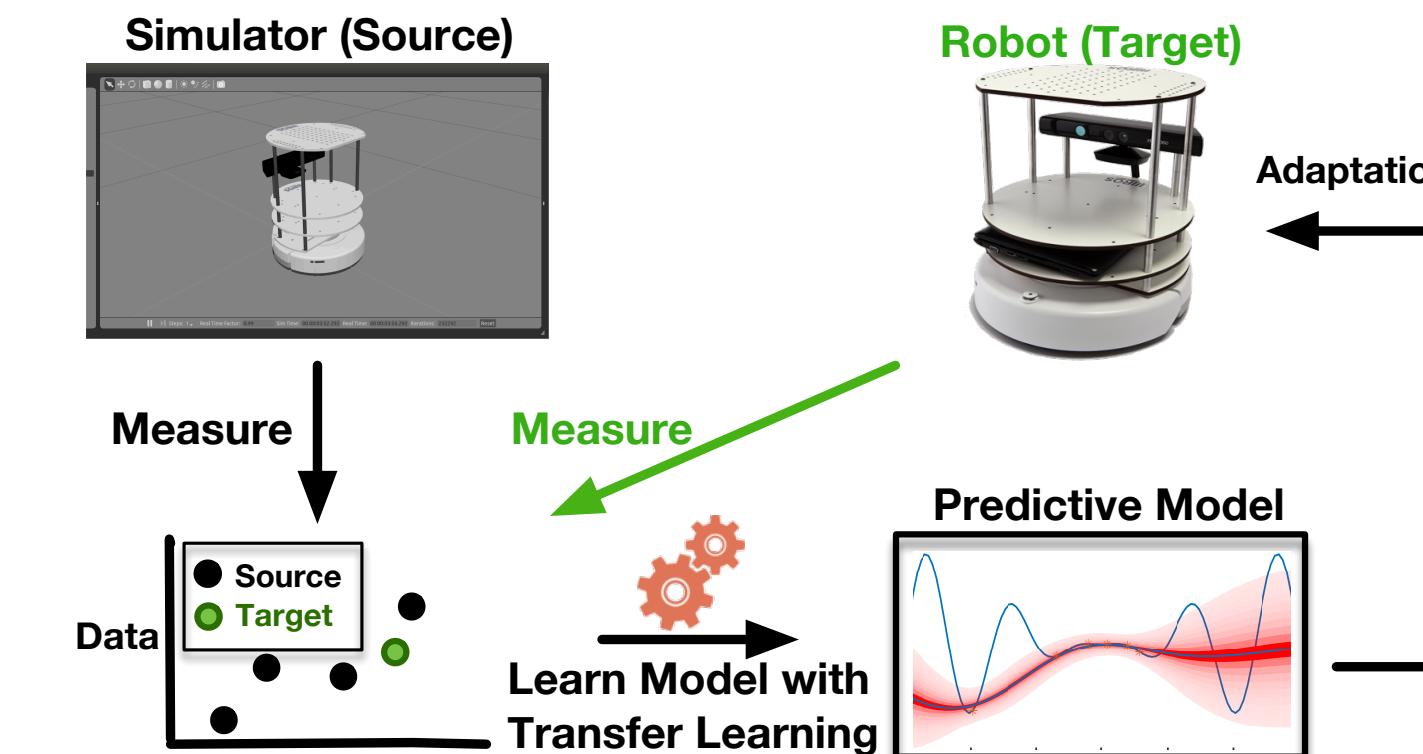
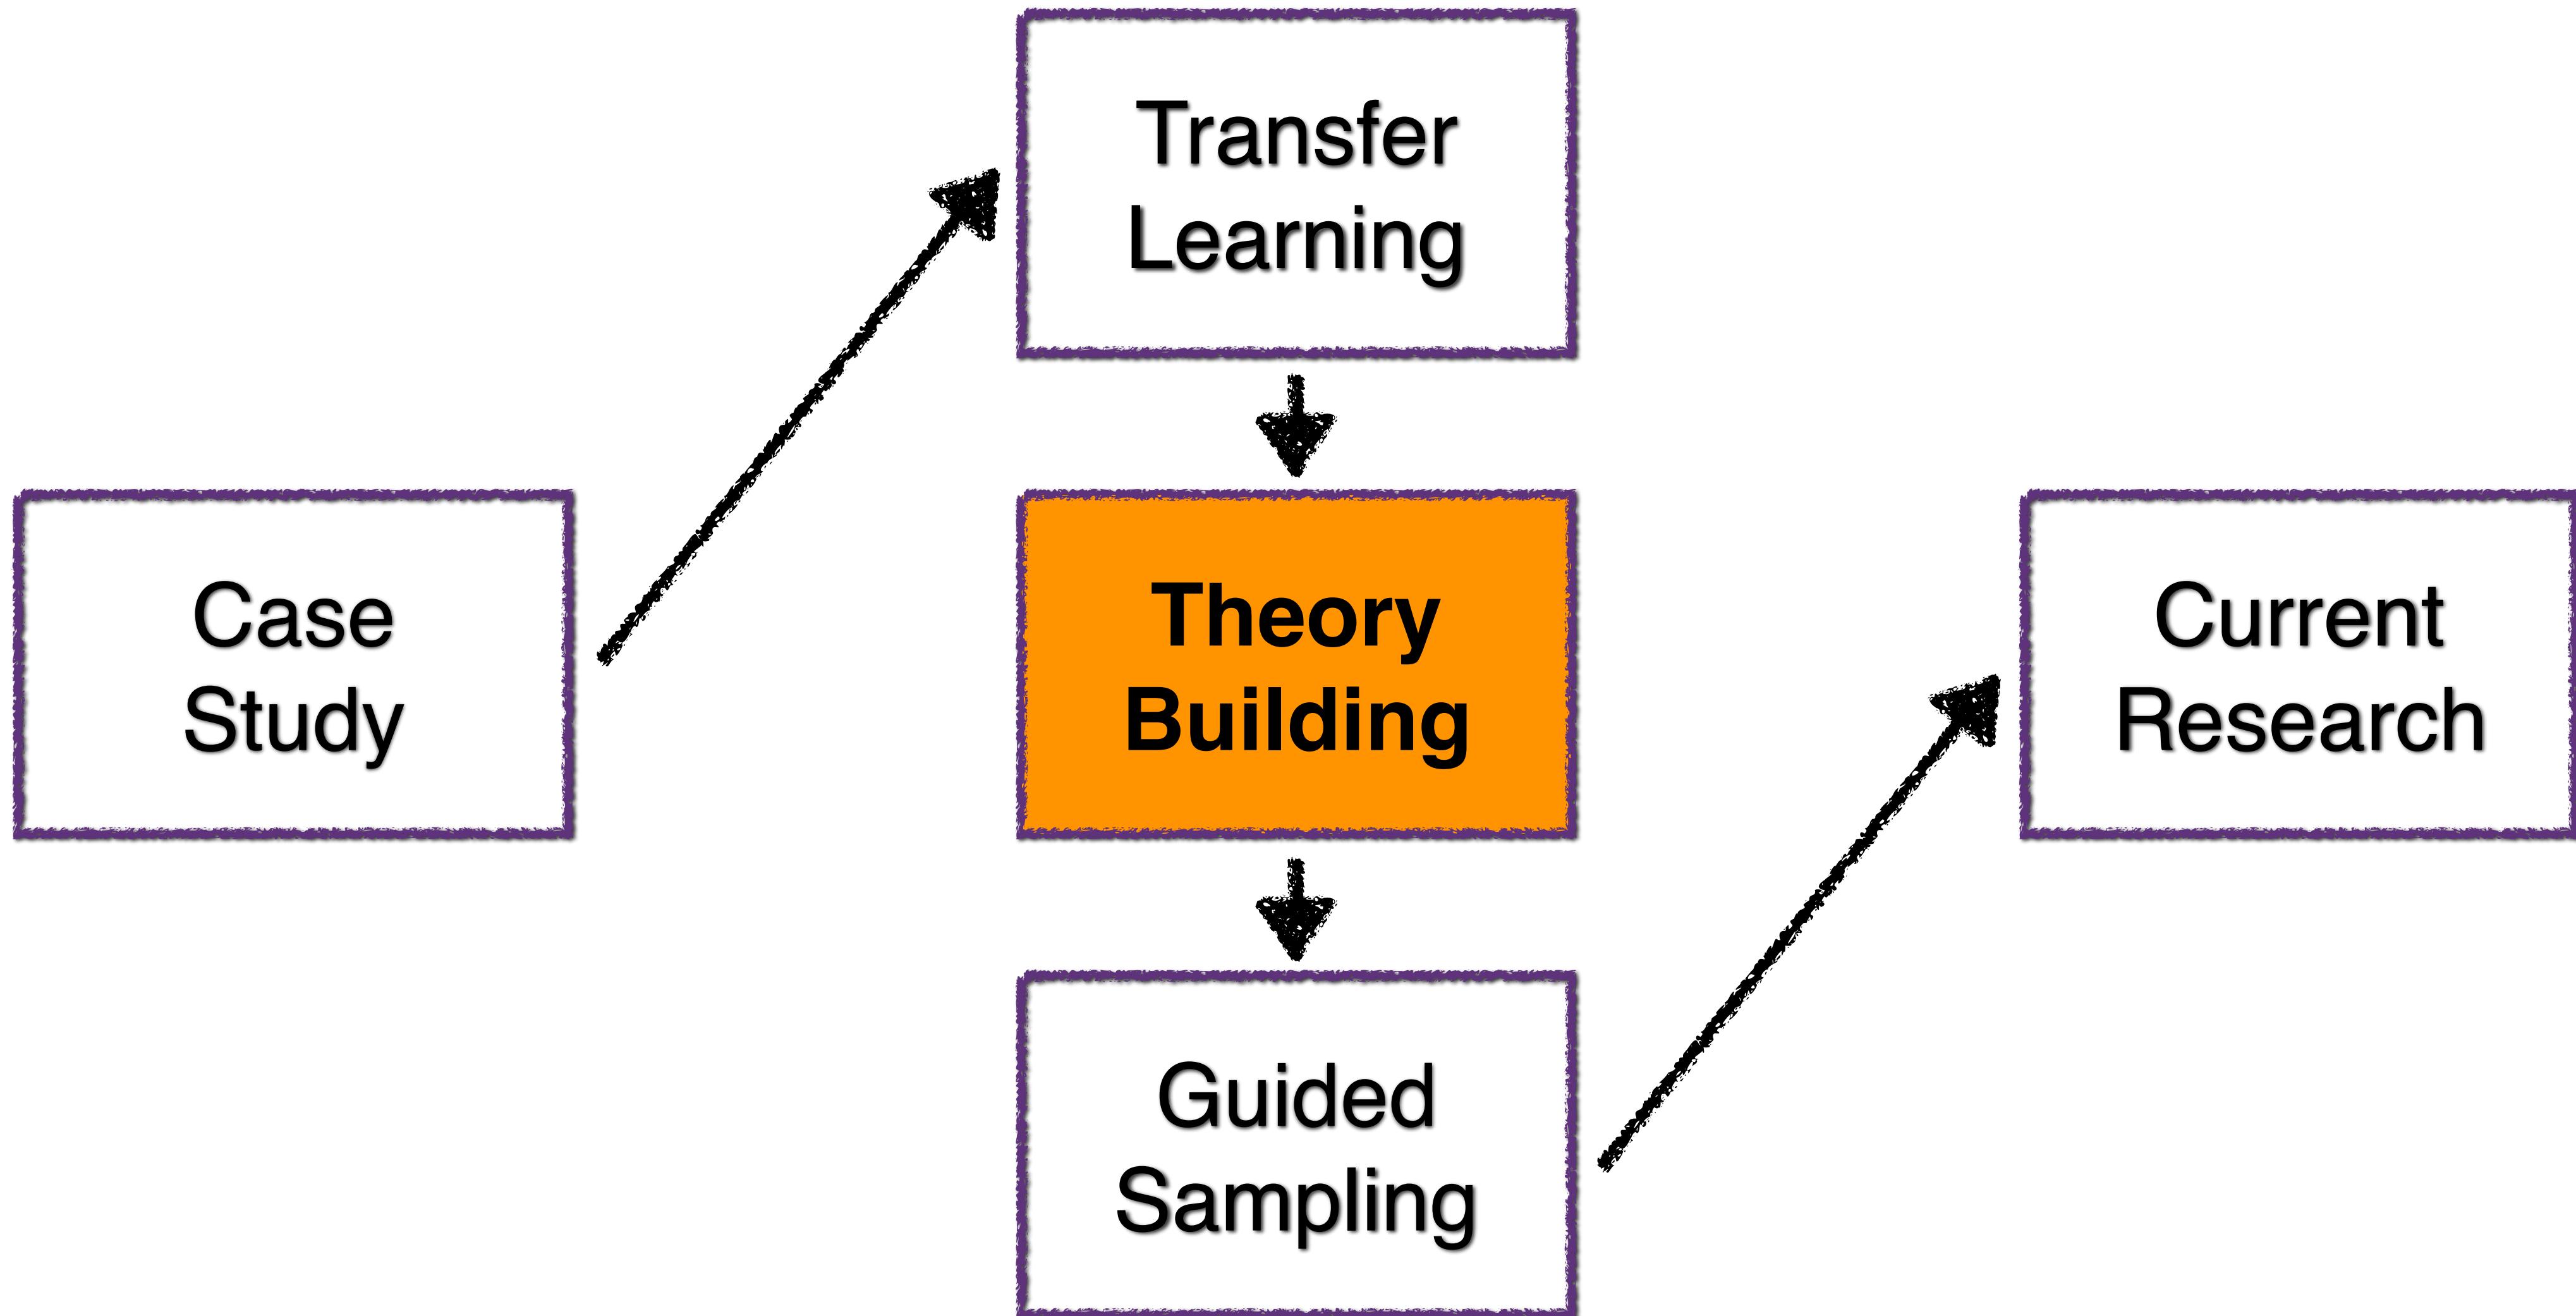


Fig. 1: Transfer learning for performance model learning.

Outline



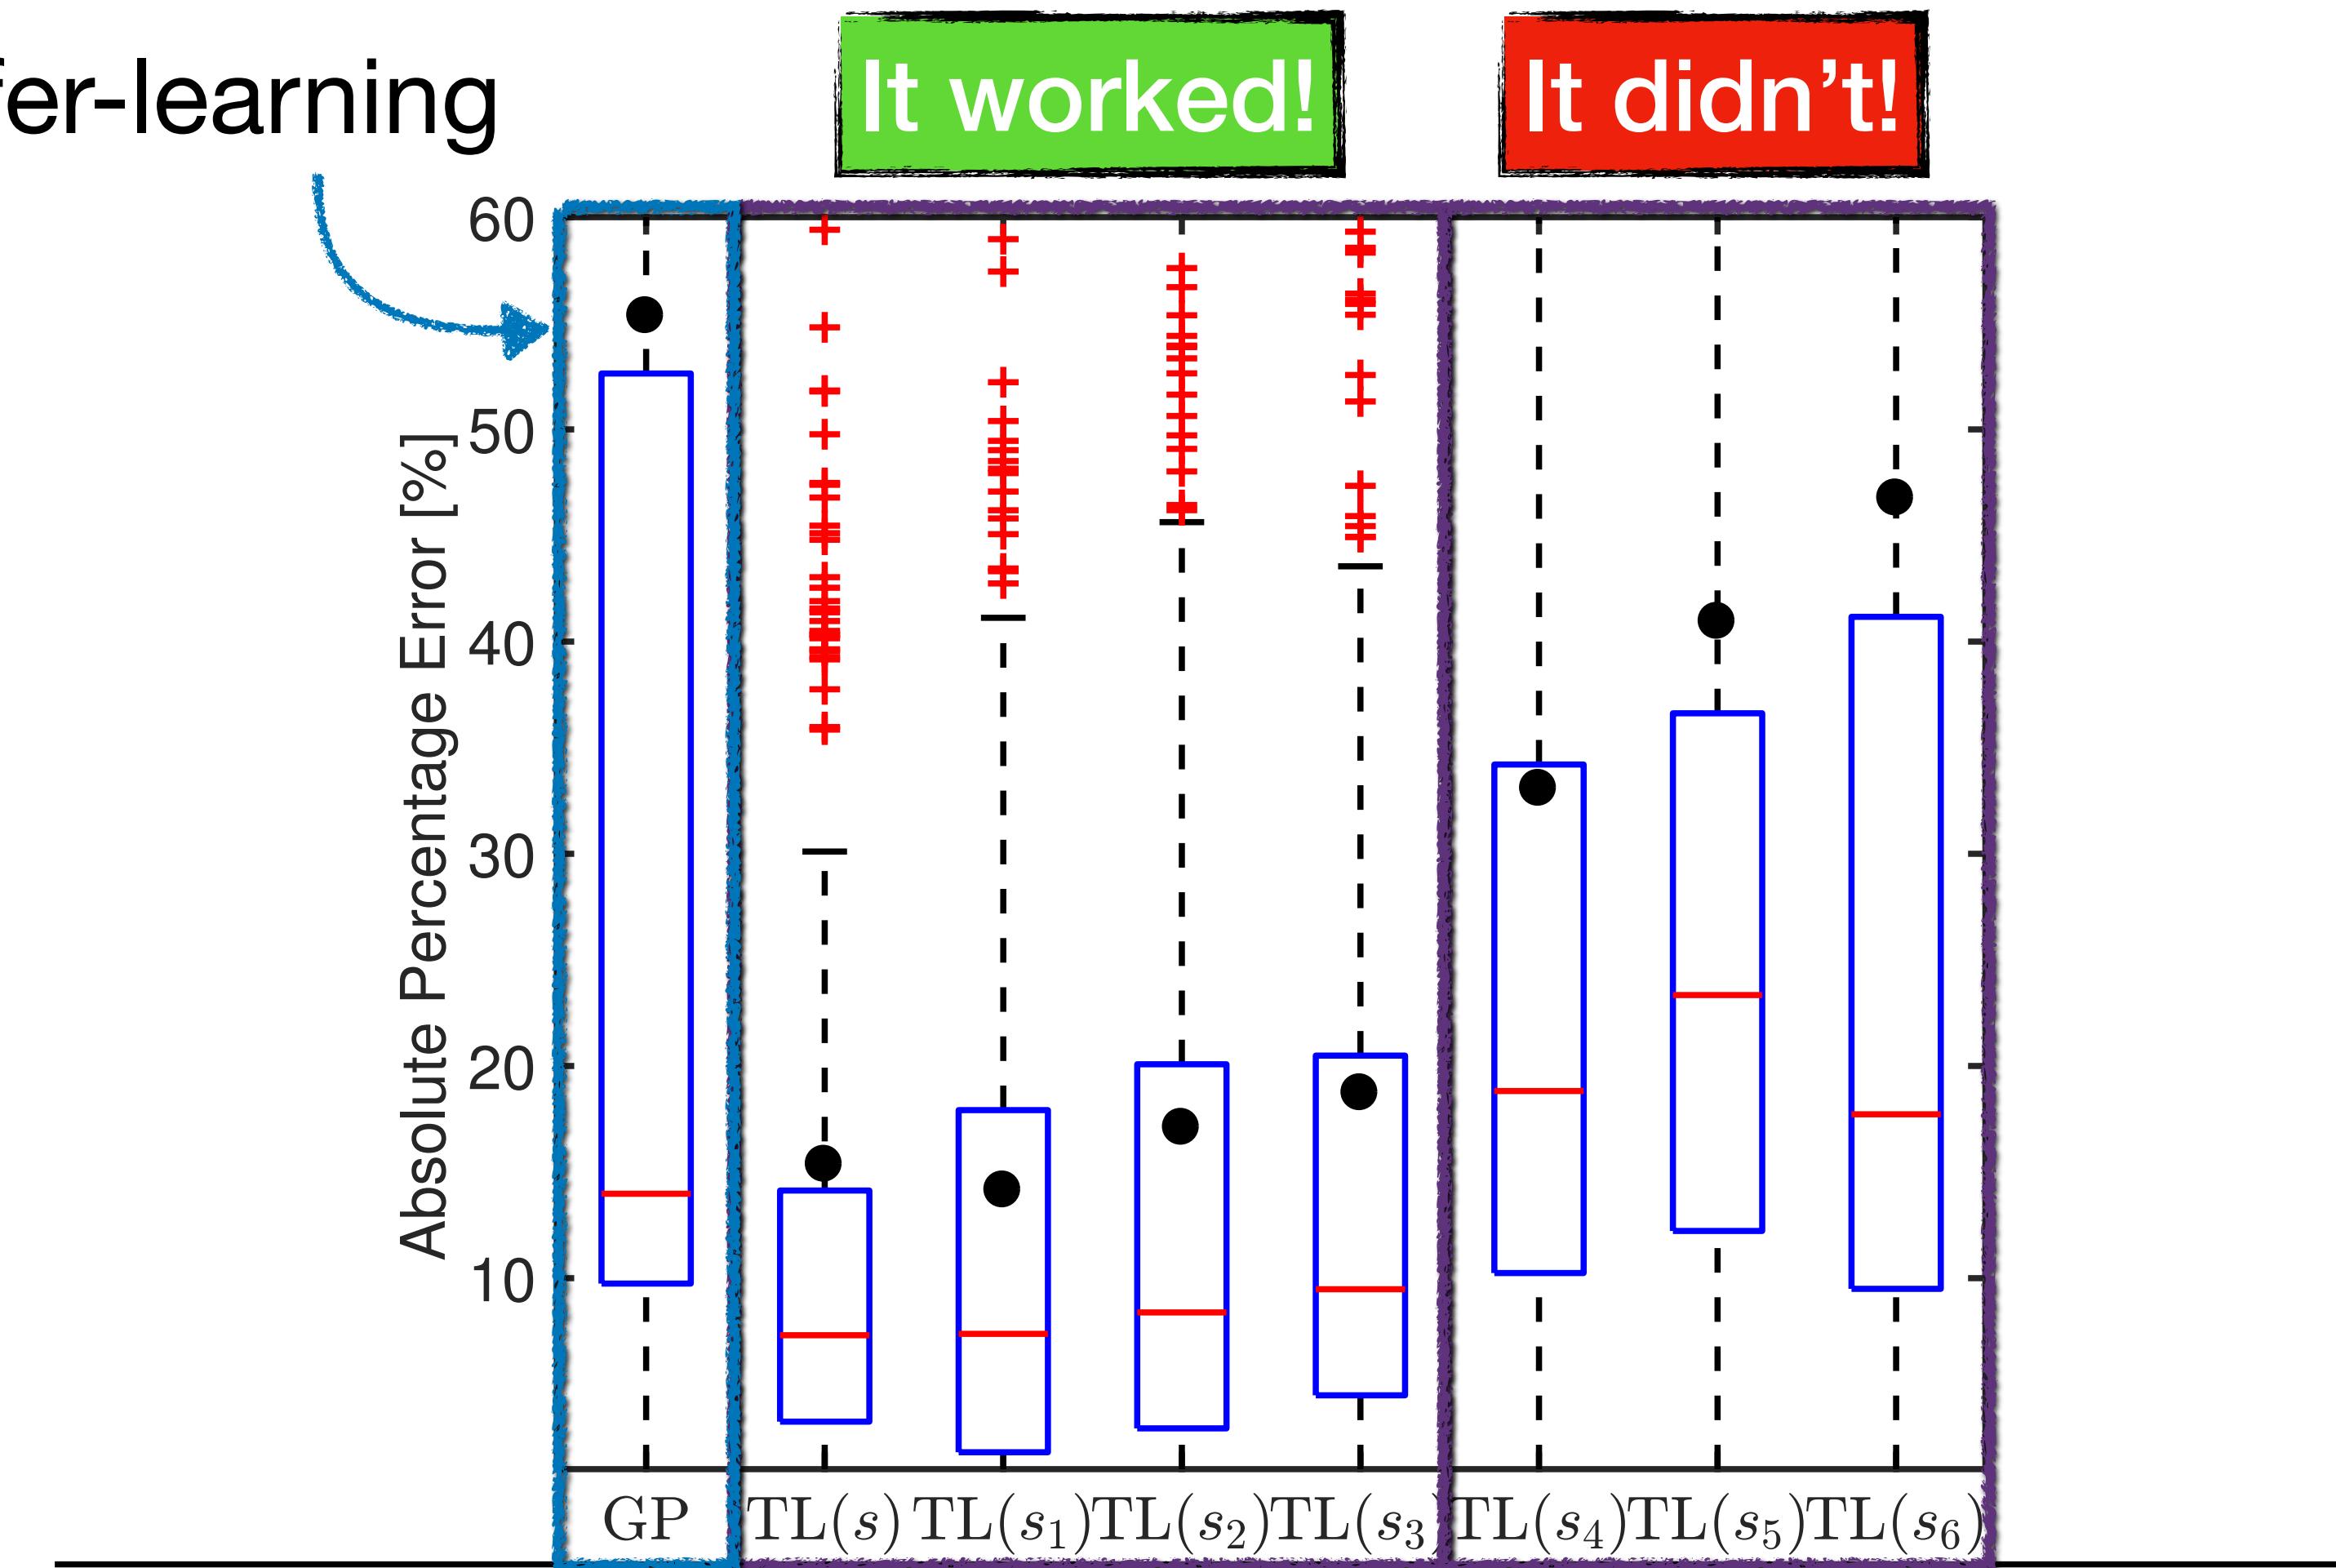
Looking further: When transfer learning goes wrong

Non-transfer-learning

It worked!

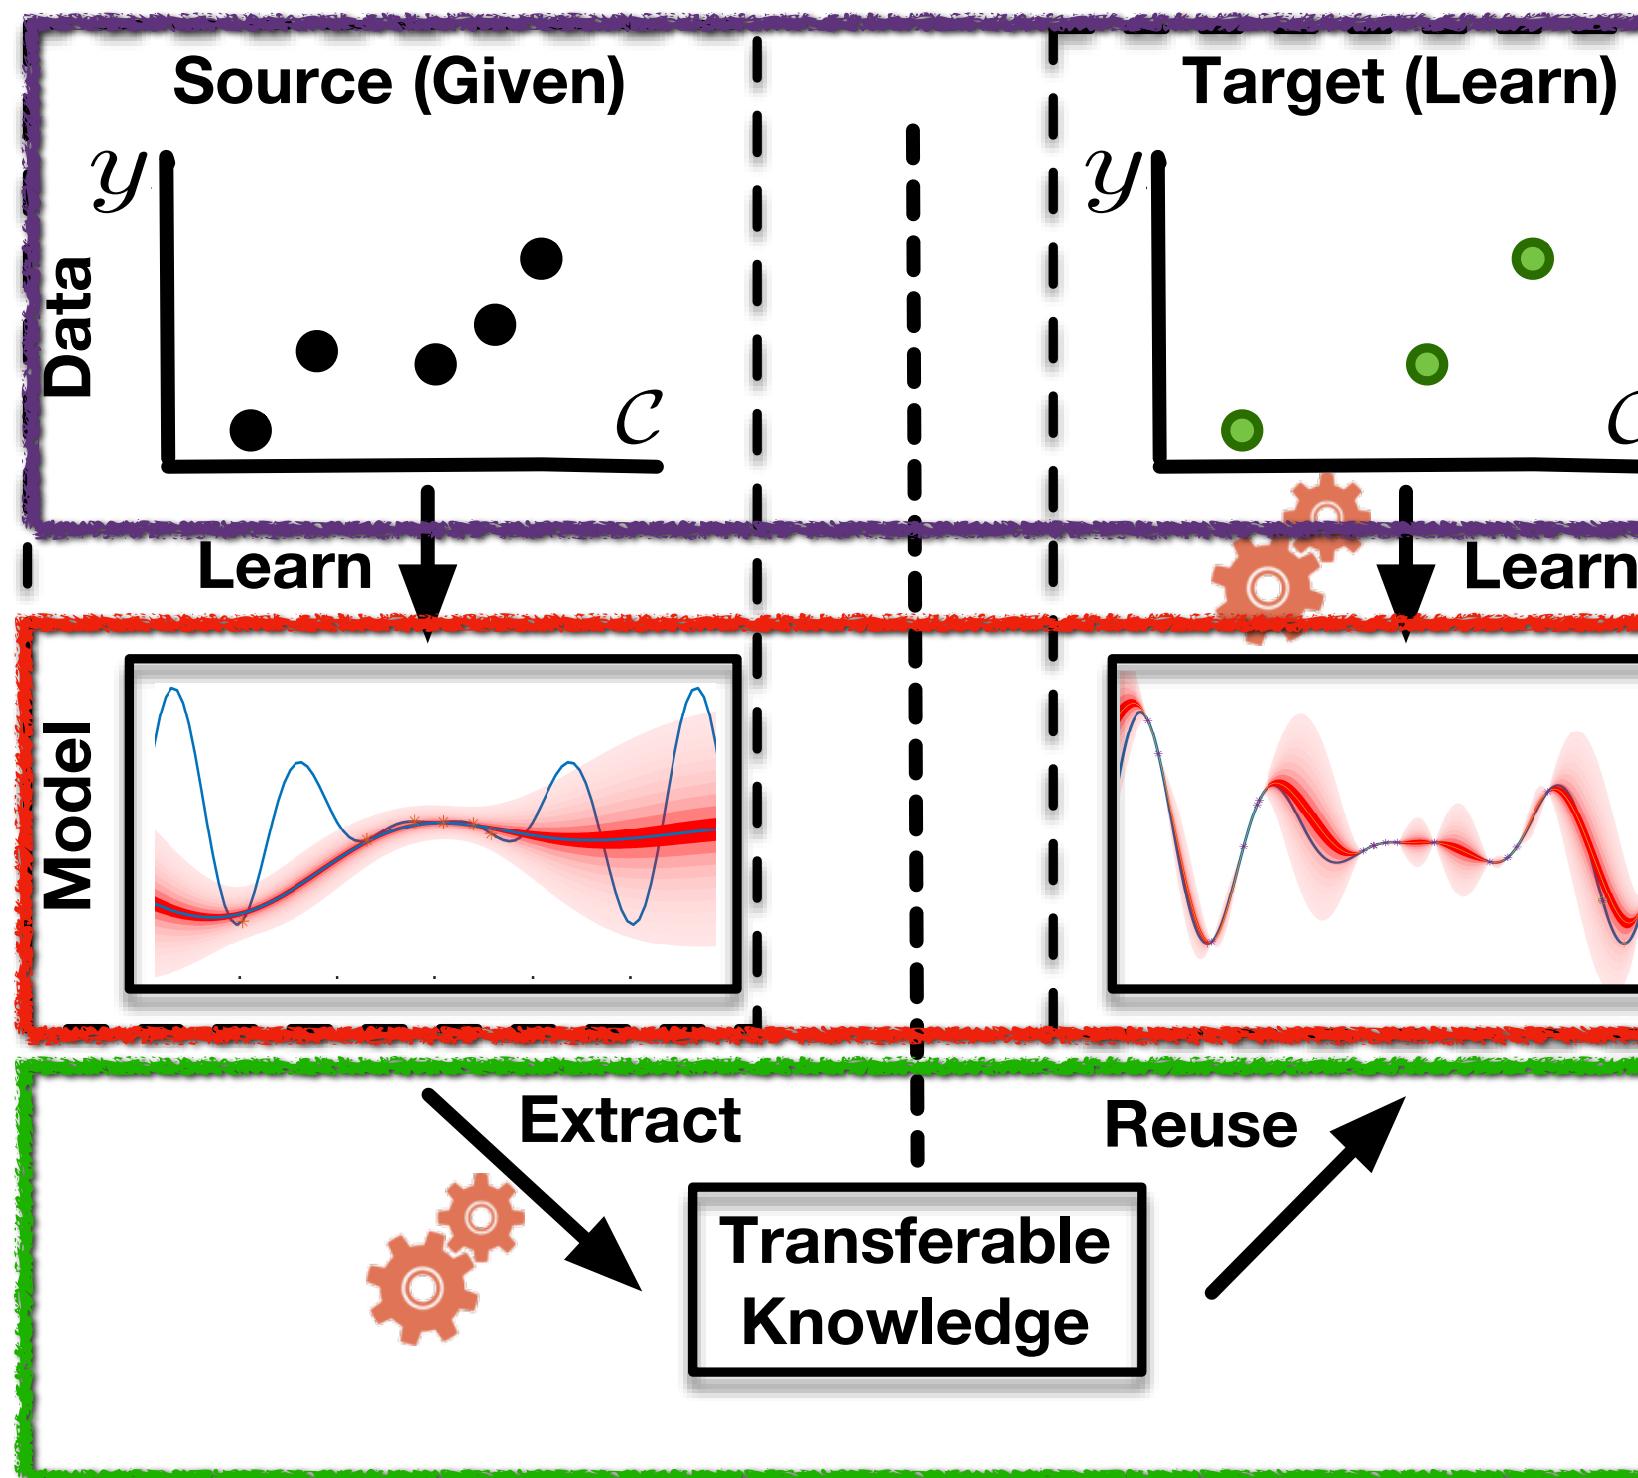
It didn't!

Insight: Predictions become more accurate when the source is **more related** to the target.



Sources	s	s_1	s_2	s_3	s_4	s_5	s_6
noise-level	0	5	10	15	20	25	30
corr. coeff.	0.98	0.95	0.89	0.75	0.54	0.34	0.19
$\mu(pe)$	15.34	14.14	17.09	18.71	33.06	40.93	46.75

Key question: Can we develop a theory to explain when transfer learning works?



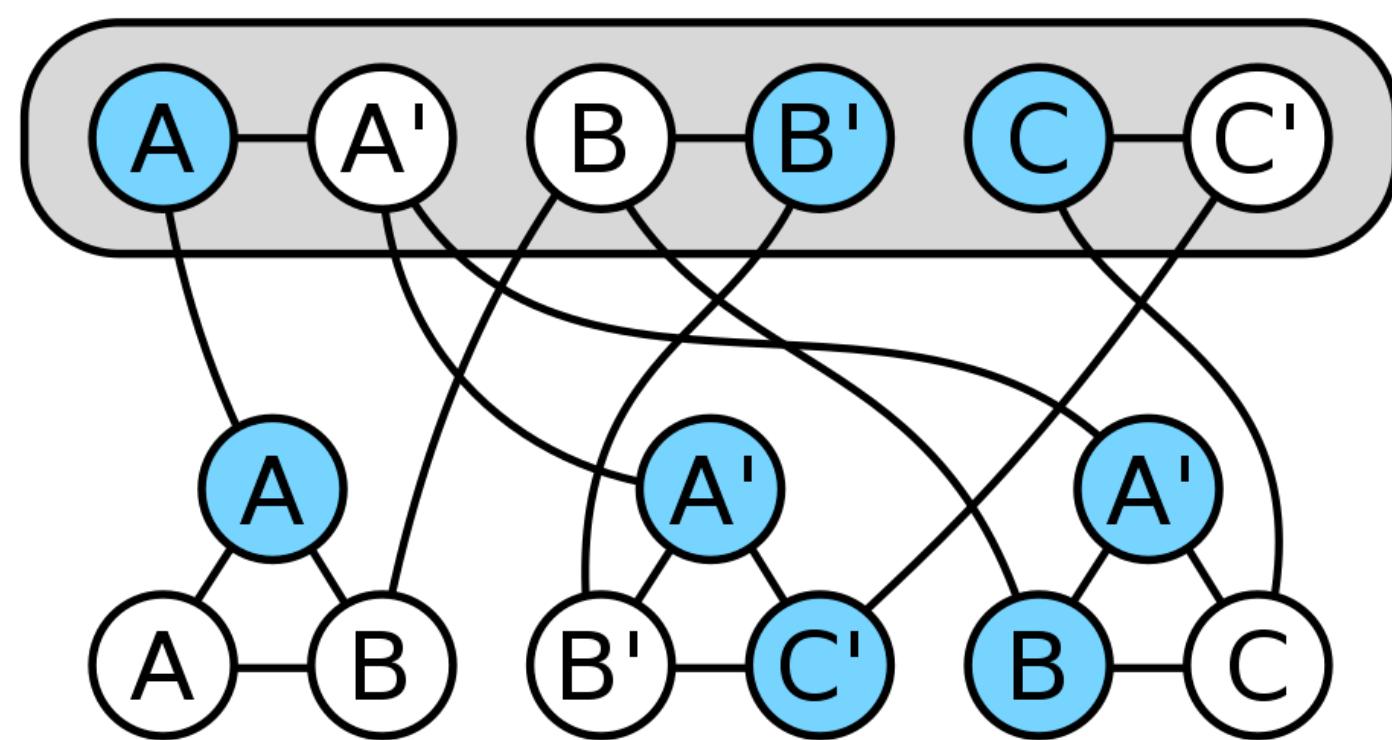
Q1: How source and target are “related”?

Q2: What characteristics are preserved?

Q3: What are the actionable insights?

Our empirical study: We looked at different highly-configurable systems to gain insights

$$(A \vee B) \wedge (\neg A \vee \neg B \vee \neg C) \wedge (\neg A \vee B \vee C)$$



SPEAR (SAT Solver)

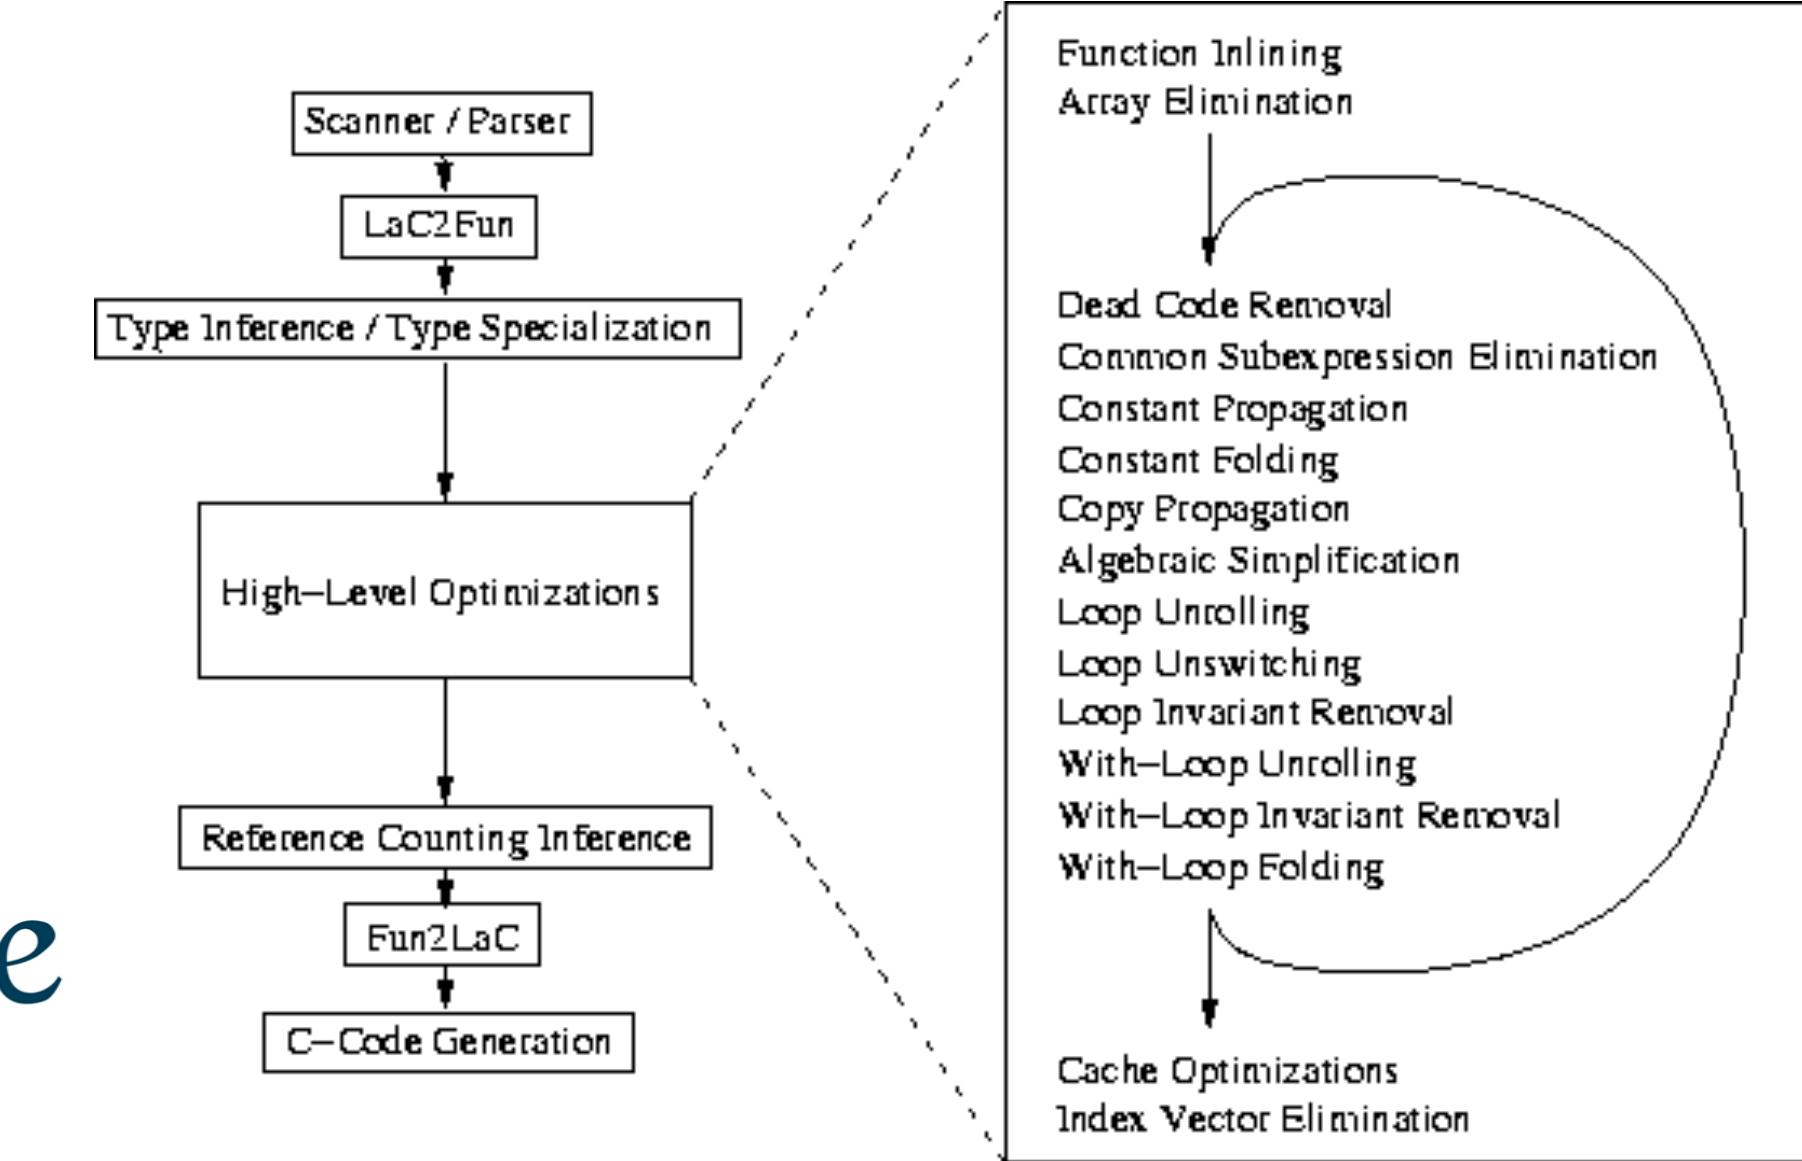
Analysis time

14 options

X264 (video encoder)

Encoding time

16 options



SQLite (DB engine)

Query time

14 options

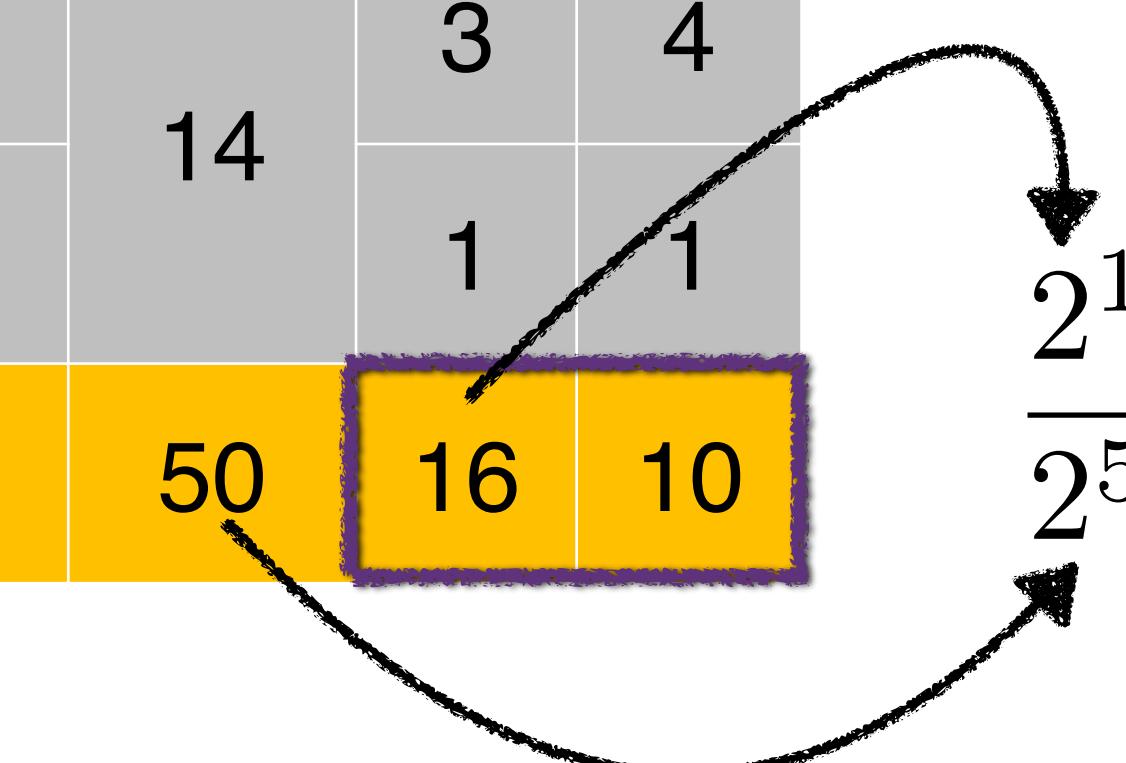
SaC (Compiler)

Execution time

50 options

Influential options and interactions are preserved across environments

Soft	Environmental change	Severity	Dim	t-test	
x264	Version	Large	16	12	10
	Hardware/workload/ver	V Large		8	9
SQLite	write-seq -> write-batch	V Large	14	3	4
	read-rand -> read-seq	Medium		1	1
SaC	Workload	V Large	50	16	10


$$\frac{2^{16}}{2^{50}} = 0.00000000058$$

We only need to explore part of the space:

Implication: Avoid wasting budget on non-informative part of configuration space and focusing where it matters.

Details: [ASE '17]

Transfer Learning for Performance Modeling of Configurable Systems: An Exploratory Analysis

Pooyan Jamshidi
Carnegie Mellon University, USA

Norbert Siegmund
Bauhaus-University Weimar, Germany

Miguel Velez, Christian Kästner
Akshay Patel, Yuvraj Agarwal
Carnegie Mellon University, USA

Abstract—Modern software systems provide many configuration options which significantly influence their non-functional properties. To understand and predict the effect of configuration options, several sampling and learning strategies have been proposed, albeit often with significant cost to cover the highly dimensional configuration space. Recently, transfer learning has been applied to reduce the effort of constructing performance models by transferring knowledge about performance behavior across environments. While this line of research is promising to learn more accurate models at a lower cost, it is unclear why and when transfer learning works for performance modeling. To shed light on when it is beneficial to apply transfer learning, we conducted an empirical study on four popular software systems, varying software configurations and environmental conditions, such as hardware, workload, and software versions, to identify the key knowledge pieces that can be exploited for transfer learning. Our results show that in small environmental changes (e.g., homogeneous workload change), by applying a linear transformation to the performance model, we can understand the performance behavior of the target environment, while for severe environmental changes (e.g., drastic workload change) we

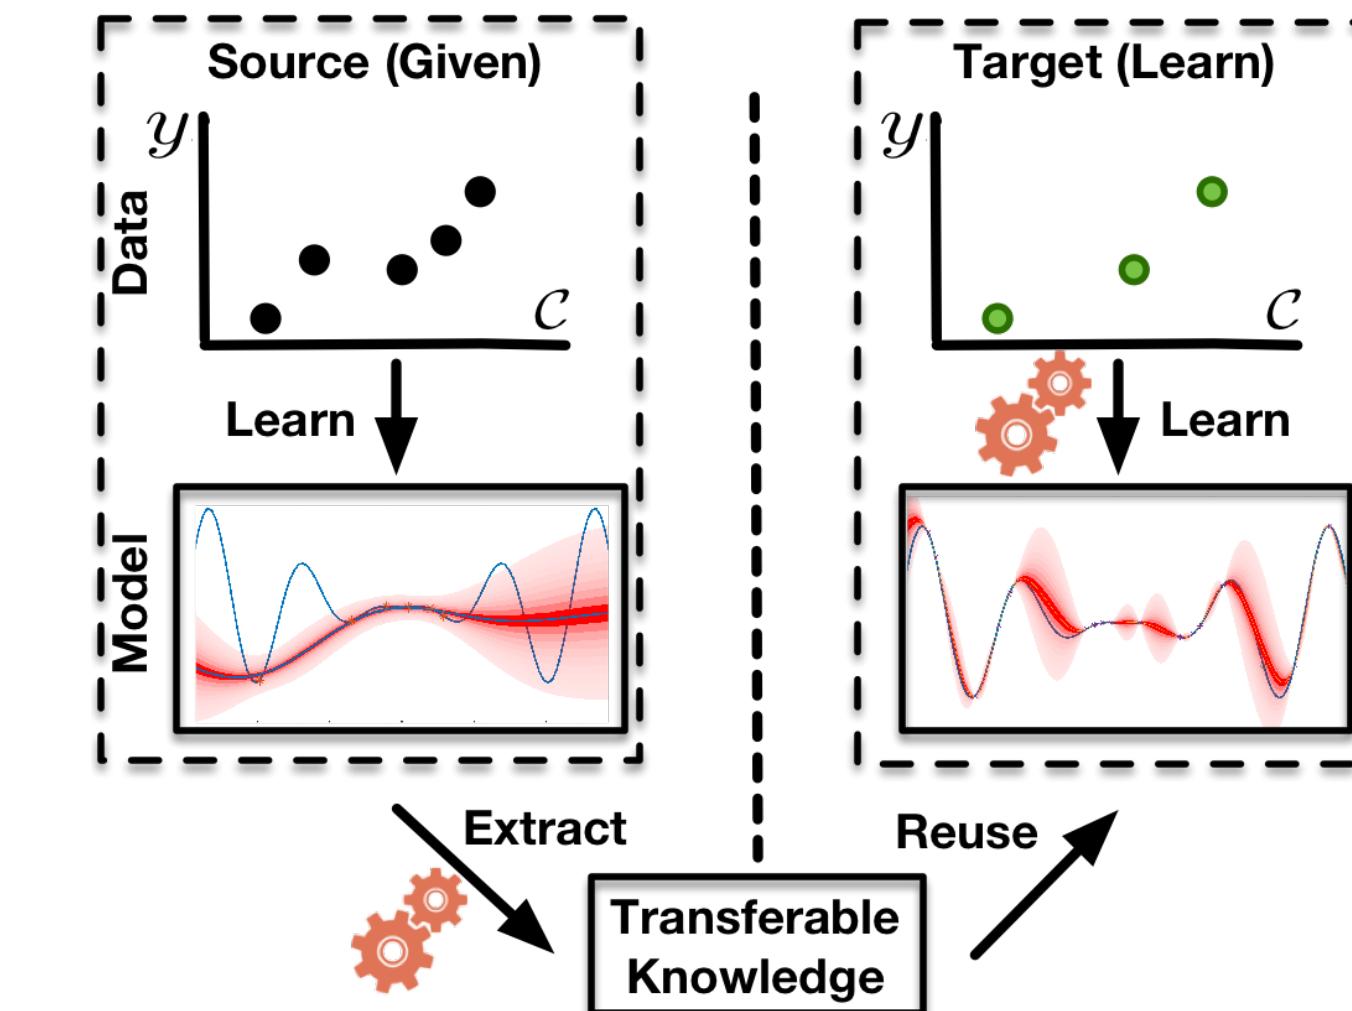
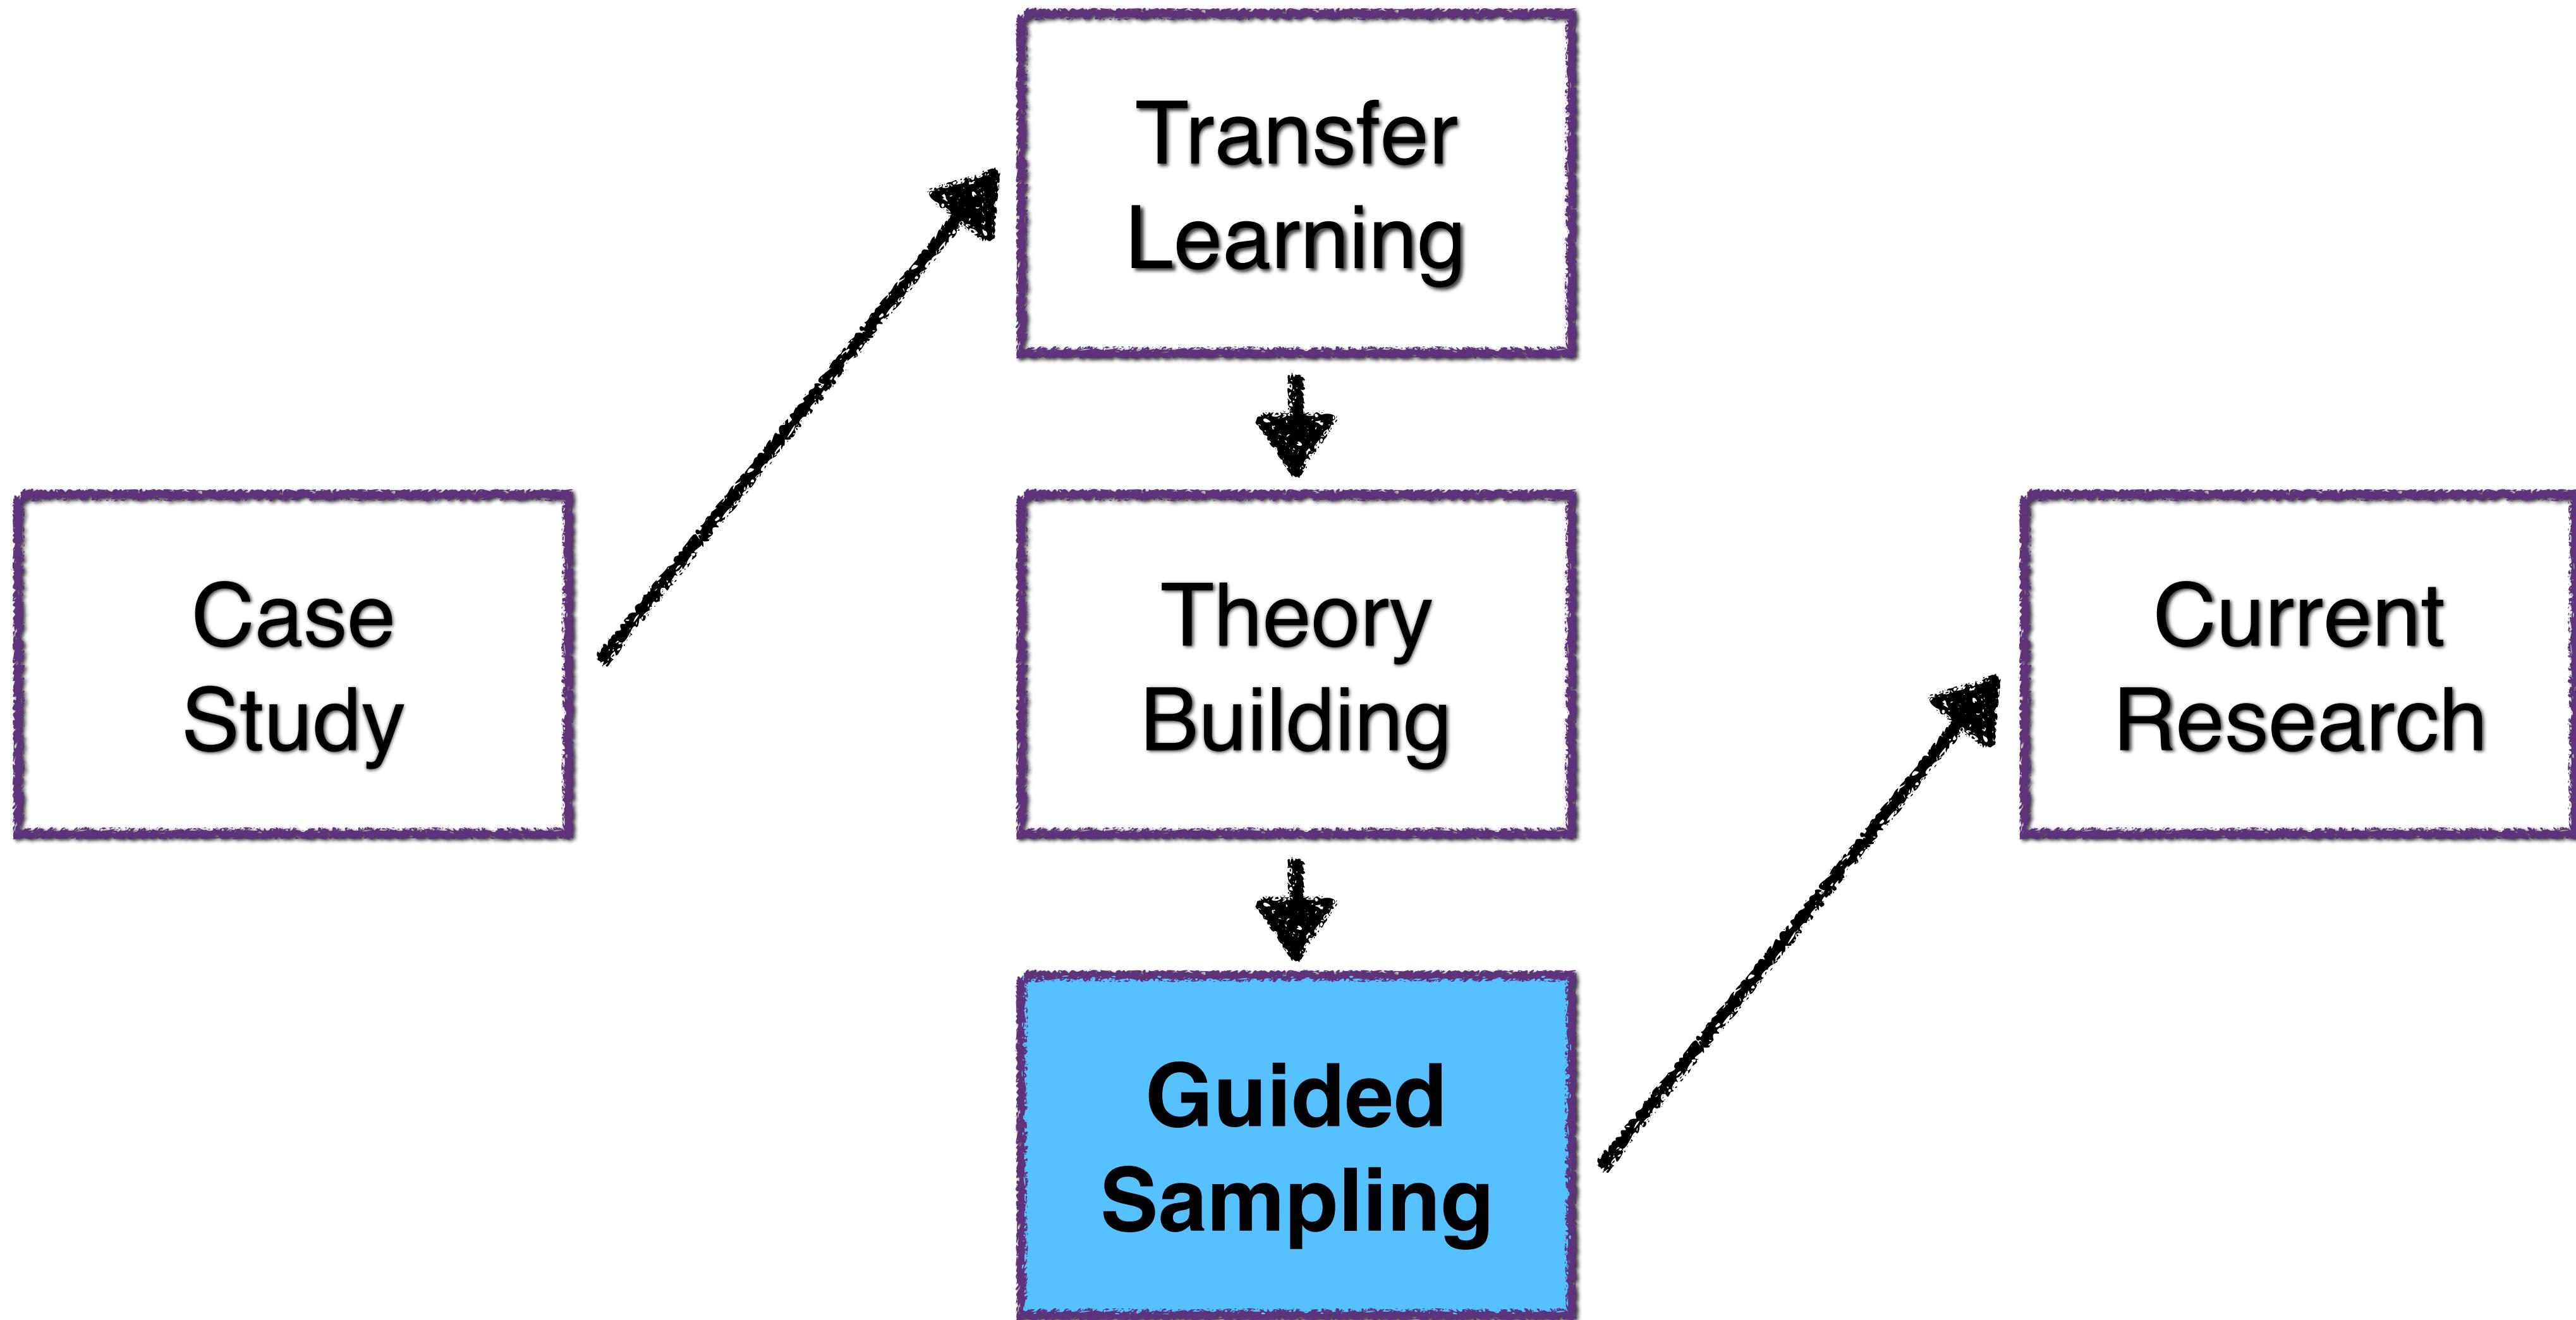
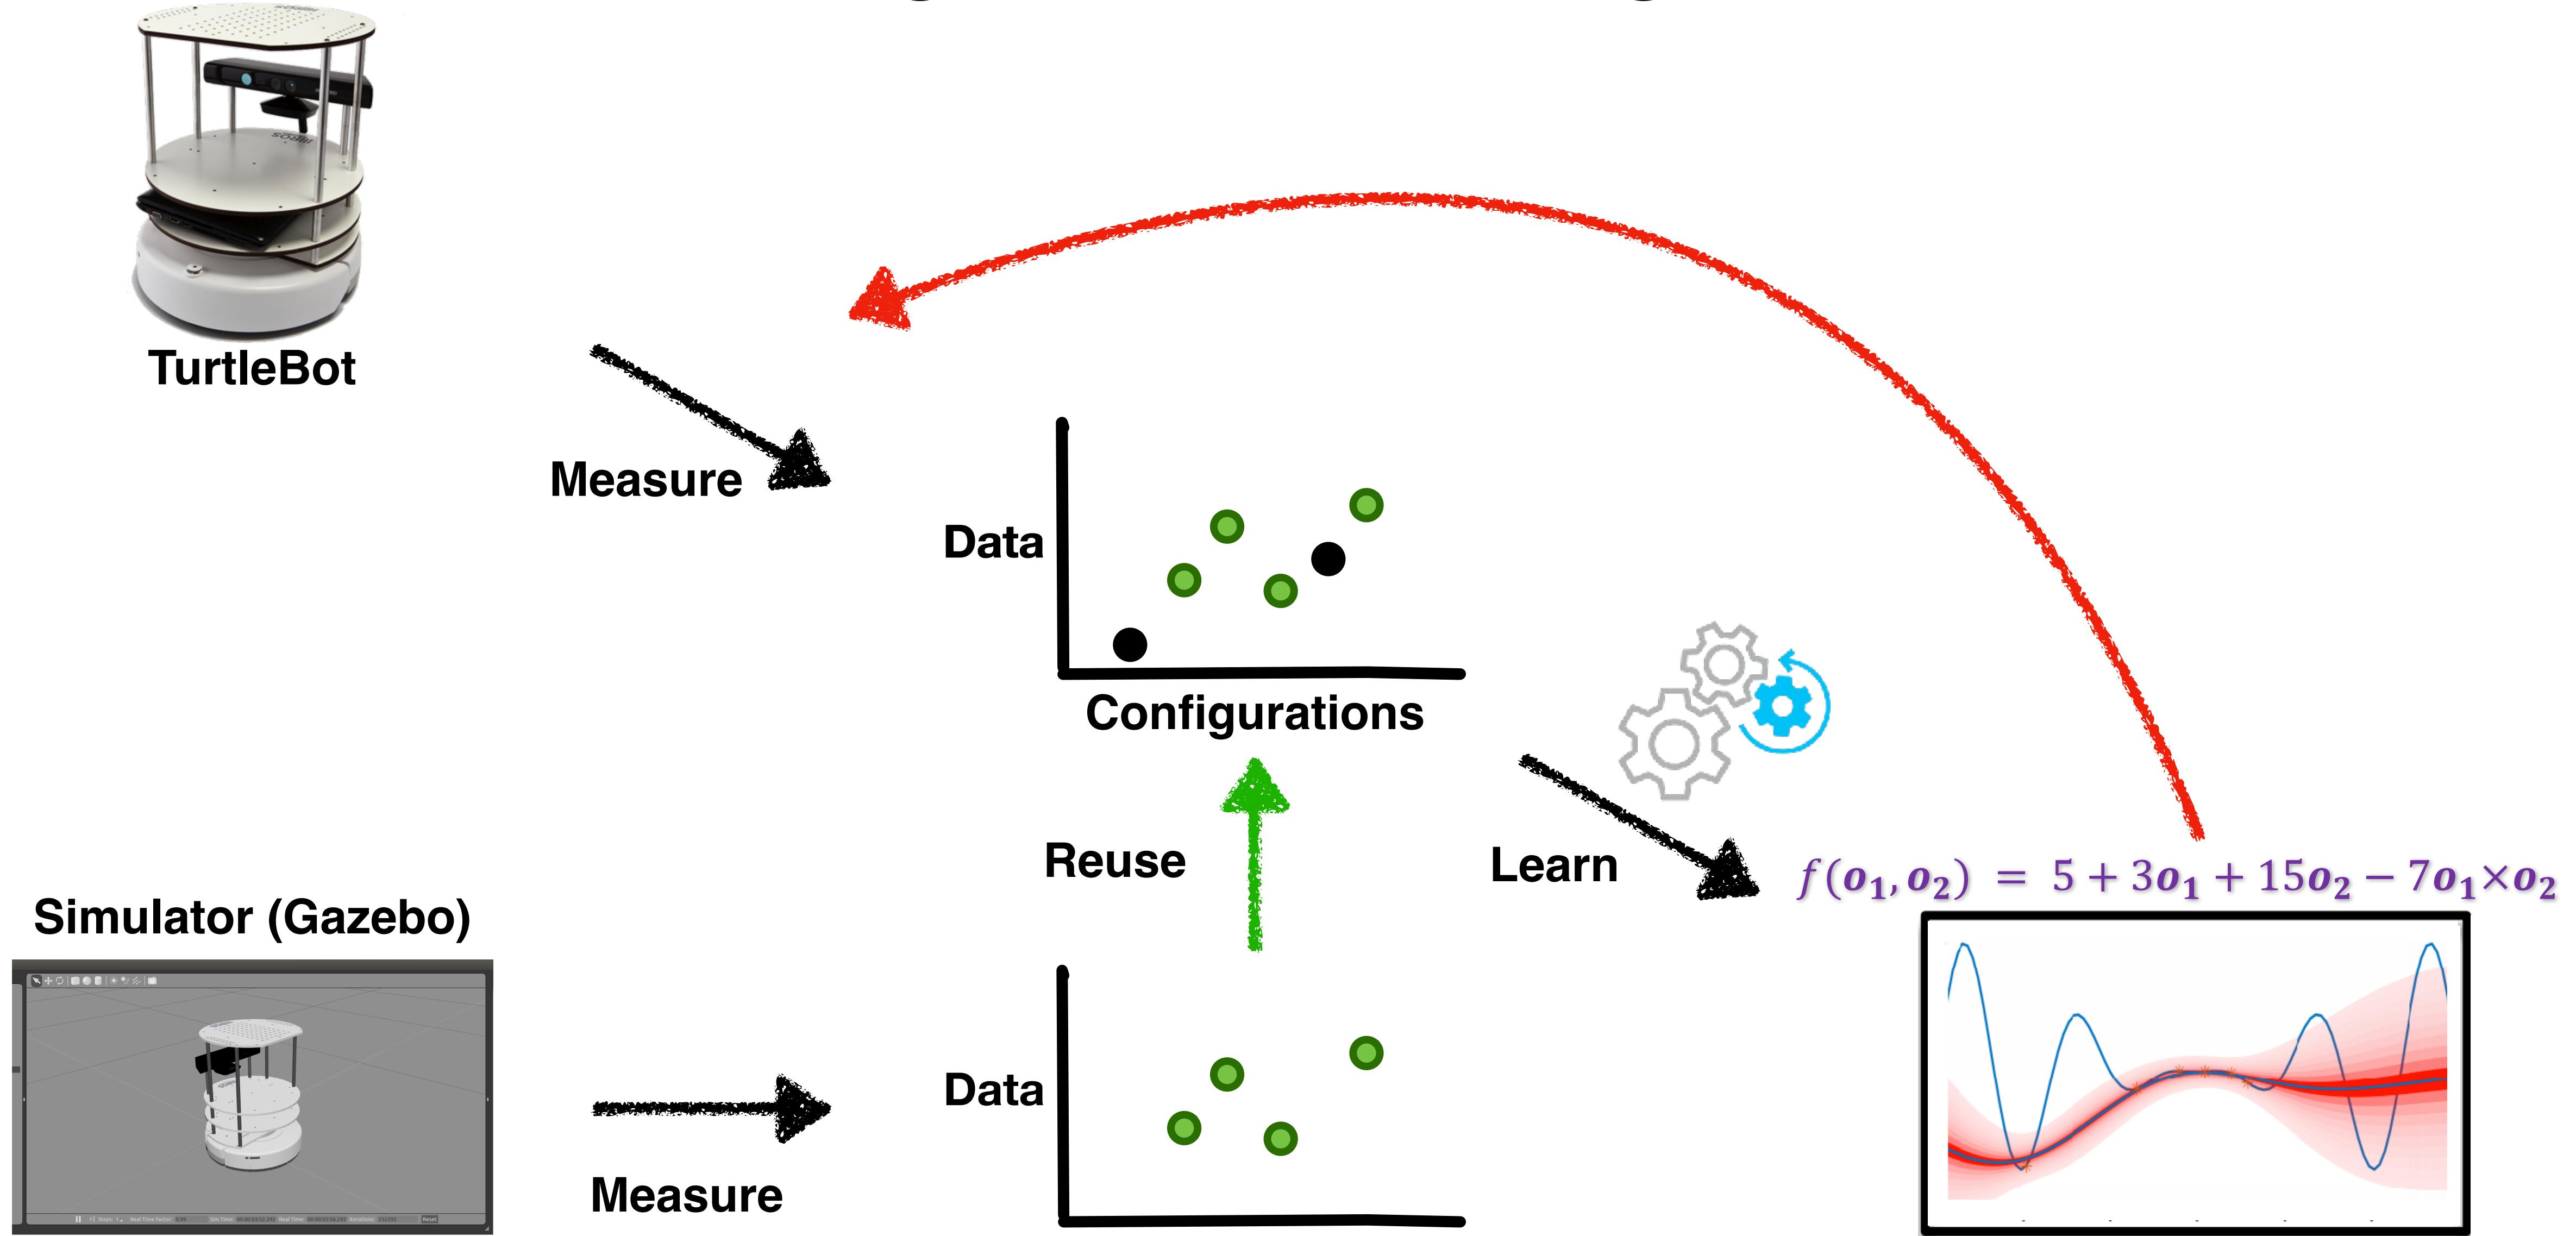


Fig. 1: Transfer learning is a form of machine learning that takes advantage of transferable knowledge from source to learn an accurate, reliable, and less costly model for the target environment.

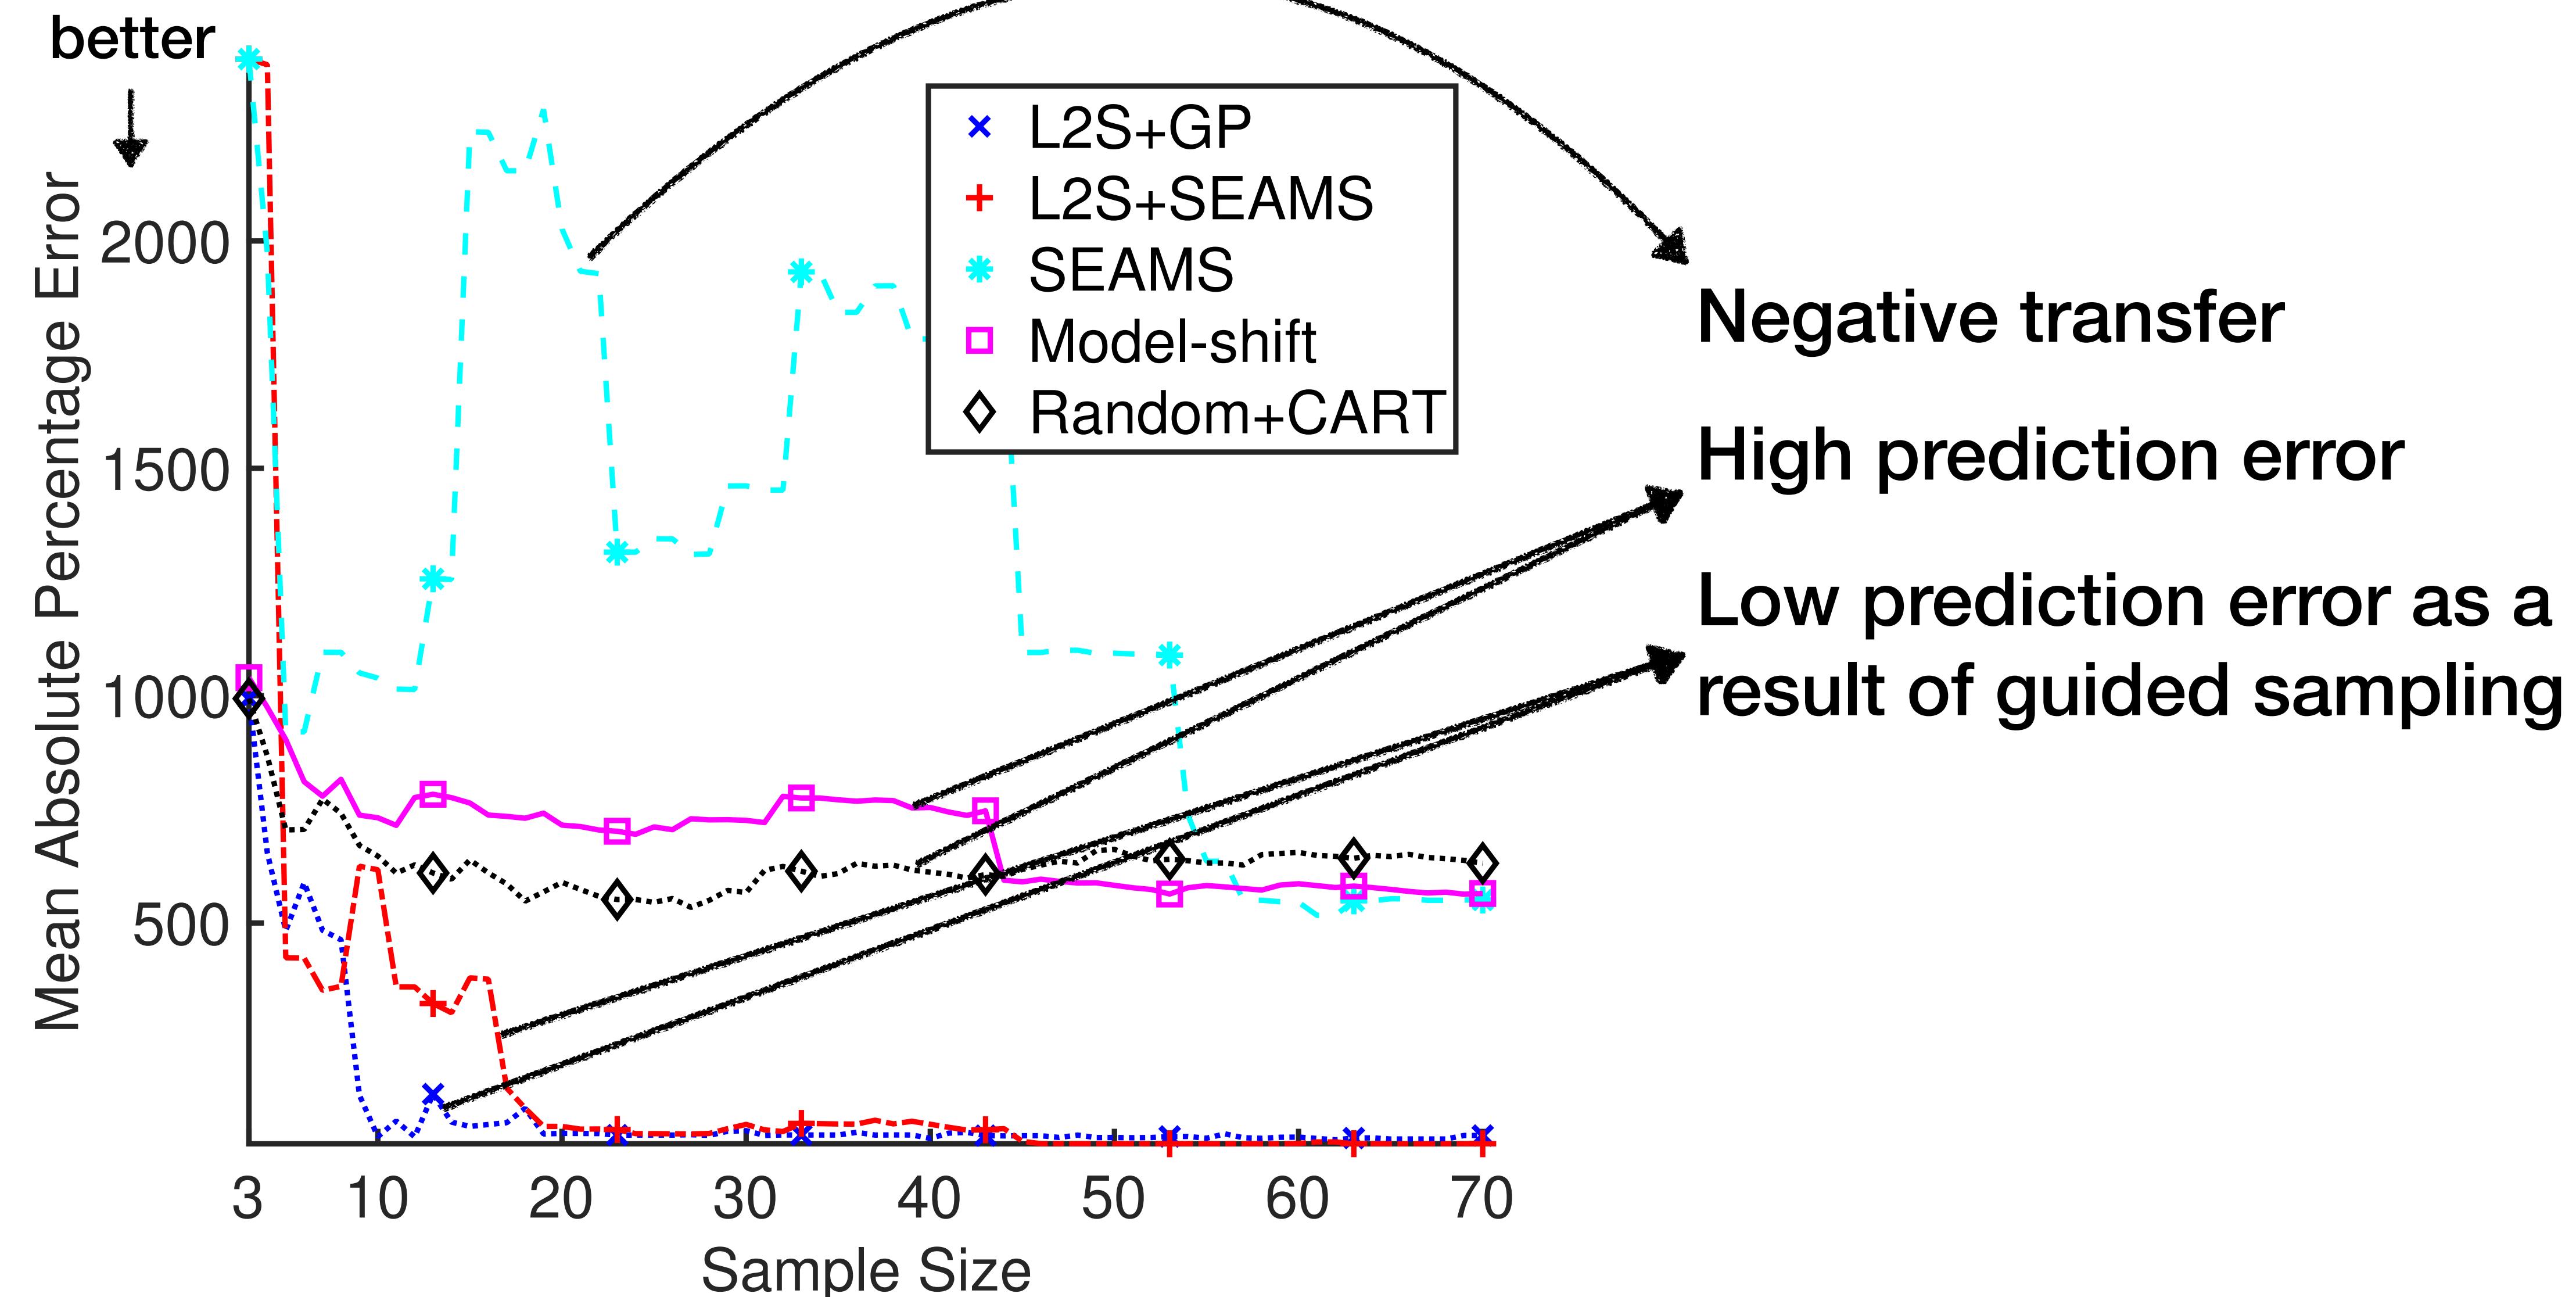
Outline



Insights from our empirical study lead to the development of a guided sampling



Simple transfer learning do not work in severe changes



Details: [FSE '18]

Learning to Sample: Exploiting Similarities across Environments to Learn Performance Models for Configurable Systems

Pooyan Jamshidi
University of South Carolina
USA

Miguel Velez
Christian Kästner
Carnegie Mellon University
USA

Norbert Siegmund
Bauhaus-University Weimar
Germany

ABSTRACT

Most software systems provide options that allow users to tailor the system in terms of functionality and qualities. The increased flexibility raises challenges for understanding the configuration space and the effects of options and their interactions on performance and other non-functional properties. To identify how options and interactions affect the performance of a system, several sampling and learning strategies have been recently proposed. However, existing approaches usually assume a fixed environment (hardware, workload, software release) such that learning has to be repeated once the environment changes. Repeating learning and measurement for each environment is expensive and often practically infeasible. Instead, we pursue a strategy that transfers knowledge across environments but sidesteps heavyweight and expensive transfer-

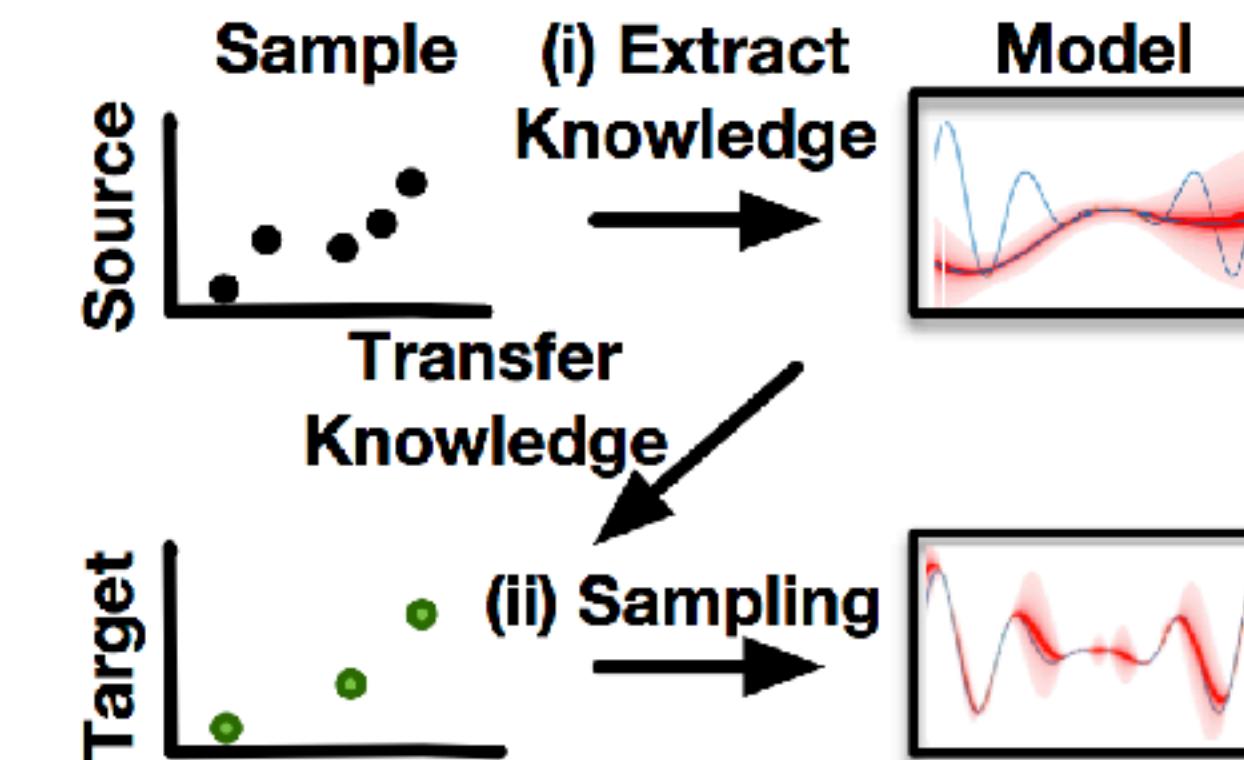
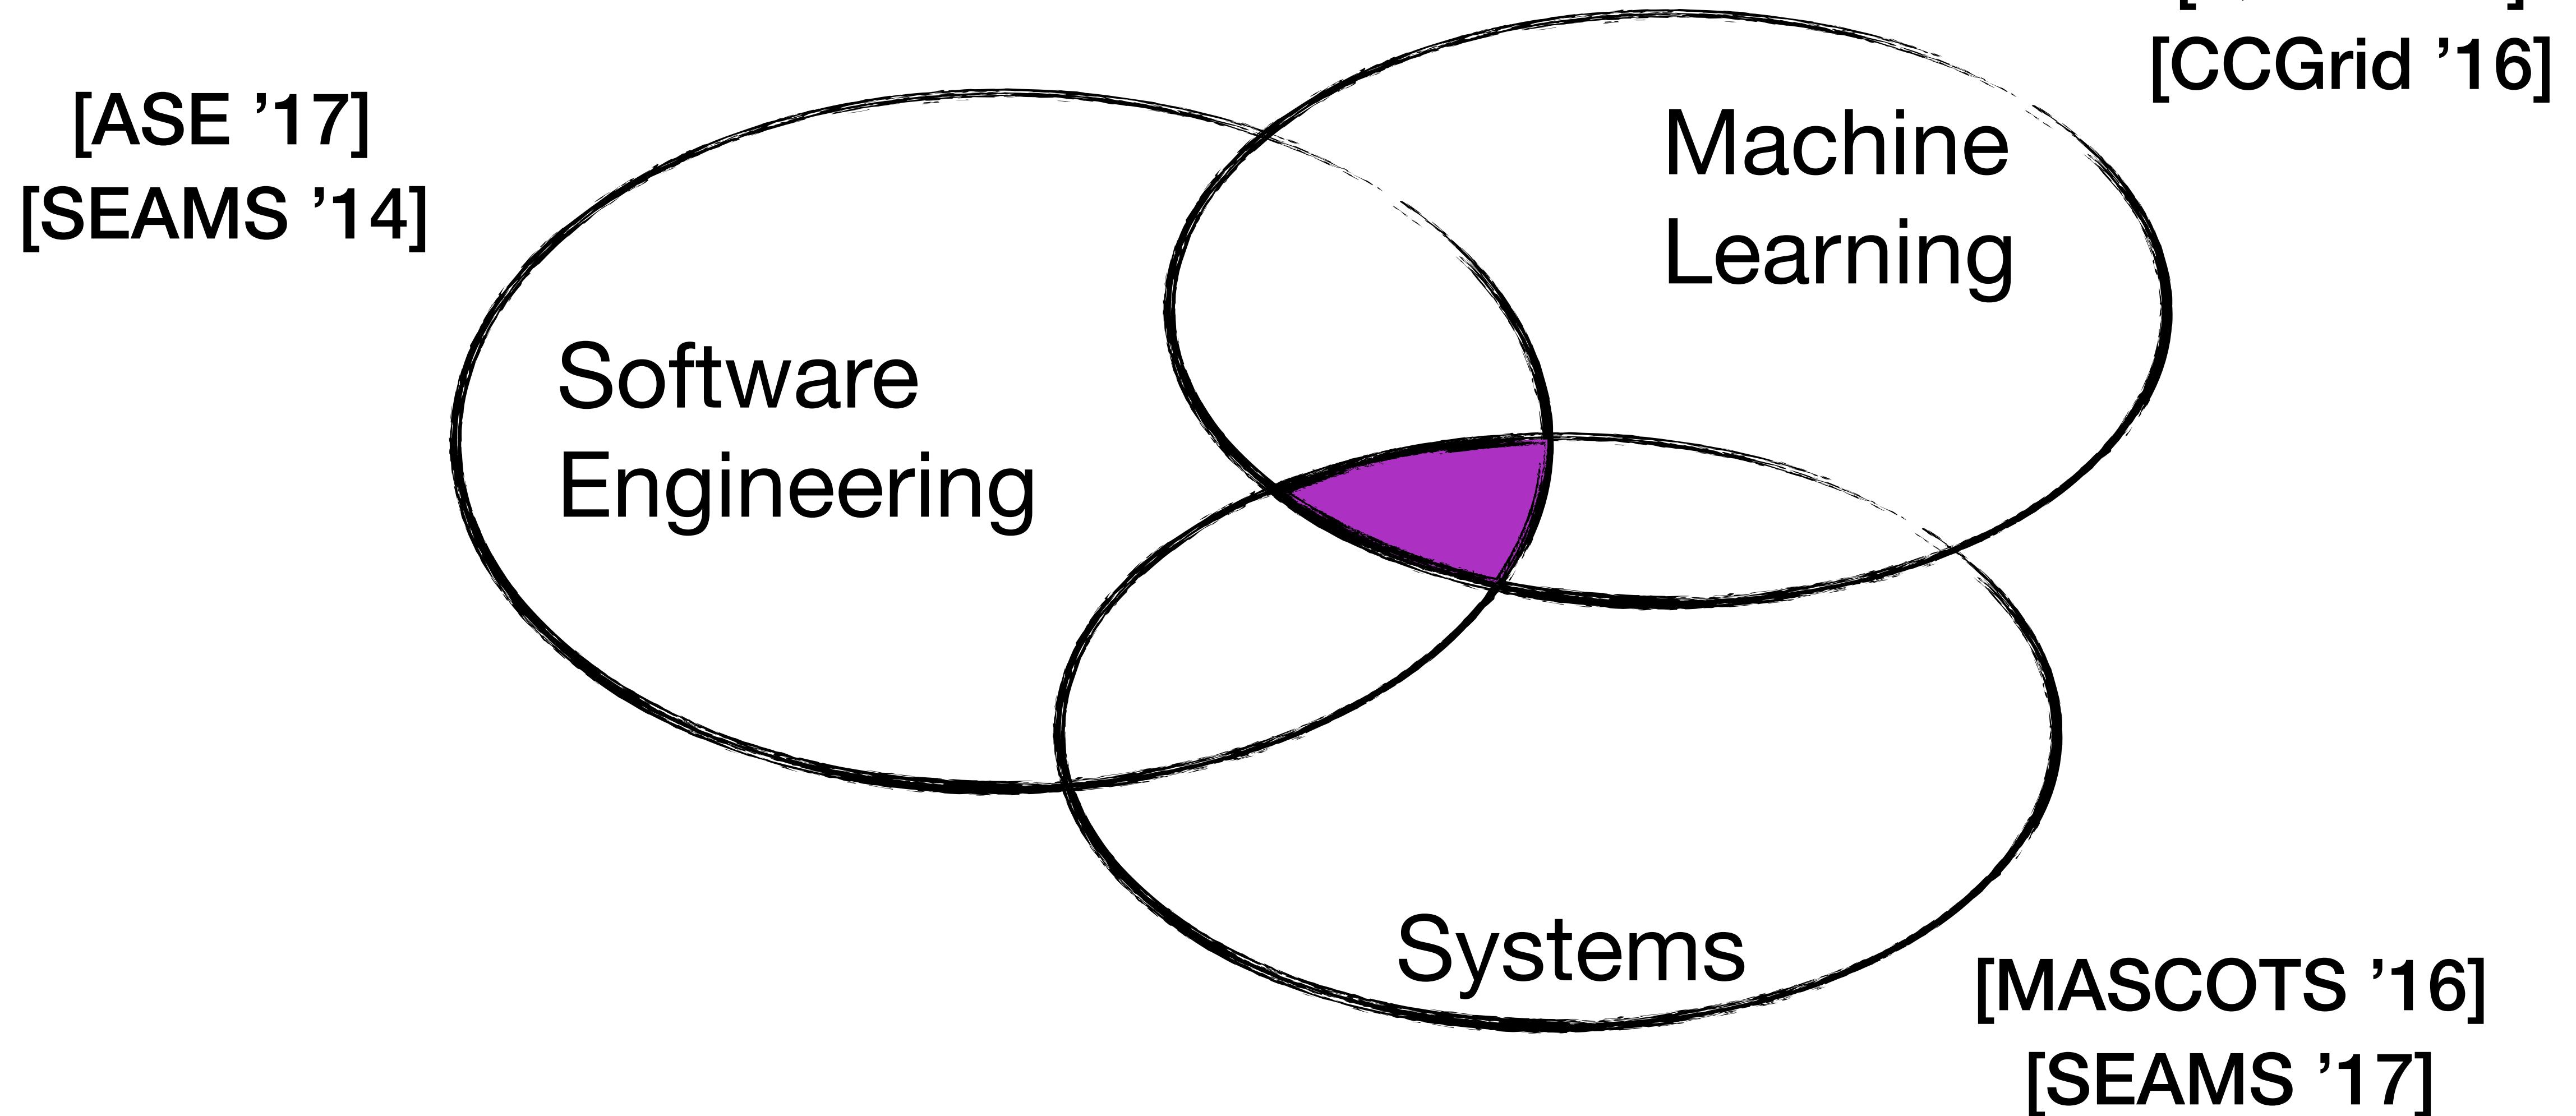
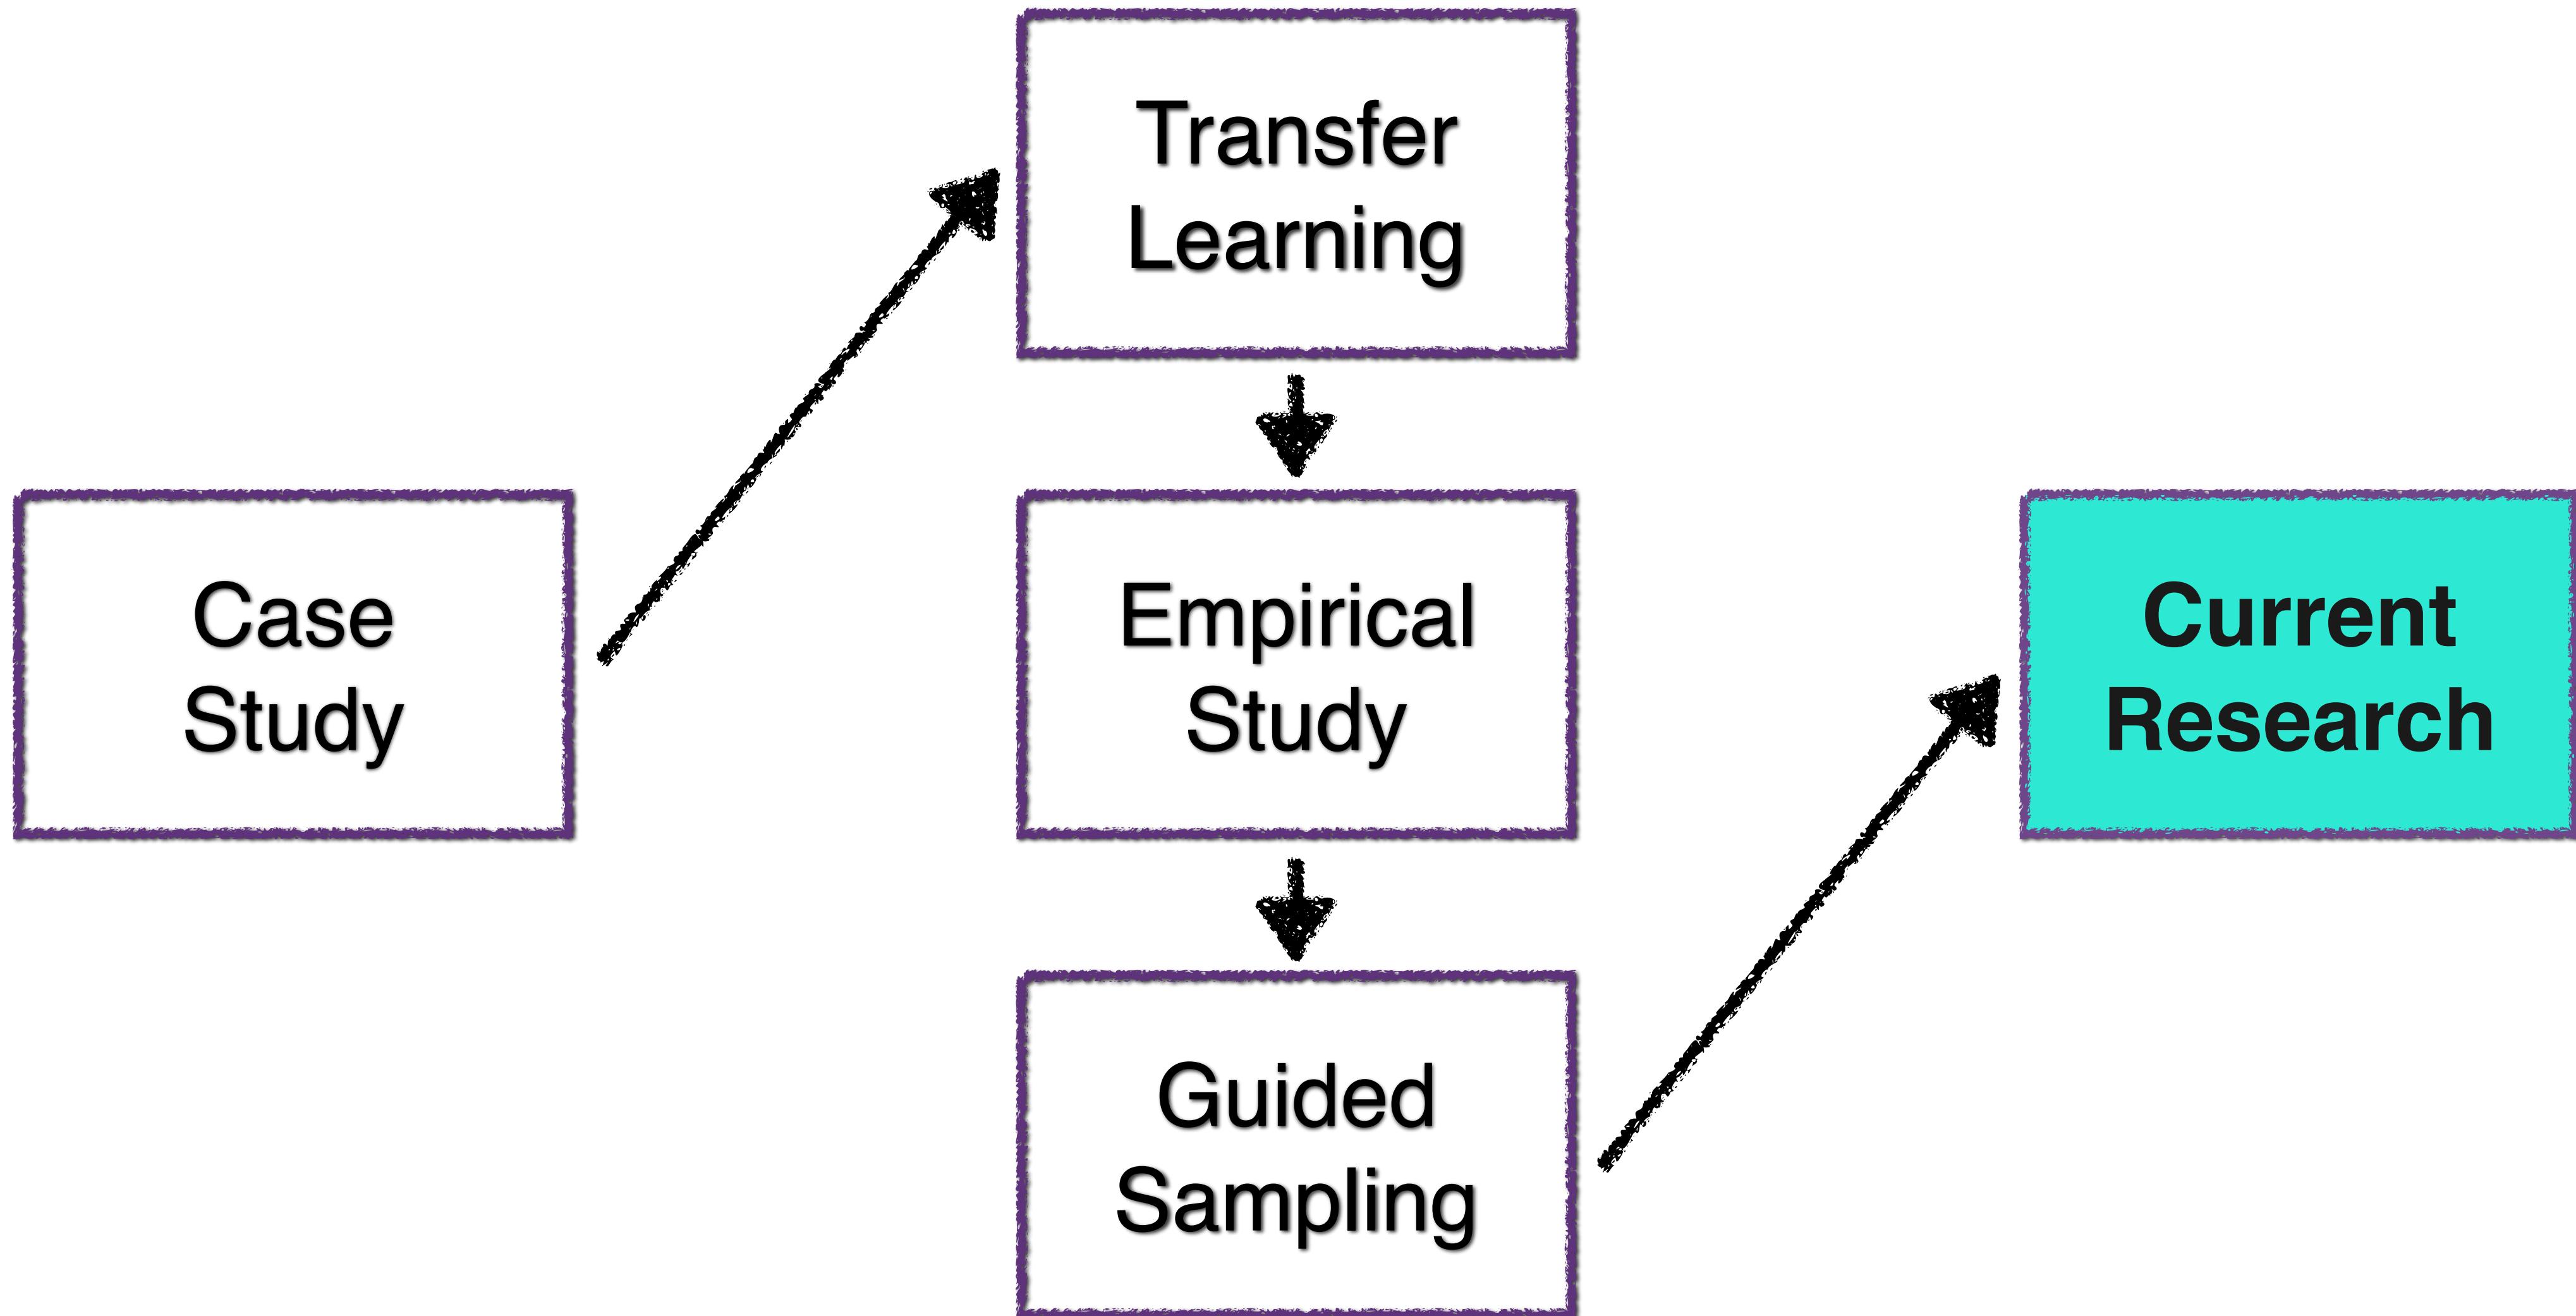


Figure 1: L2S performs guided sampling employing the knowledge extracted from a source environment.

Research interests



Outline

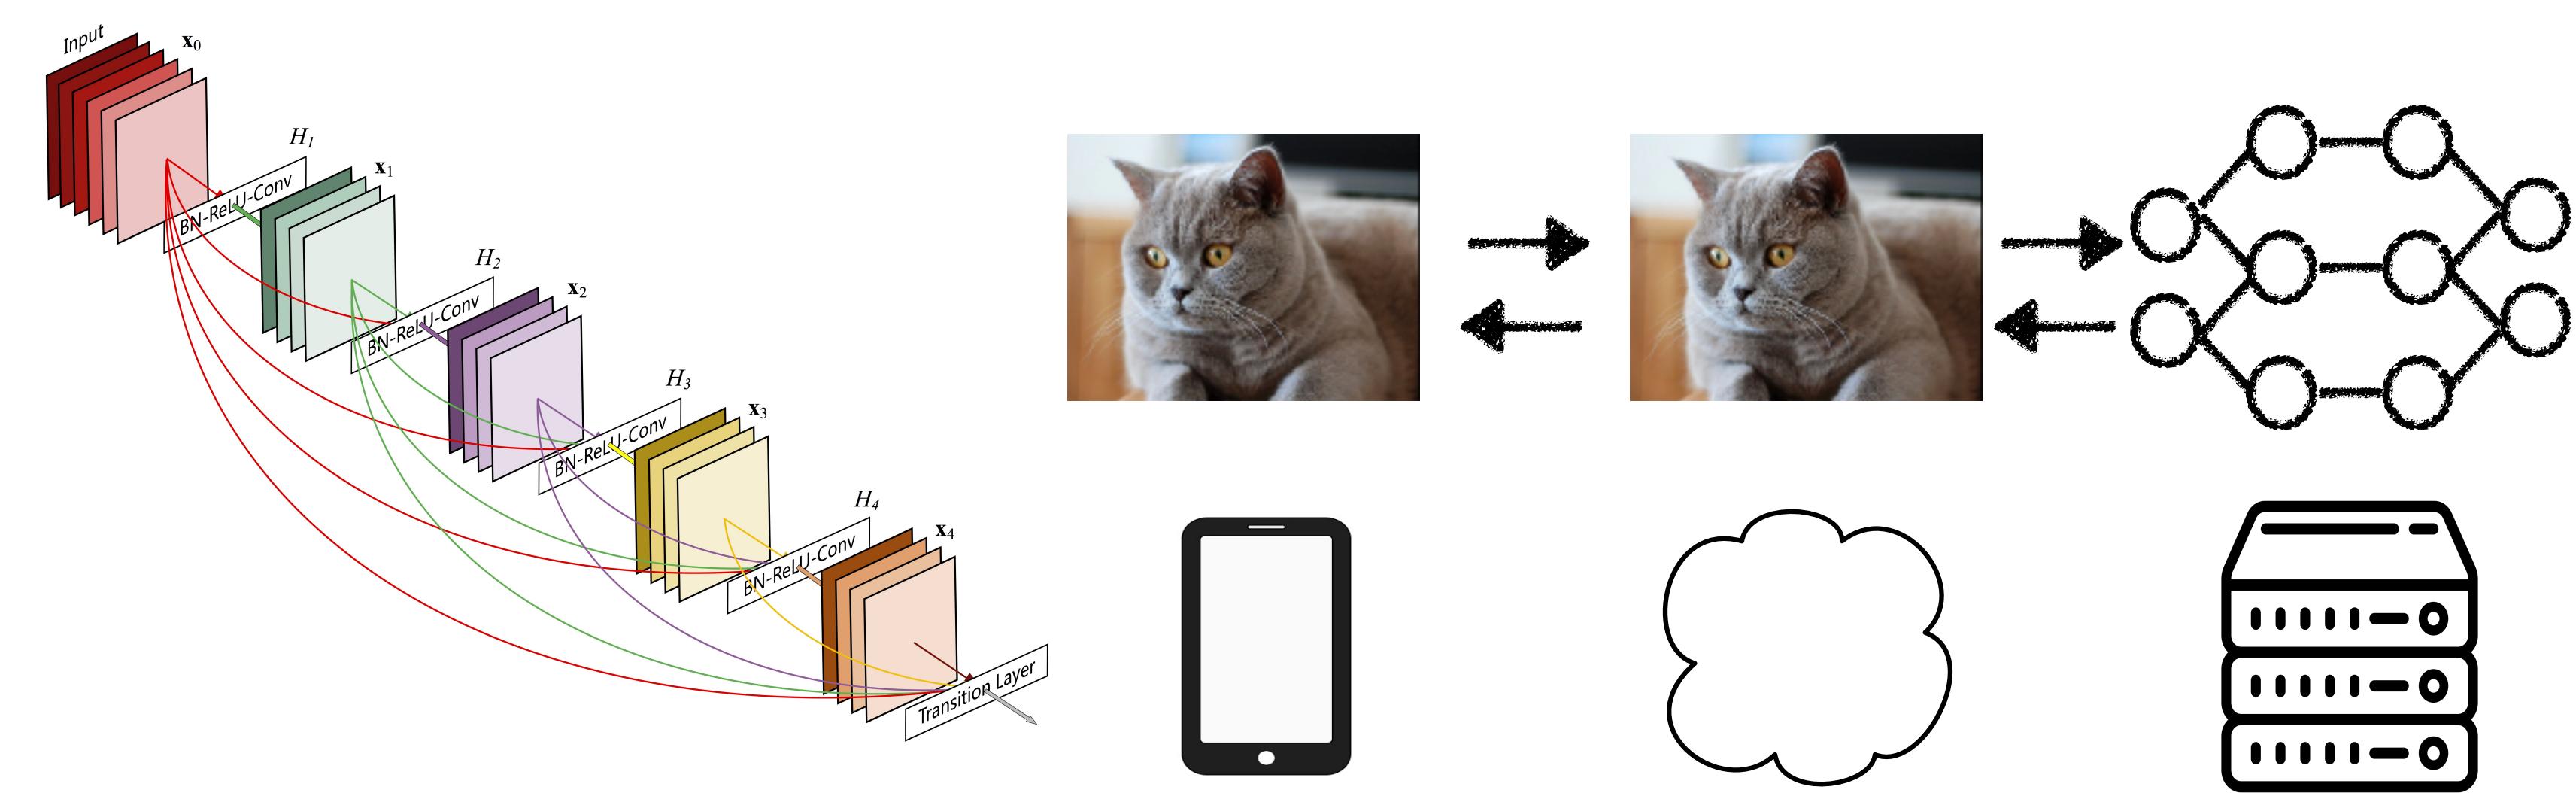
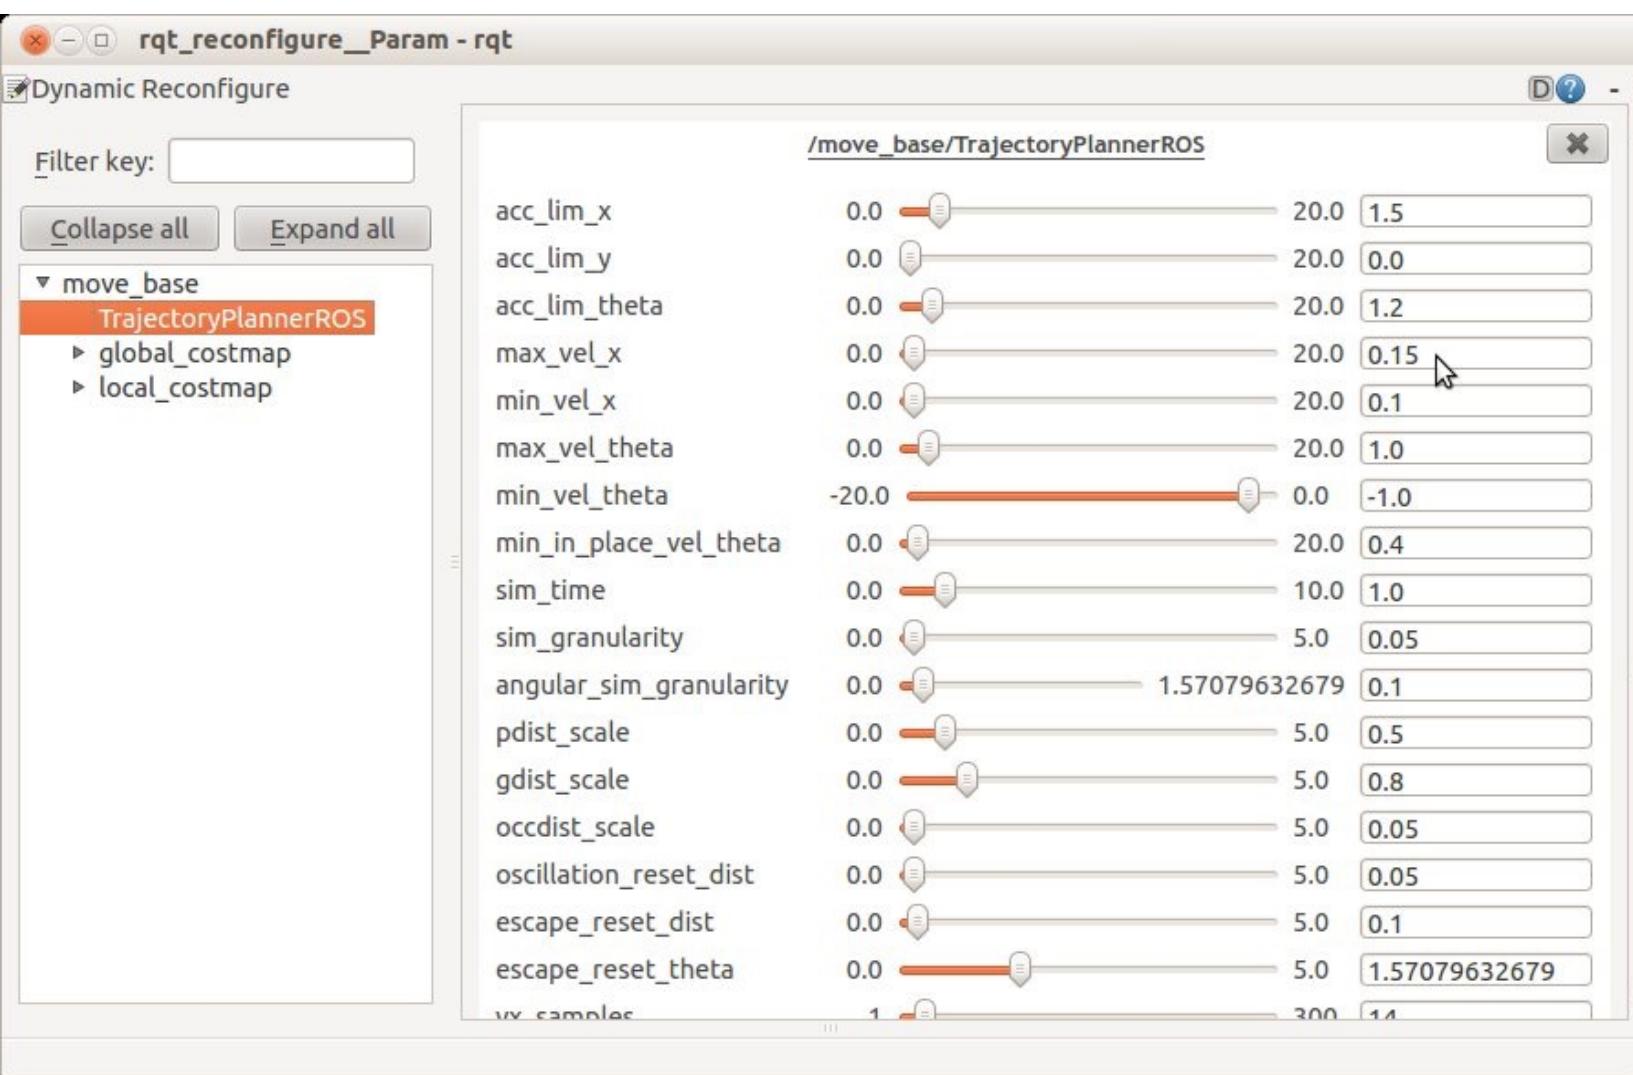




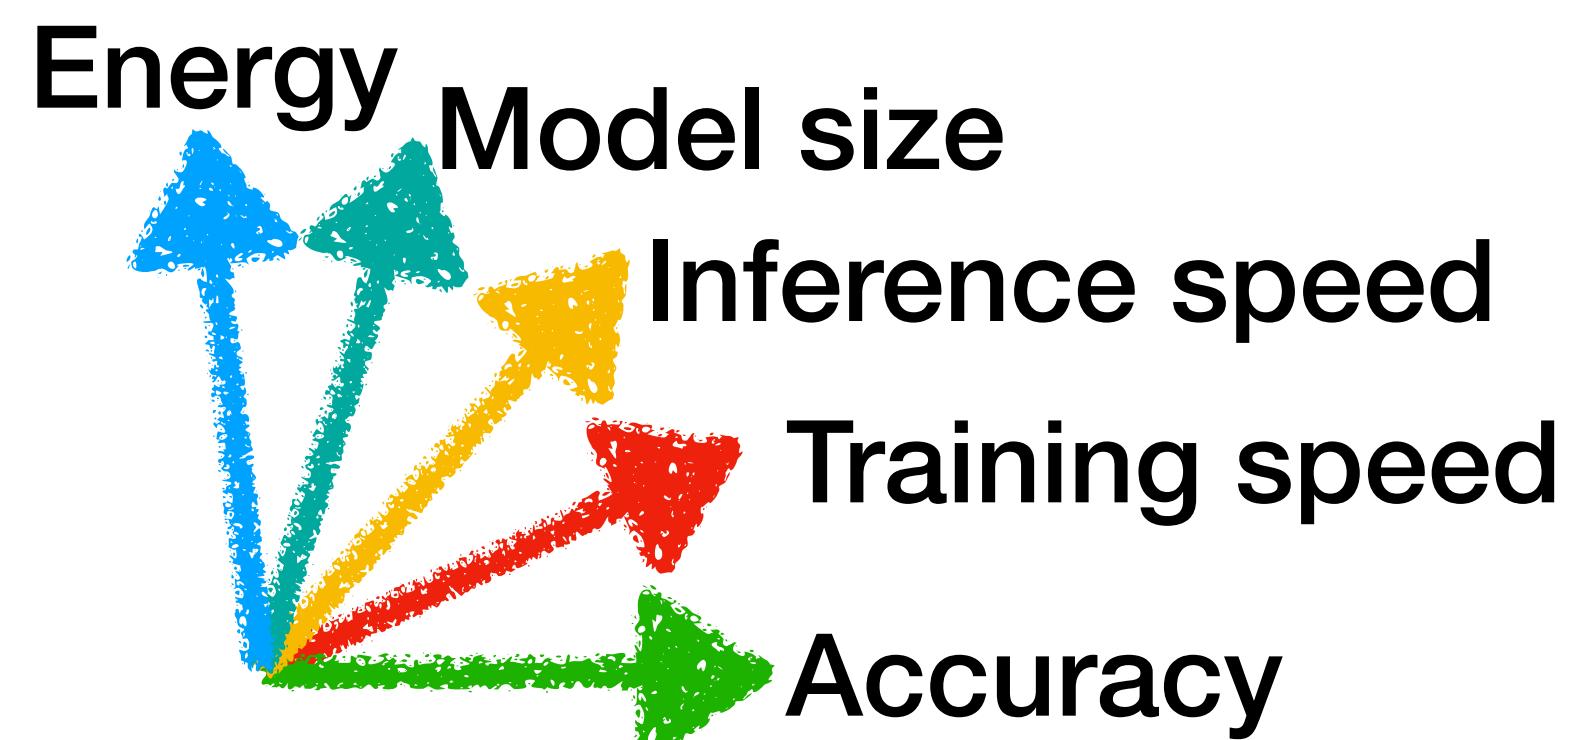
What will the software systems
of the future look like?

Software 2.0

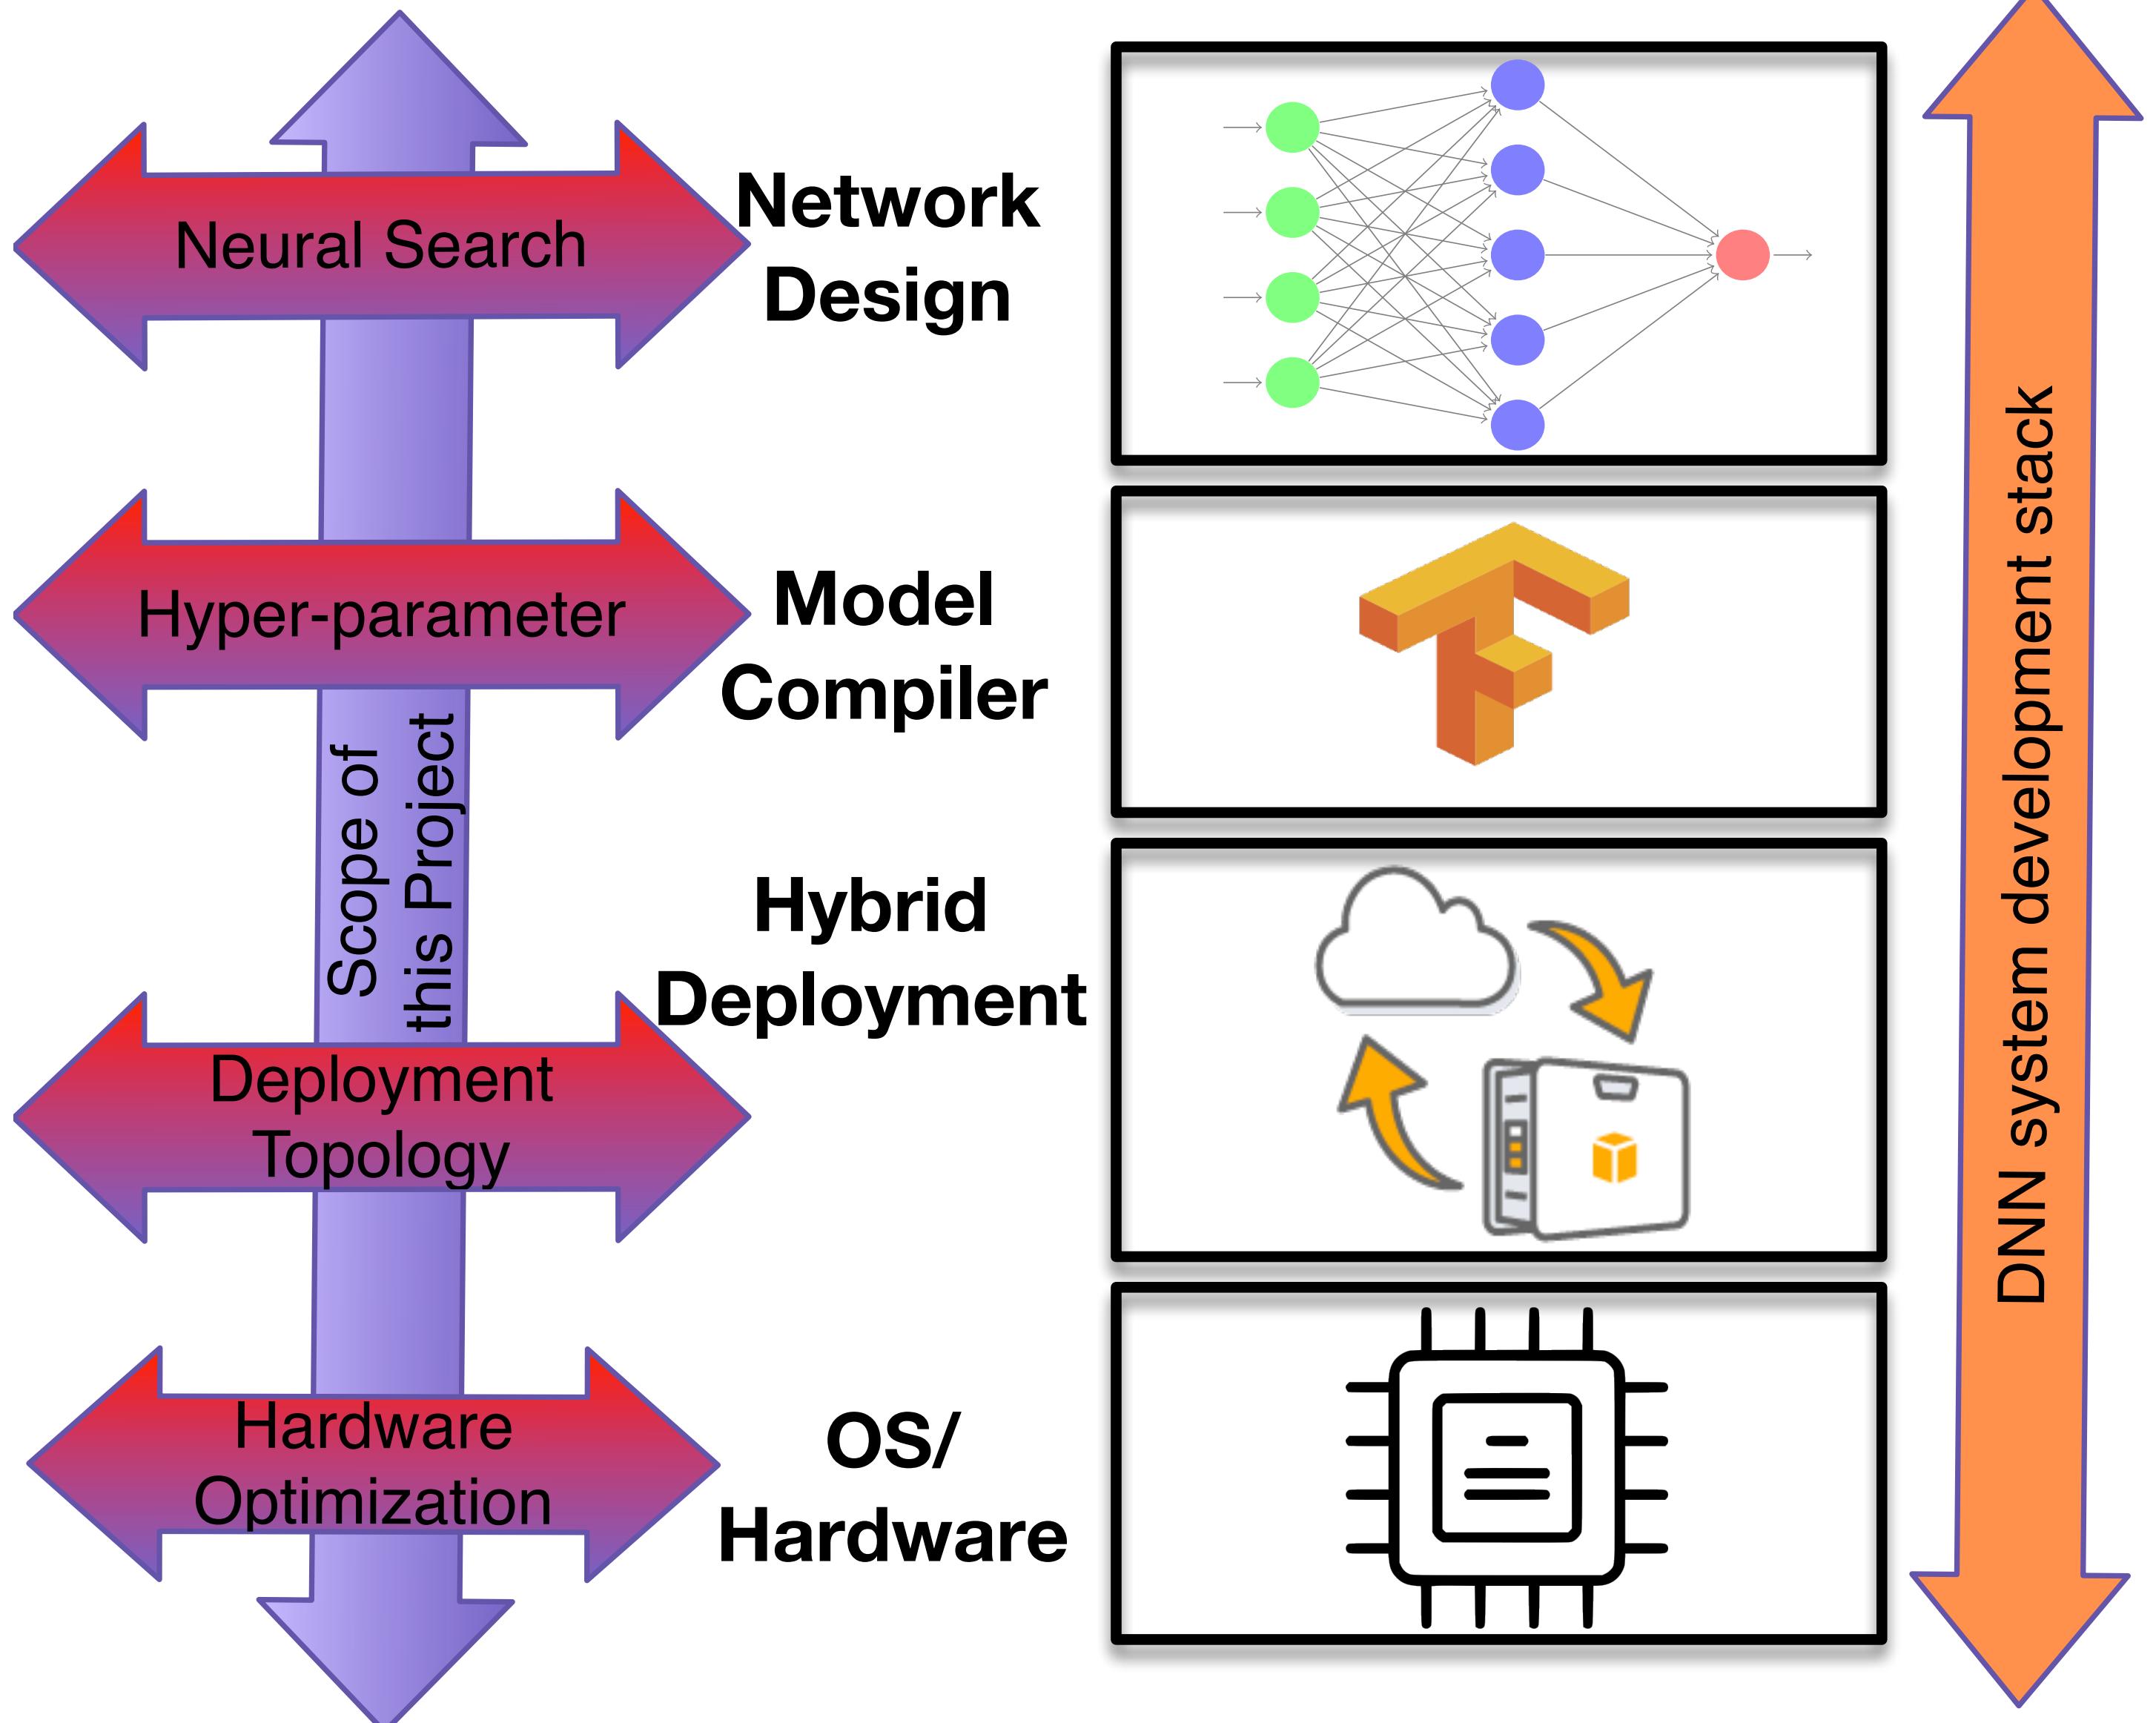
Increasingly **customized** and **configurable**



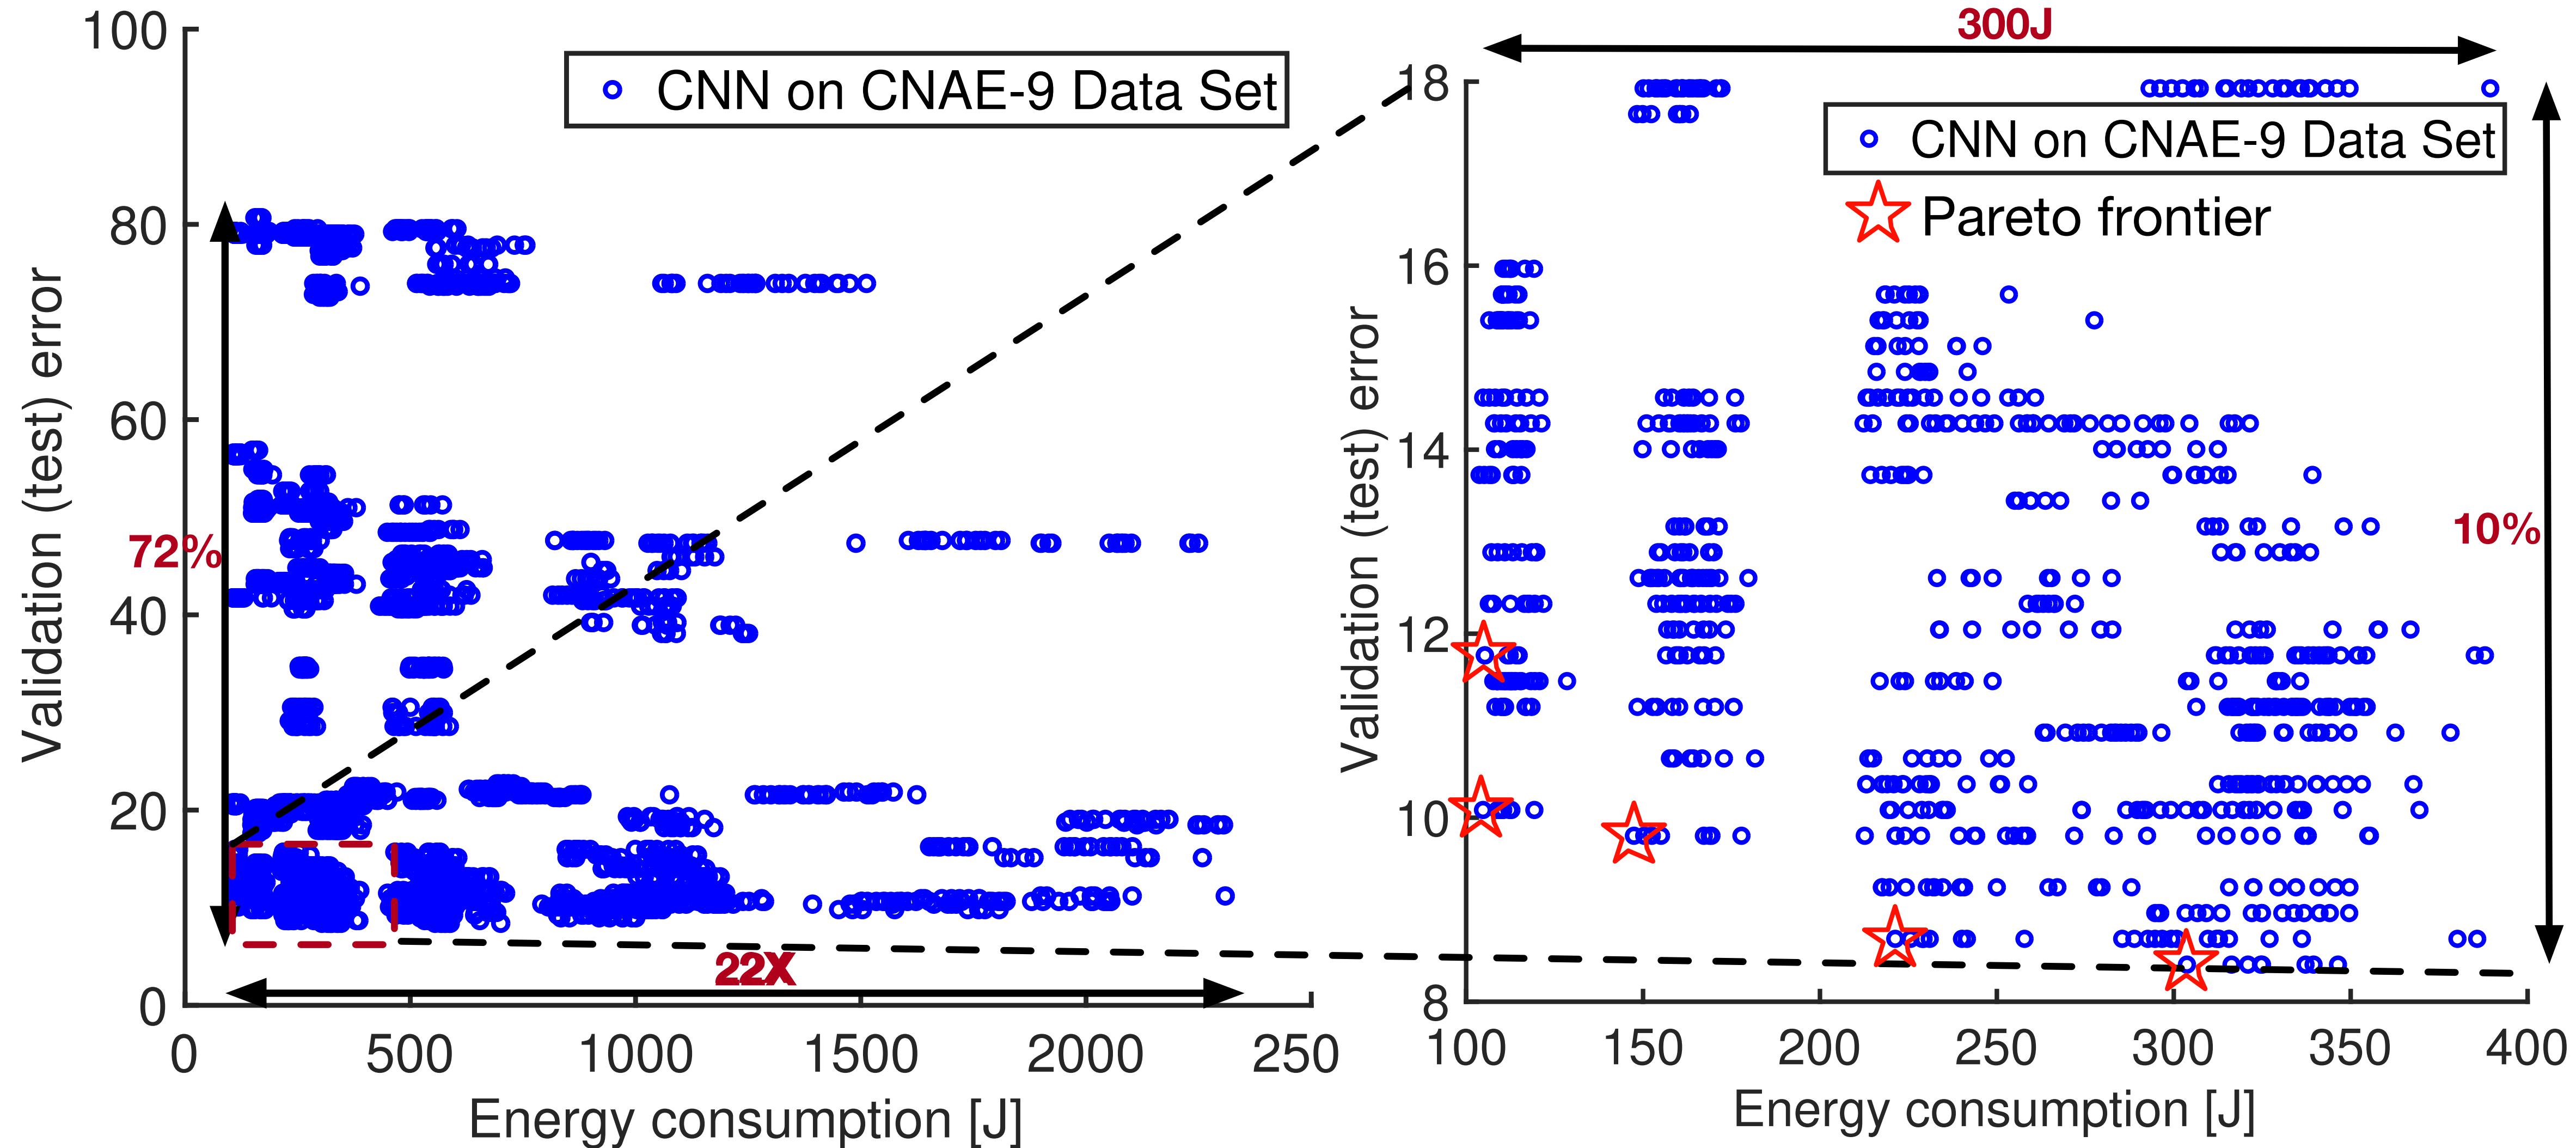
Increasingly **competing objectives**



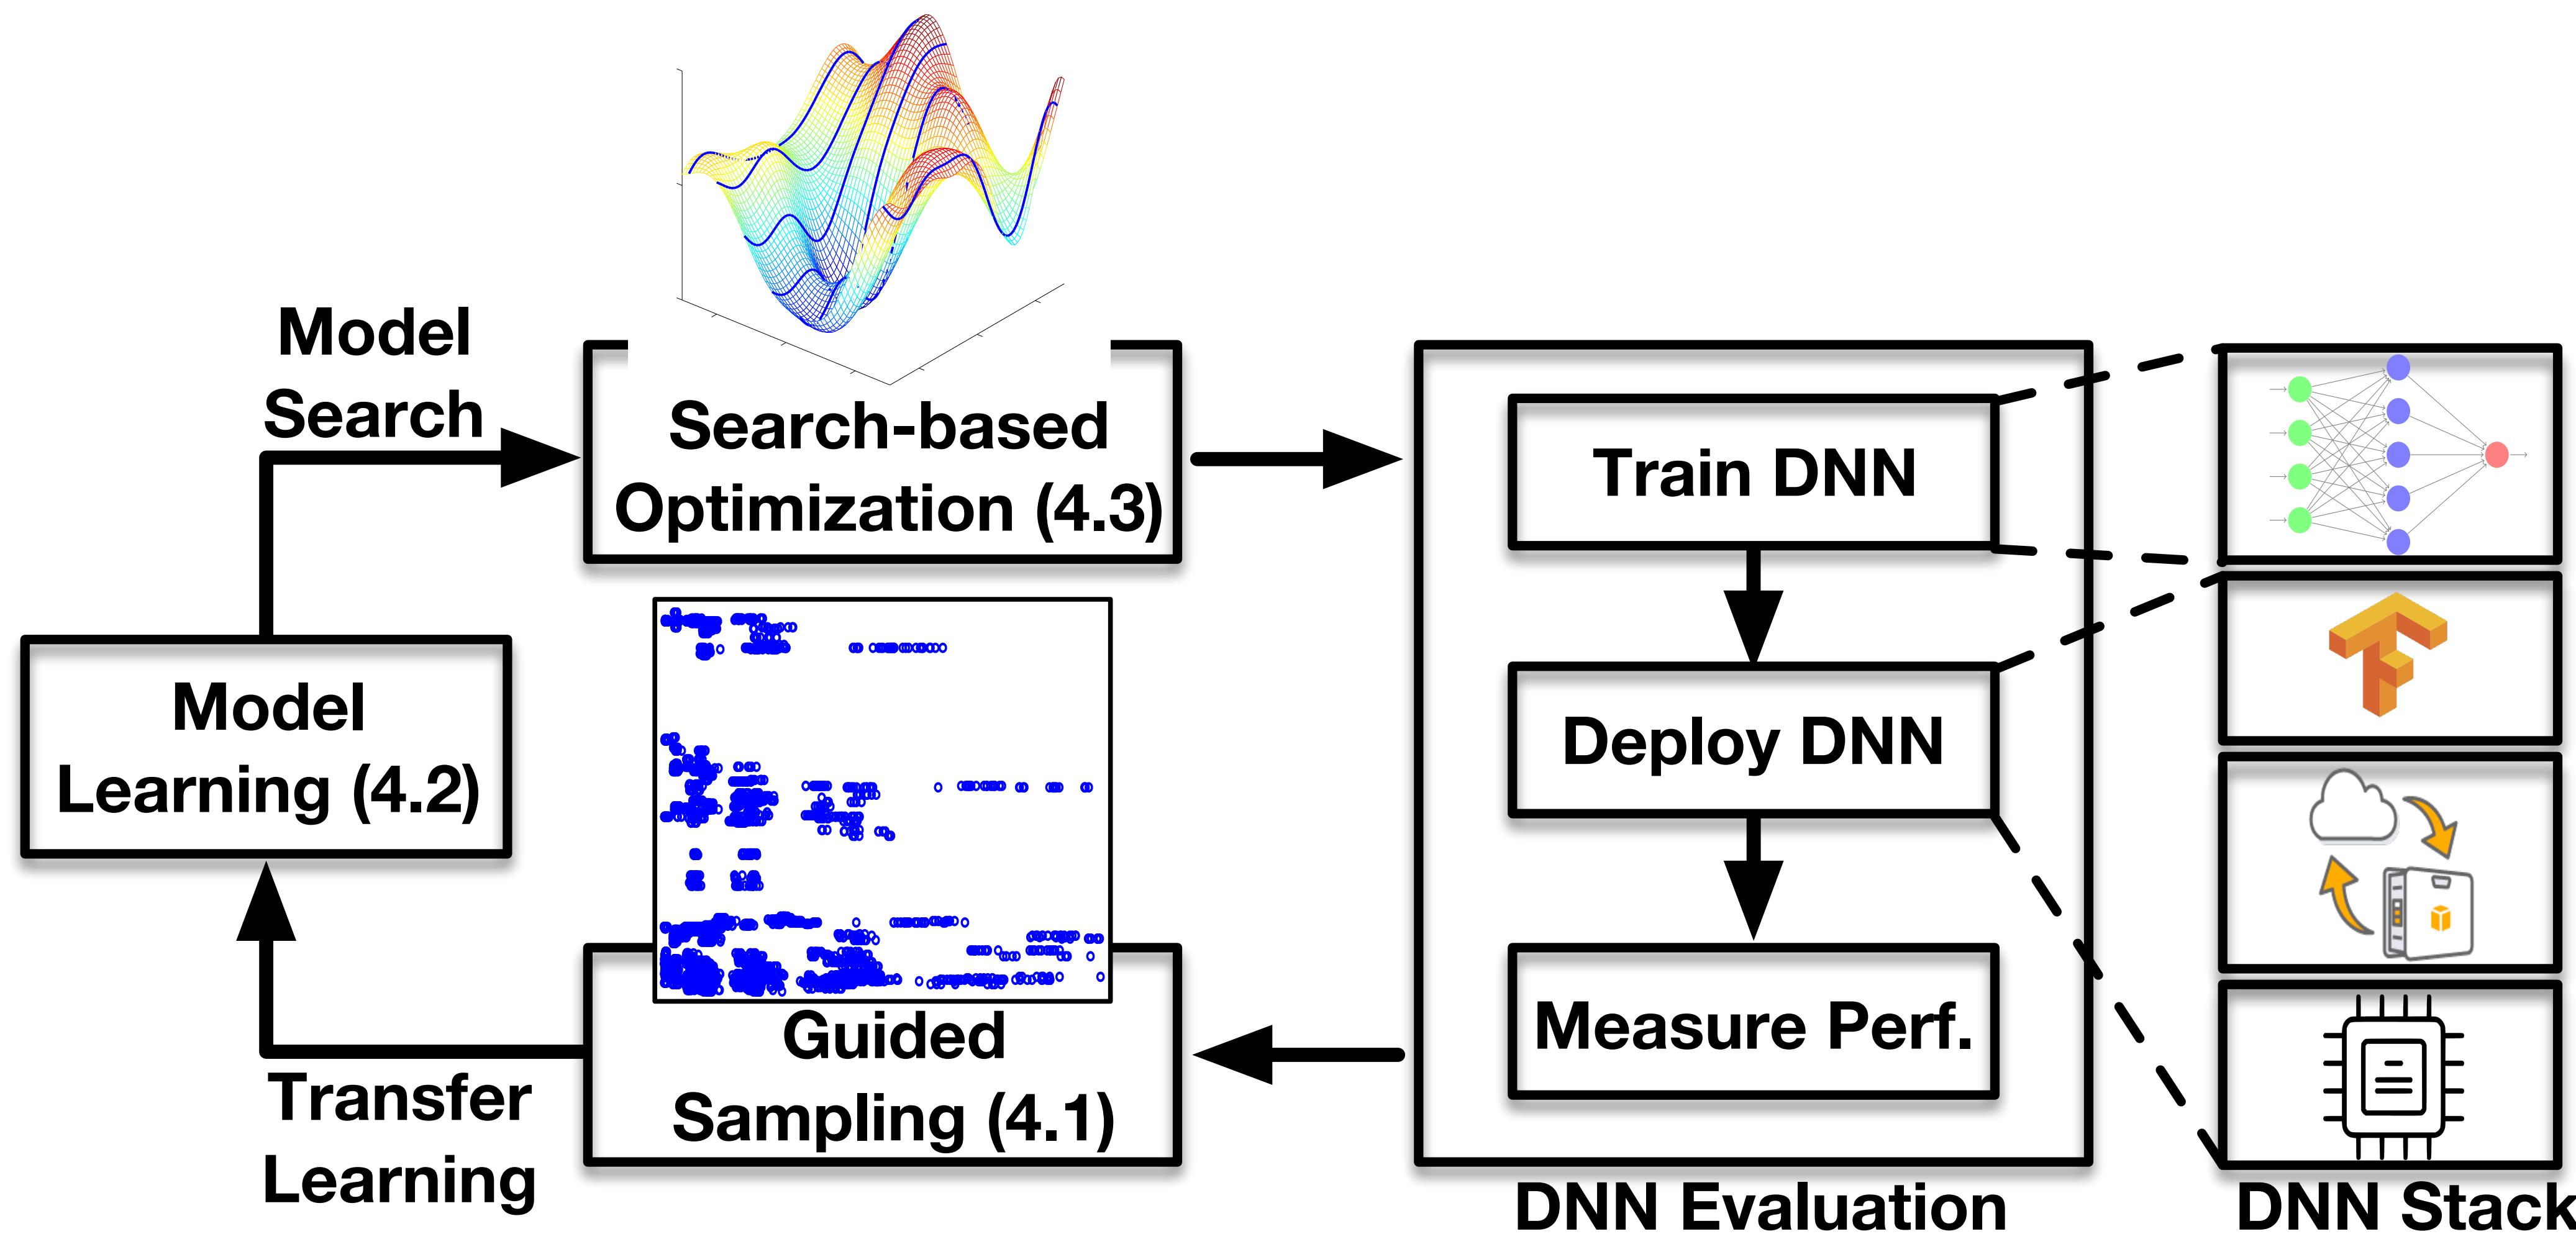
Deep neural network as a highly configurable system



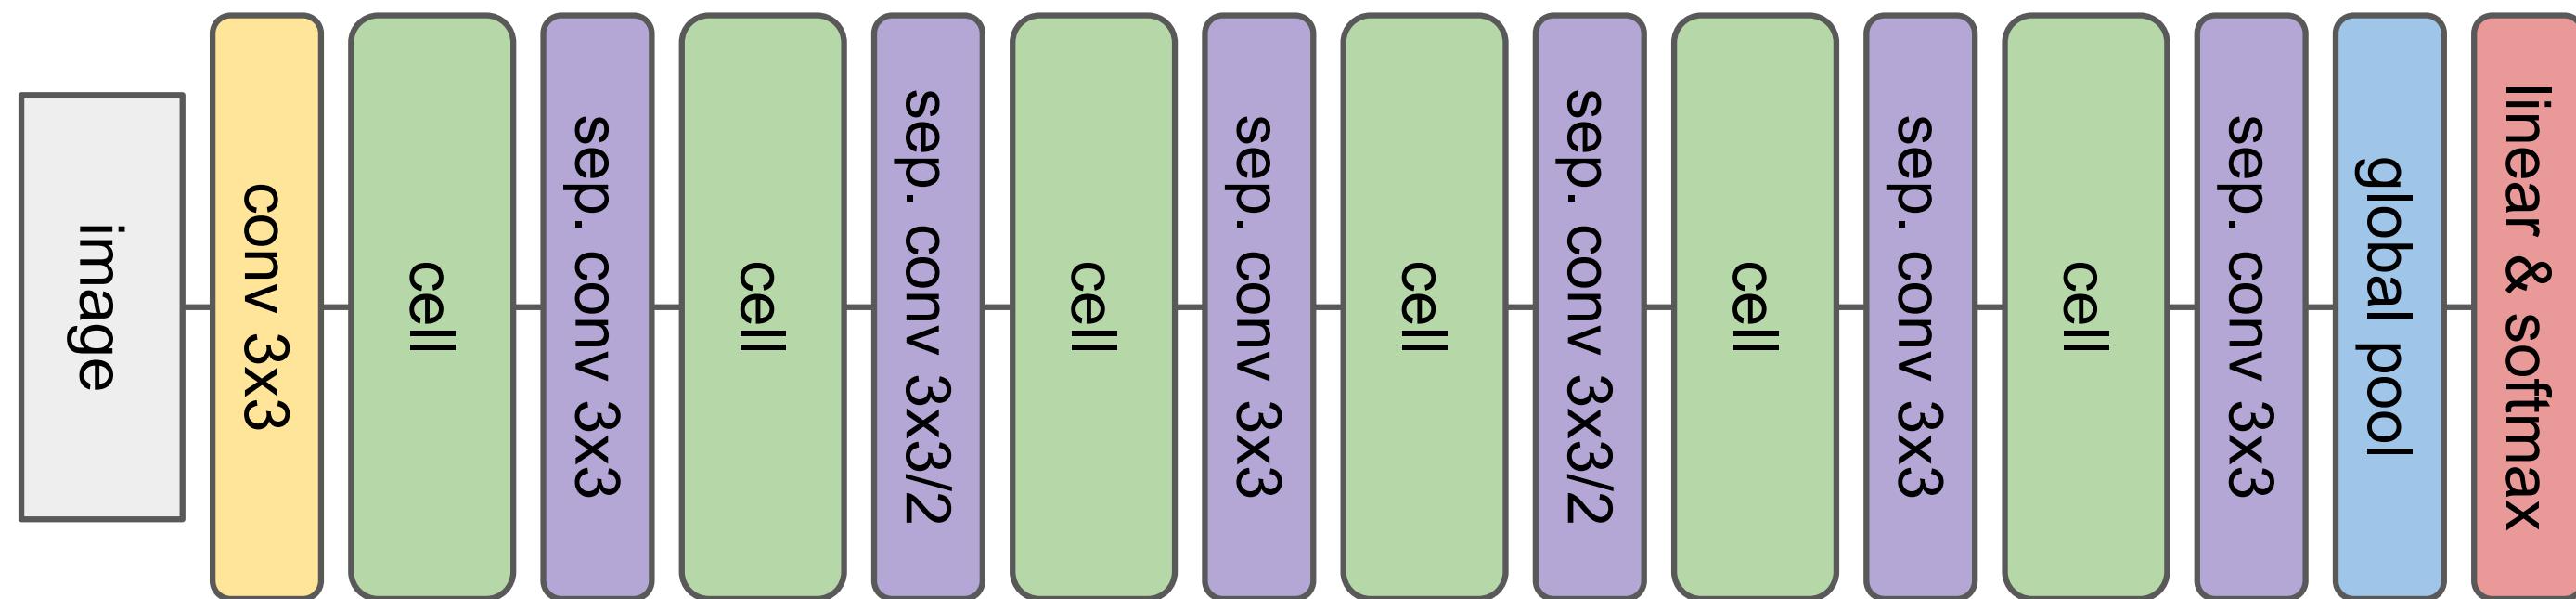
We found many configuration with the same accuracy while having drastically different energy demand



Optimizing Energy Consumption of Deep Neural Networks



Exploring the design space of deep networks



Optimal Architecture
(Yesterday)

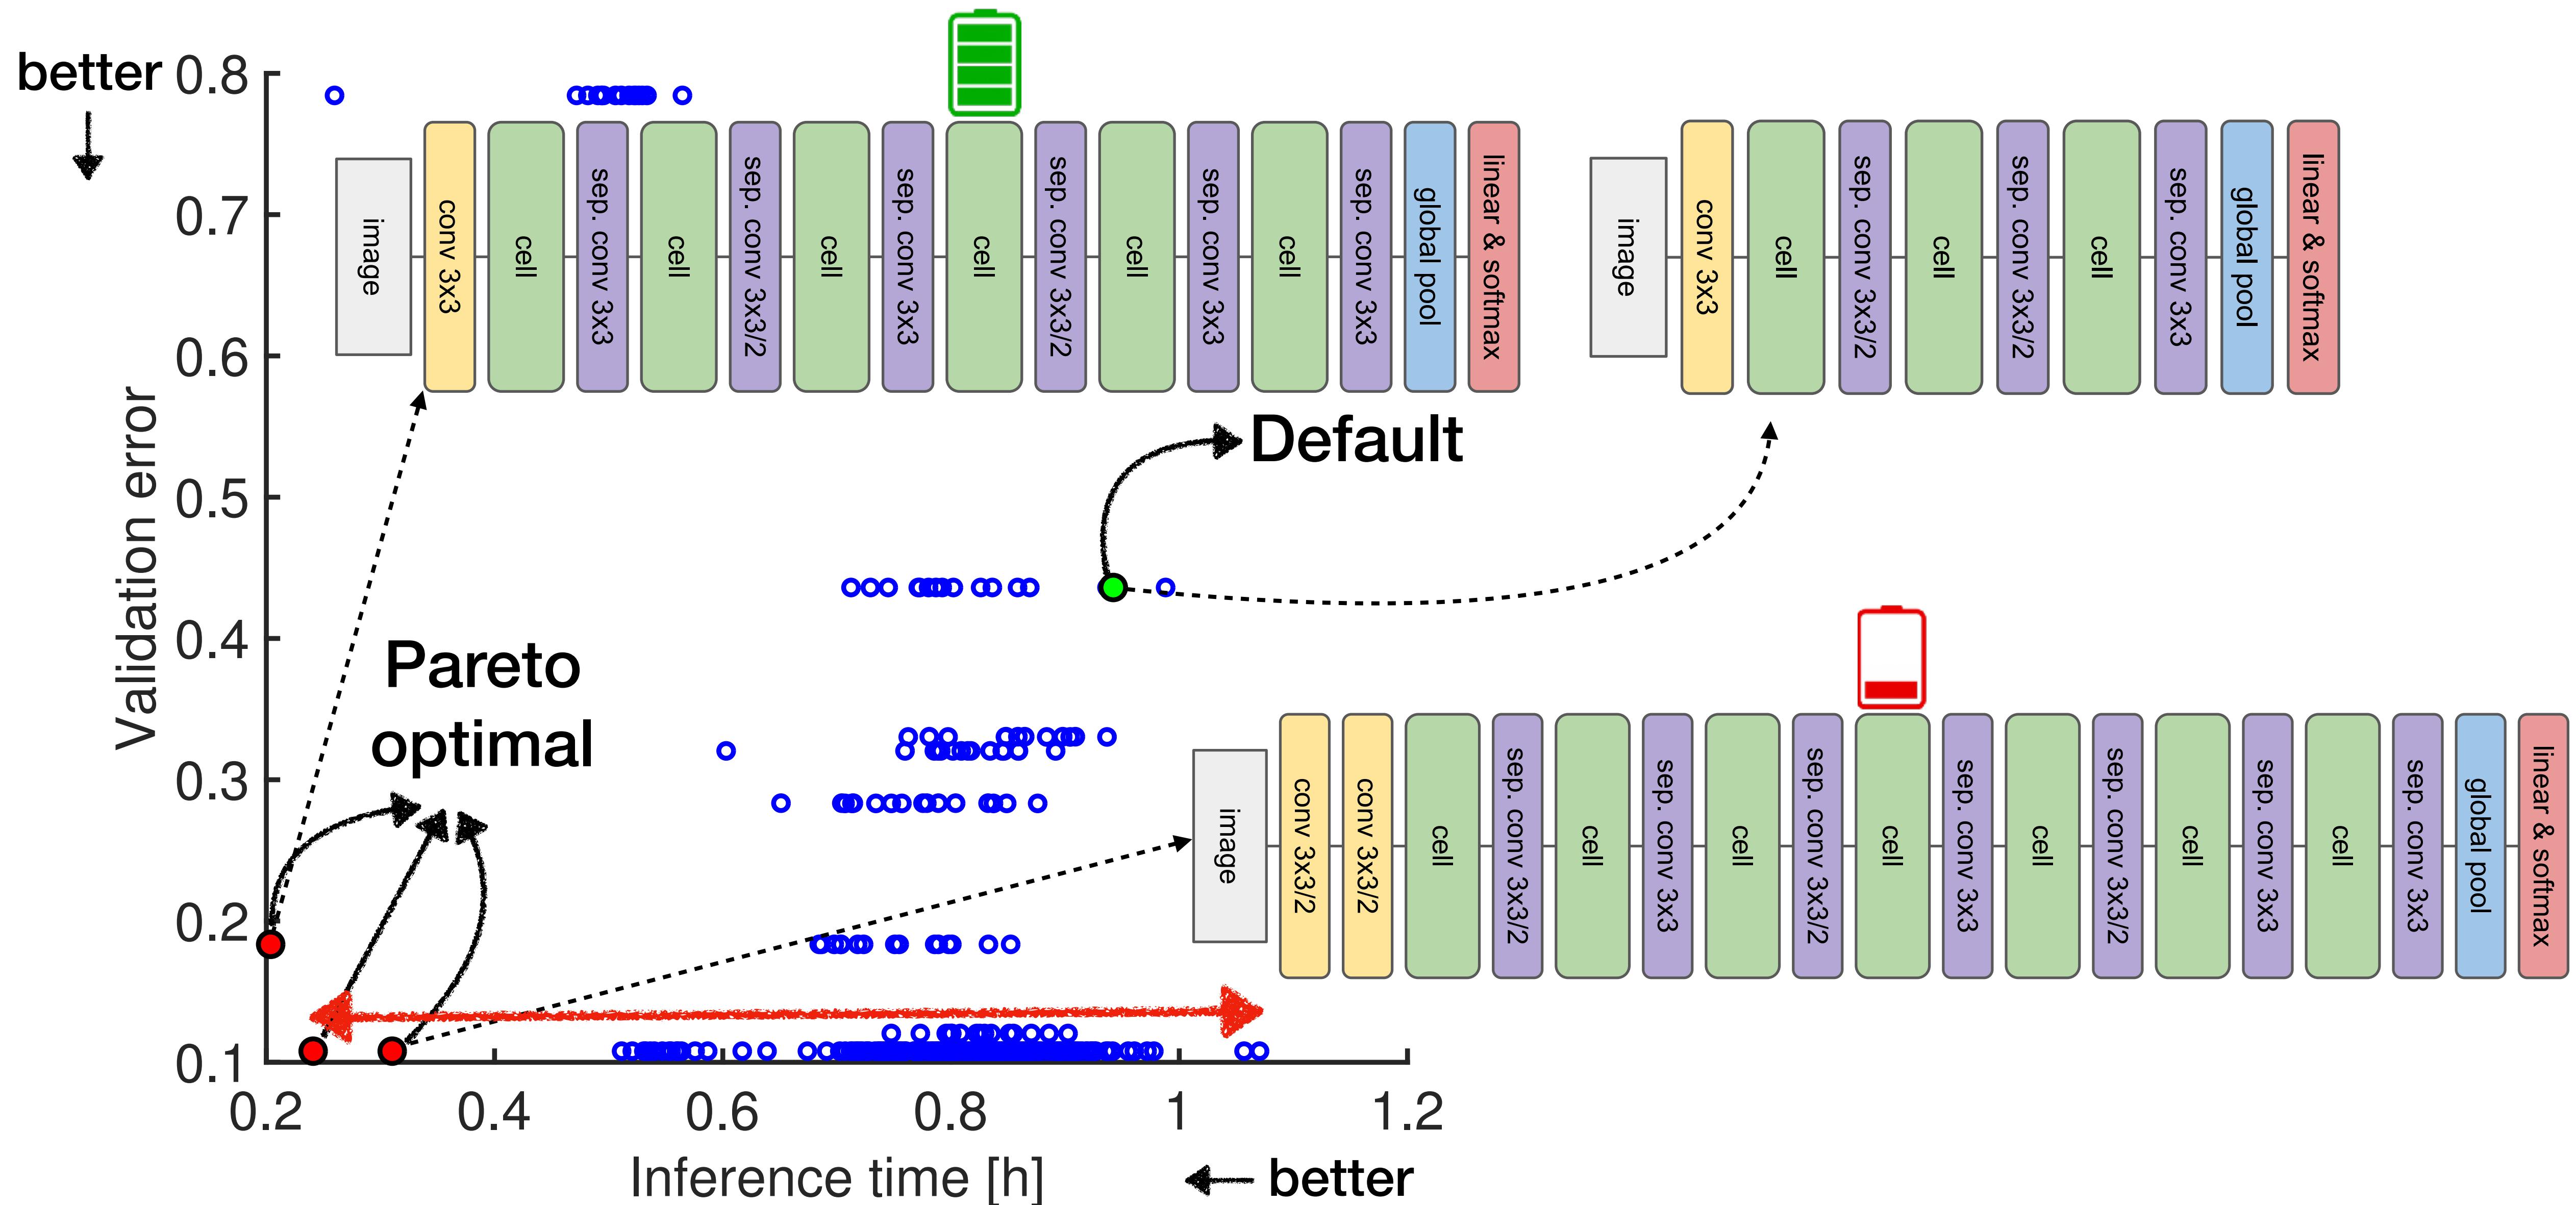


New Fraud Pattern

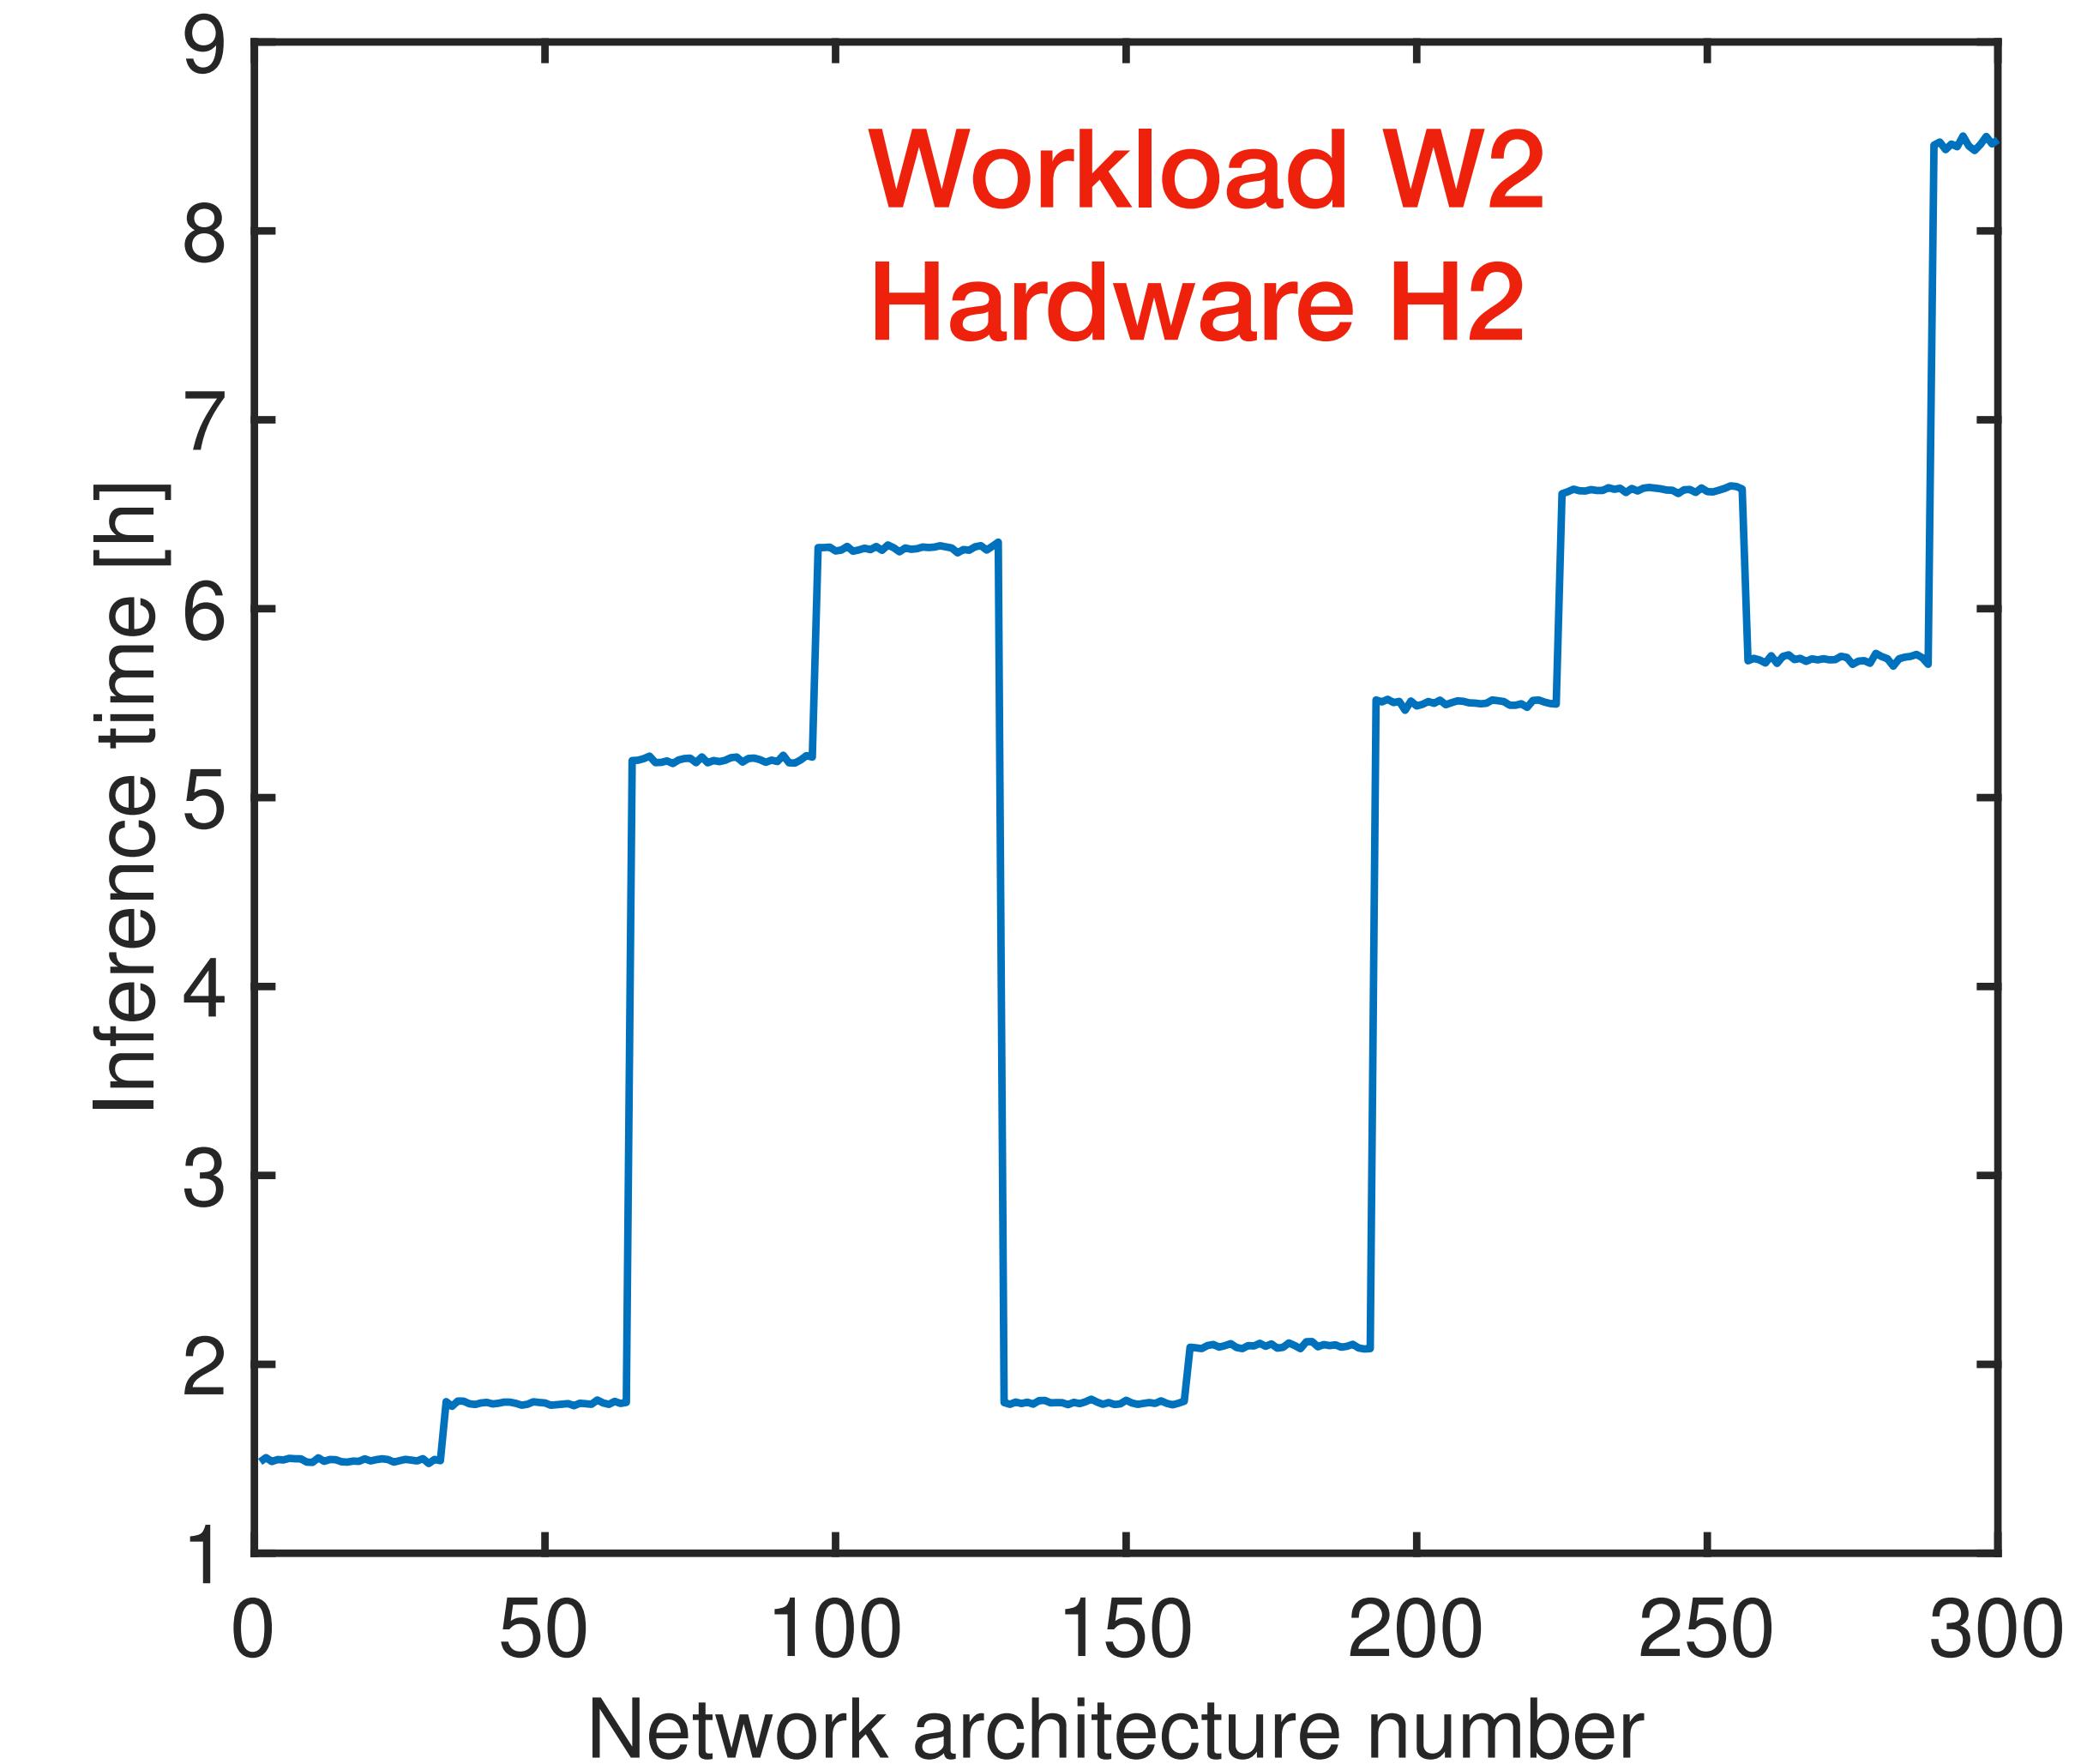
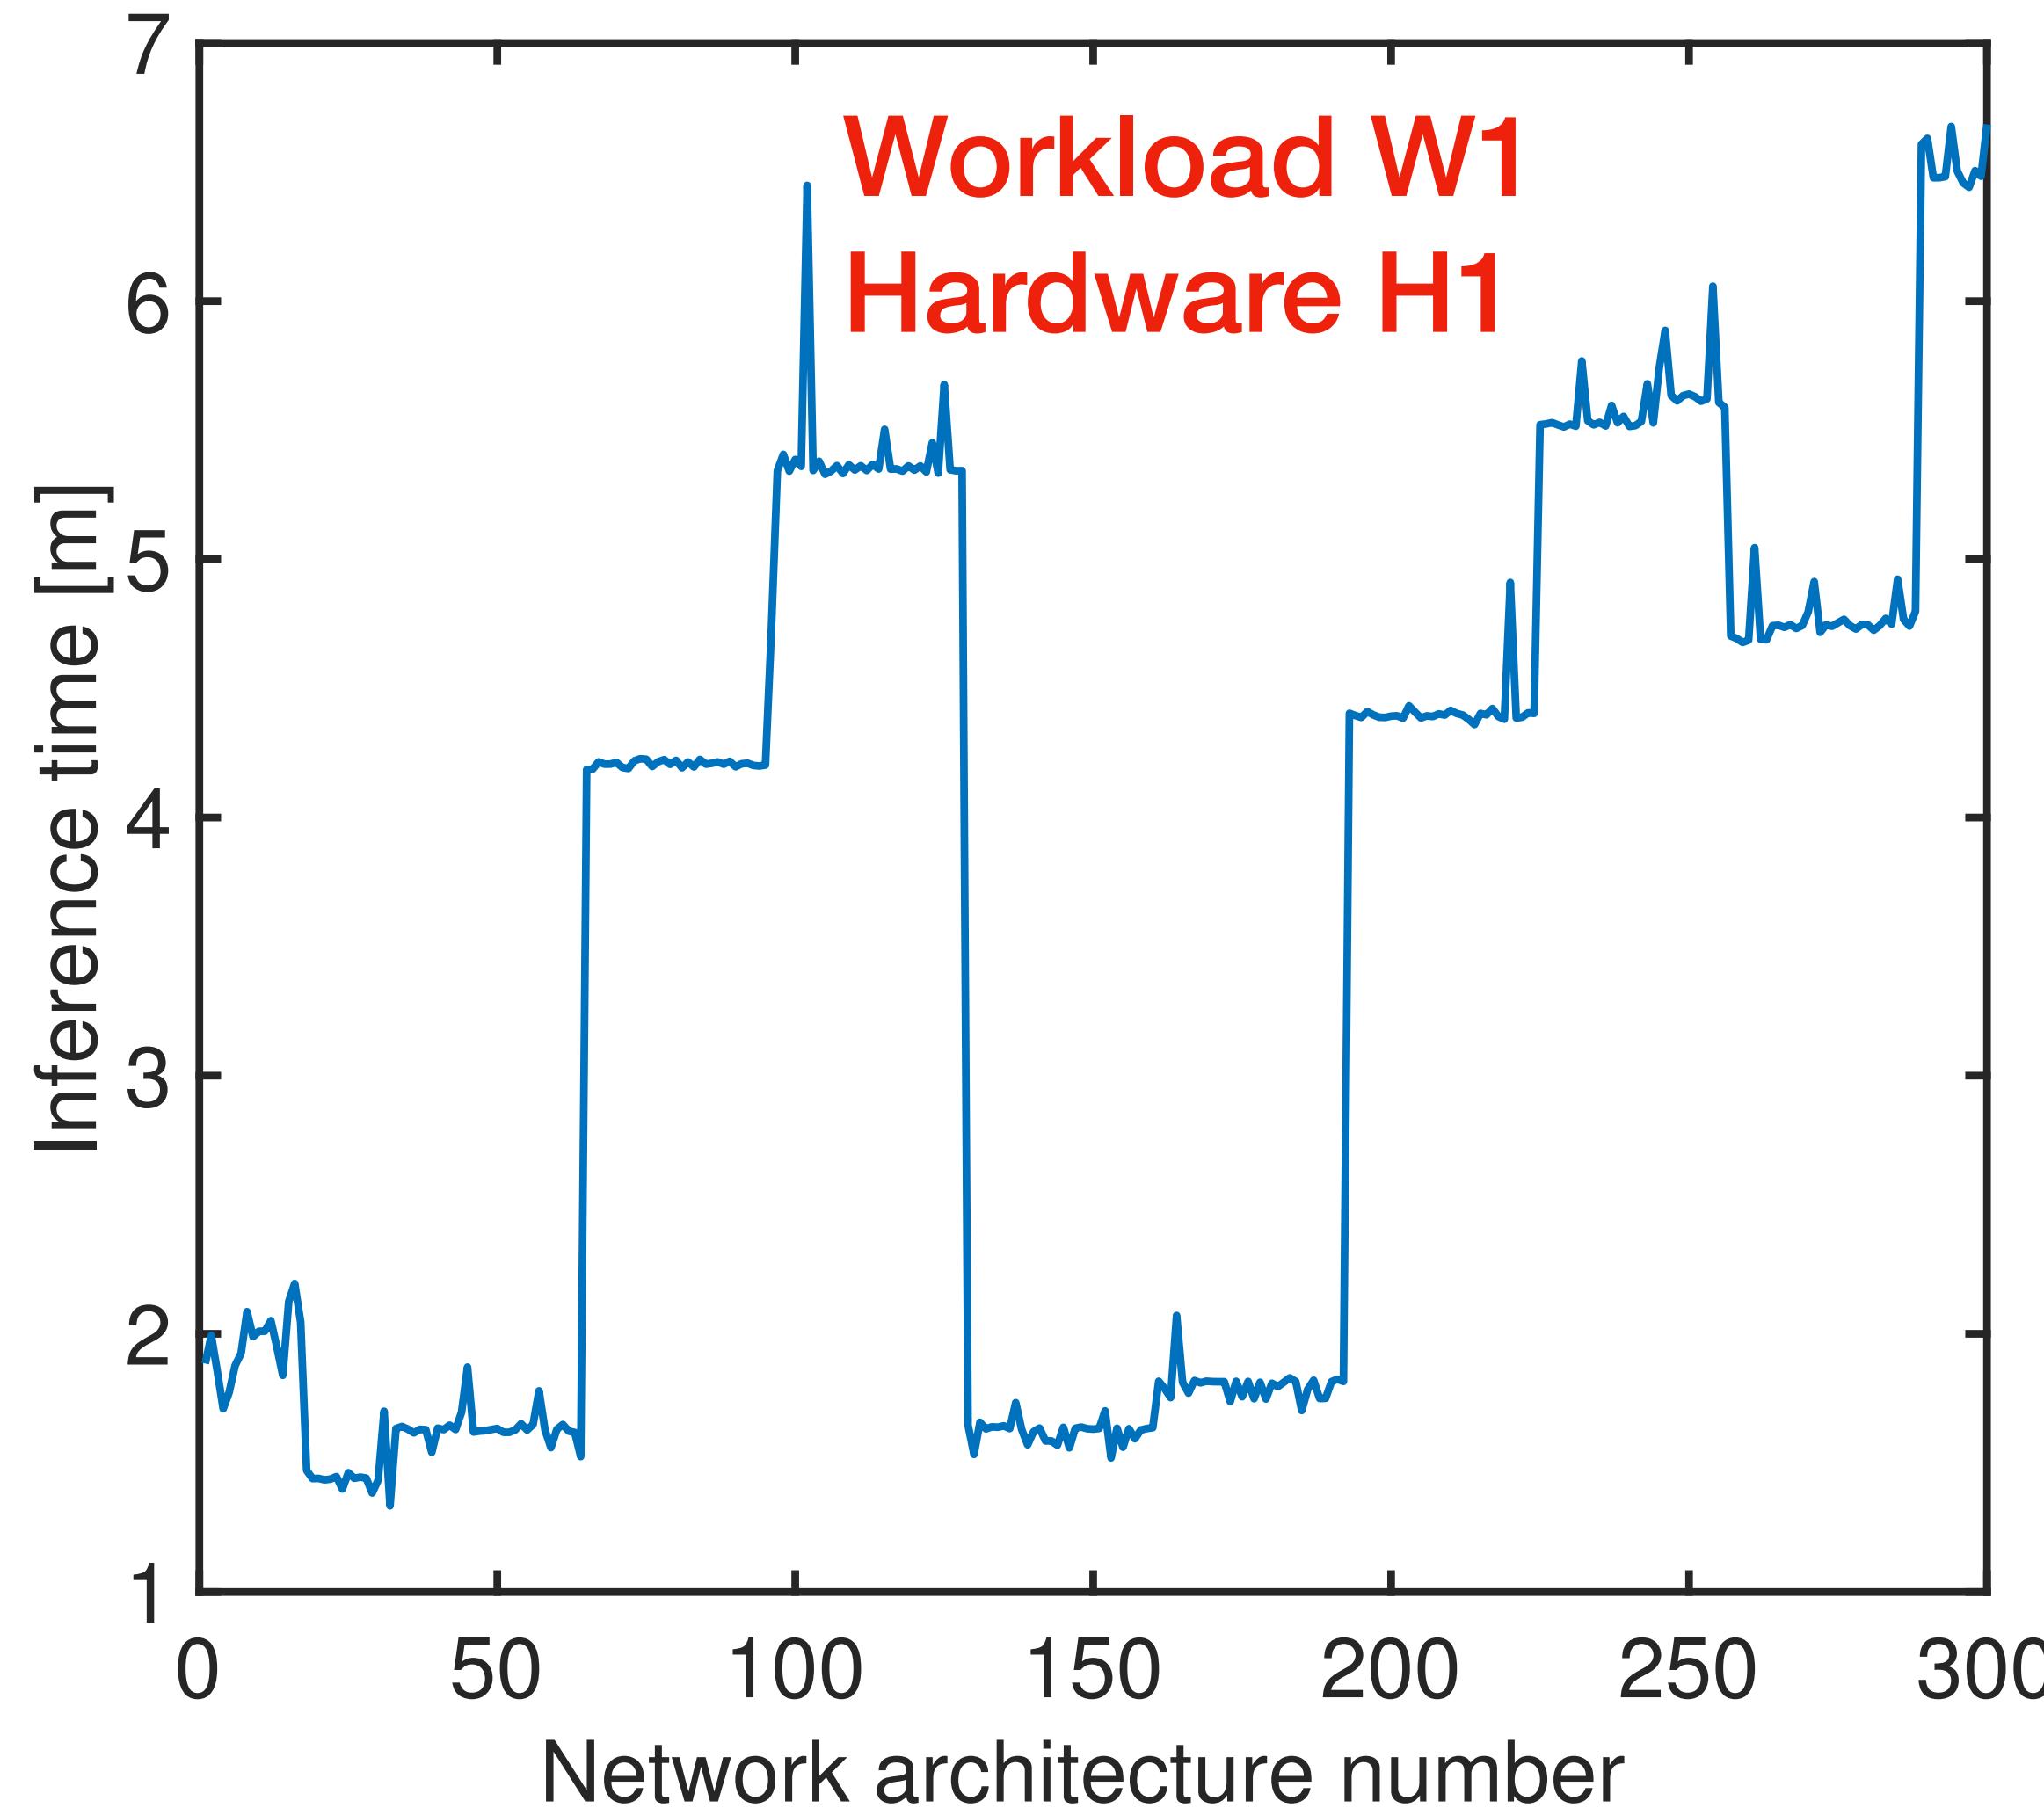


Optimal Architecture
(Today)

Exploring the design space of deep networks



Insight: Learn a model on a cheaper workload to explore the expensive workload faster



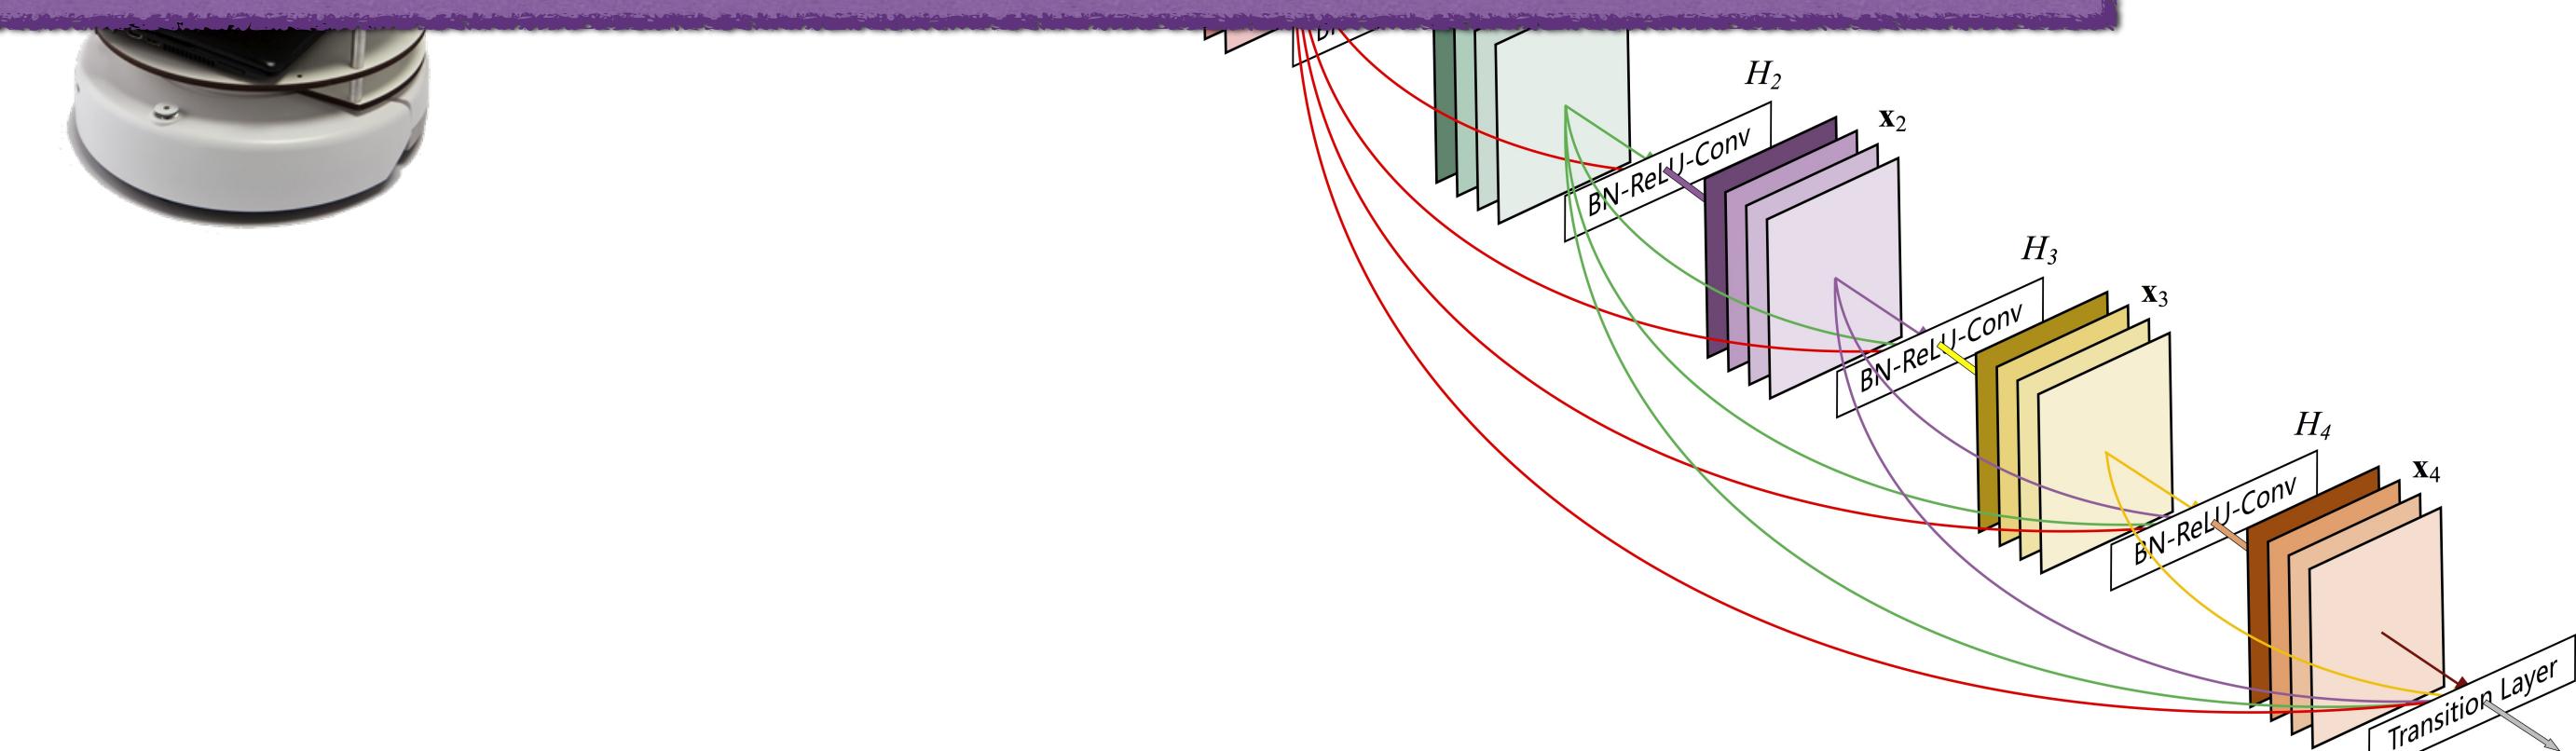
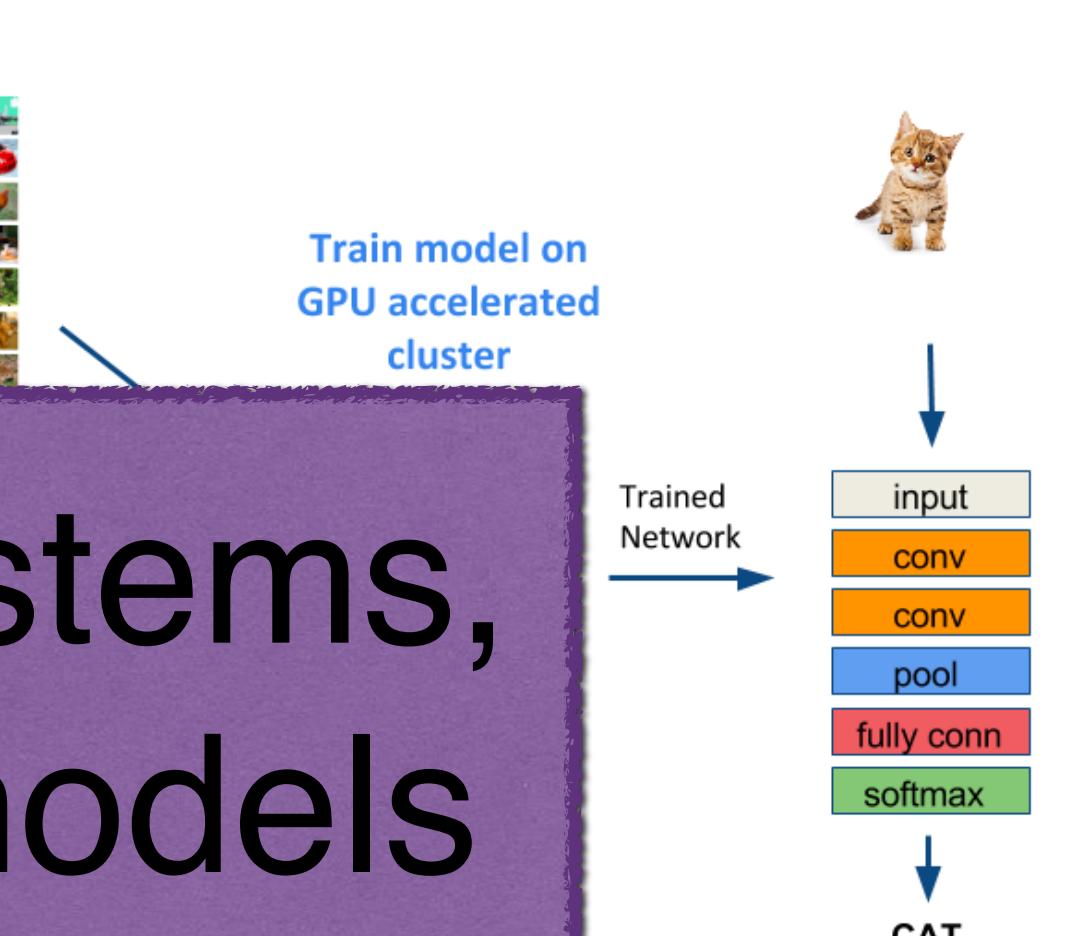
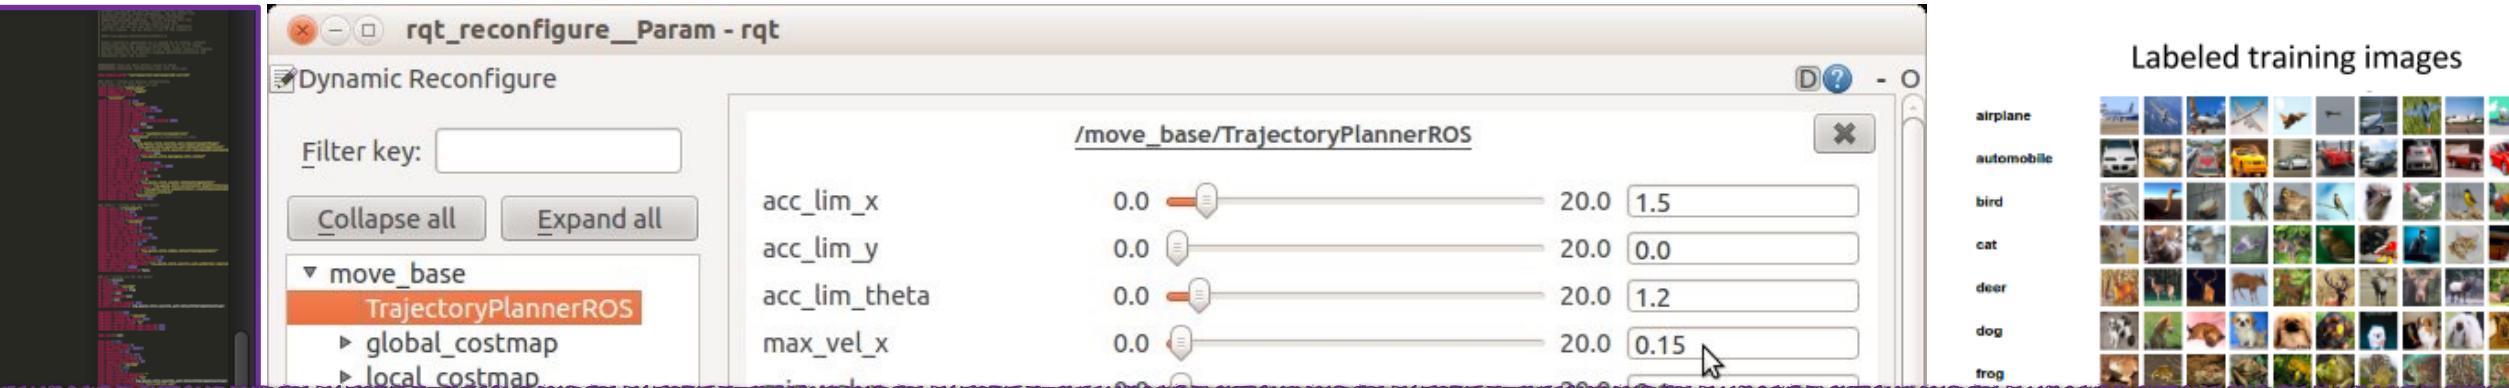
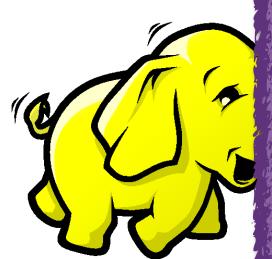
Interested?

Get in touch!

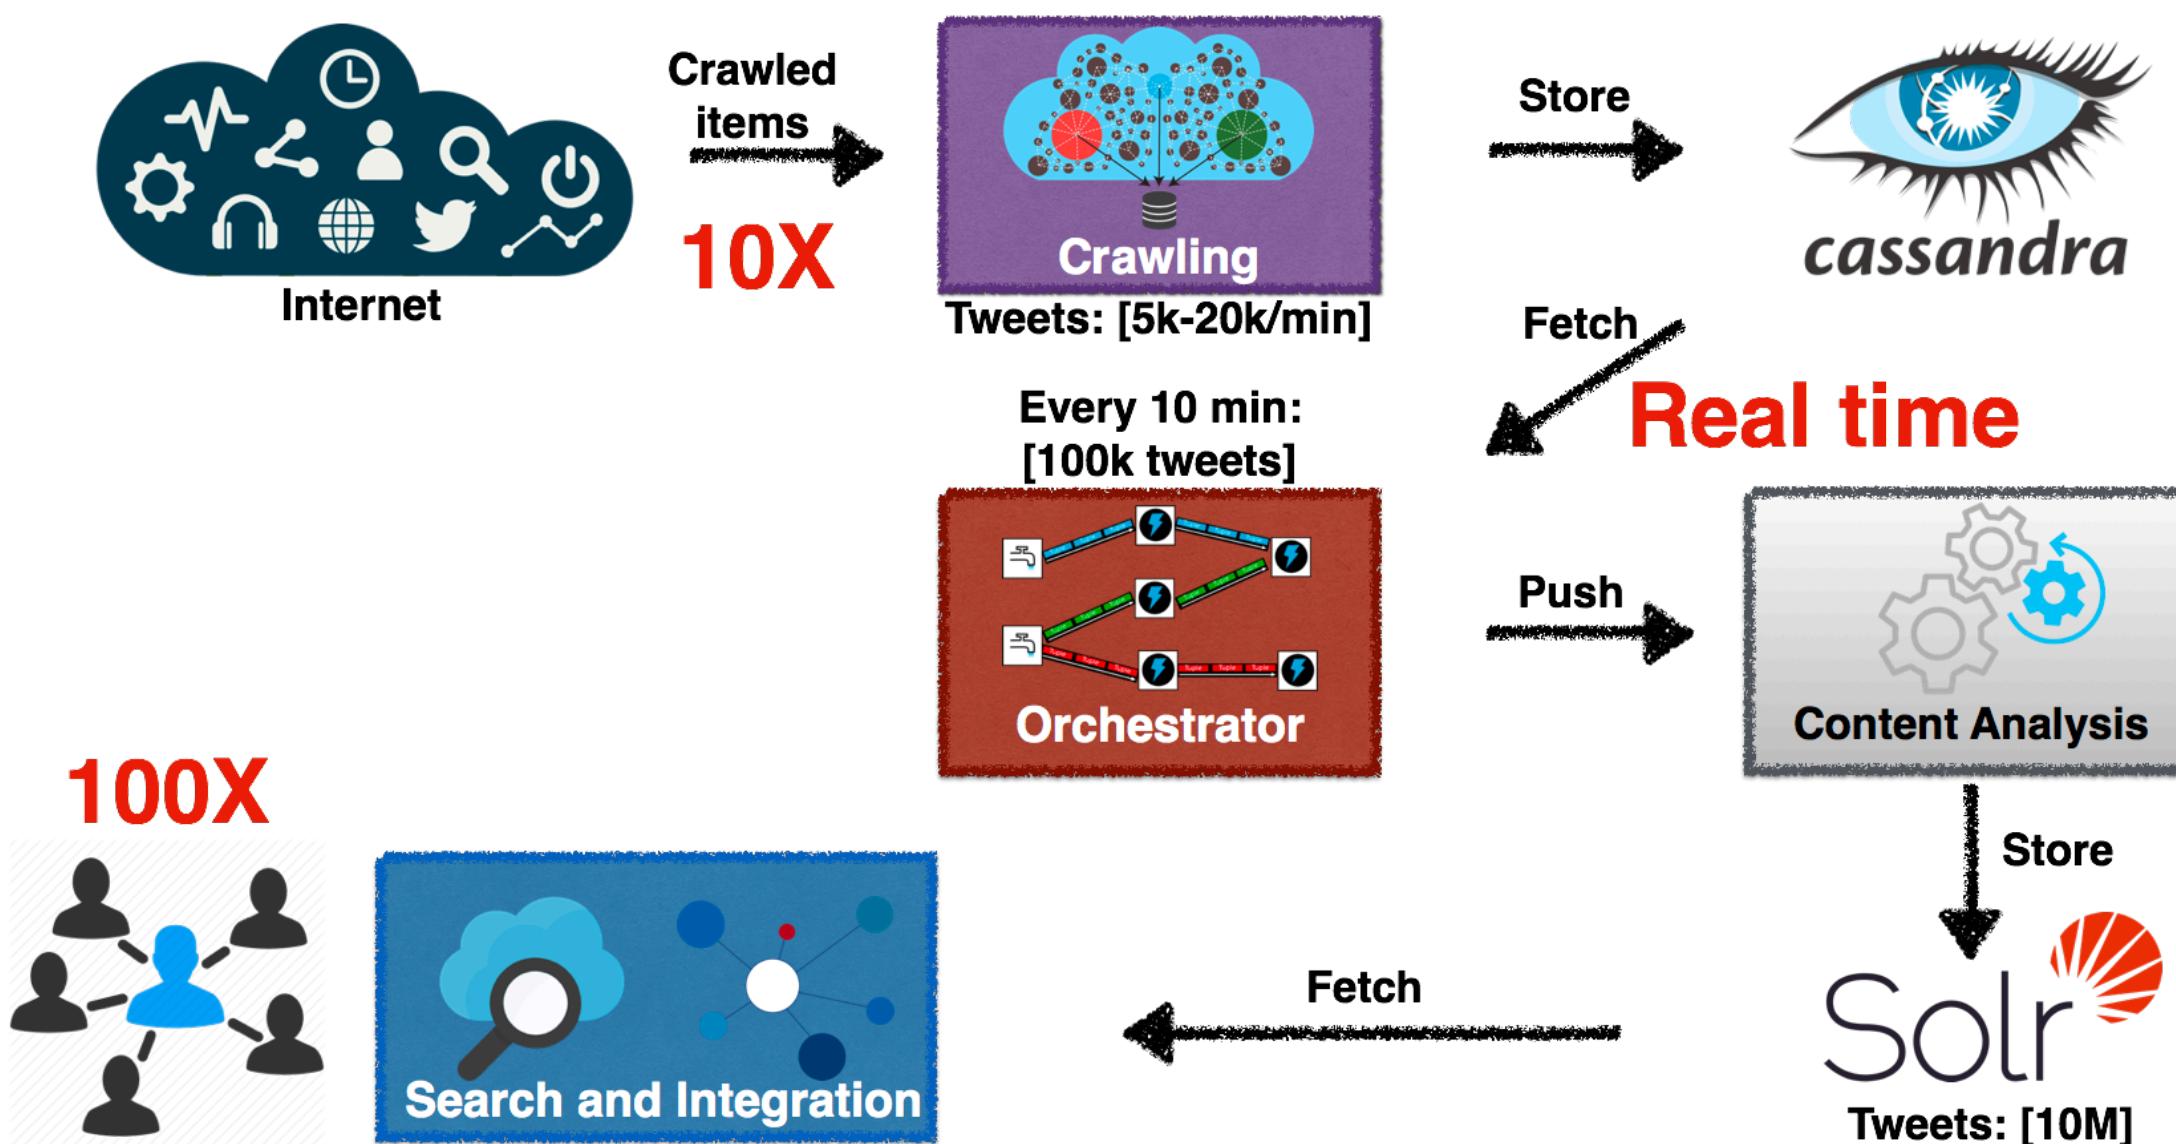
built Many systems are now configurable

```
102  
103 drpc.port: 3772  
104 drpc.worker.threads: 64  
105 drpc.max_buffer_size: 1048576  
106 drpc.queue.size: 128  
107 drpc.invocations.port: 3773  
108 drpc.invocations.threads: 64  
109 drpc.request.timeout.secs: 600  
110 drpc.childopts: "-Xmx768m"  
111 drpc.http.port: 3774  
112 drpc.https.port: -1  
113 drpc.https.keystore.pass:  
114 drpc.https.keystore.type:  
115 drpc.http.creds.plugin:  
116 drpc.authorizer.acl.file:  
117 drpc.authorizer.acl.strict:  
118  
119 transactional.zookeeper.  
transactional.zookeeper.  
transactional.zookeeper.  
120  
121  
122  
123 ## blobstore configs  
124 supervisor.blobstore.class:  
125 supervisor.blobstore.download:  
126 supervisor.blobstore.download:  
127 supervisor.localizer.class:  
128 supervisor.localizer.clear:
```

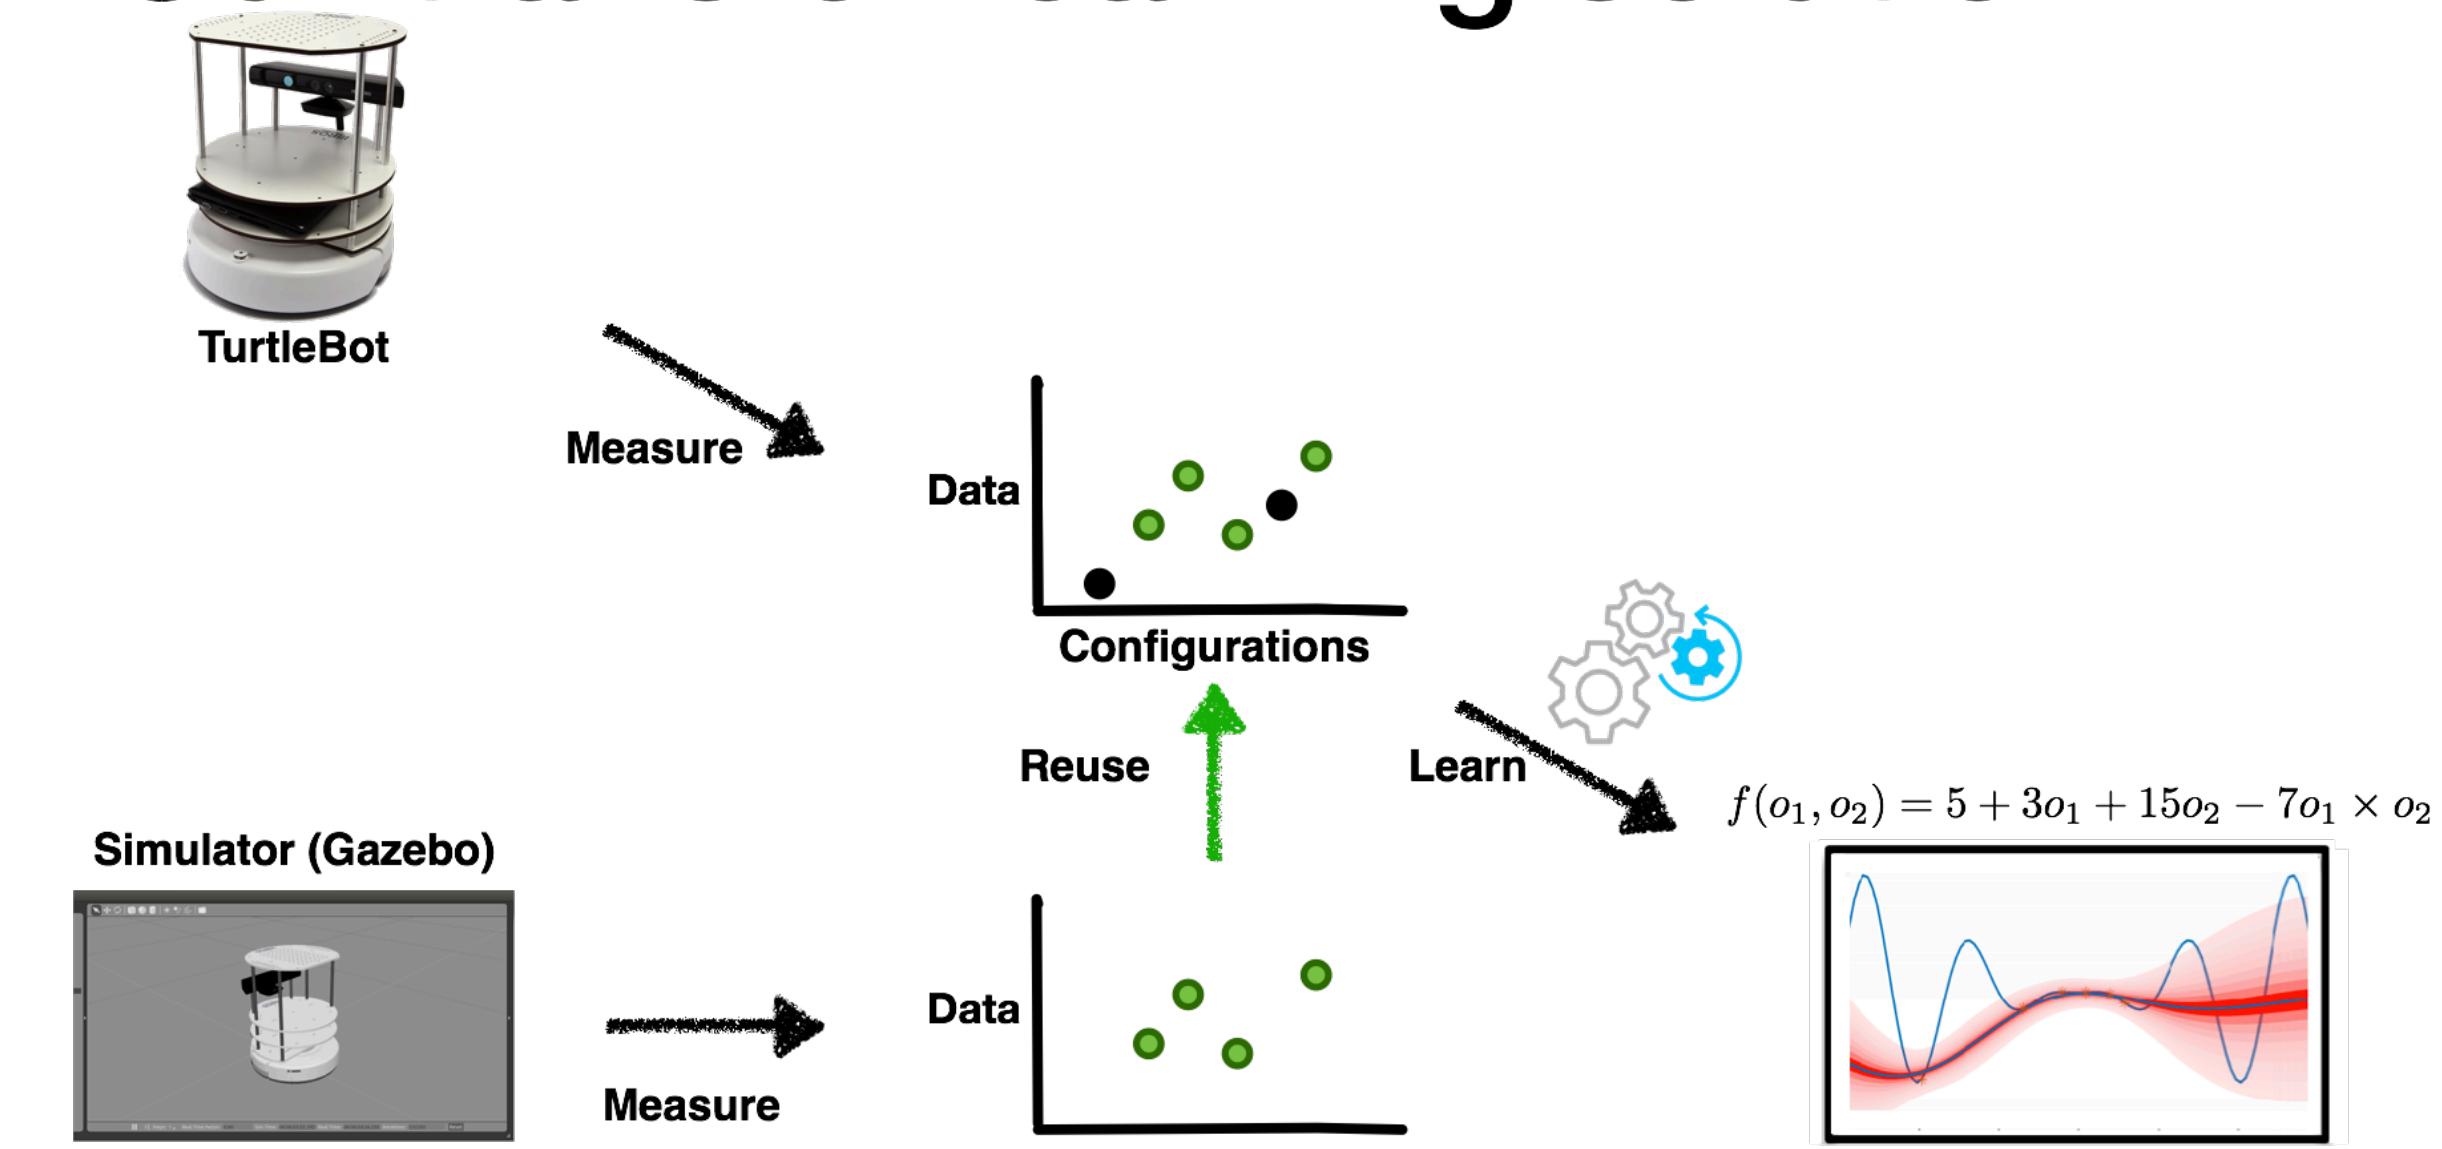
Given the ever growing configurable systems,
how can we enable learning practical models
that scale well and provide reliable predictions
for exploring the configuration space?



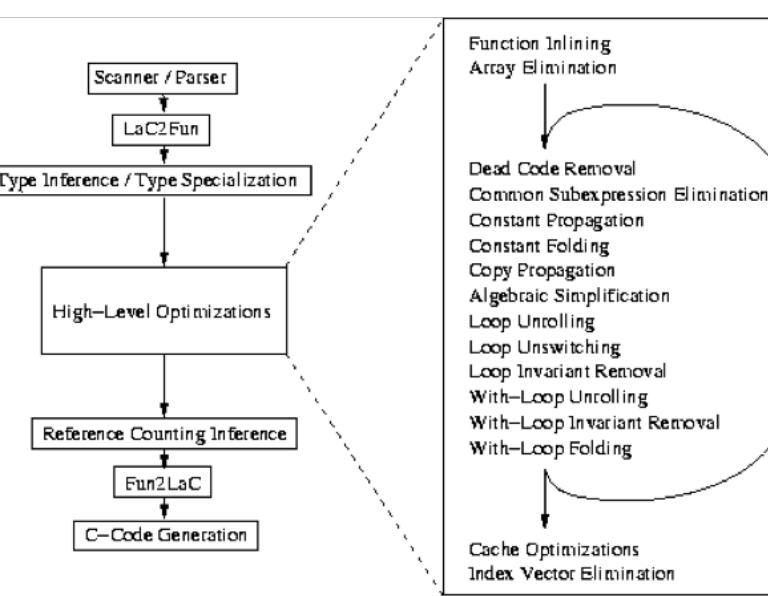
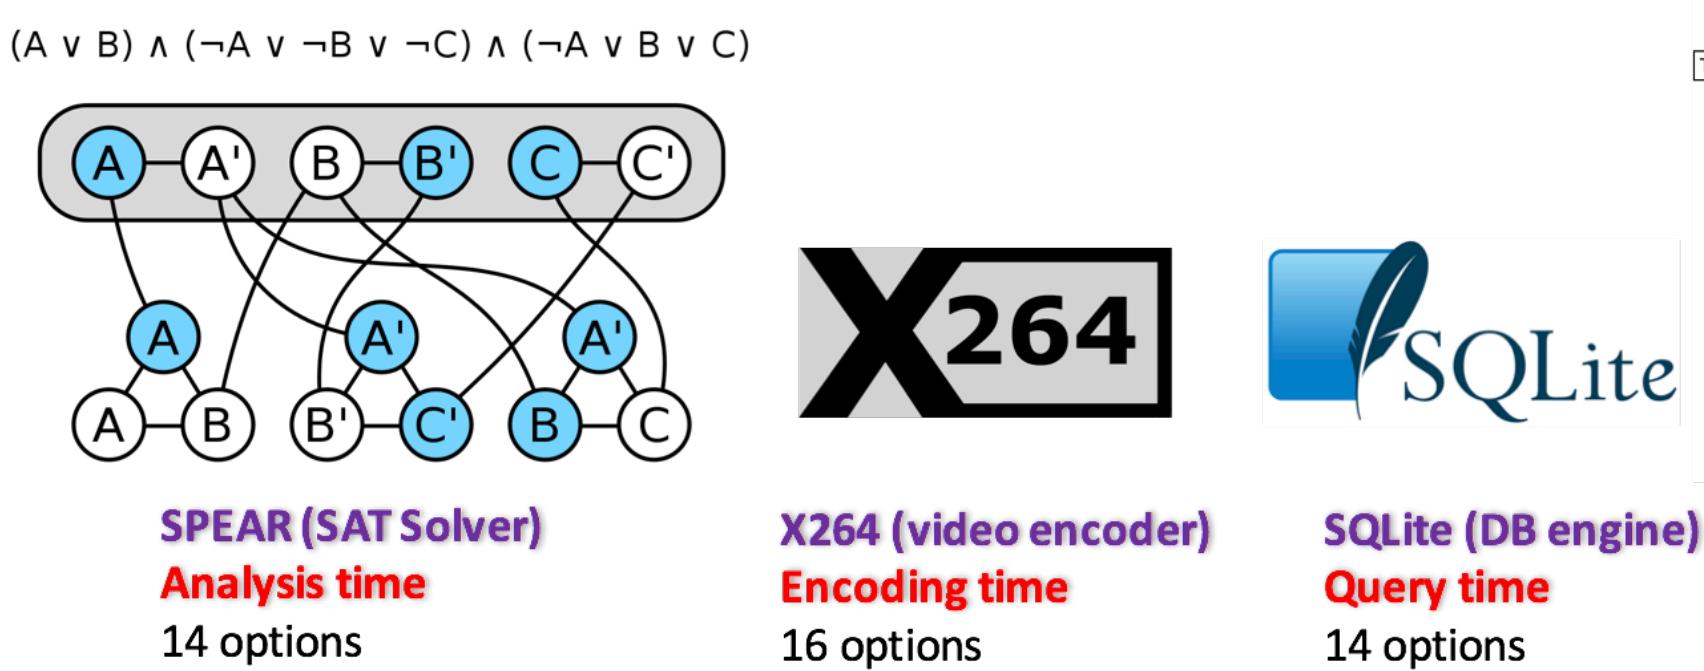
Challenges



Our transfer learning solution



Our empirical study: We looked at different highly-configurable systems to gain insights



Exploring the design space of deep networks

